

Mining for Dark Matter substructure: Learning from lenses without a likelihood

Johann Brehmer

New York University

DMWG seminar, RIKEN

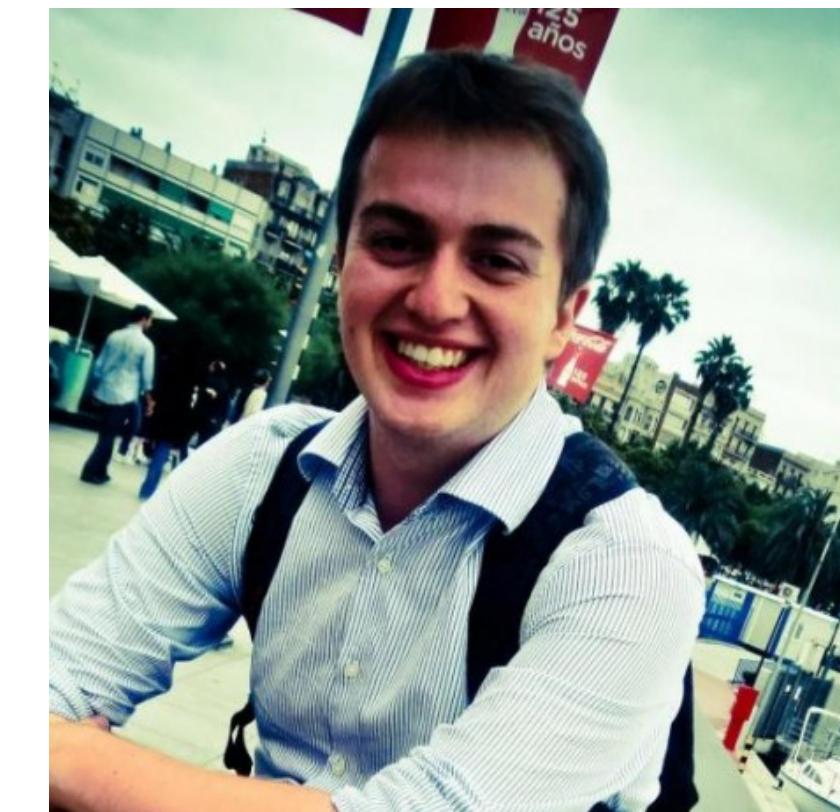
February 17, 2020



Siddharth Mishra-Sharma



Joeri Hermans



Gilles Louppe



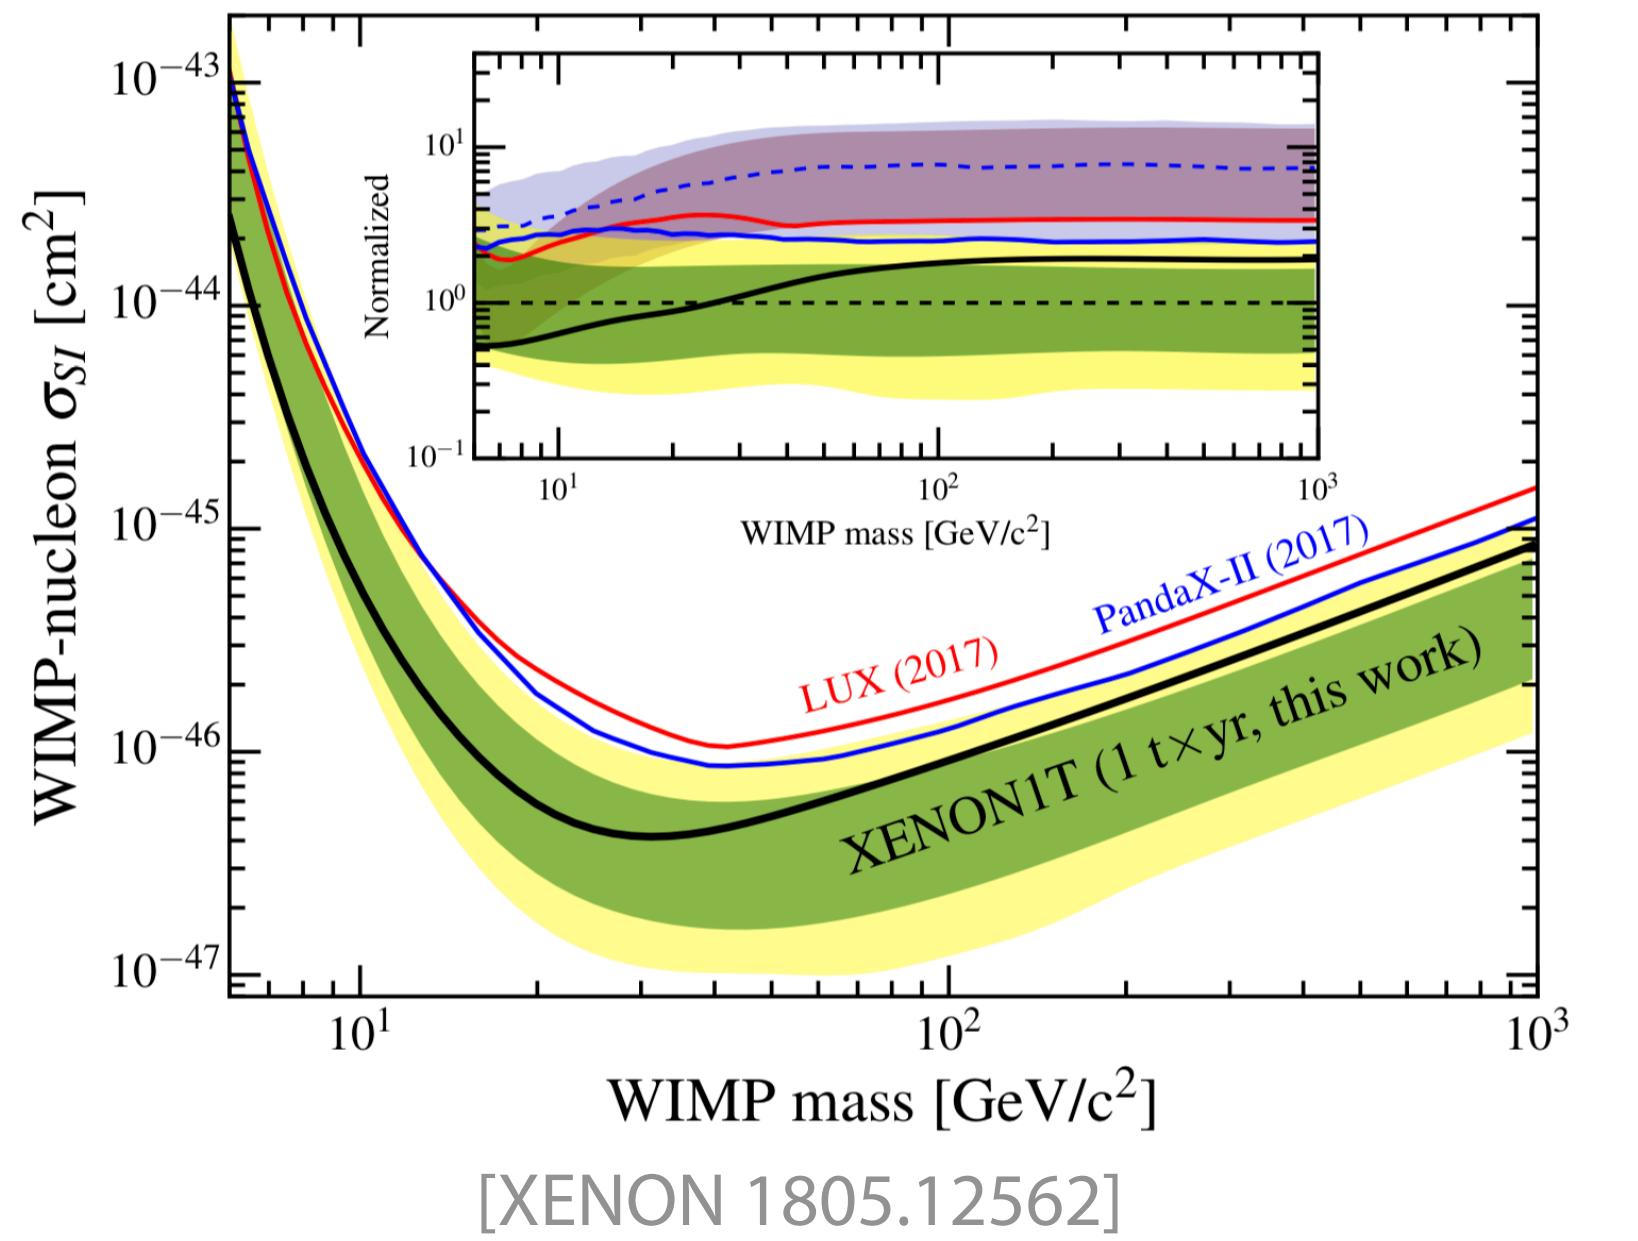
Kyle Cranmer



The SCAILFIN Project
scailfin.github.io

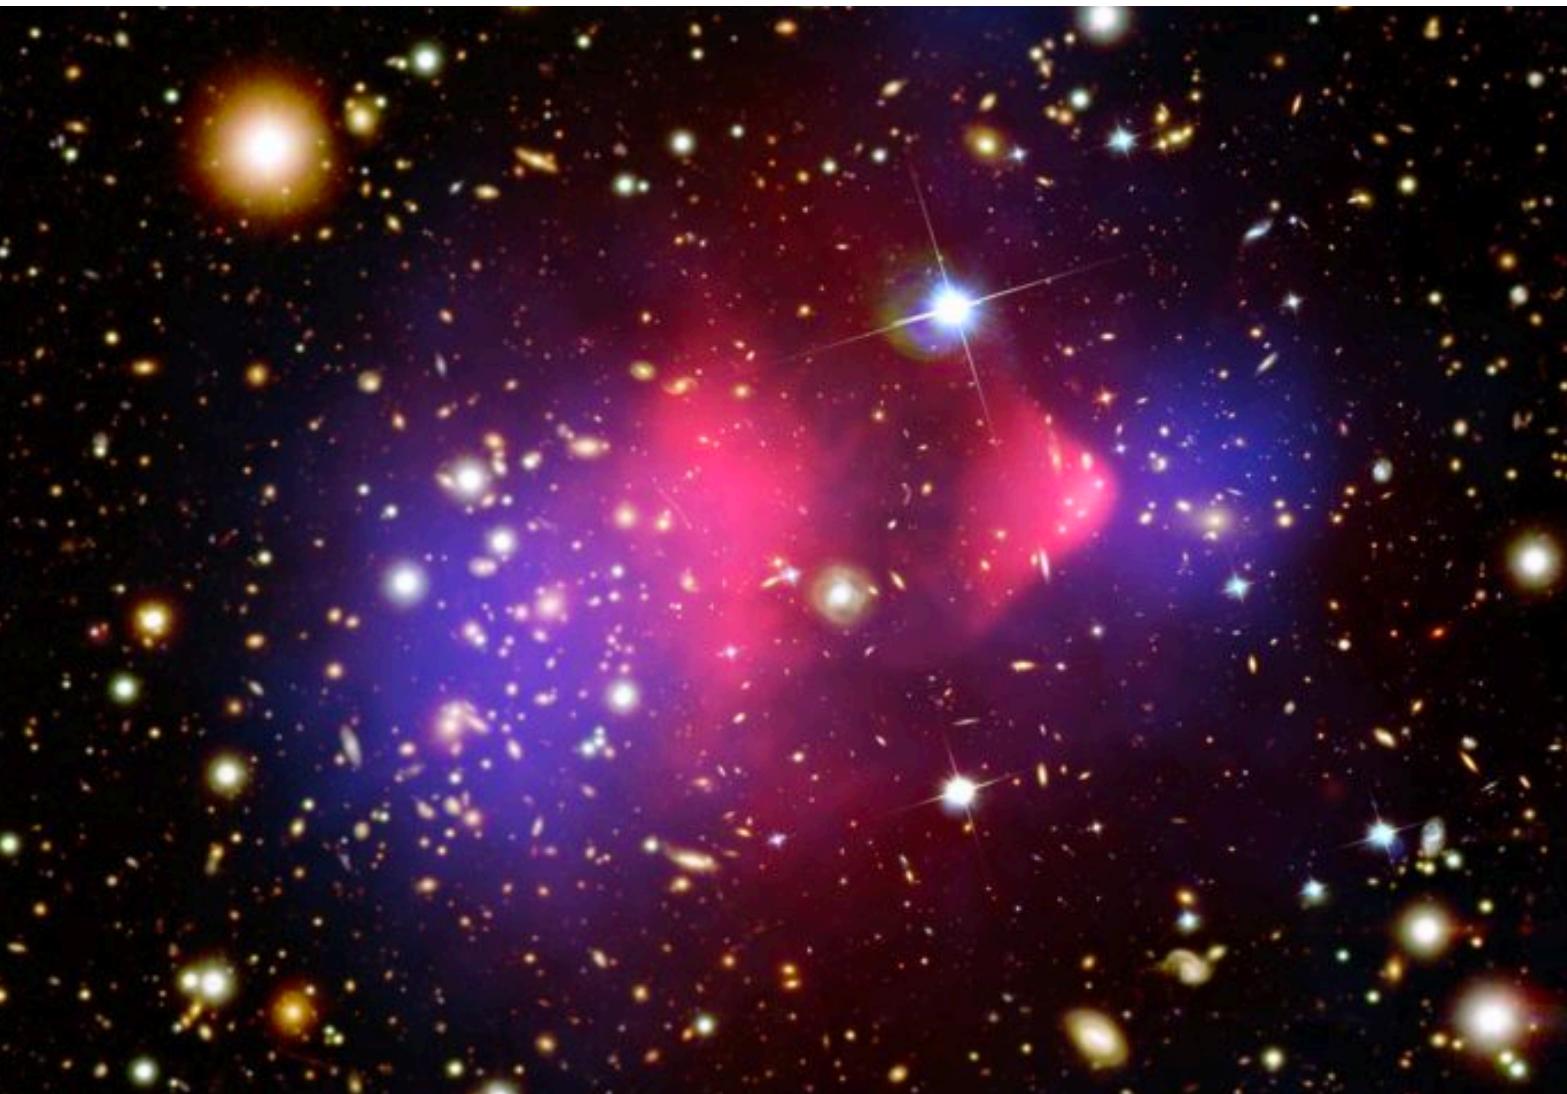
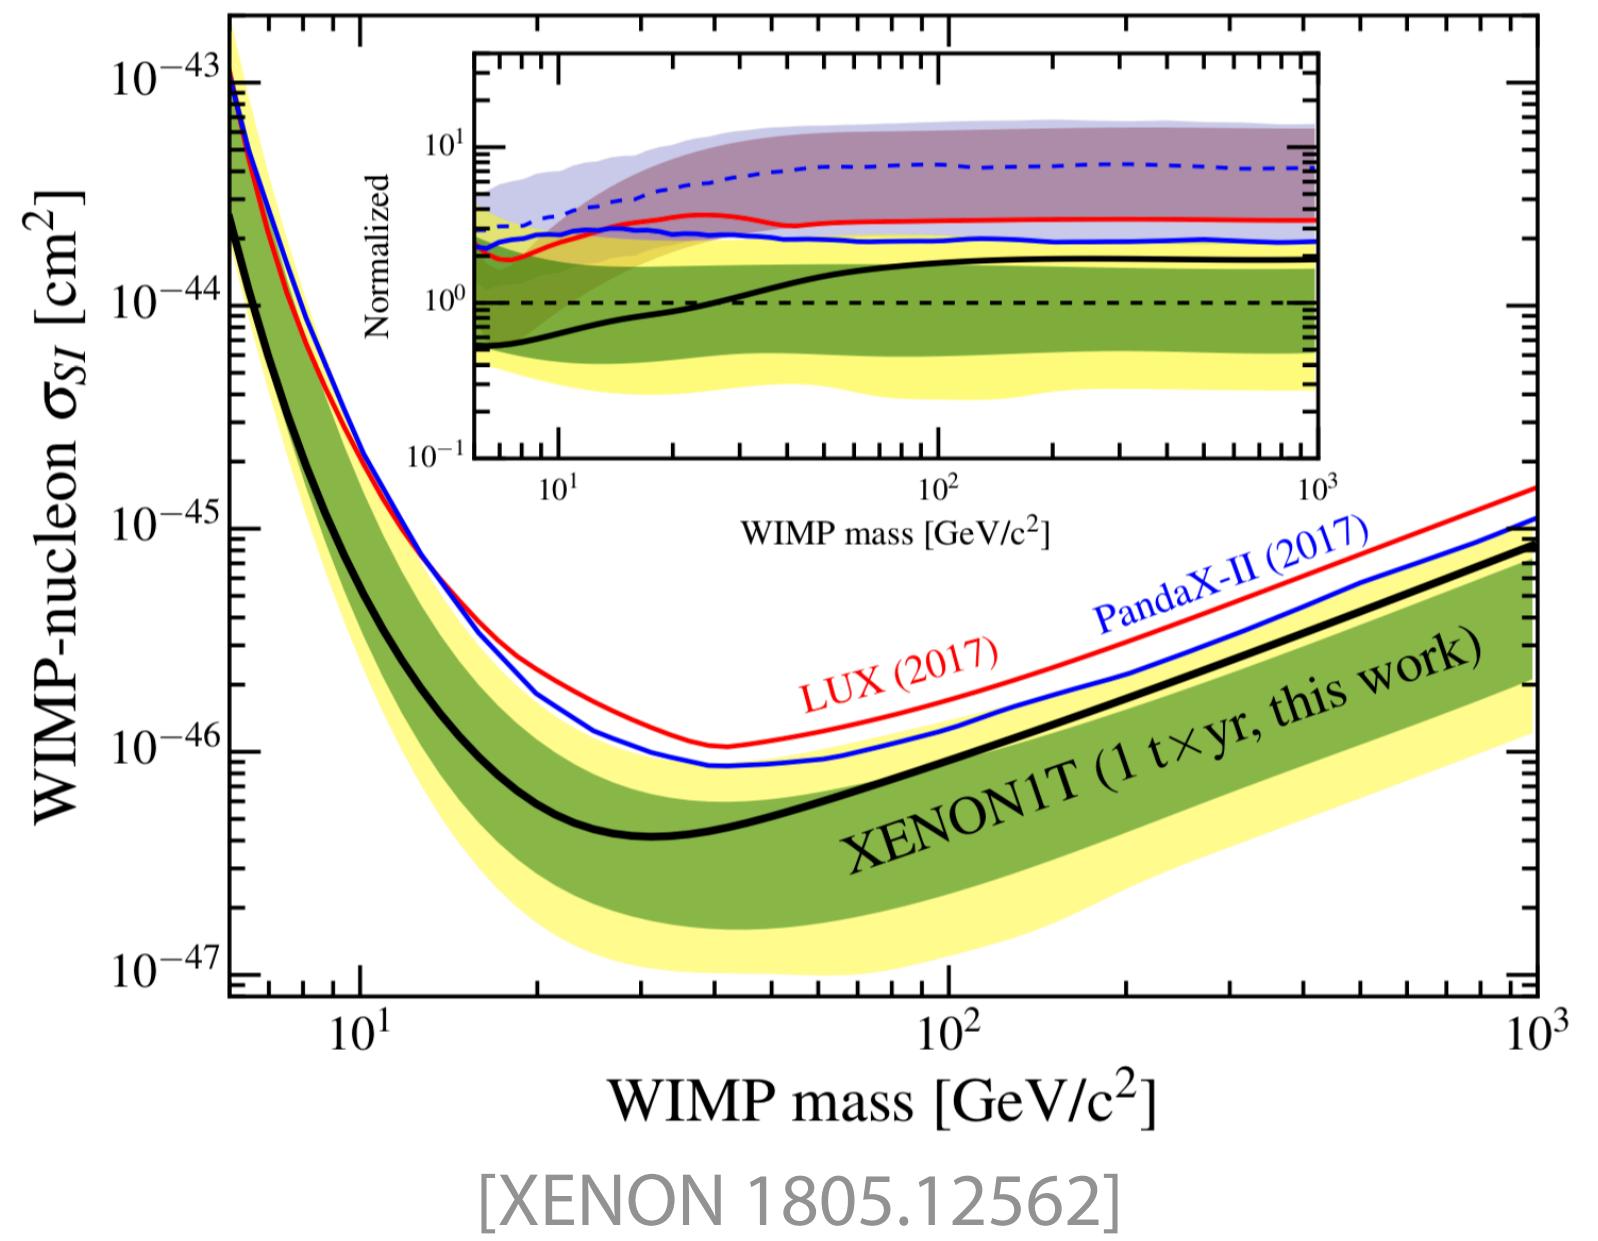


Dark matter



Non-gravitational interactions:
no evidence so far

Dark matter

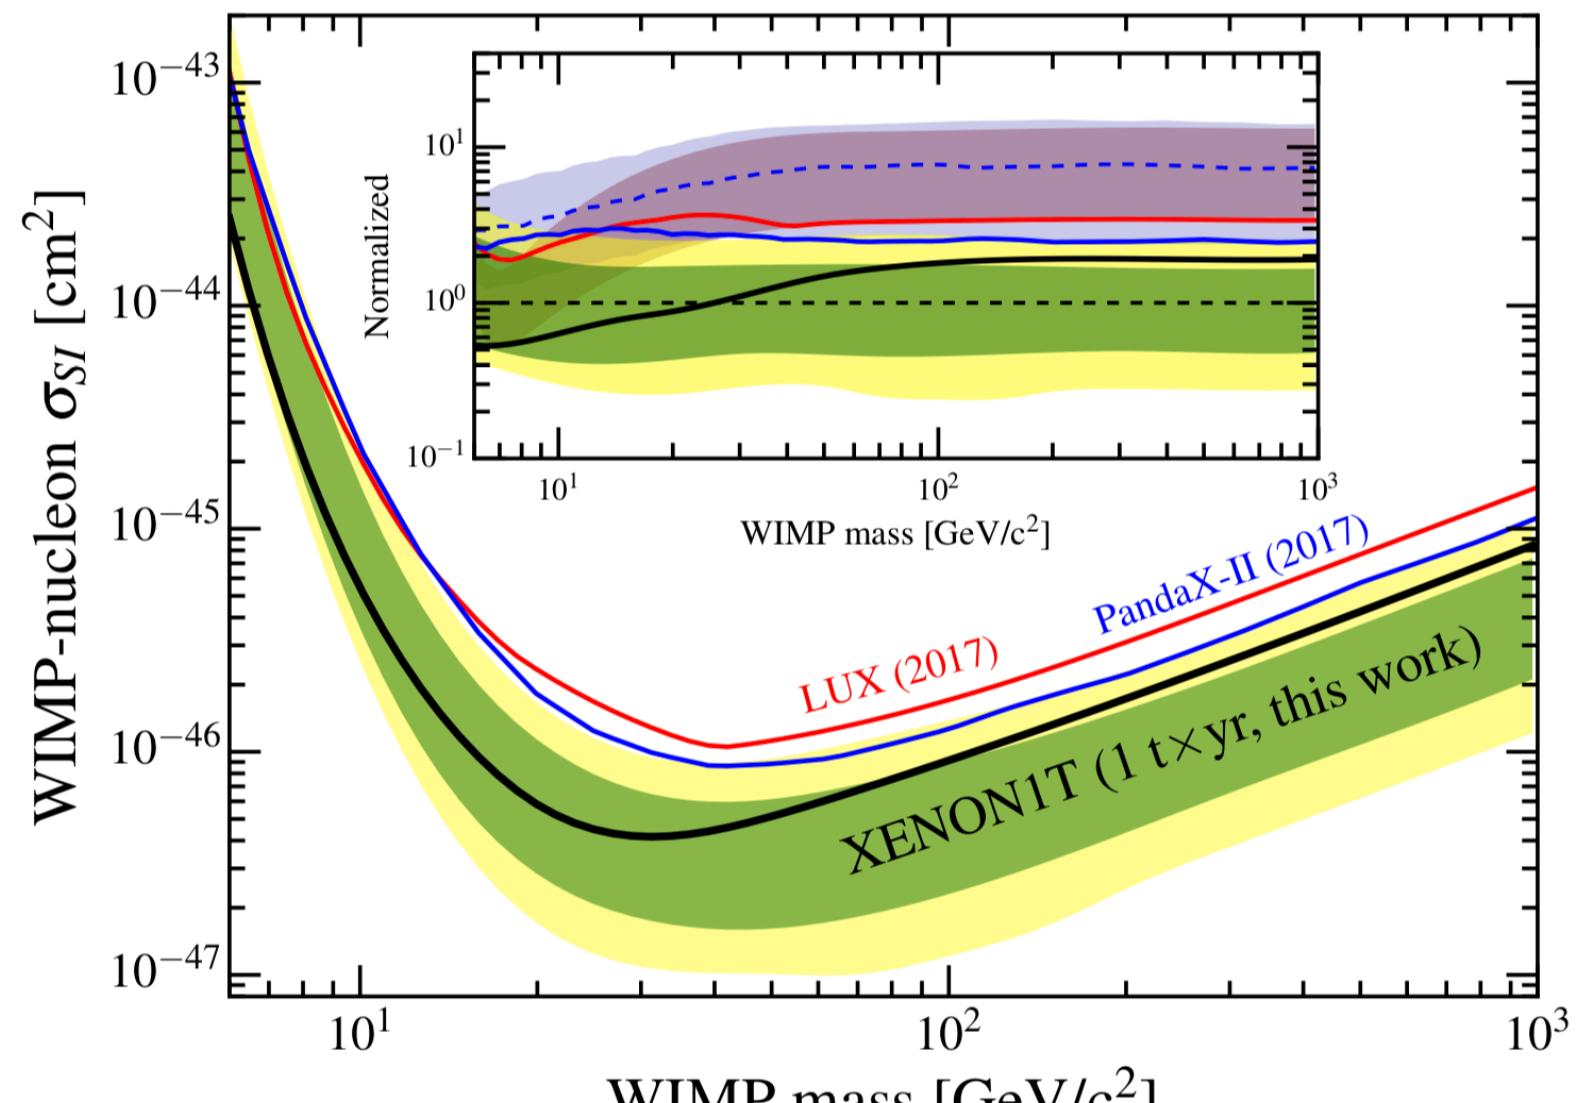


[NASA]

Non-gravitational interactions:
no evidence so far

Gravitational interactions at **large length scales**: plenty of evidence, consistent with Λ CDM

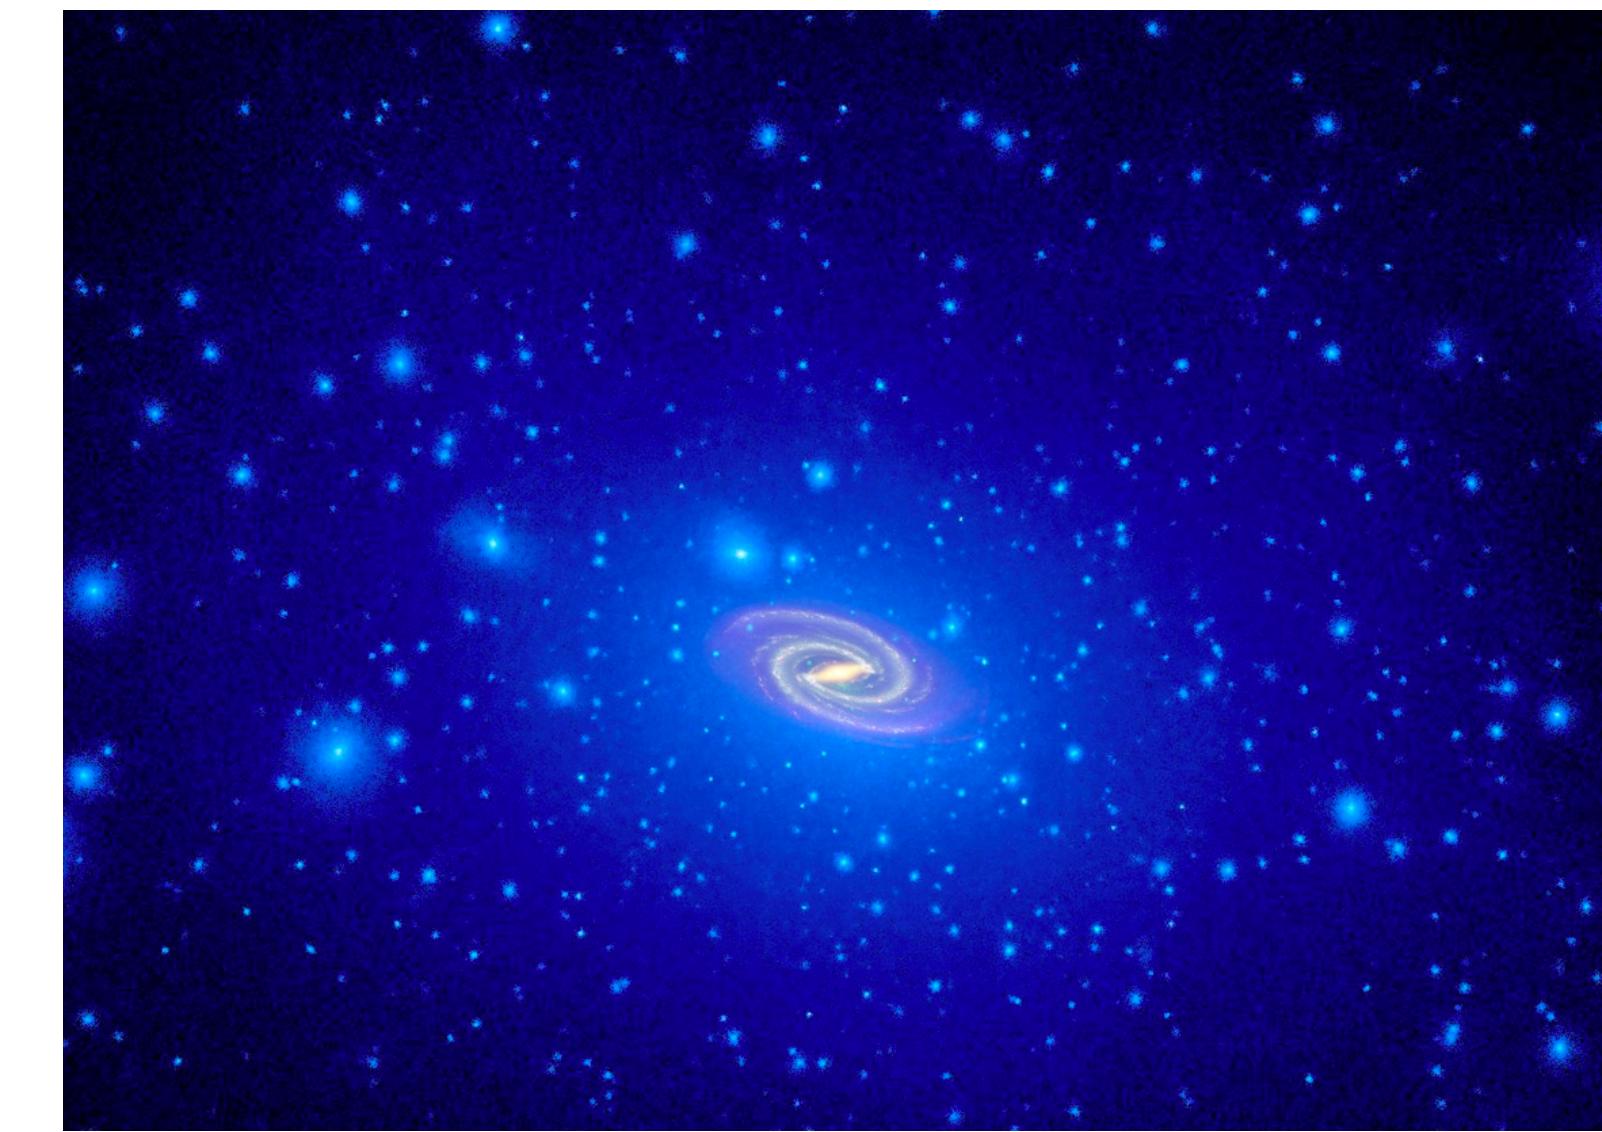
Dark matter



[XENON 1805.12562]



[NASA]



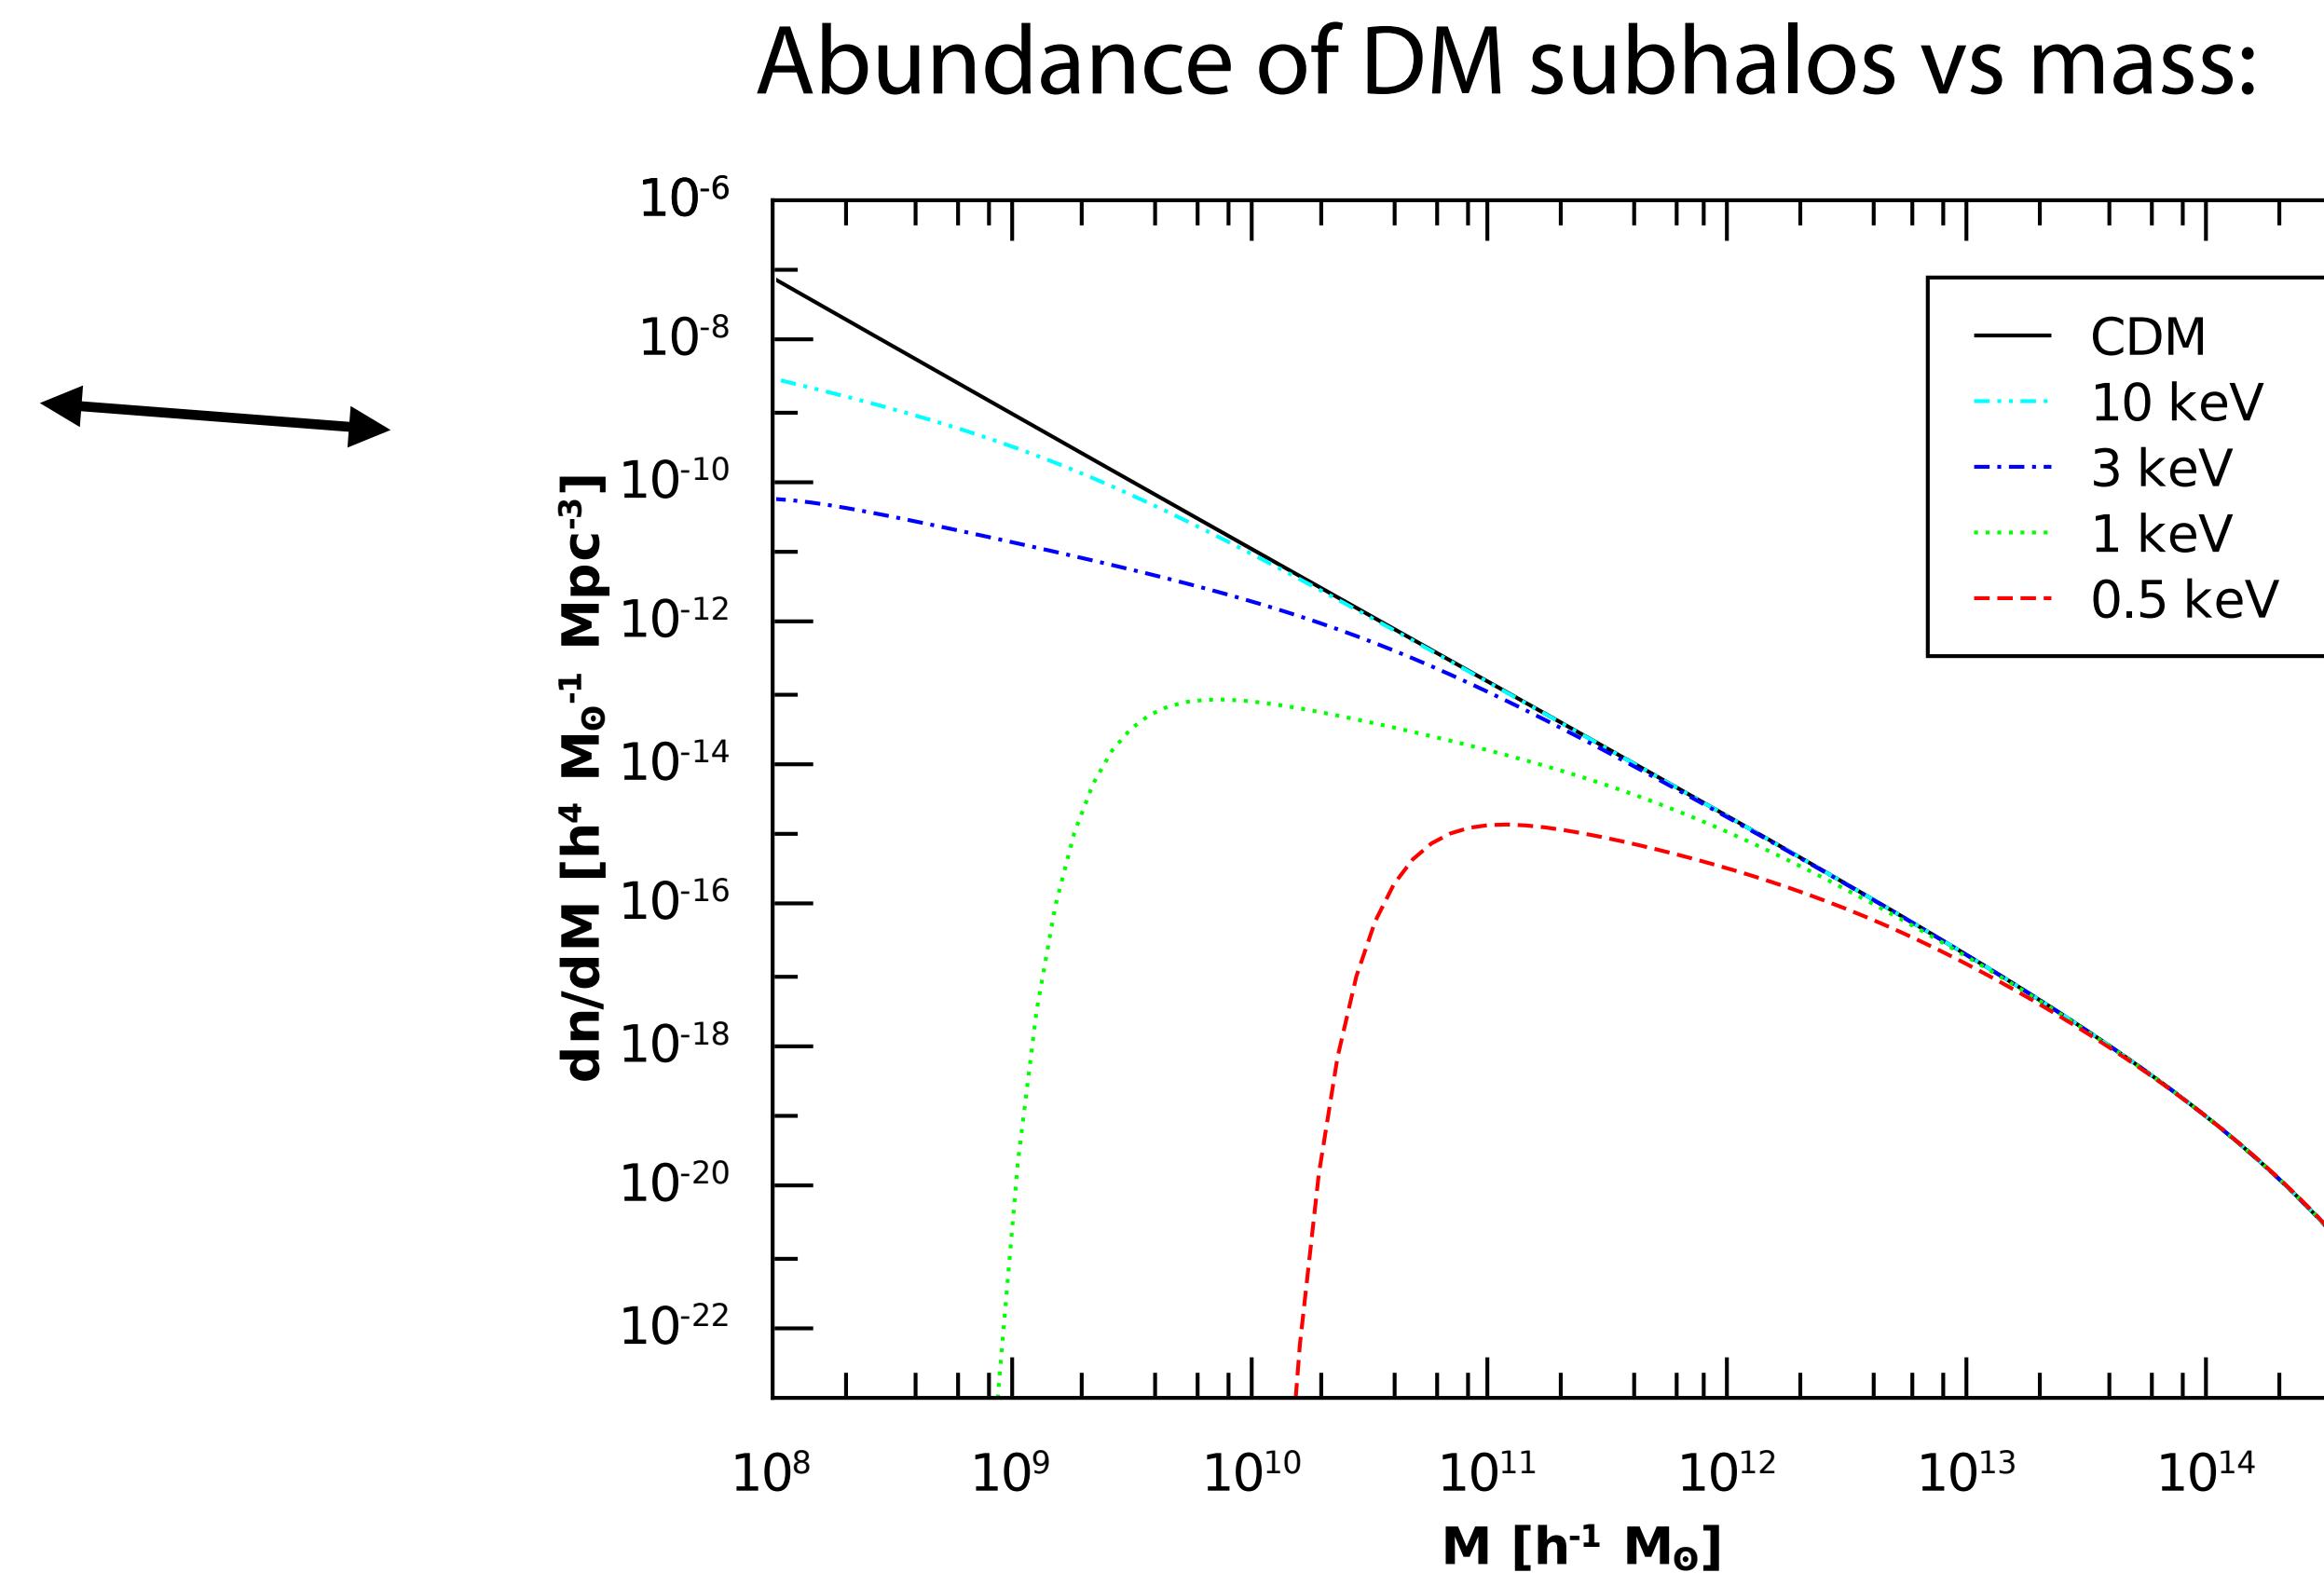
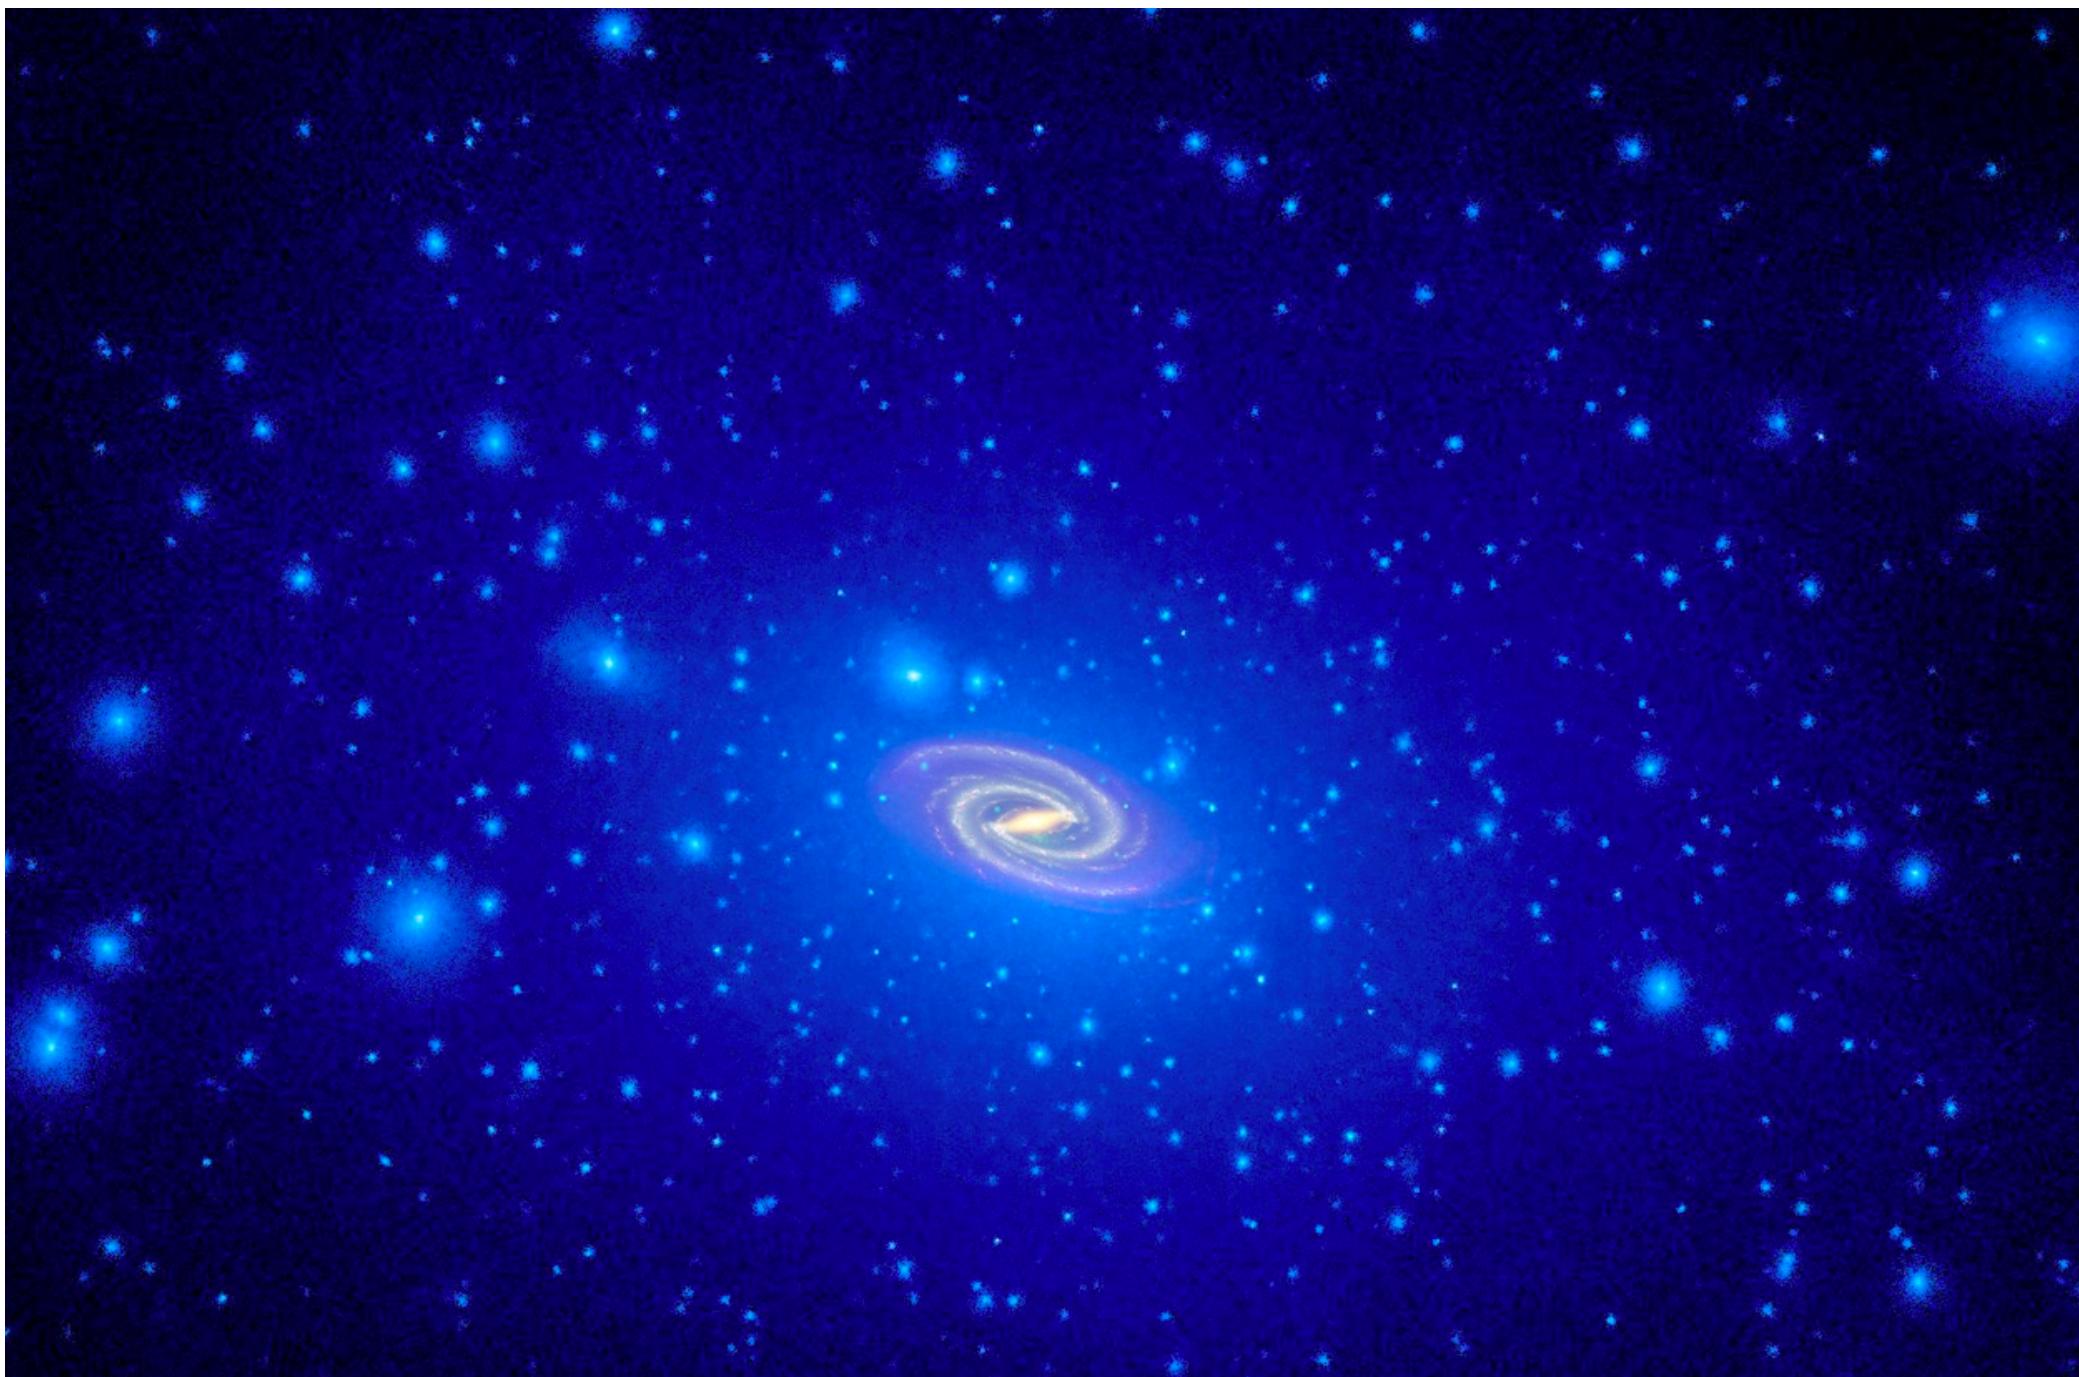
[T. Brown, J.Tumlinson]

Non-gravitational interactions:
no evidence so far

Gravitational interactions at **large length scales**: plenty of evidence, consistent with Λ CDM

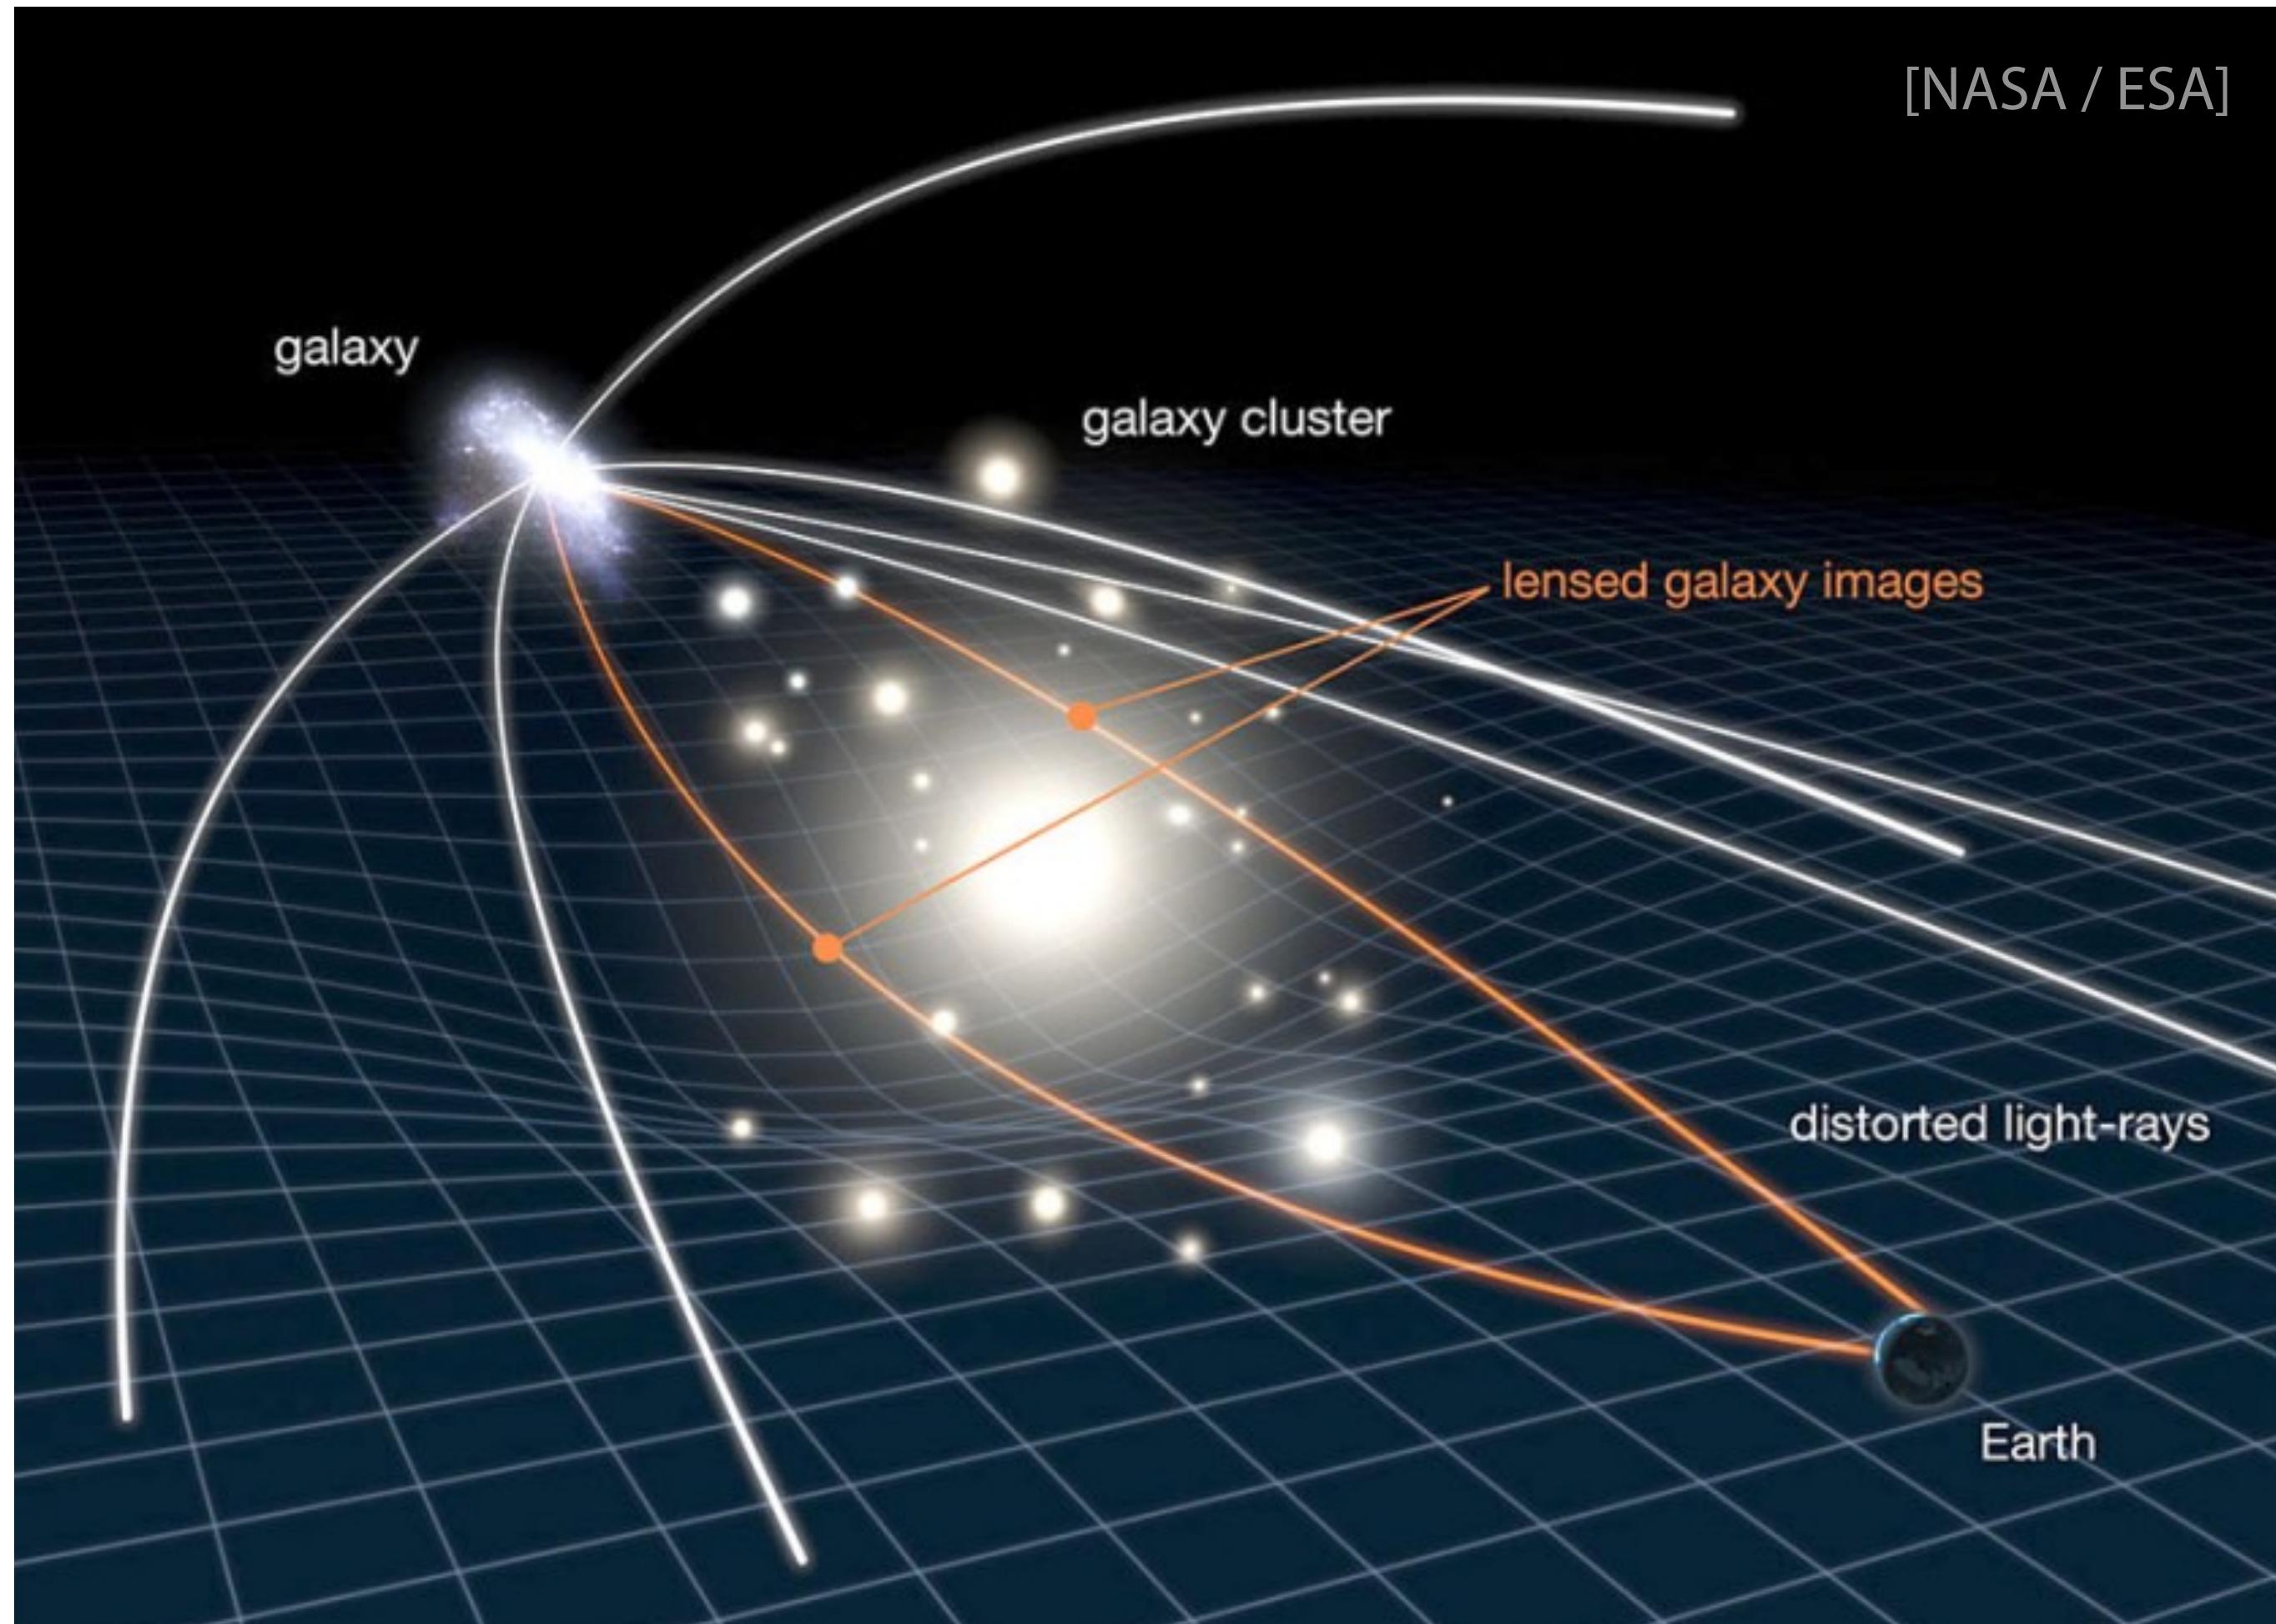
Gravitational interactions at **small scales (DM substructure)**: beginning to be probed, many models predict deviations from Λ CDM

DM small-scale structure is a probe of its particle nature



[R. Dunstan et al 1109.6291]

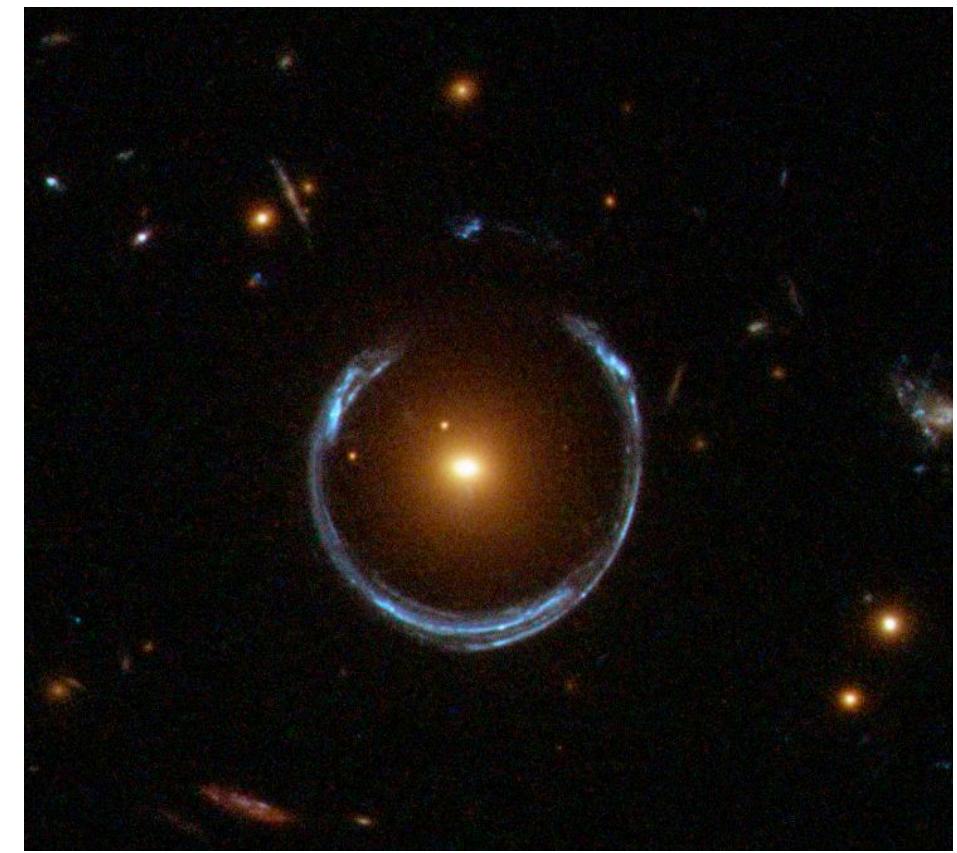
Strong gravitational lensing



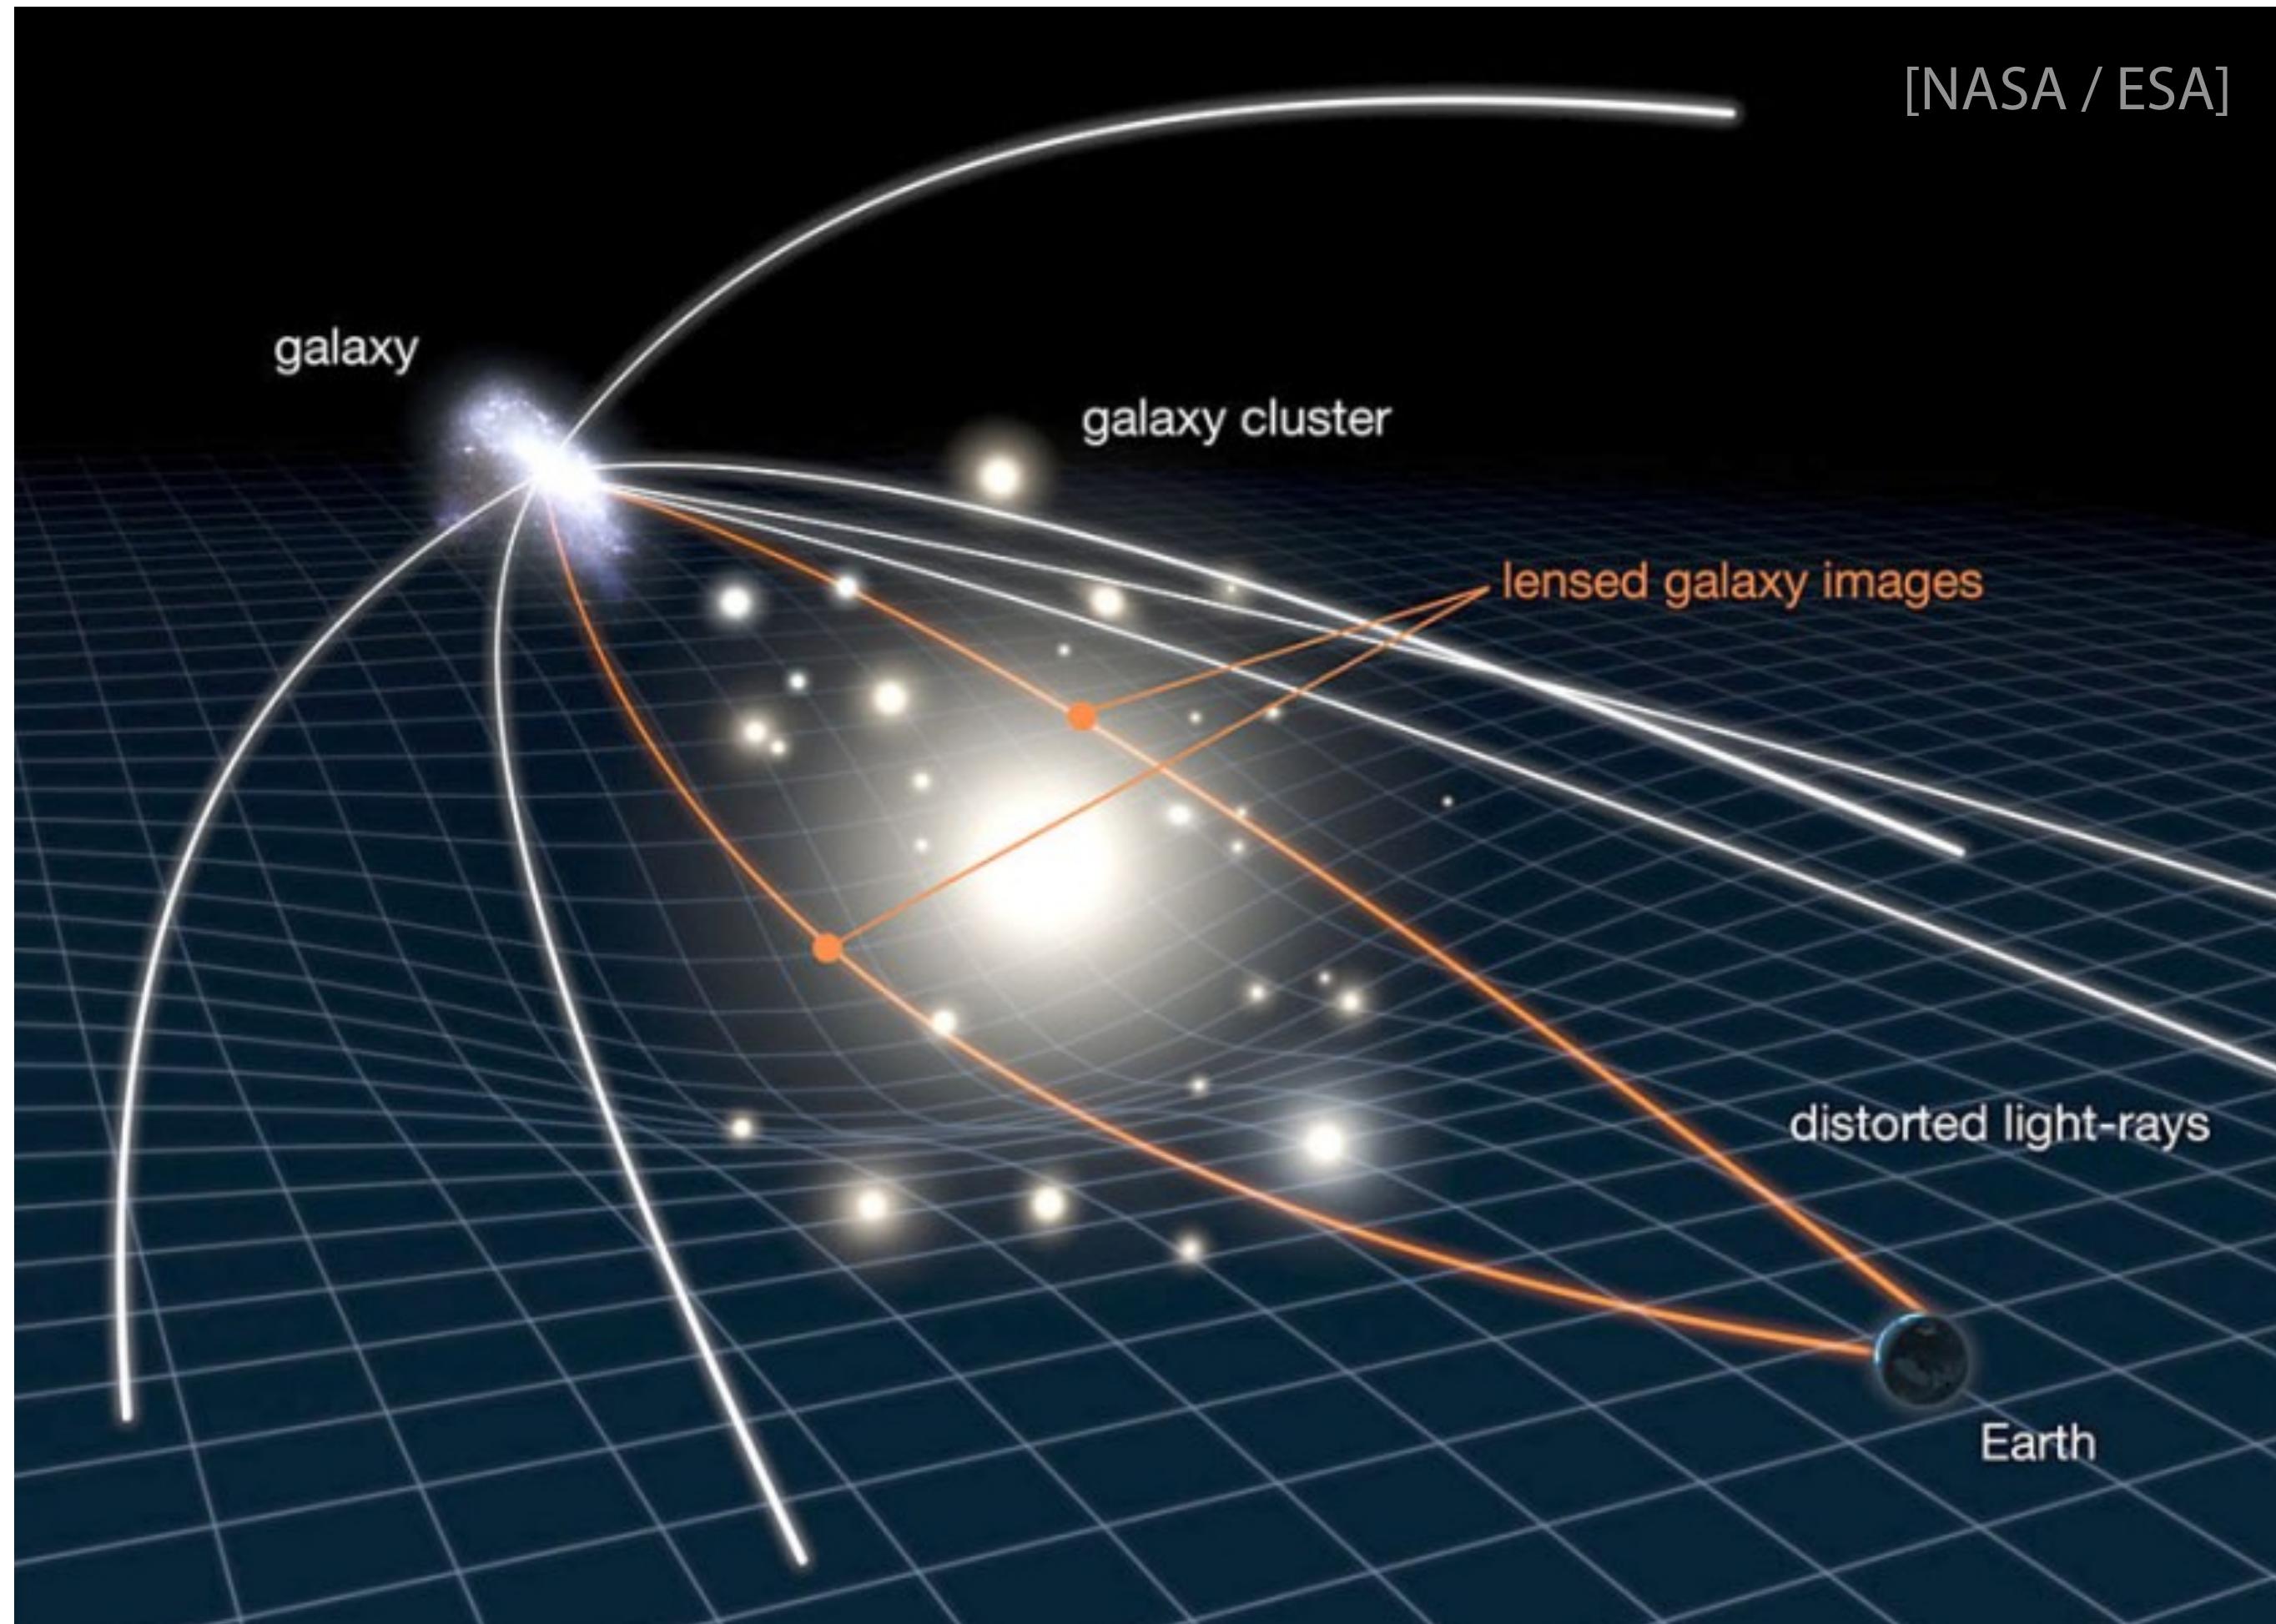
Multiple images
of quasars



Extended arcs
from galaxies



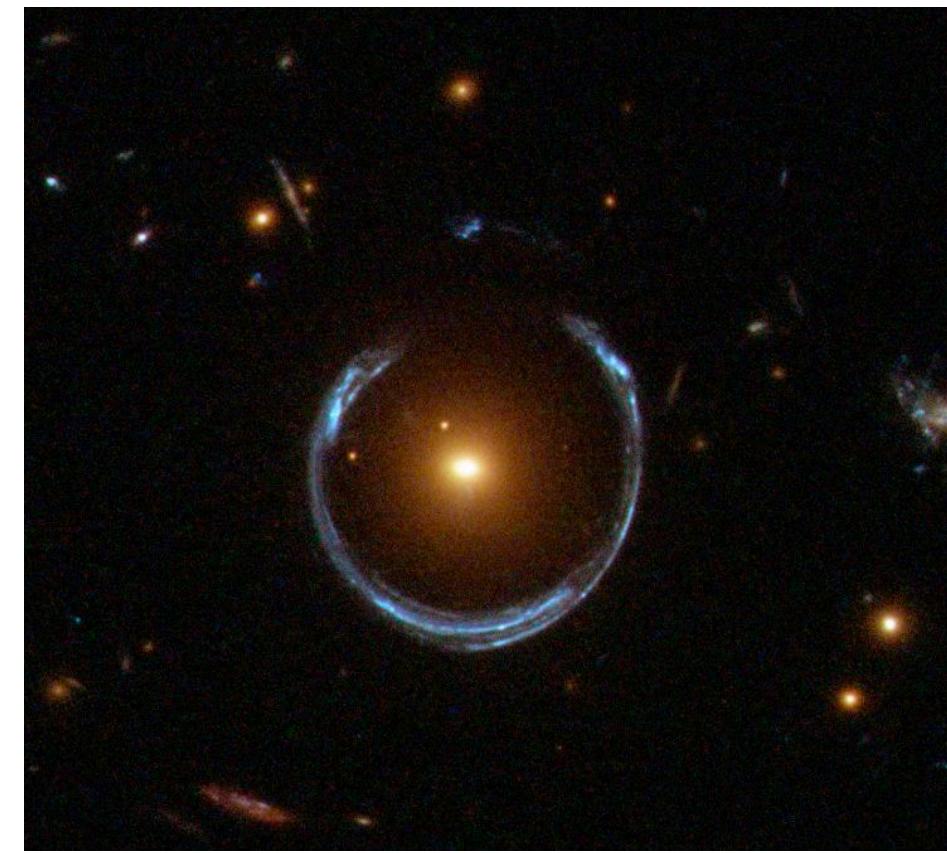
Strong gravitational lensing



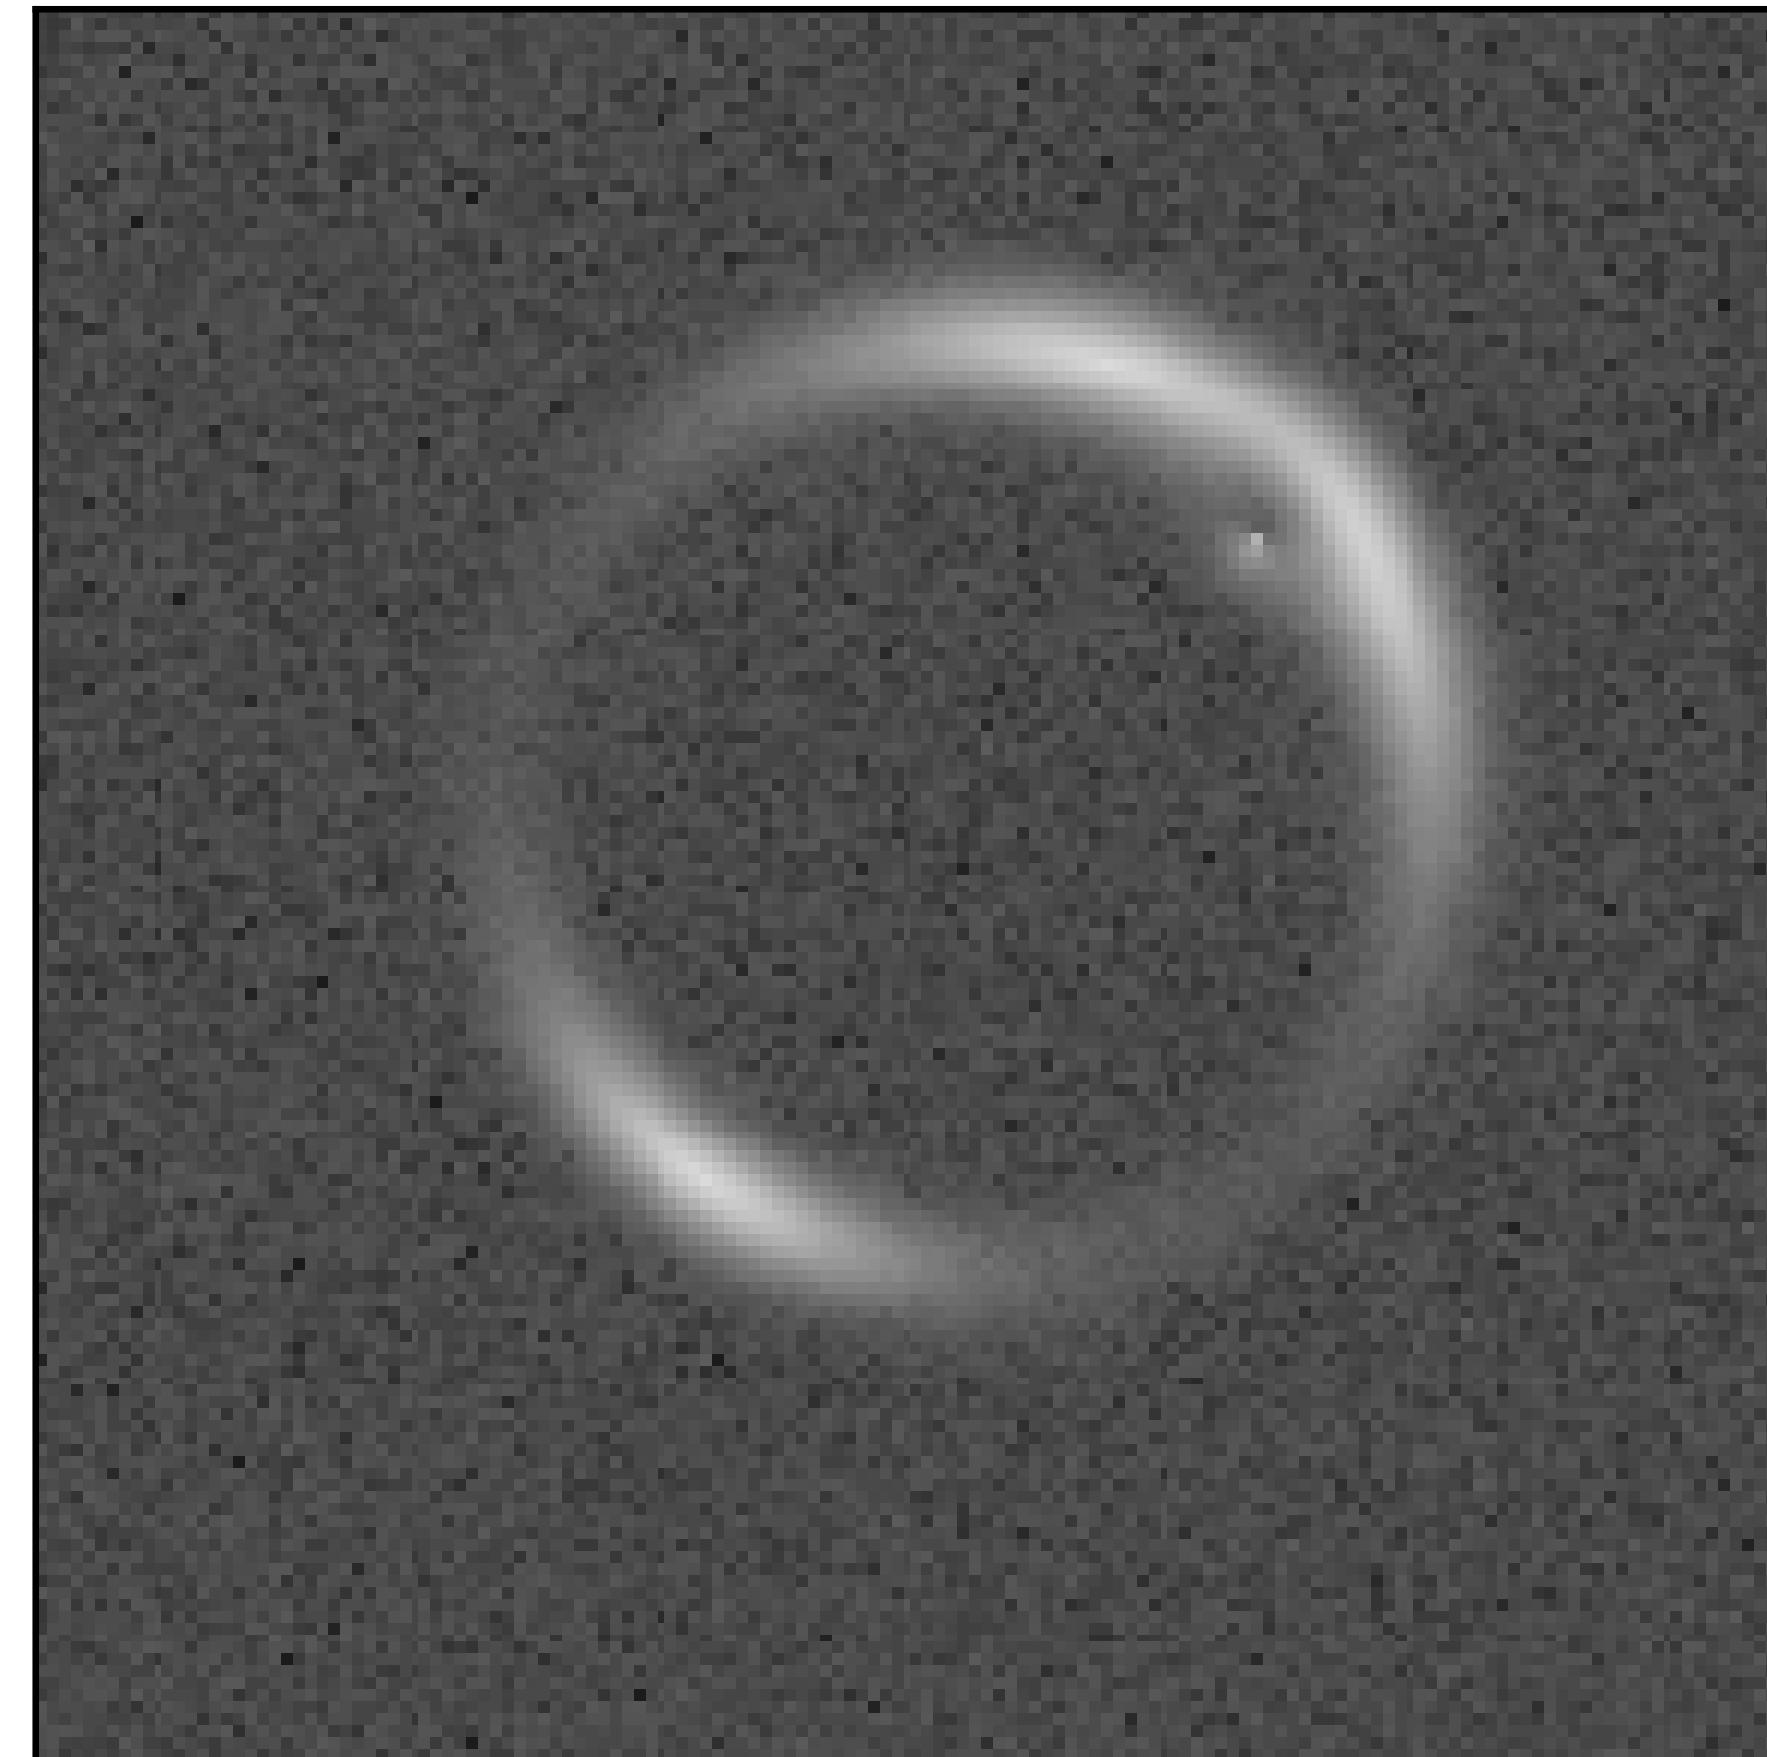
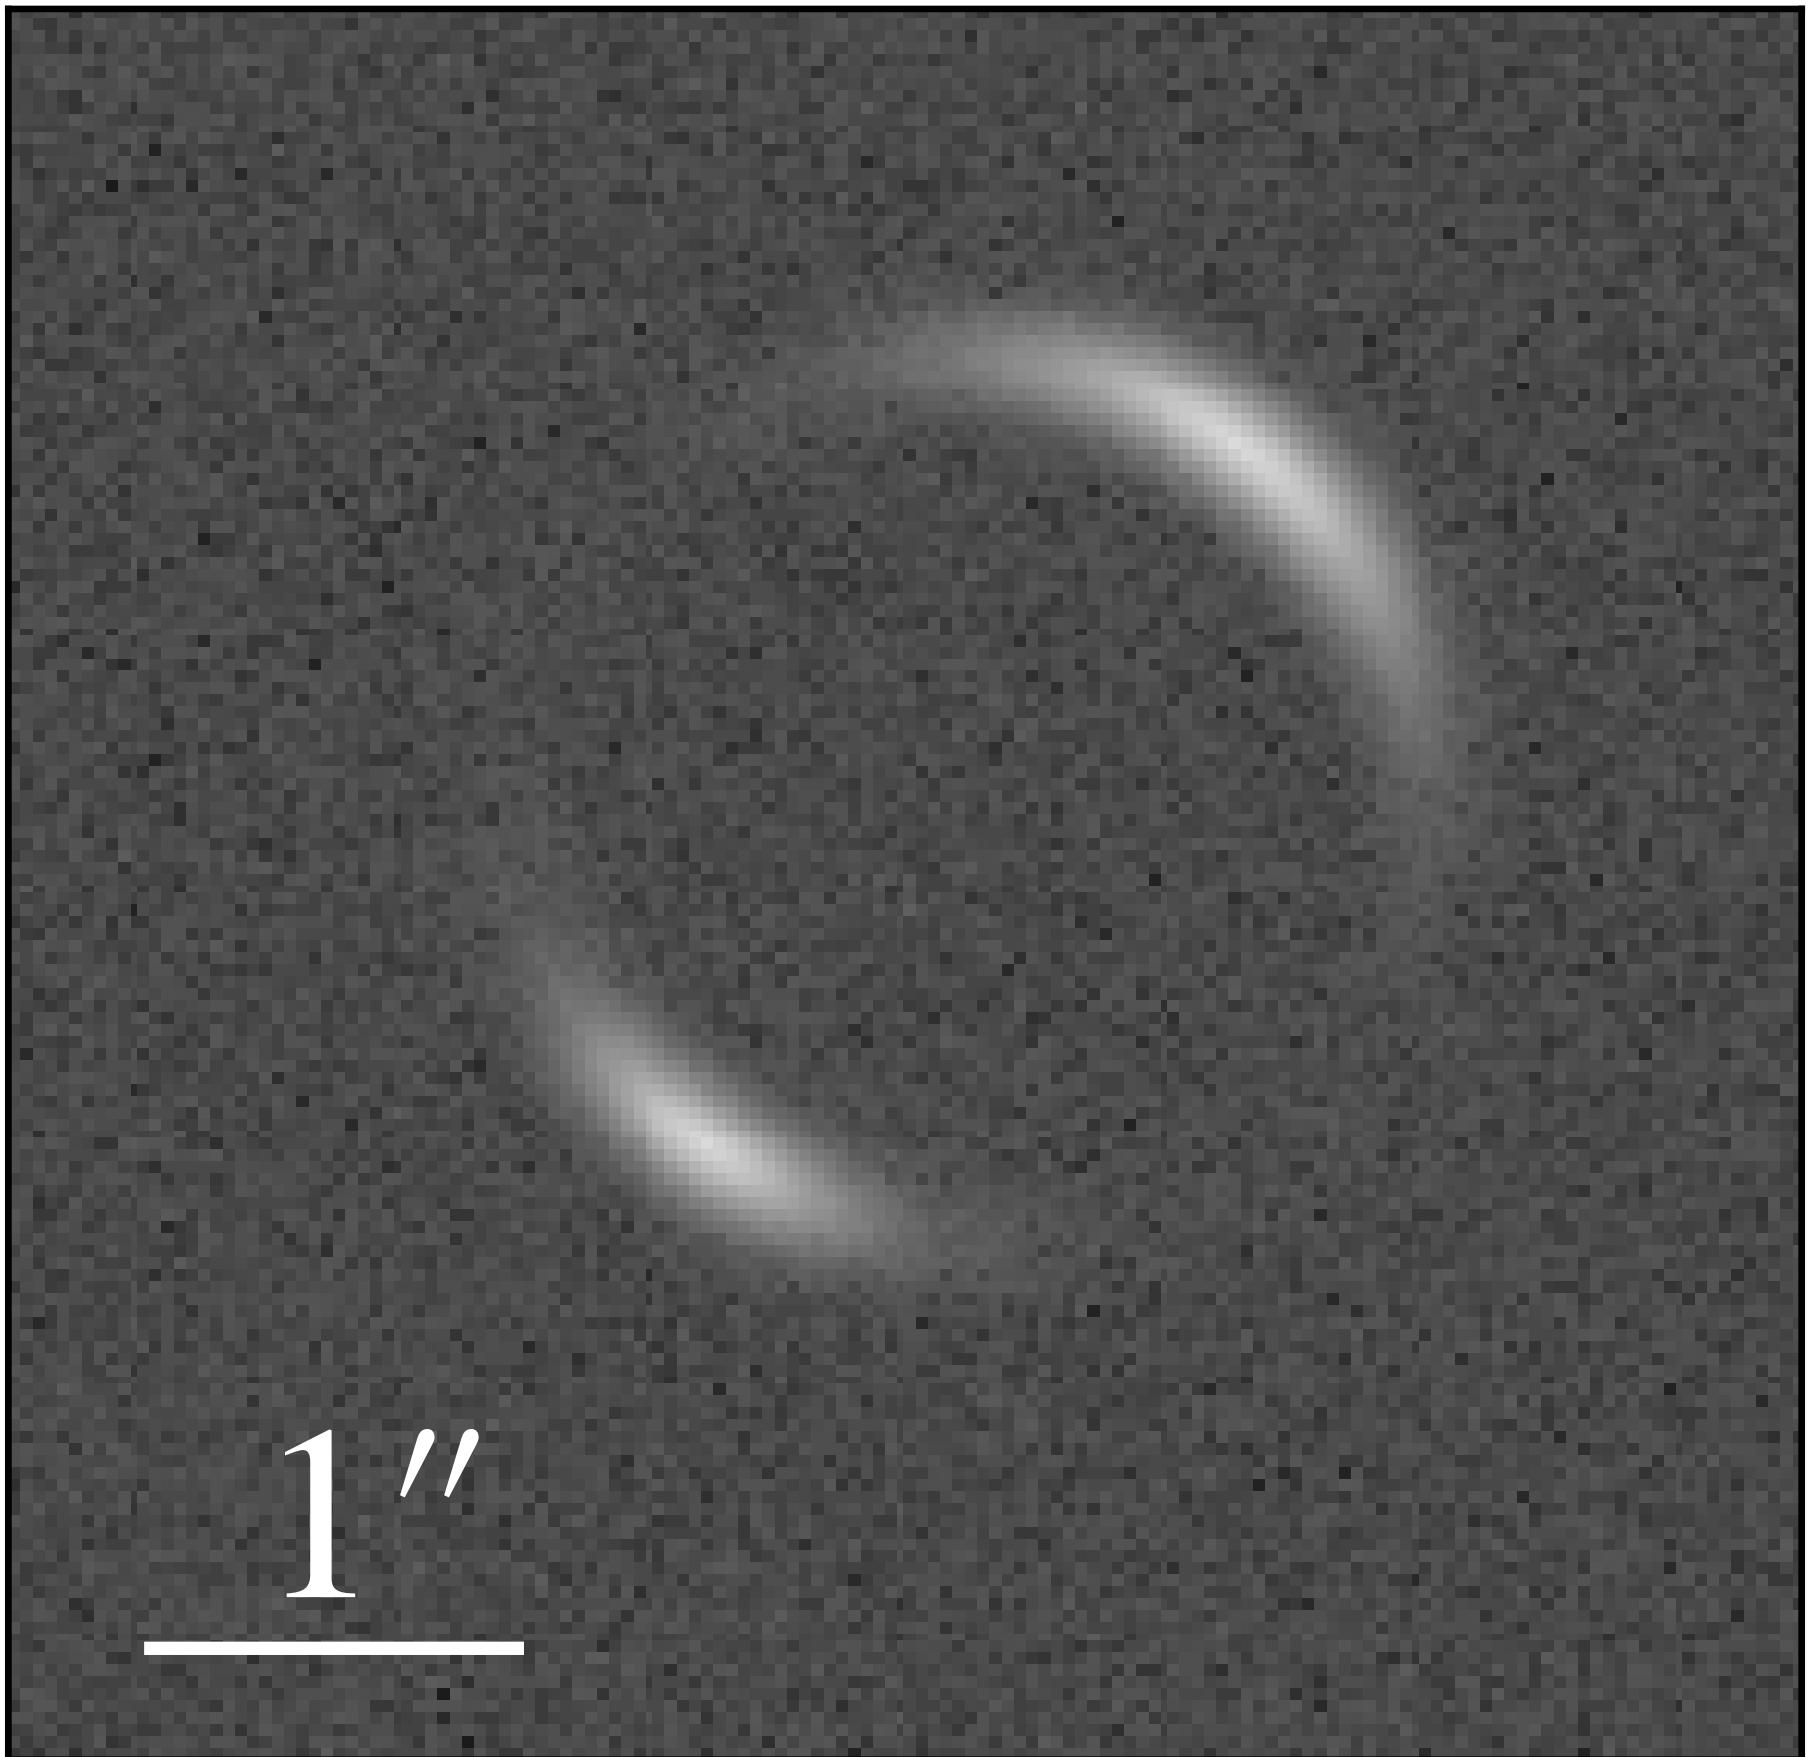
Multiple images
of quasars



Extended arcs
from galaxies

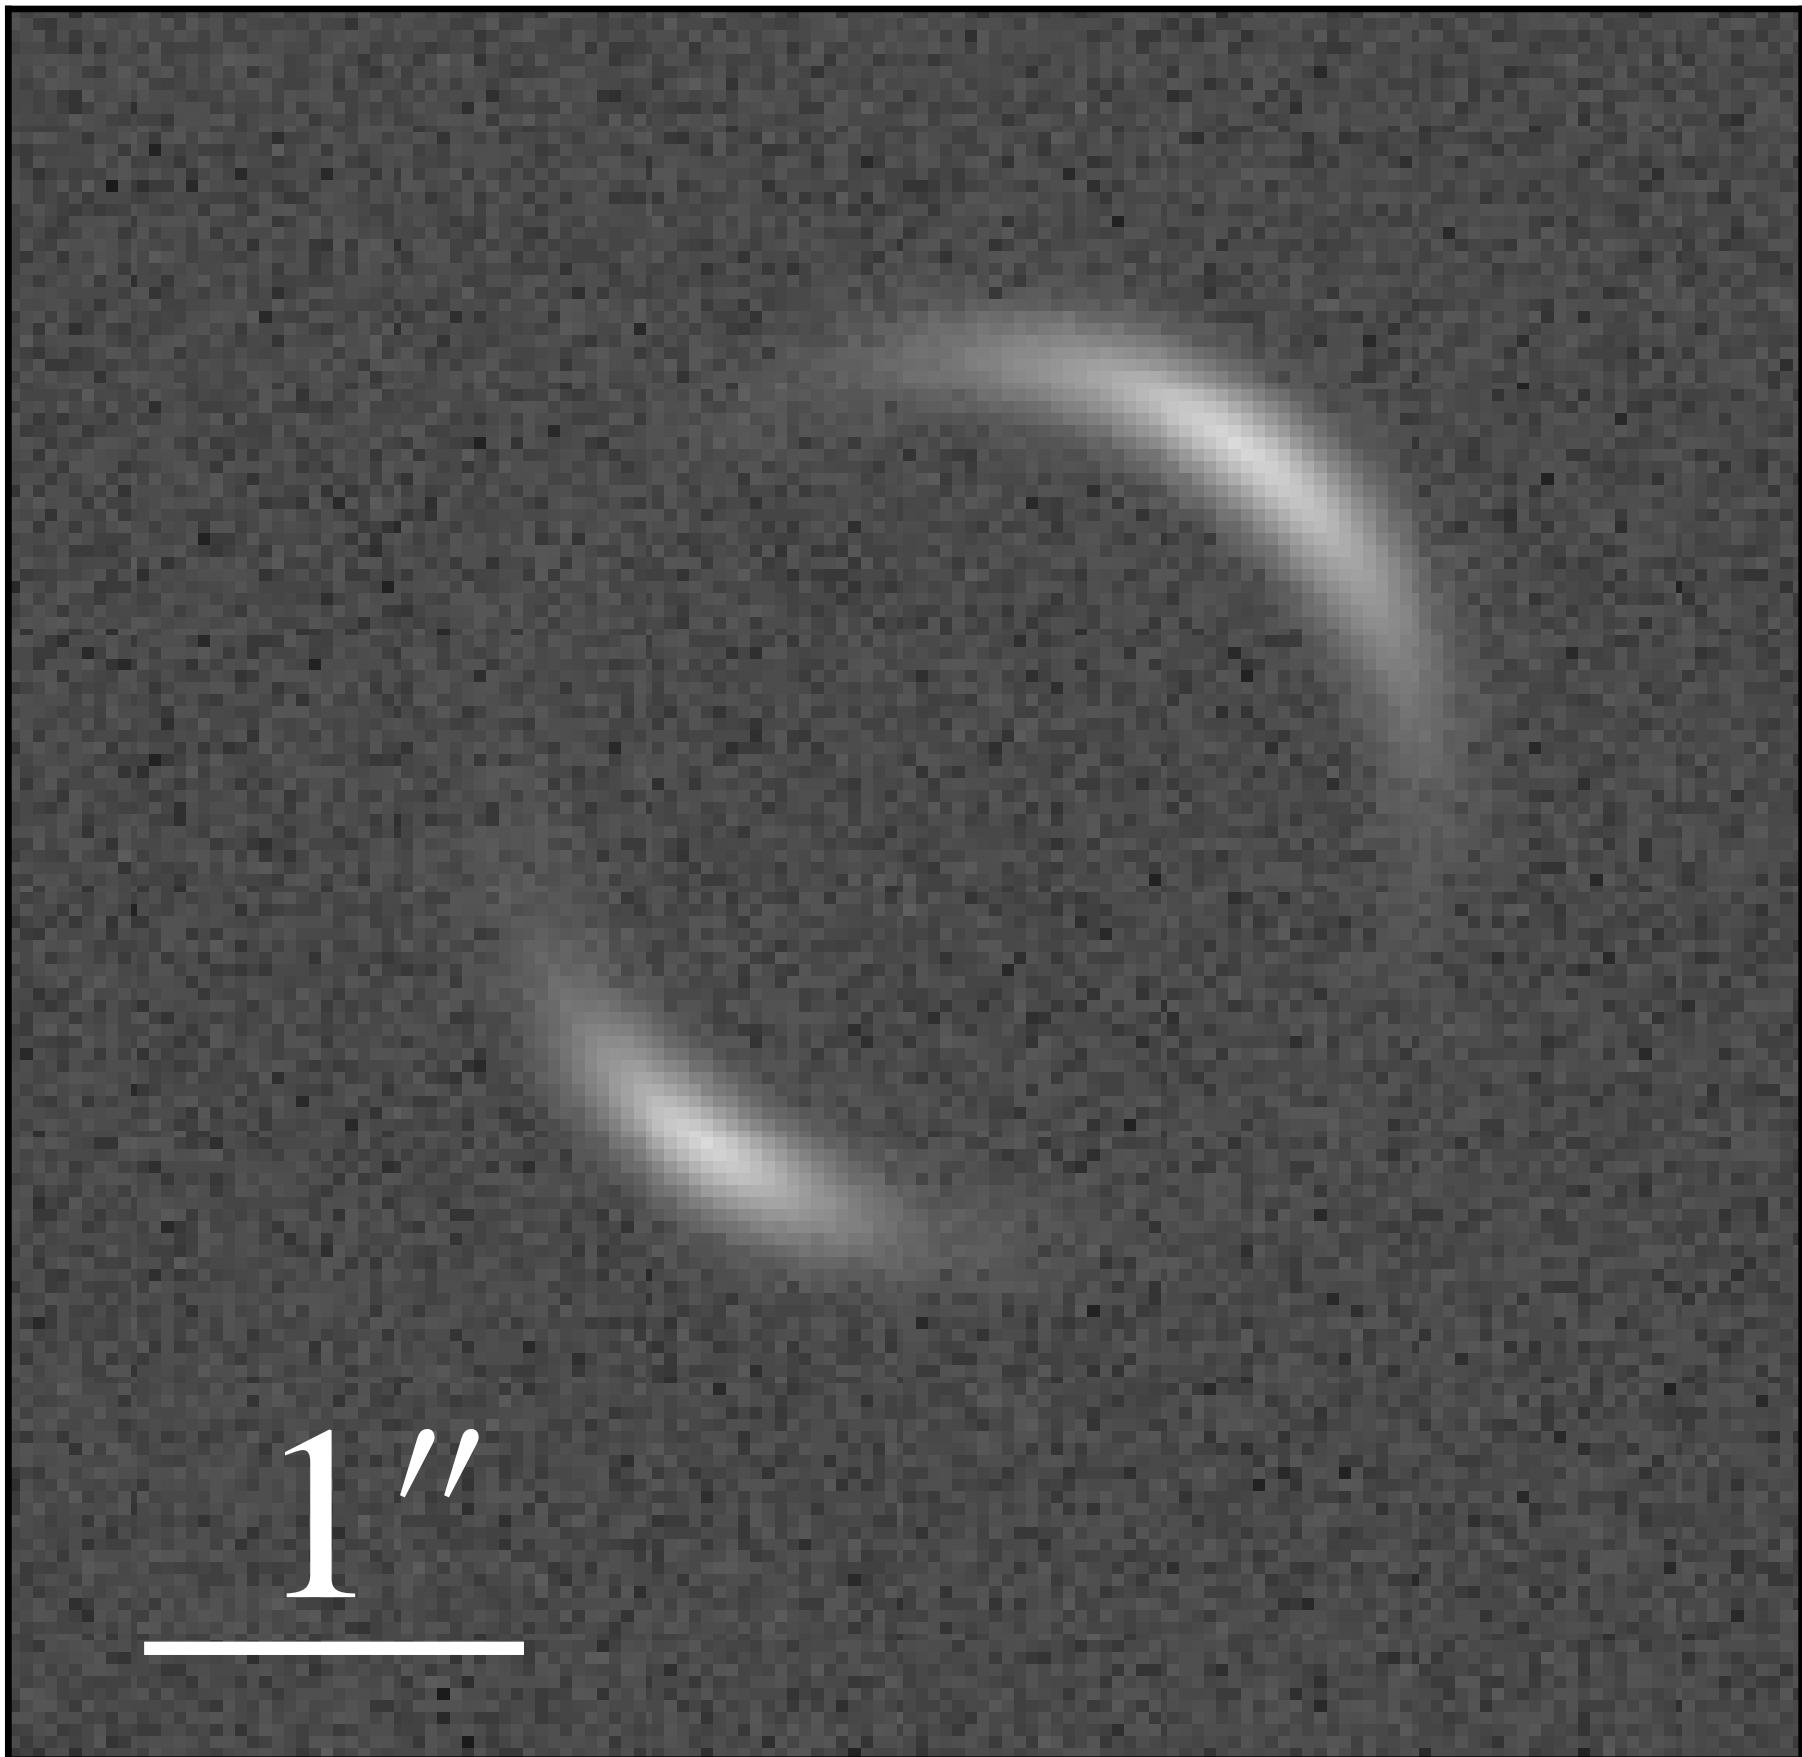


Subhalos affect strong lensing

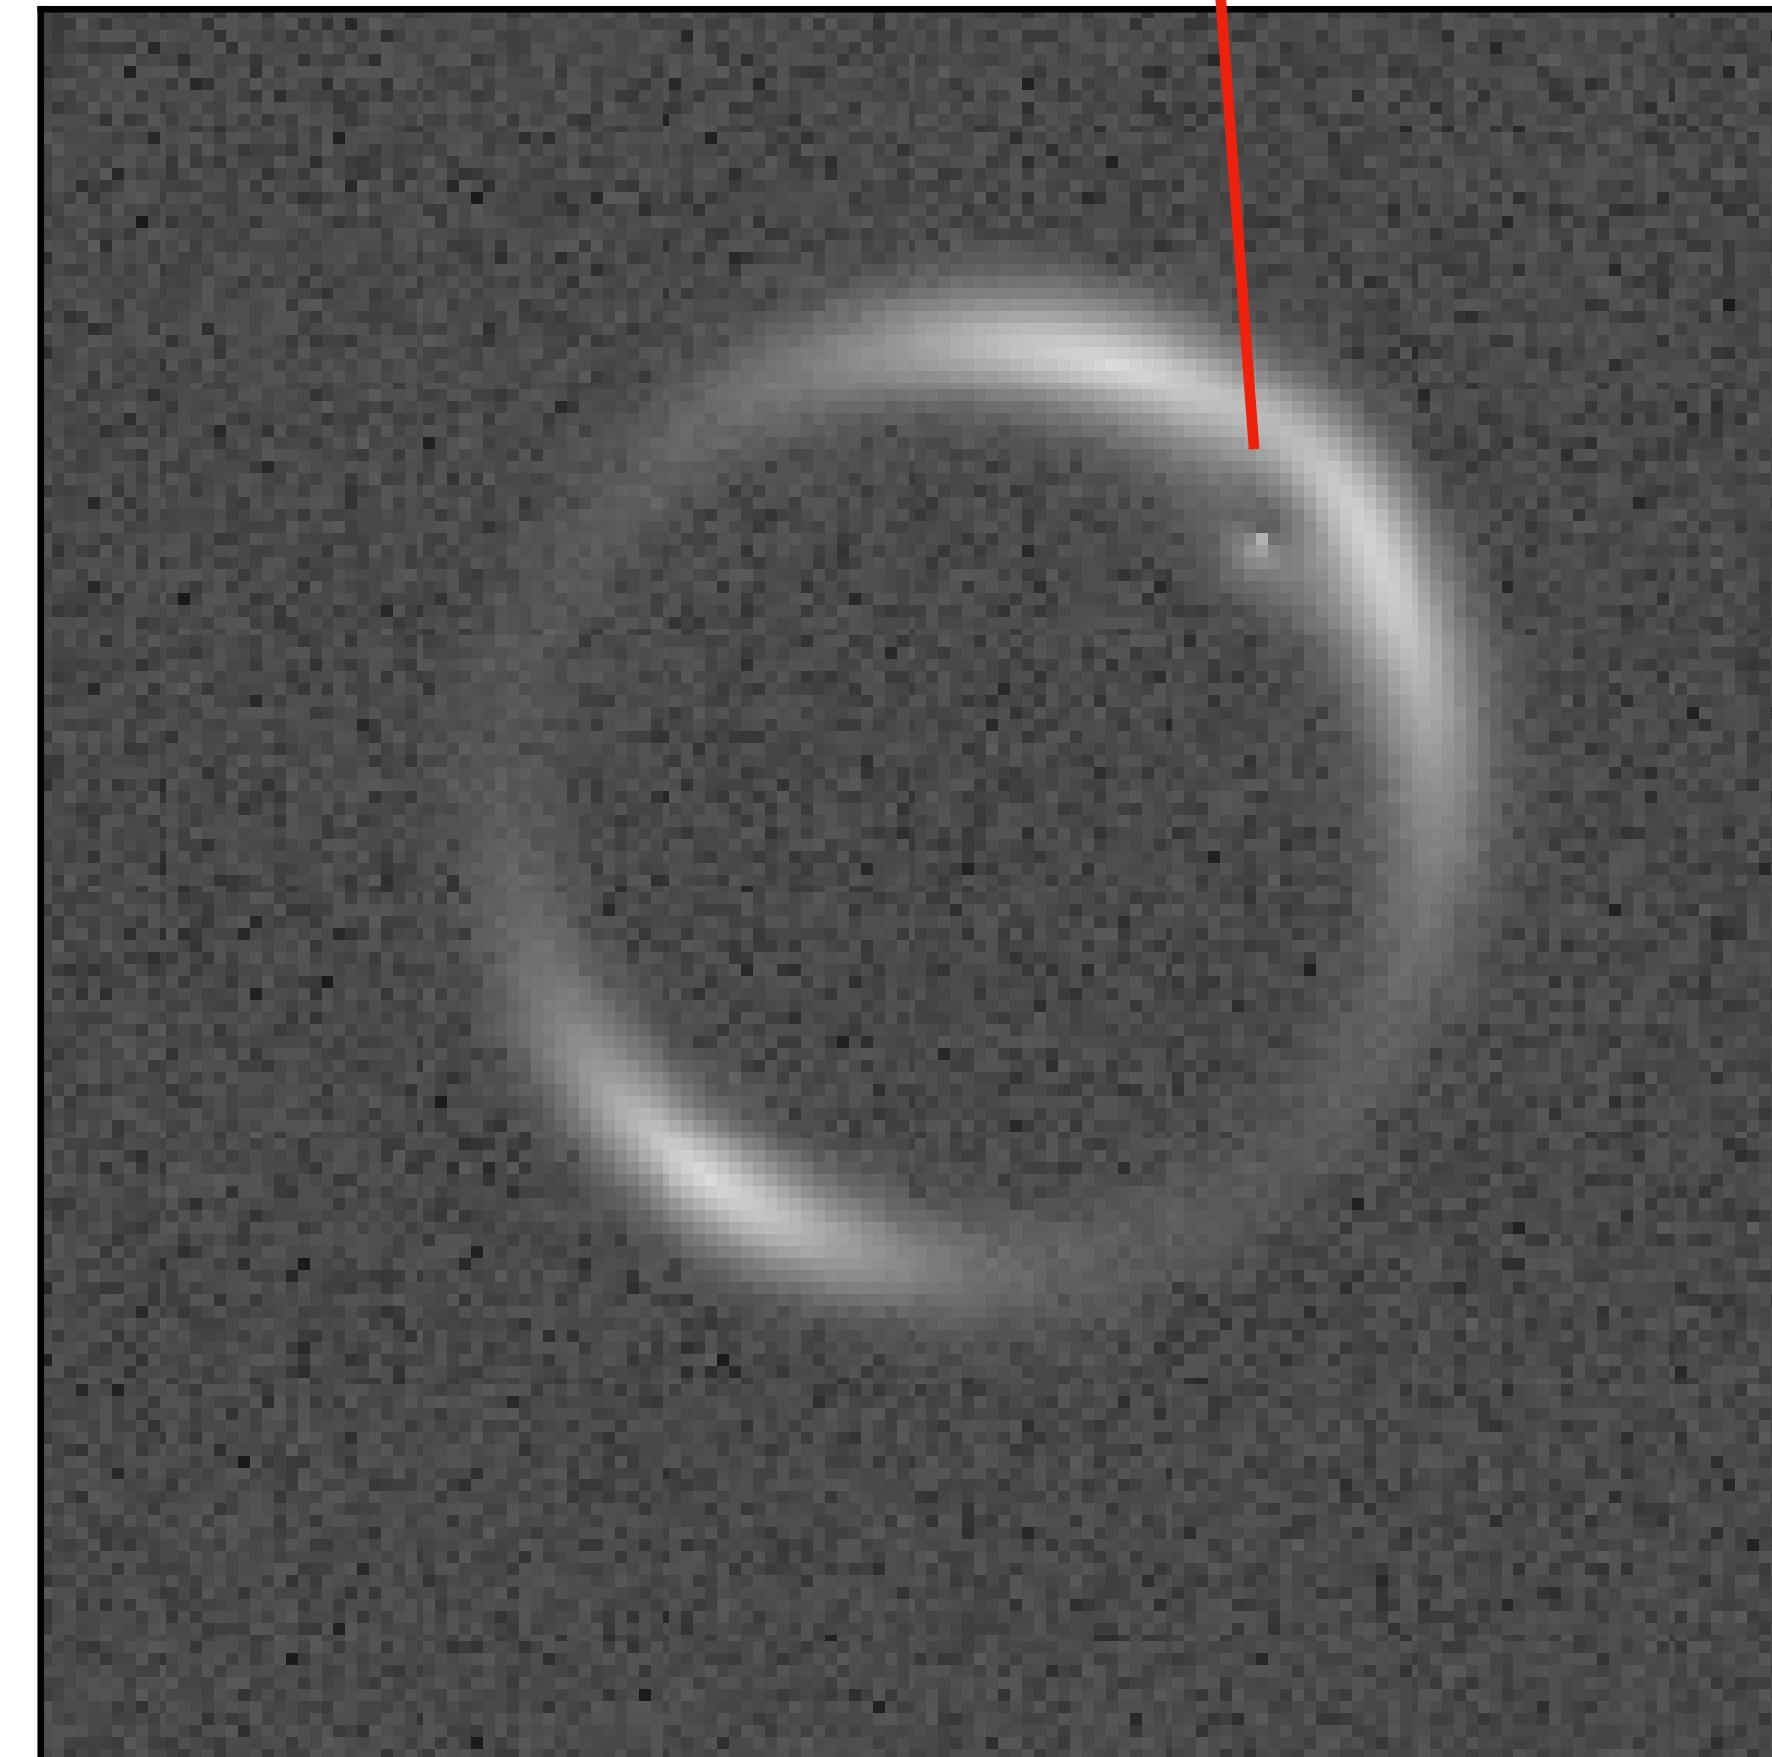


Subhalos affect strong lensing

Smooth halo only

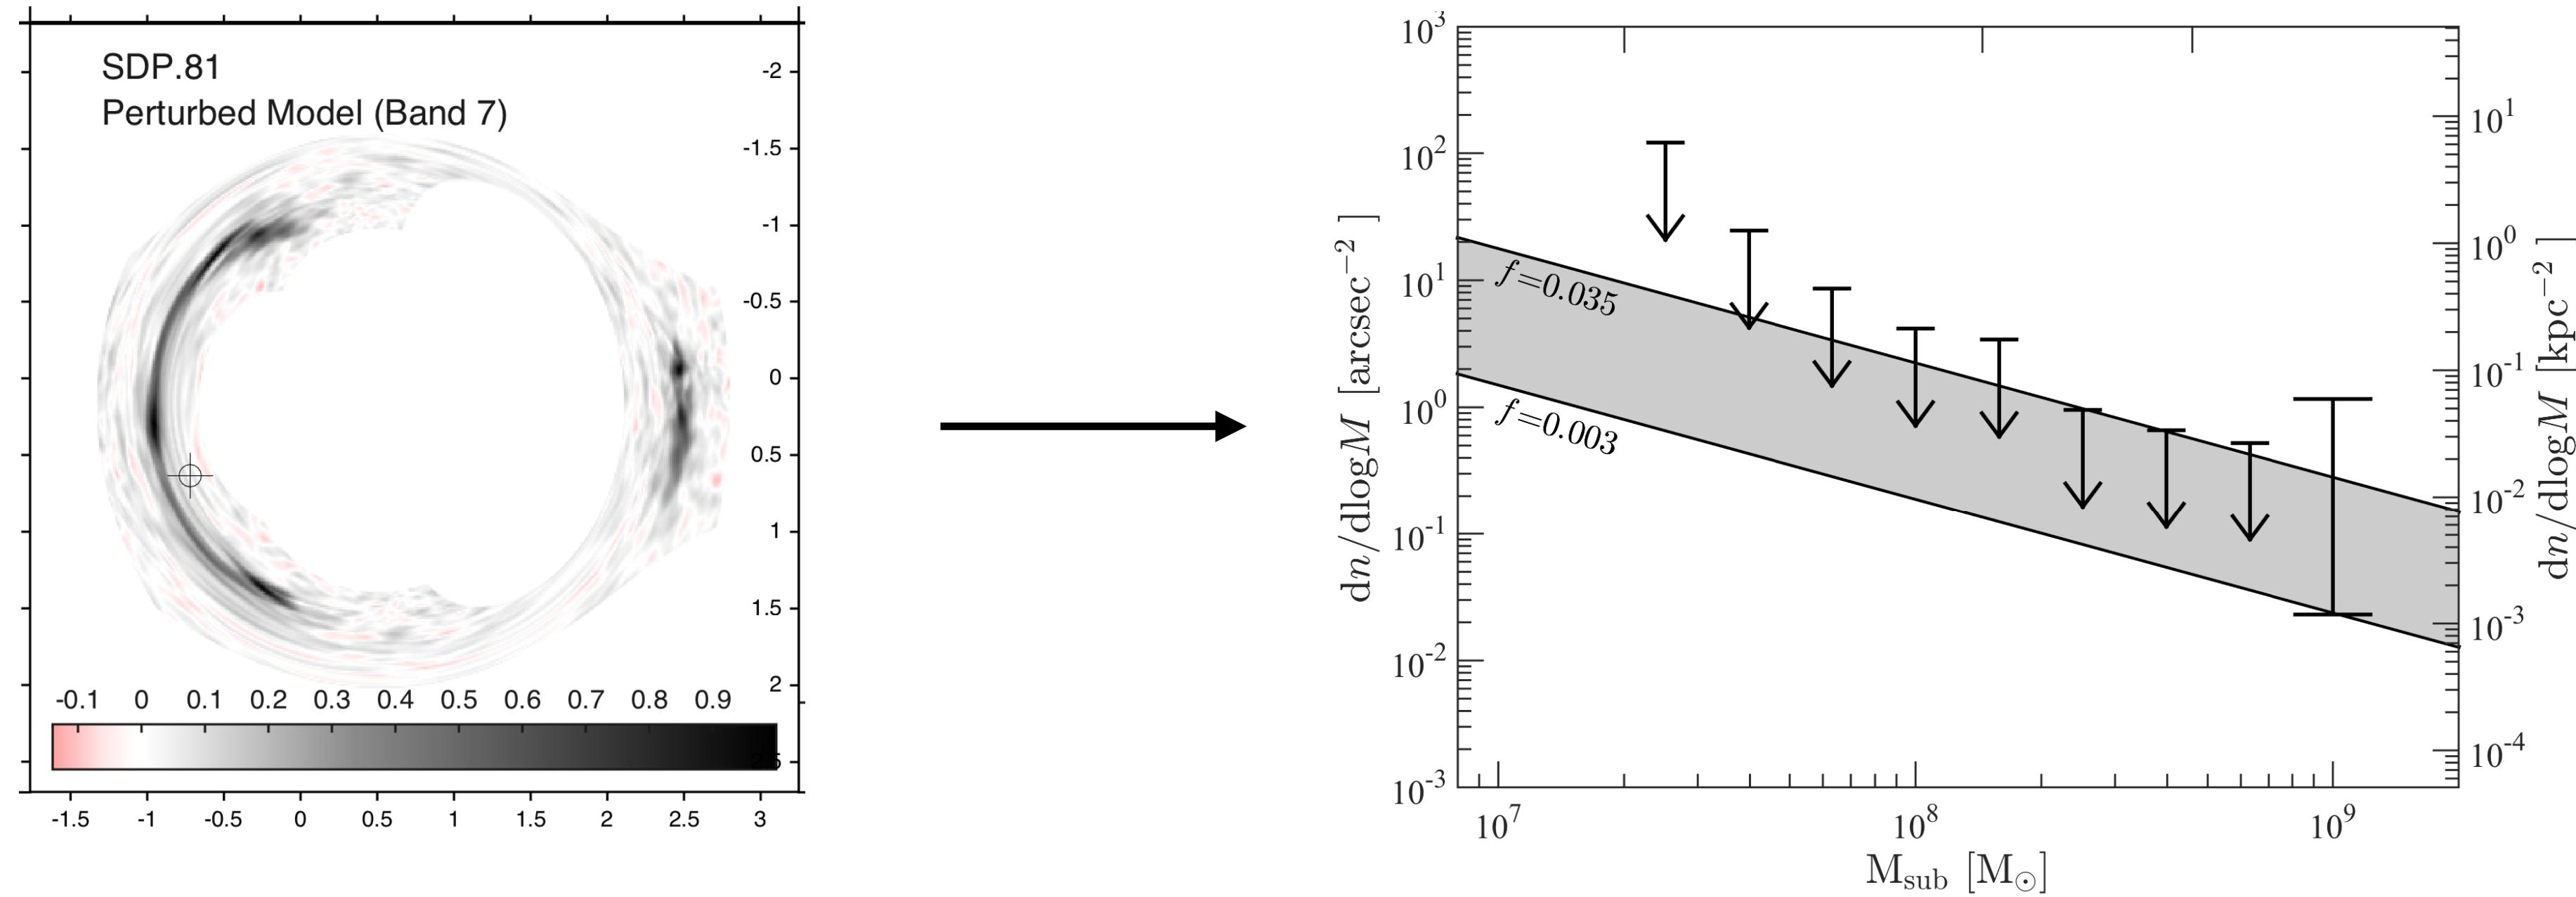


Smooth halo + **subhalo**



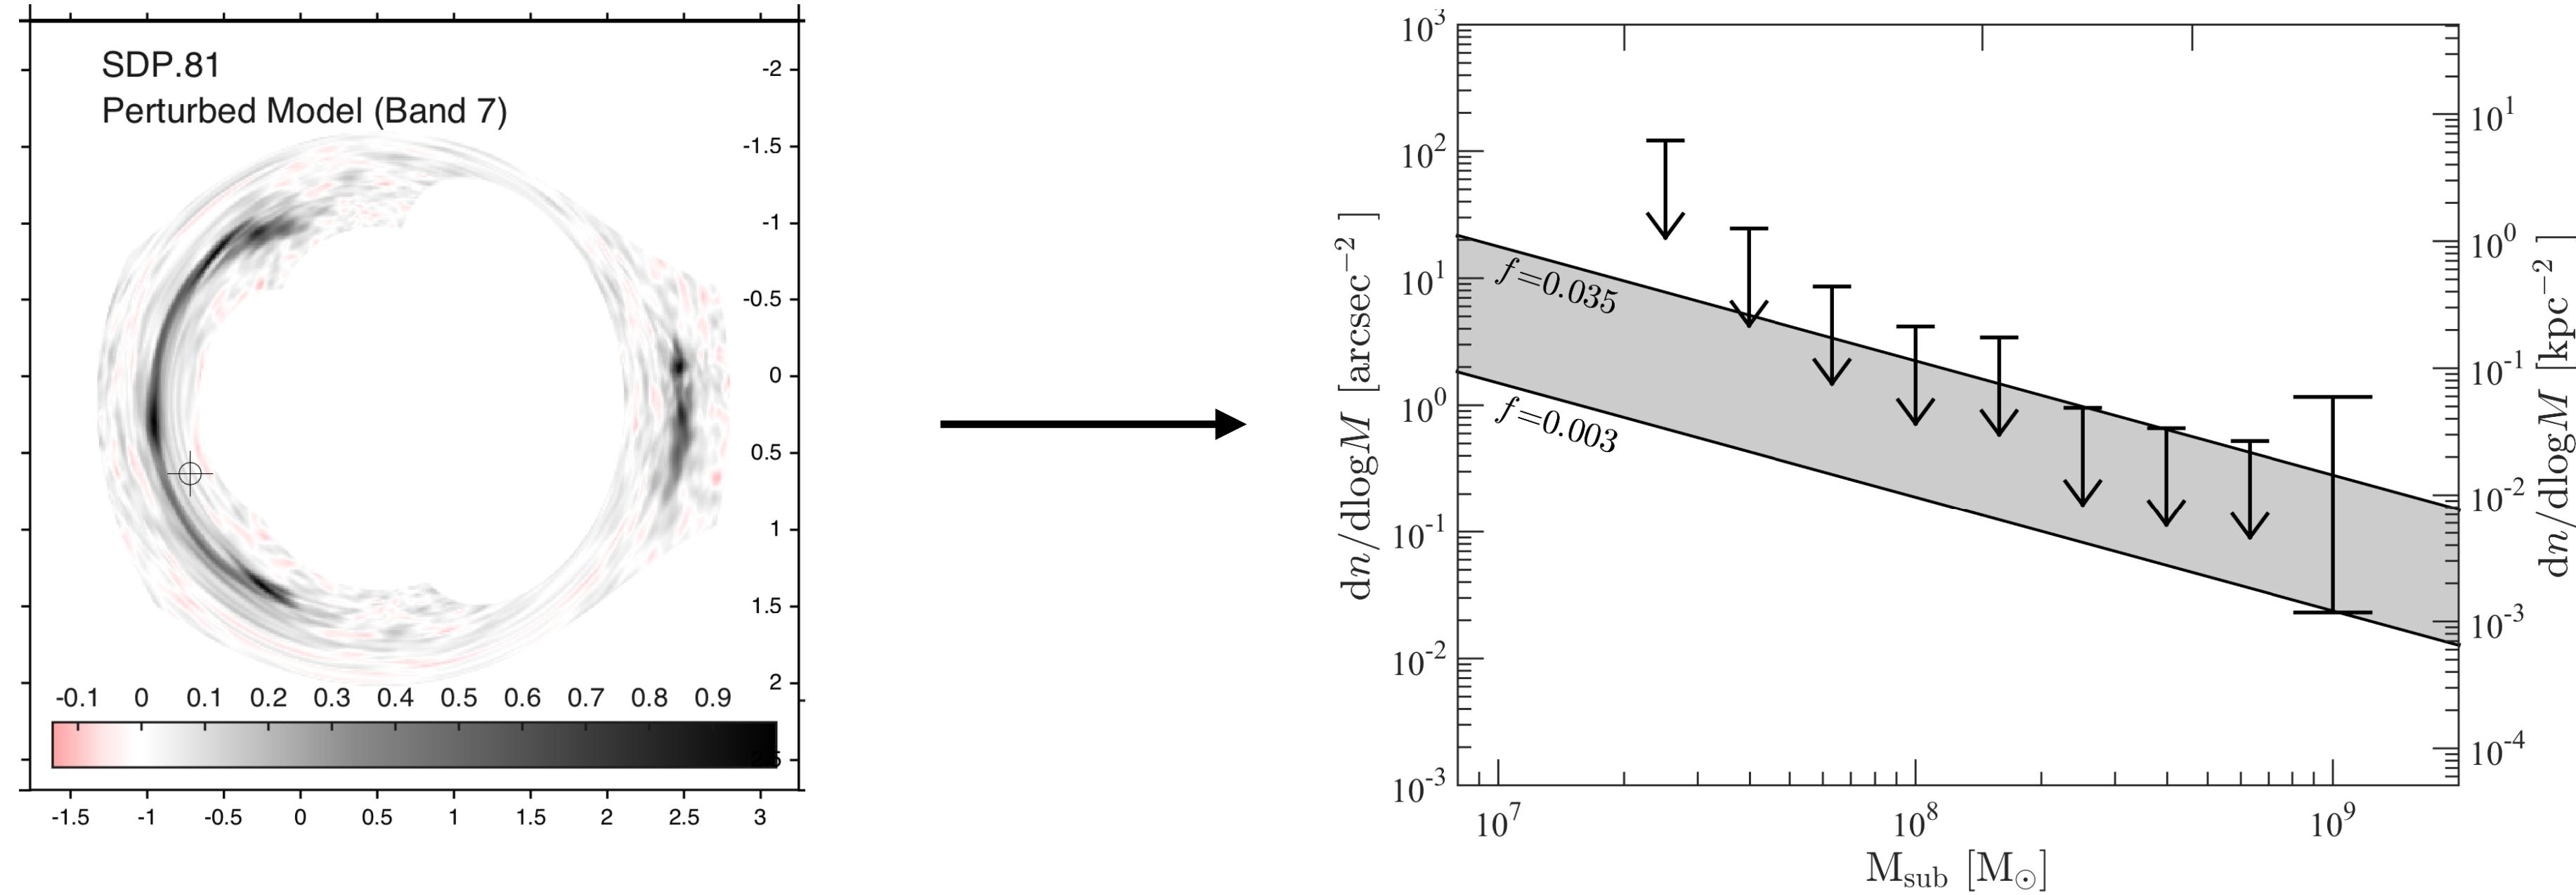
Inference clump by clump

- Individual heavy subhalos leave a relatively clear signal and can be identified [Hezaveh et al 1601.01388]



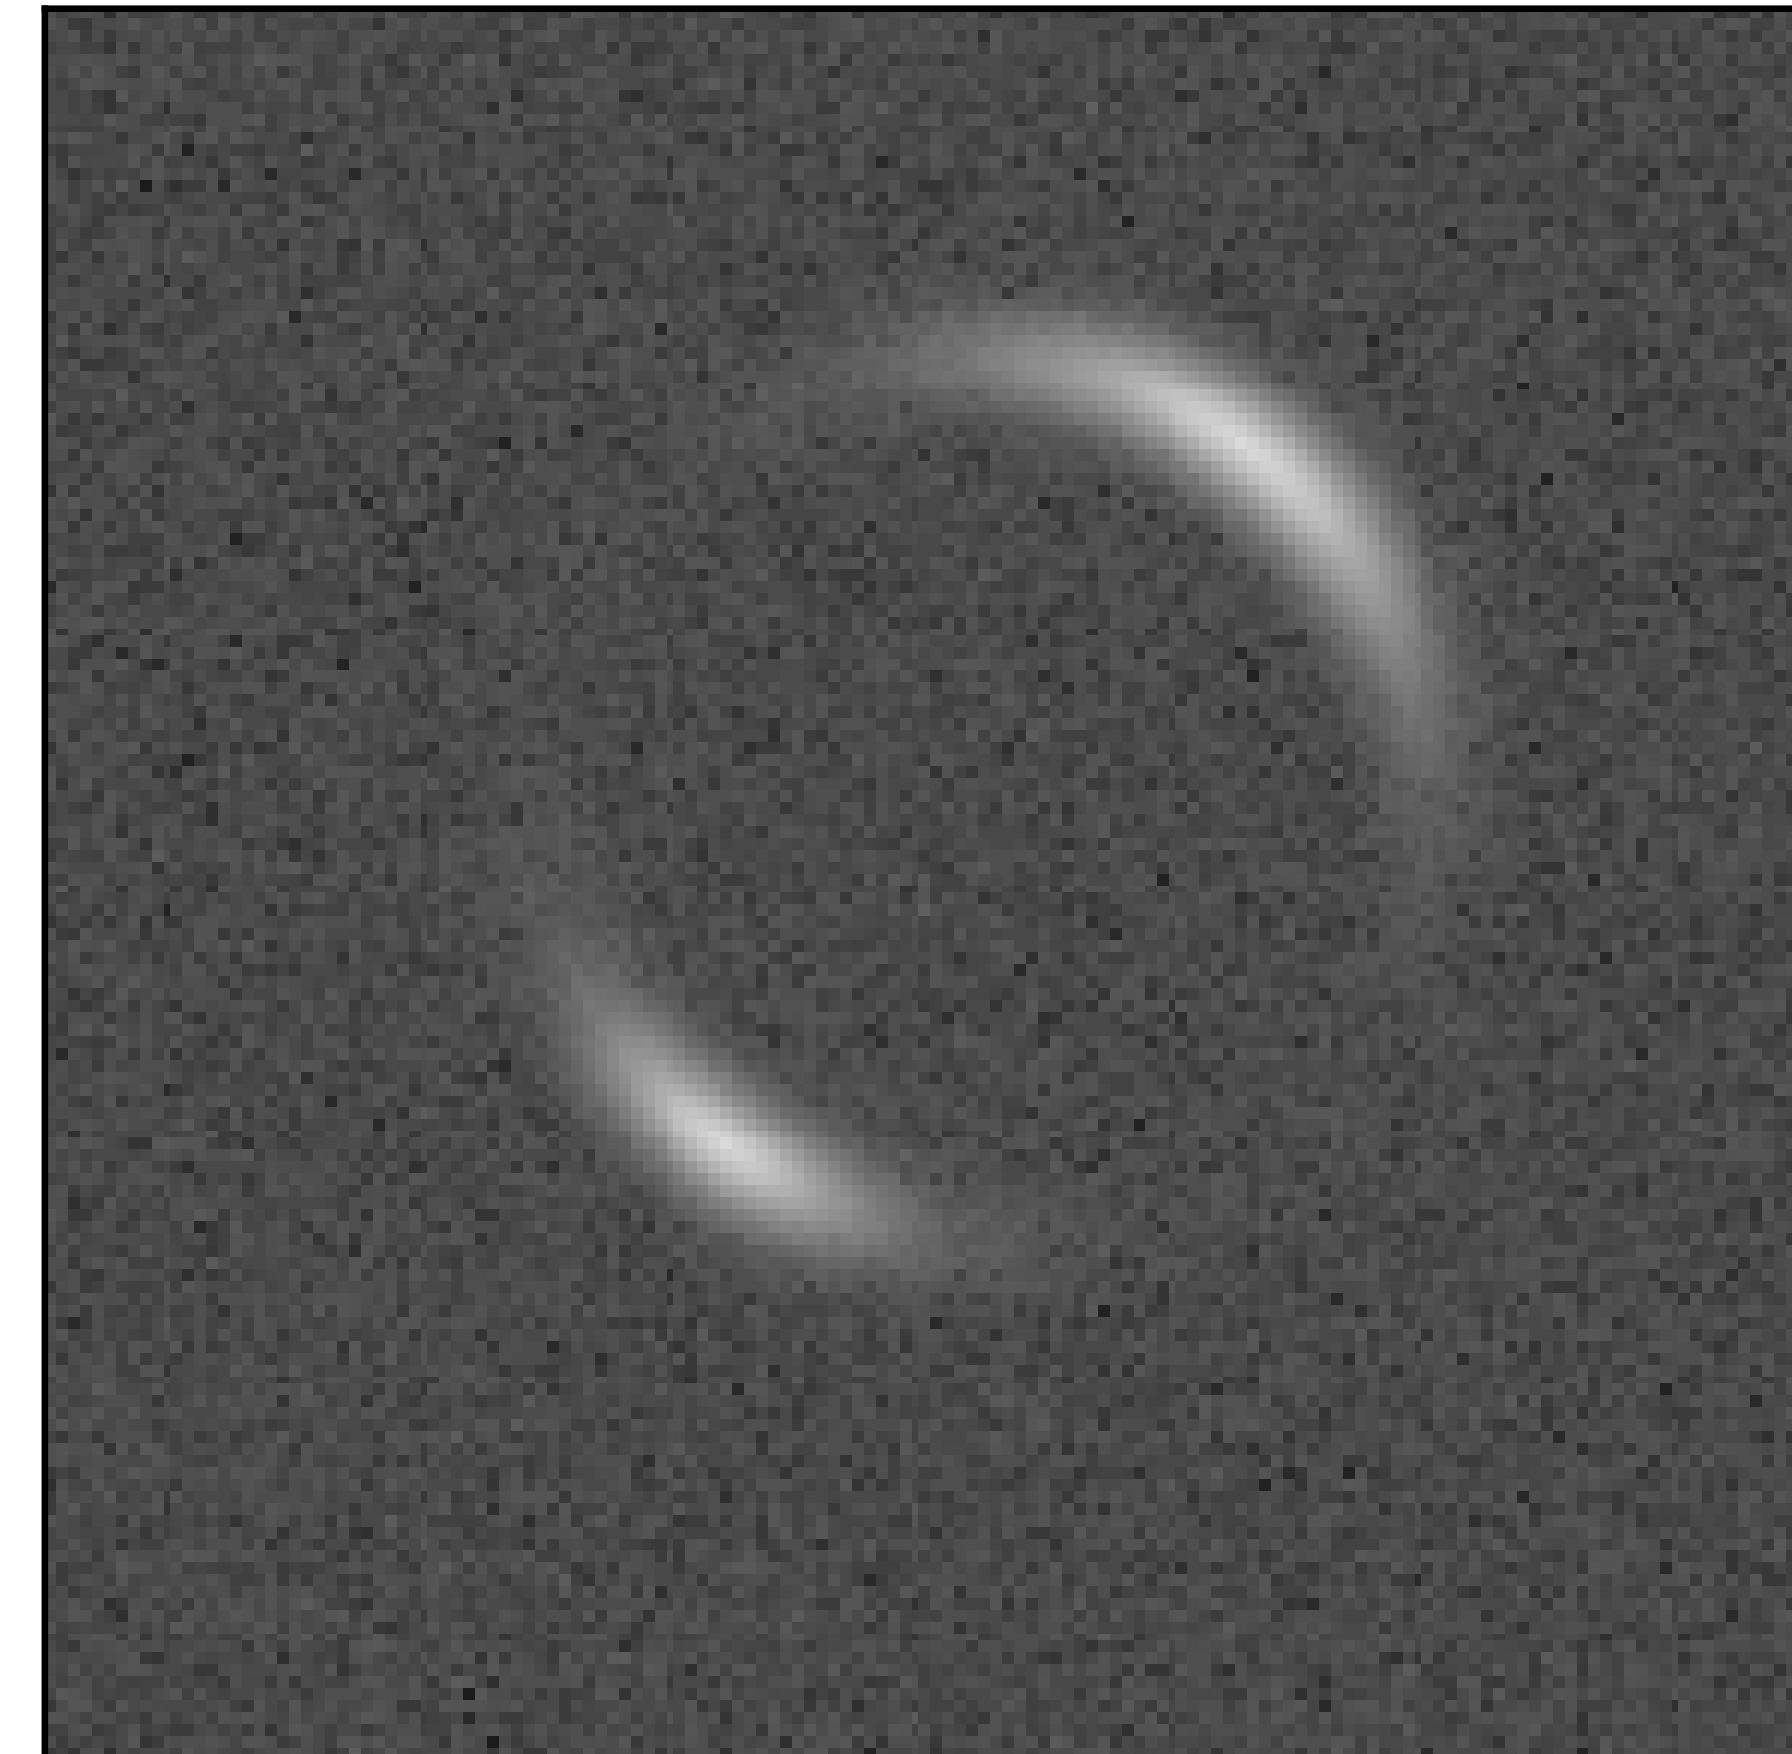
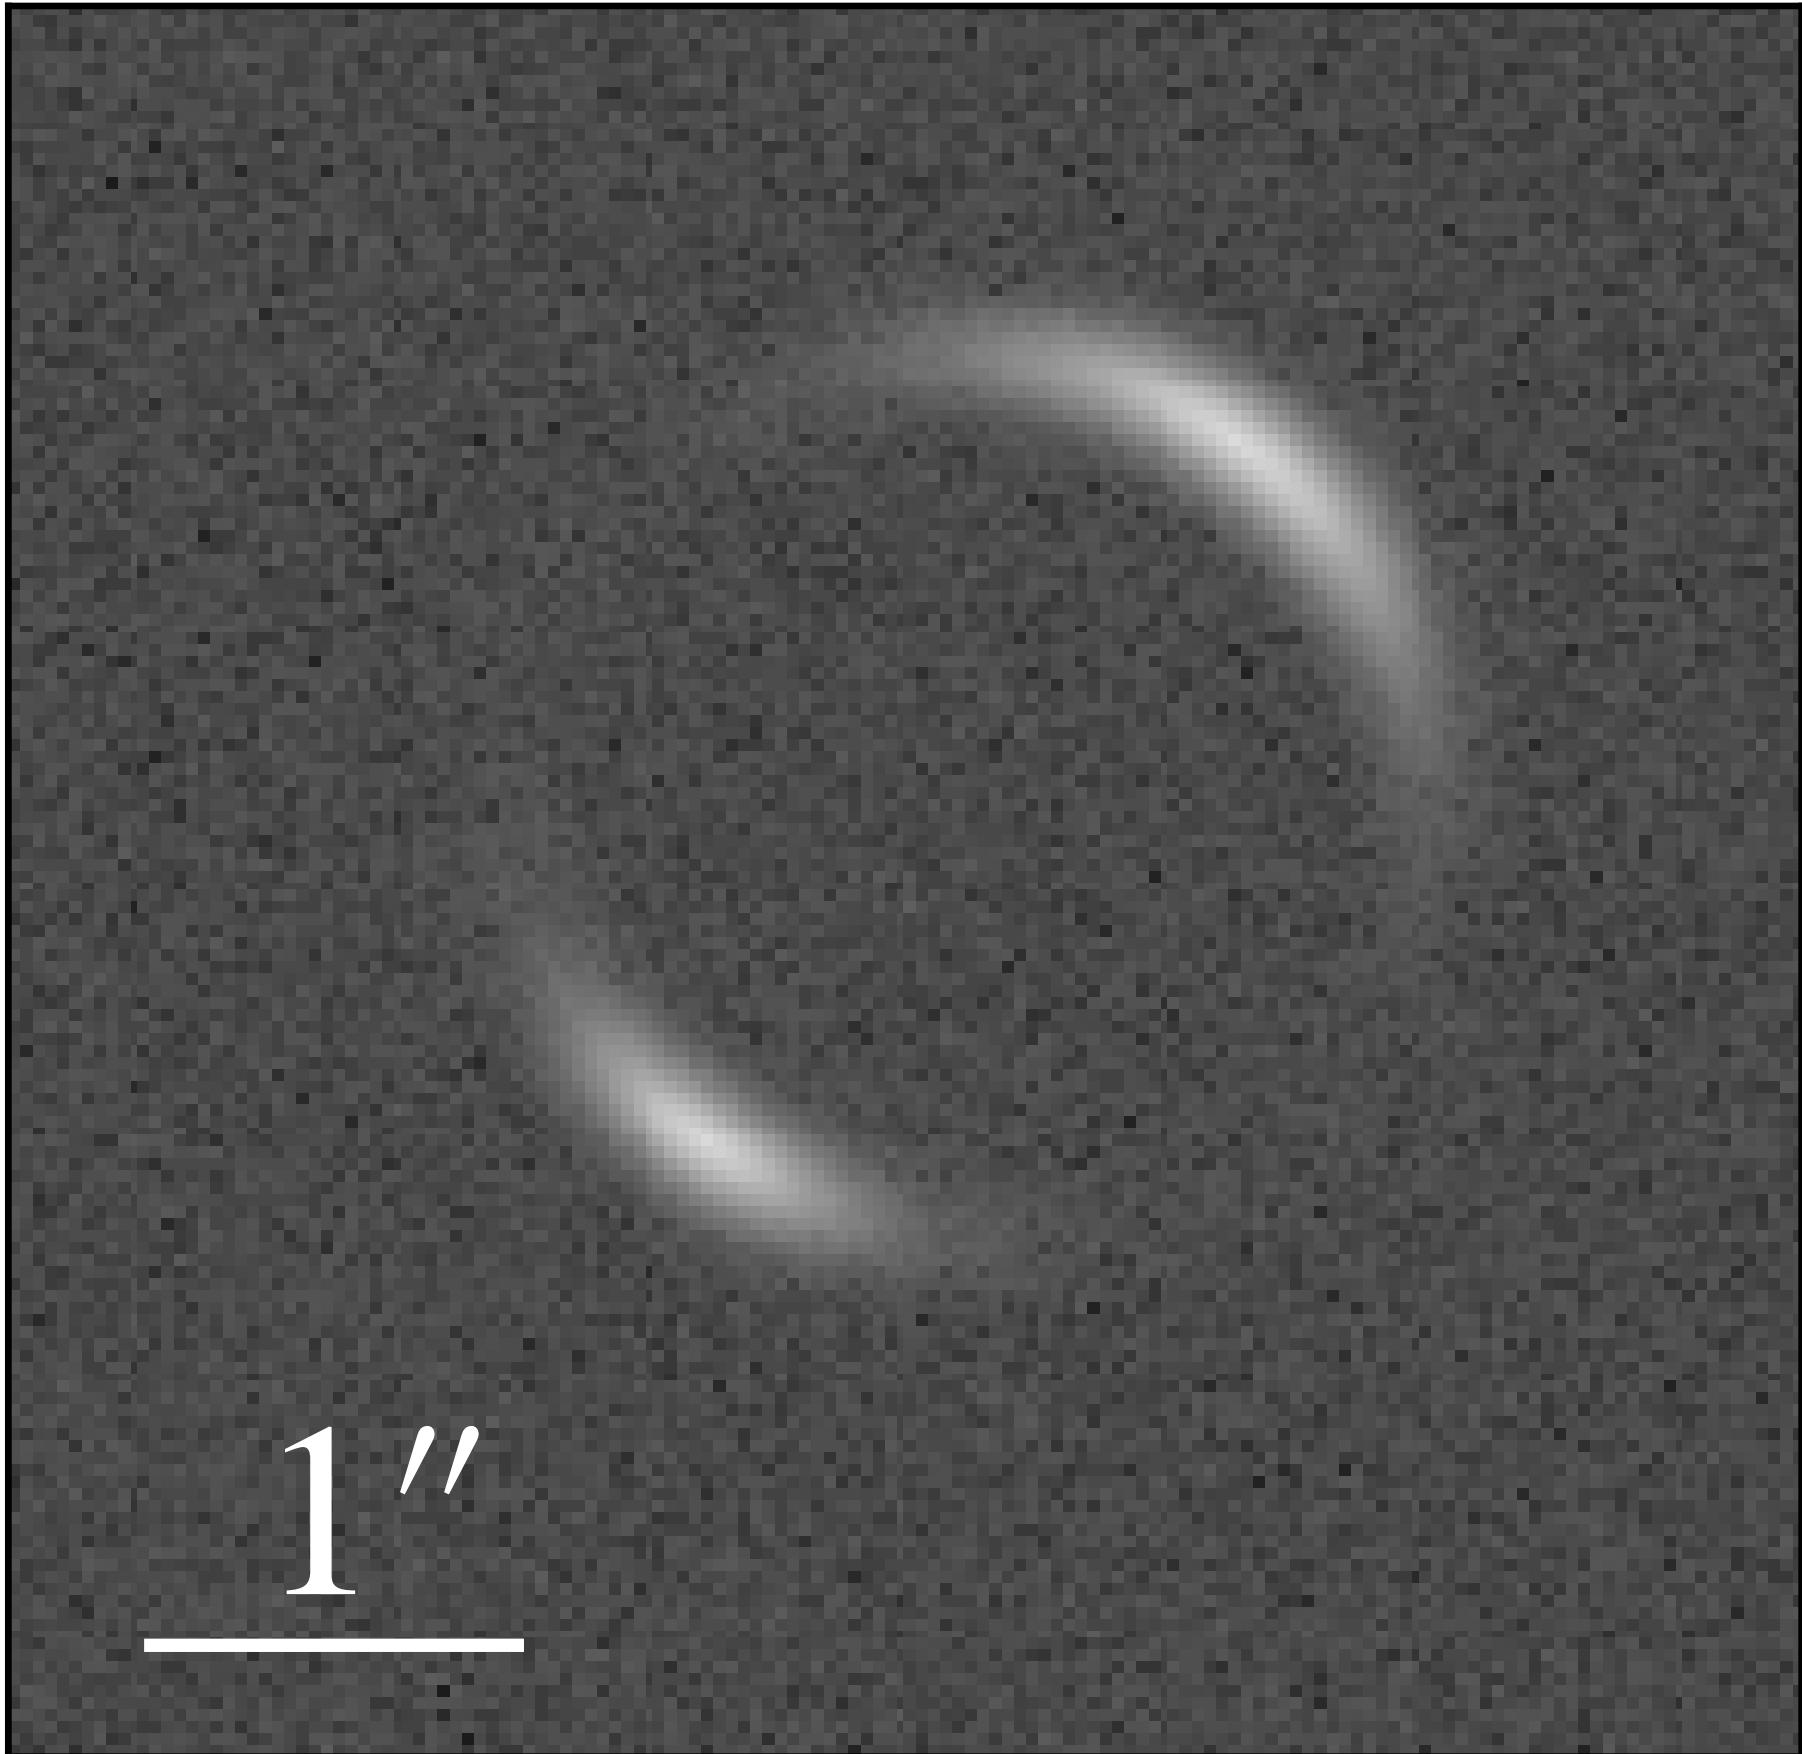
Inference clump by clump

- Individual heavy subhalos leave a relatively clear signal and can be identified [Hezaveh et al 1601.01388]



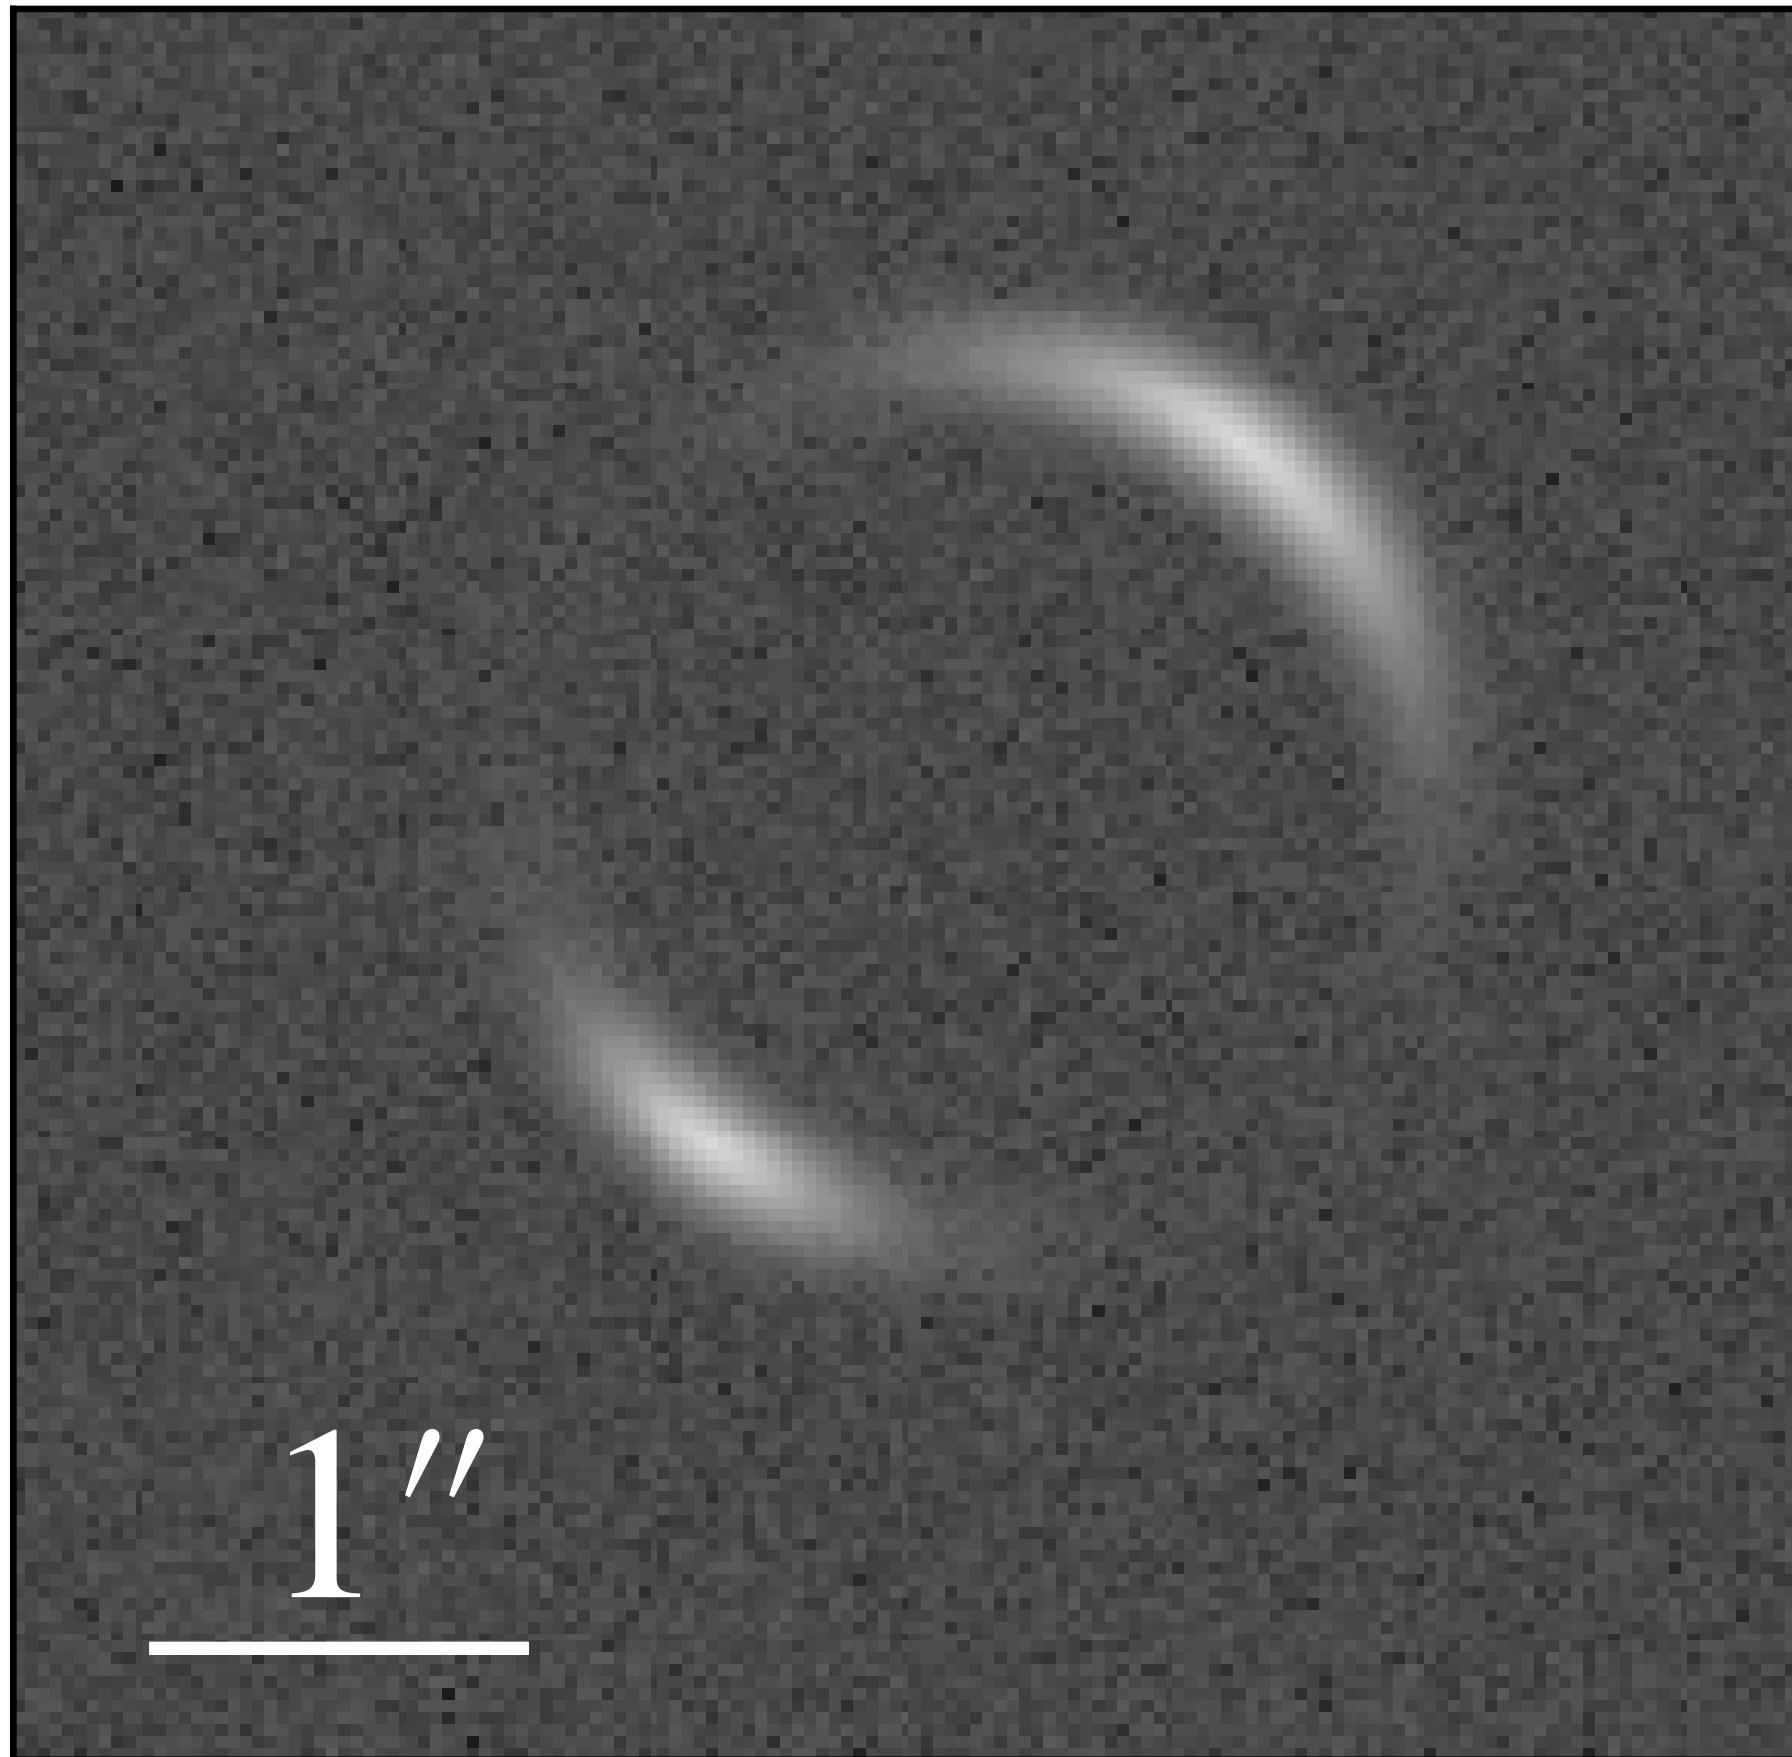
- Transdimensional inference allows simultaneously inferring multiple individual subhalos and population-level parameters [B. Brewer et al 1508.00662; T. Daylan et al 1706.06111]
- Works for subhalos $\sim 10^9 M_{\odot}$, scales poorly to many lighter subhalos (or many observed lenses)

Many light subhalos collectively affect strong lensing

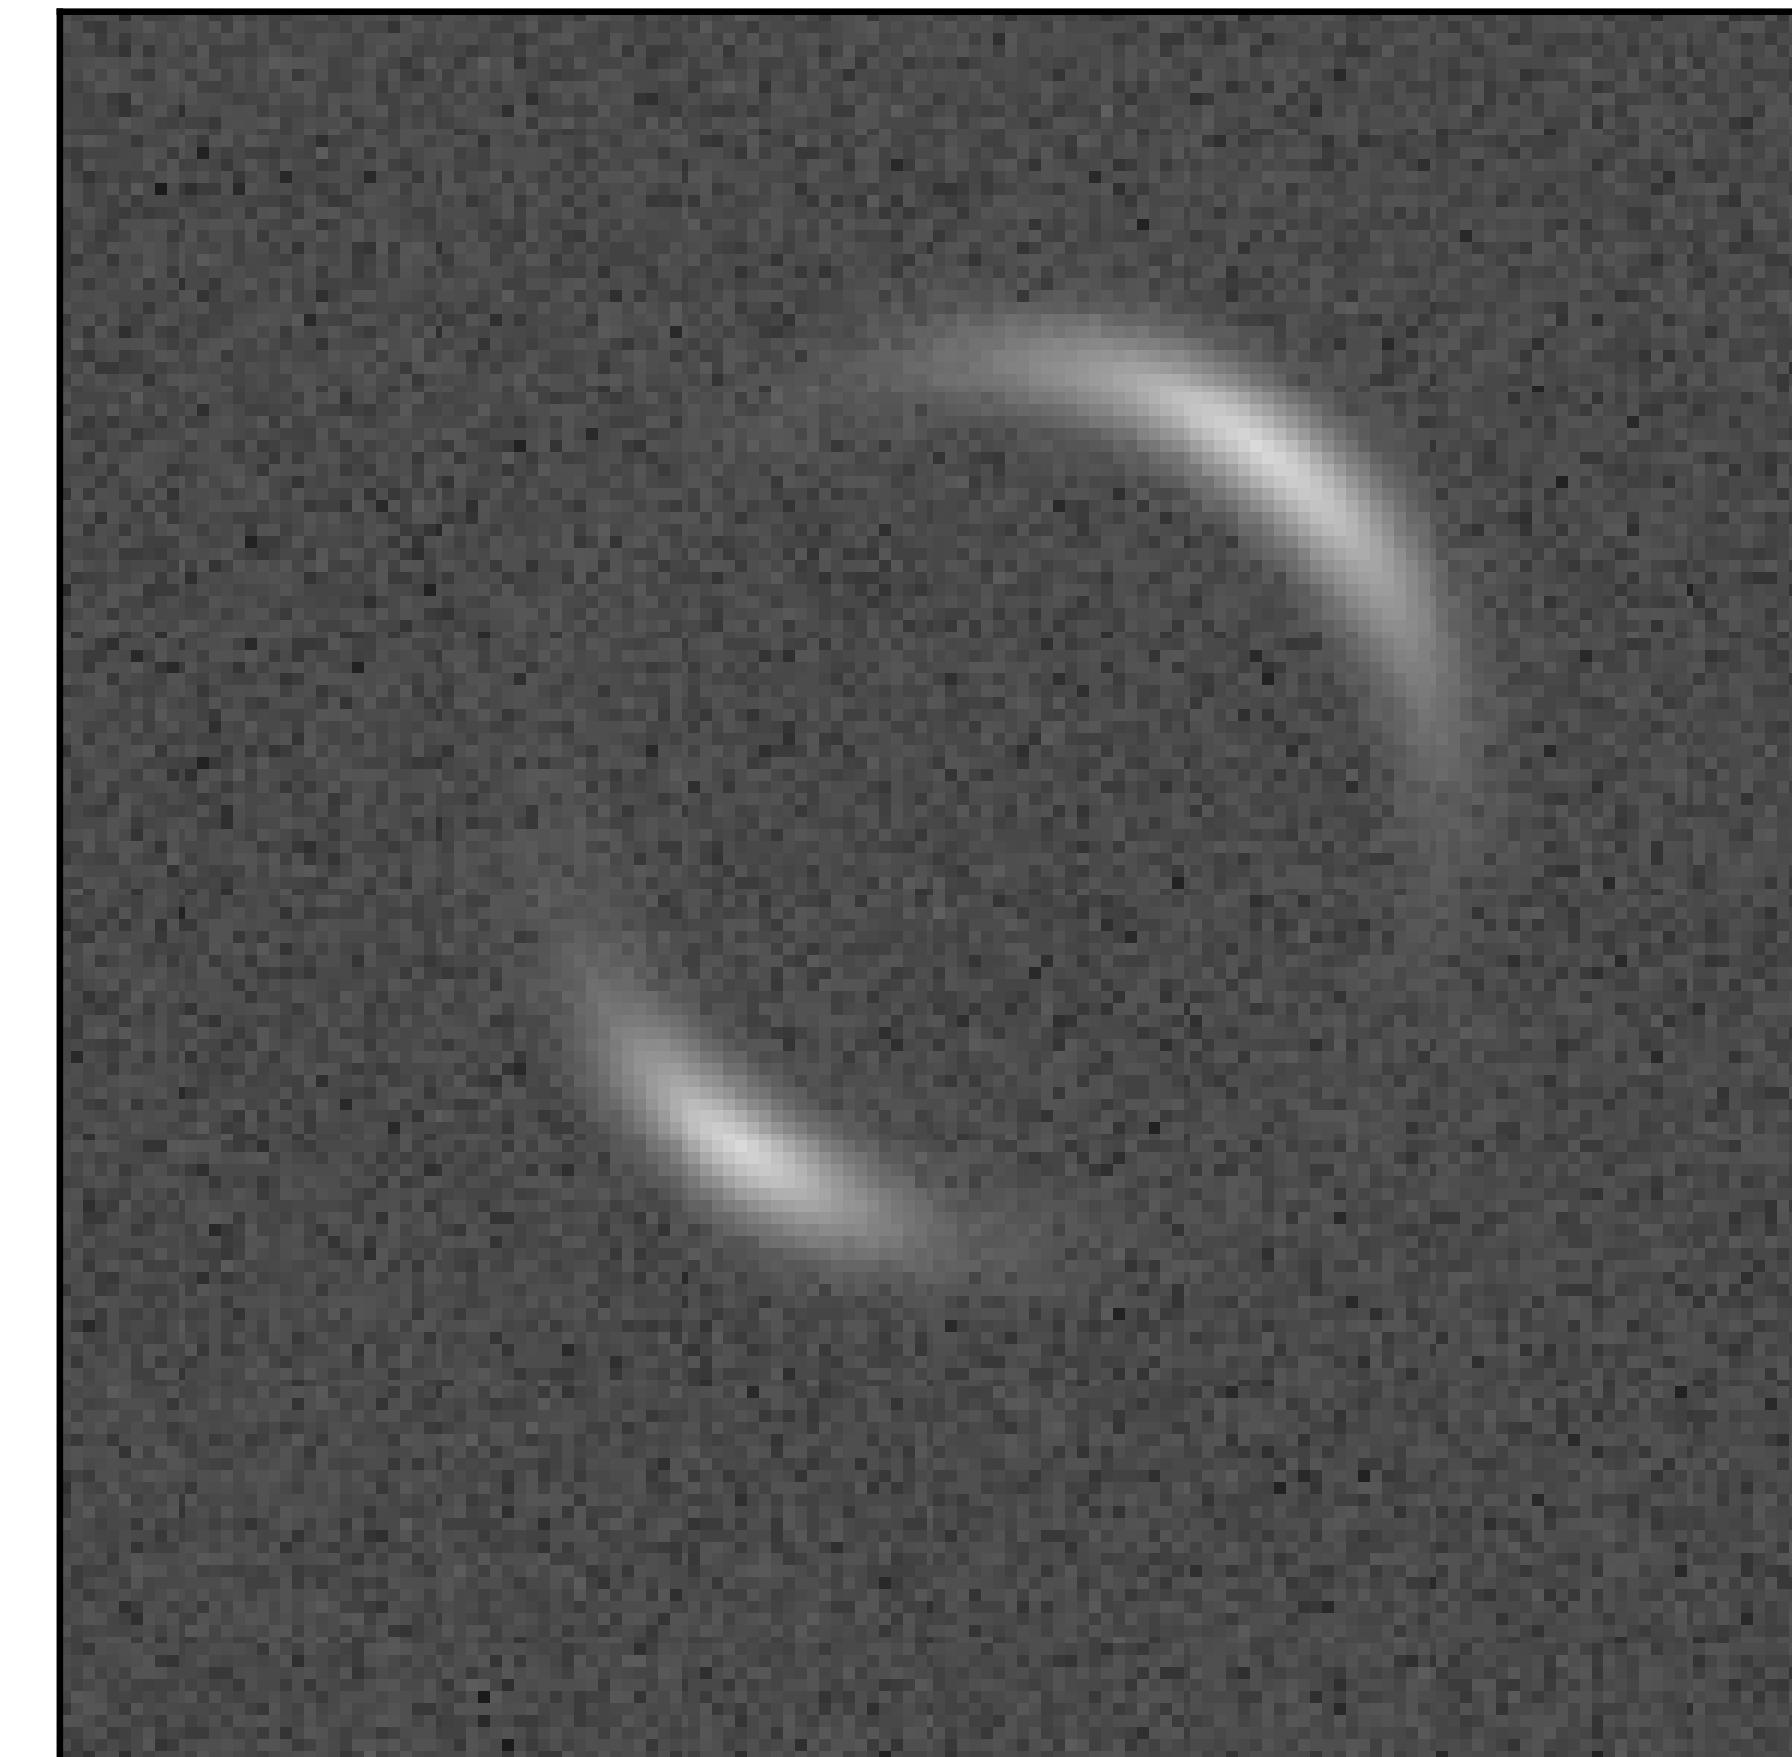


Many light subhalos collectively affect strong lensing

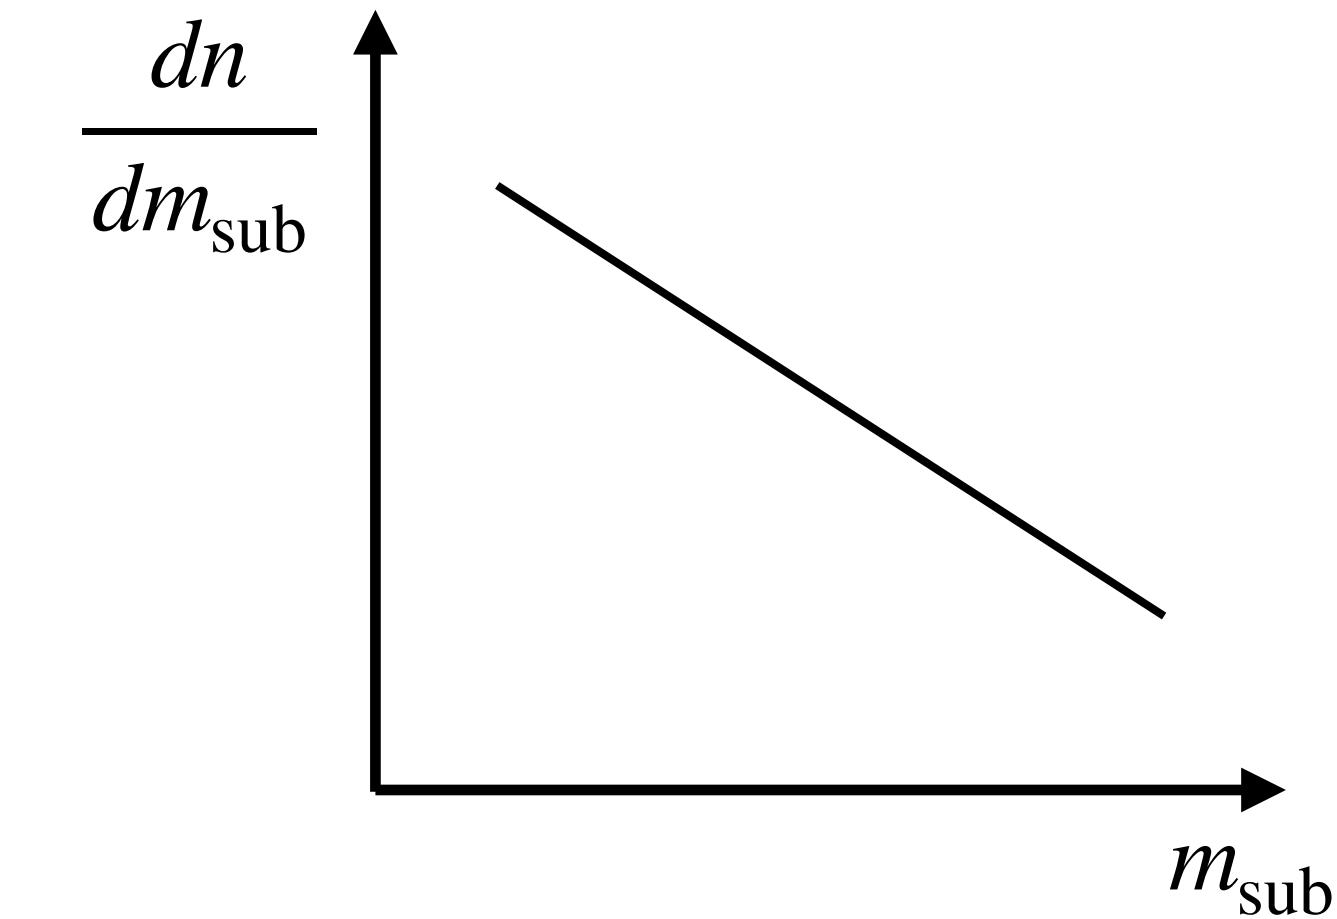
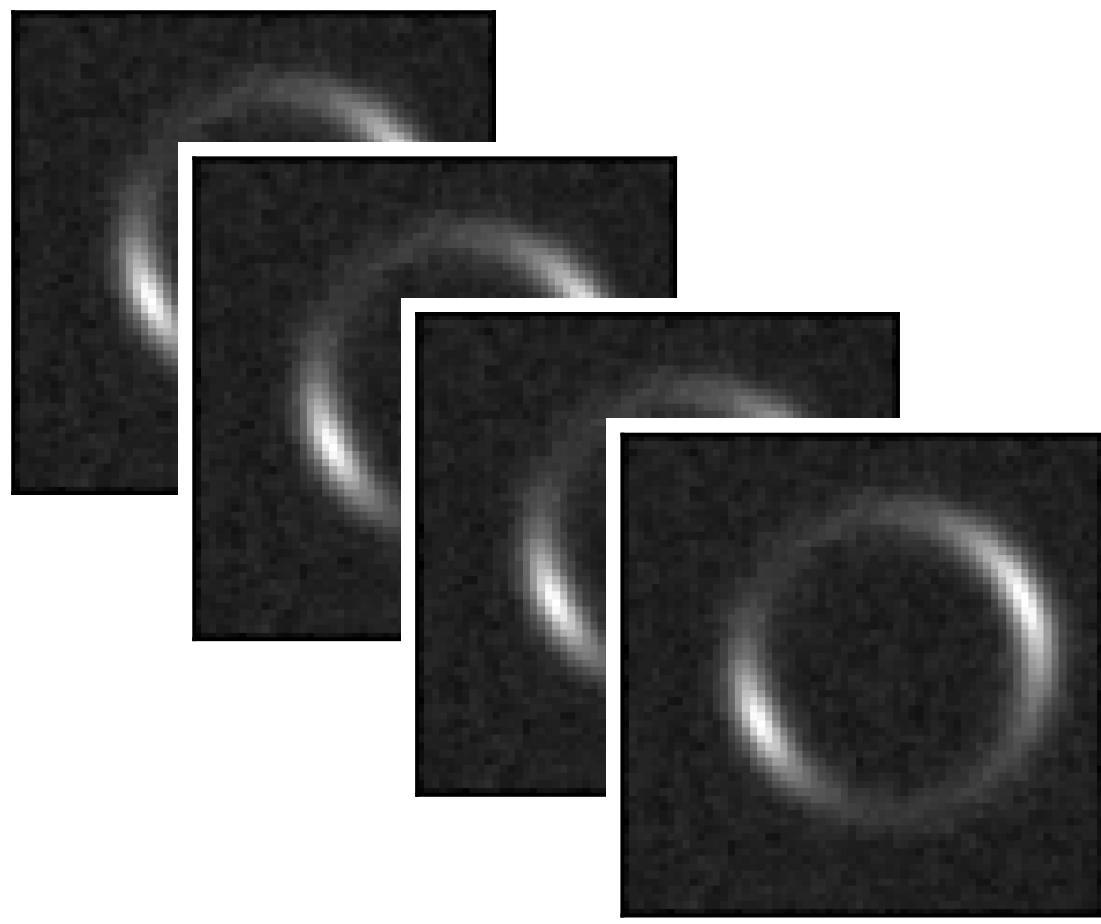
Smooth halo only



Smooth halo + subhalos



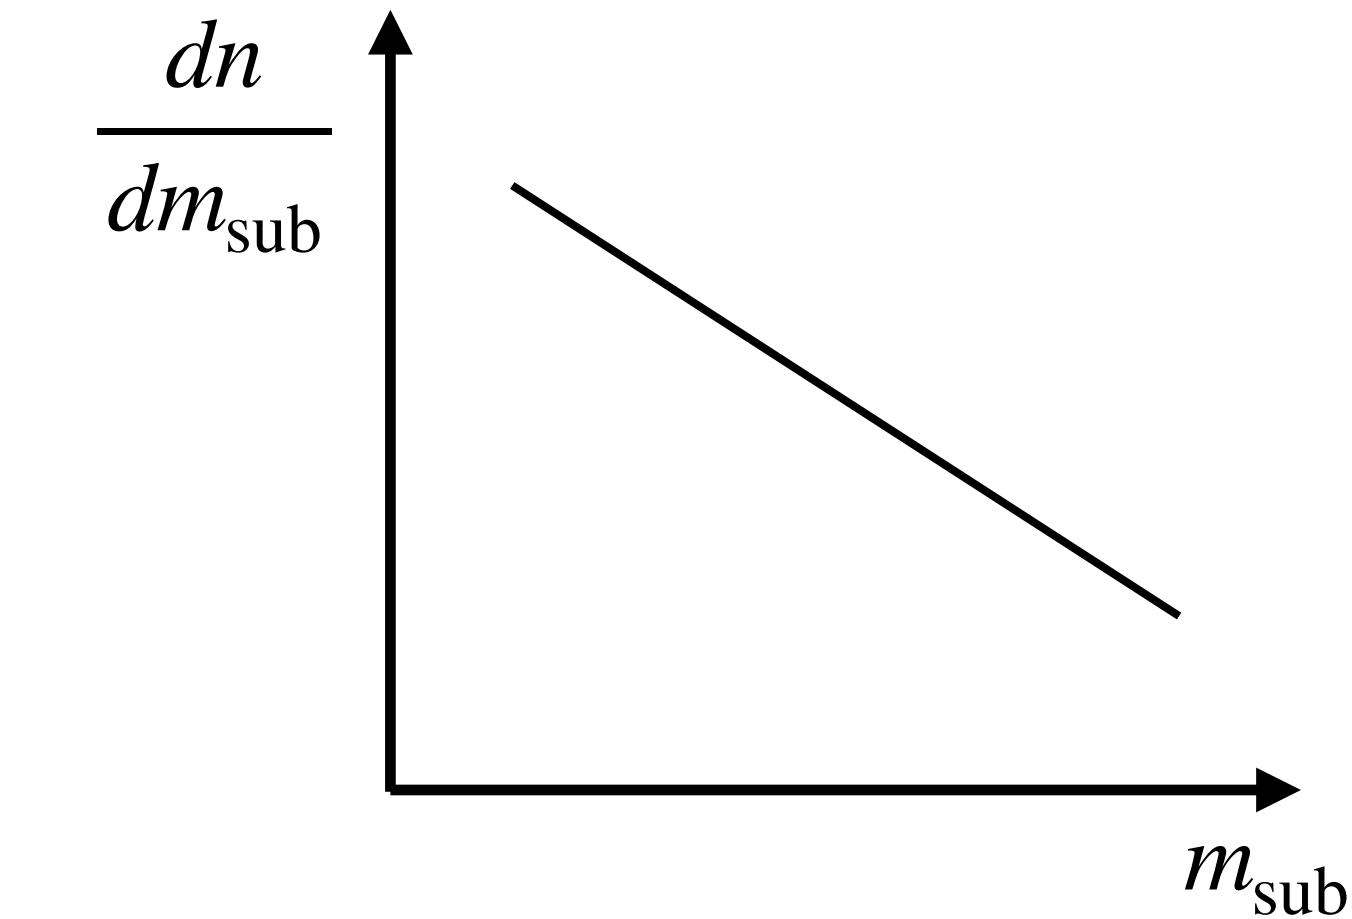
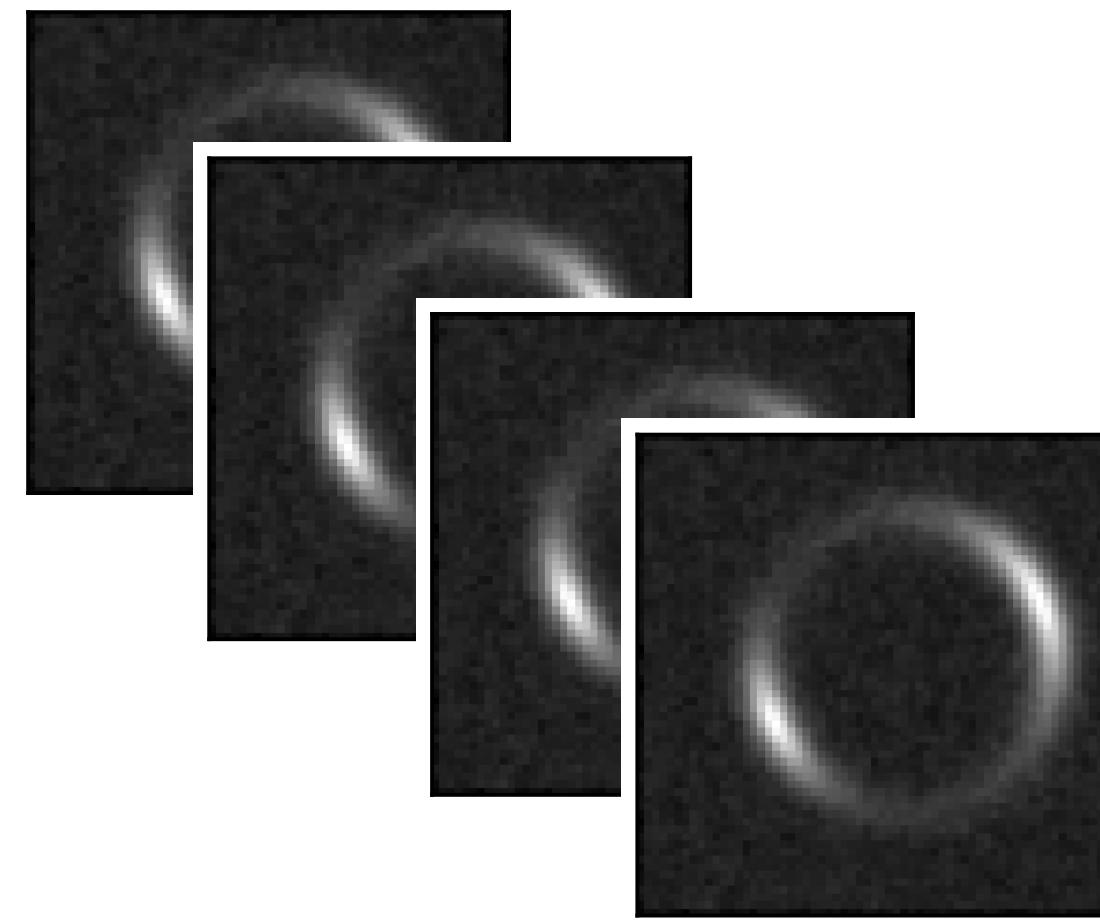
Goal: Scalable inference for small subhalos



Future surveys (LSST, Euclid) are expected to deliver large samples of galaxy-galaxy strong lenses [Collett et al 1507.02657]

Goal: infer subhalo mass distribution through collective effects of many light subhalos

Goal: Scalable inference for small subhalos



Future surveys (LSST, Euclid) are expected to deliver large samples of galaxy-galaxy strong lenses [Collett et al 1507.02657]

Goal: infer subhalo mass distribution through collective effects of many light subhalos

- ⇒ Need inference technique that
- scales to many lenses
 - can deal with a large number of subhalos
 - captures subtle effects in high-dimensional image data

A key challenge is that
we cannot calculate the likelihood function.

Modeling strong lensing

Model parameters

Subhalo mass
function parameters

$$\theta$$

Modeling strong lensing

Model parameters

Subhalo mass
function parameters

θ

Latent variables

Source / host halo
properties

z_{source} z_{lens}

Subhalo number,
masses, positions

n_{sub} $\{z_{\text{sub } i}\}$

$$p(z_{\text{source}}) \quad p(z_{\text{lens}}) \quad p(n_{\text{sub}}|\theta) \prod_i^{n_{\text{sub}}} p(z_{\text{sub } i}|\theta)$$

Modeling strong lensing

Model parameters

Subhalo mass
function parameters

$$\theta$$

Latent variables

Source / host halo
properties

$$z_{\text{source}}$$

Subhalo number,
masses, positions

$$n_{\text{sub}} \quad \{z_{\text{sub } i}\}$$

Observables

Observed lens image

$$x$$

$$p(z_{\text{source}}) \quad p(z_{\text{lens}}) \quad p(n_{\text{sub}}|\theta) \prod_i^{n_{\text{sub}}} p(z_{\text{sub } i}|\theta)$$

$$p(x|z_{\text{source}}, z_{\text{lens}}, \{z_{\text{sub } i}\})$$

Modeling strong lensing

Model parameters

Subhalo mass
function parameters

$$\theta$$

Latent variables

Source / host halo
properties

$$z_{\text{source}}$$

$$z_{\text{lens}}$$

Subhalo number,
masses, positions

$$n_{\text{sub}}$$

$$\{z_{\text{sub } i}\}$$

Observables

Observed lens image

$$x$$

We can easily write a simulator that samples from

$$p(x, z | \theta) = p(z_{\text{source}}) \quad p(z_{\text{lens}}) \quad p(n_{\text{sub}} | \theta) \prod_i^{n_{\text{sub}}} p(z_{\text{sub } i} | \theta) \quad p(x | z_{\text{source}}, z_{\text{lens}}, \{z_{\text{sub } i}\})$$



Prediction (simulation)

Modeling strong lensing

Model parameters

Subhalo mass
function parameters

$$\theta$$

Latent variables

Source / host halo
properties

$$z_{\text{source}}$$

$$z_{\text{lens}}$$

Subhalo number,
masses, positions

$$n_{\text{sub}}$$

$$\{z_{\text{sub } i}\}$$

Observables

Observed lens image

$$x$$

The key quantity for inference is the marginal likelihood

$$p(x|\theta) = \int dz_{\text{source}} \int dz_{\text{lens}} \sum_{n_{\text{sub}}} \int d^{n_{\text{sub}}} z_{\text{sub}} \quad p(z_{\text{source}}) \quad p(z_{\text{lens}}) \quad p(n_{\text{sub}}|\theta) \prod_i^{n_{\text{sub}}} p(z_{\text{sub } i}|\theta) \quad p(x|z_{\text{source}}, z_{\text{lens}}, \{z_{\text{sub } i}\})$$



Inference

Modeling strong lensing

Model parameters

Subhalo mass
function parameters

$$\theta$$

Latent variables

Source / host halo
properties

$$z_{\text{source}}$$

$$z_{\text{lens}}$$

Subhalo number,
masses, positions

$$n_{\text{sub}}$$

$$\{z_{\text{sub } i}\}$$

Observables

Observed lens image

$$x$$

The key quantity for inference is the marginal likelihood

$$p(x|\theta) = \boxed{\int dz_{\text{source}} \int dz_{\text{lens}} \sum_{n_{\text{sub}}} \int d^{n_{\text{sub}}} z_{\text{sub}} p(z_{\text{source}}) p(z_{\text{lens}}) p(n_{\text{sub}}|\theta) \prod_i^n p(z_{\text{sub } i}|\theta)} p(x|z_{\text{source}}, z_{\text{lens}}, \{z_{\text{sub } i}\})$$

Unfortunately, it is intractable:

we cannot compute the integral over this huge latent space. This is a major challenge for inference.

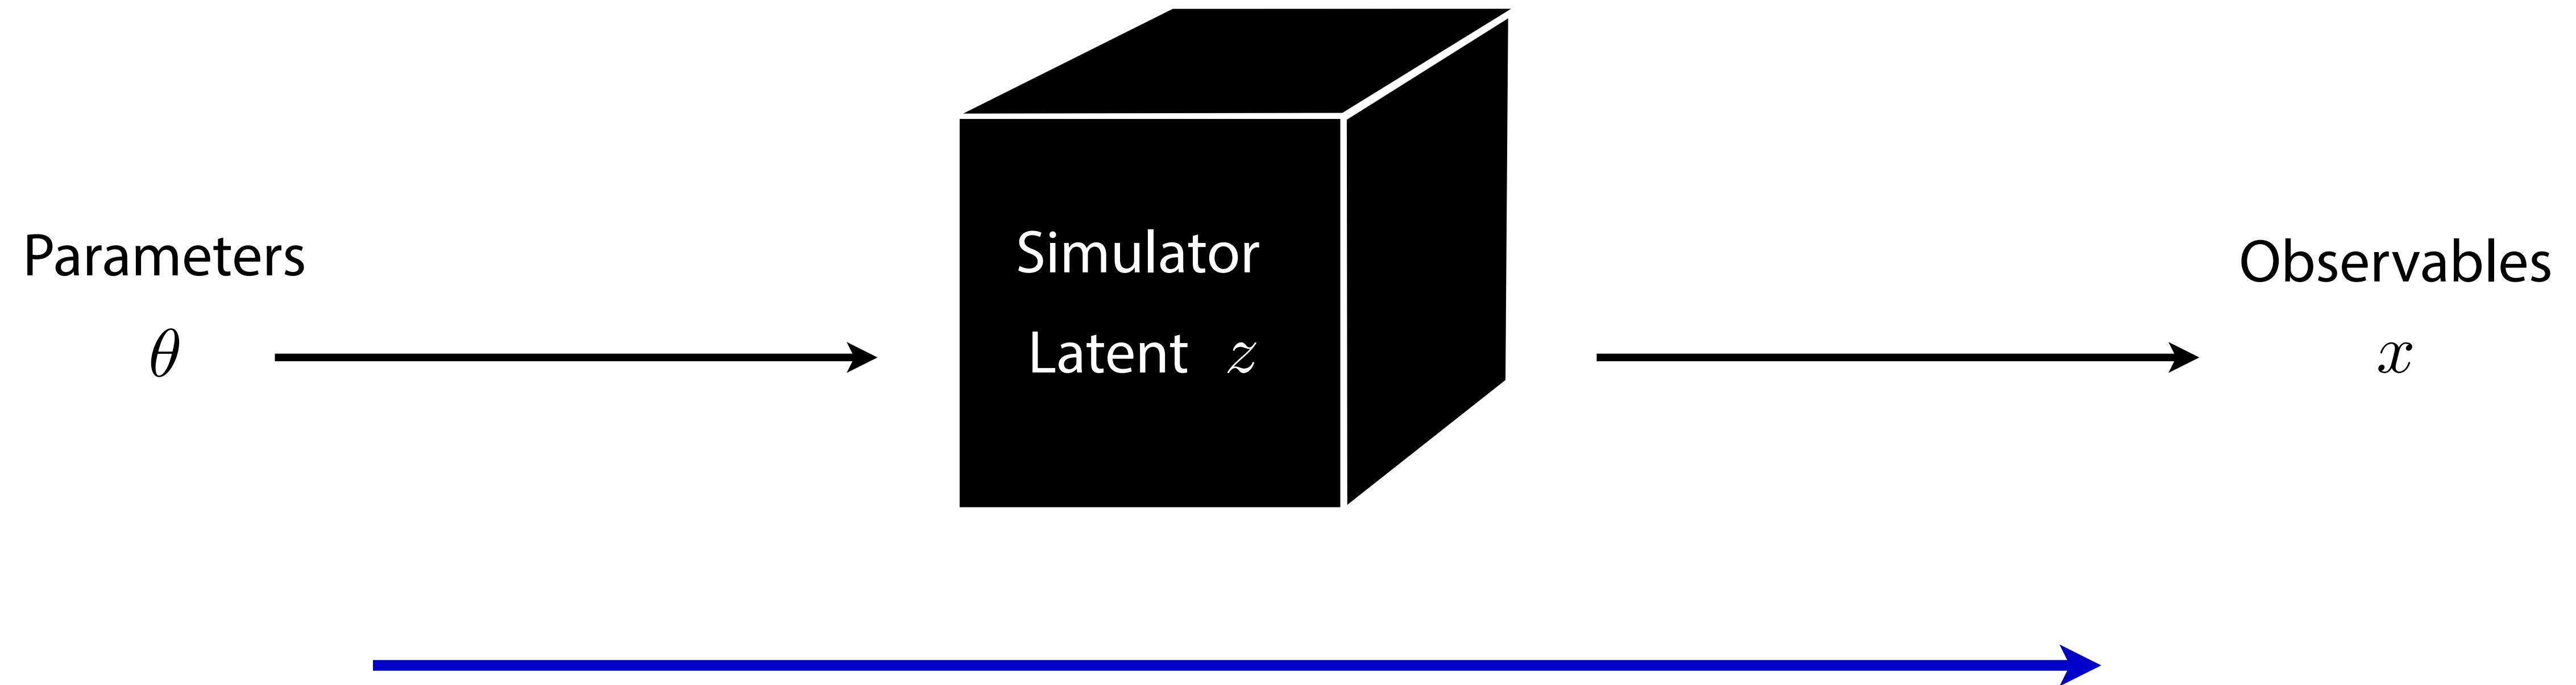
Similarly, MCMC based on $p(x, z|\theta)$ is extremely inefficient.

(More subtle: We cannot sample from $p(z|x, \theta)$ efficiently.)



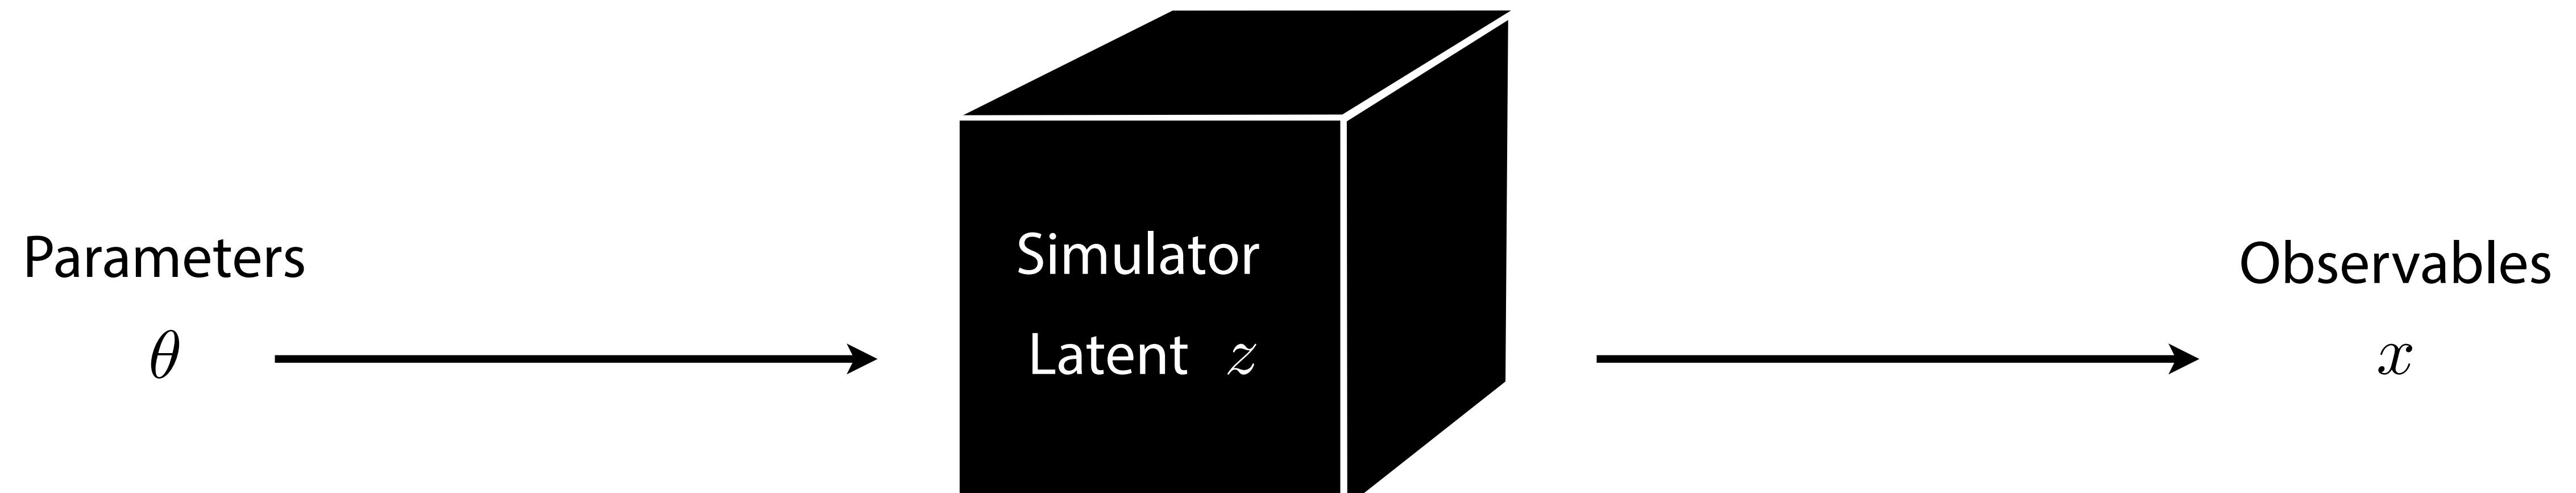
Inference

Simulation-based (“likelihood-free”) inference



- Prediction:
- Well-understood mechanistic model
 - Simulator can generate samples $x \sim p(x|\theta)$

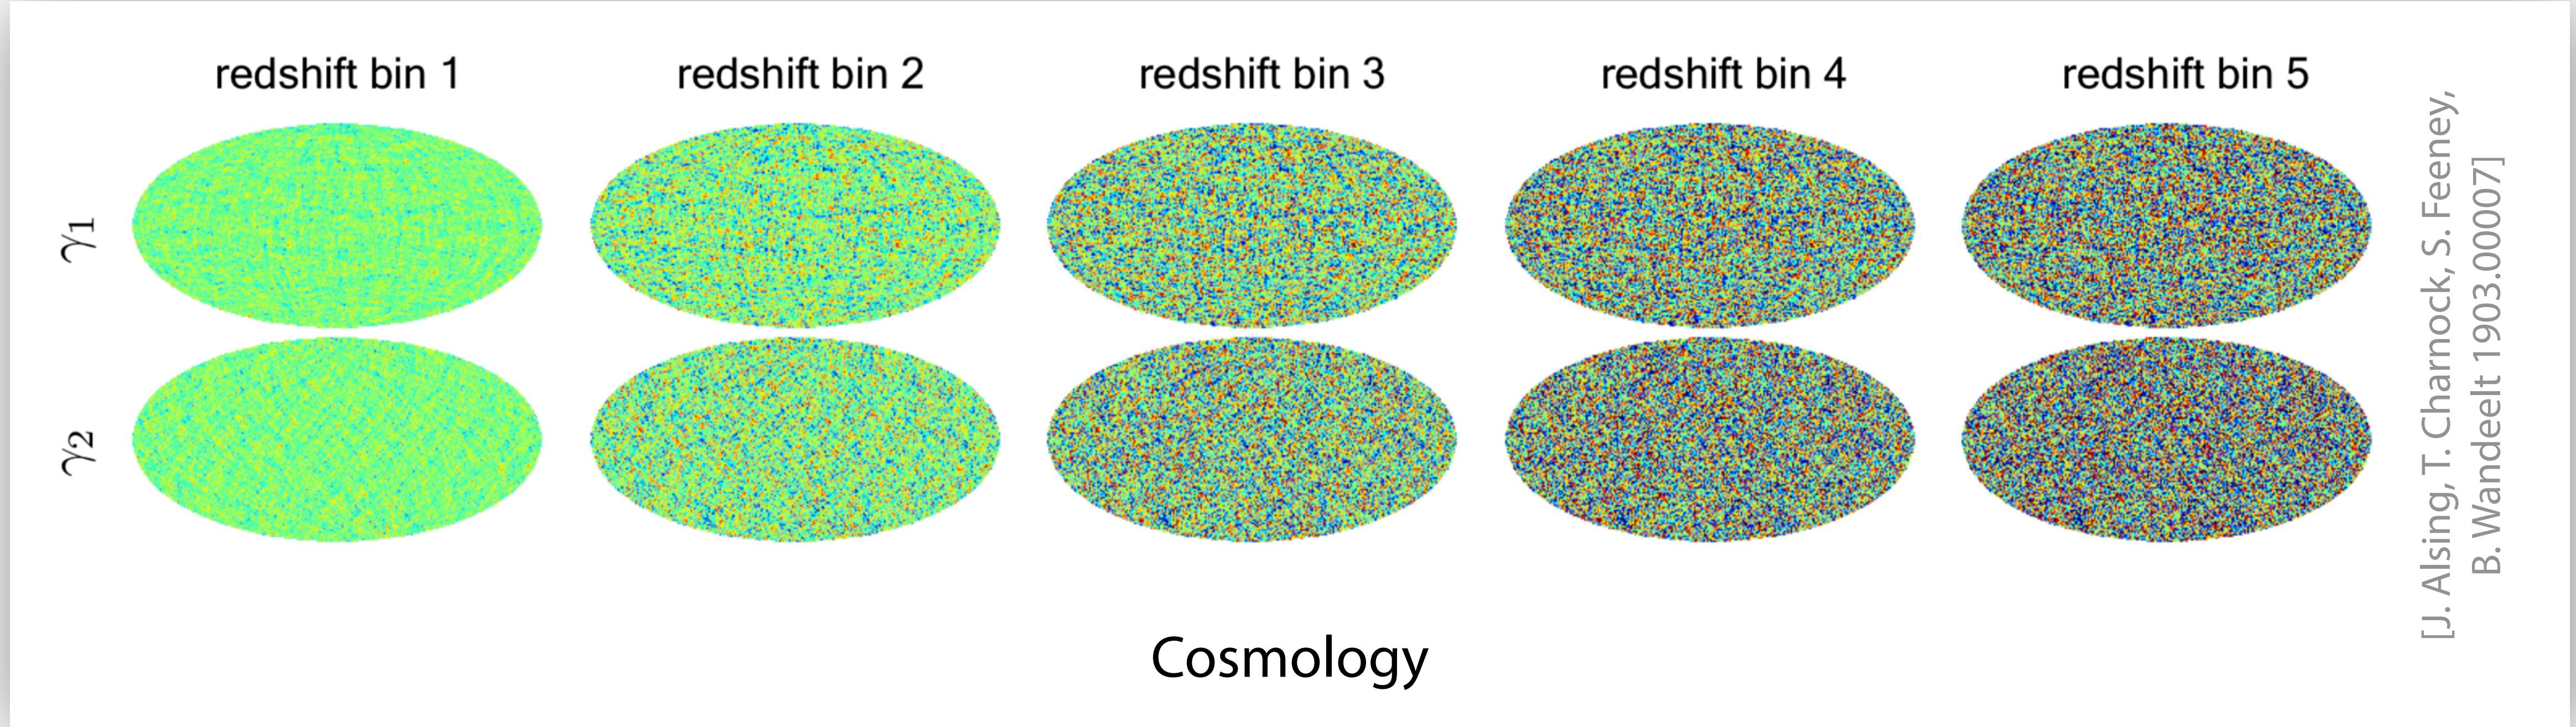
Simulation-based (“likelihood-free”) inference



- Prediction:
- Well-understood mechanistic model
 - Simulator can generate samples $x \sim p(x|\theta)$

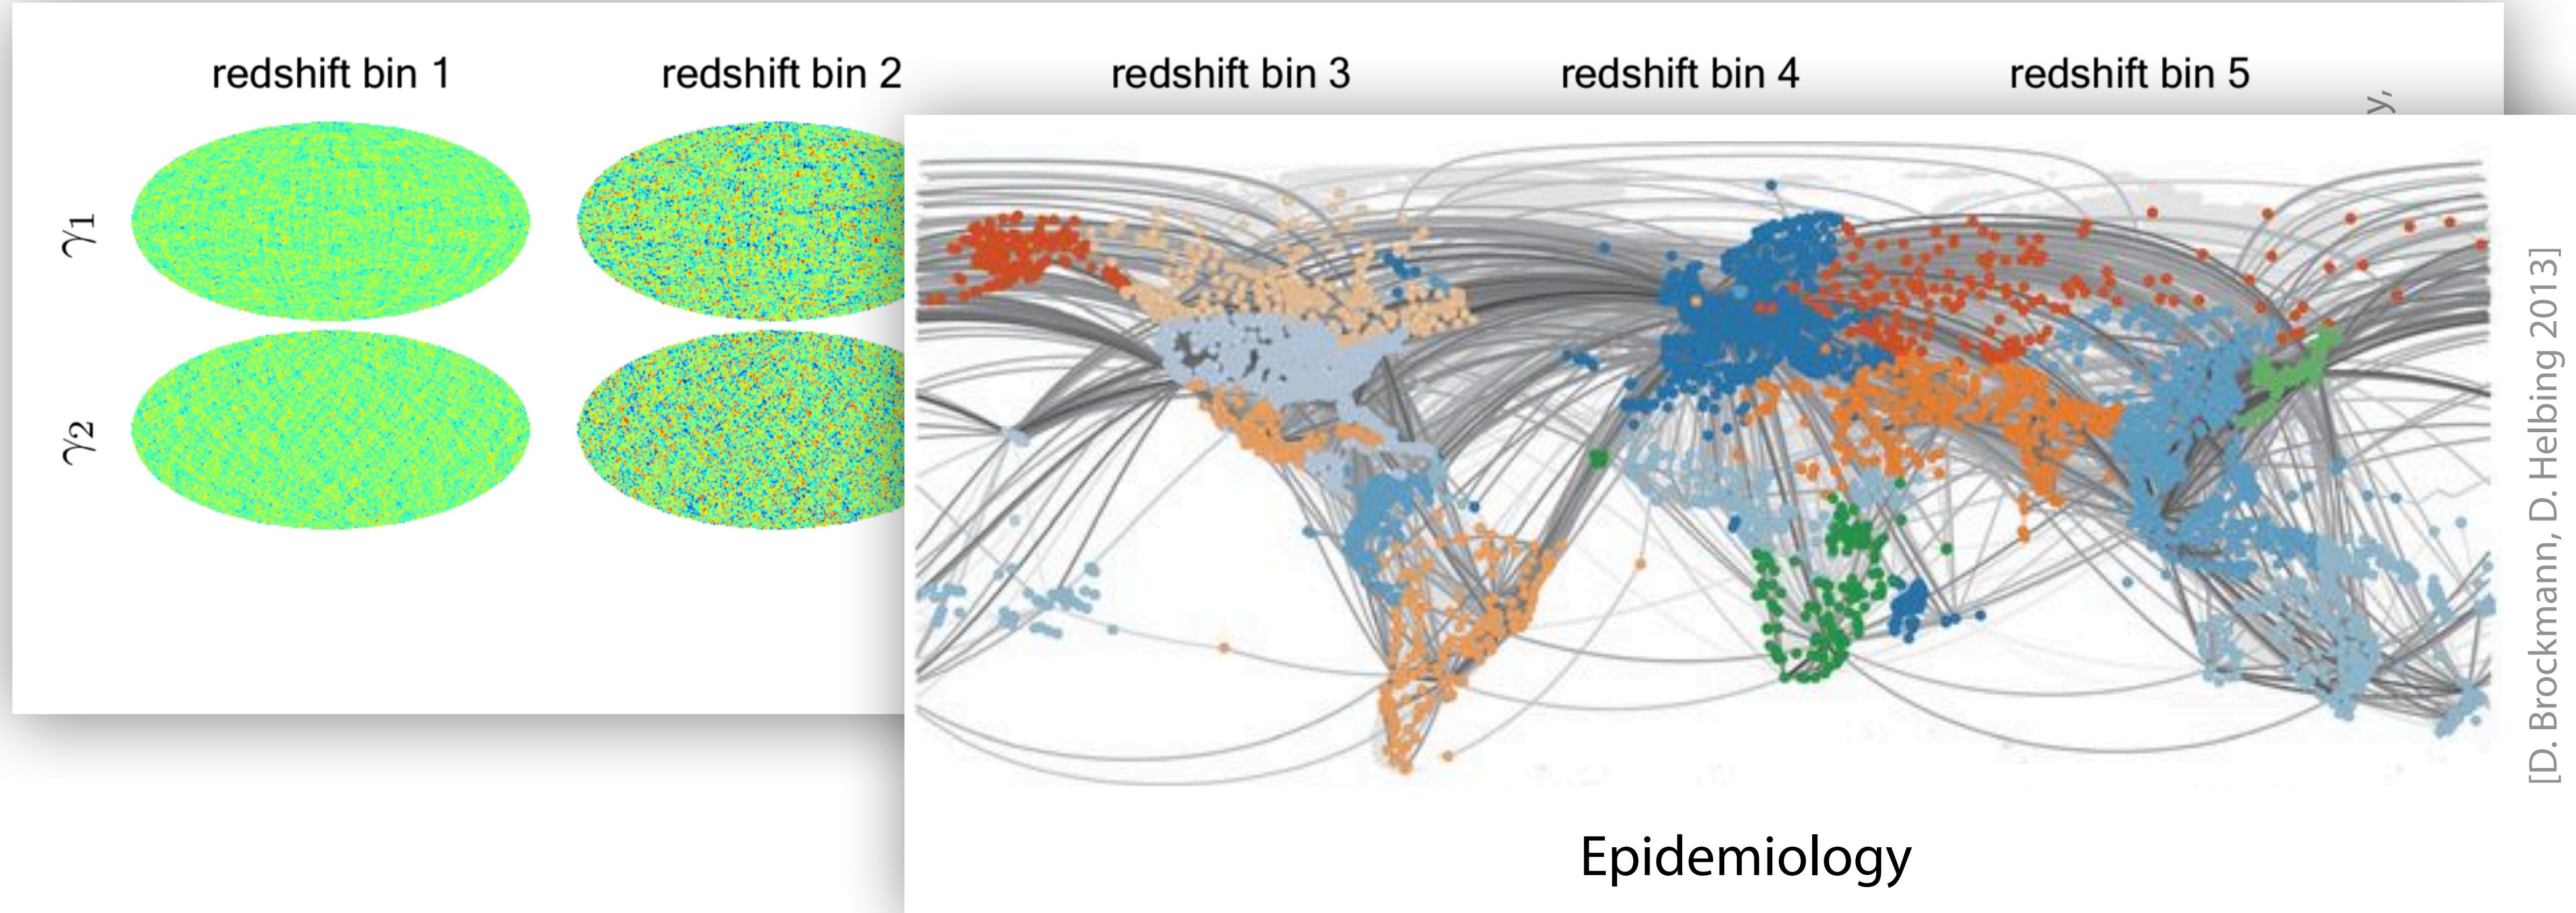
- Inference:
- Likelihood $p(x|\theta) = \int dz p(x, z|\theta)$ is intractable
 - Inference is challenging

Simulation-based inference appears all over science

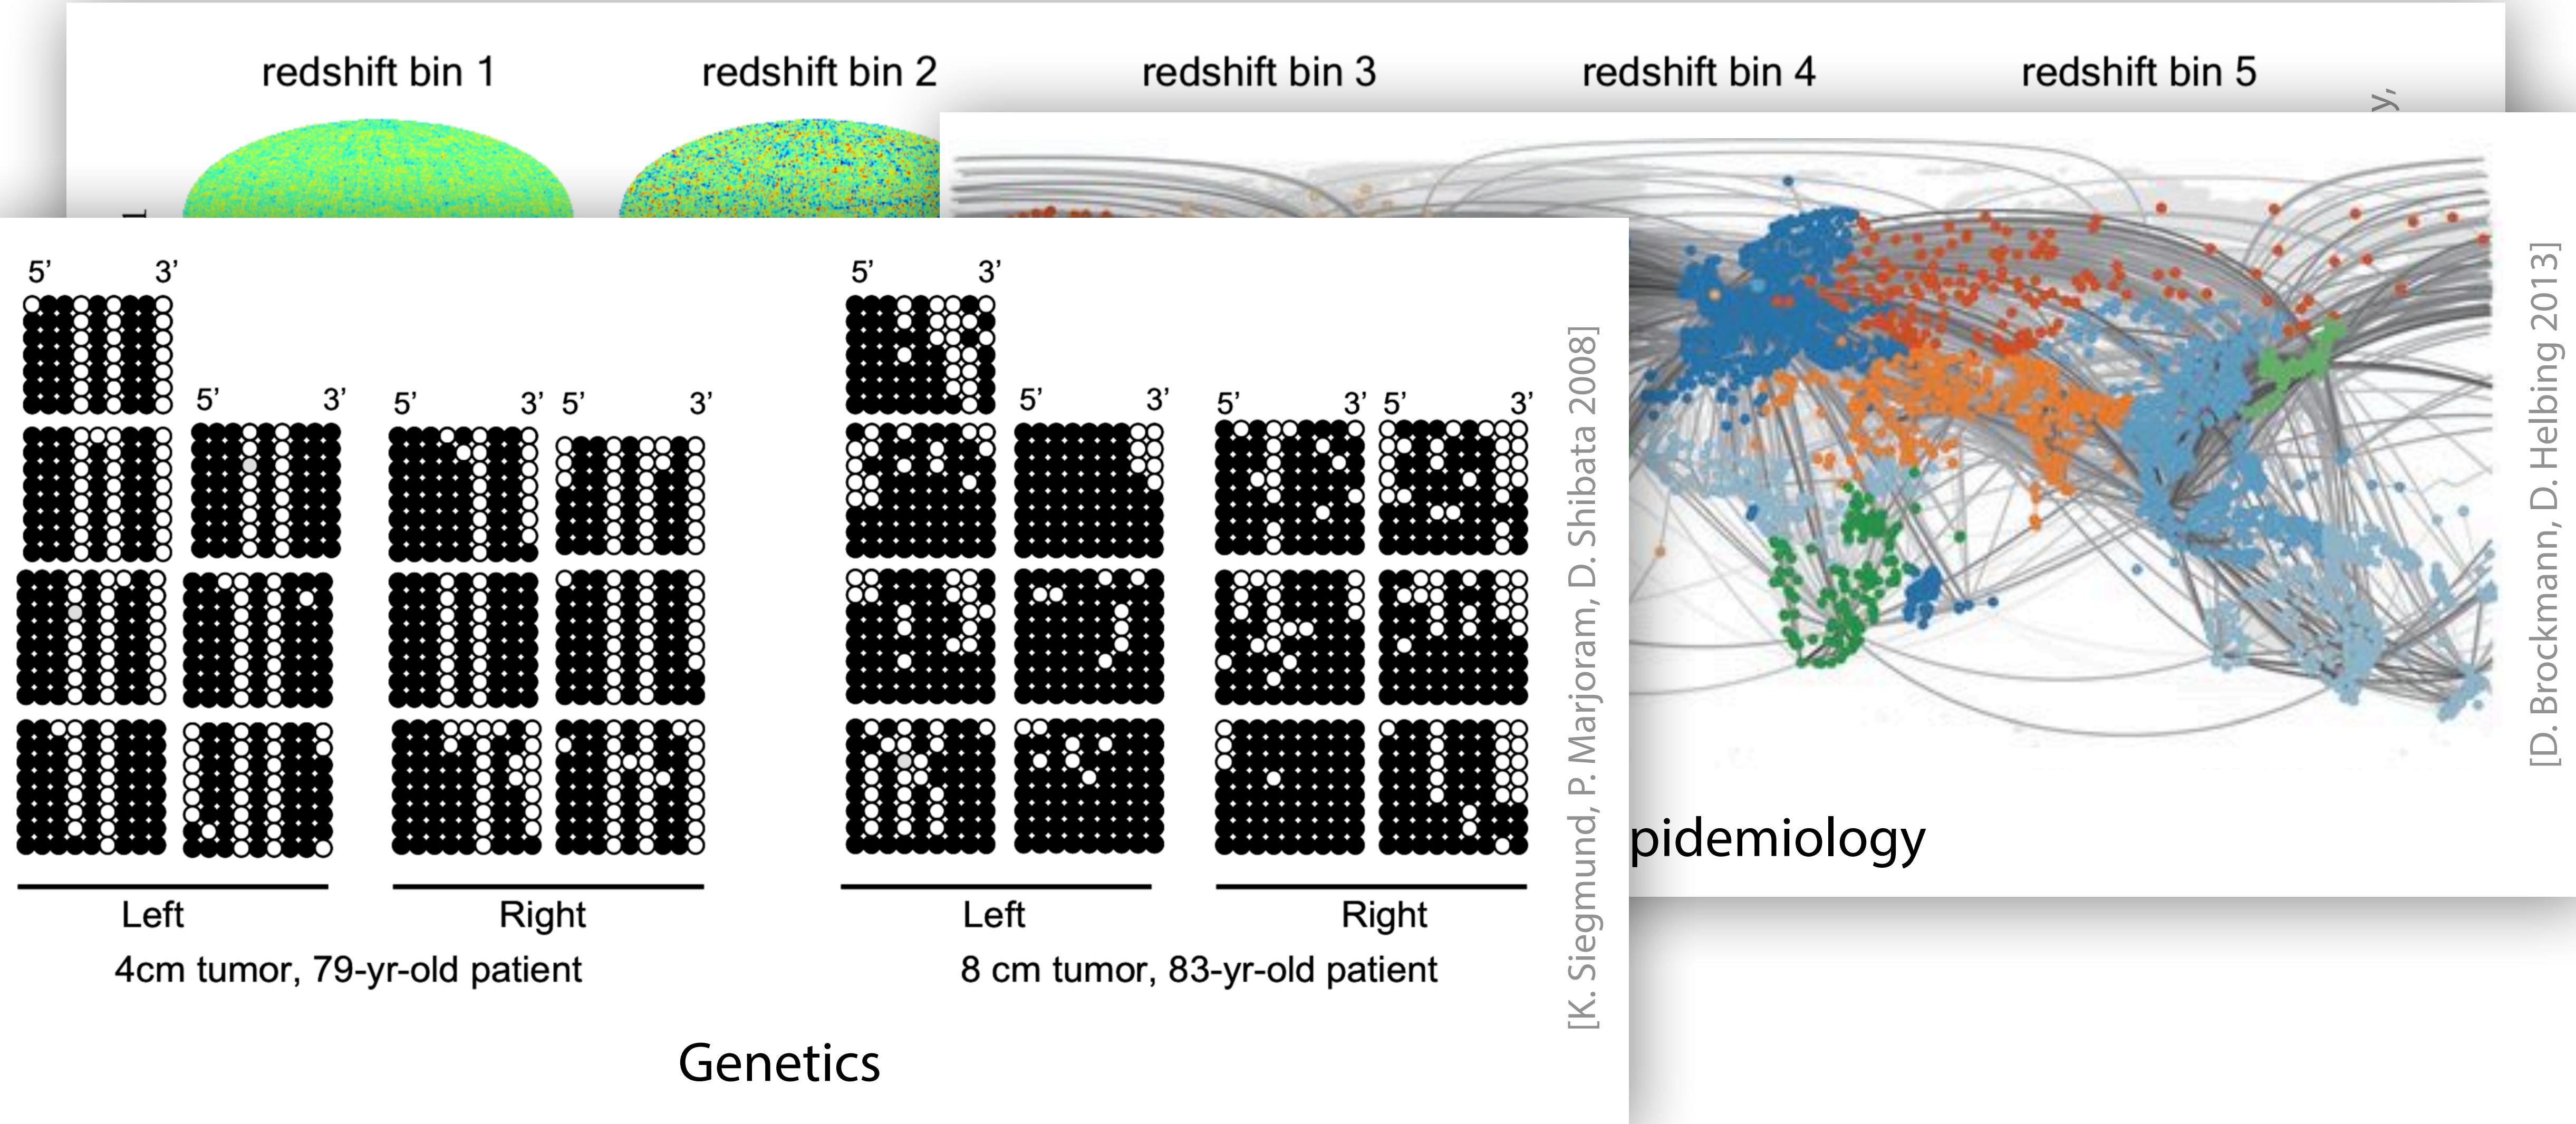


[J. Alsing, T. Charnock, S. Feeney,
B. Wandelt 1903.00007]

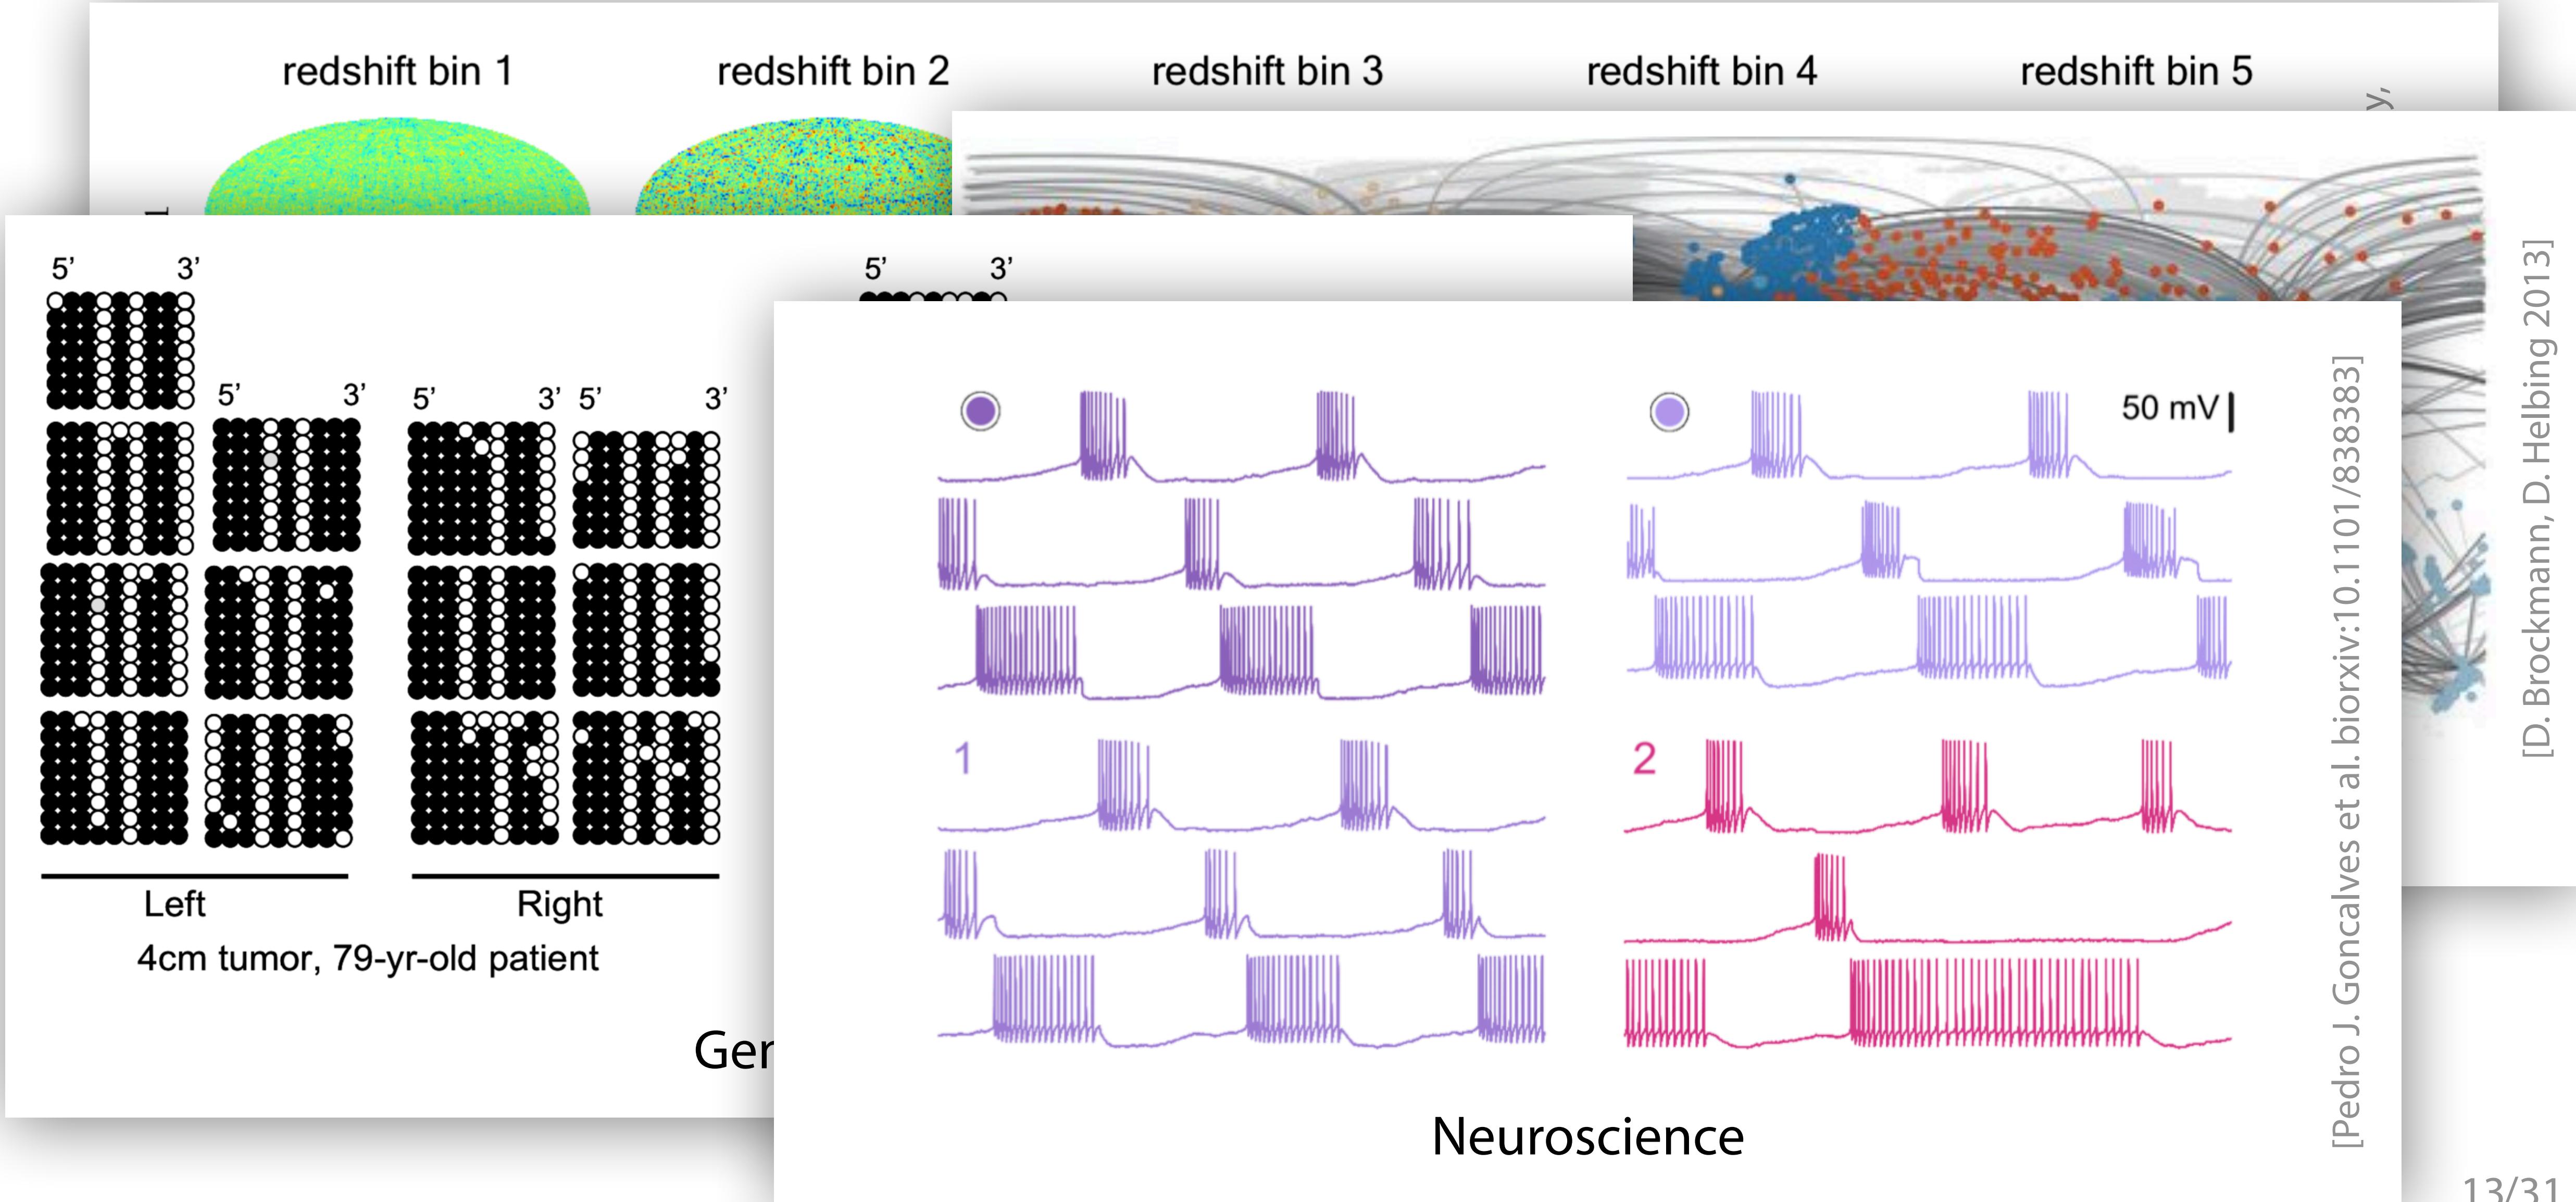
Simulation-based inference appears all over science



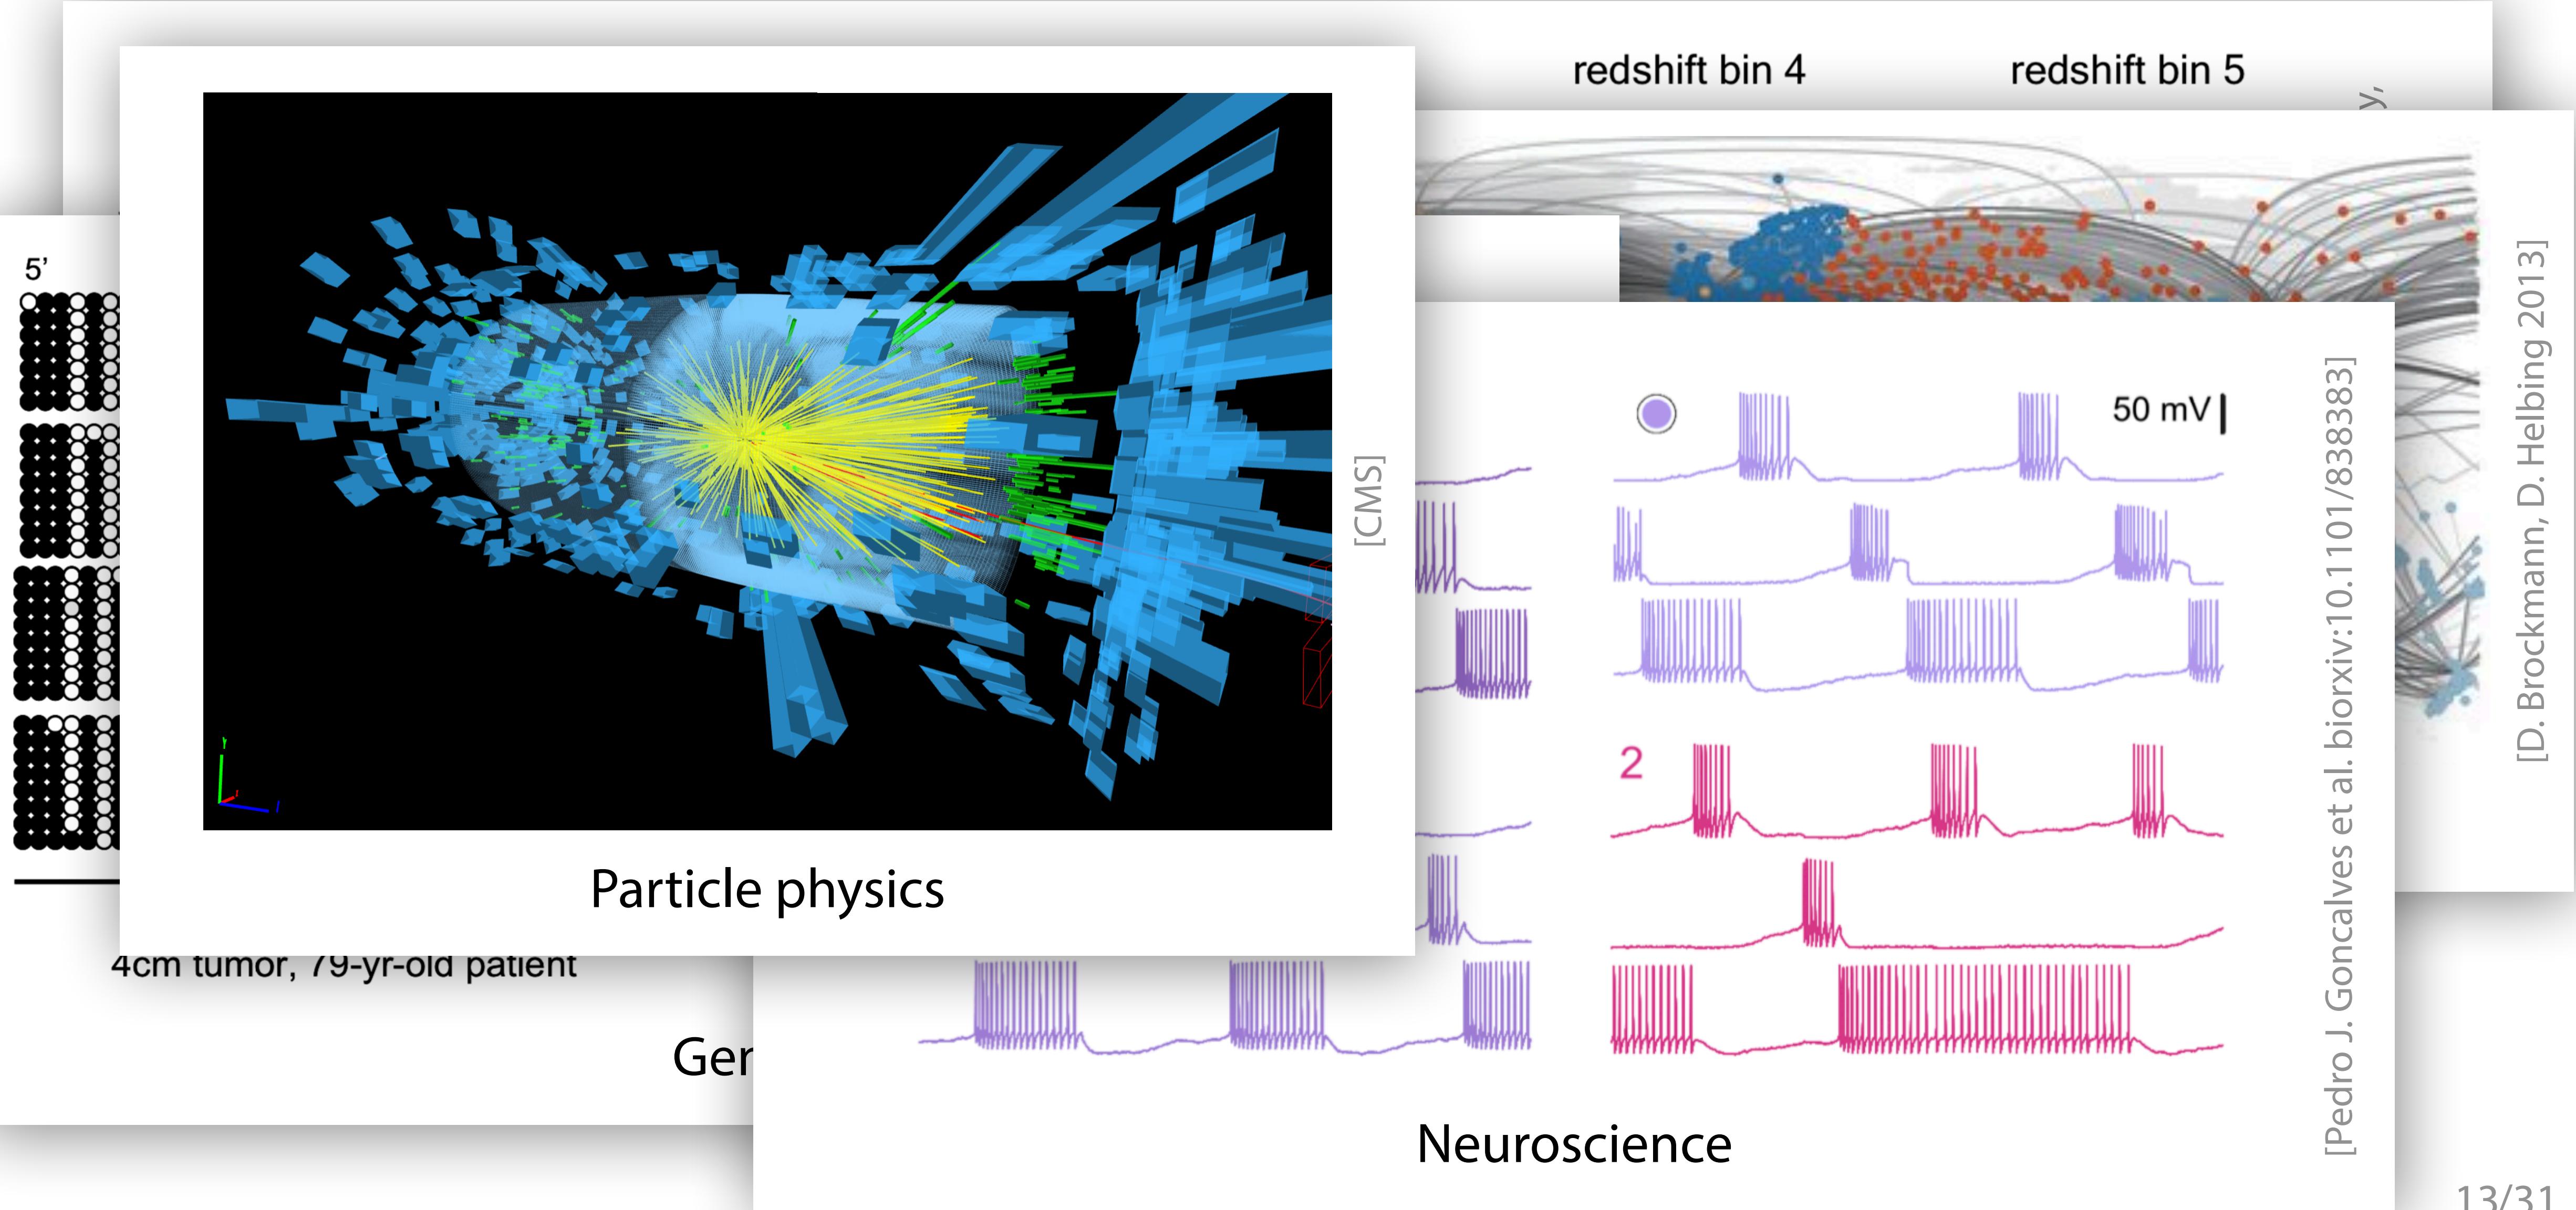
Simulation-based inference appears all over science



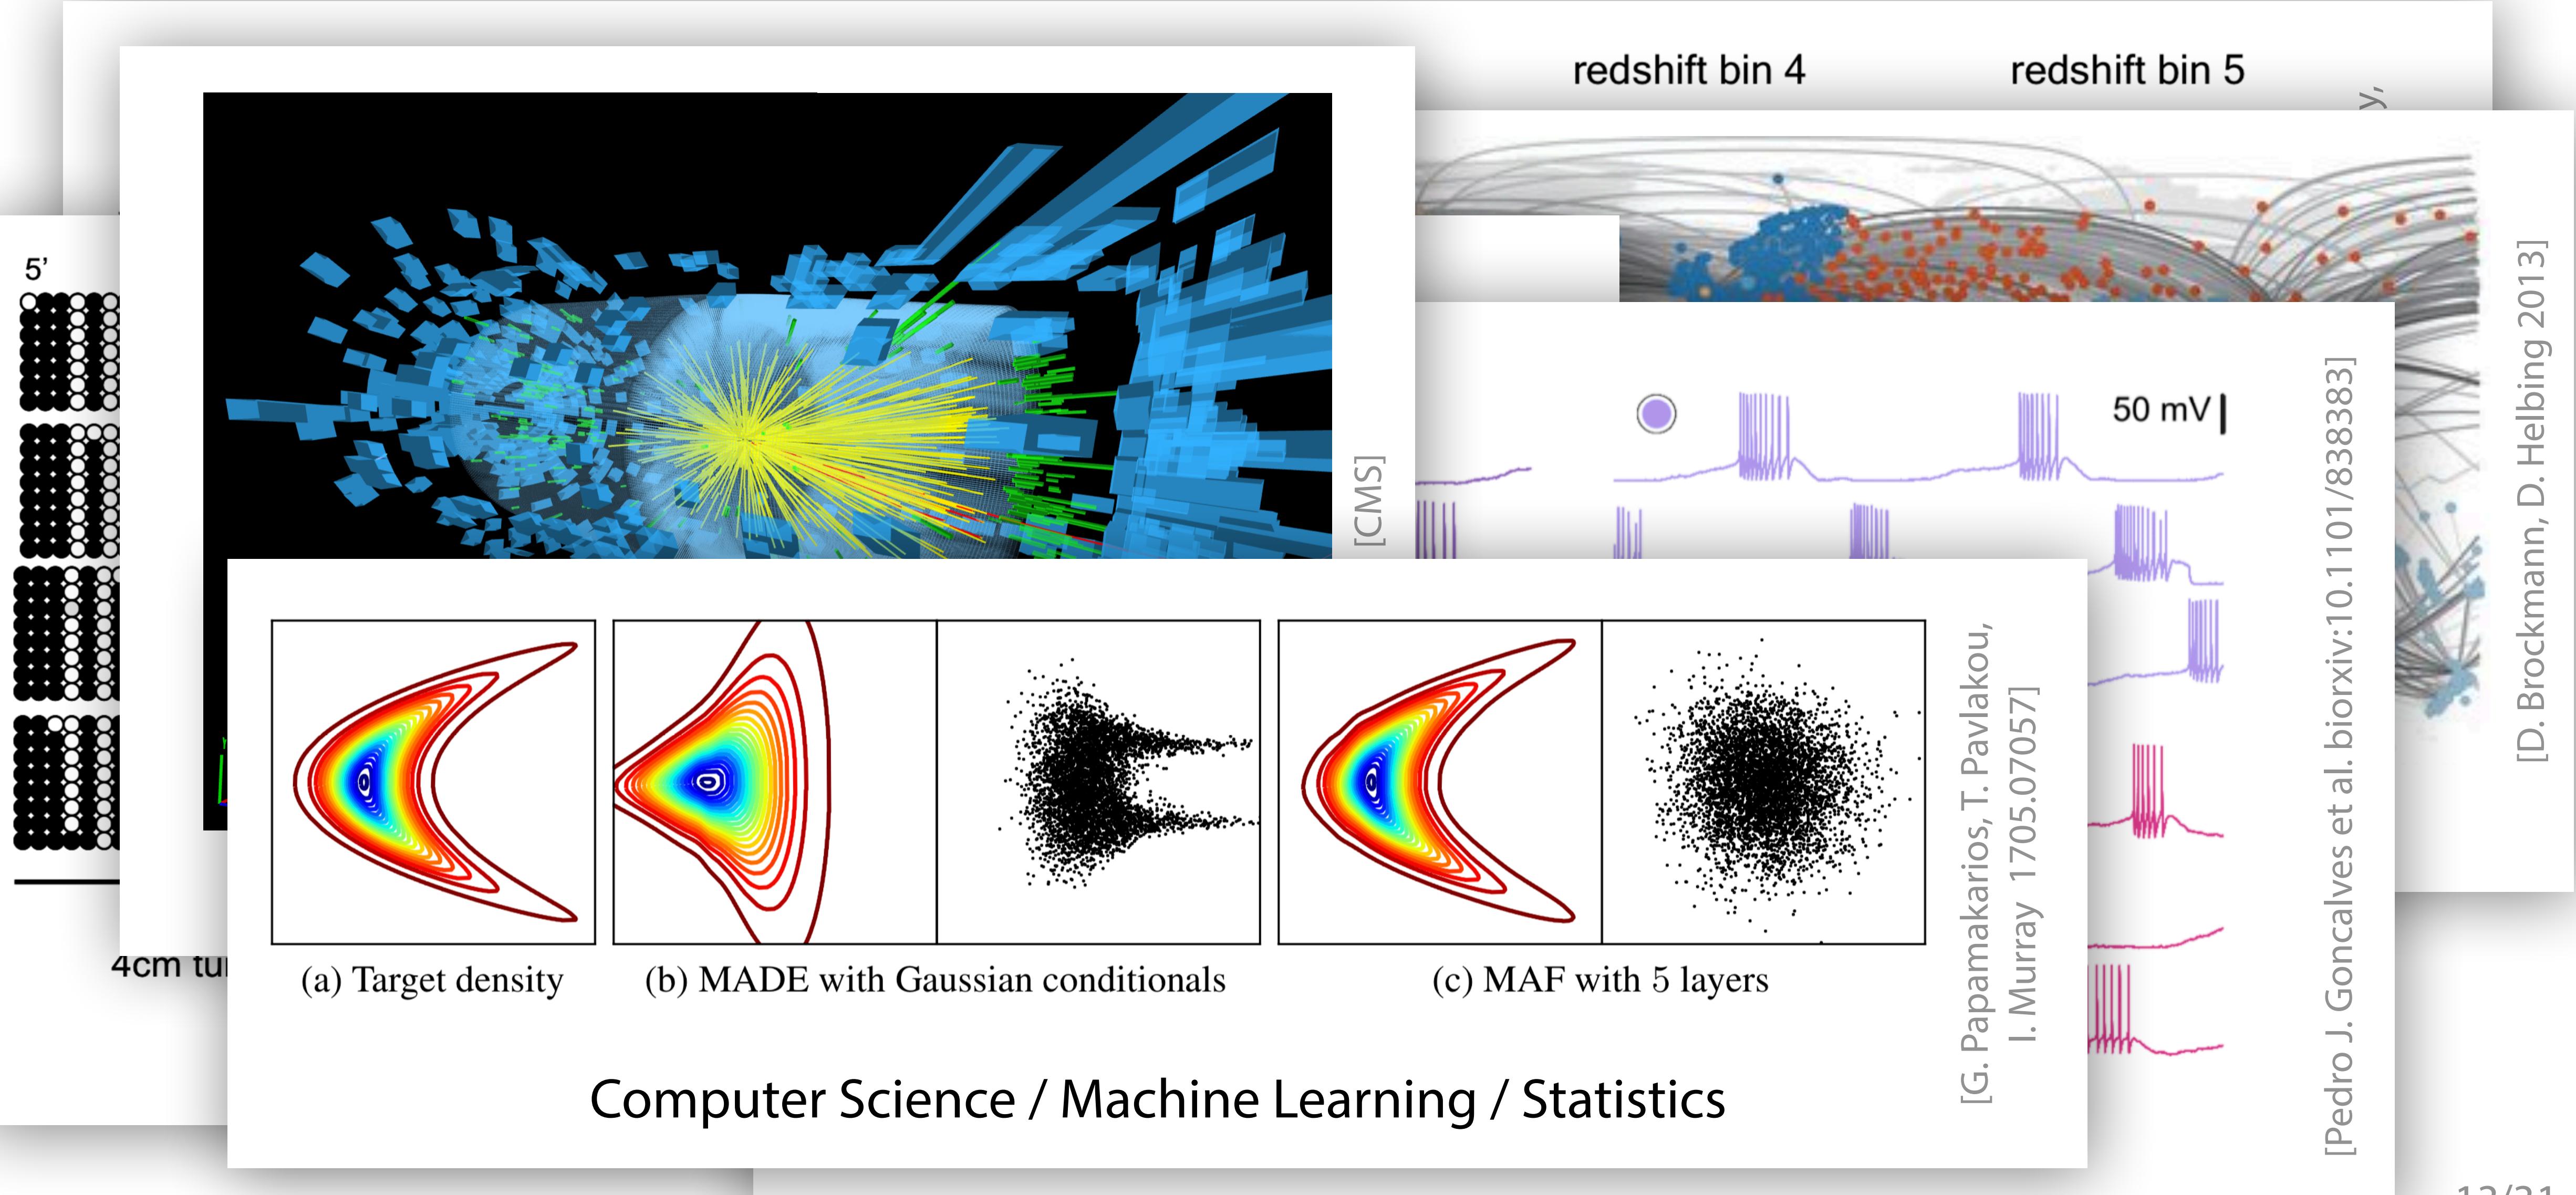
Simulation-based inference appears all over science



Simulation-based inference appears all over science

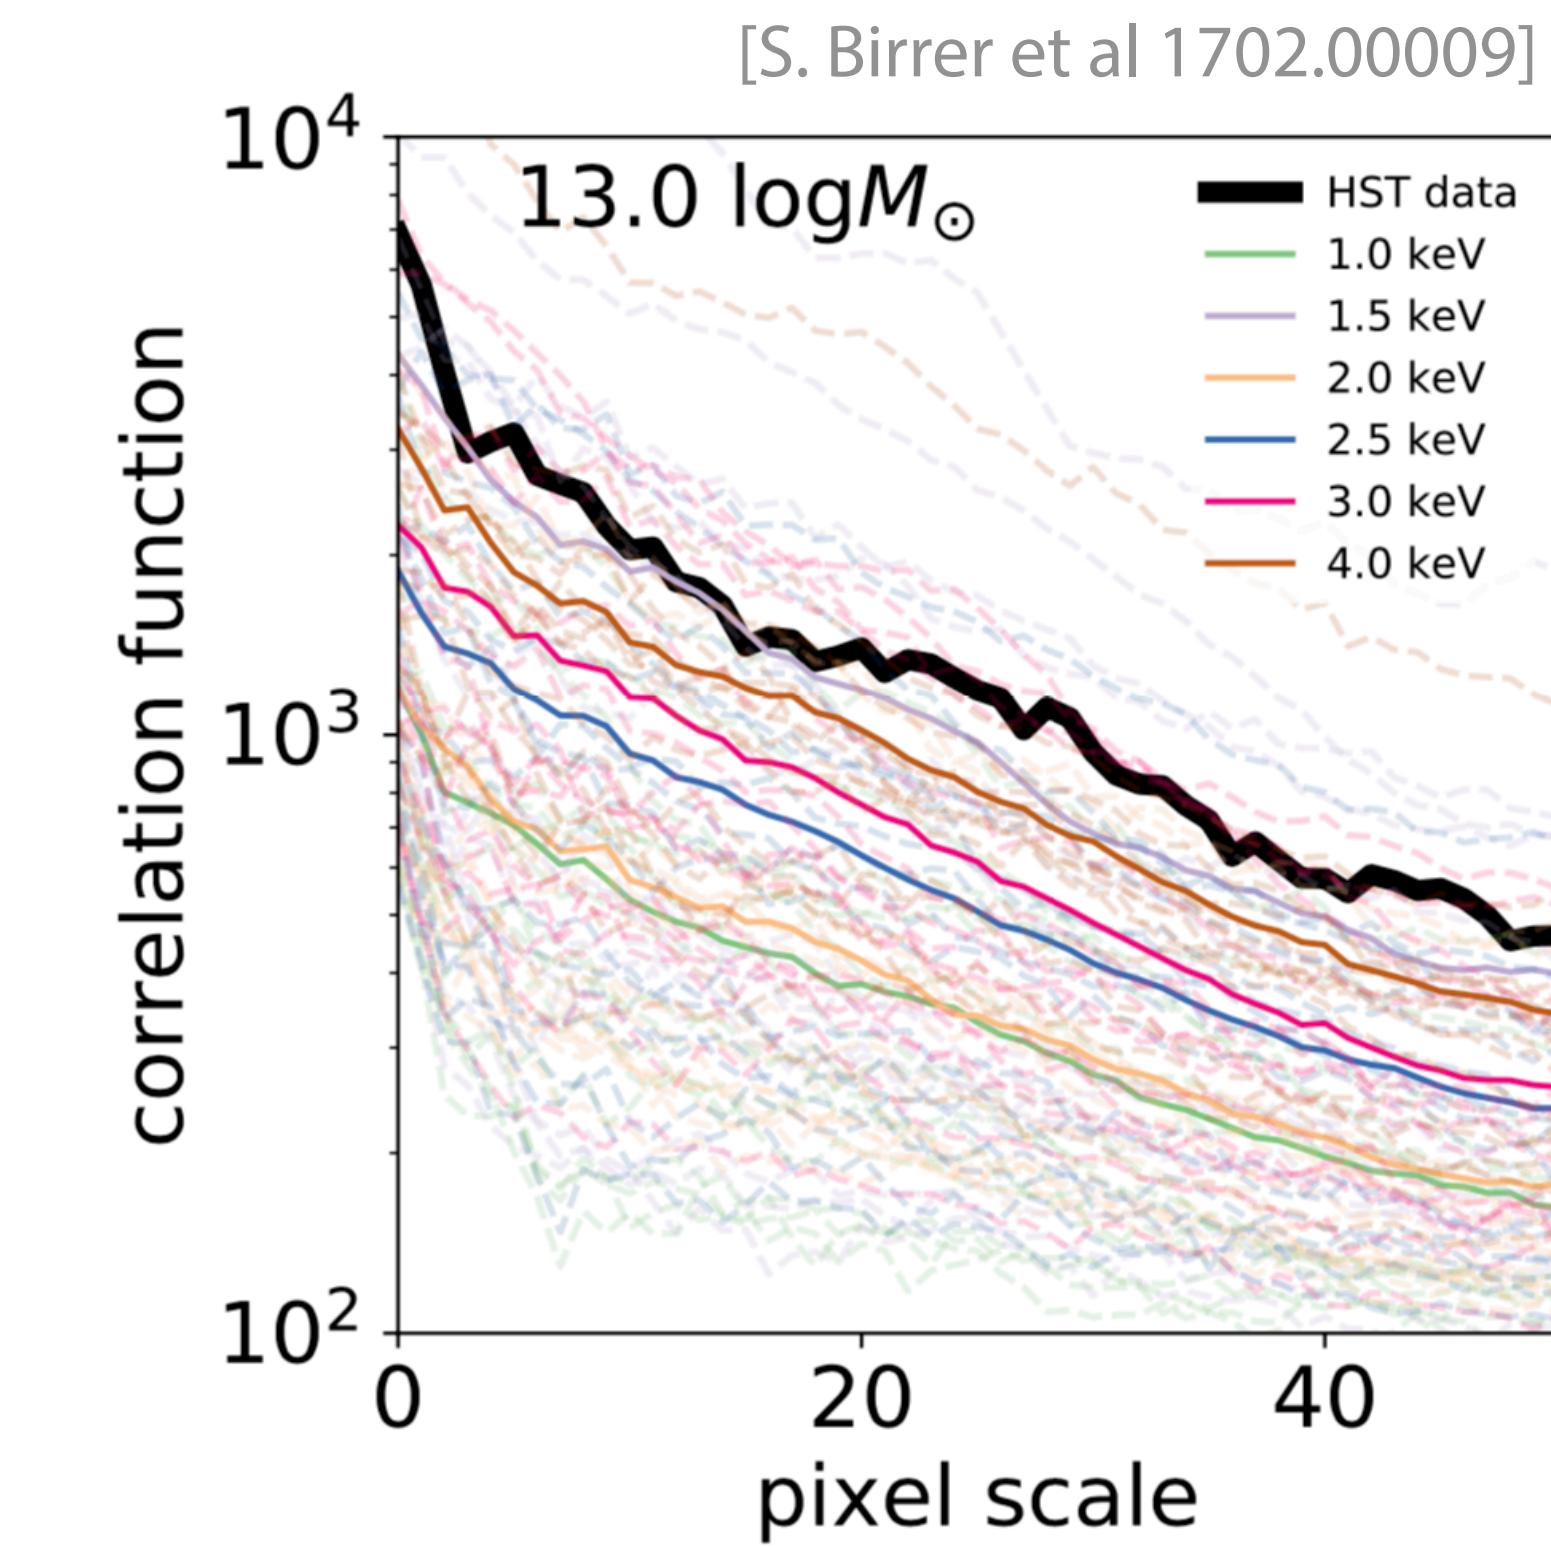
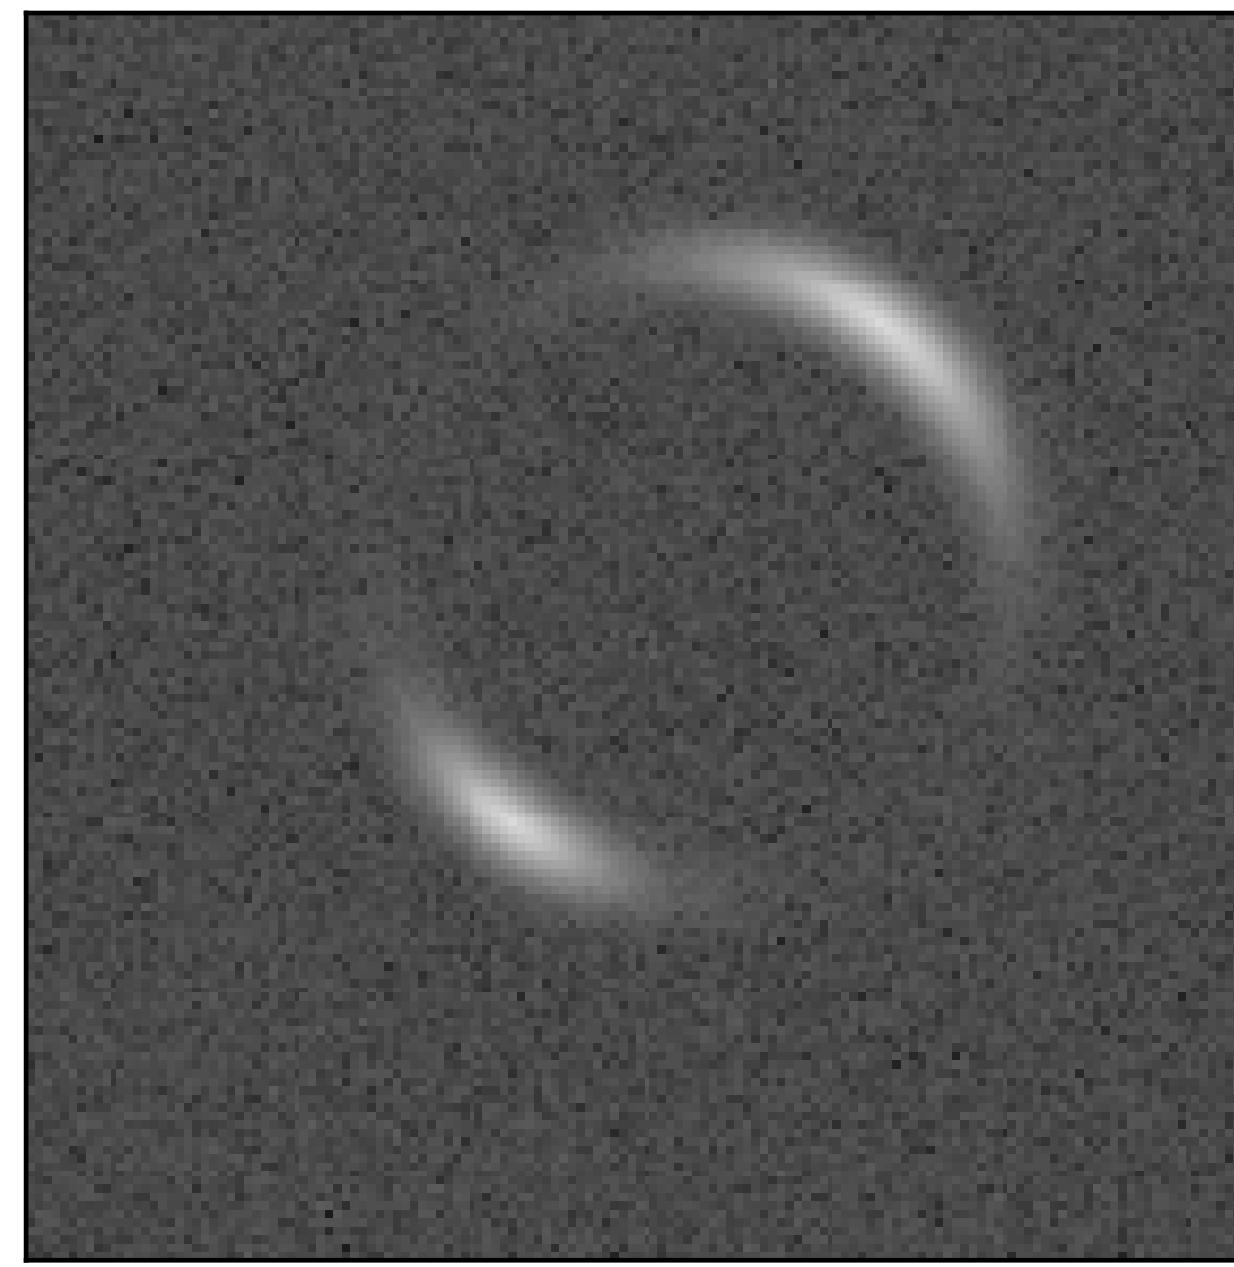


Simulation-based inference appears all over science



Traditionally, inference is made possible
with summary statistics.

Solve it with summary statistics

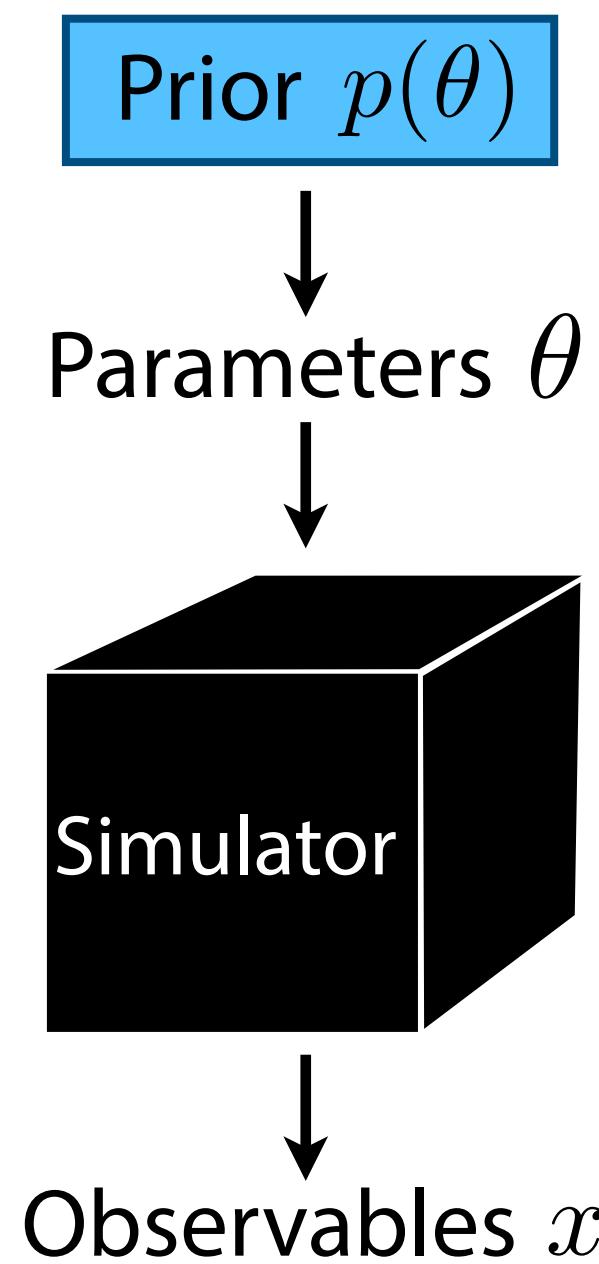


High-dimensional observable data x
(pixel values)

Low-dimensional summary statistics x'
(e.g. correlation functions)

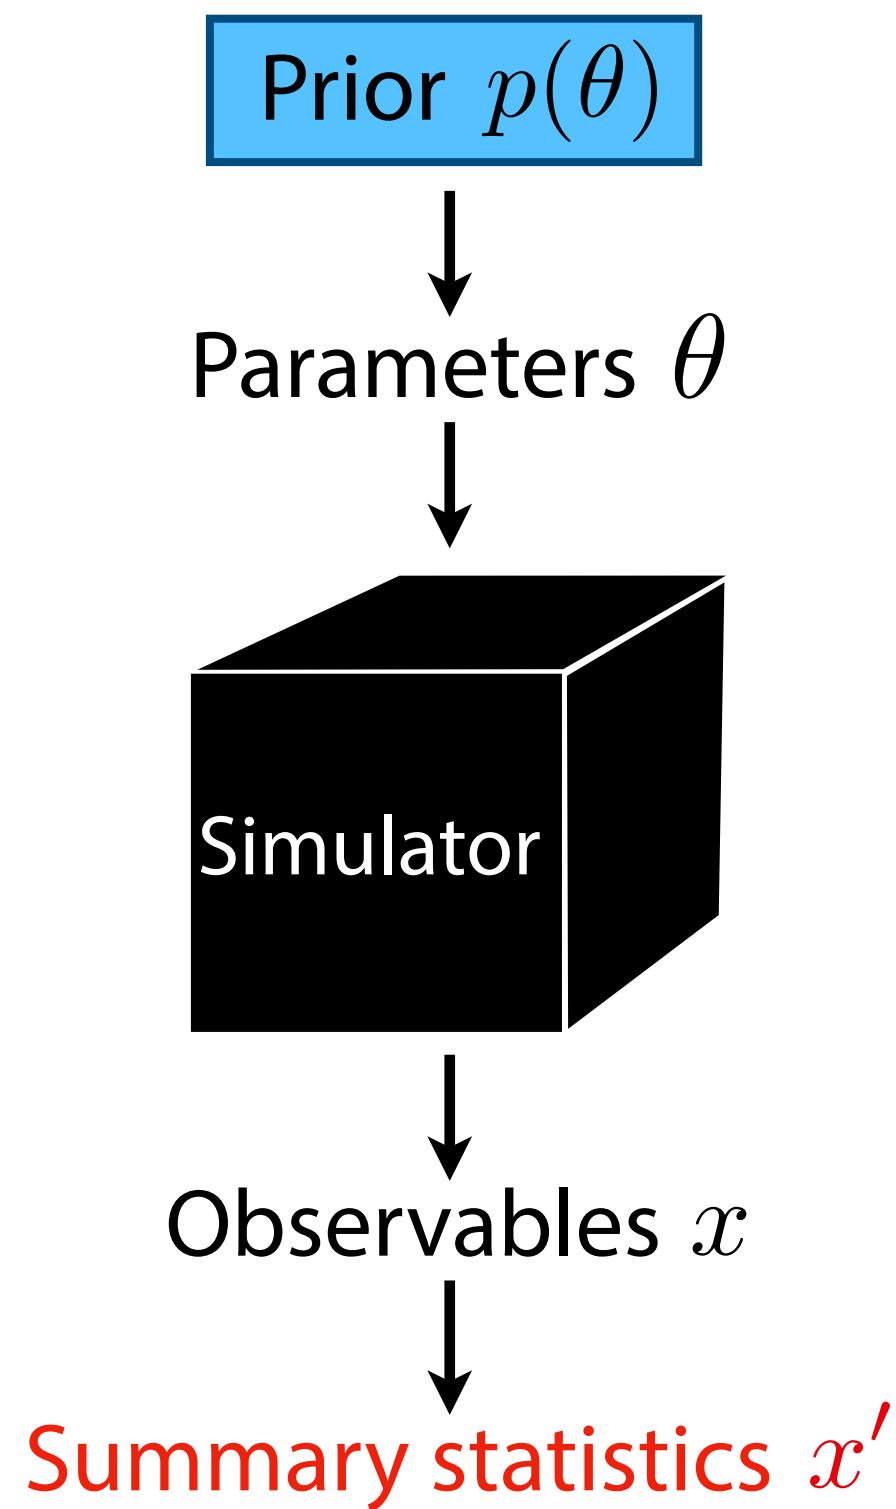
Approximate Bayesian Computation (ABC)

[D. Rubin 1984]



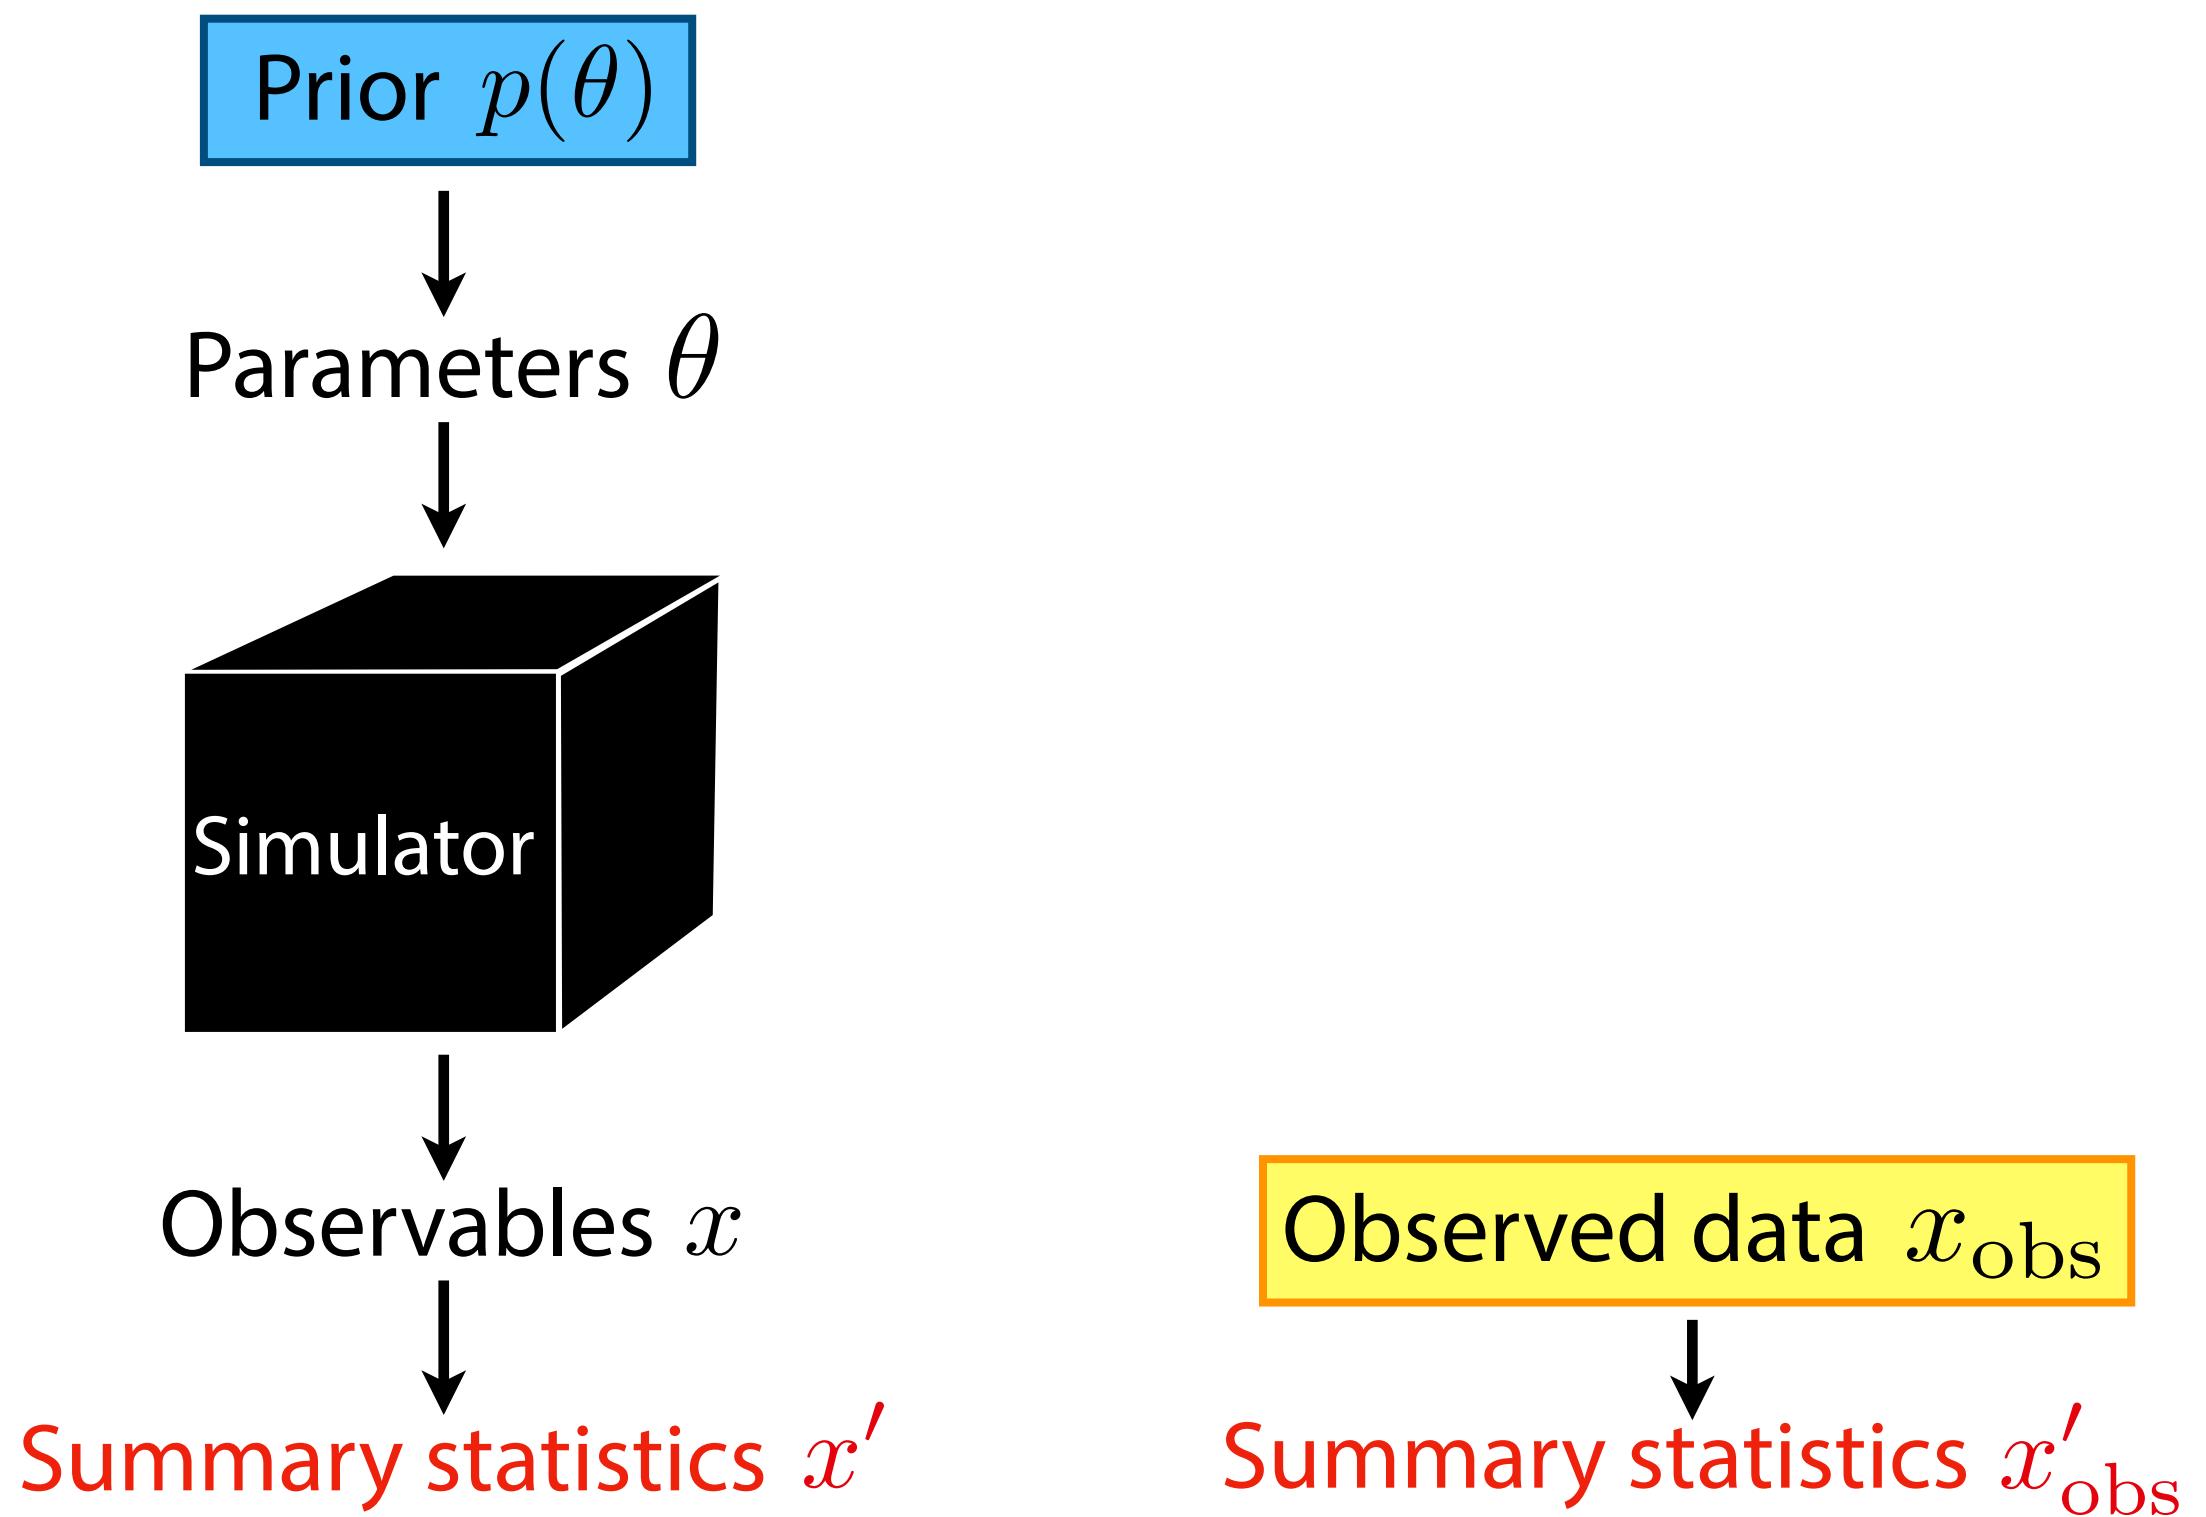
Approximate Bayesian Computation (ABC)

[D. Rubin 1984]



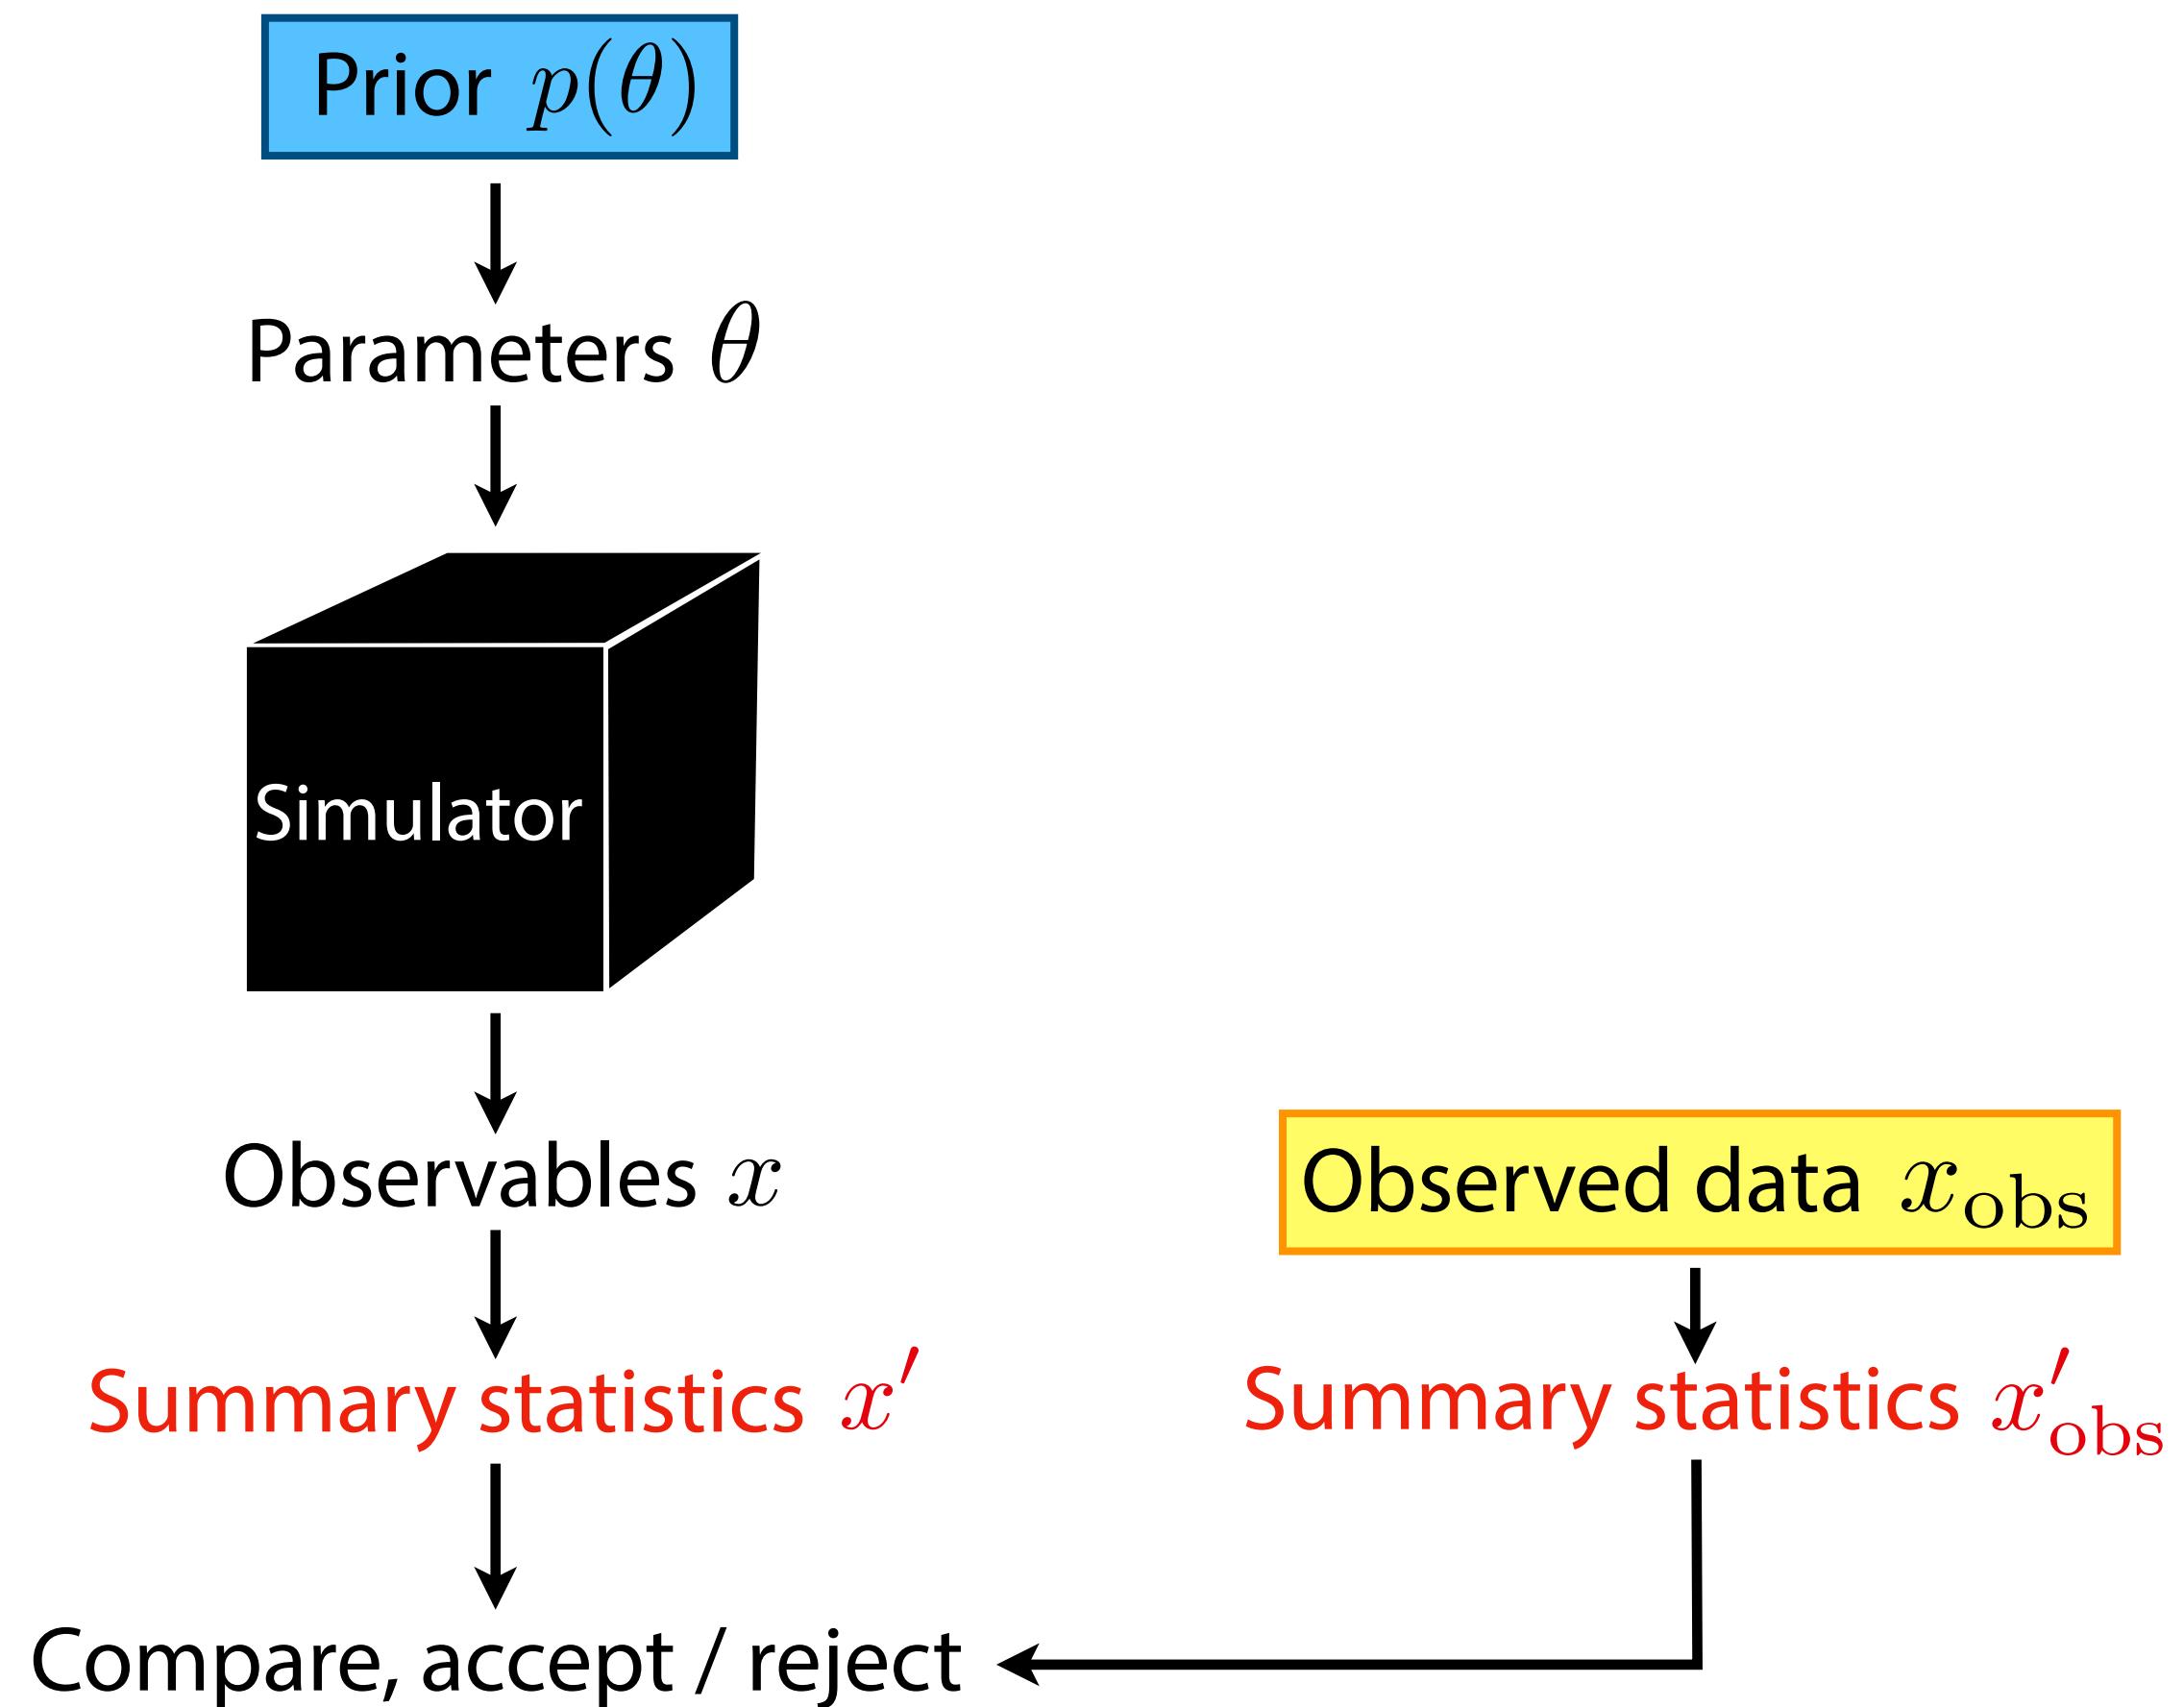
Approximate Bayesian Computation (ABC)

[D. Rubin 1984]



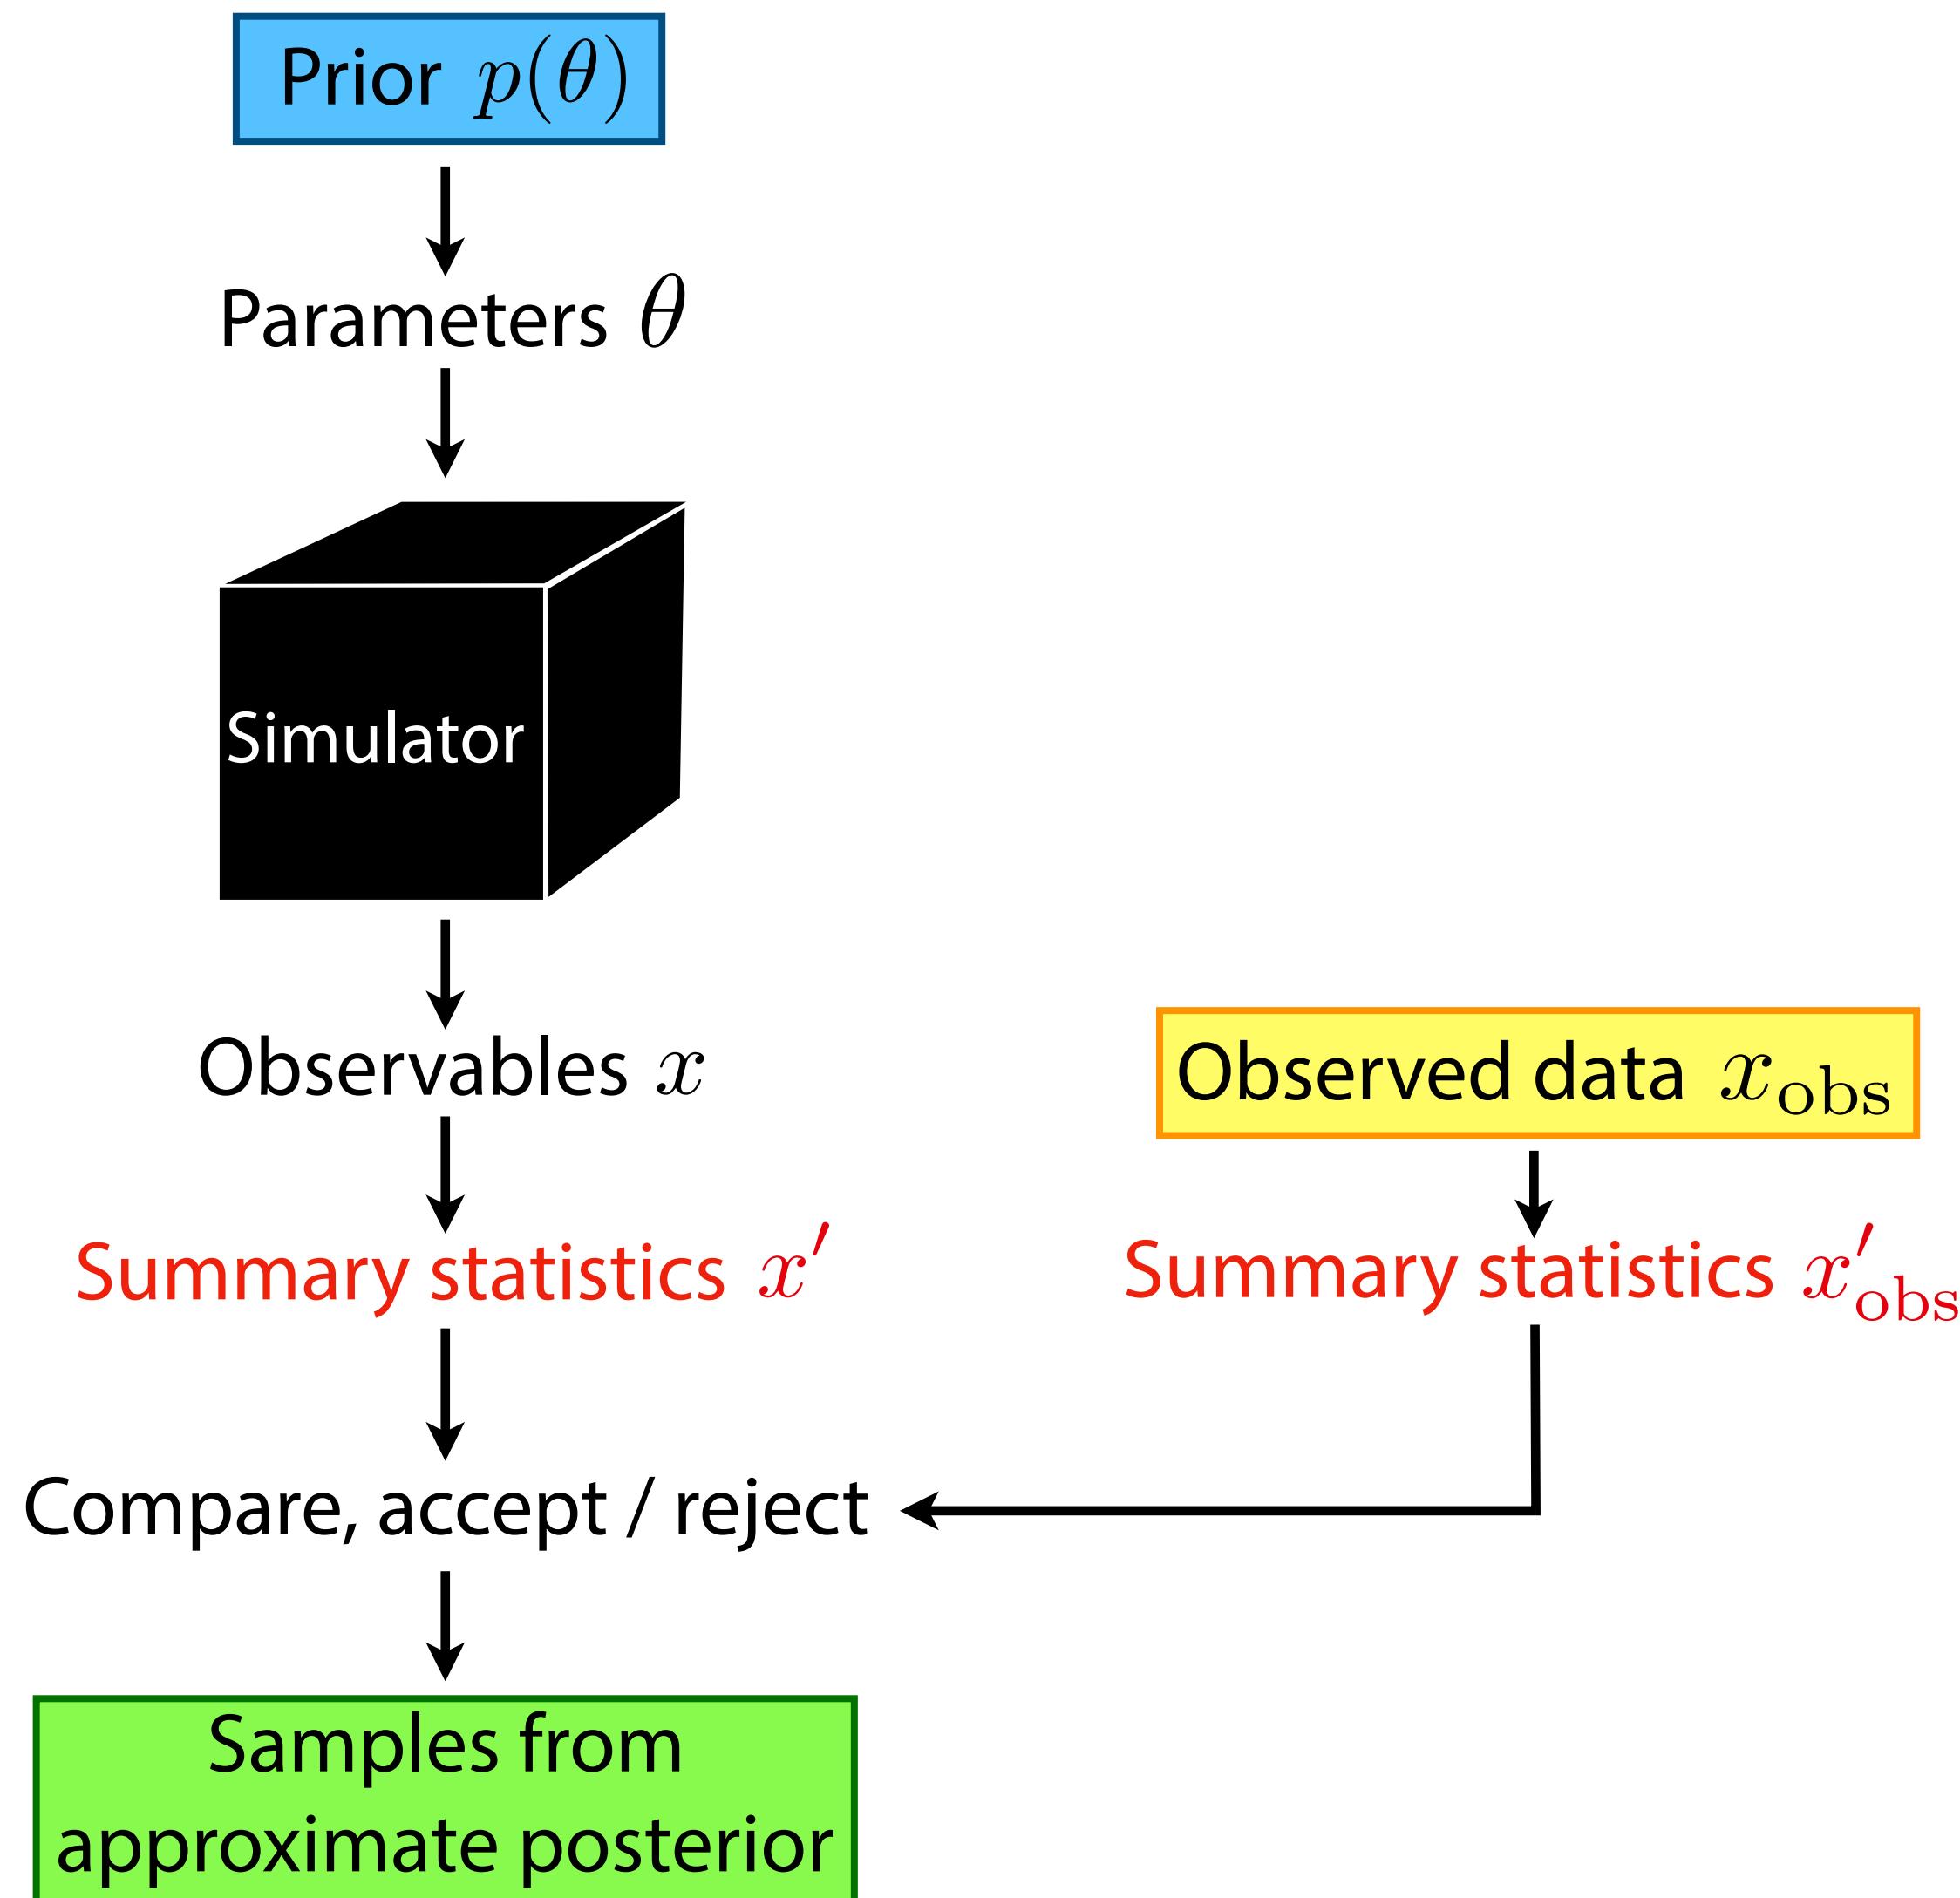
Approximate Bayesian Computation (ABC)

[D. Rubin 1984]



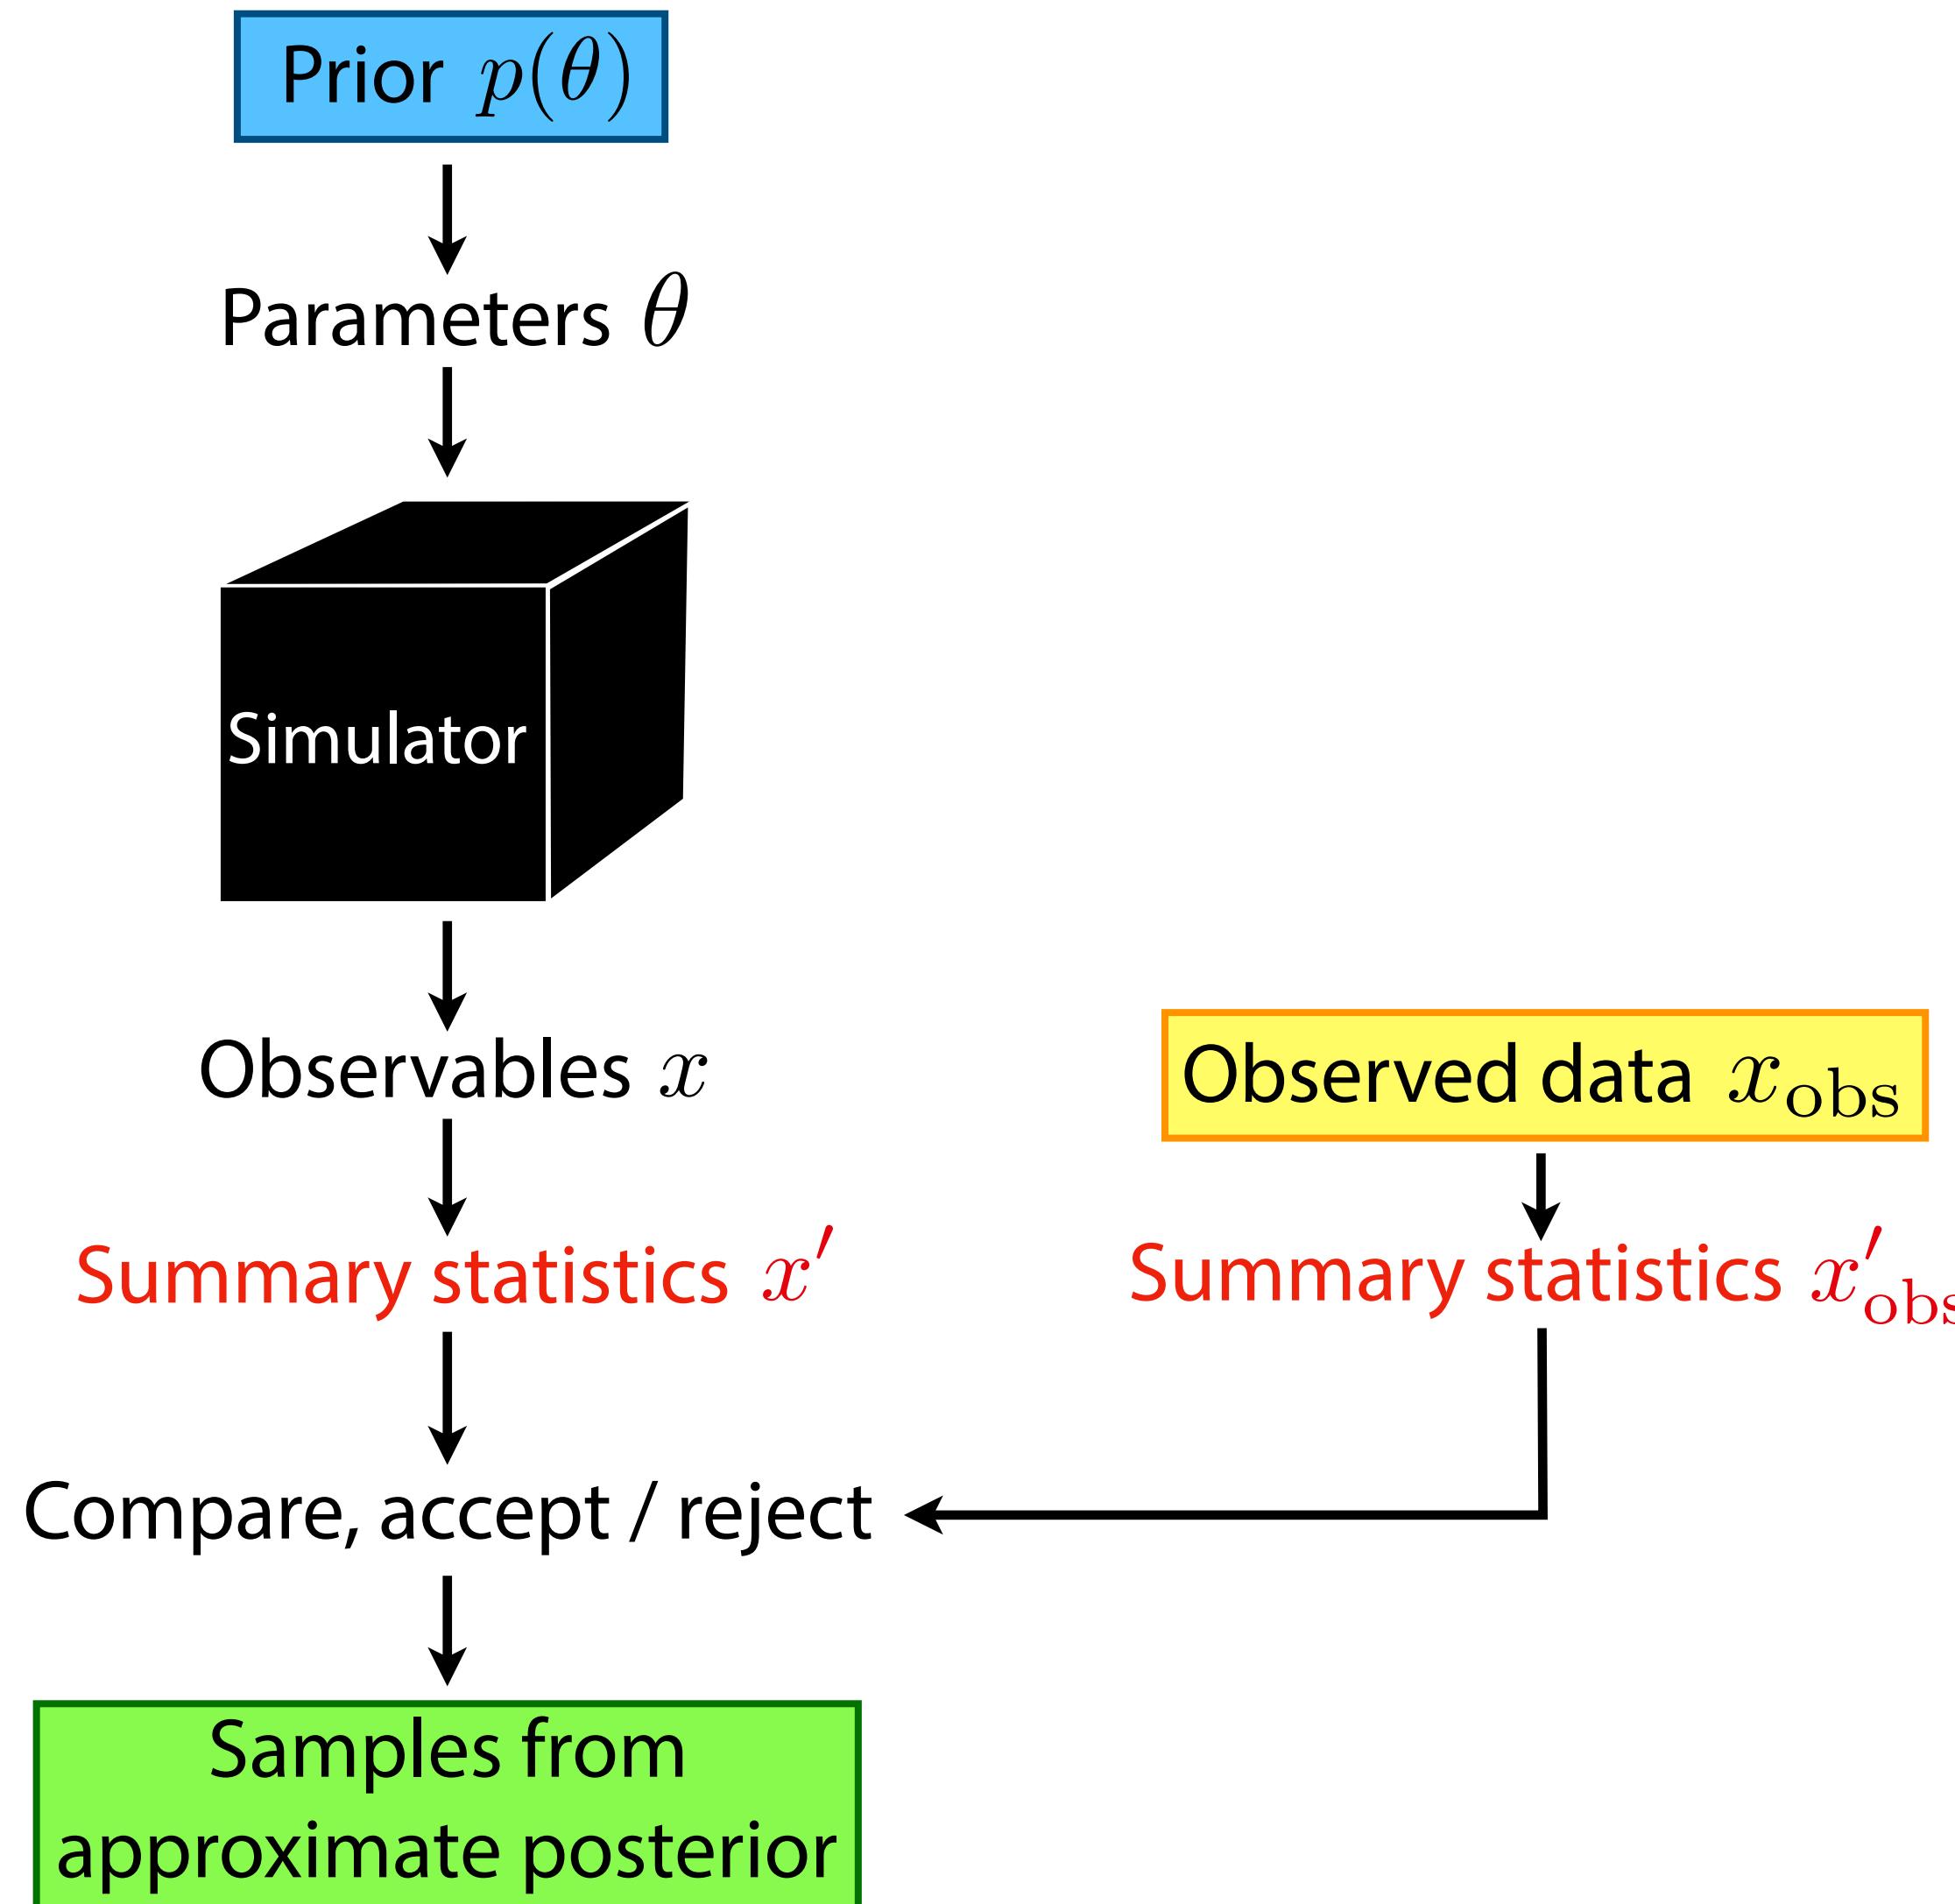
Approximate Bayesian Computation (ABC)

[D. Rubin 1984]



Approximate Bayesian Computation (ABC)

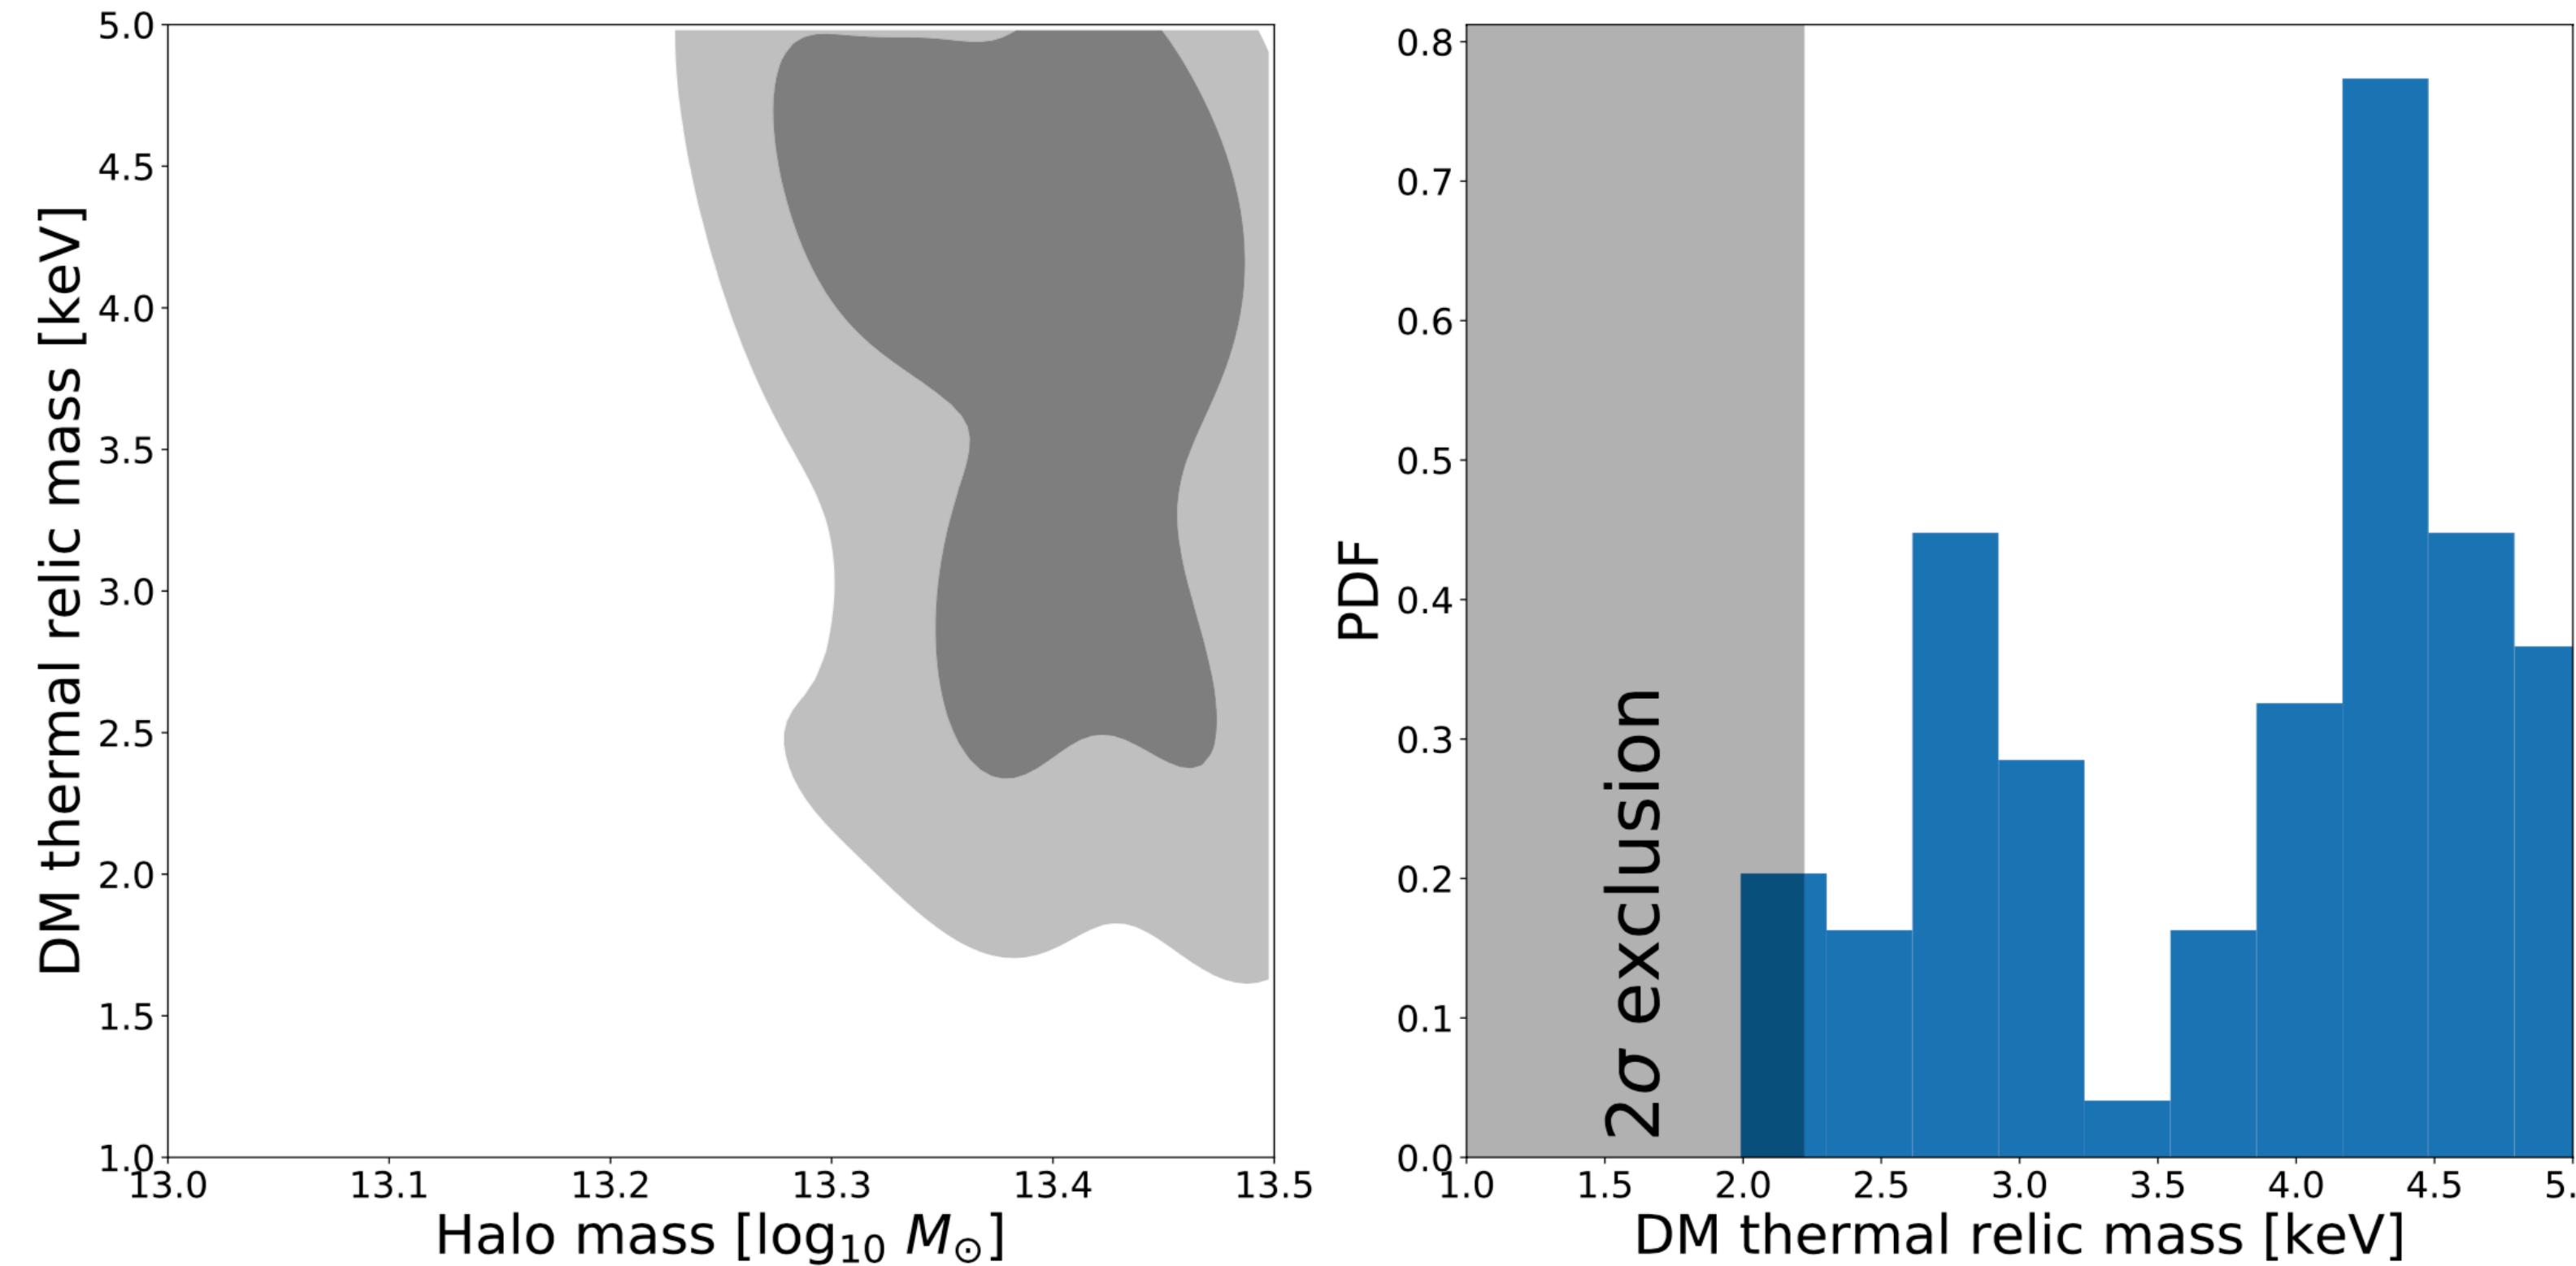
[D. Rubin 1984]



- Reduction to summary statistics makes inference possible, but throws away information
- Rejection algorithm can be very sample inefficient

Substructure ABC

[S. Birrer et al 1702.00009]



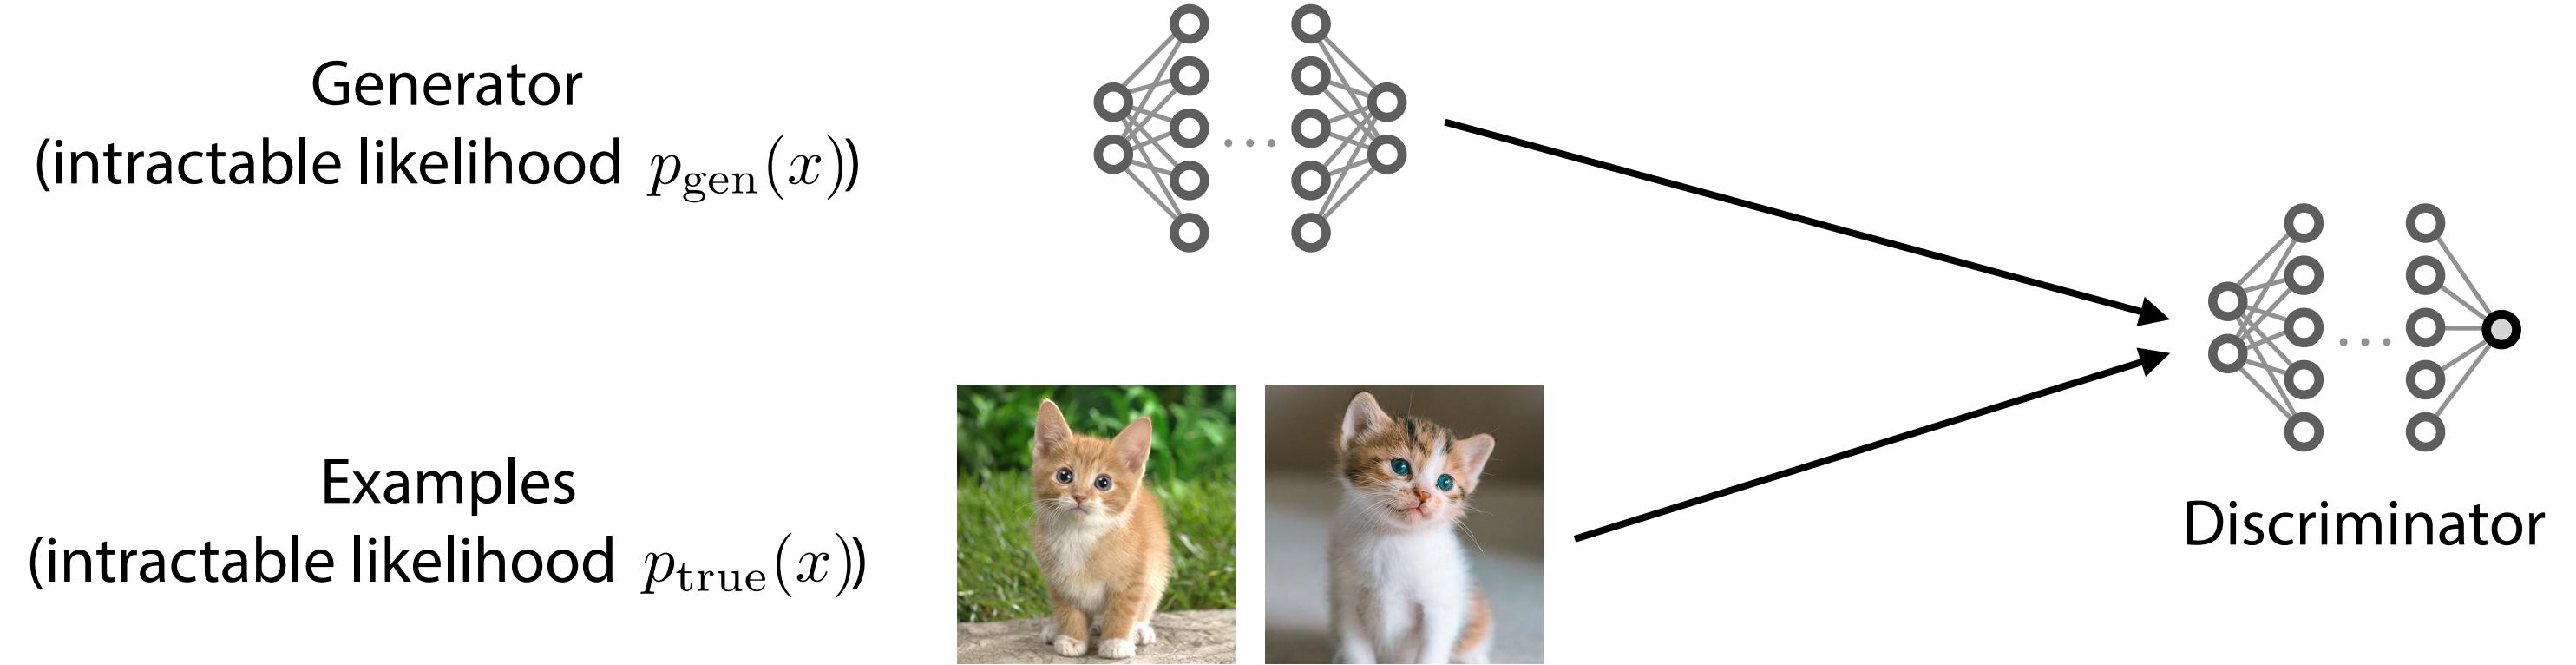
- Successful proof of principle application on HST data (RXJ1131-1231)
- Summary statistics based on correlation functions: results may be overly conservative
- Poor sample efficiency of ABC rejection: “posterior not converged”

Now, we can also use machine learning
to enable inference.

[K. Cranmer, JB, G. Louppe 1911.01429]

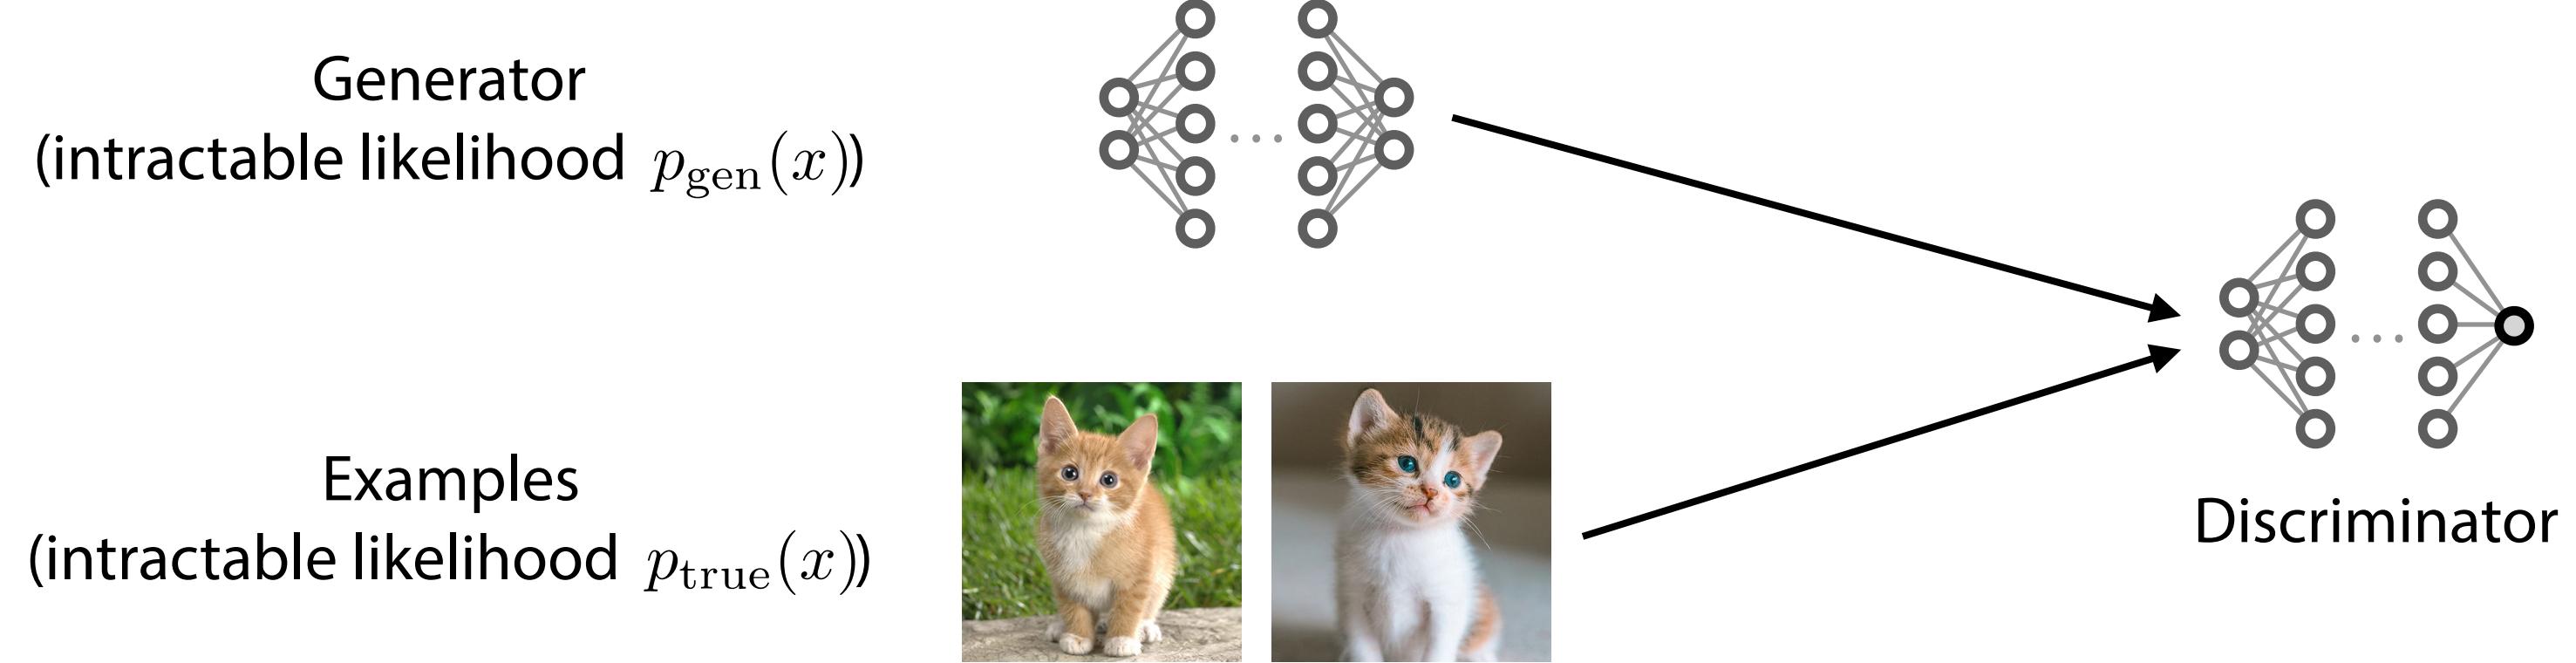
The likelihood ratio trick

- Generative Adversarial Networks (GANs):



The likelihood ratio trick

- Generative Adversarial Networks (GANs):



[I. Goodfellow et al. 1406.2661]

Discriminator learns decision function

$$s(x) \rightarrow \frac{p_{\text{true}}(x)}{p_{\text{gen}}(x) + p_{\text{true}}(x)}$$

The likelihood ratio trick

- Generative Adversarial Networks (GANs)

Generator
(intractable likelihood p_g)

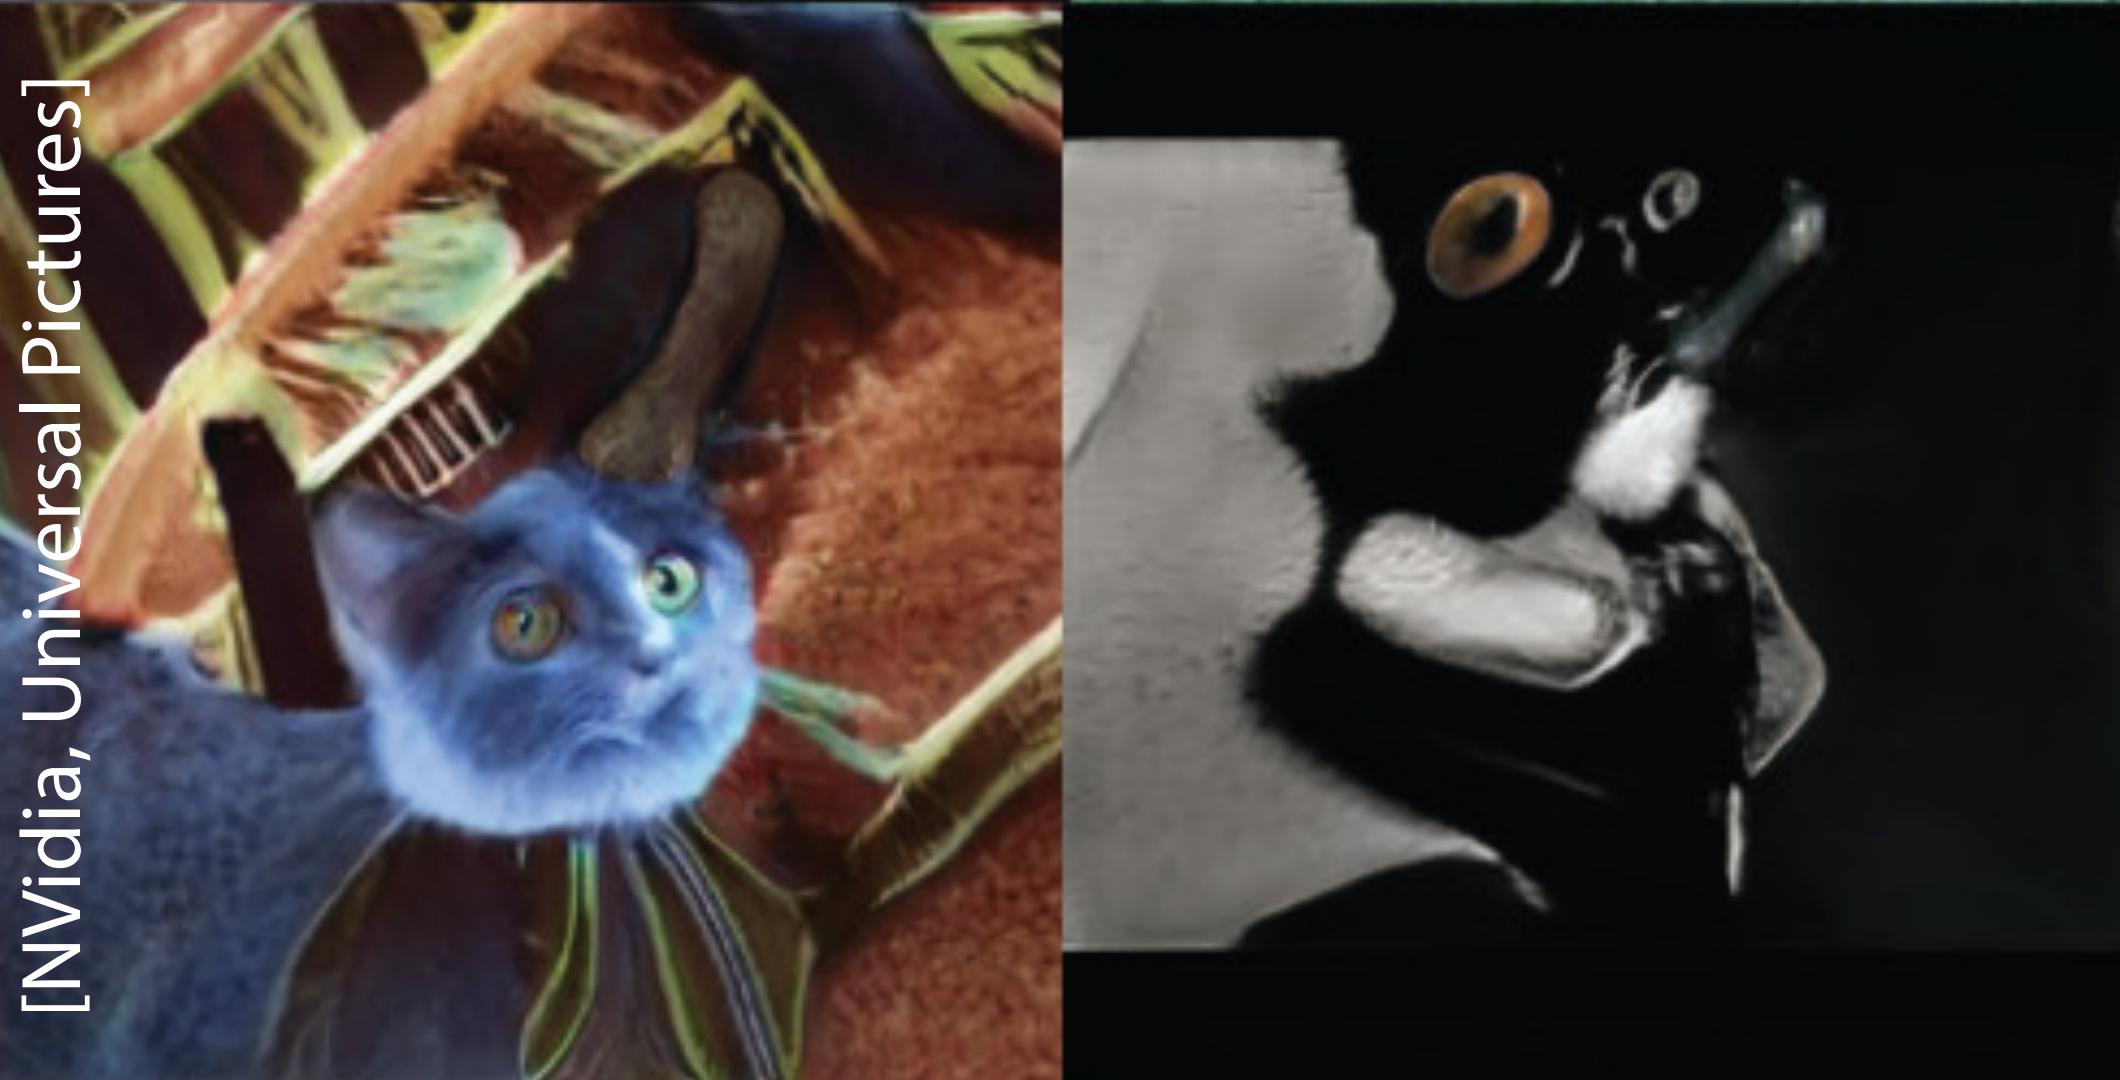


Examples
(intractable likelihood p_t)



[Goodfellow et al. 1406.2661]

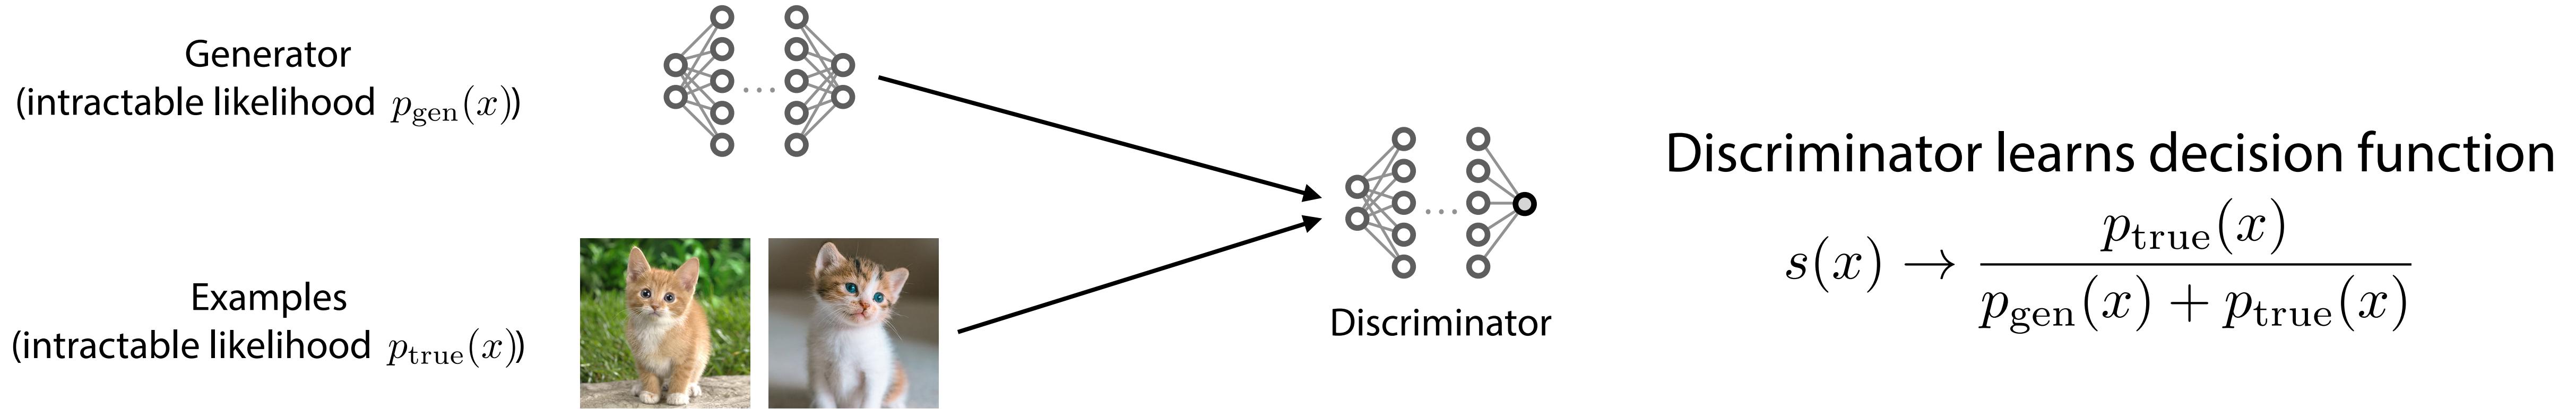
decision function
$$\frac{p_{\text{true}}(x)}{p_{\text{true}}(x) + p_{\text{generated}}(x)}$$



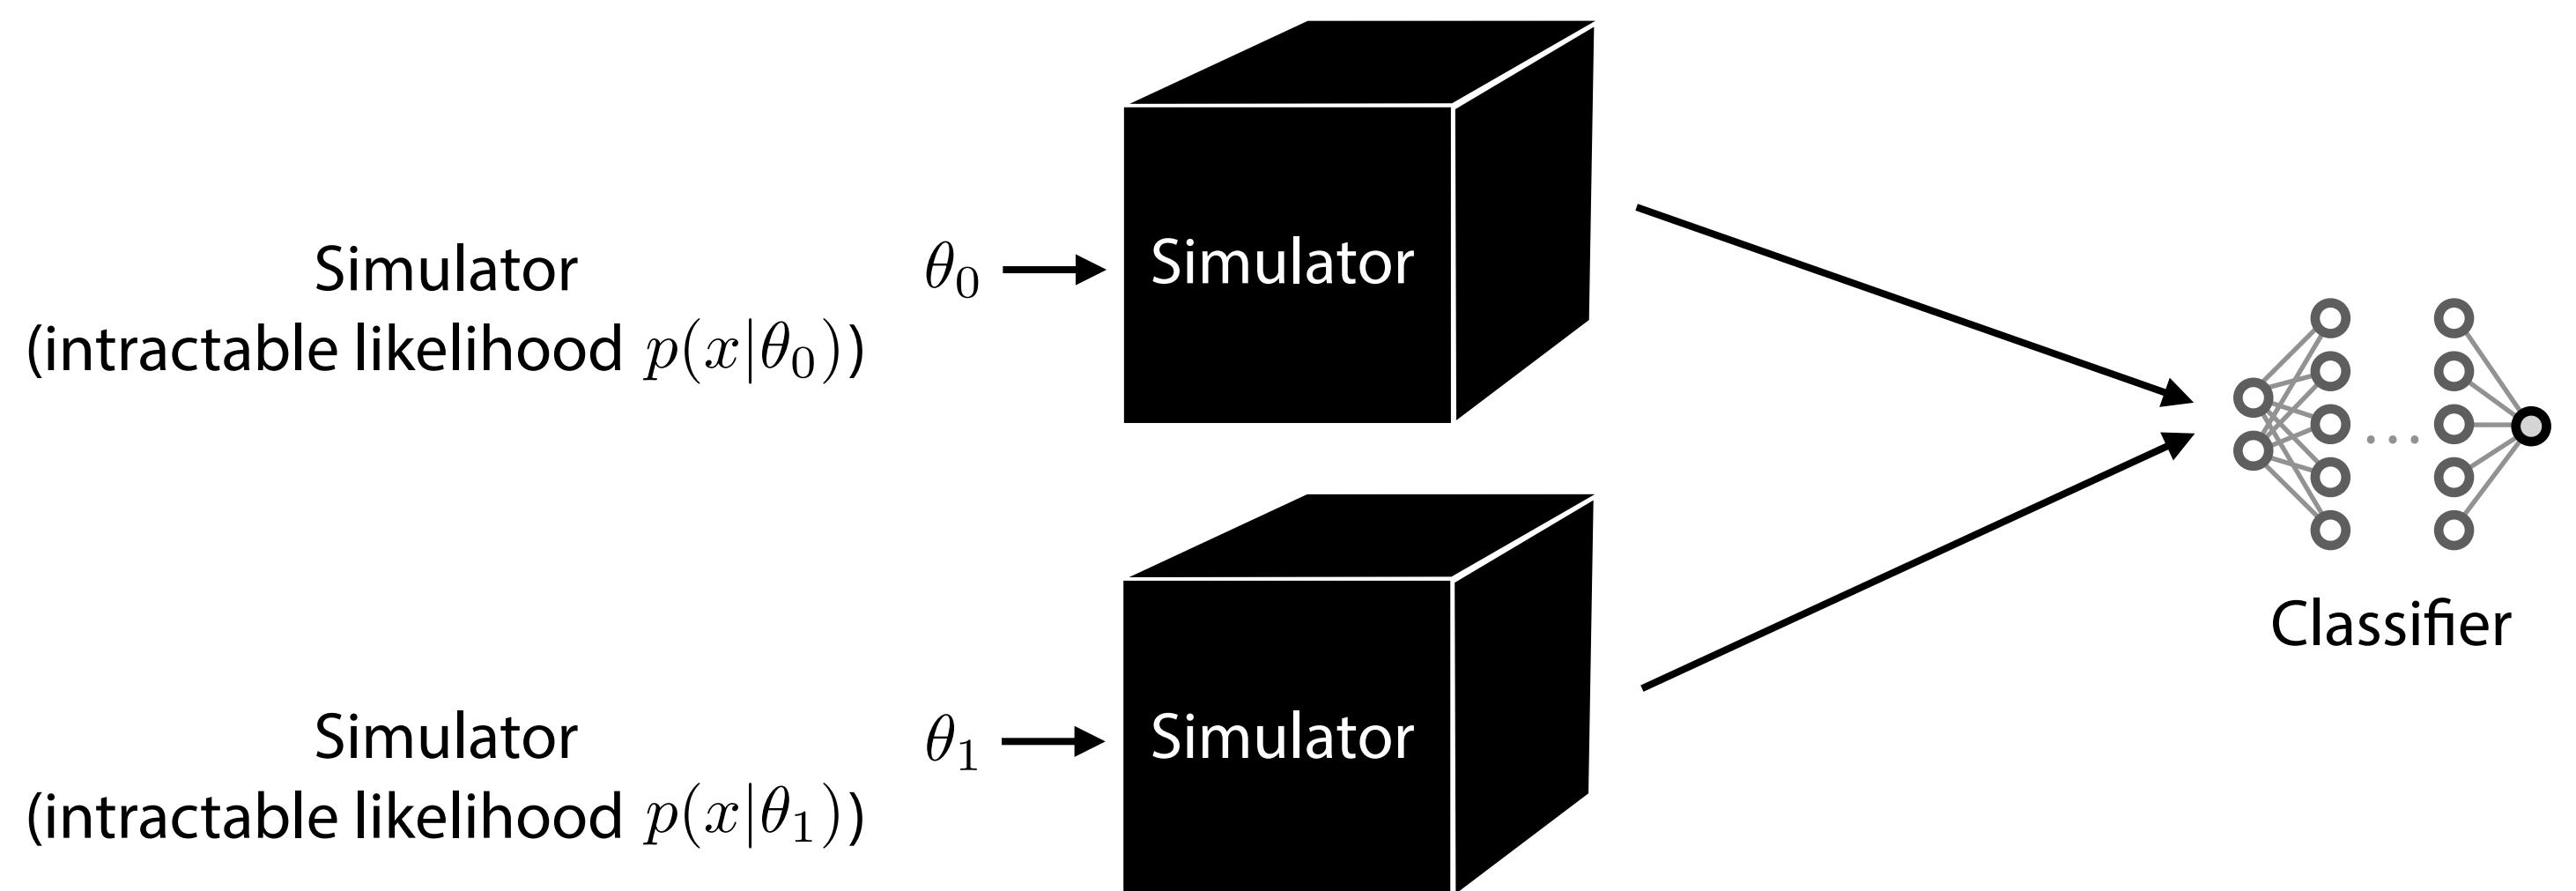
[Nvidia, Universal Pictures]

The likelihood ratio trick

- Generative Adversarial Networks (GANs):

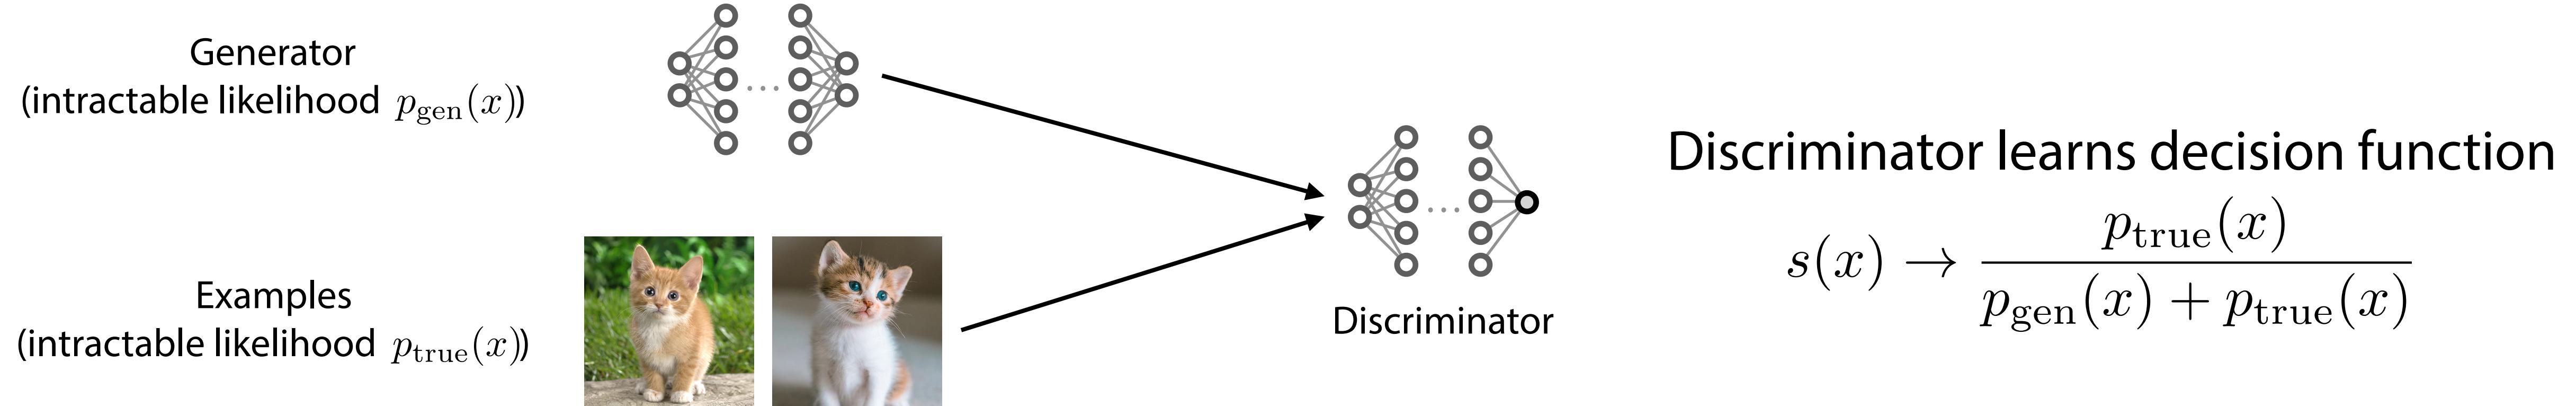


- Similarly, we can train a classifier between two sets of simulated samples

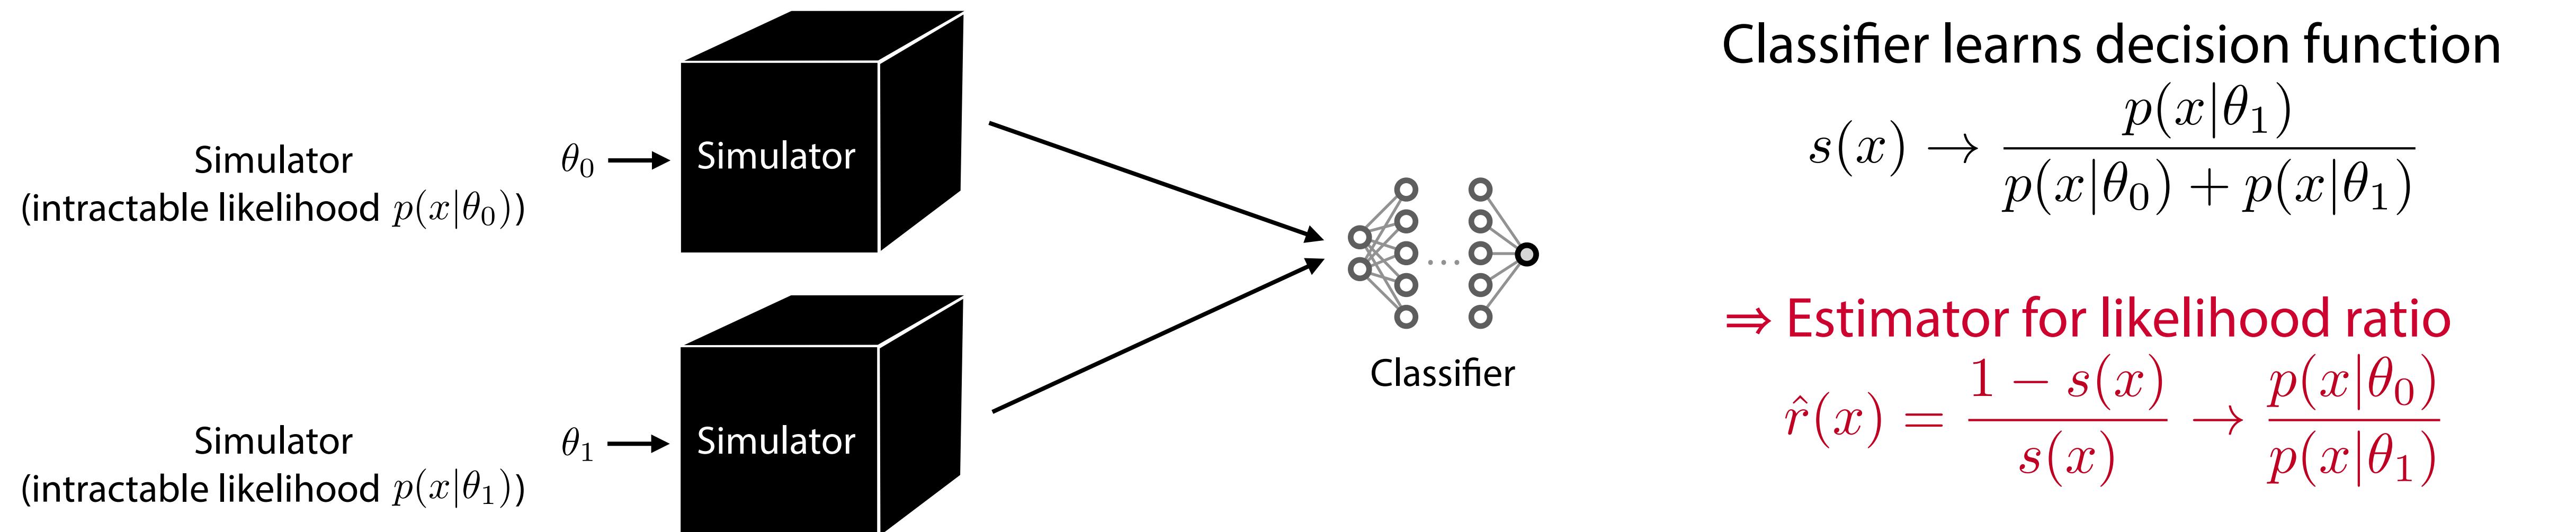


The likelihood ratio trick

- Generative Adversarial Networks (GANs):

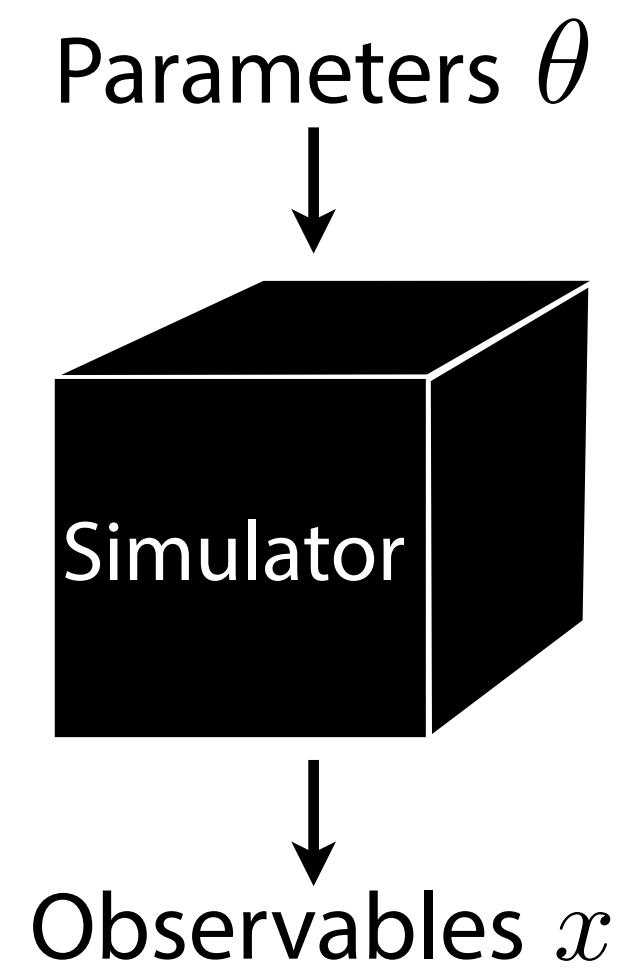


- Similarly, we can train a classifier between two sets of simulated samples



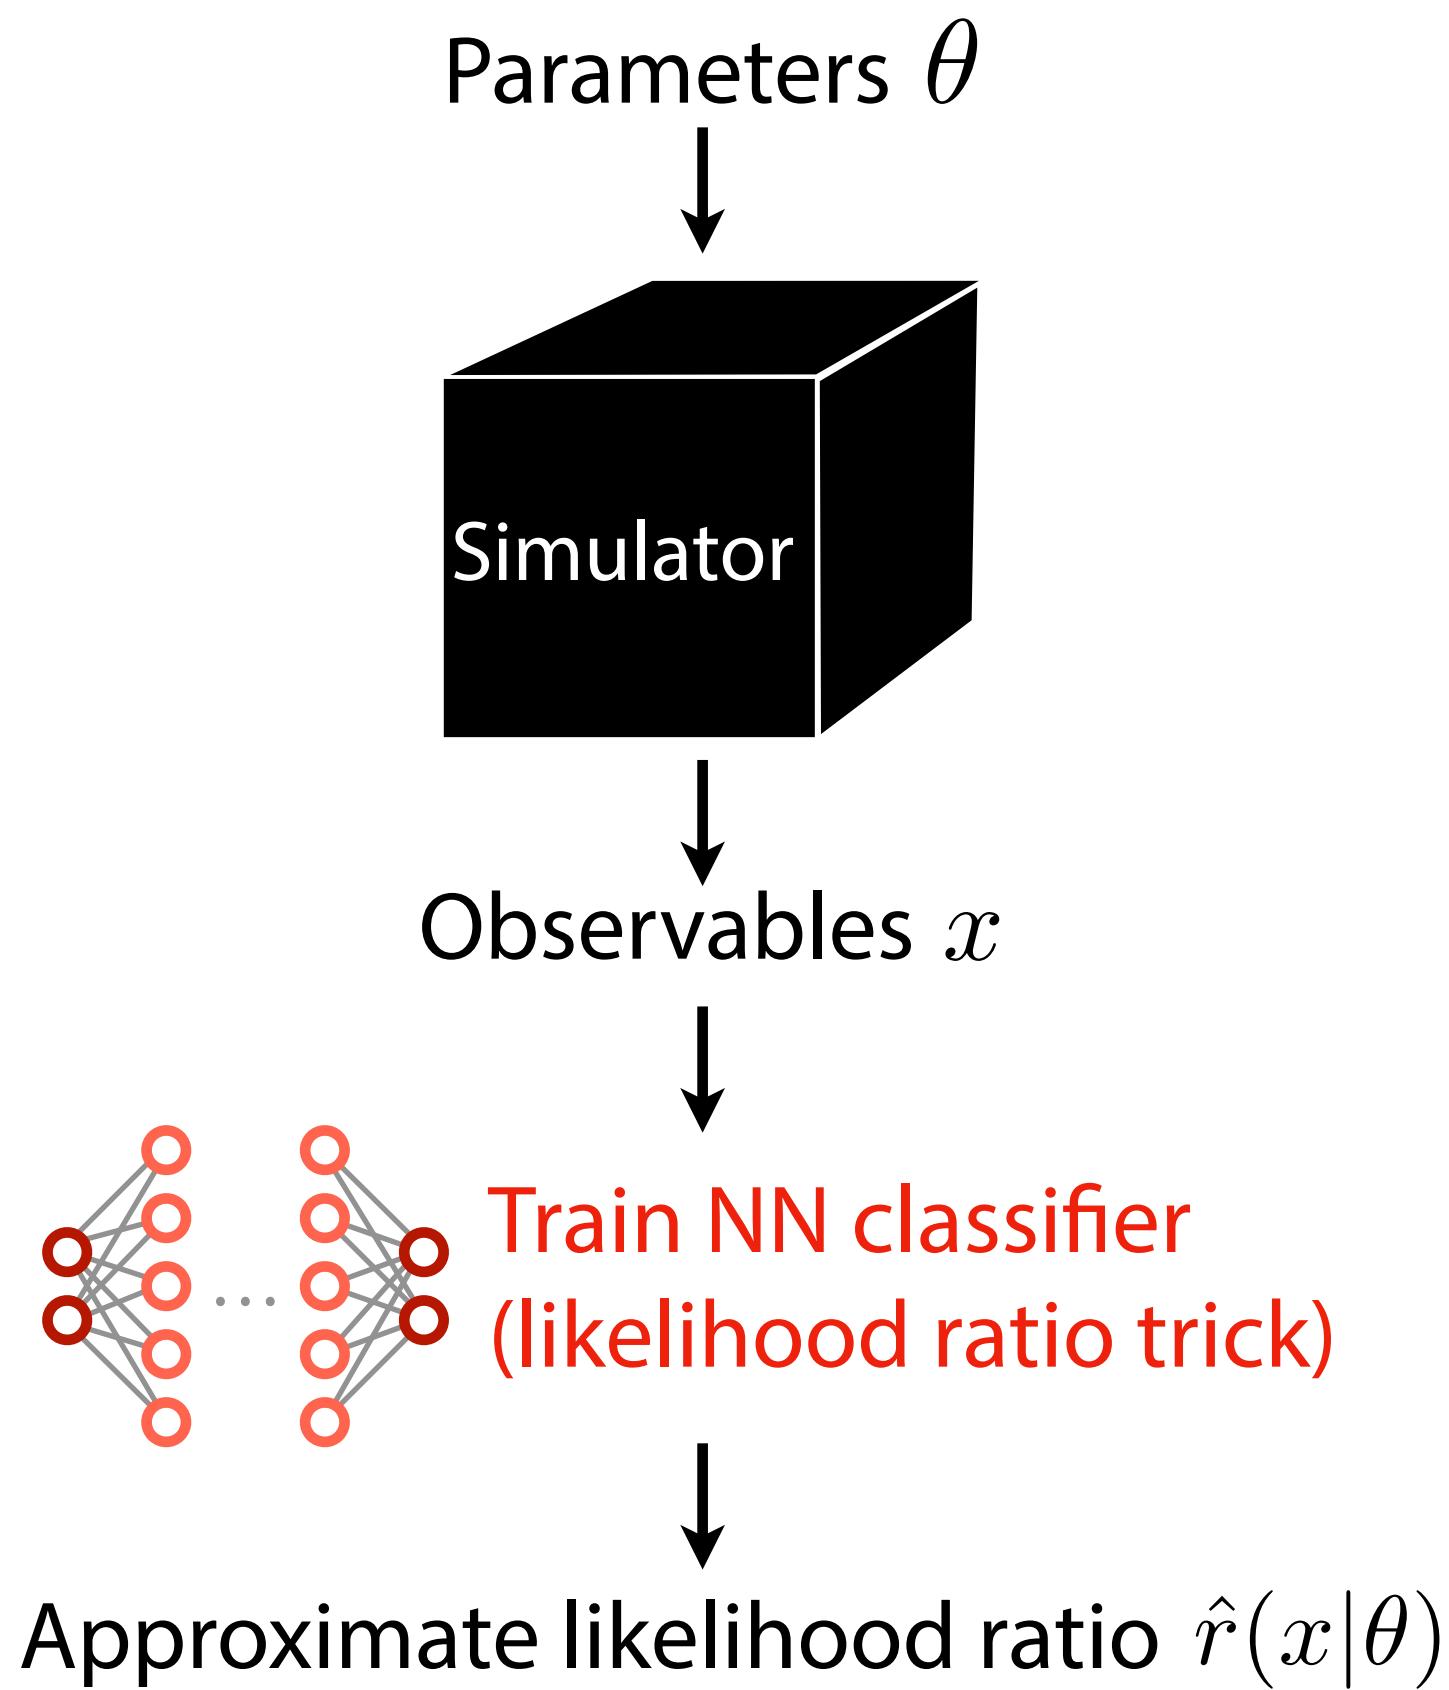
Inference by likelihood ratio trick

[K. Cranmer J. Pavez, G. Louppe 1506.02169]



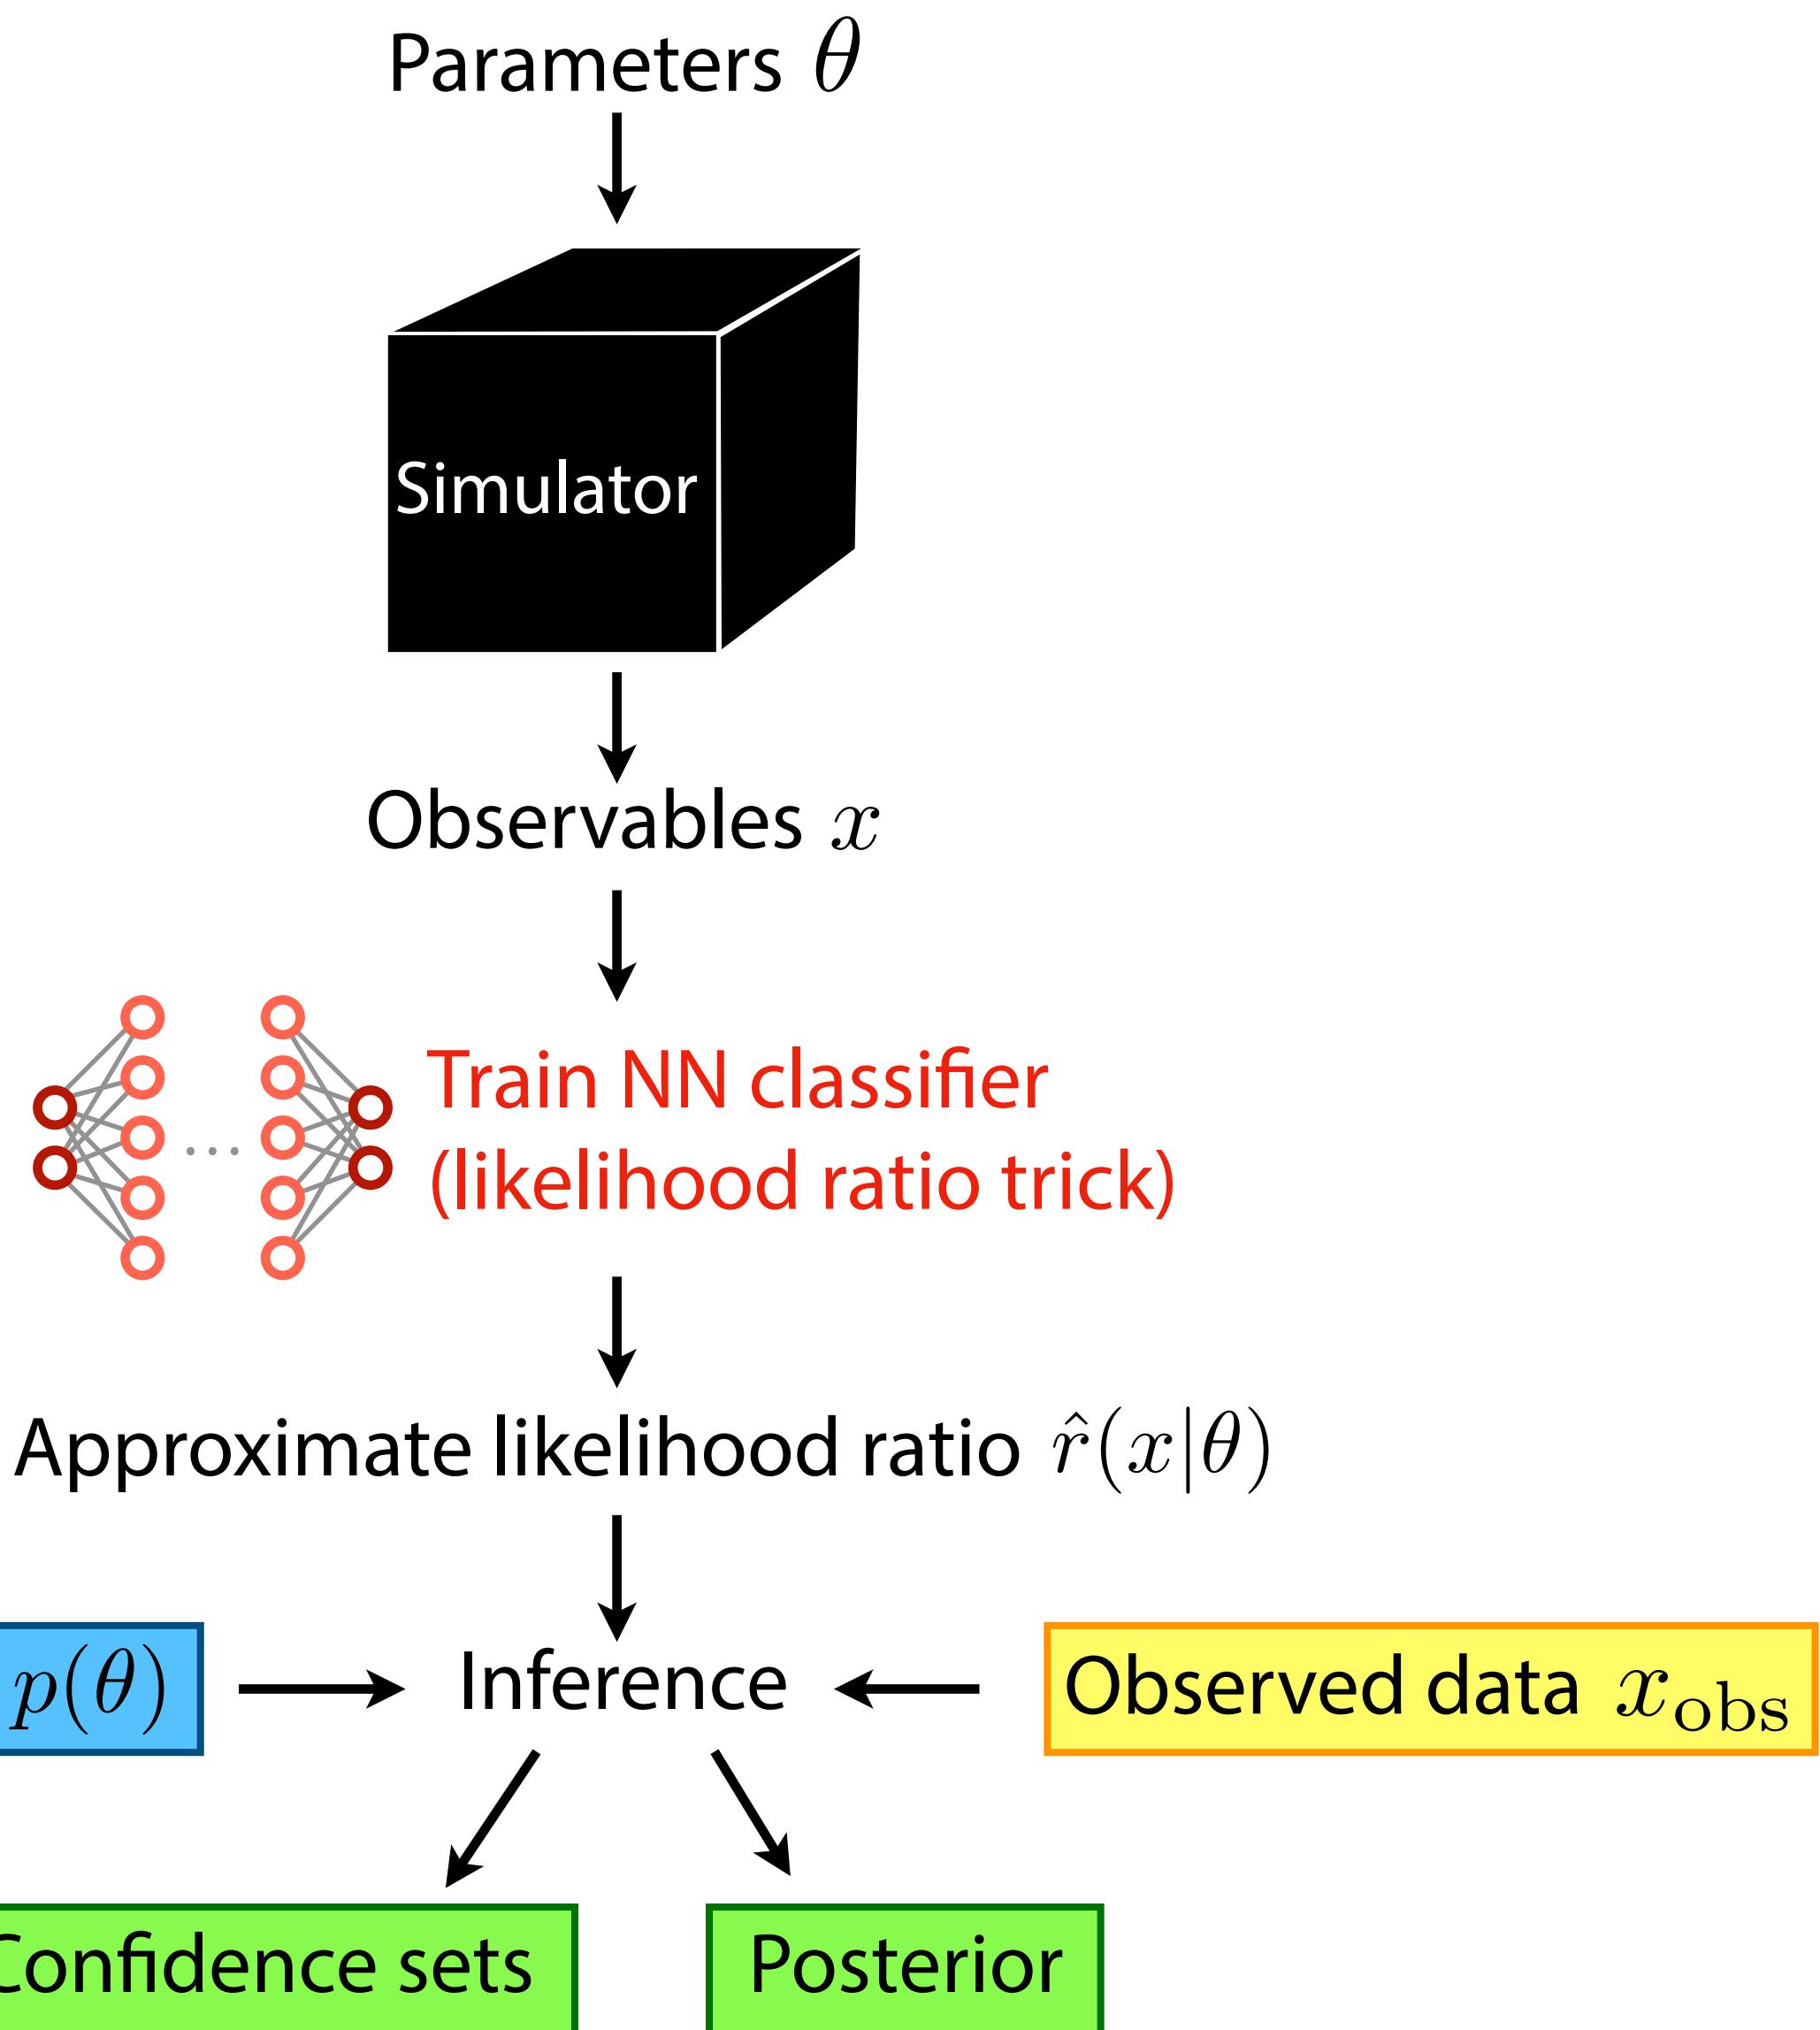
Inference by likelihood ratio trick

[K. Cranmer J. Pavez, G. Louppe 1506.02169]



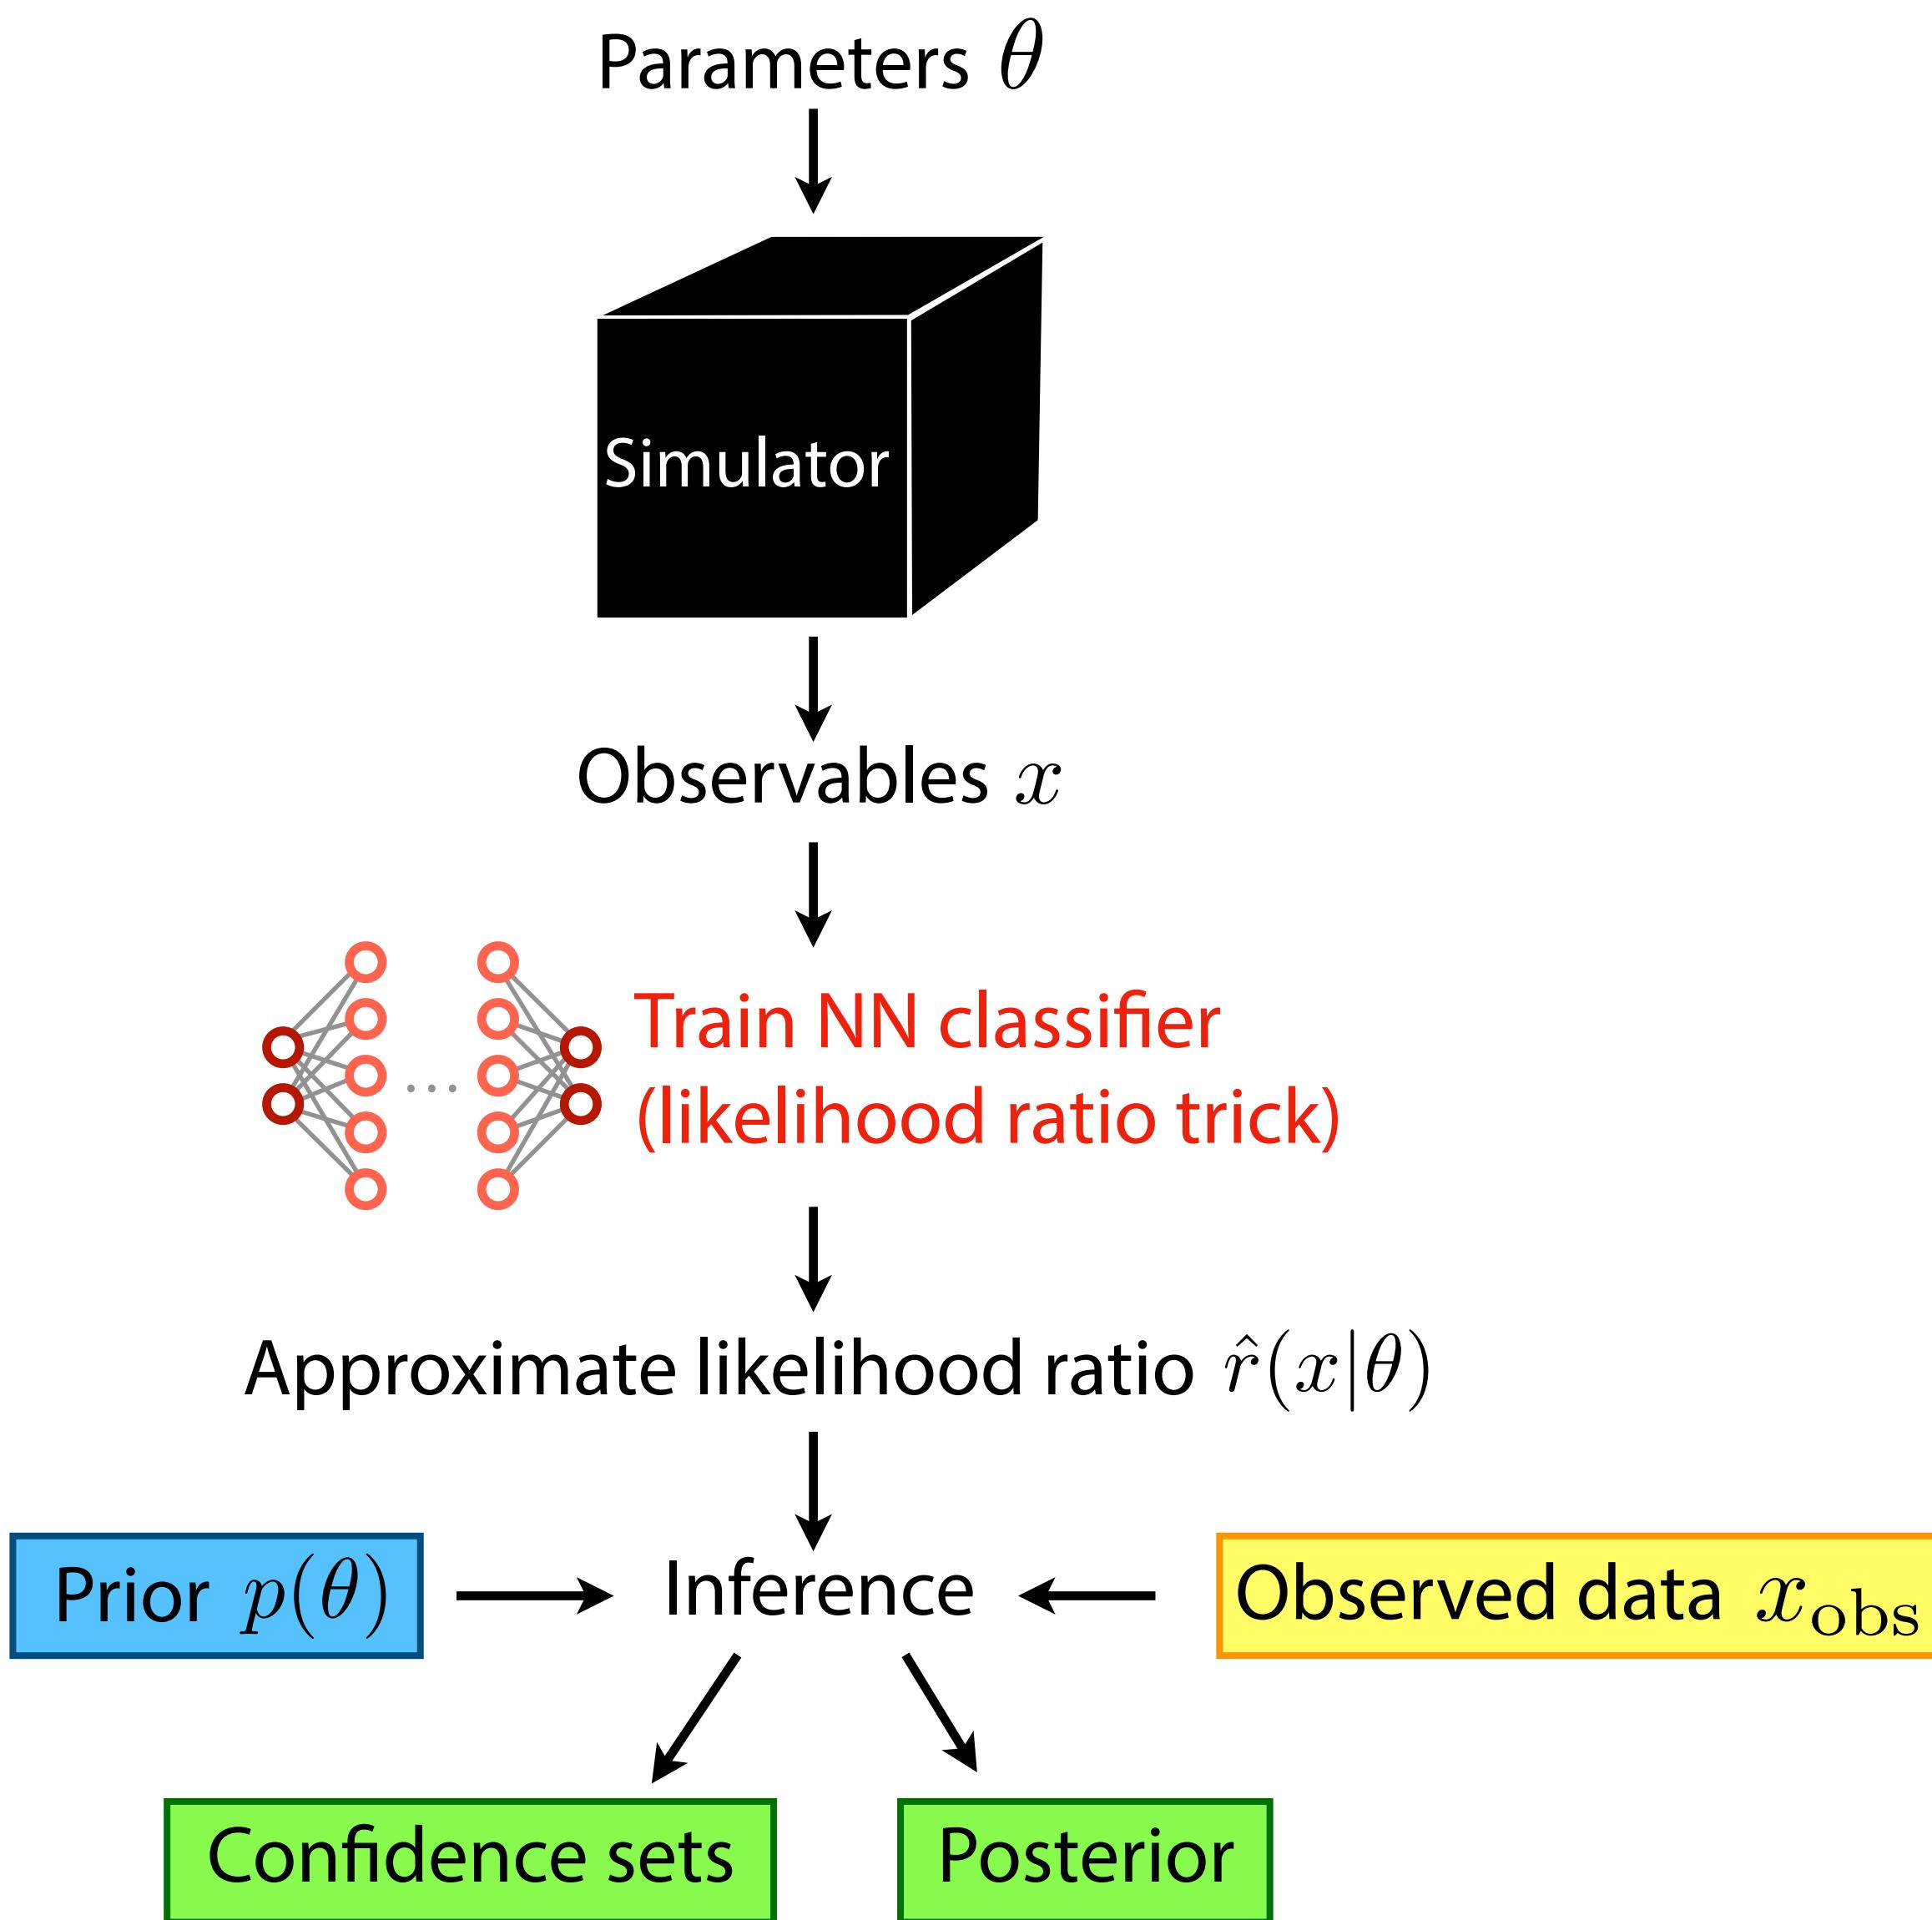
Inference by likelihood ratio trick

[K. Cranmer J. Pavez, G. Louppe 1506.02169]



Inference by likelihood ratio trick

[K. Cranmer J. Pavez, G. Louppe 1506.02169]

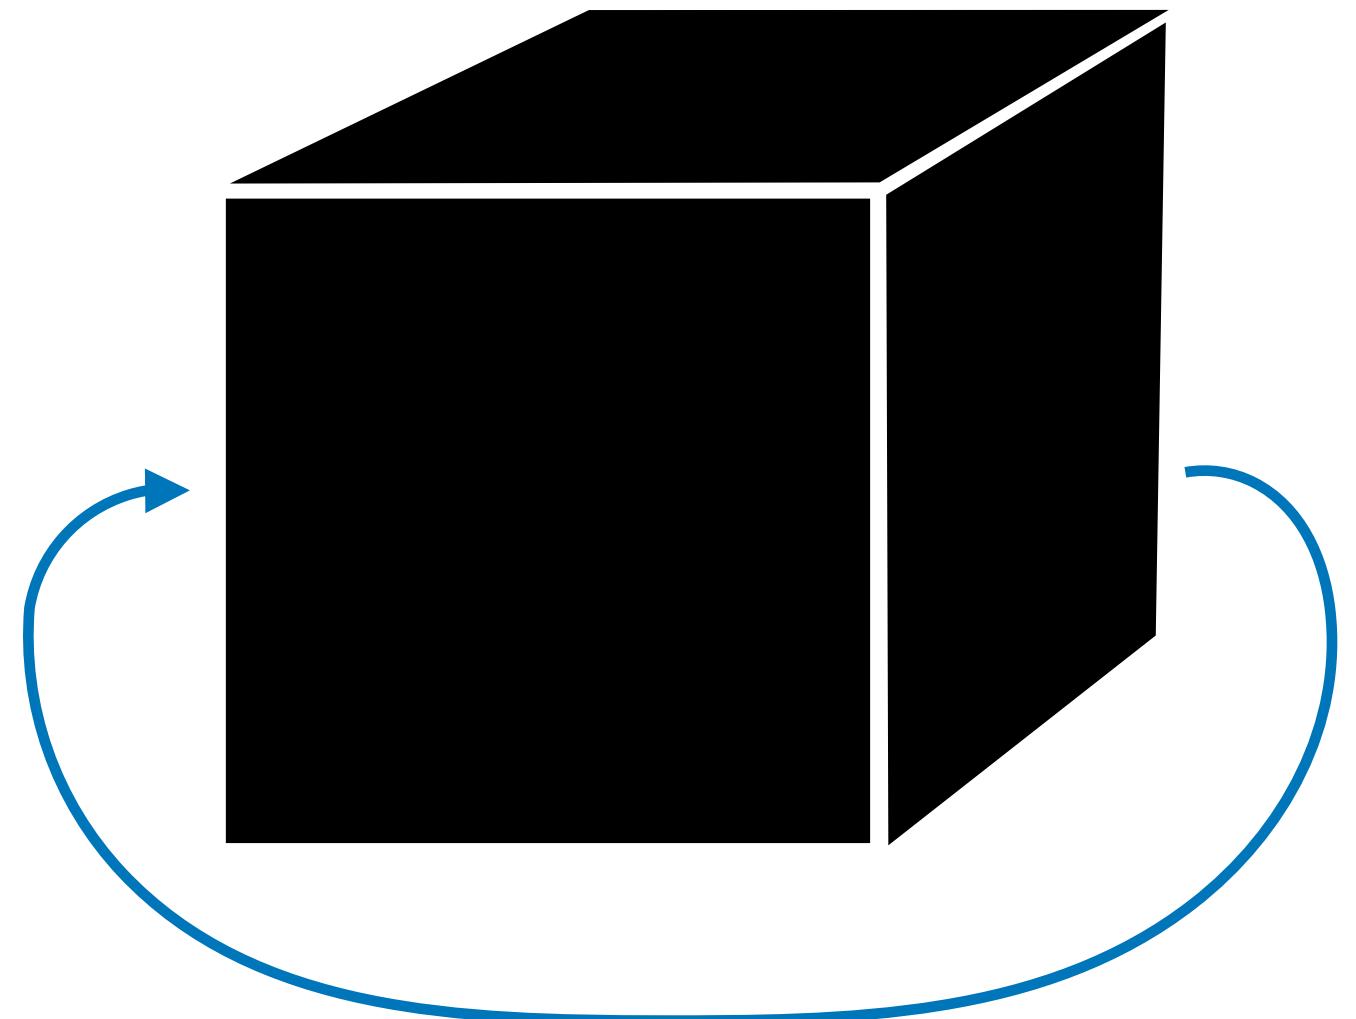


- Instead of the intractable likelihood of the simulator, we use the likelihood ratio learned by a NN from simulated data
- Scales well to high-dimensional data (no compression to summary stats necessary)
- Amortized: After upfront simulation + training phase, inference is efficient for new data or prior
- Alternative approaches: train NN to learn likelihood or posterior

[review: K. Cranmer, JB, G. Louppe 1911.01429]

Improving sample efficiency

[K. Cranmer, JB, G. Louppe 1911.01429]

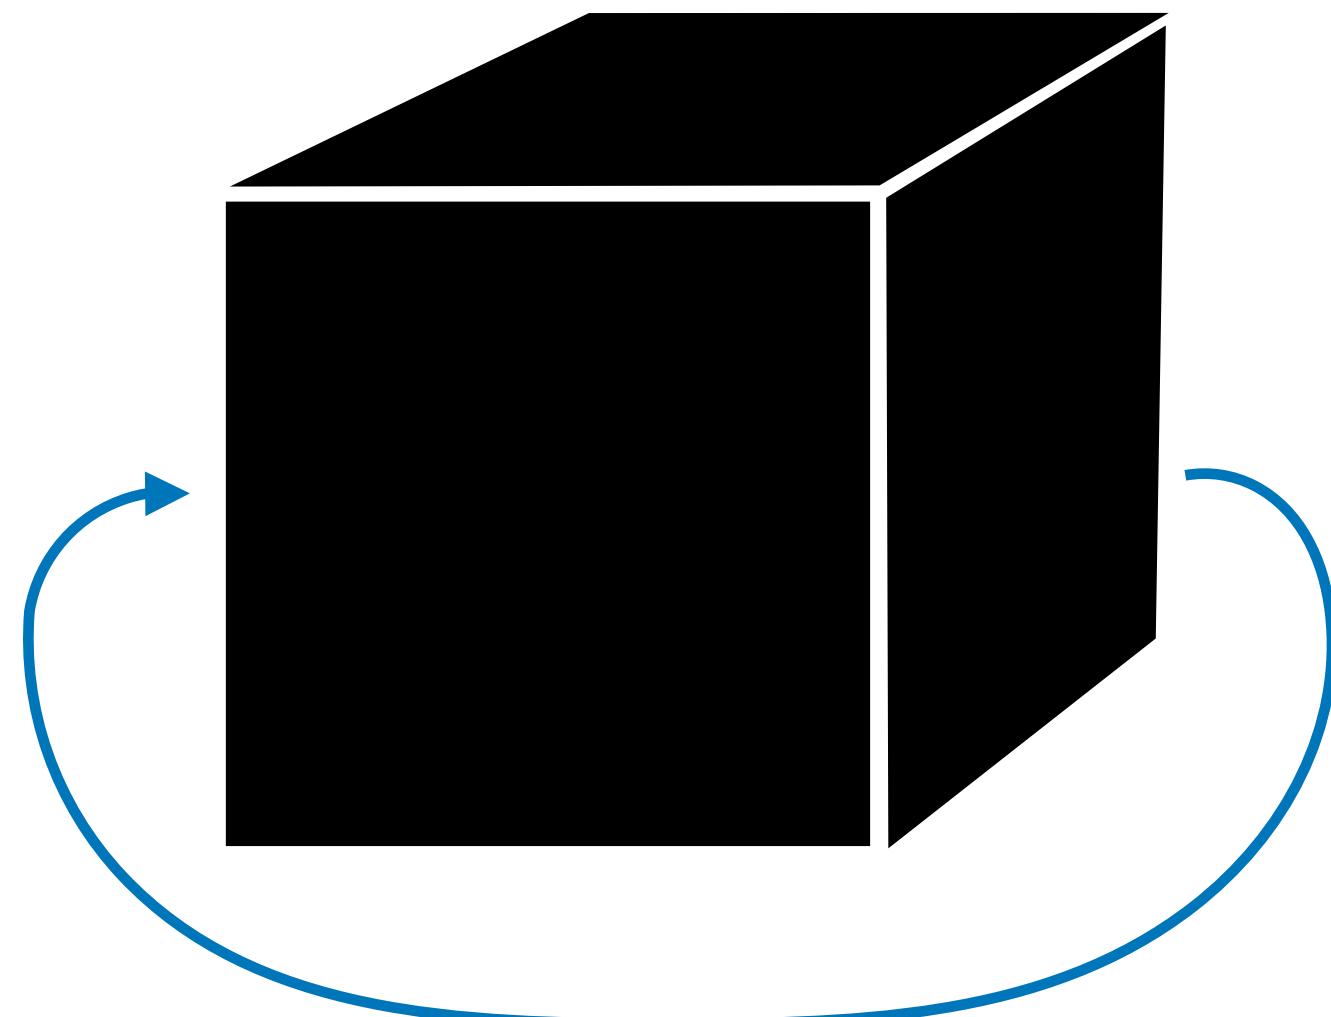


Active learning:

Iteratively guide simulator to
important parameter regions
based on past results

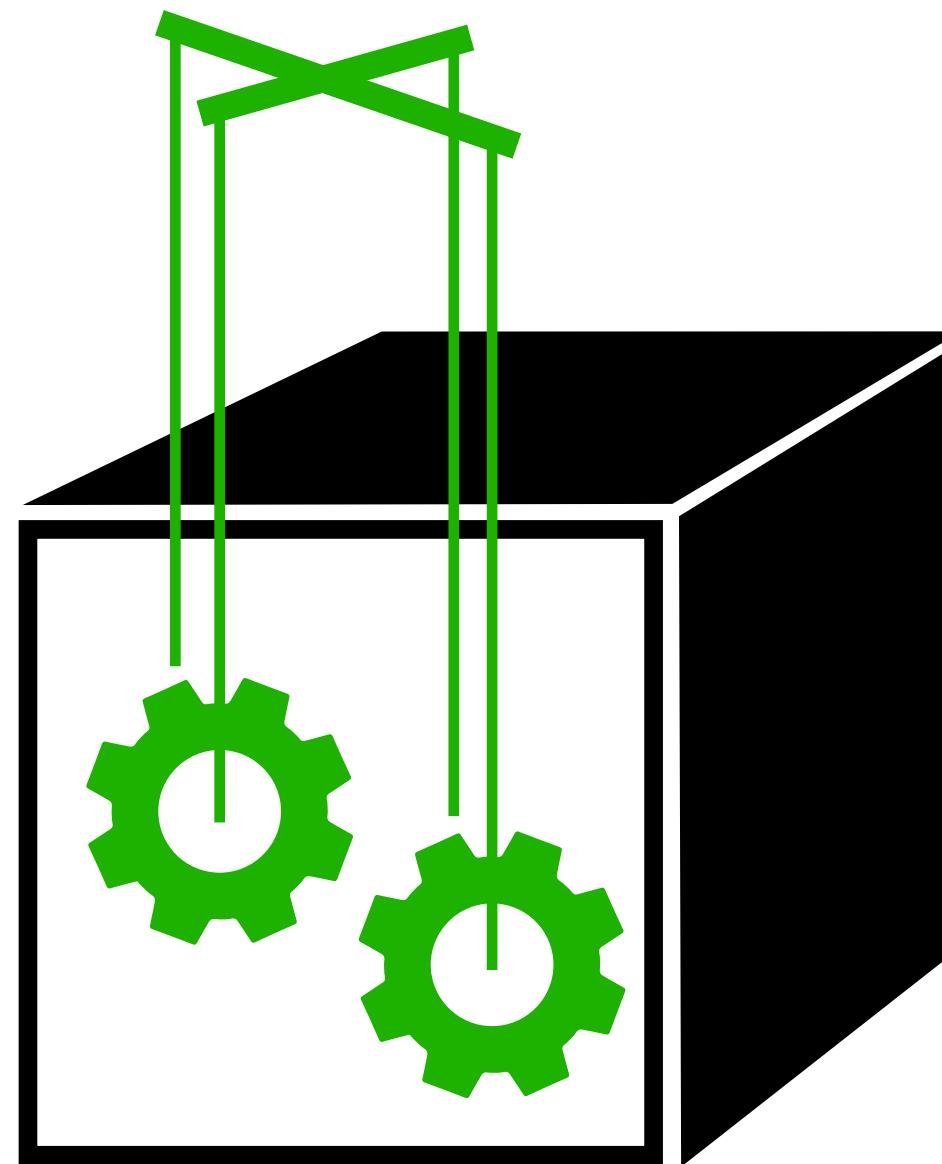
Improving sample efficiency

[K. Cranmer, JB, G. Louppe 1911.01429]



Active learning:

Iteratively guide simulator to important parameter regions based on past results

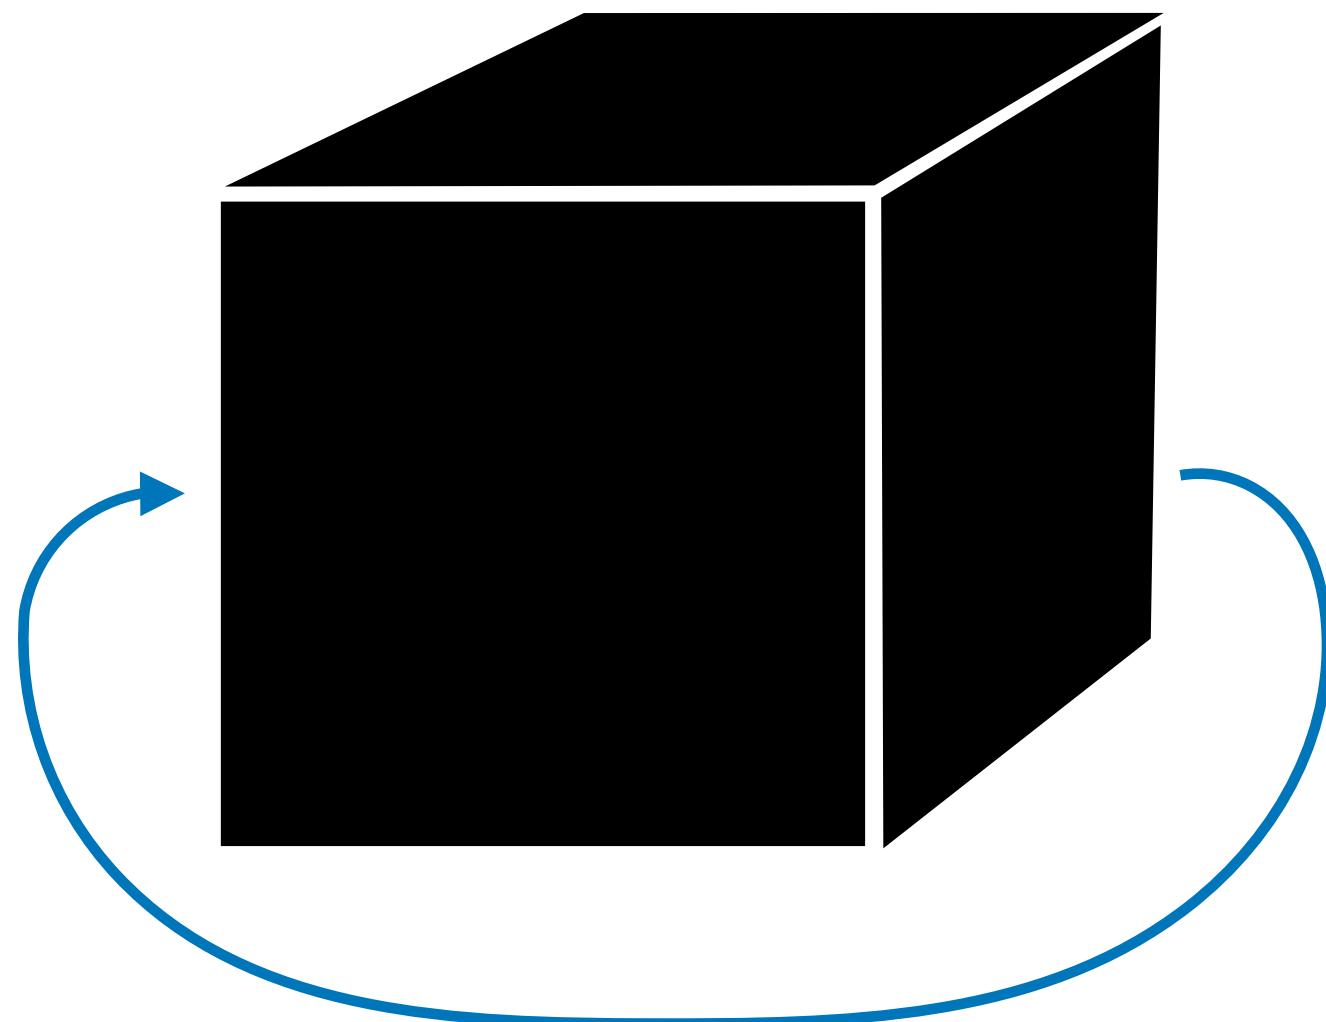


Probabilistic programming:

Explicitly write simulator as probabilistic program, with ability to condition execution trace on observations

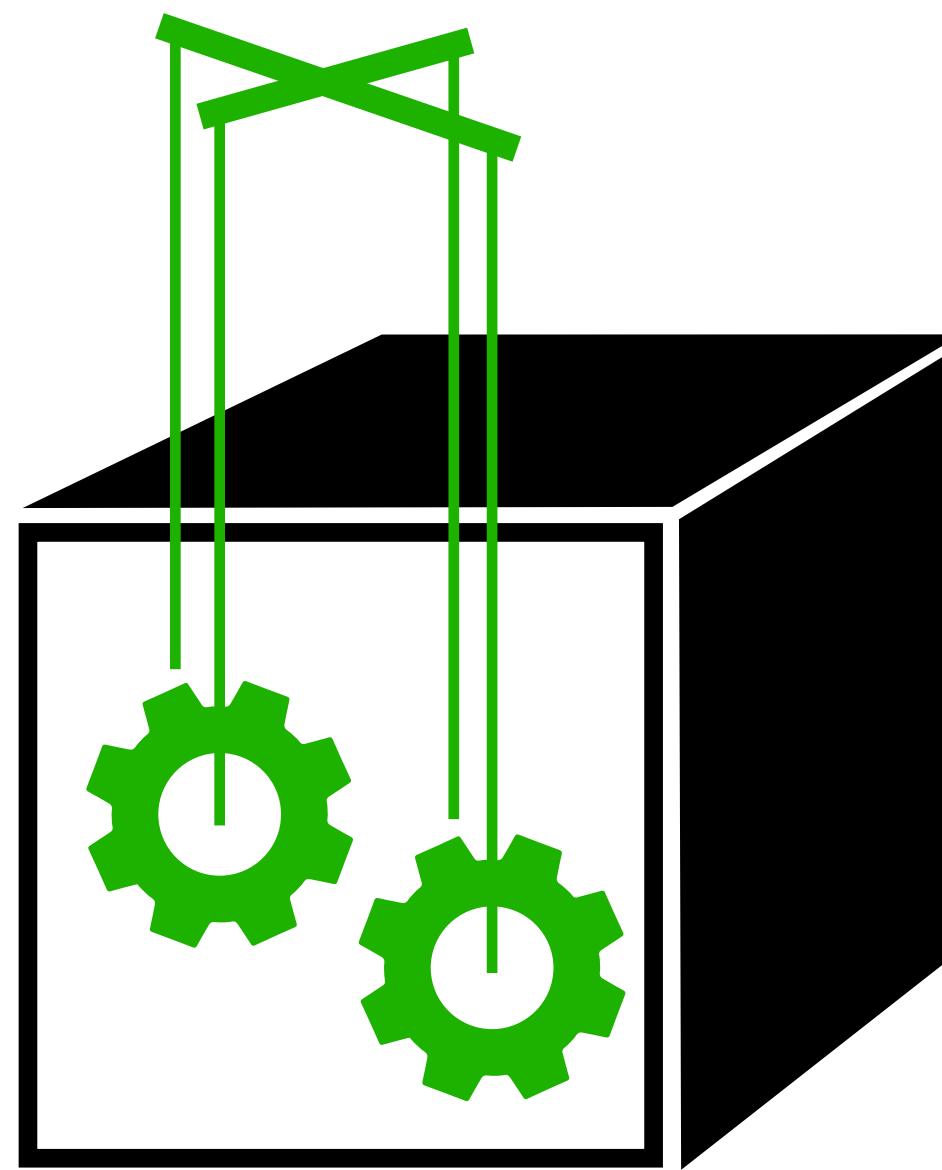
Improving sample efficiency

[K. Cranmer, JB, G. Louppe 1911.01429]



Active learning:

Iteratively guide simulator to important parameter regions based on past results



Probabilistic programming:

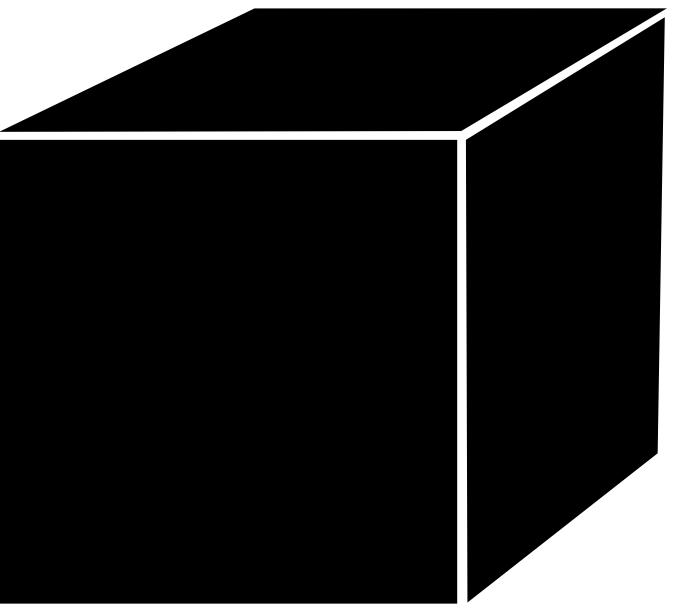
Explicitly write simulator as probabilistic program, with ability to condition execution trace on observations



Mining gold:

Extract and leverage information from the simulator that characterizes the latent process

Improving robustness to mismodelling

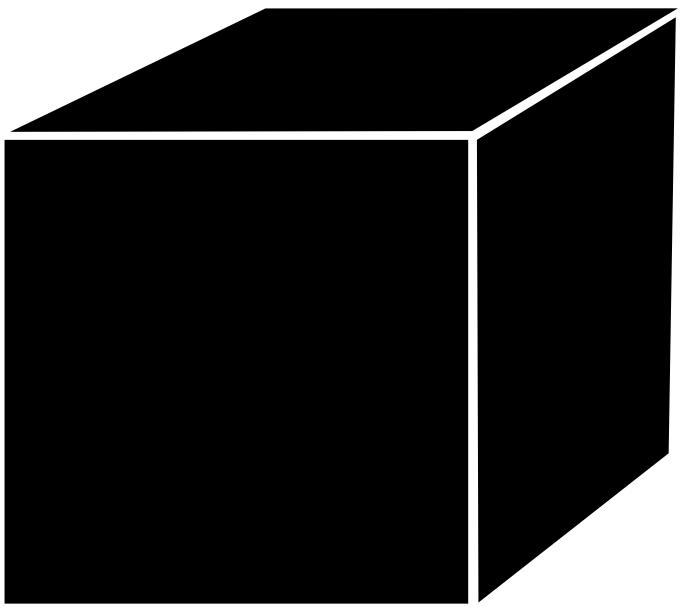


Don't fully trust the simulator?

- Nuisance parameters to model systematic uncertainties
- Methods learn dependence both on parameters of interest and nuisance parameters
- Profile / marginalize over nuisance parameters
- Alternatively: Robustness to nuisance with adversarial training

[G. Louppe, M. Kagan, K. Cranmer 1611.01046]

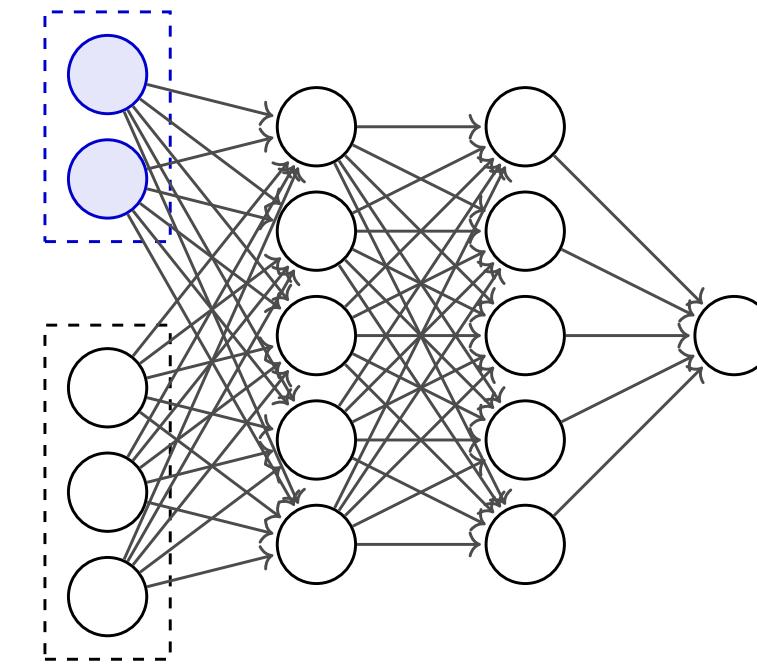
Improving robustness to mismodelling



Don't fully trust the simulator?

- Nuisance parameters to model systematic uncertainties
- Methods learn dependence both on parameters of interest and nuisance parameters
- Profile / marginalize over nuisance parameters
- Alternatively: Robustness to nuisance with adversarial training

[G. Louppe, M. Kagan, K. Cranmer 1611.01046]



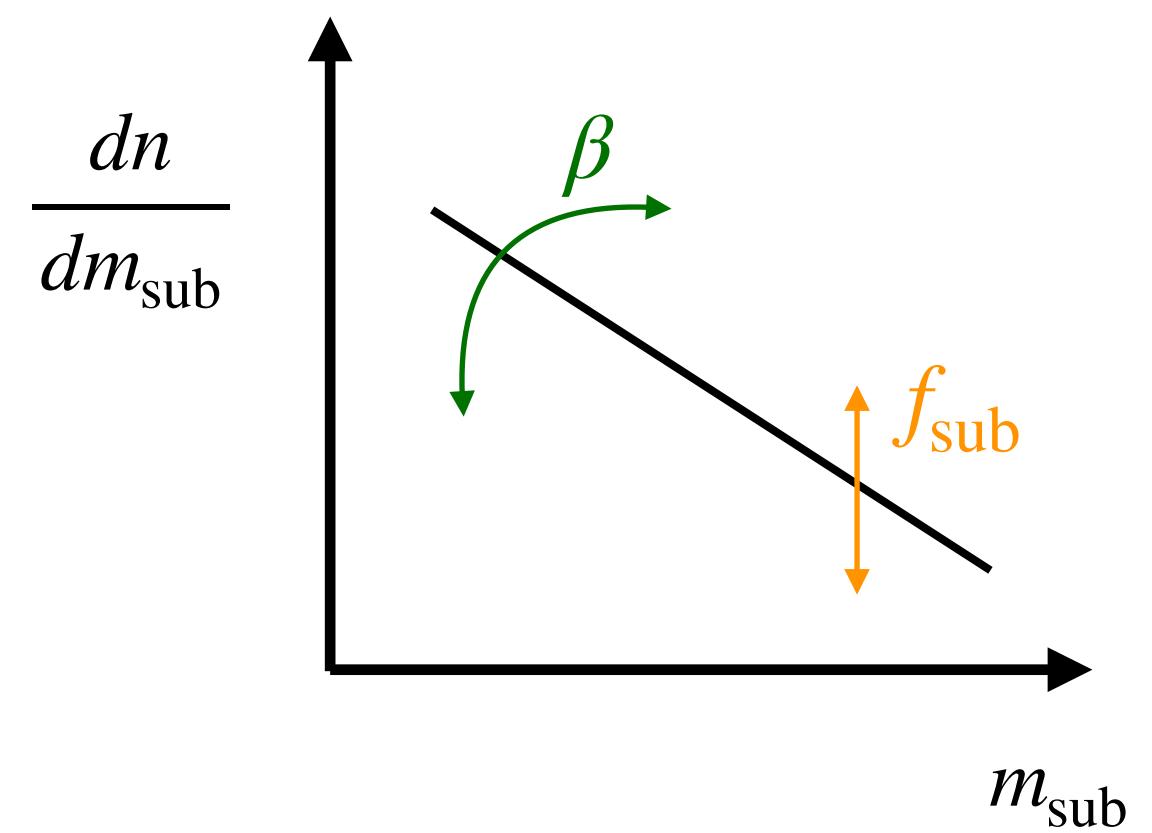
Don't blindly trust the neural network?

- Sanity checks: expectation values, "critic" tests
- Calibration / Neyman construction with toys
(badly trained network can lead to suboptimal limits, but not to wrong limits)

This method lets us infer Dark Matter substructure
from strong lensing observations.

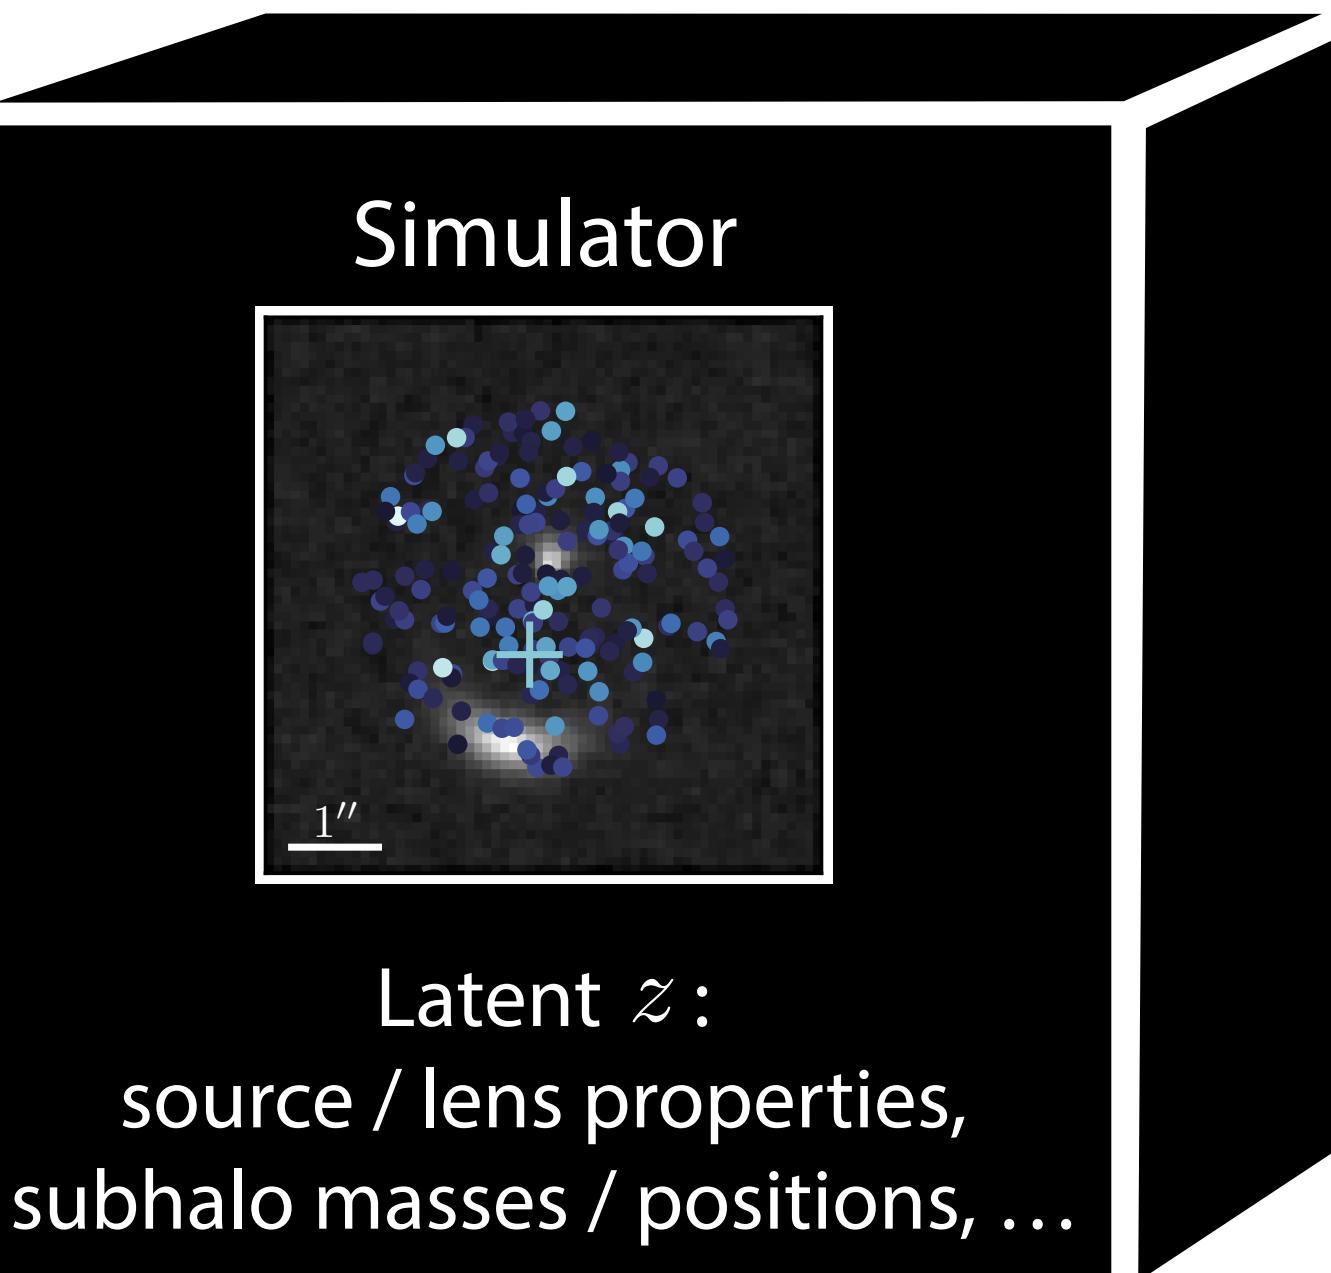
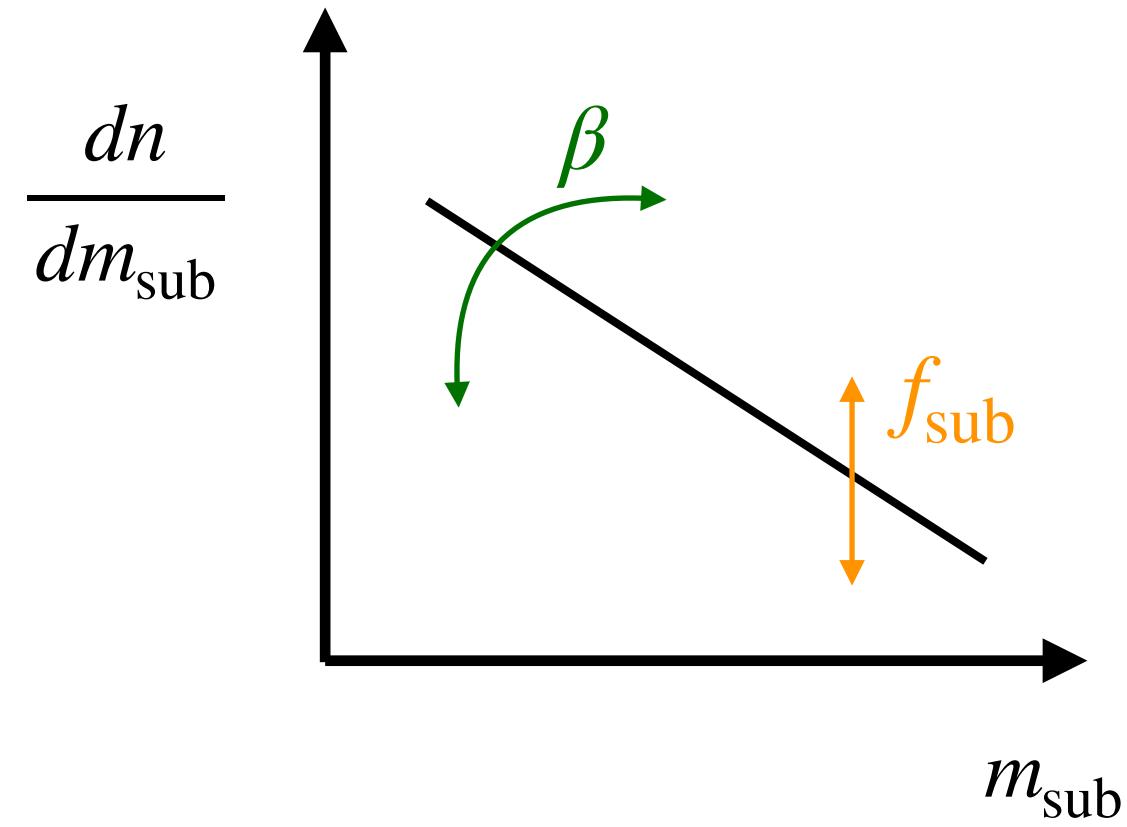
Overview

2 parameters $\theta = (\beta, f_{\text{sub}})$



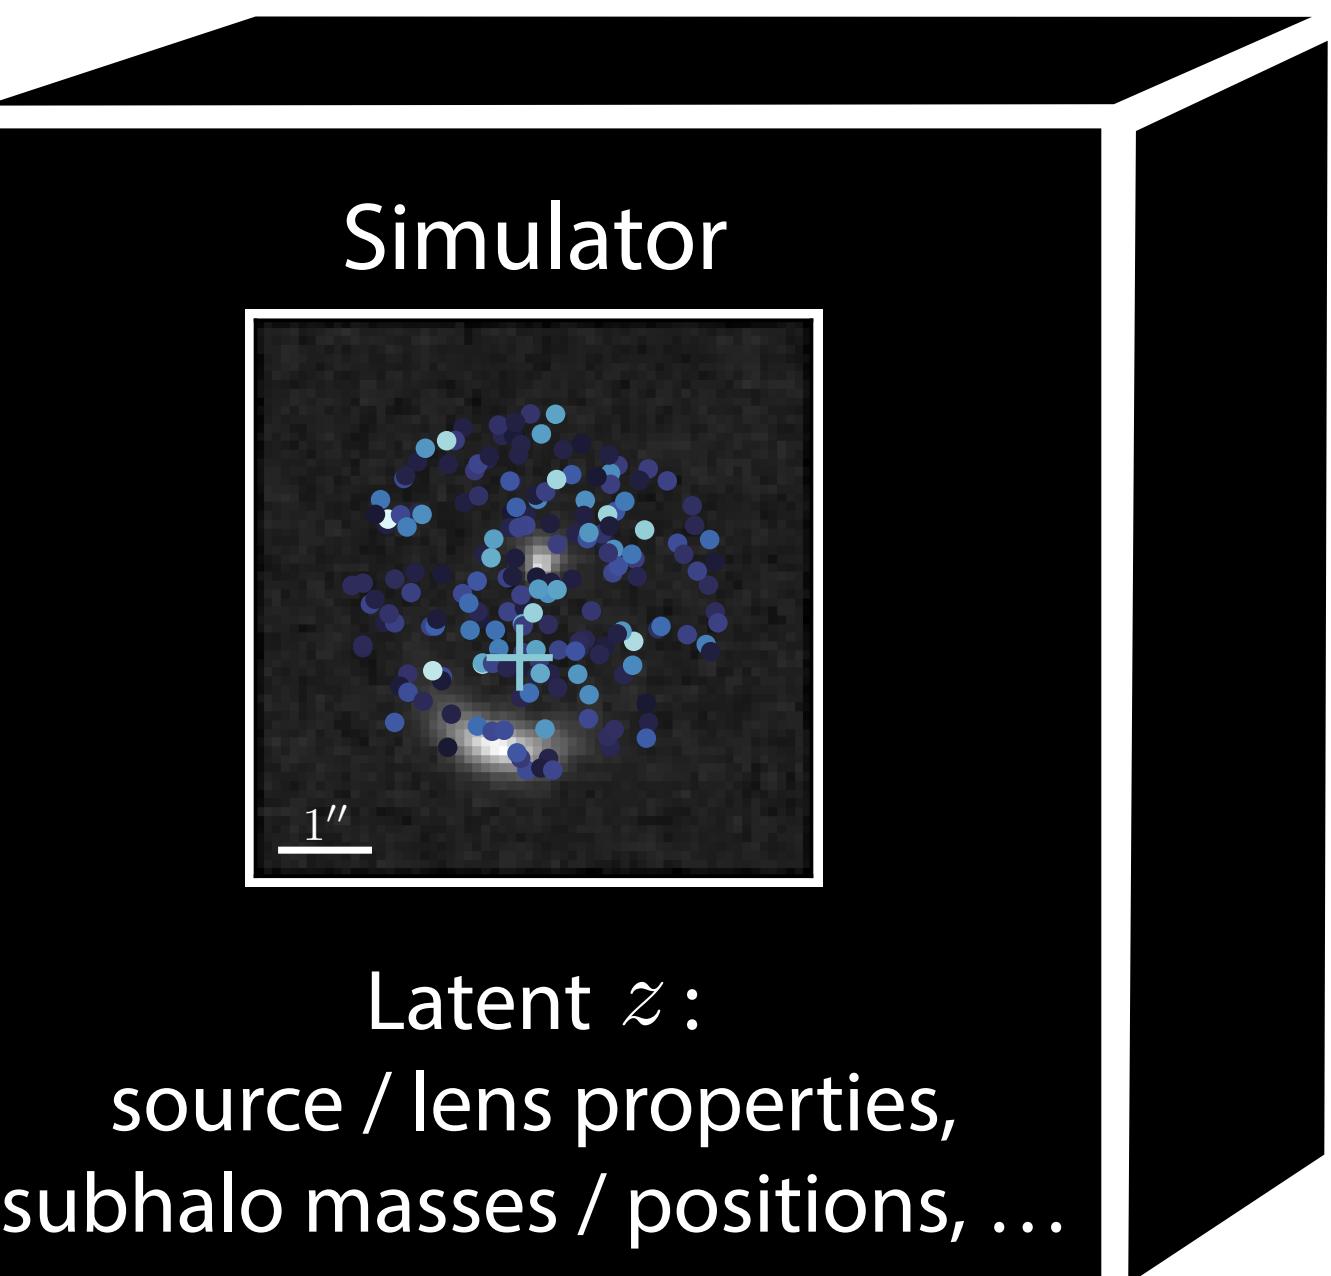
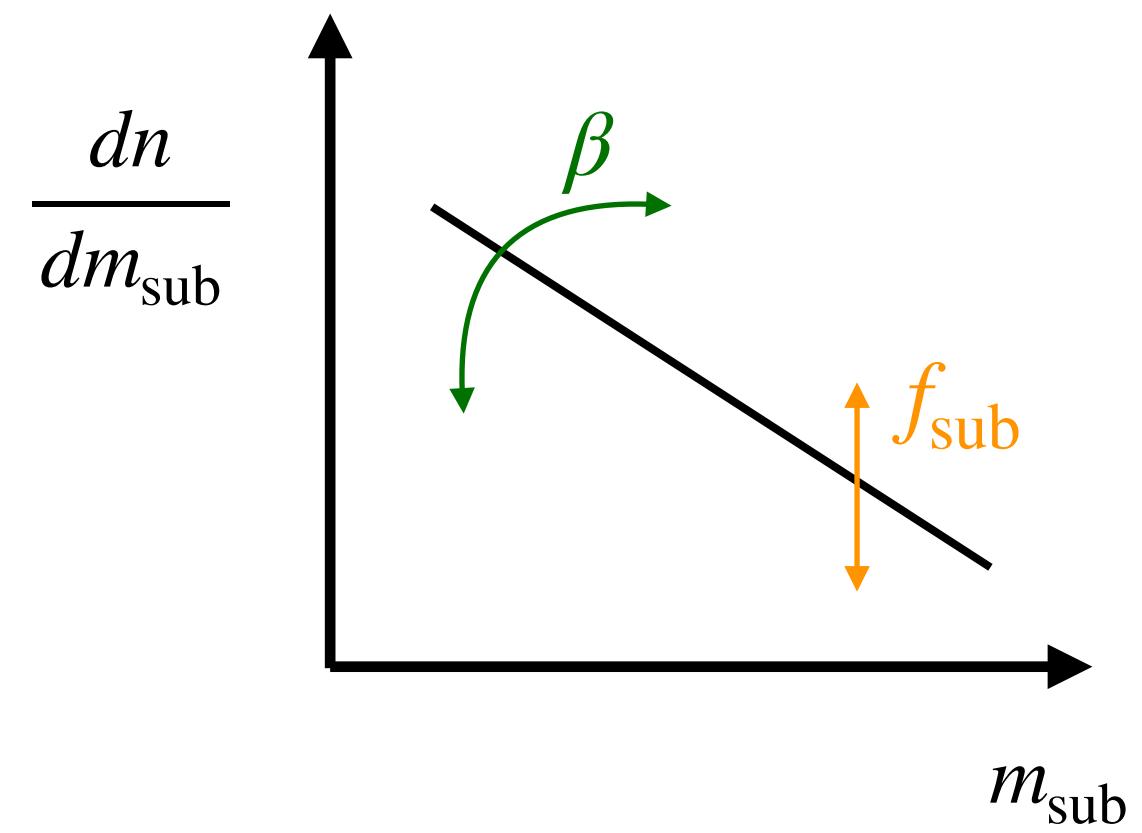
Overview

2 parameters $\theta = (\beta, f_{\text{sub}})$

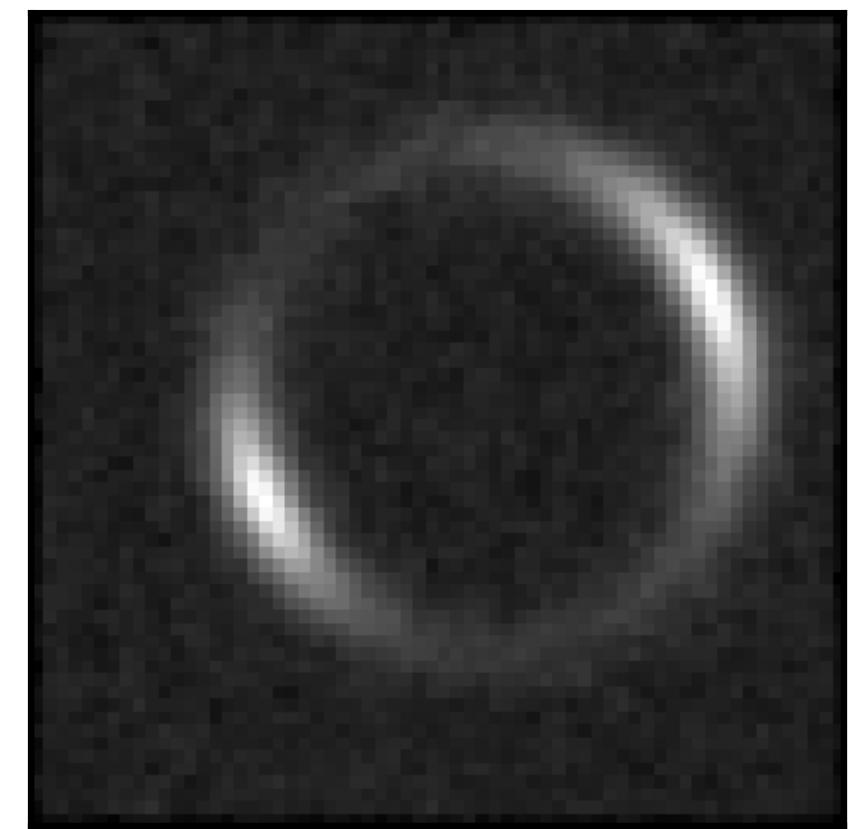


Overview

2 parameters $\theta = (\beta, f_{\text{sub}})$

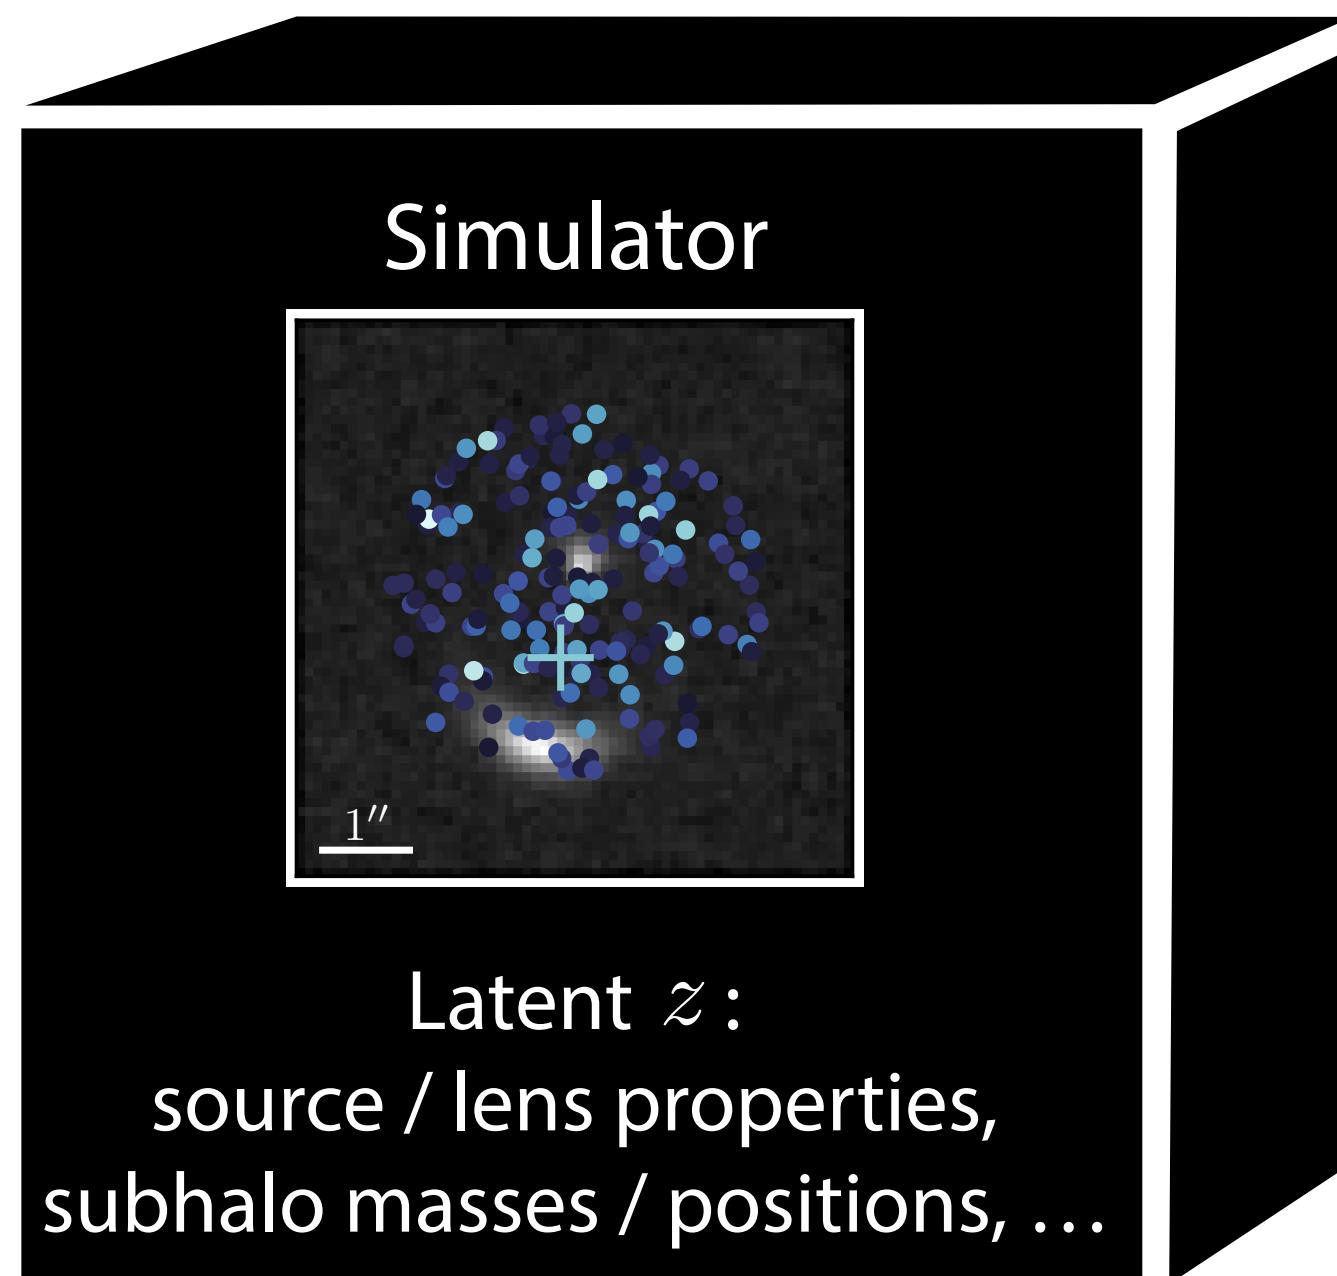
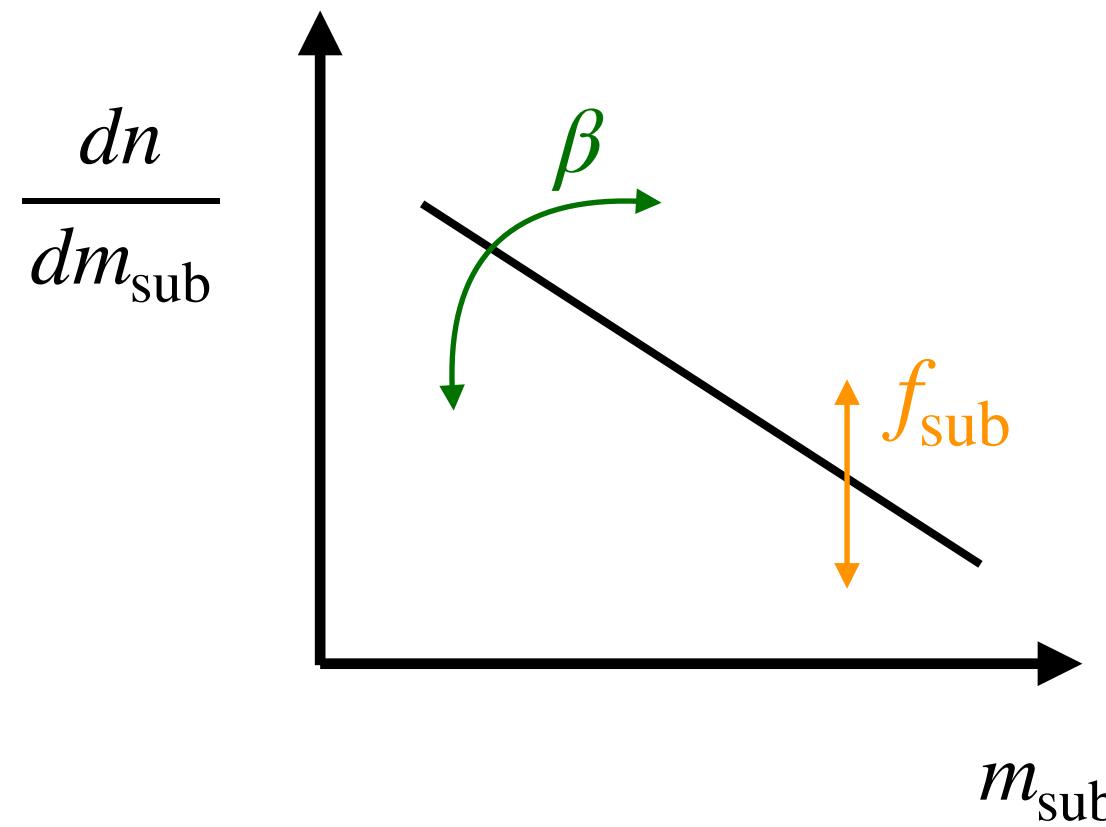


64² observables x

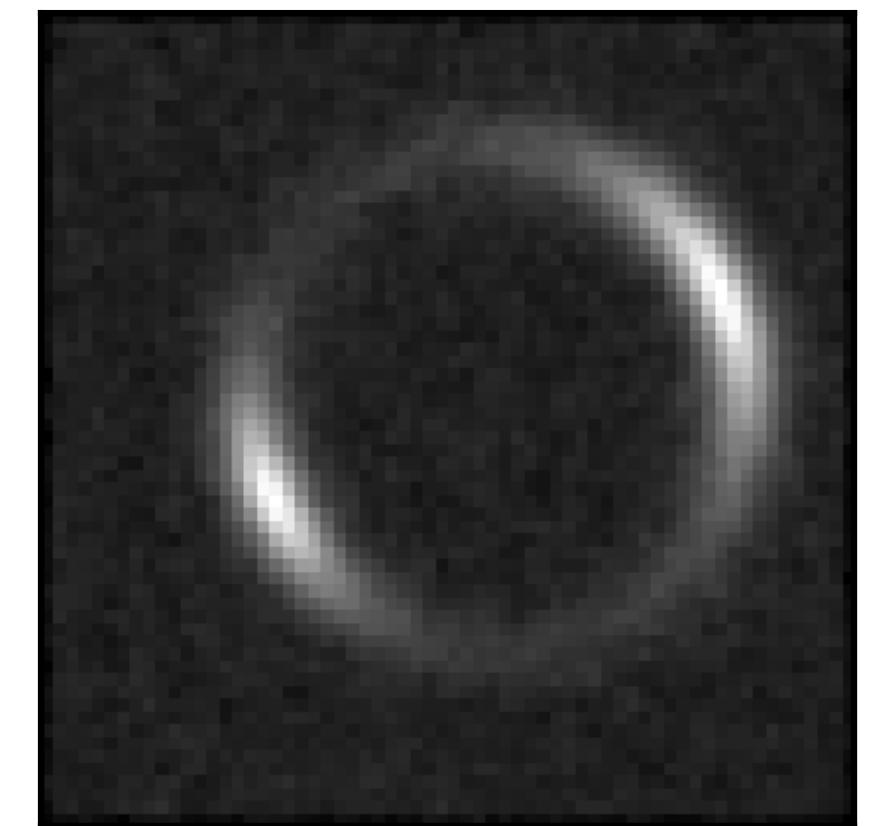


Overview

2 parameters $\theta = (\beta, f_{\text{sub}})$



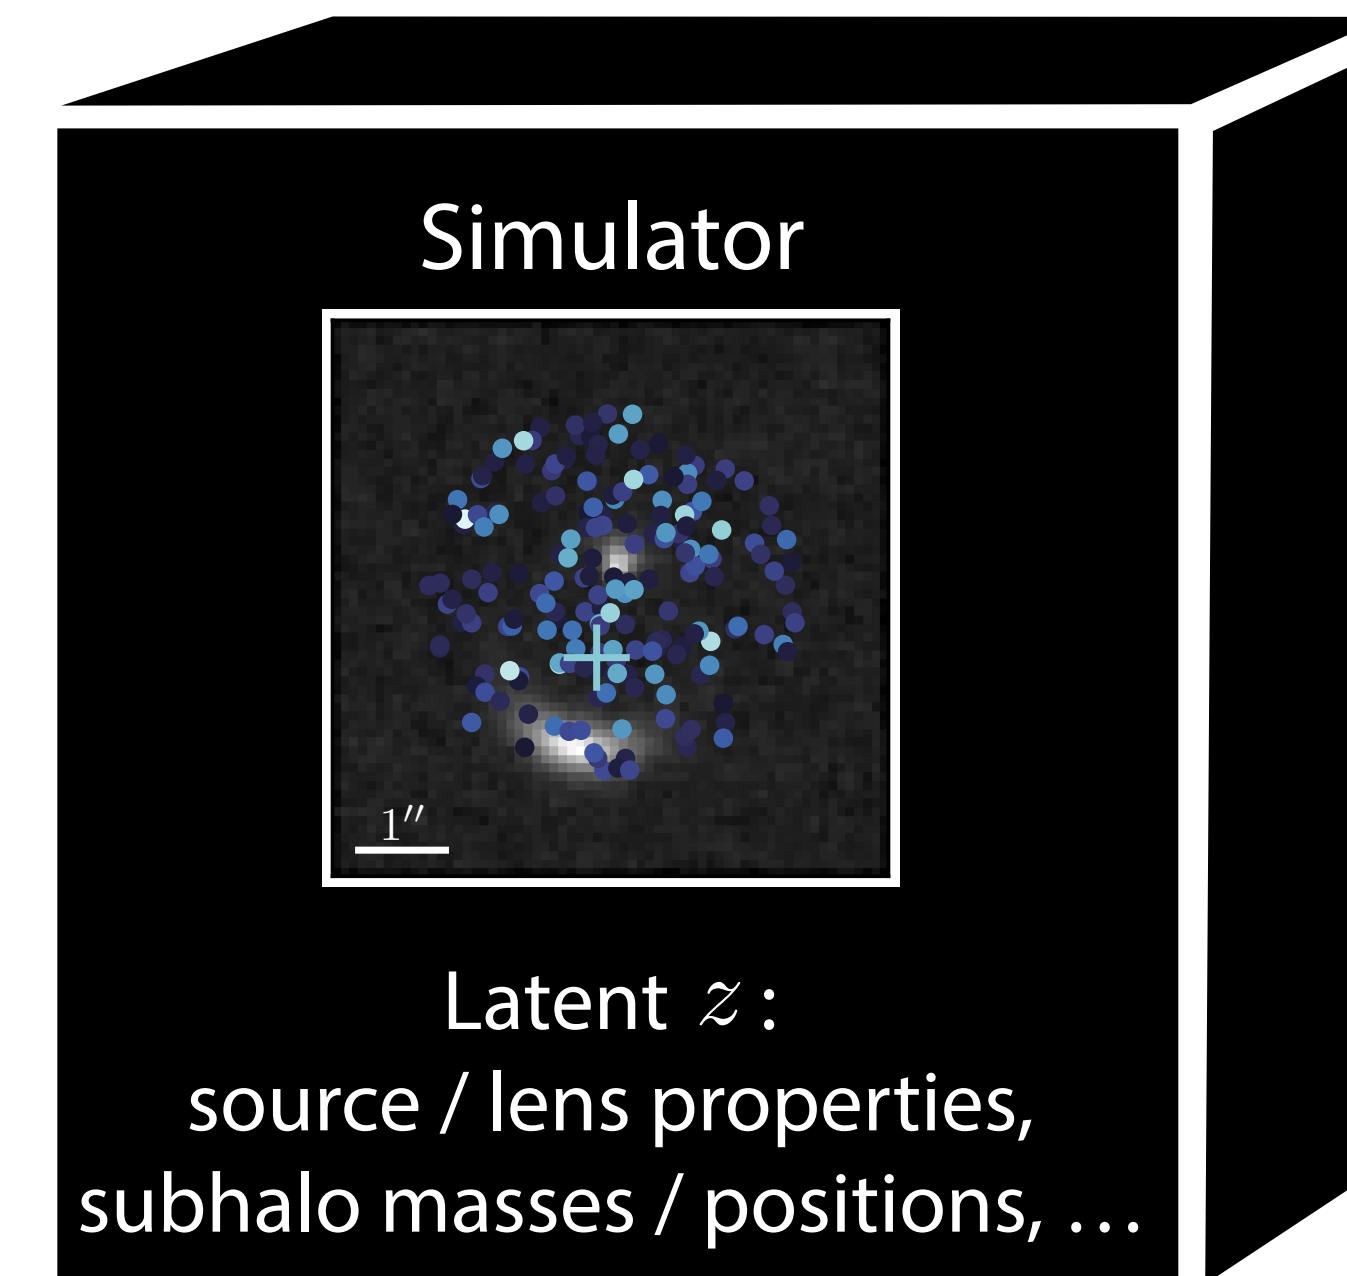
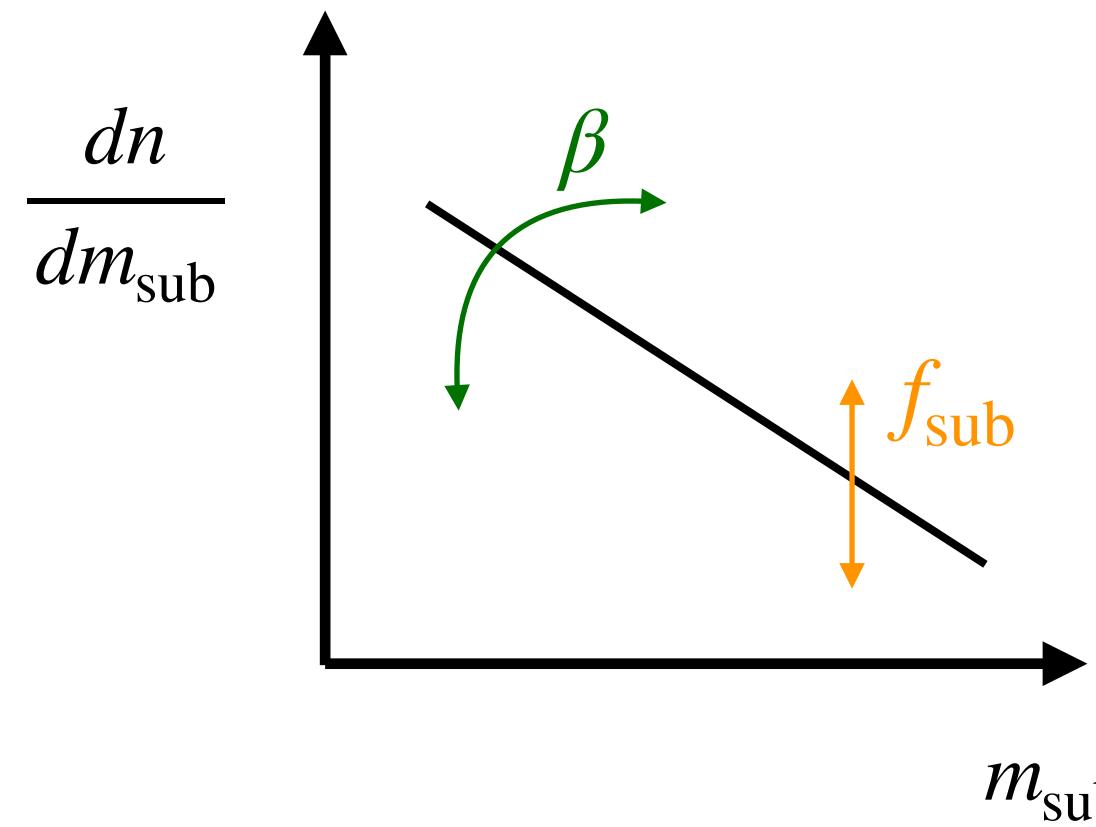
64² observables x



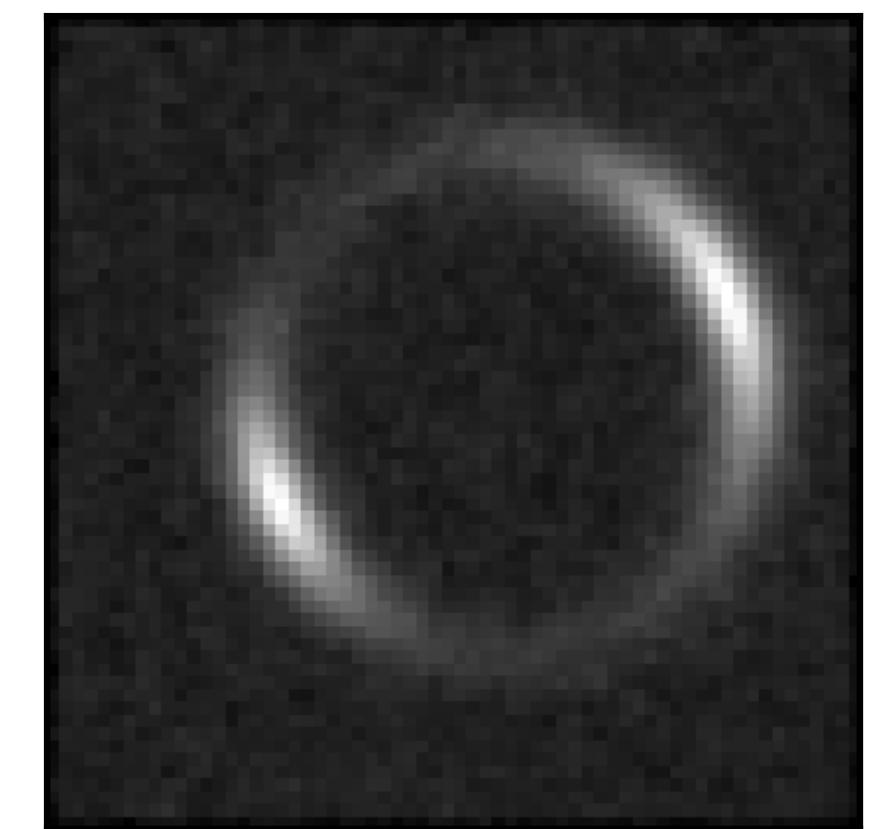
Prediction: We construct a simulator that can sample $x \sim p(x|\theta)$

Overview

2 parameters $\theta = (\beta, f_{\text{sub}})$



64² observables x

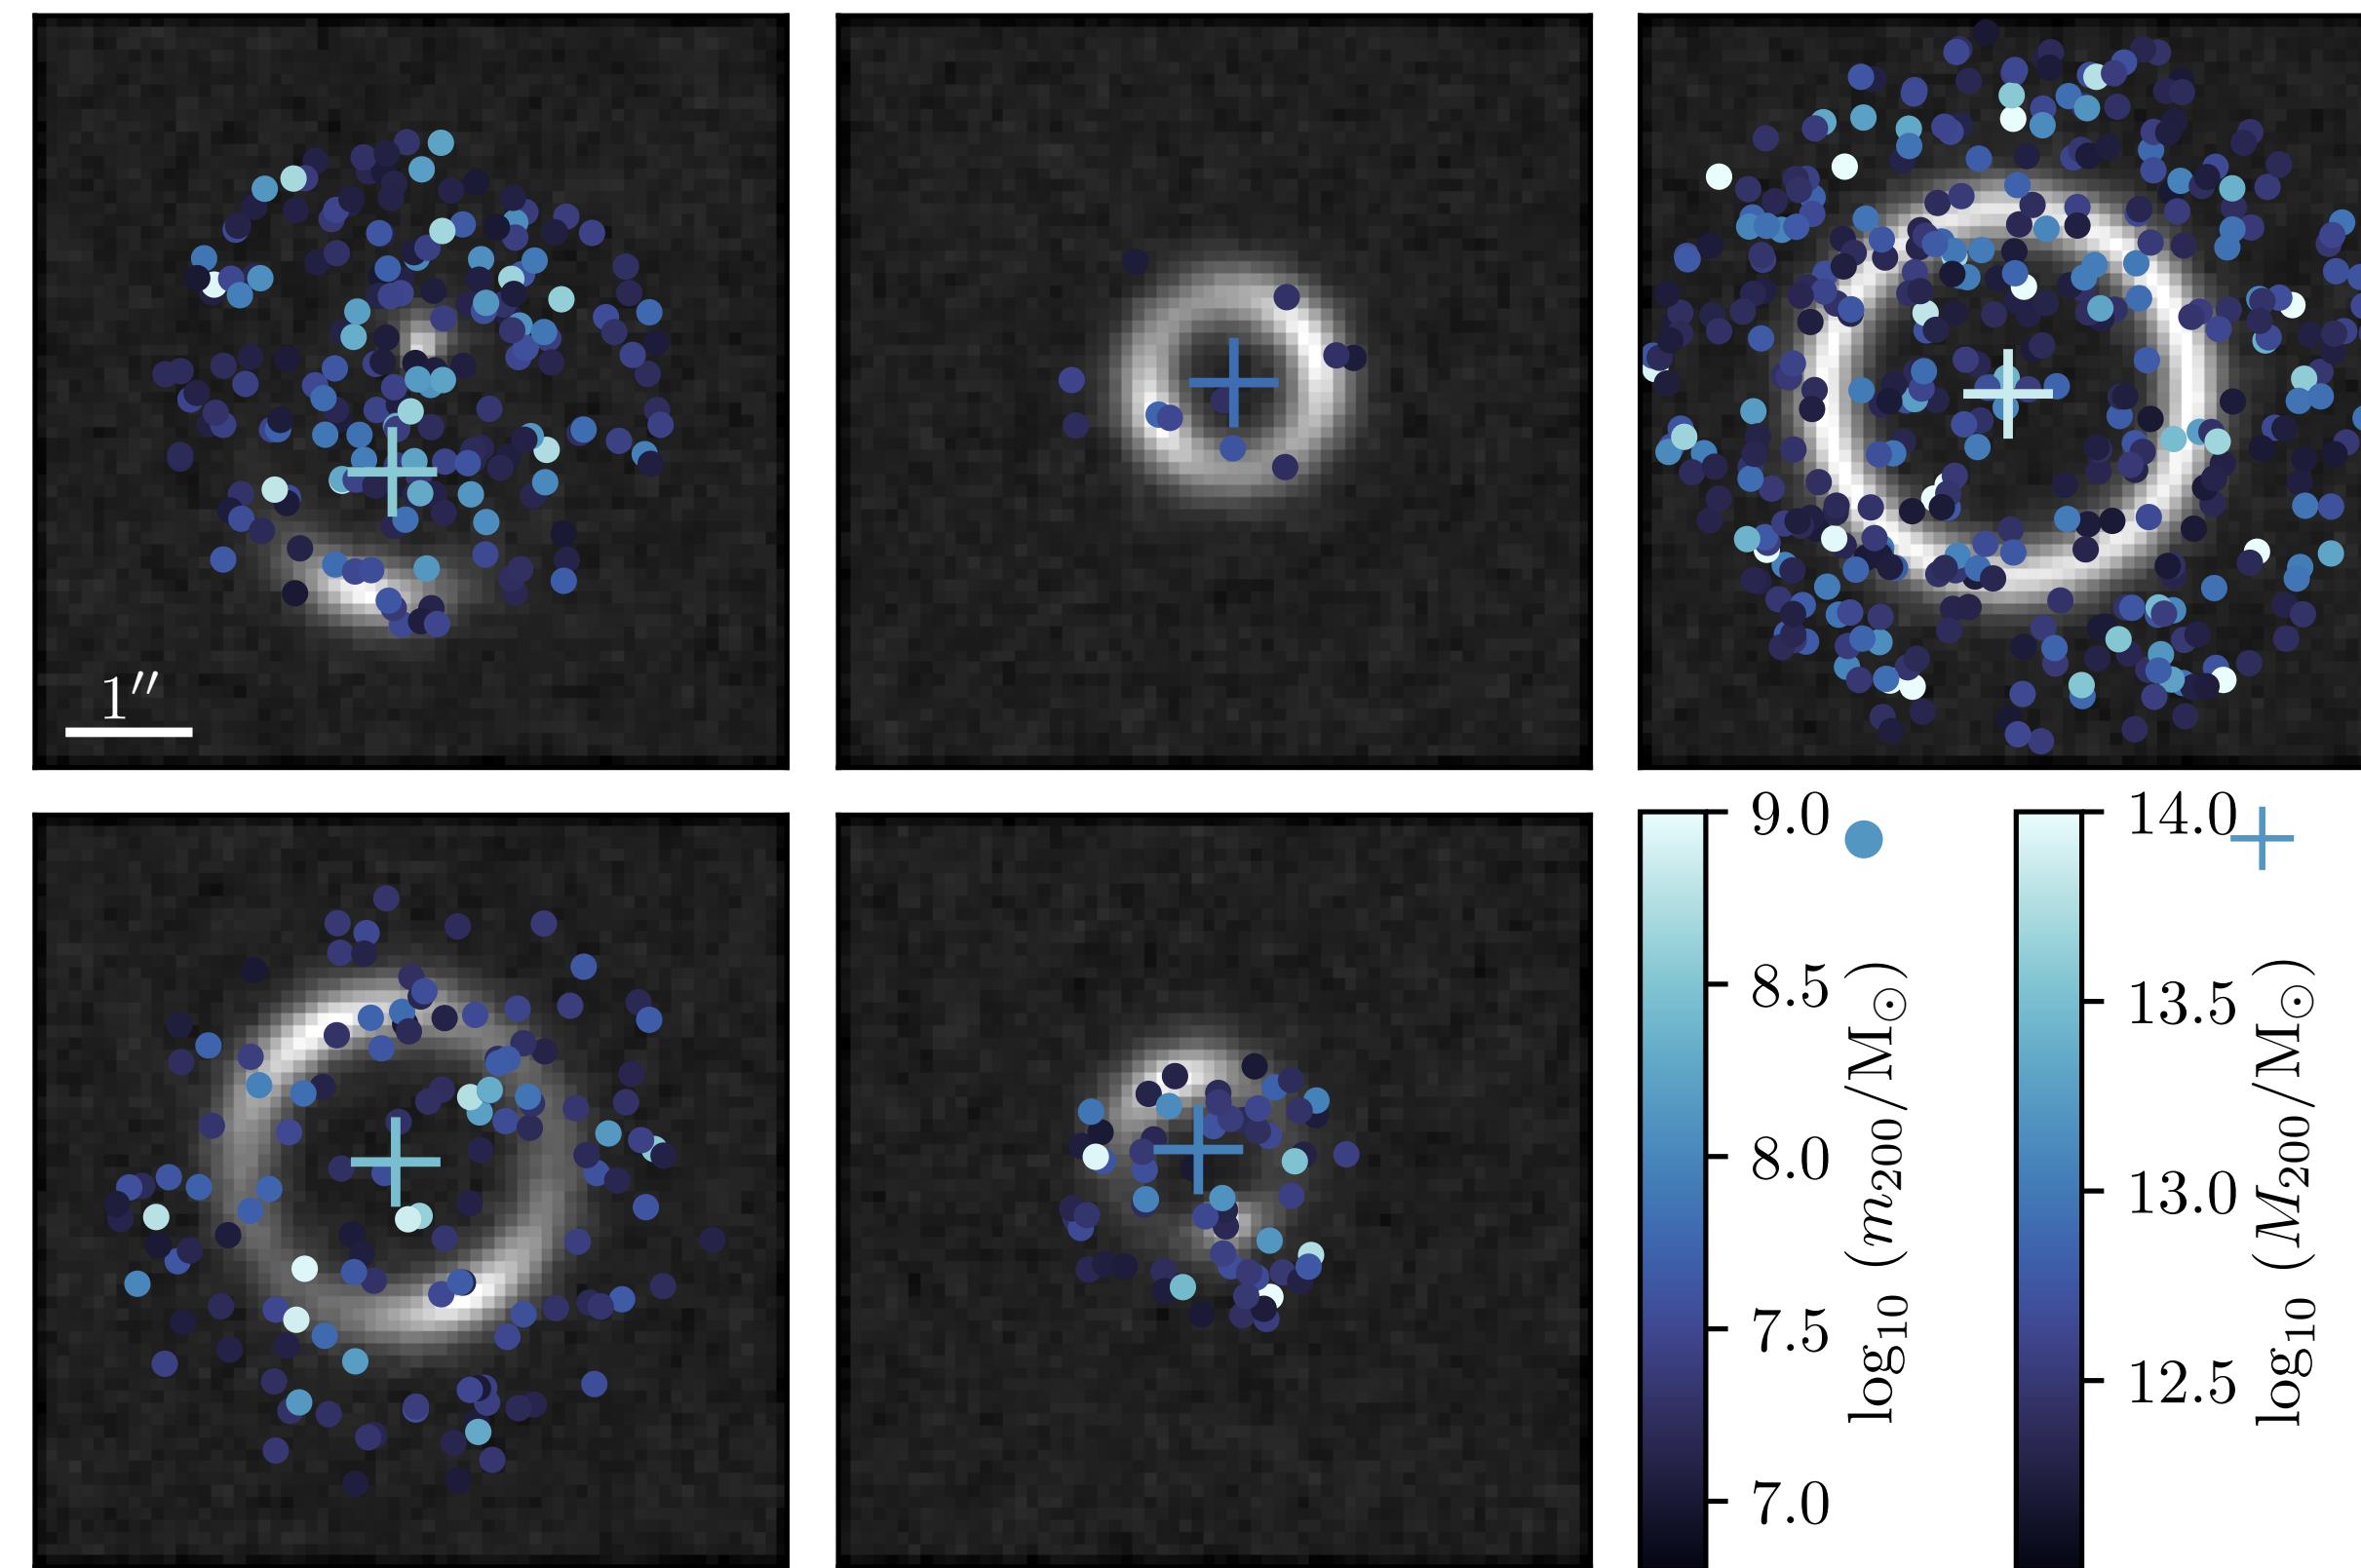


Prediction: We construct a simulator that can sample $x \sim p(x|\theta)$

Inference: We train neural likelihood ratio estimators $\hat{r}(x|\theta)$

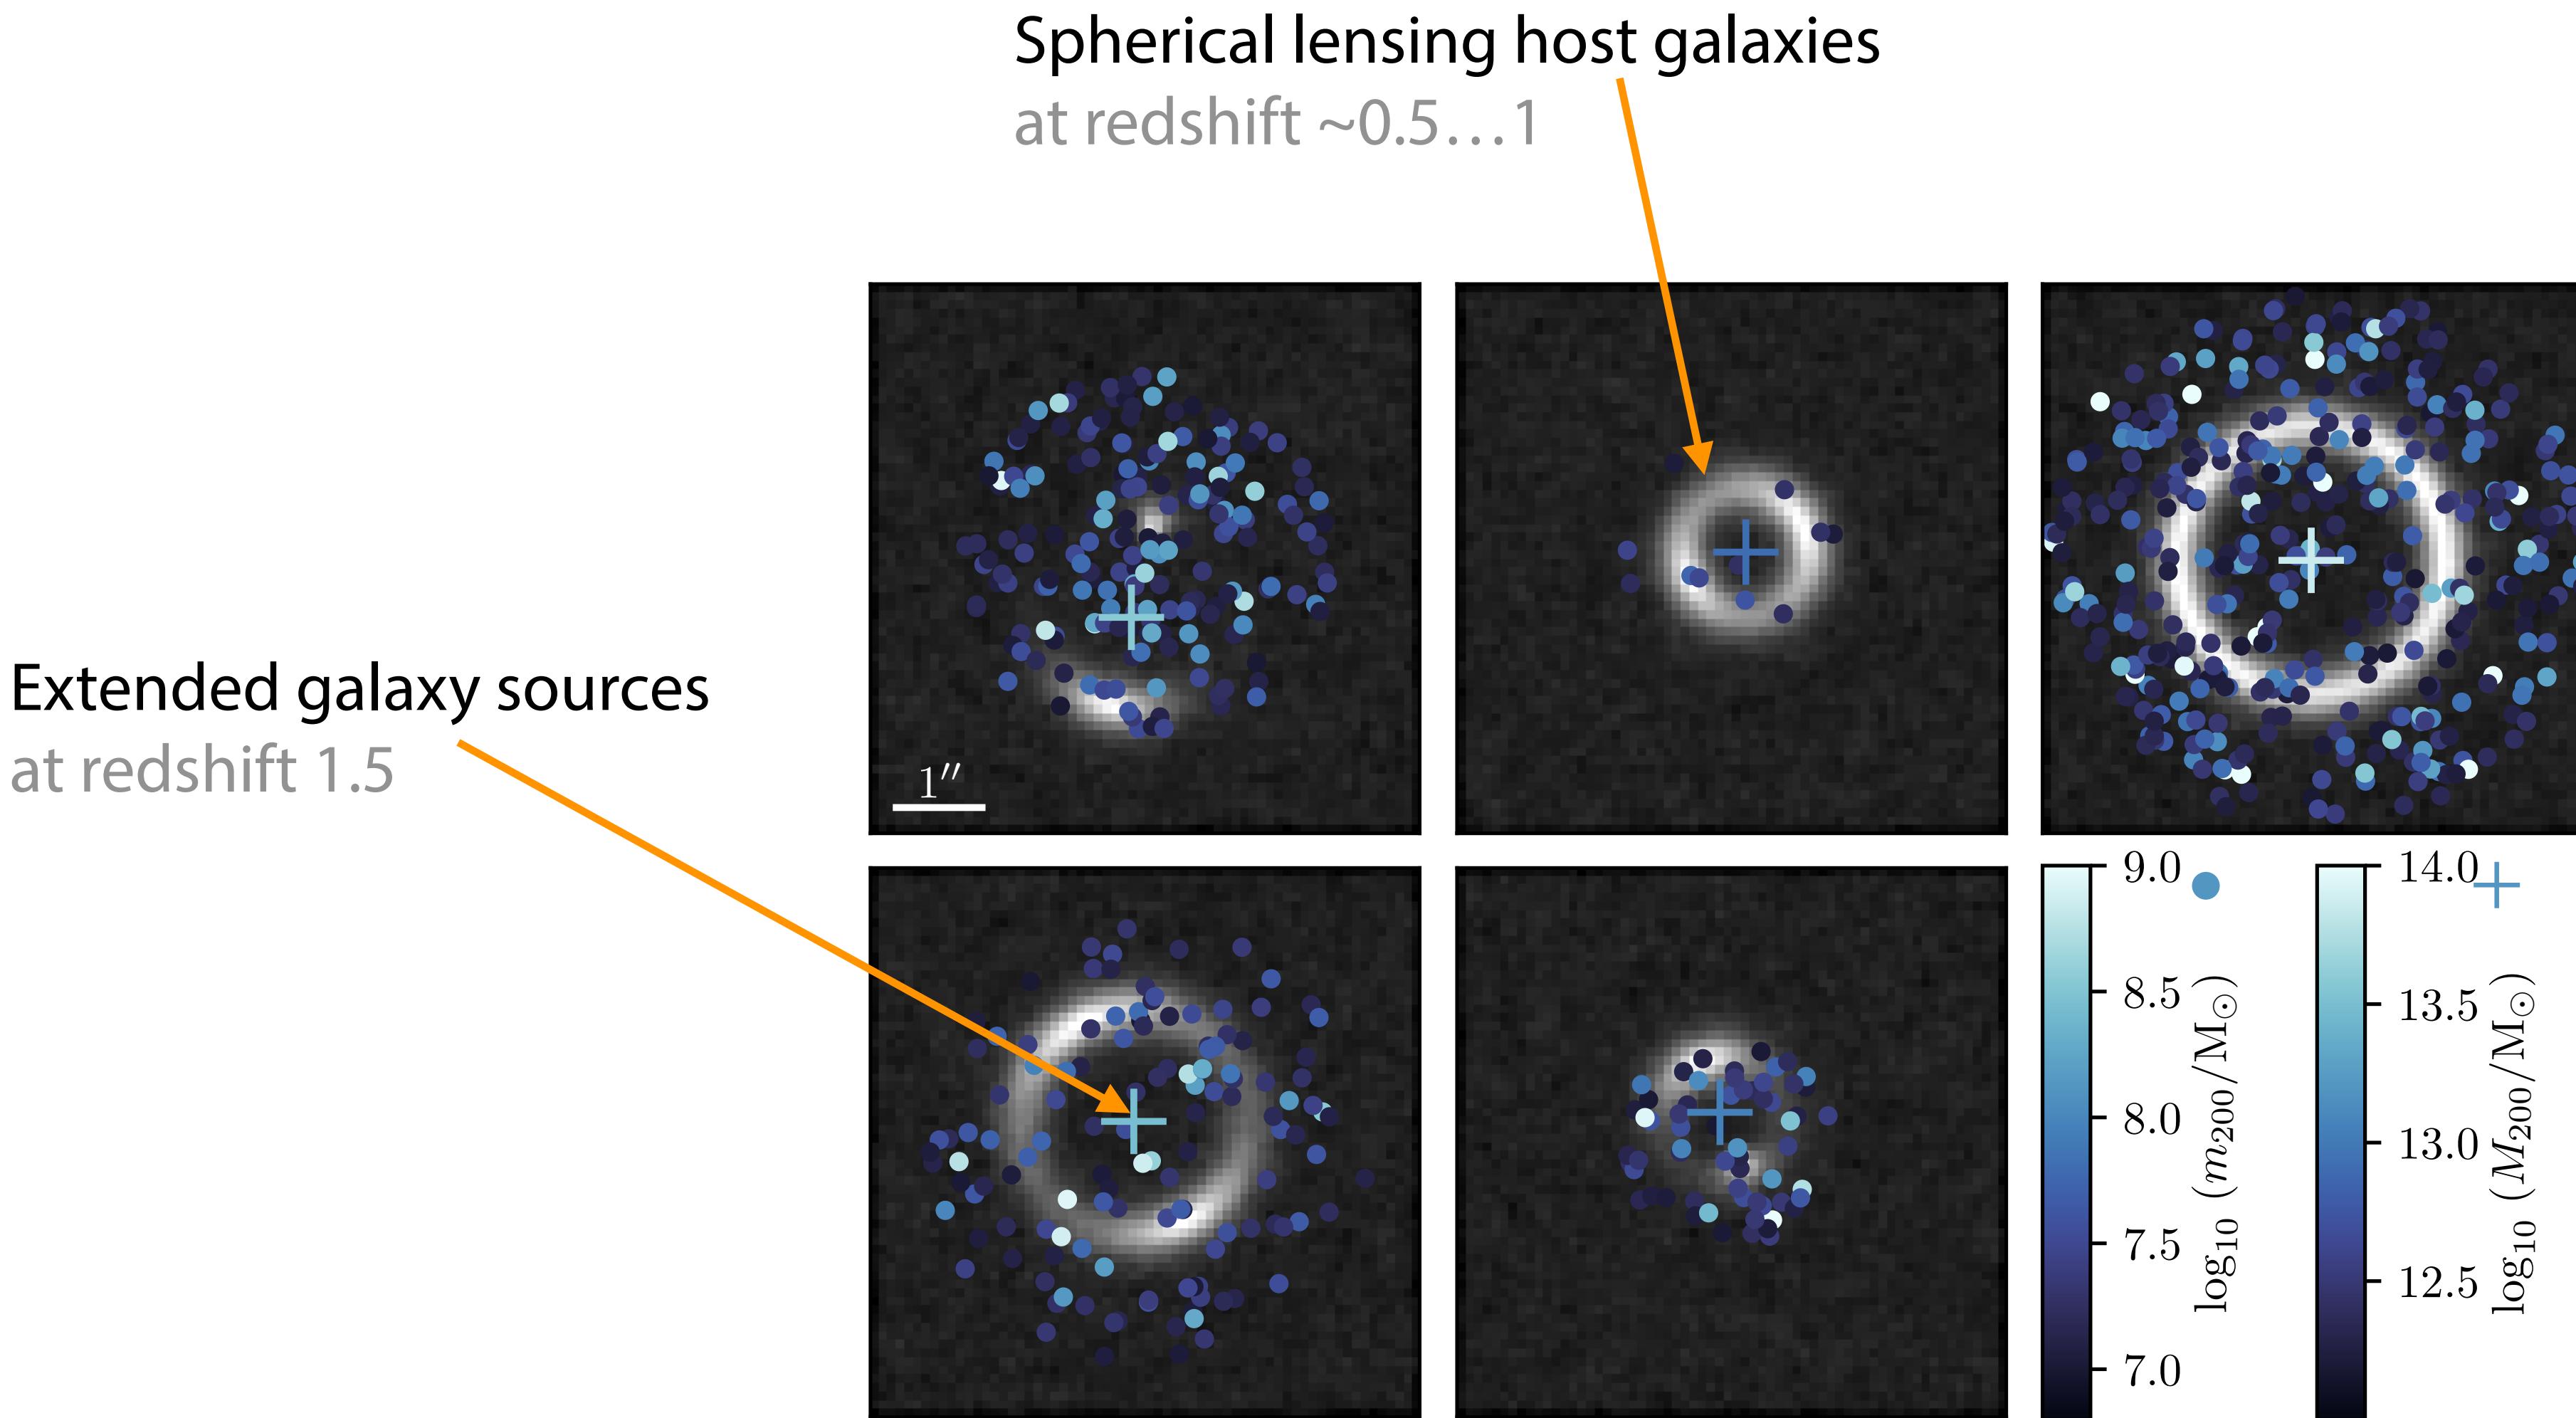
Proof-of-principle simulator

[following T. Collett 1507.02657]



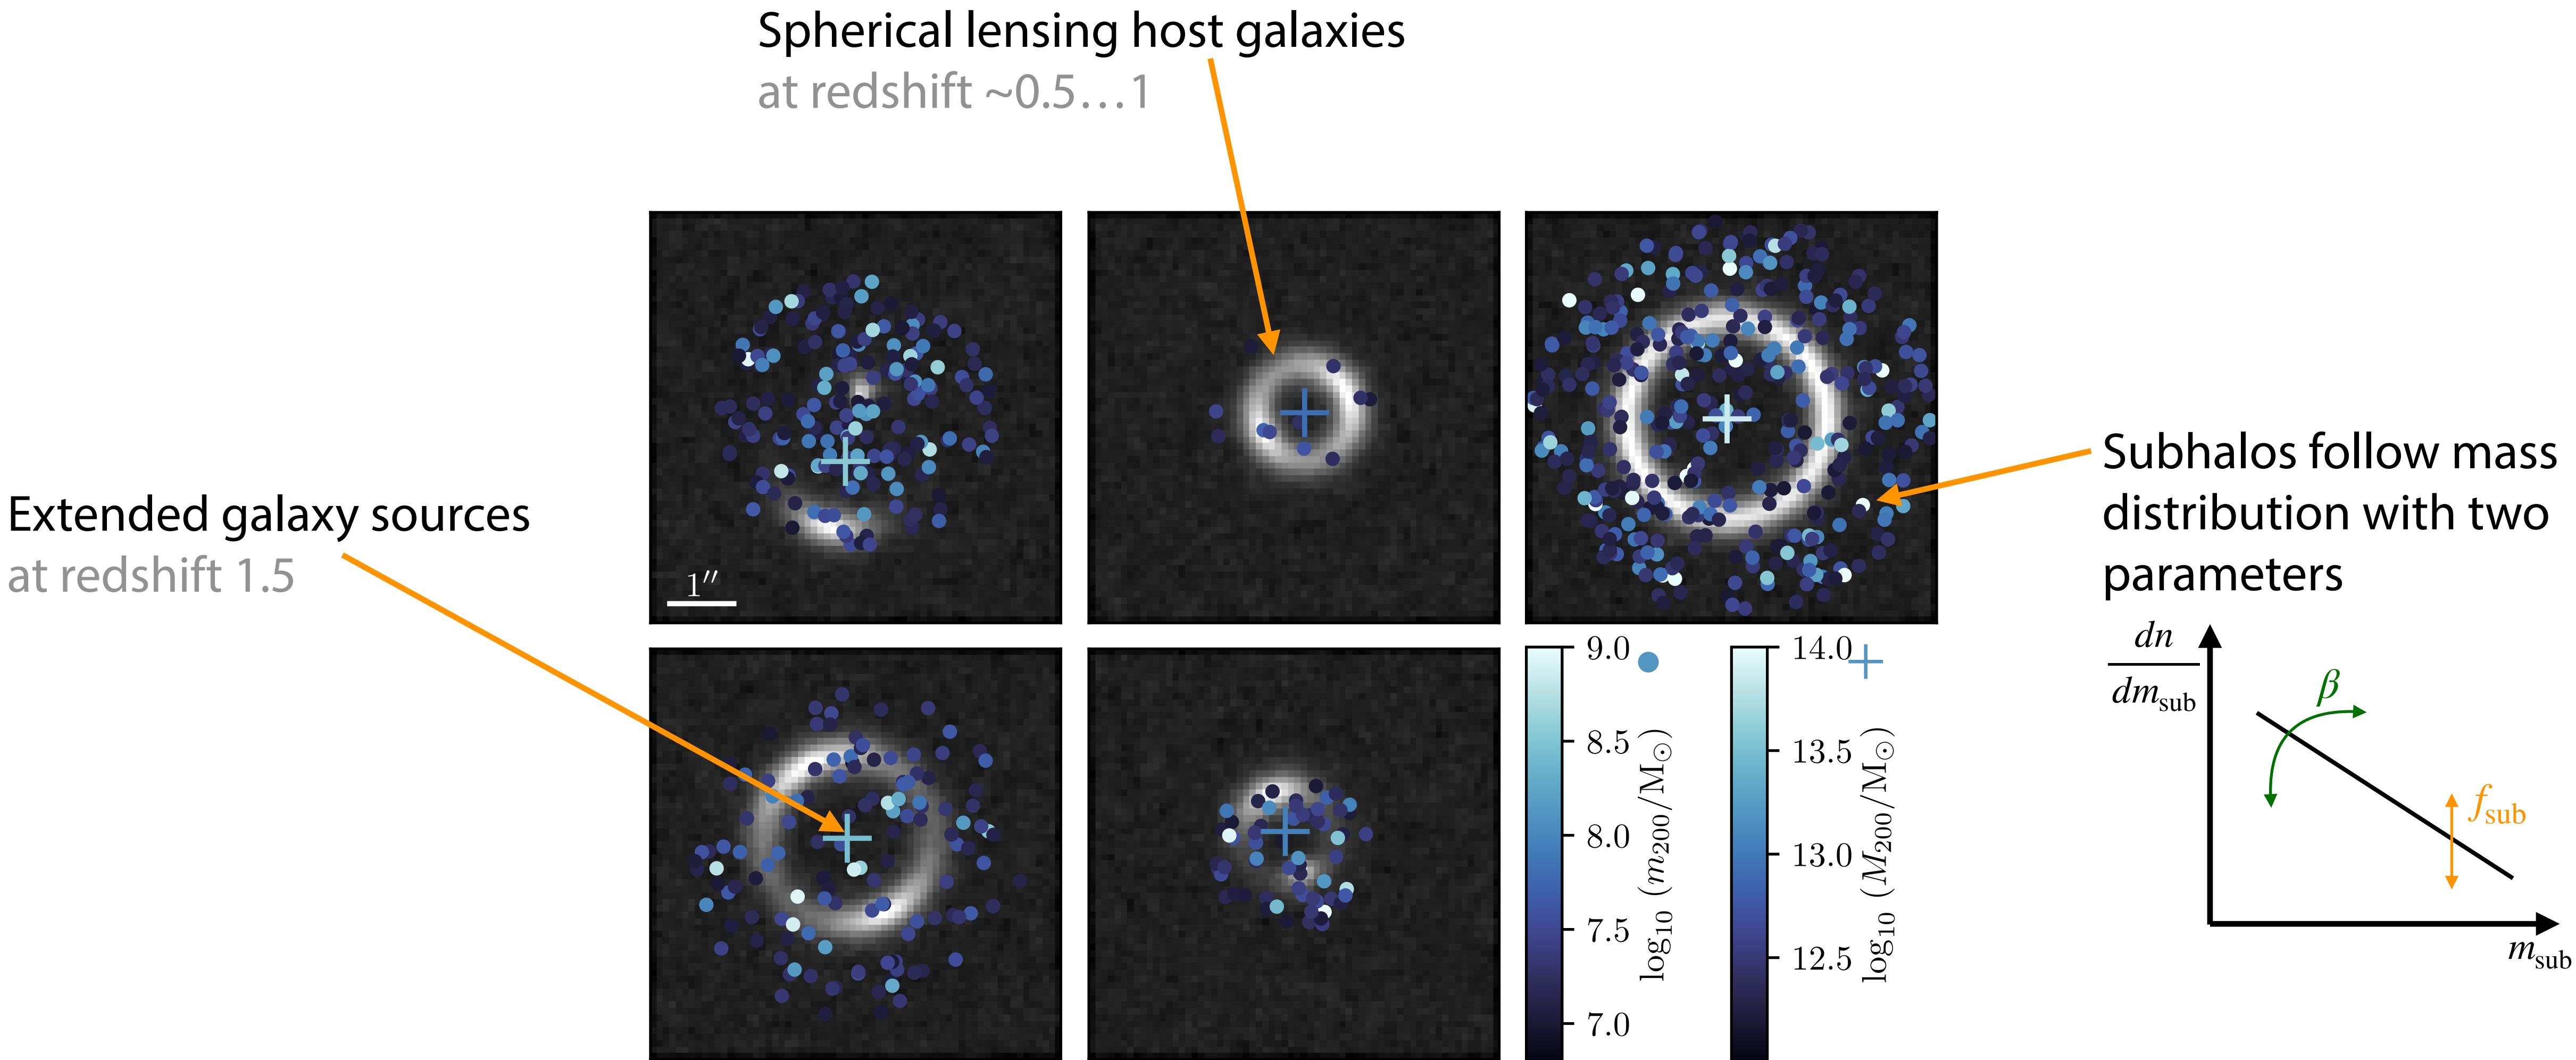
Proof-of-principle simulator

[following T. Collett 1507.02657]



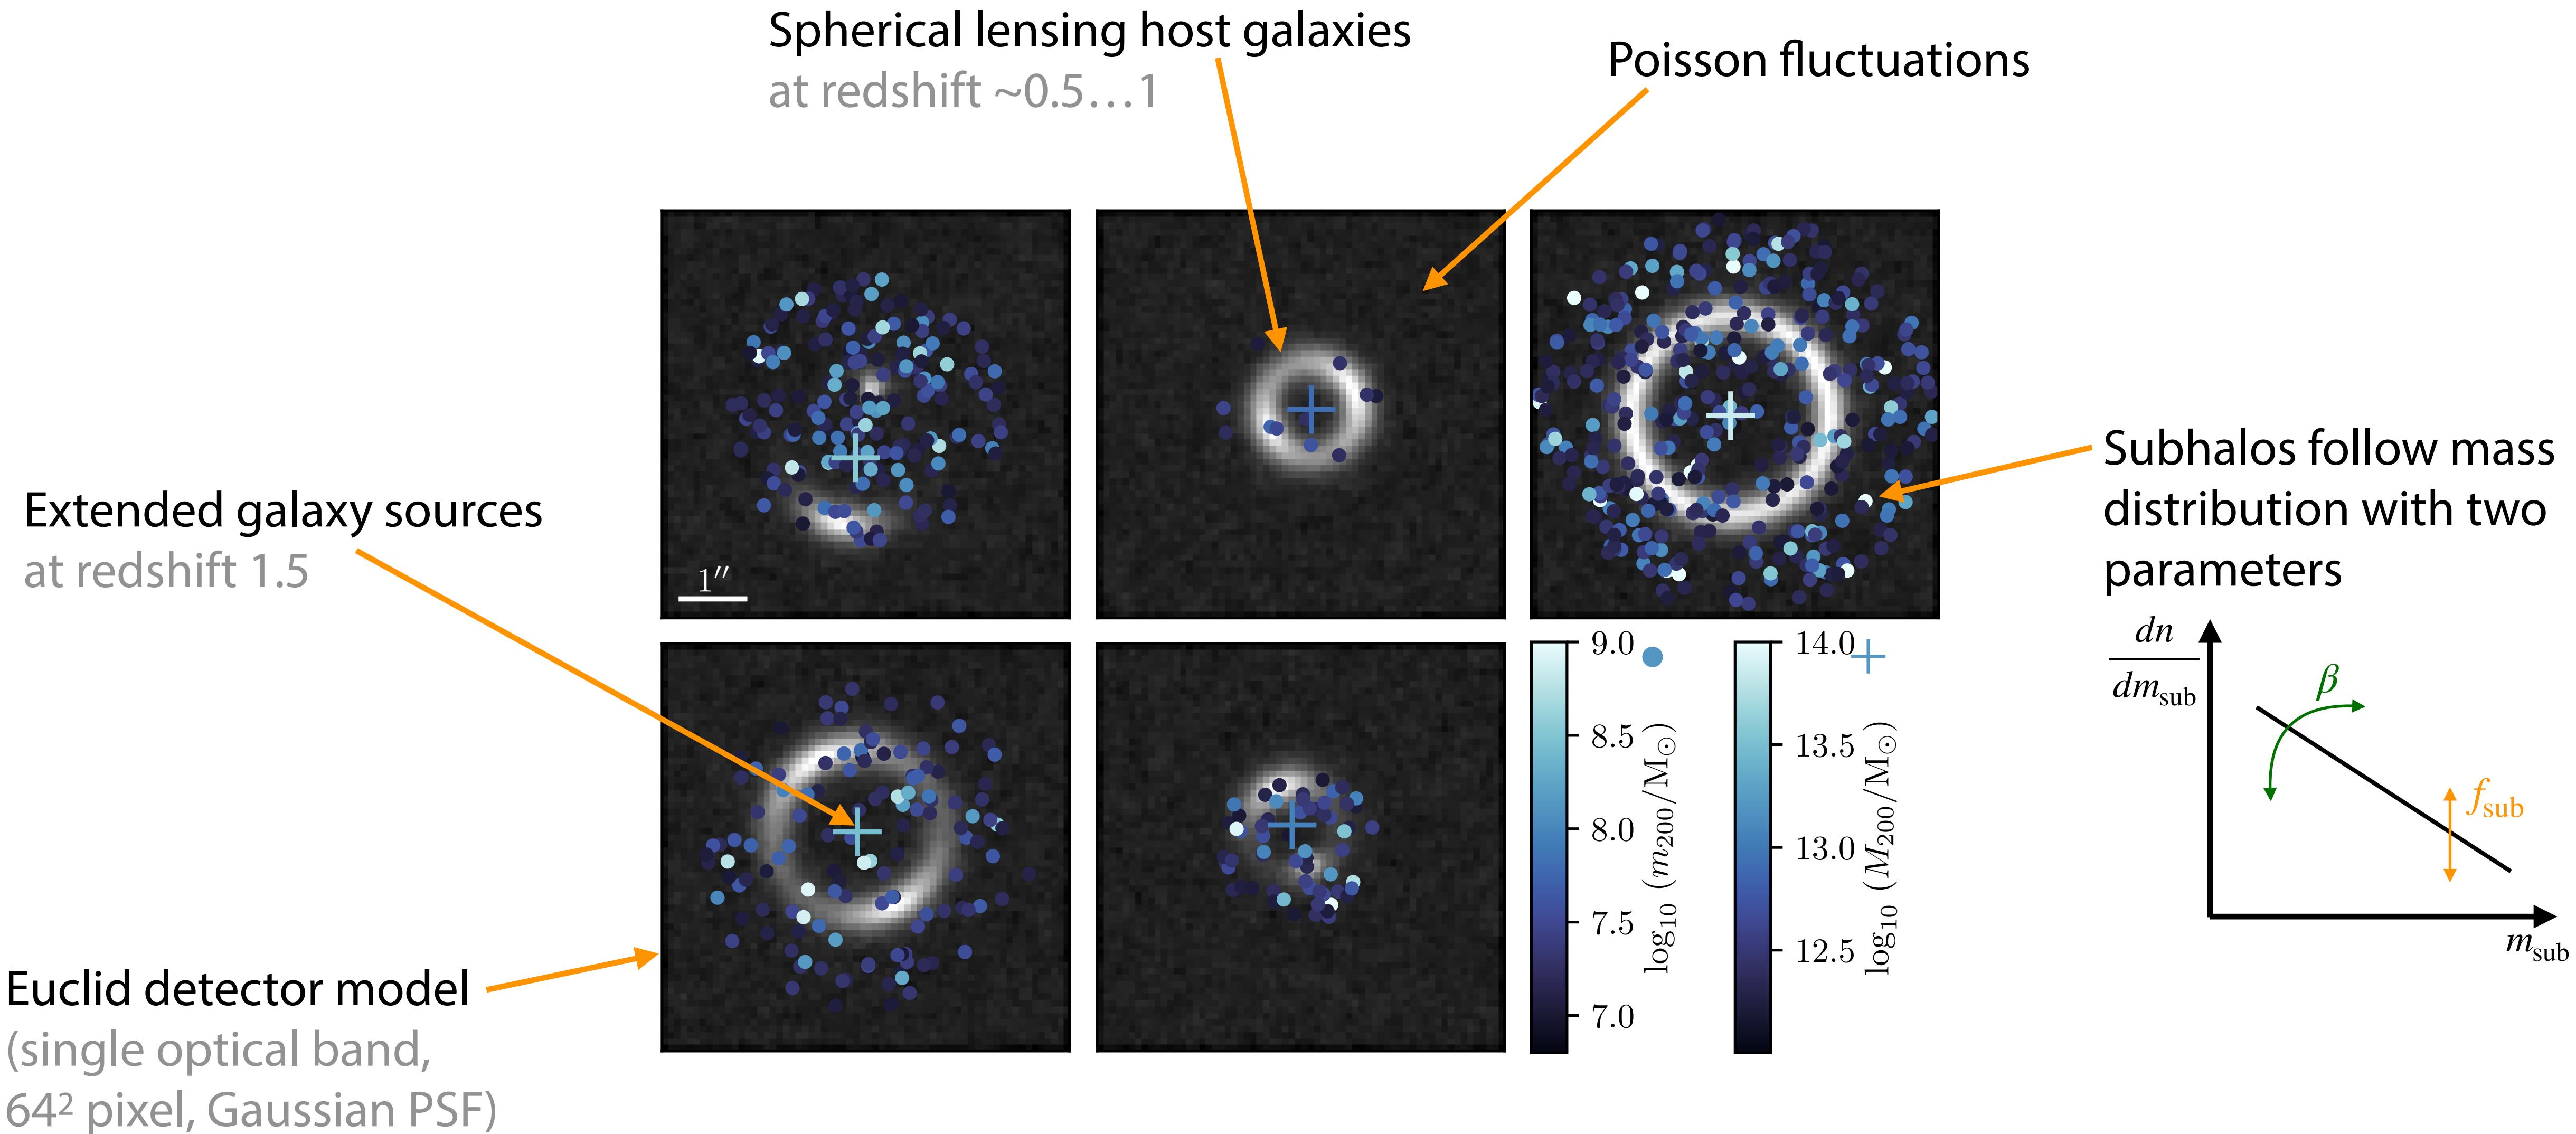
Proof-of-principle simulator

[following T. Collett 1507.02657]

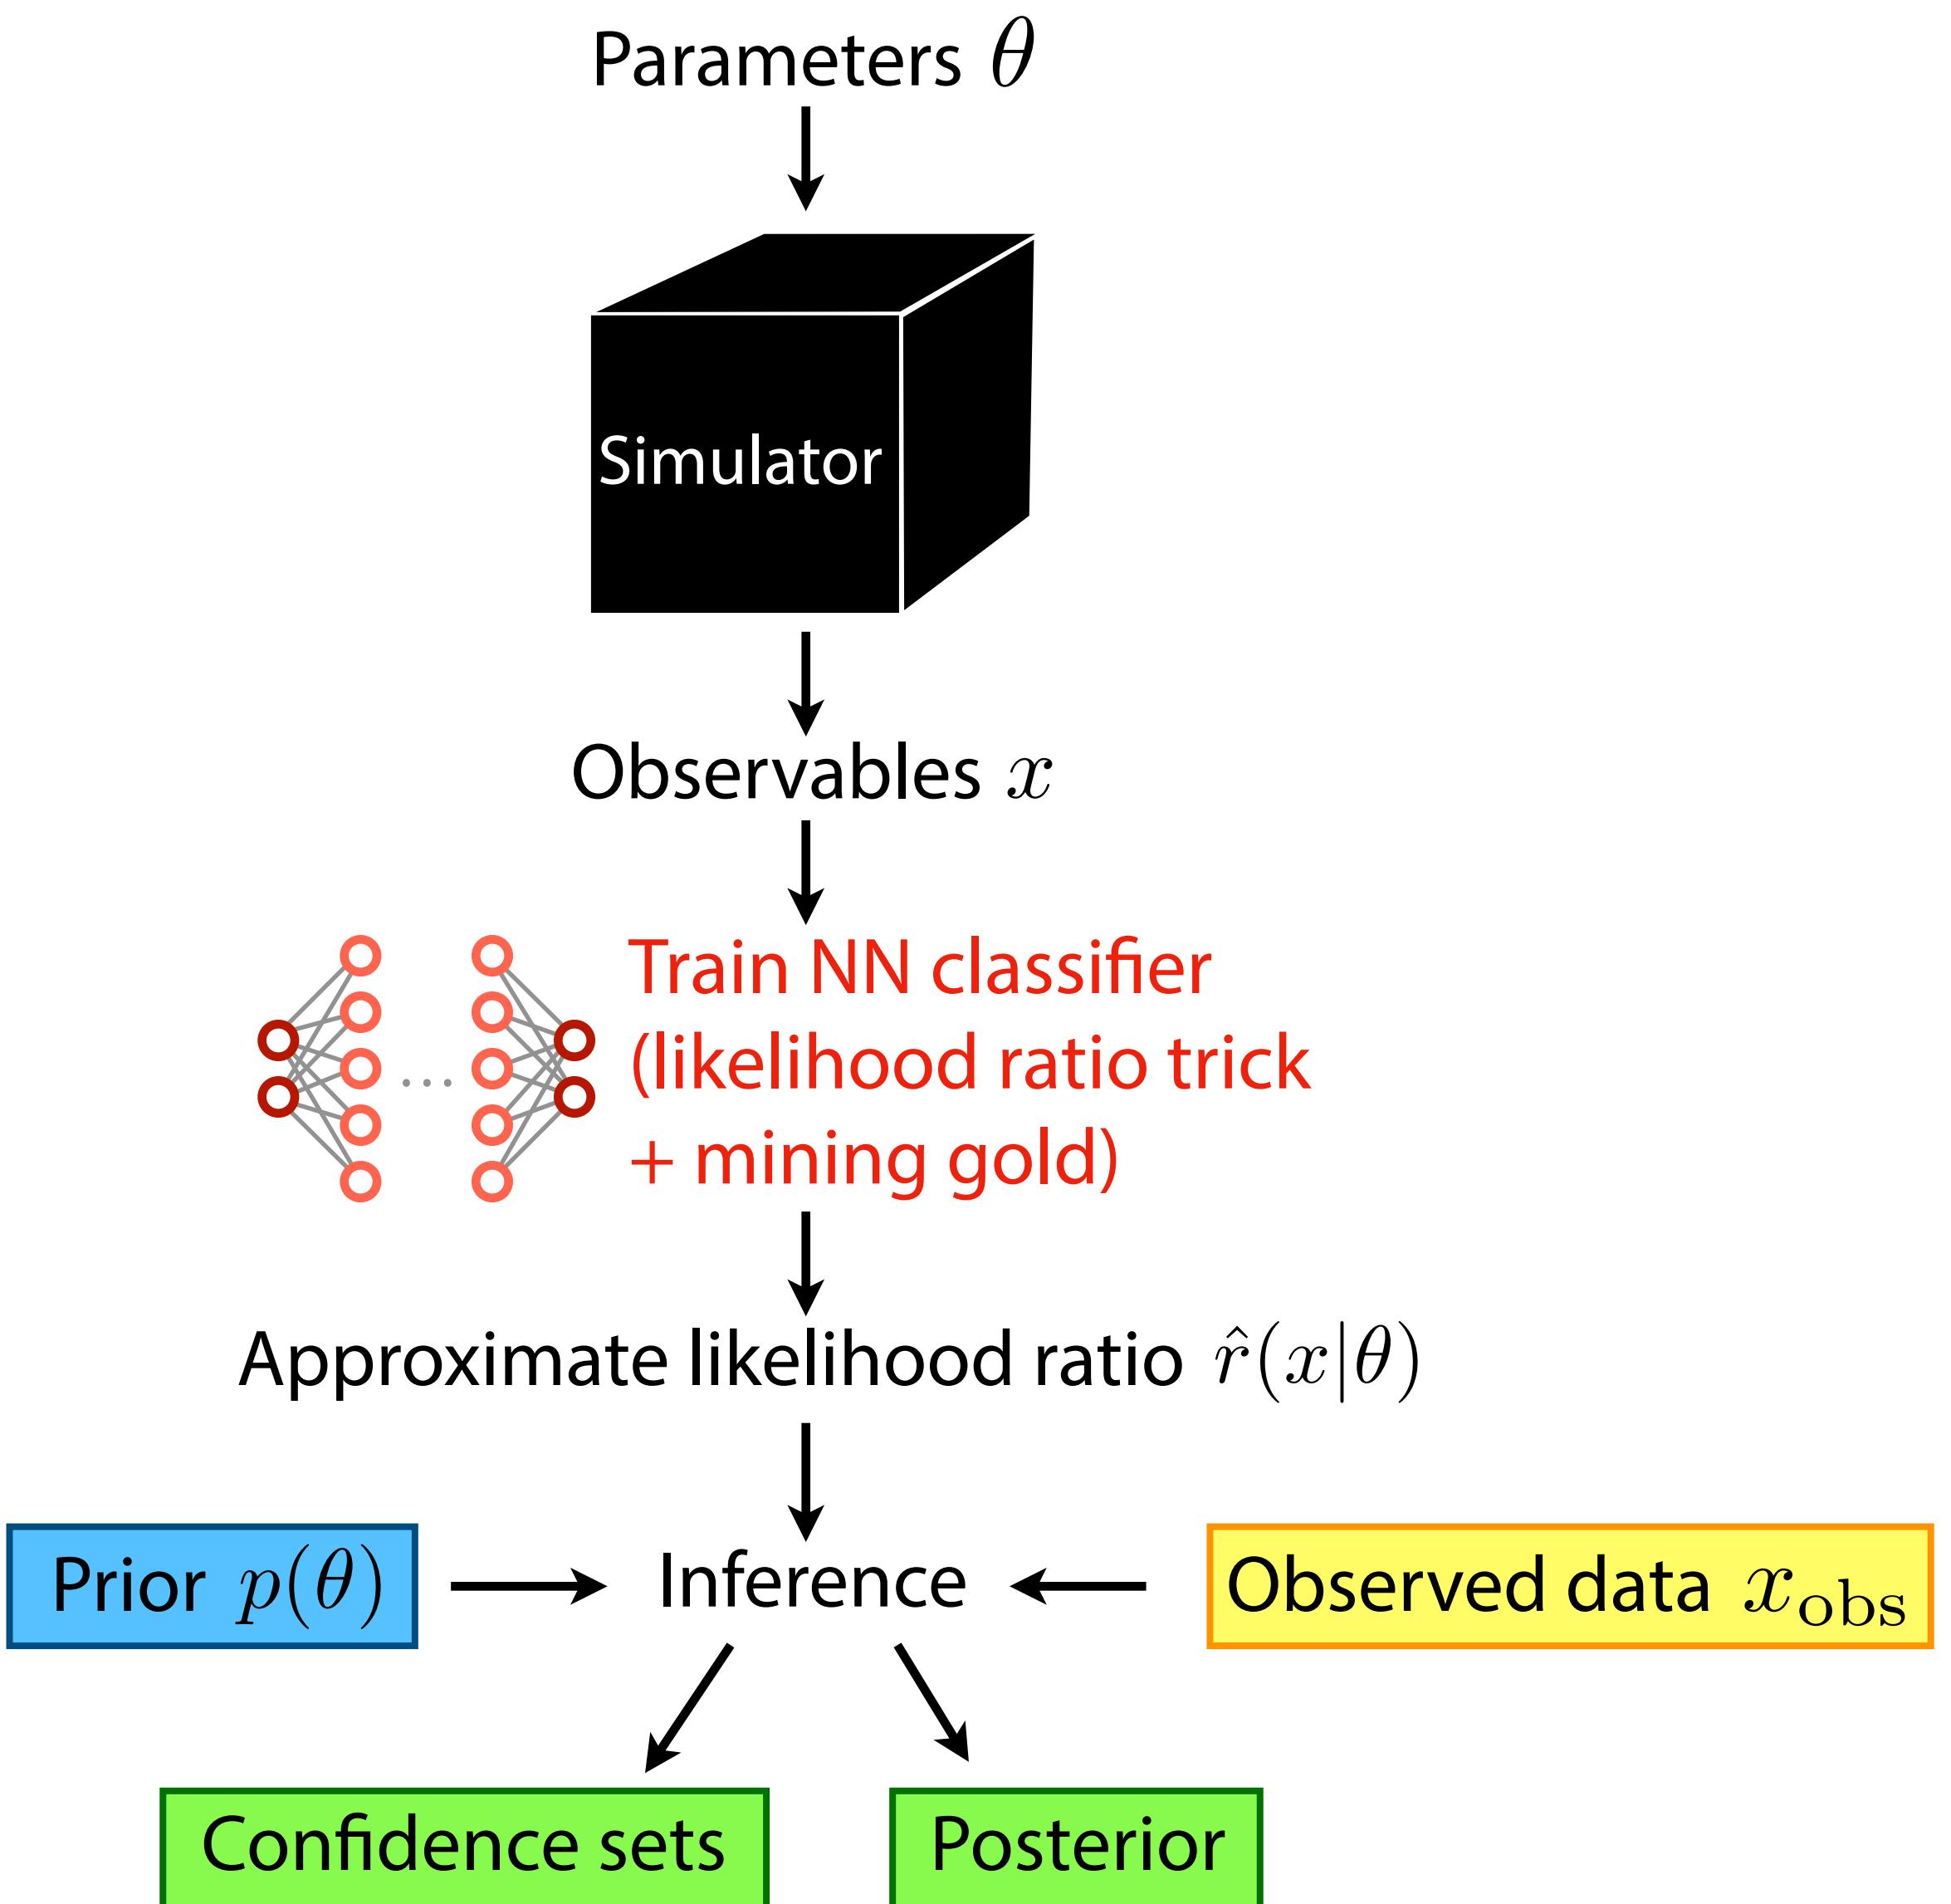


Proof-of-principle simulator

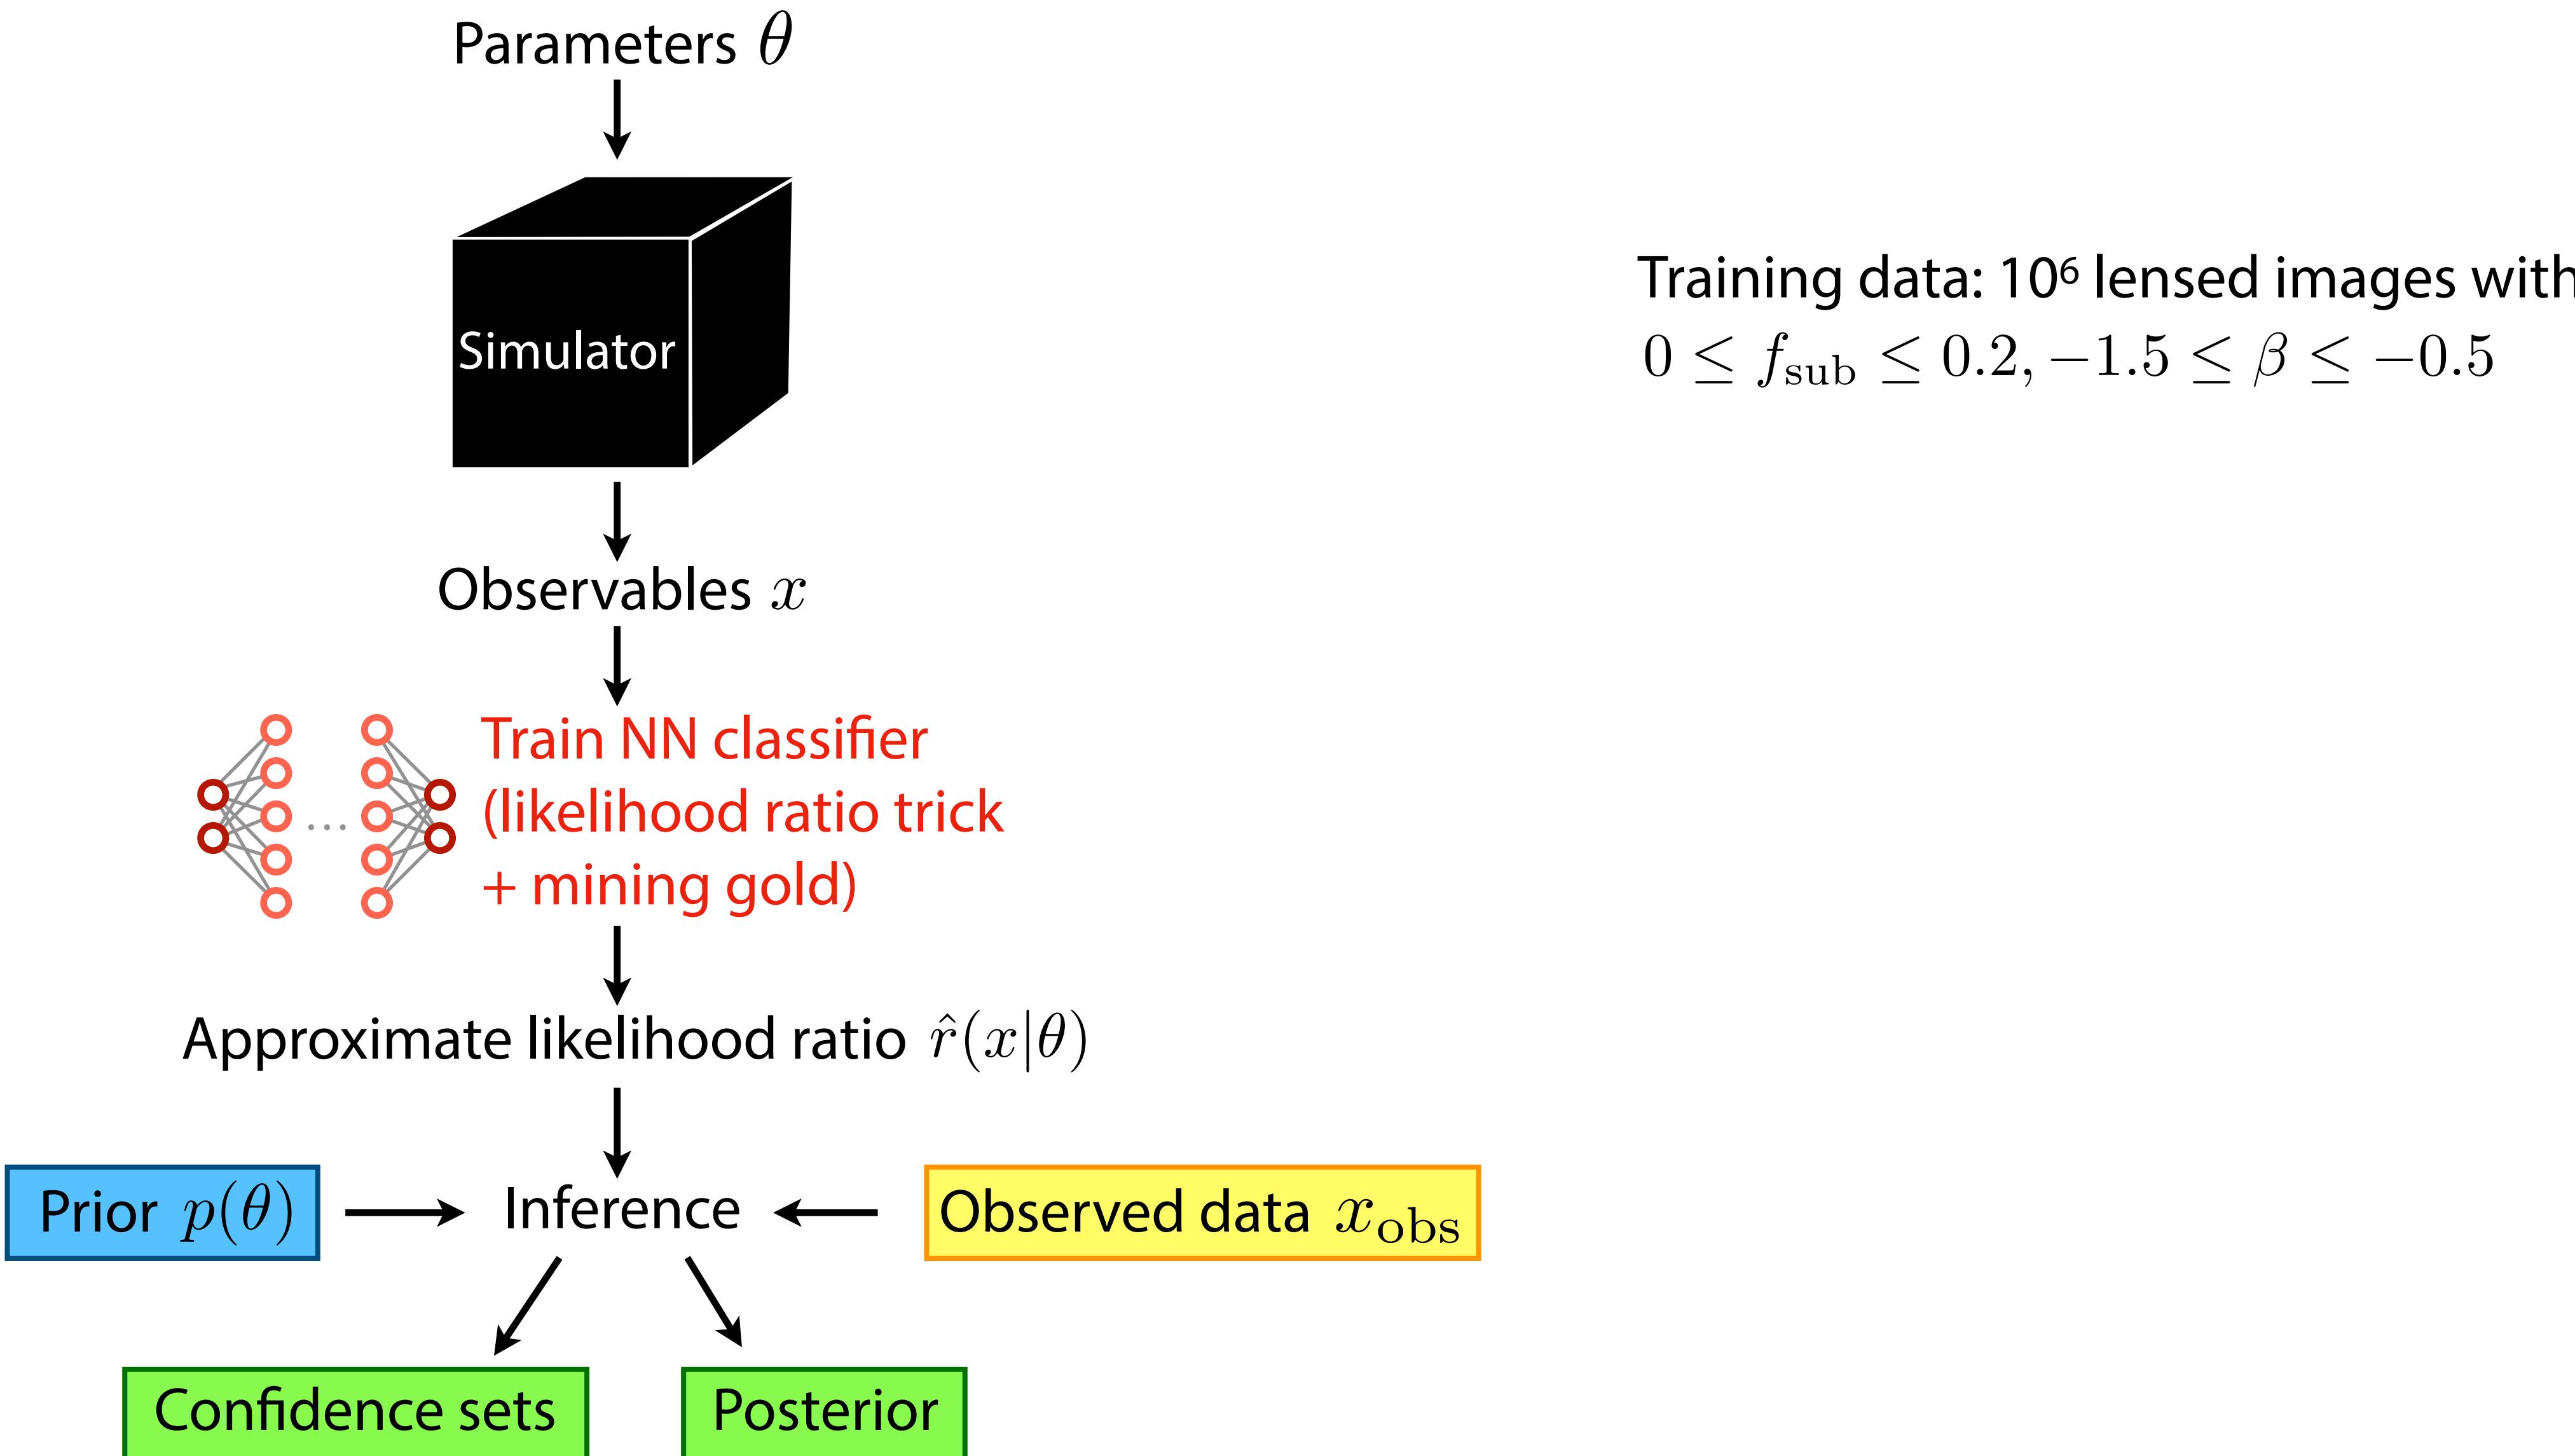
[following T. Collett 1507.02657]



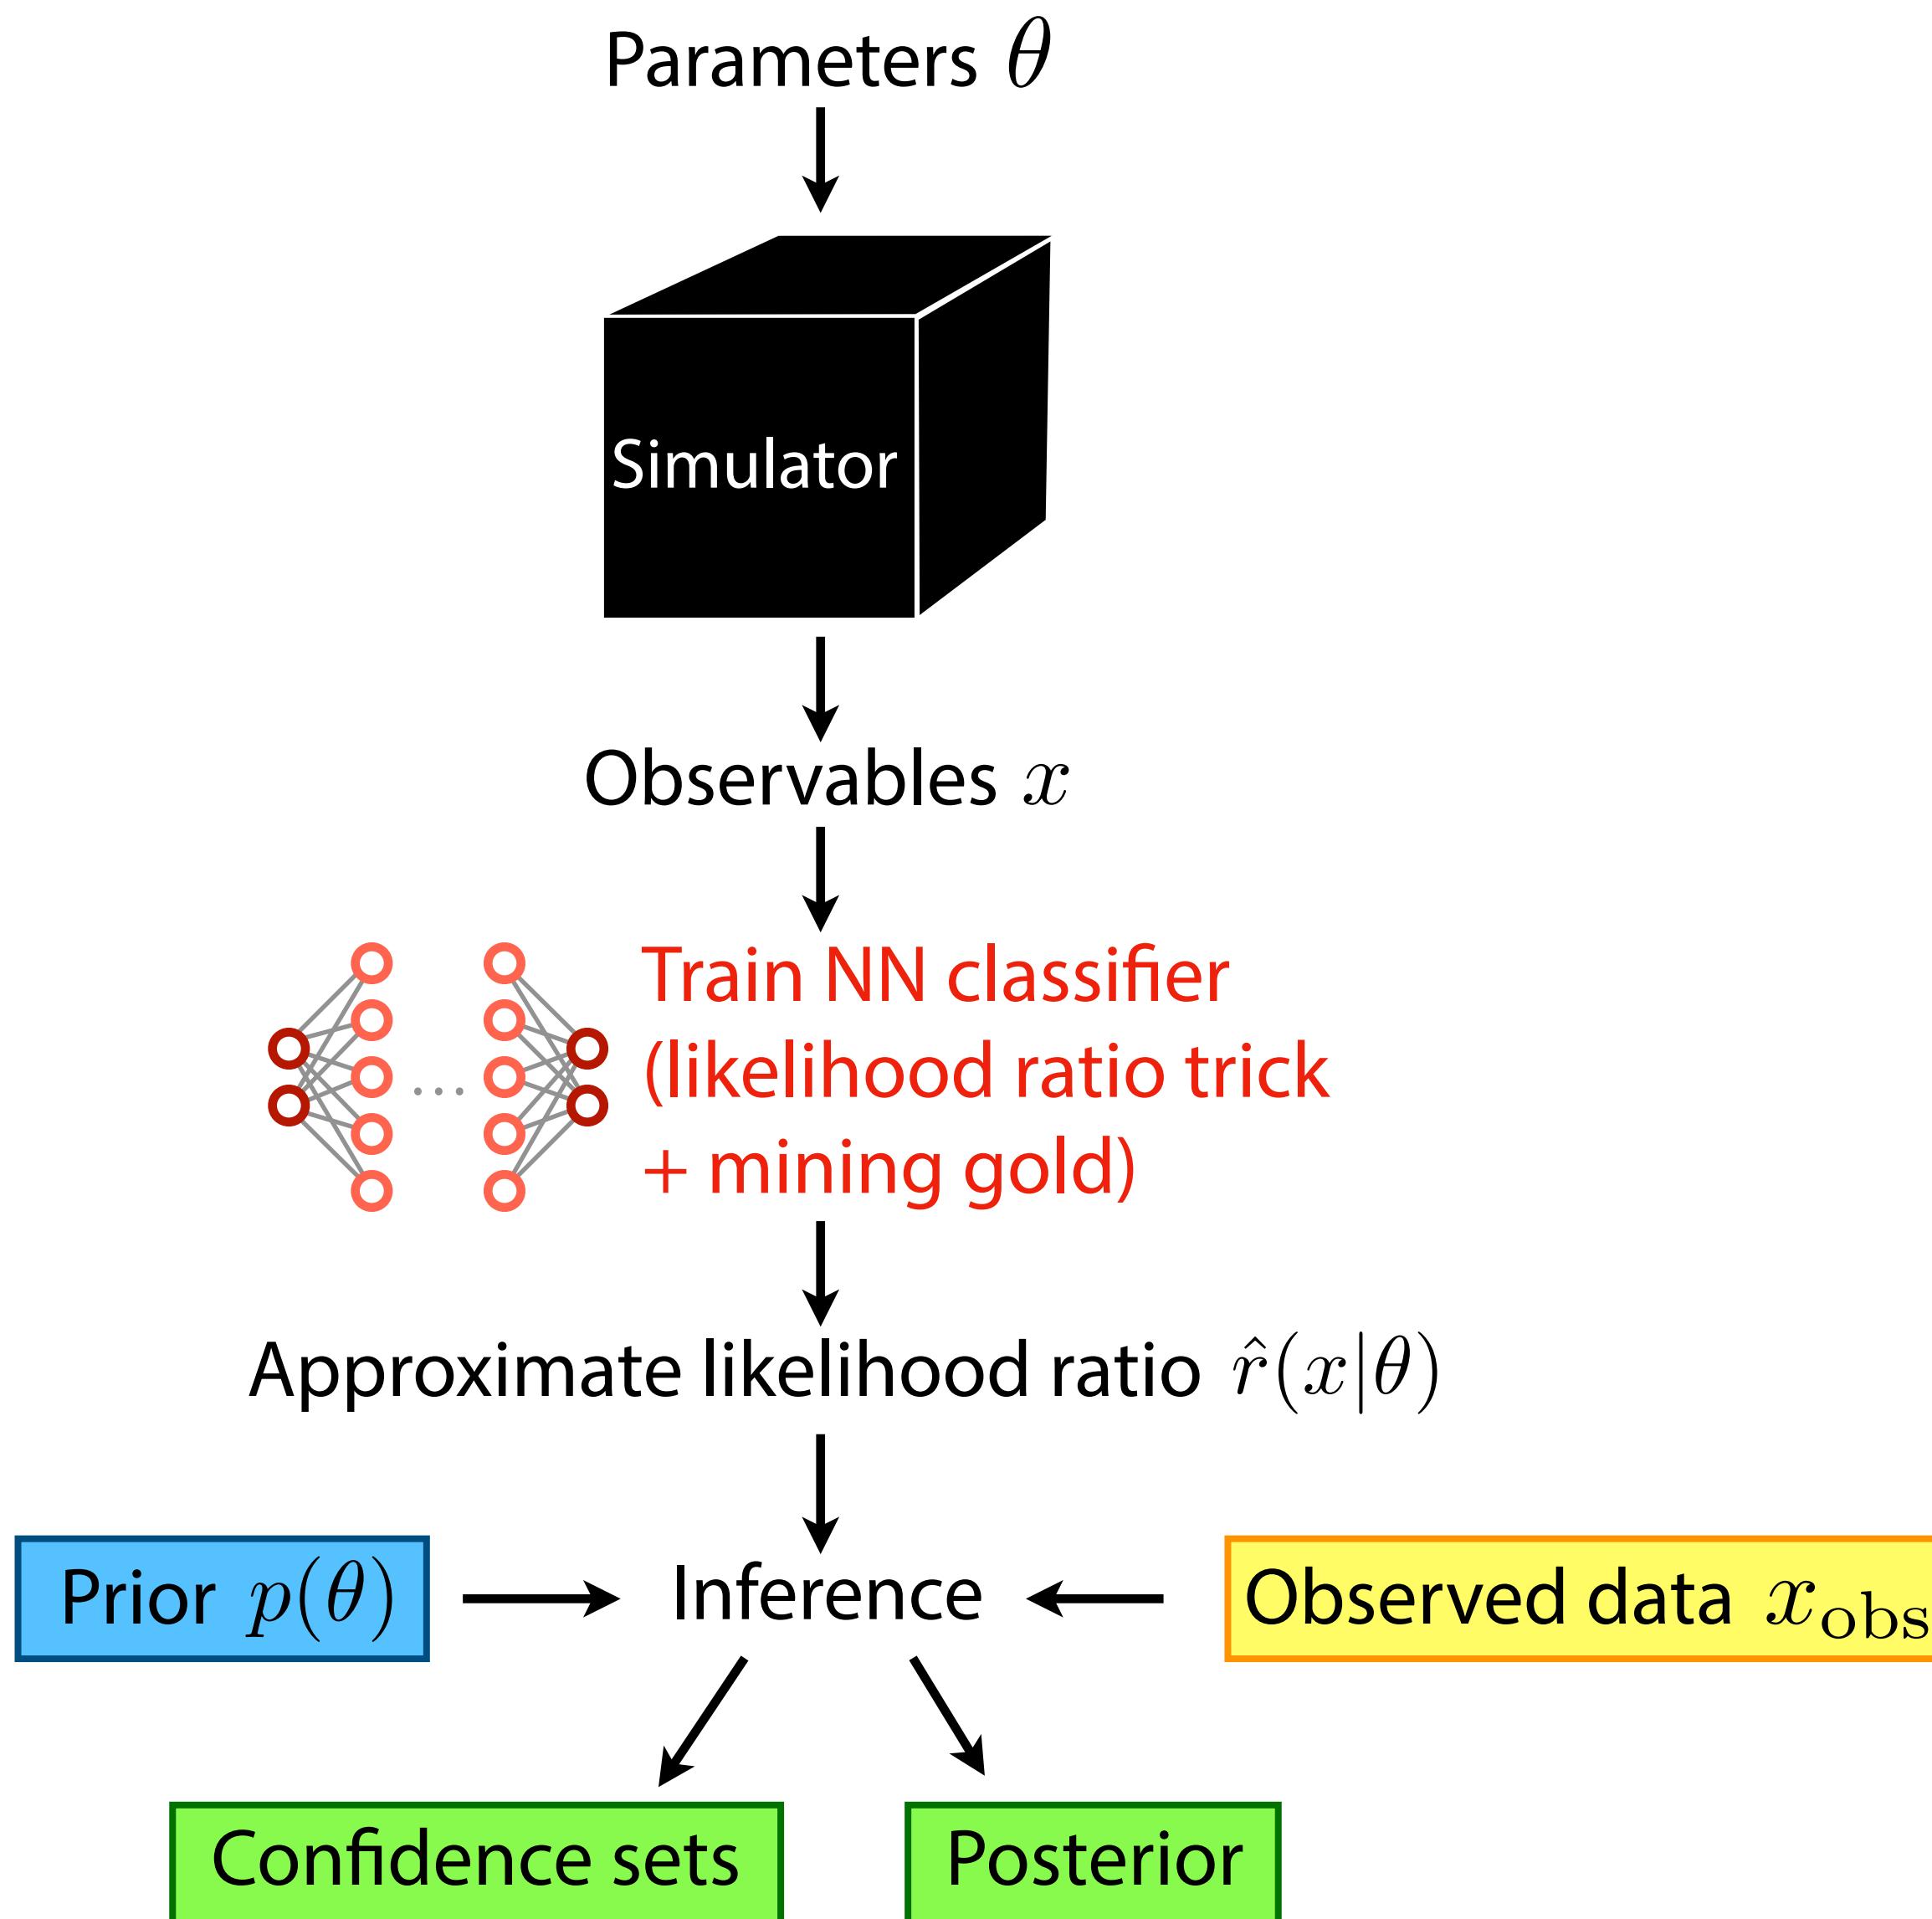
Inference setup



Inference setup



Inference setup

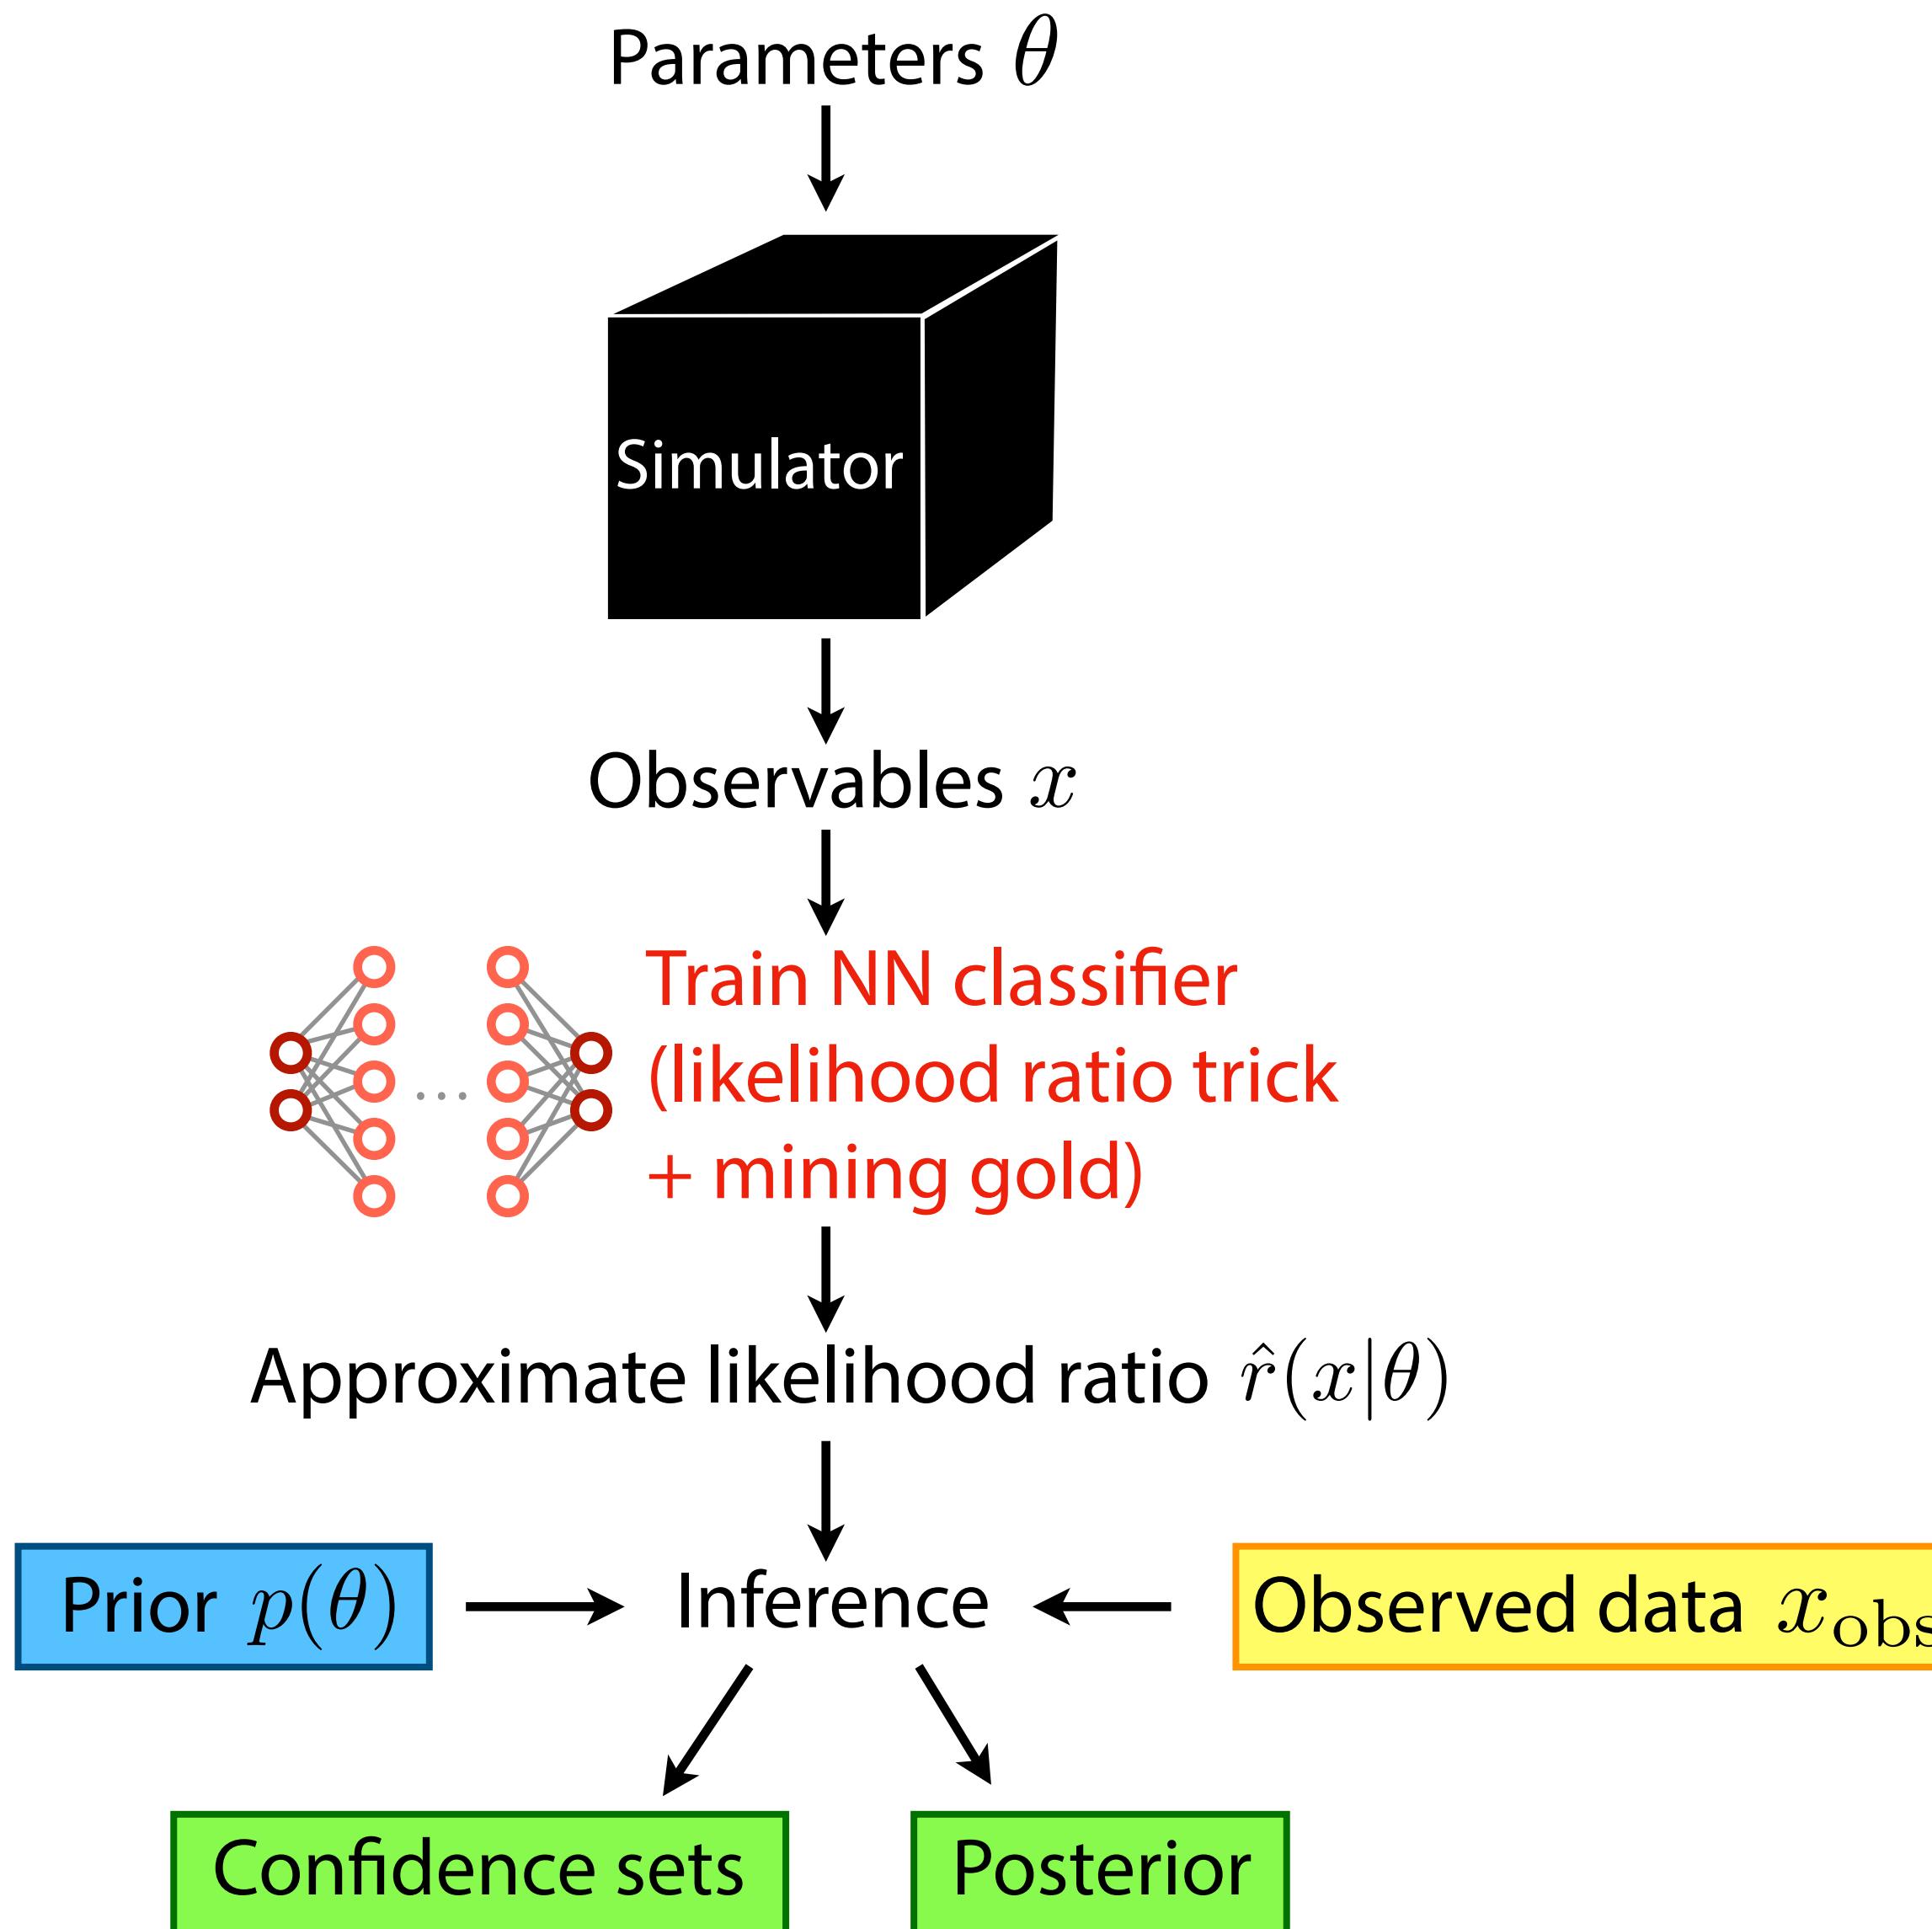


Training data: 10^6 lensed images with
 $0 \leq f_{\text{sub}} \leq 0.2, -1.5 \leq \beta \leq -0.5$

Convolutional neural network (modified ResNet-18)
trained on ALICES loss
[M. Stoye, JB, J. Pavez, G, Louppe, K. Cranmer 1808.00973]

Calibration of network output

Inference setup



Training data: 10^6 lensed images with
 $0 \leq f_{\text{sub}} \leq 0.2, -1.5 \leq \beta \leq -0.5$

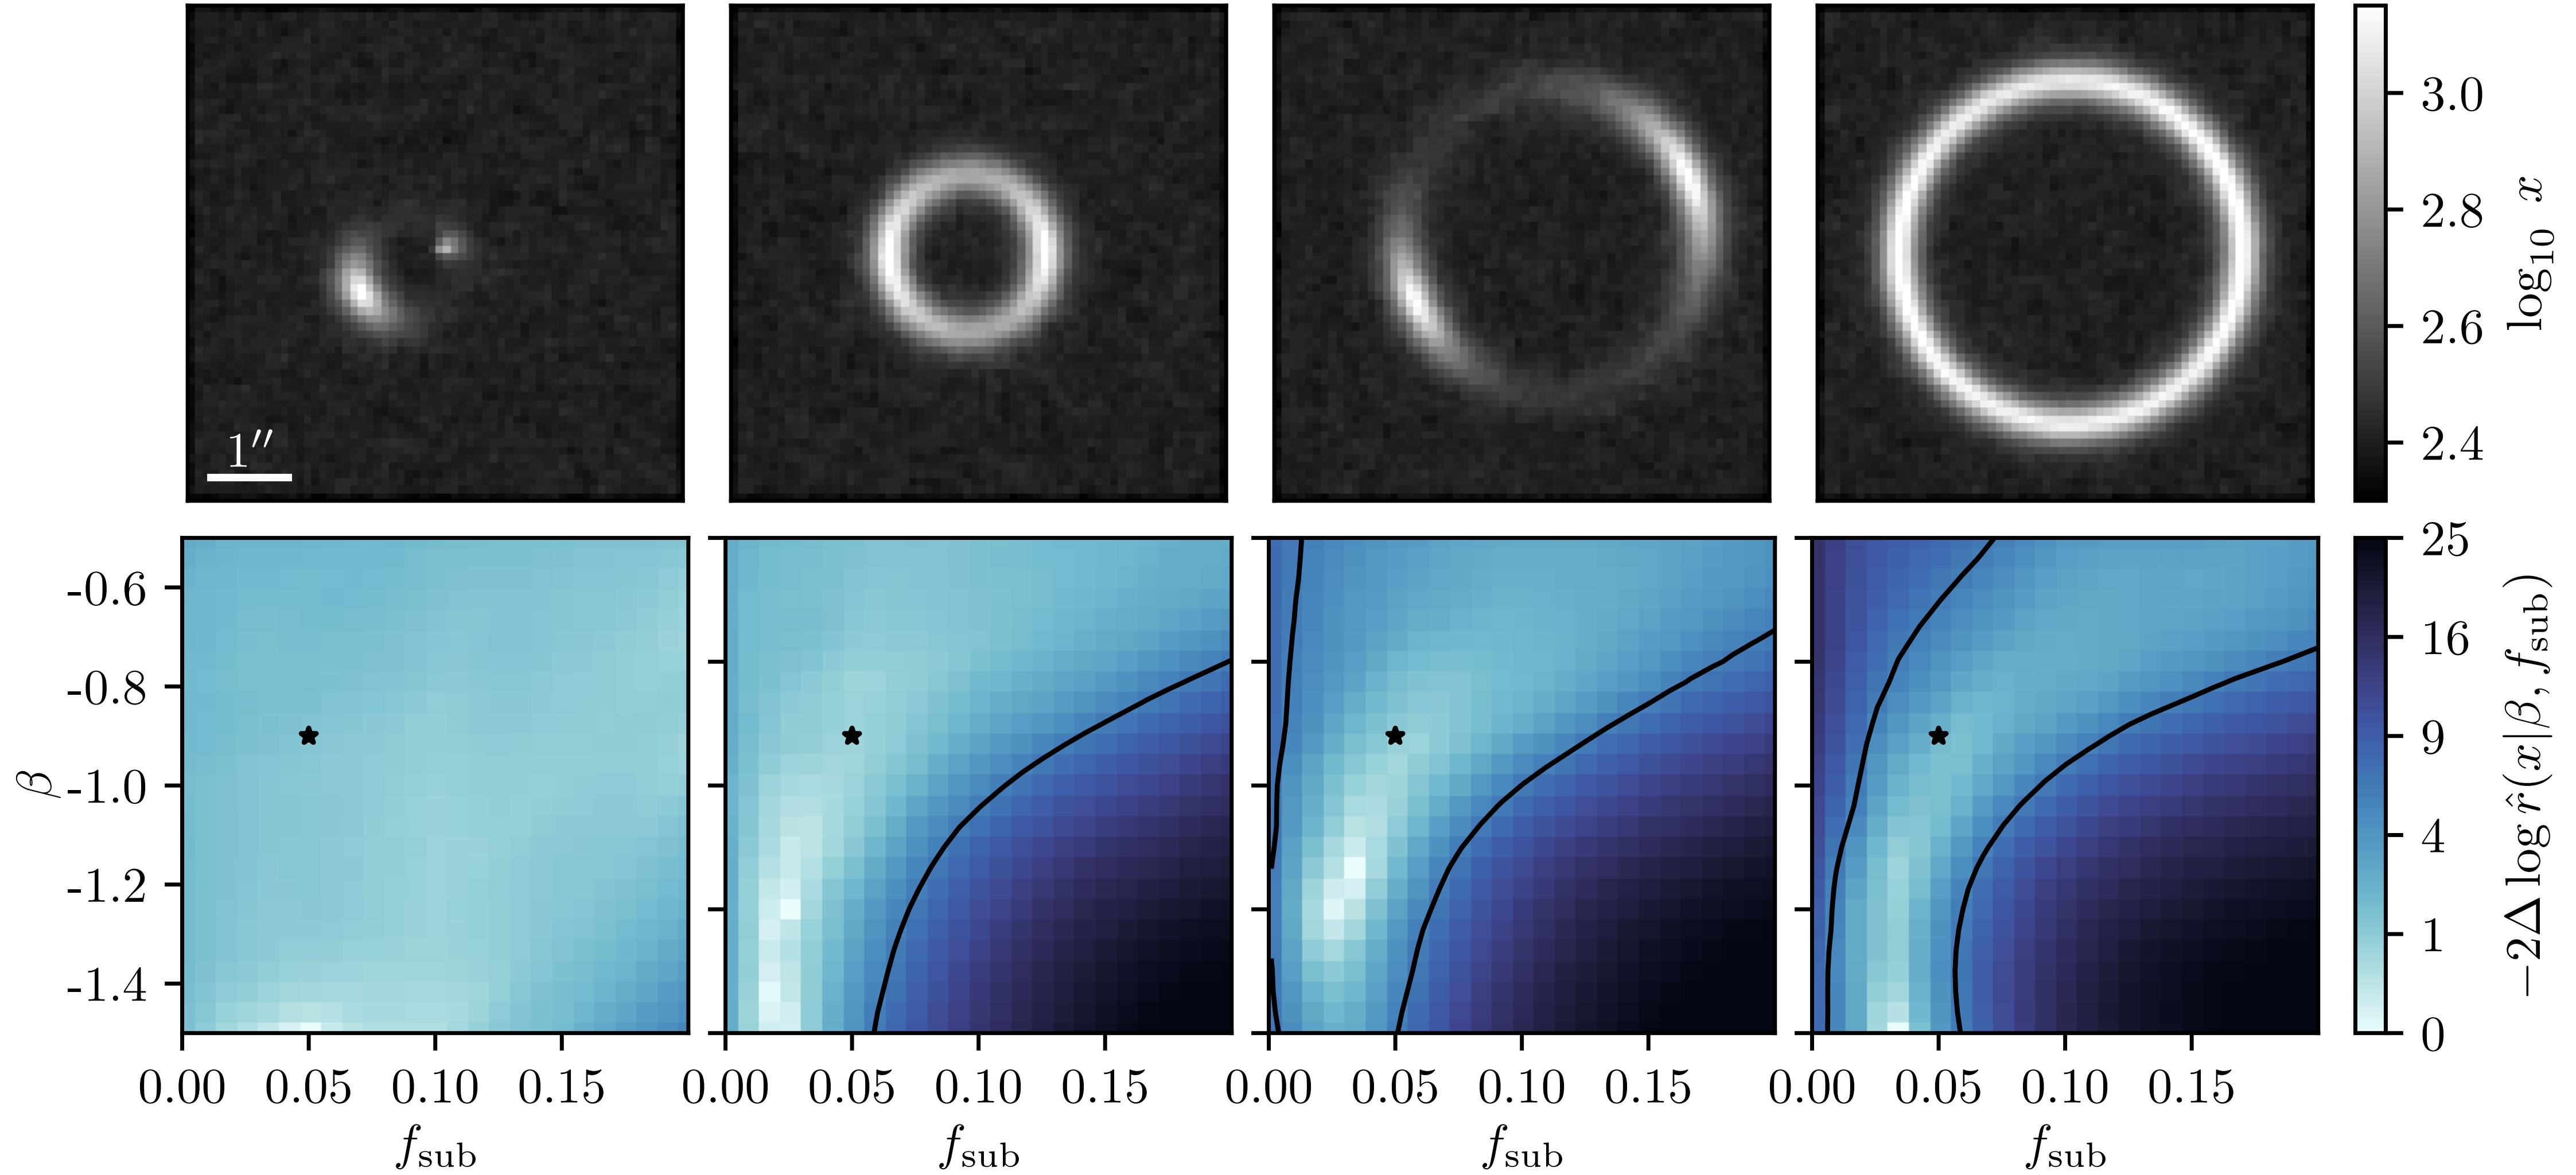
Convolutional neural network (modified ResNet-18)
trained on ALICES loss
[M. Stoye, JB, J. Pavez, G, Louppe, K. Cranmer 1808.00973]

Calibration of network output

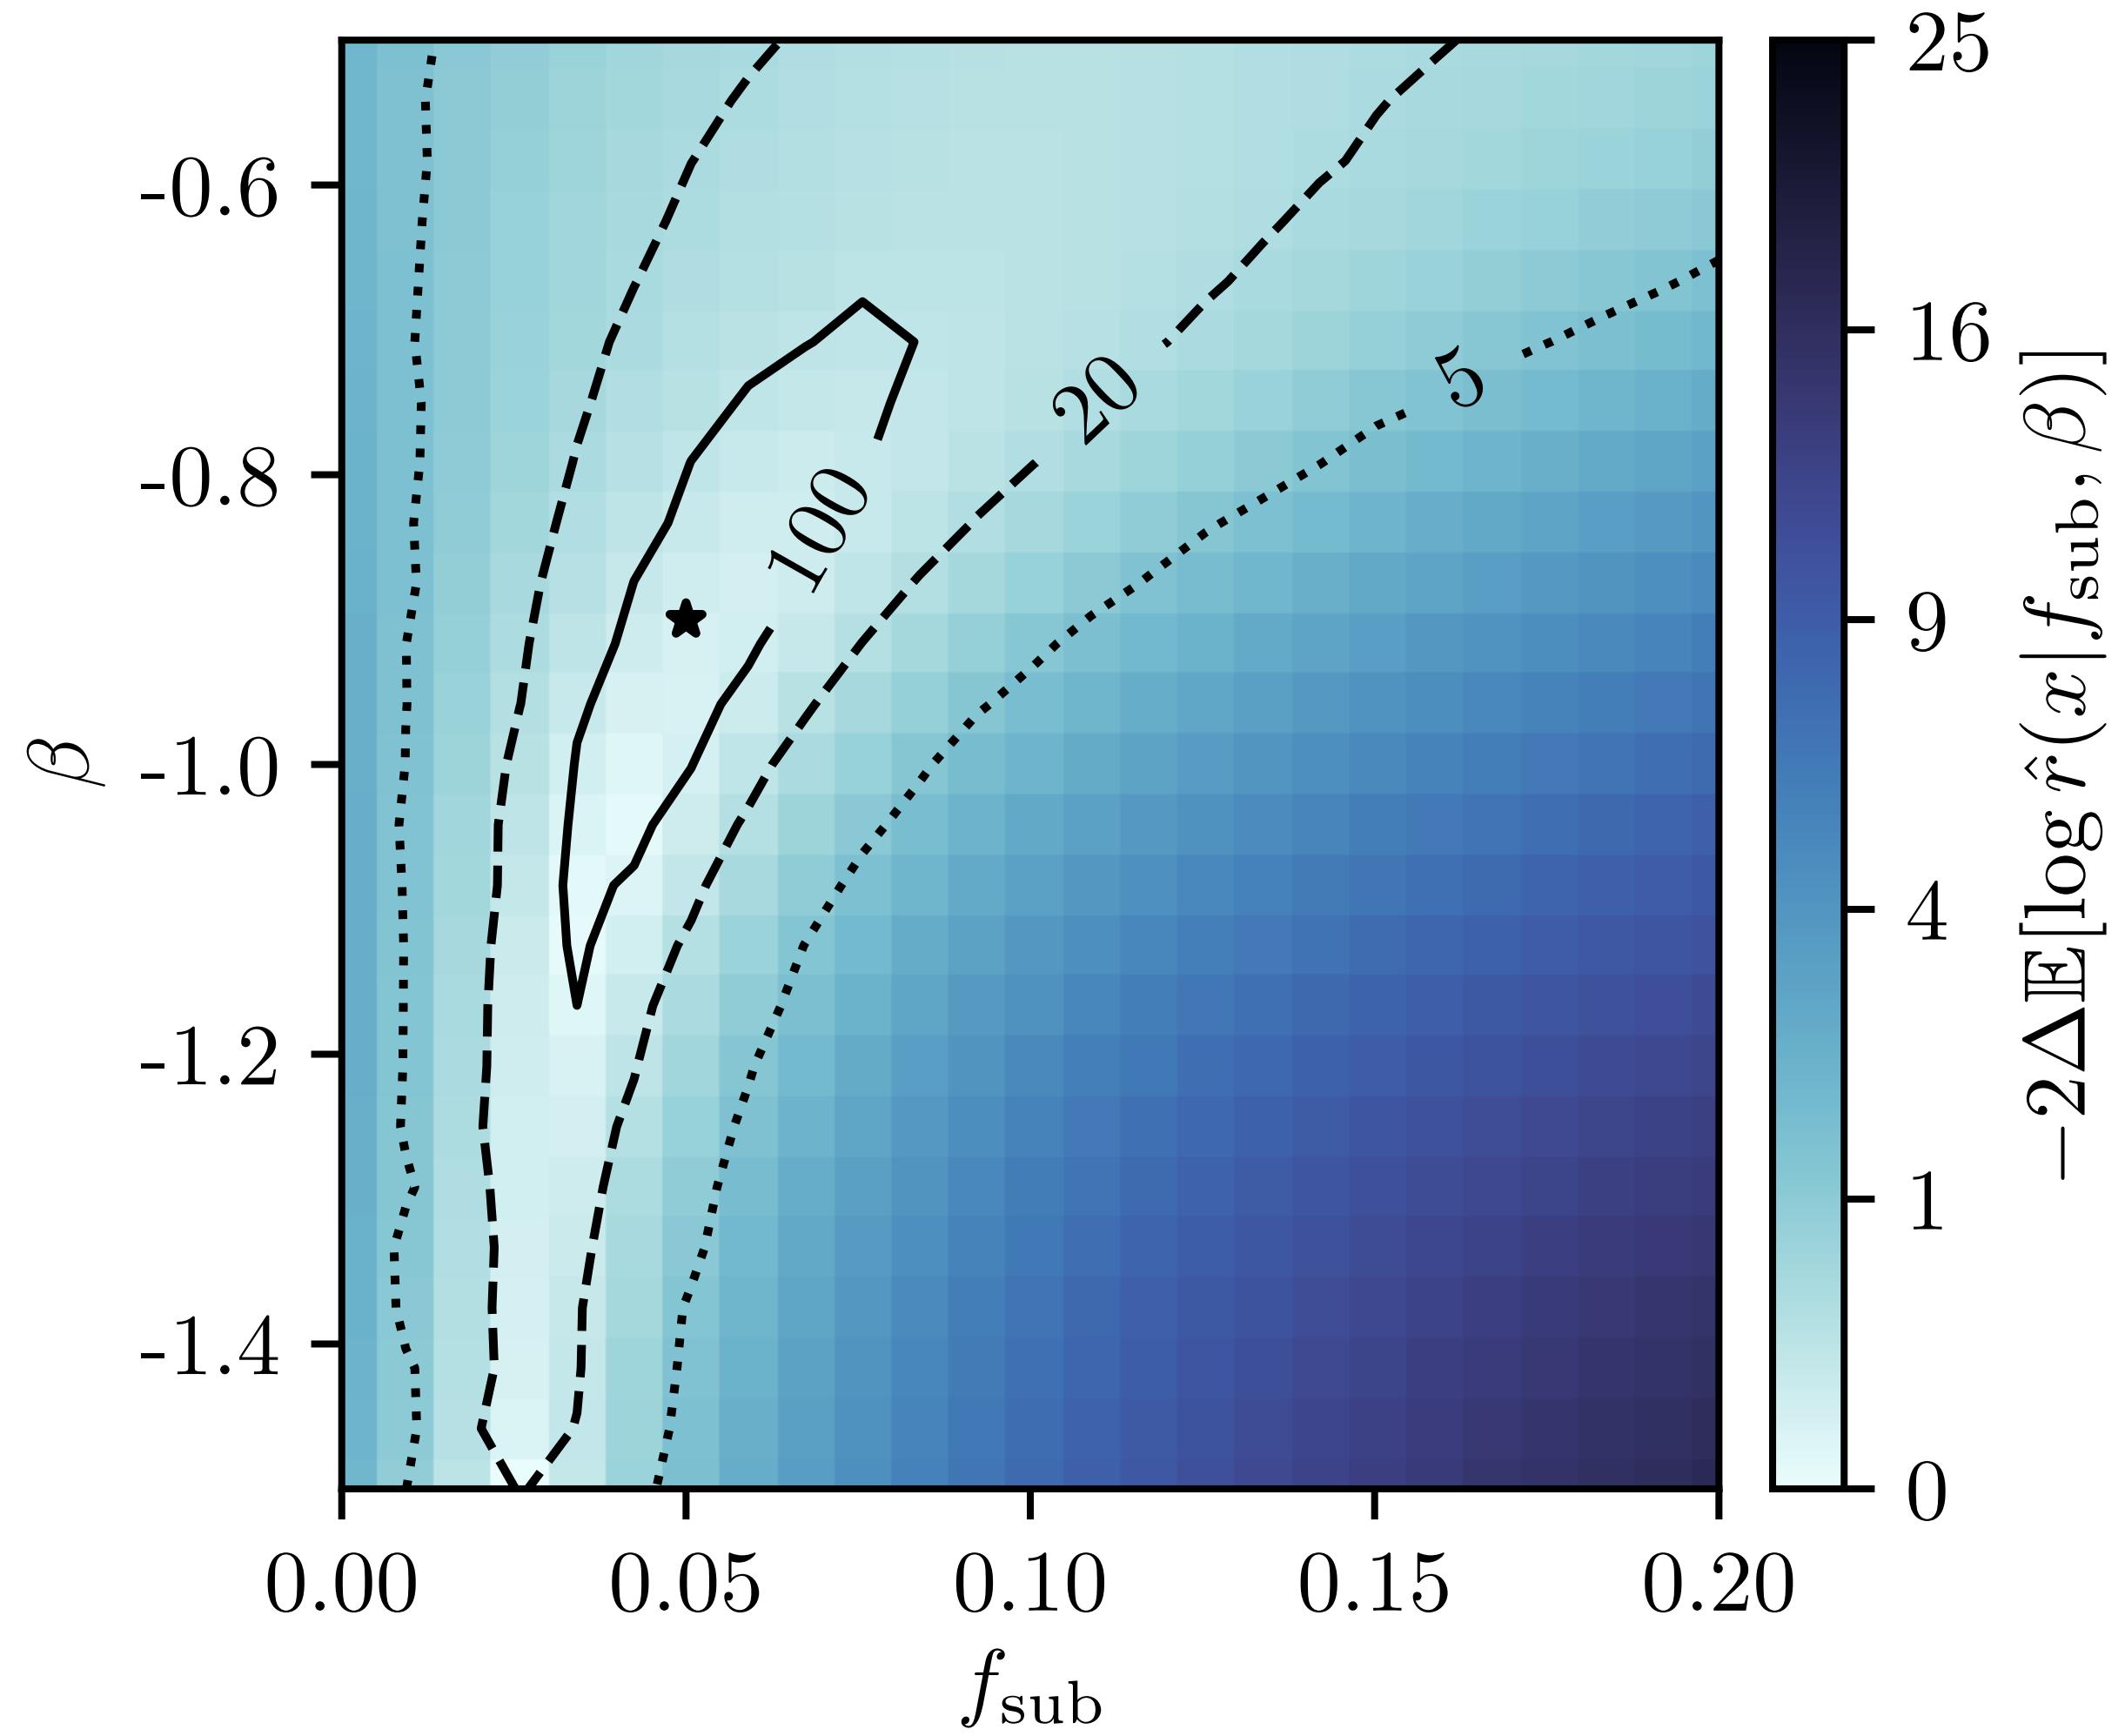
Synthetic “observed” data set: $f_{\text{sub}} = 0.05, \beta = -0.9$

Bayesian & frequentist inference

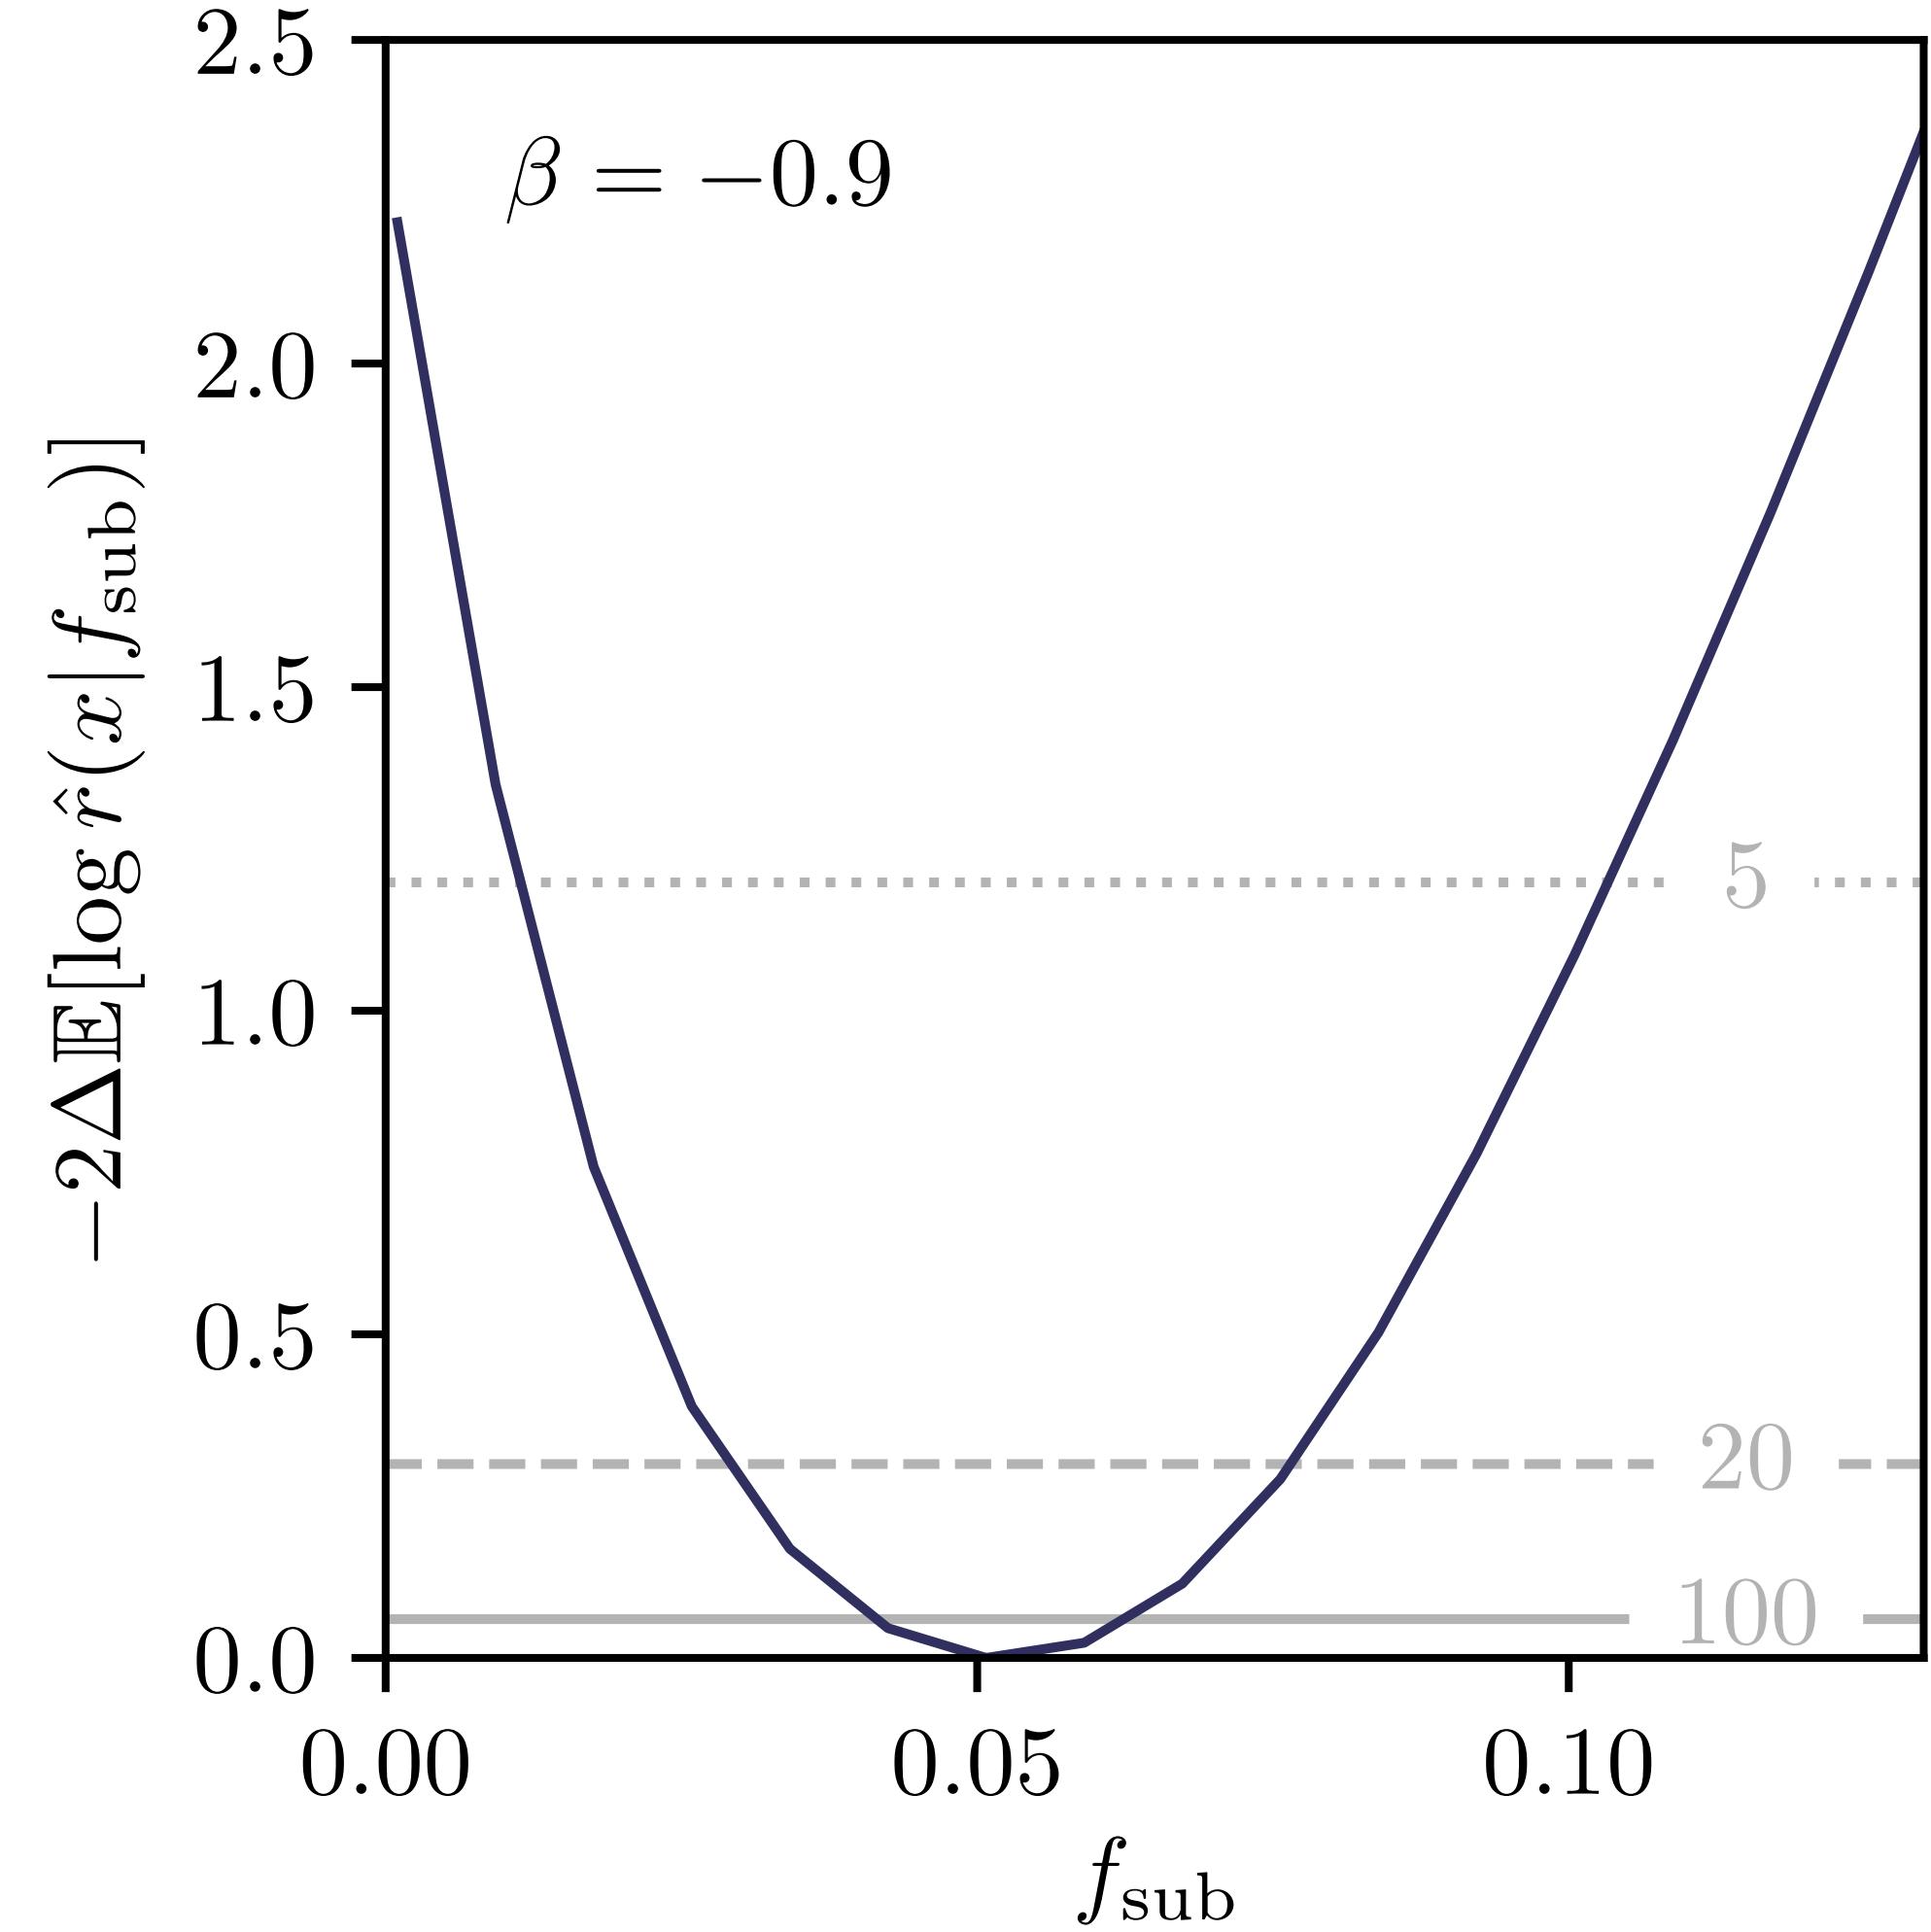
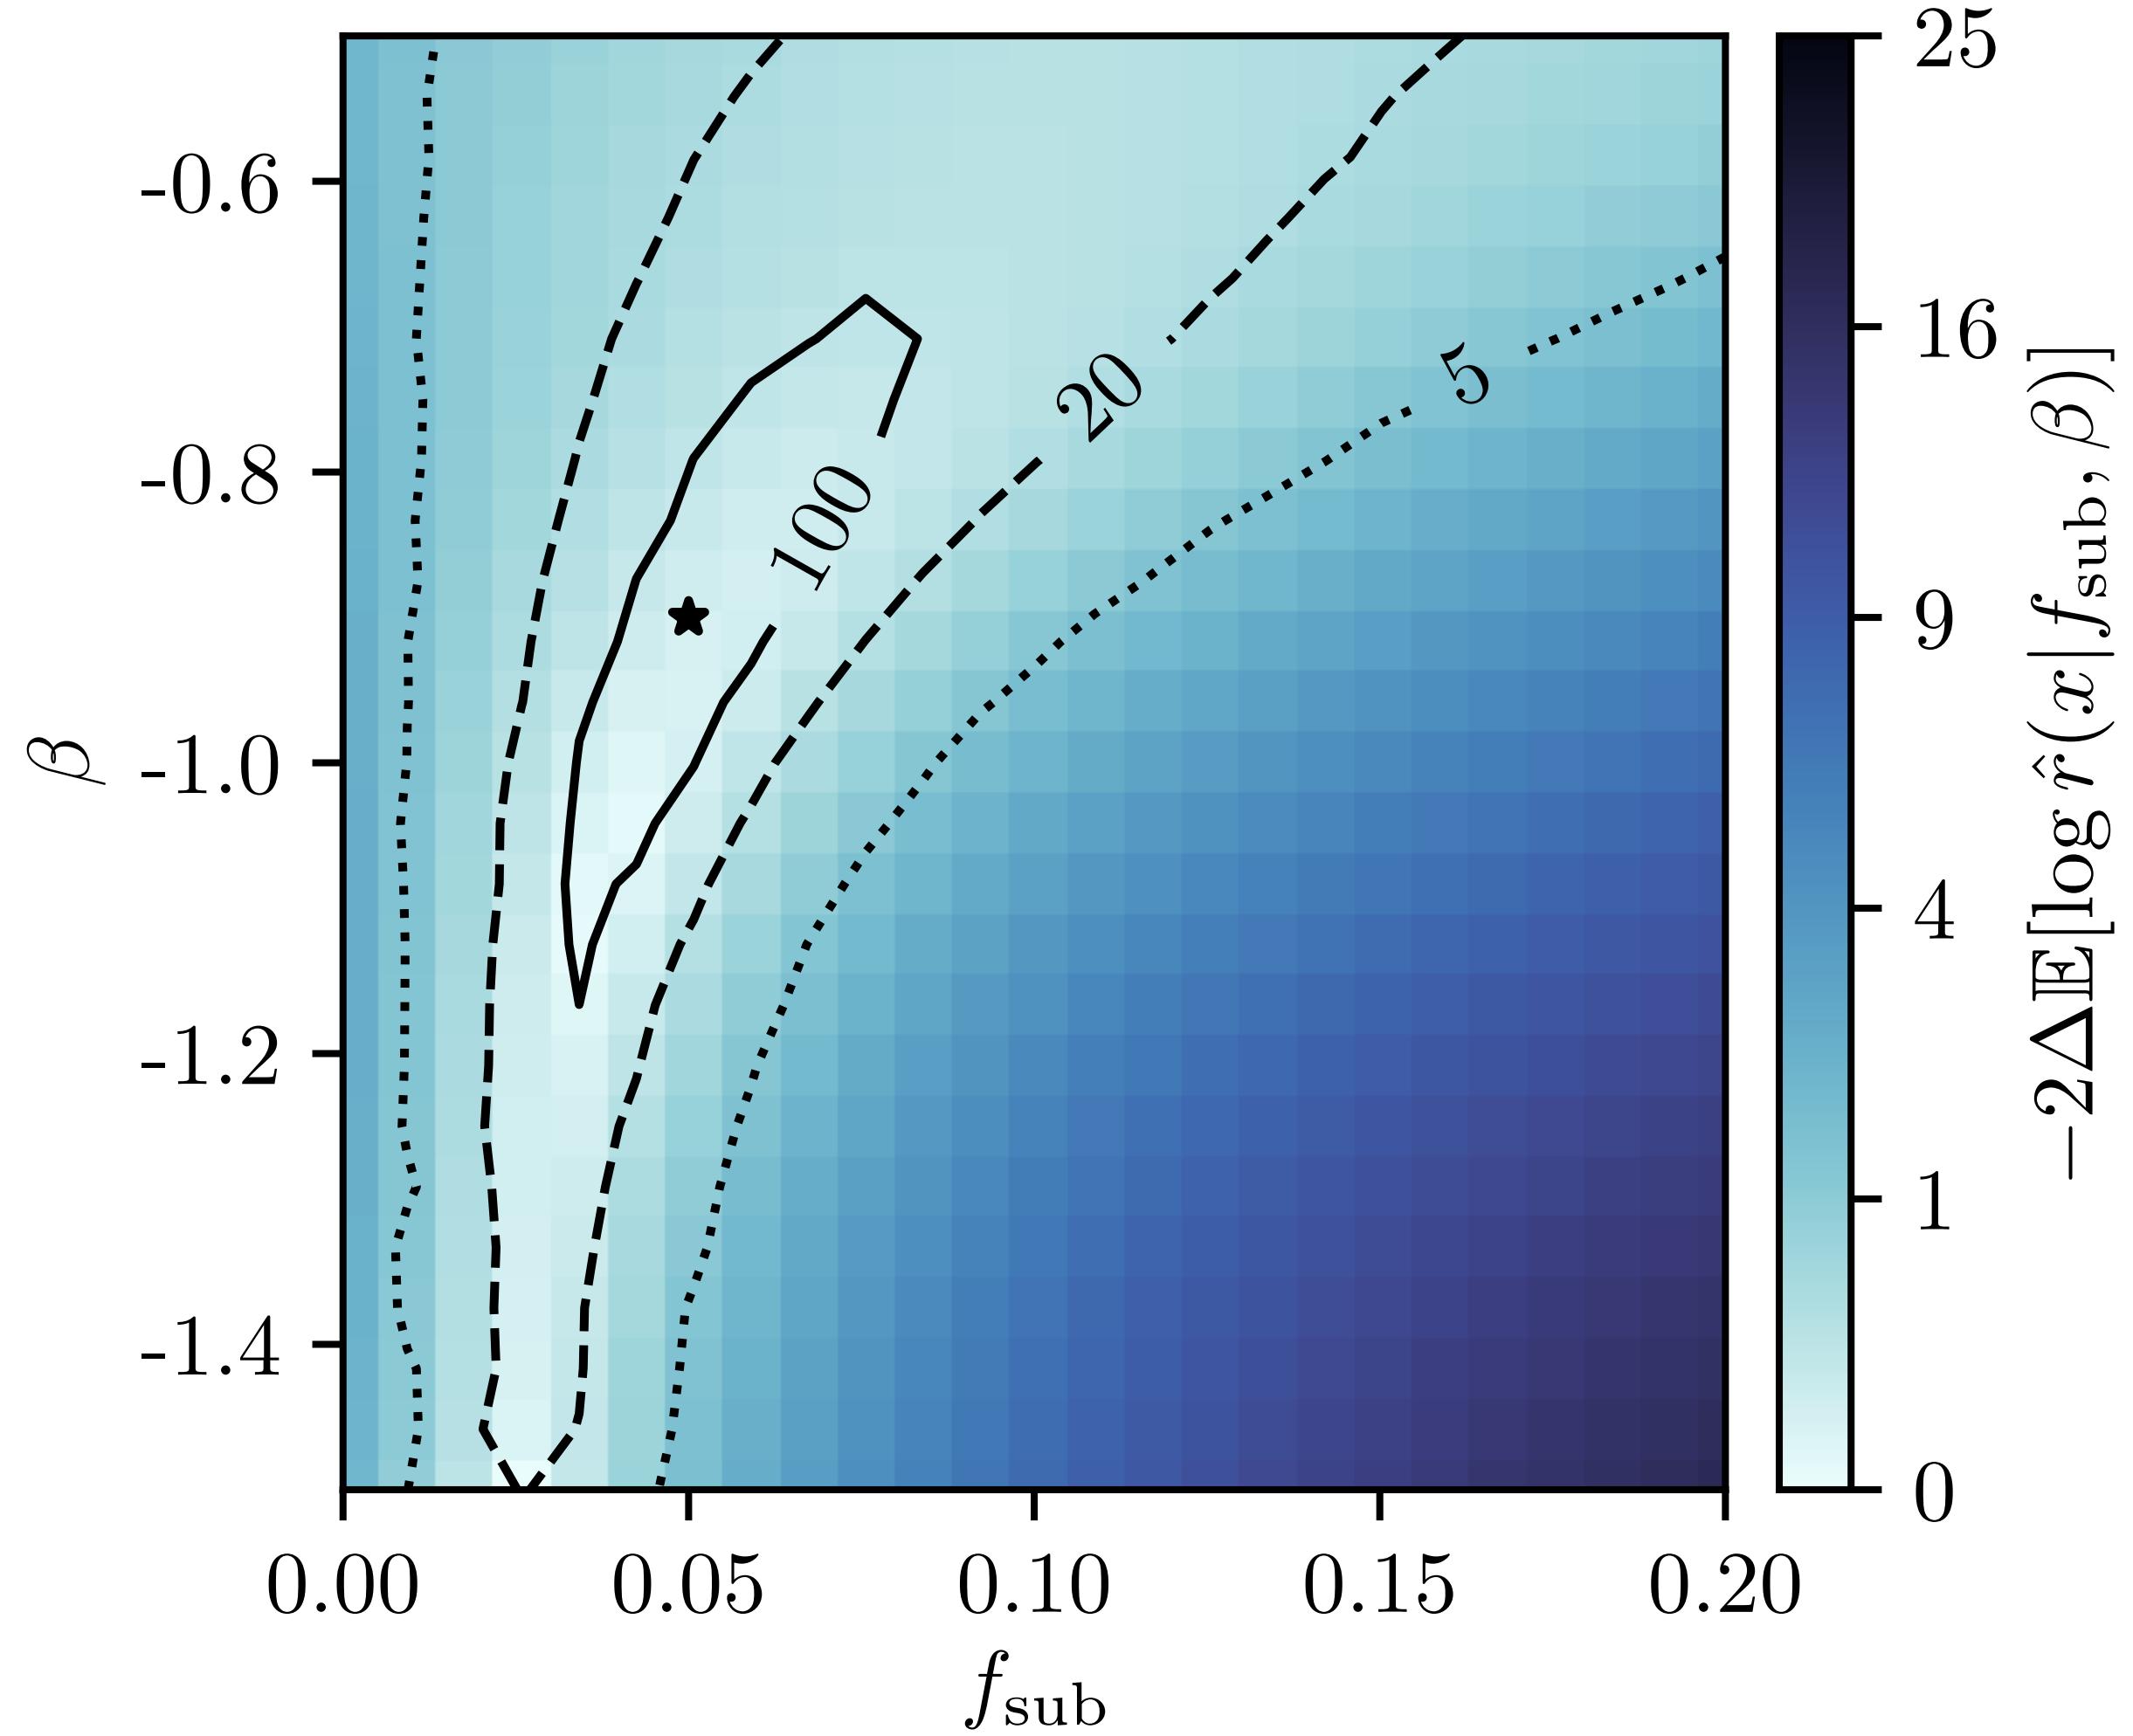
Inferring parameters from individual images



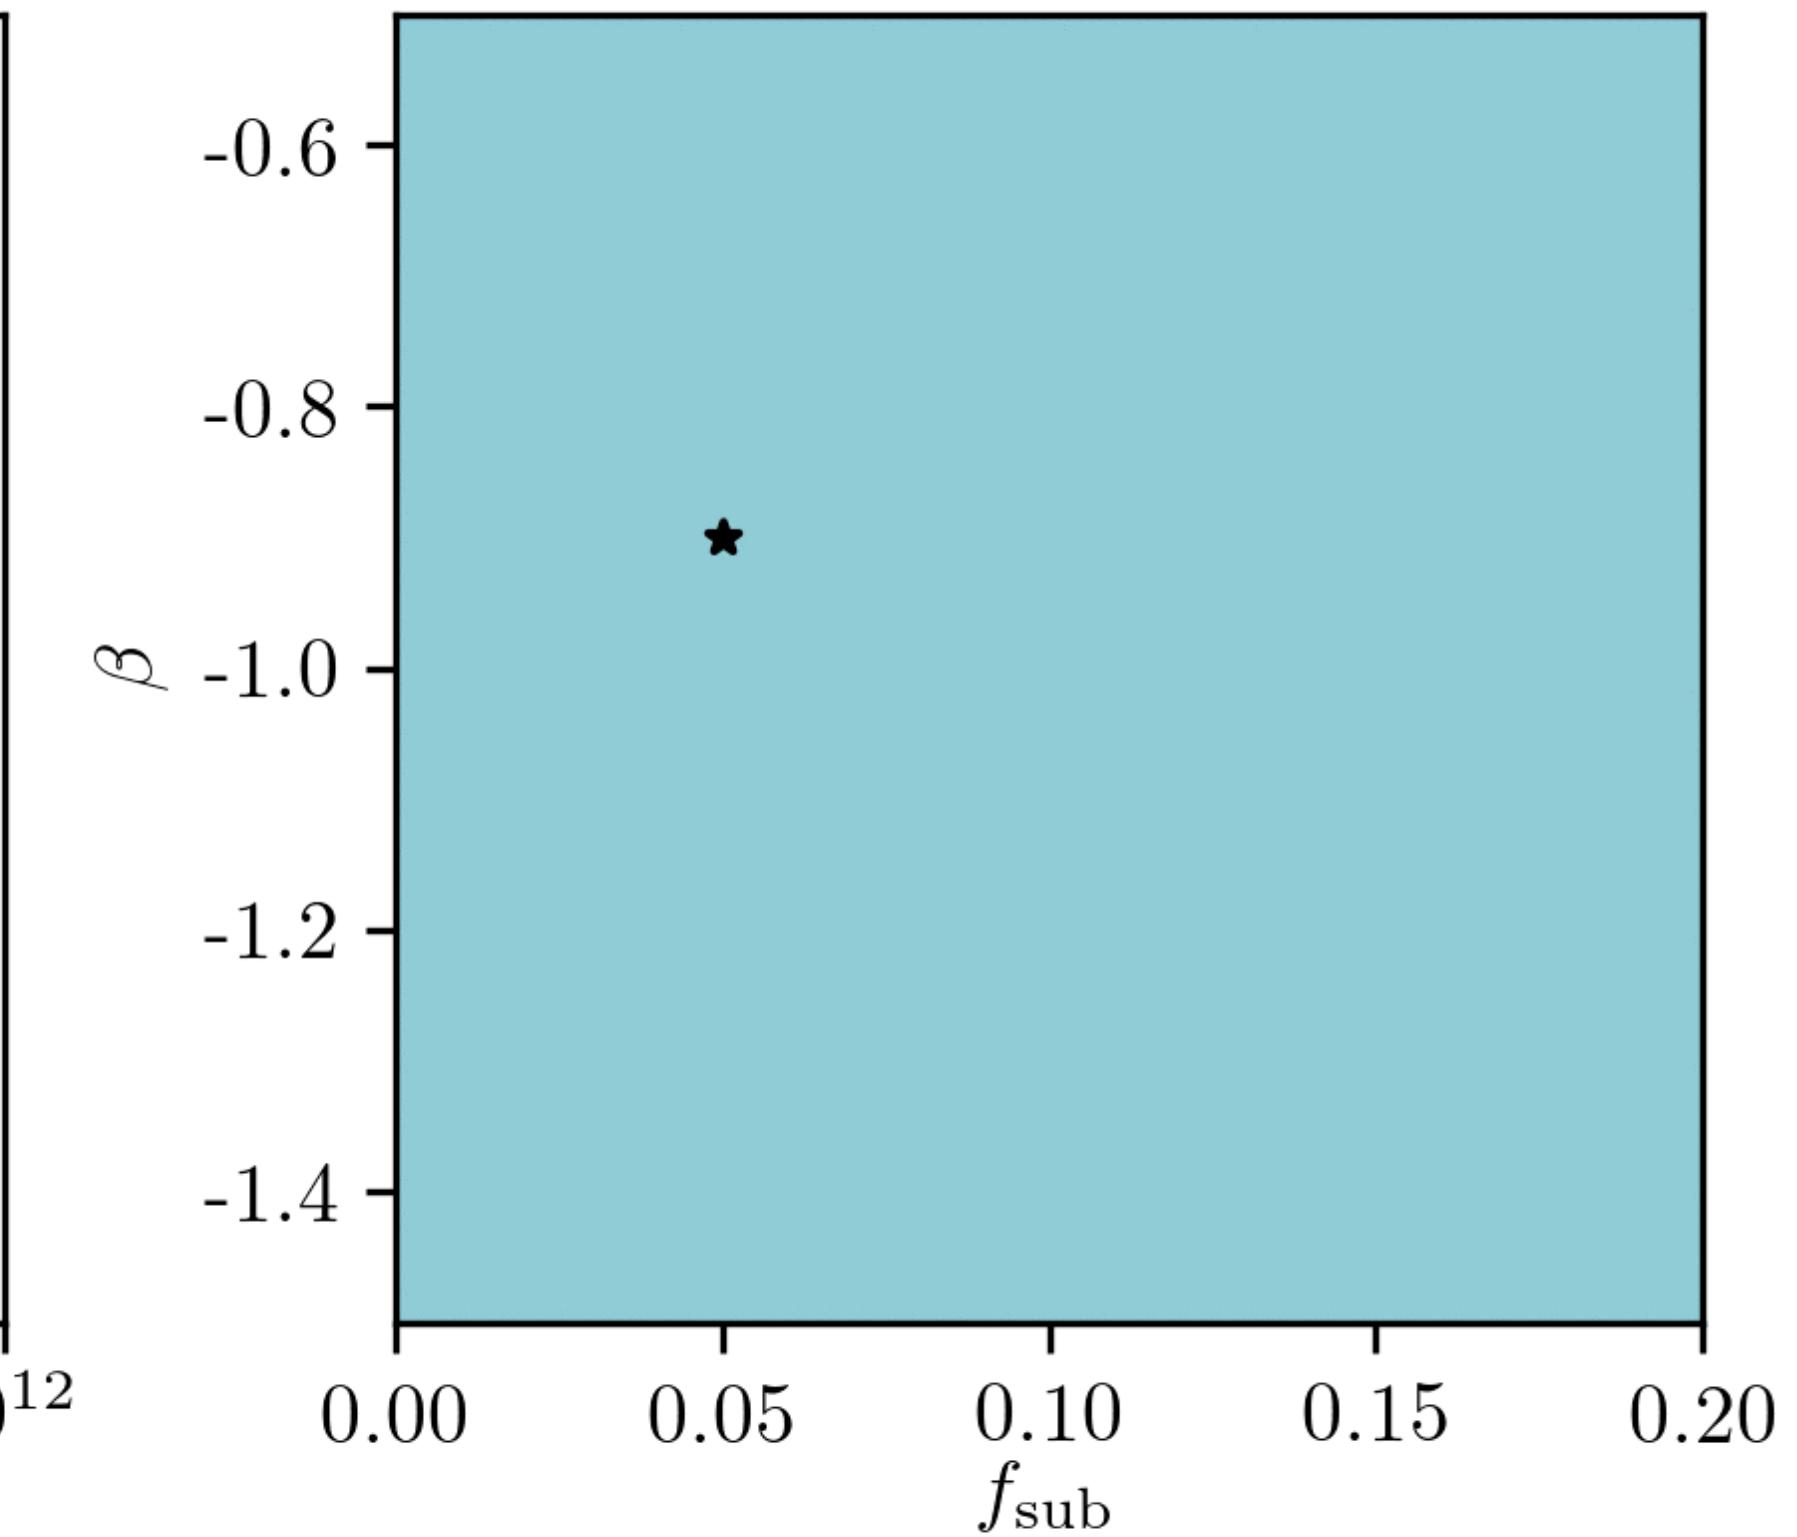
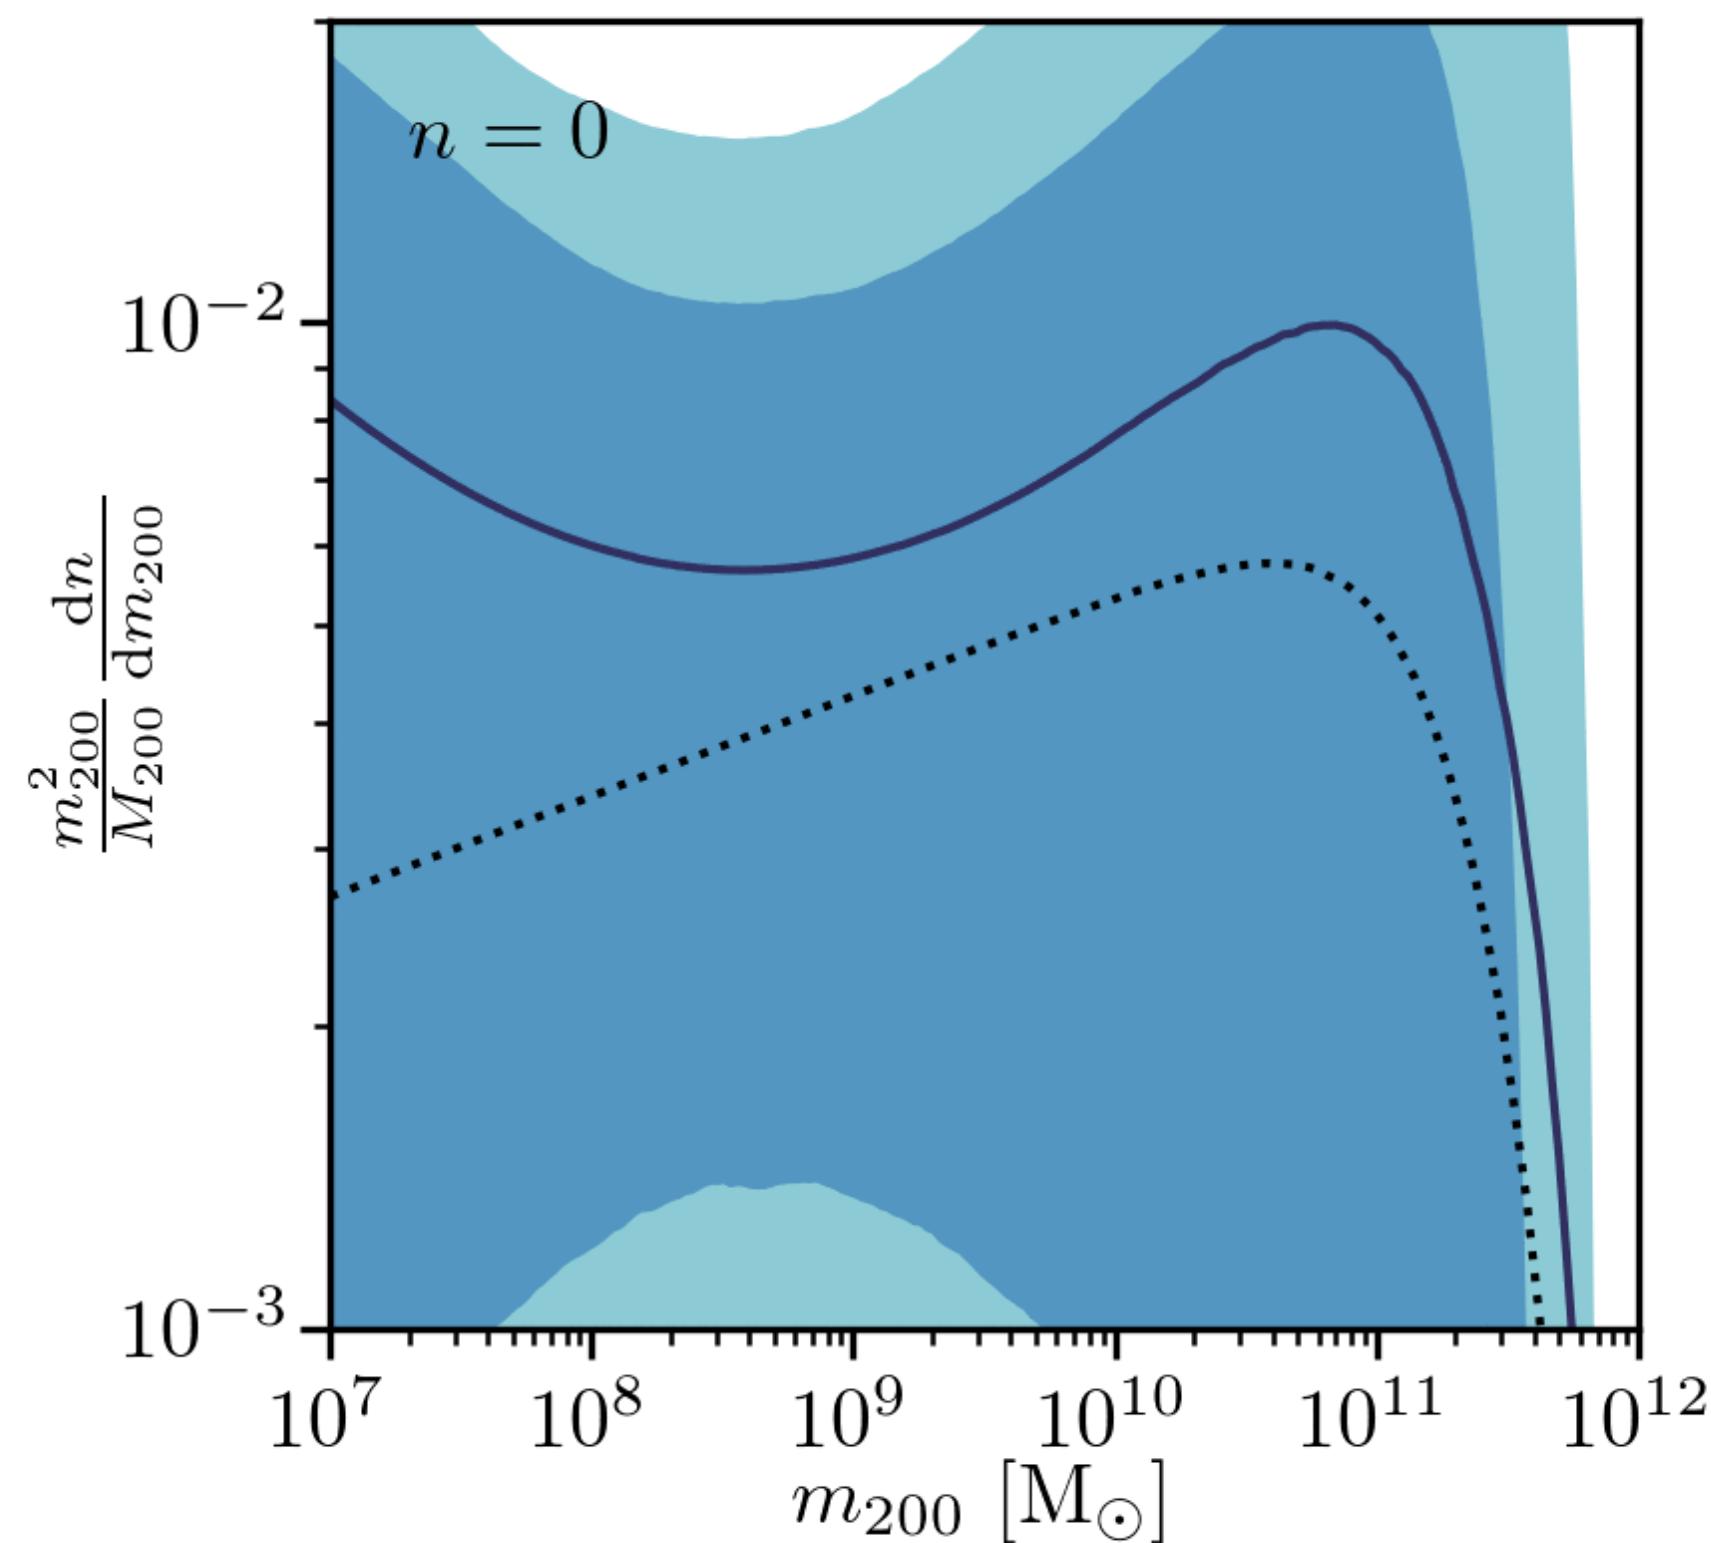
Expected likelihood ratio map



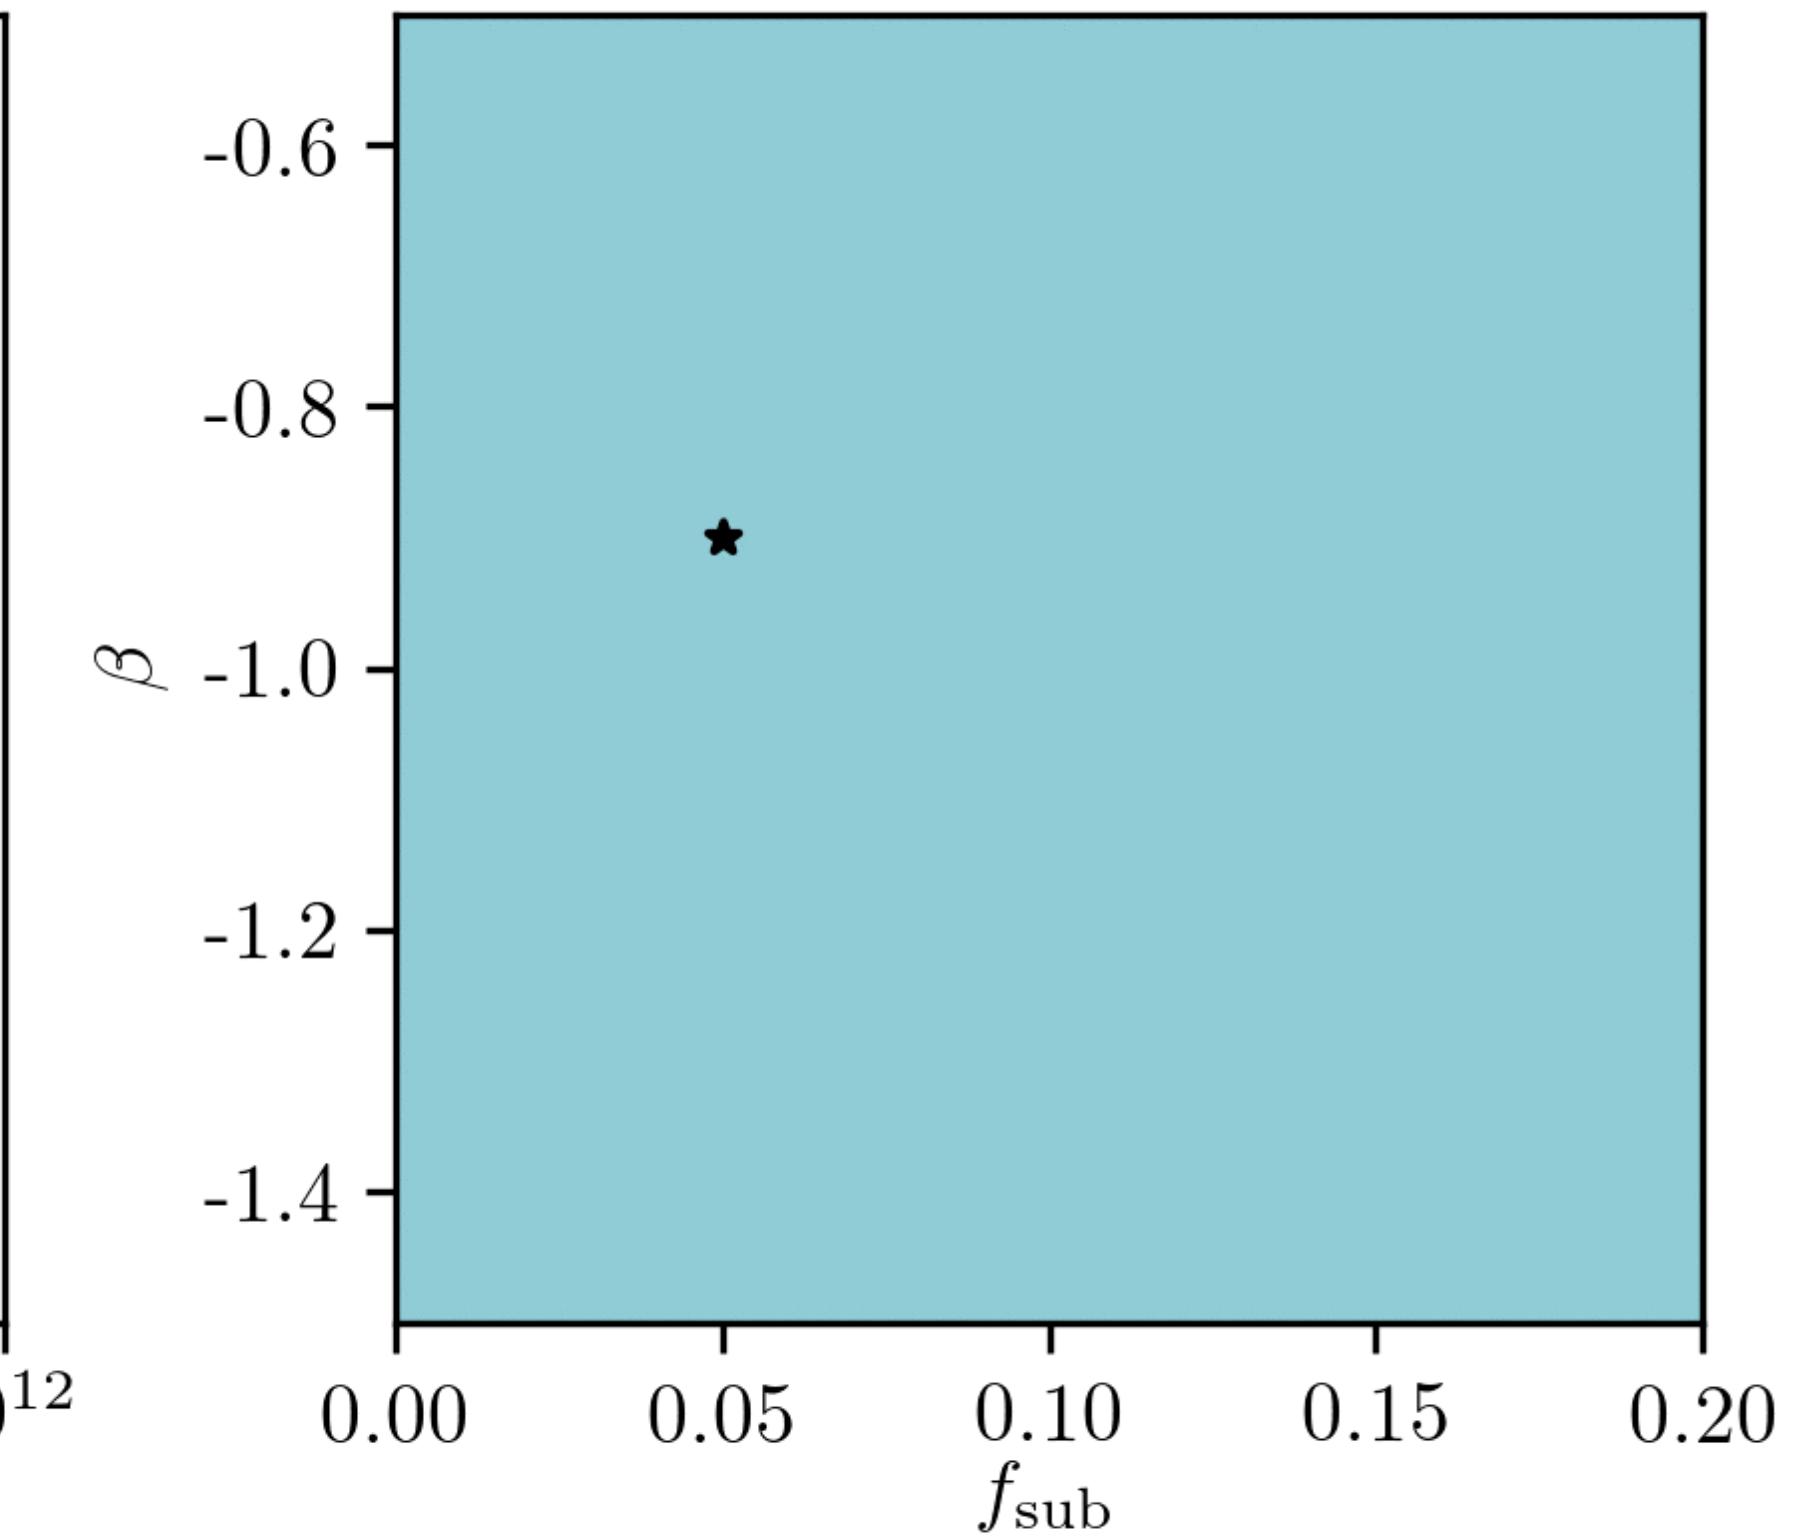
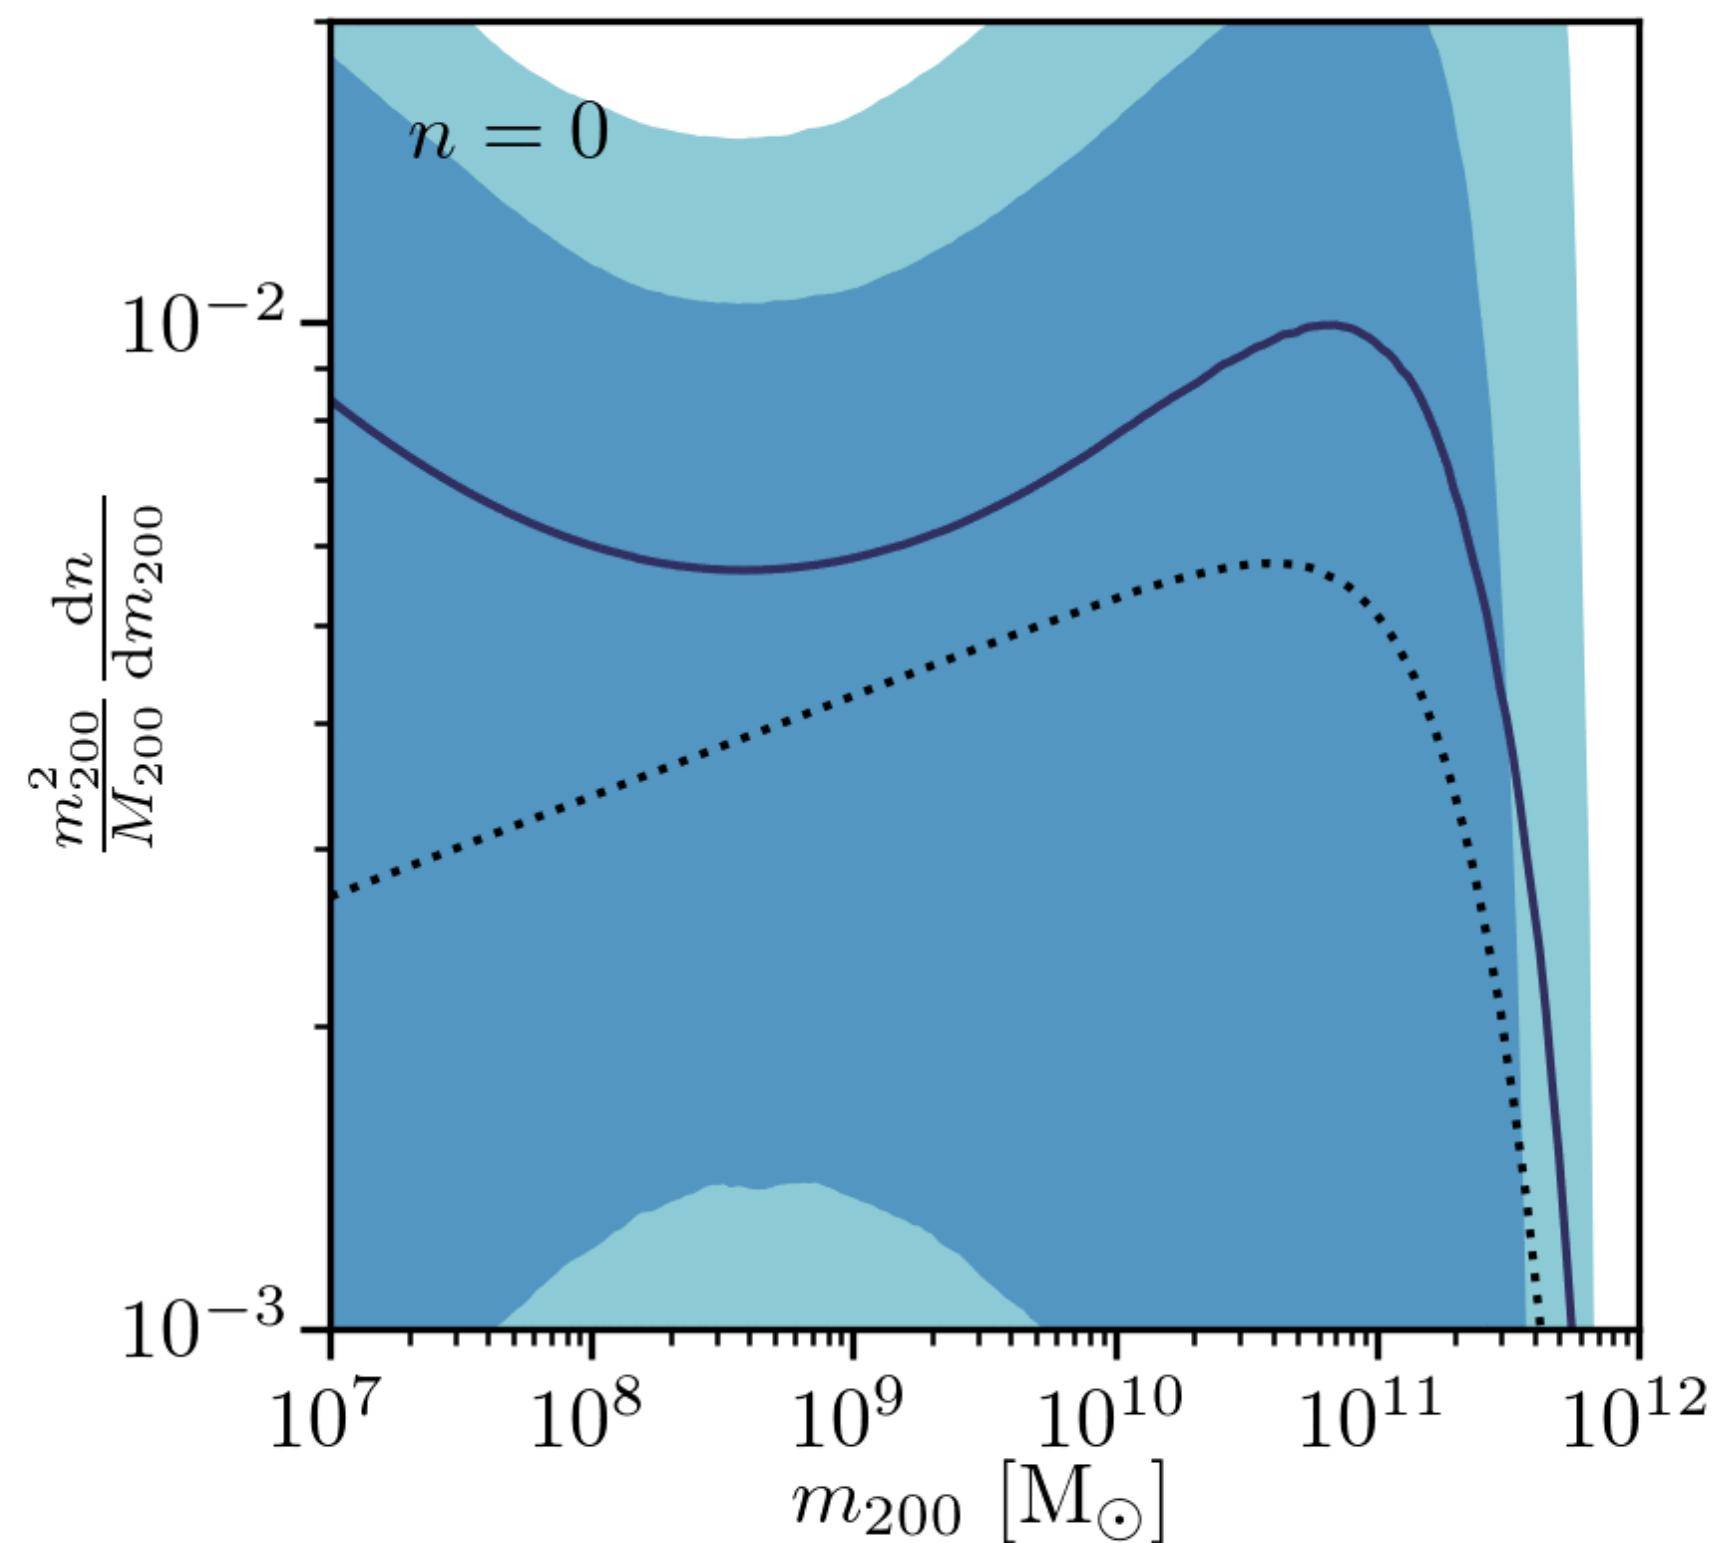
Expected likelihood ratio map



Bayesian inference

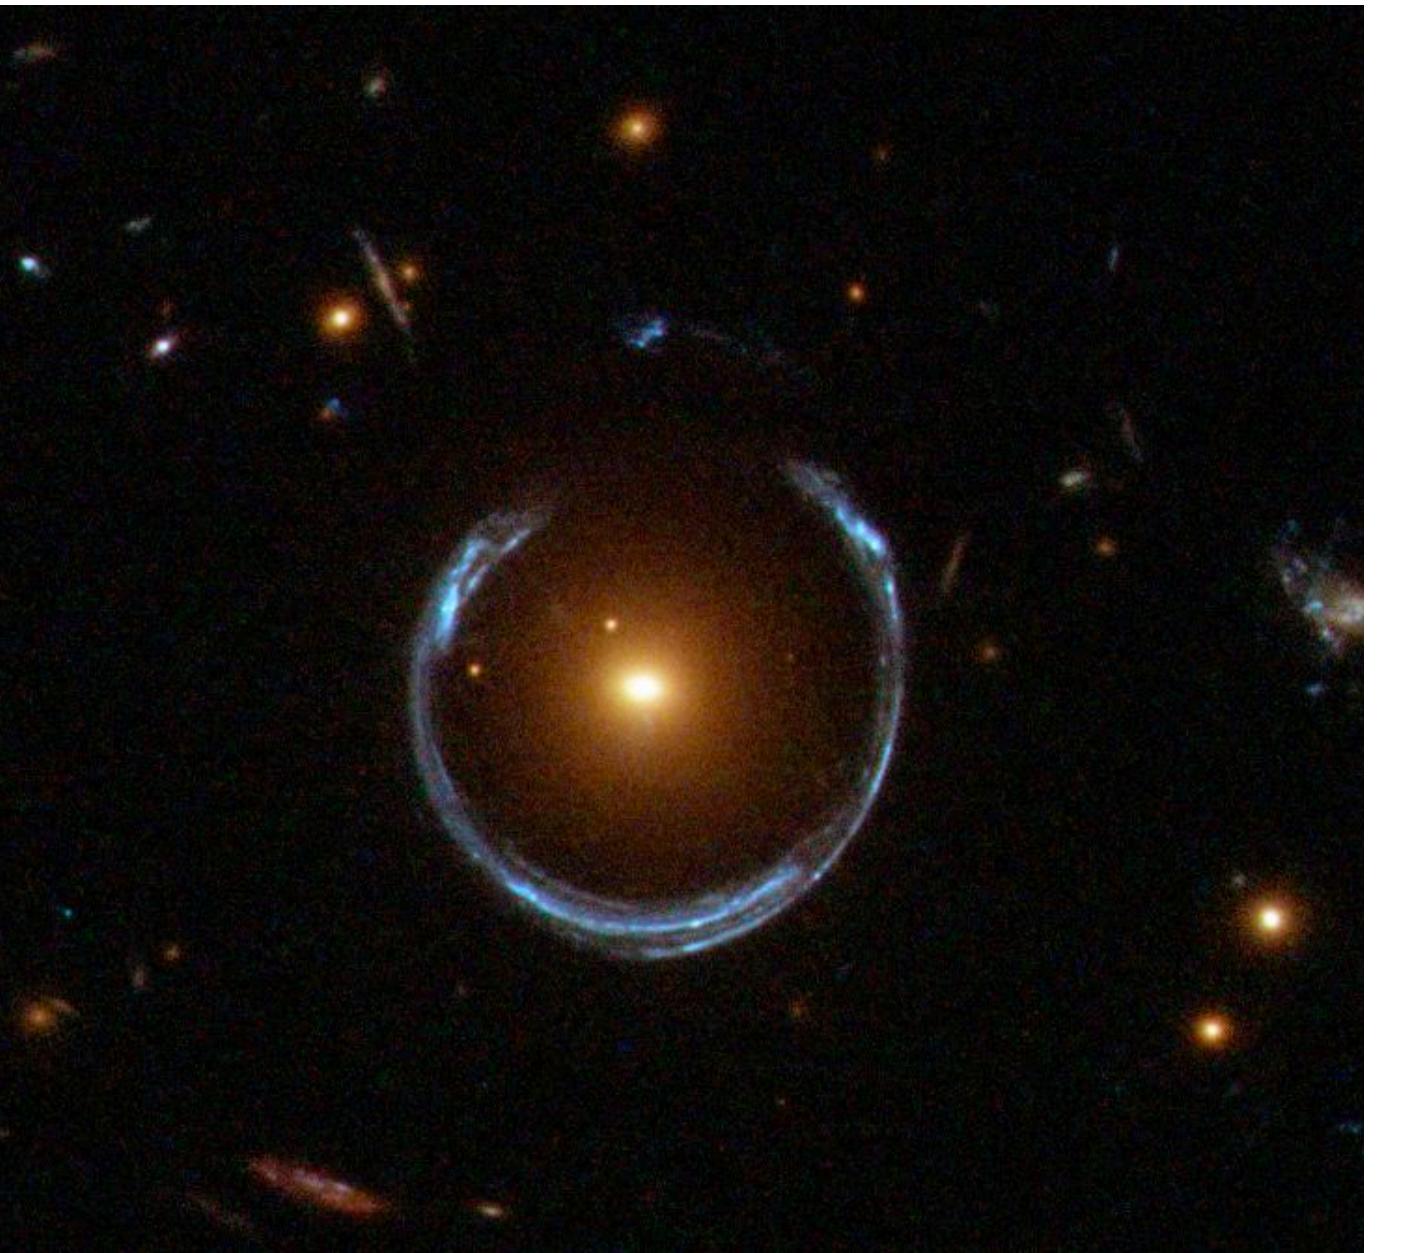


Bayesian inference



All the things we didn't do

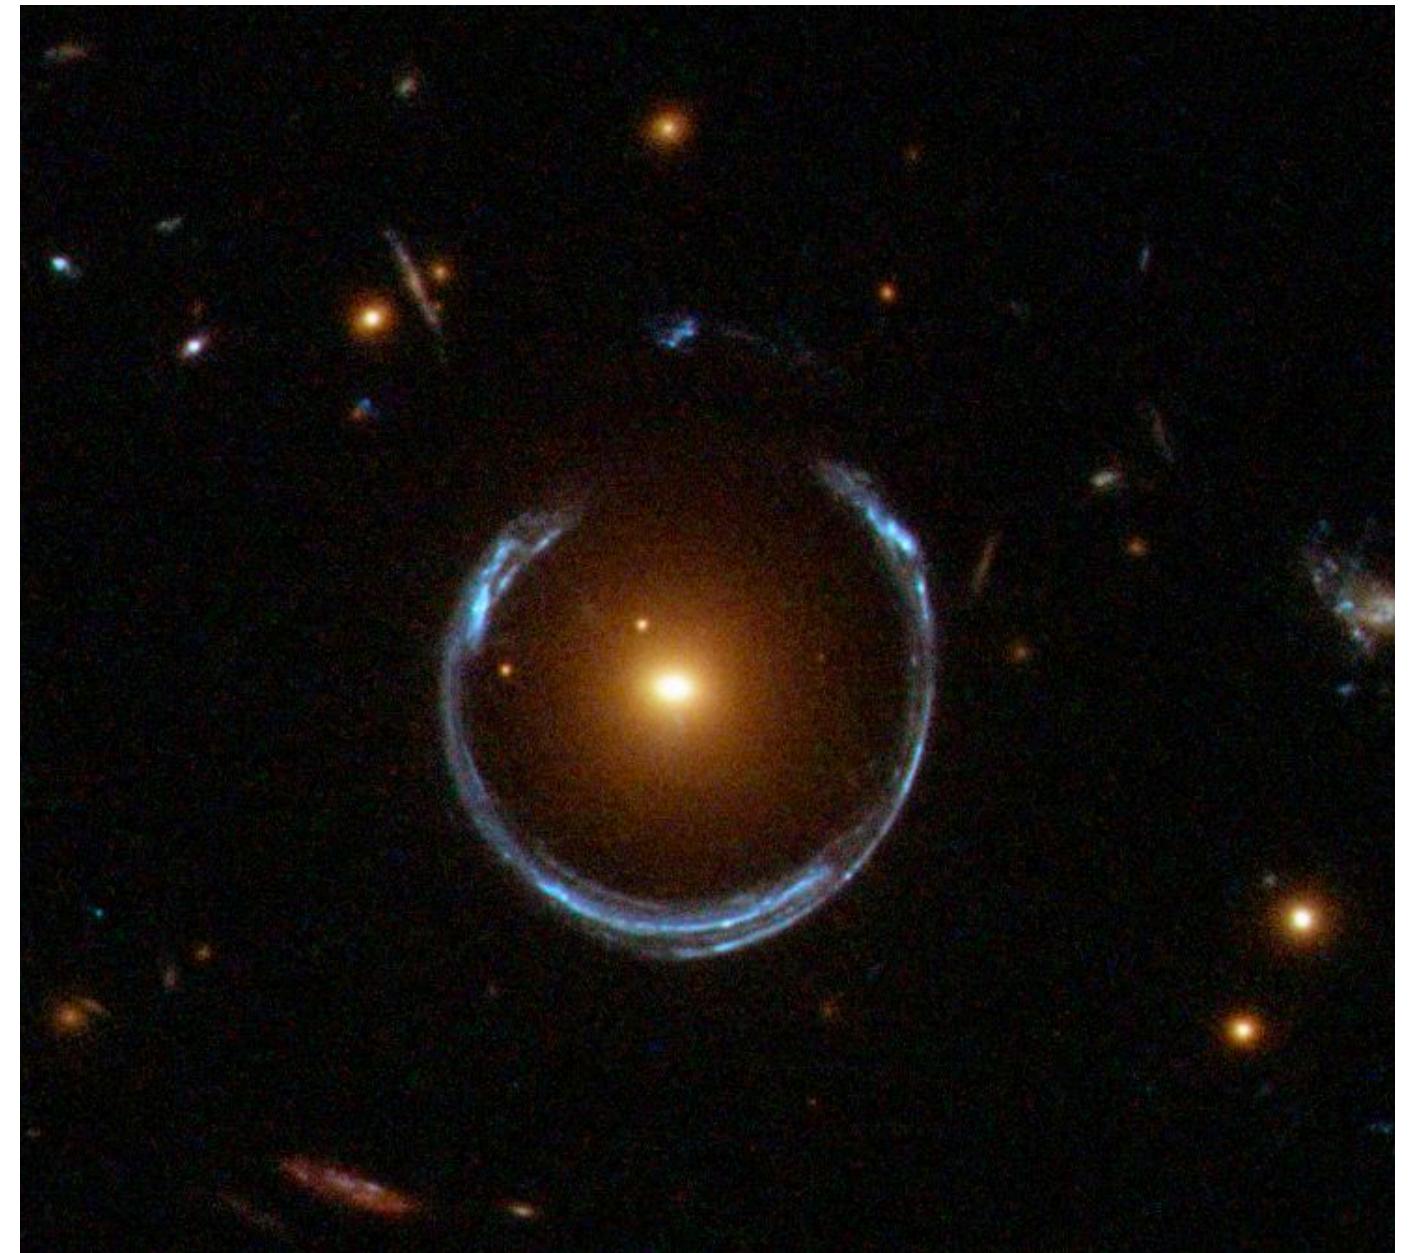
- More involved subhalo mass function
 - Warm DM with DM mass as parameter



[ESA/Hubble/NASA]

All the things we didn't do

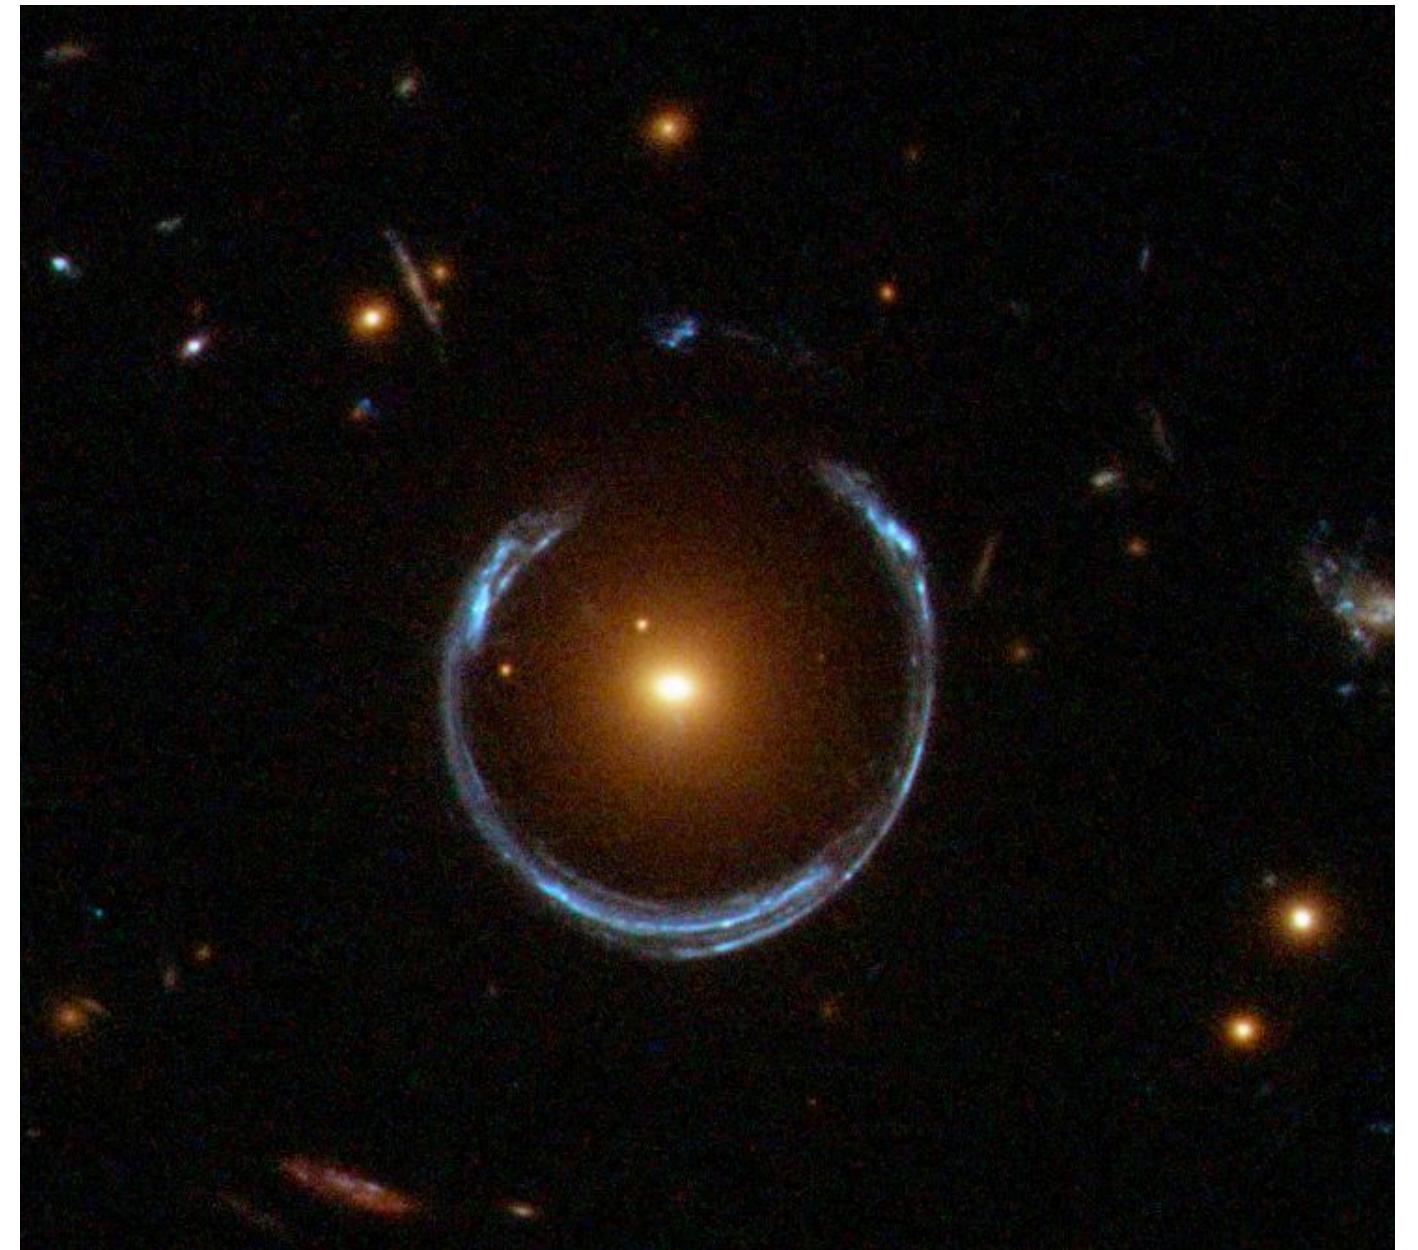
- More involved subhalo mass function
 - Warm DM with DM mass as parameter
- Realistic simulators
 - More diverse source and host galaxies (e.g. data-driven)
 - Realistic subhalo modelling (tidal disruption, redshift dependence...)
 - Line-of-sight substructure
 - Realistic observation model (variable exposure / PSF, multiple bands...)



[ESA/Hubble/NASA]

All the things we didn't do

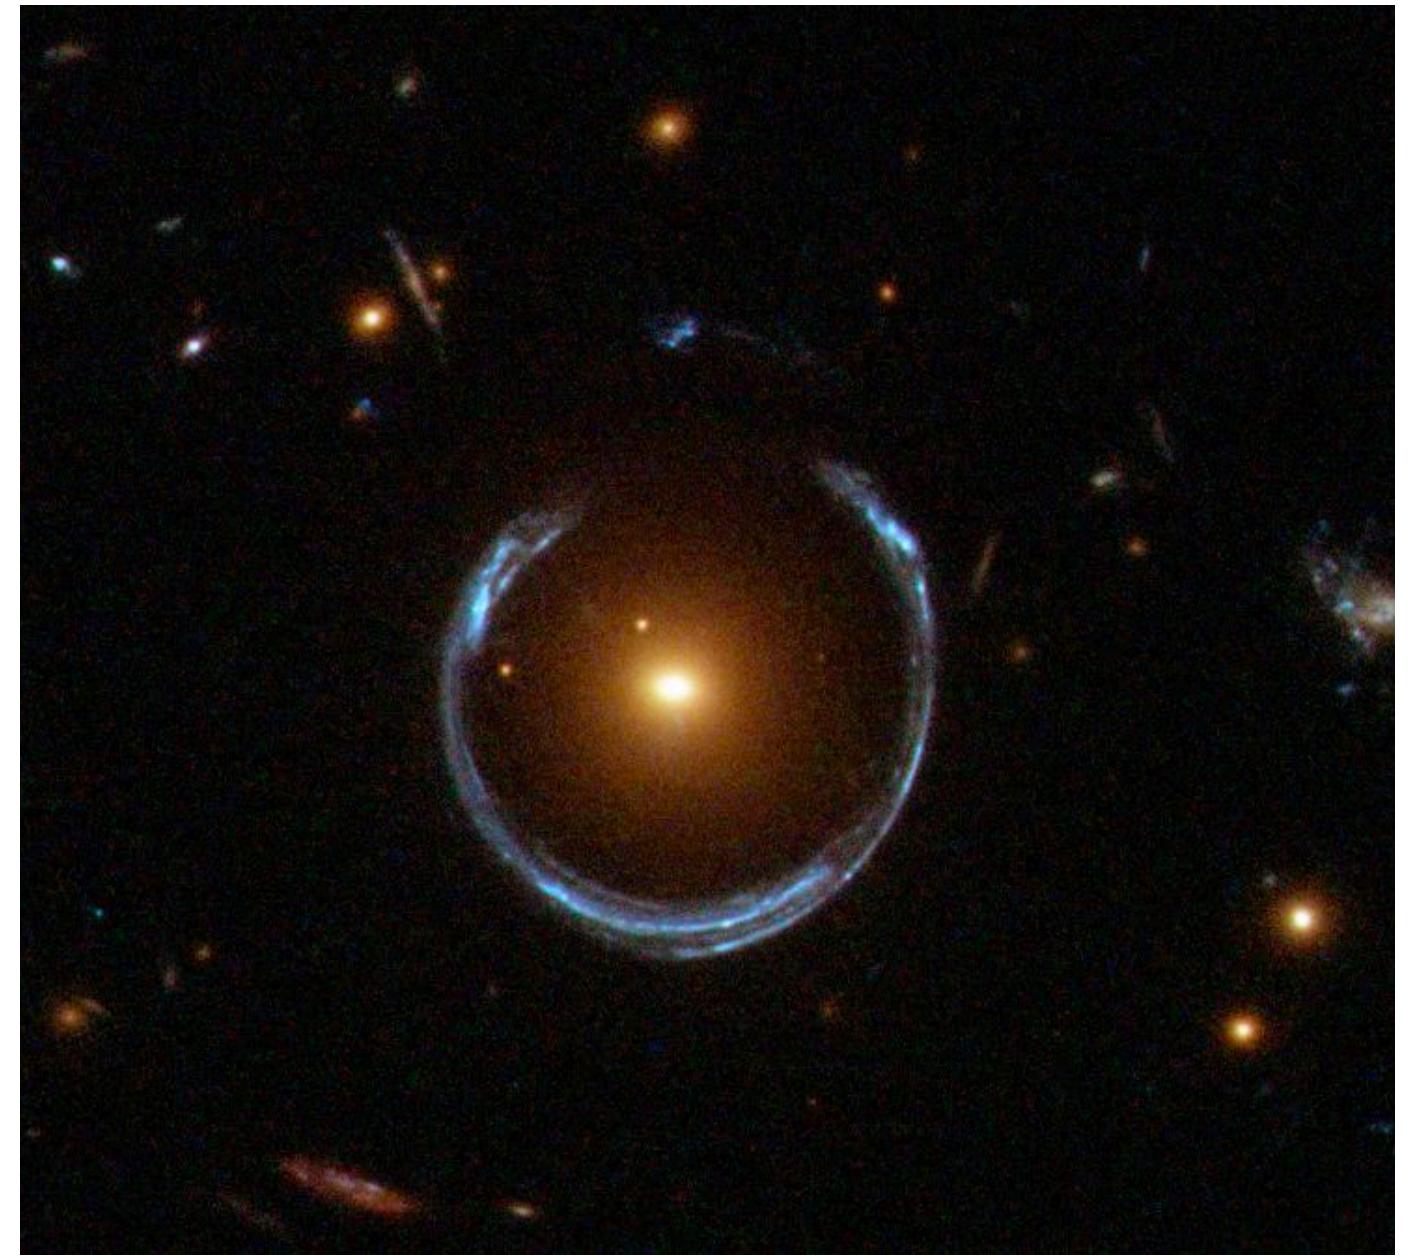
- More involved subhalo mass function
 - Warm DM with DM mass as parameter
- Realistic simulators
 - More diverse source and host galaxies (e.g. data-driven)
 - Realistic subhalo modelling (tidal disruption, redshift dependence...)
 - Line-of-sight substructure
 - Realistic observation model (variable exposure / PSF, multiple bands...)
- Use auxiliary information during inference



[ESA/Hubble/NASA]

All the things we didn't do

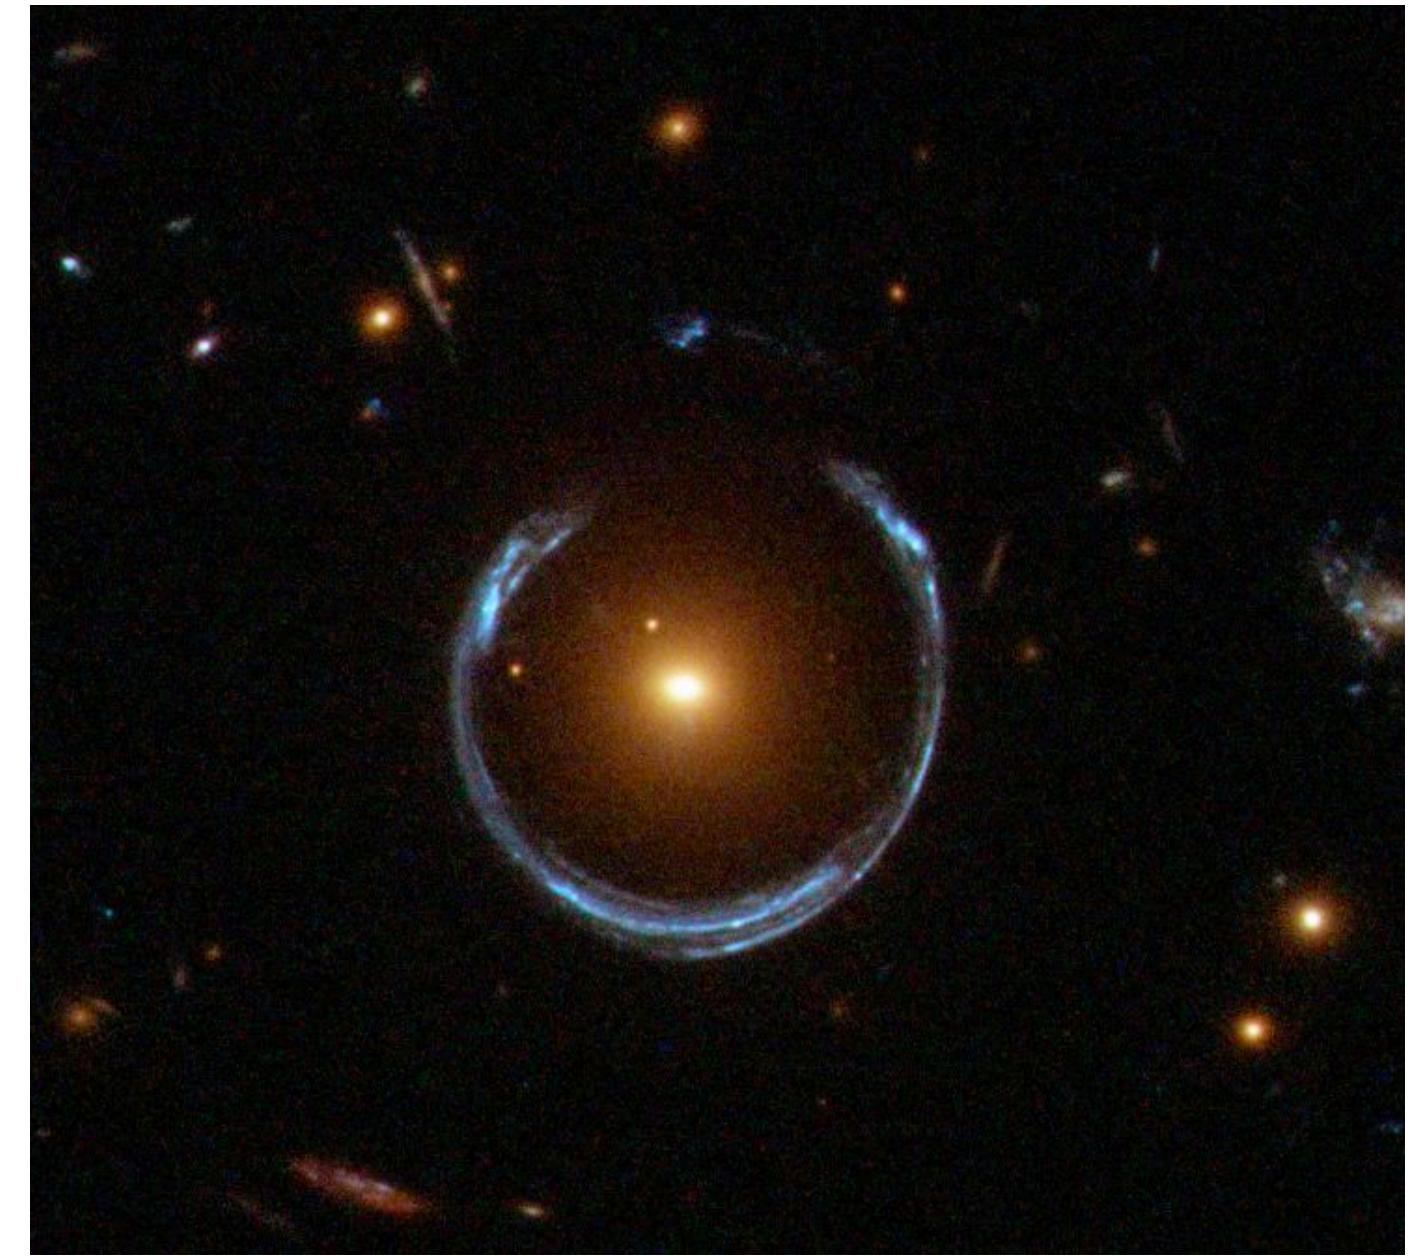
- More involved subhalo mass function
 - Warm DM with DM mass as parameter
- Realistic simulators
 - More diverse source and host galaxies (e.g. data-driven)
 - Realistic subhalo modelling (tidal disruption, redshift dependence...)
 - Line-of-sight substructure
 - Realistic observation model (variable exposure / PSF, multiple bands...)
- Use auxiliary information during inference
- Evaluation on real data



[ESA/Hubble/NASA]

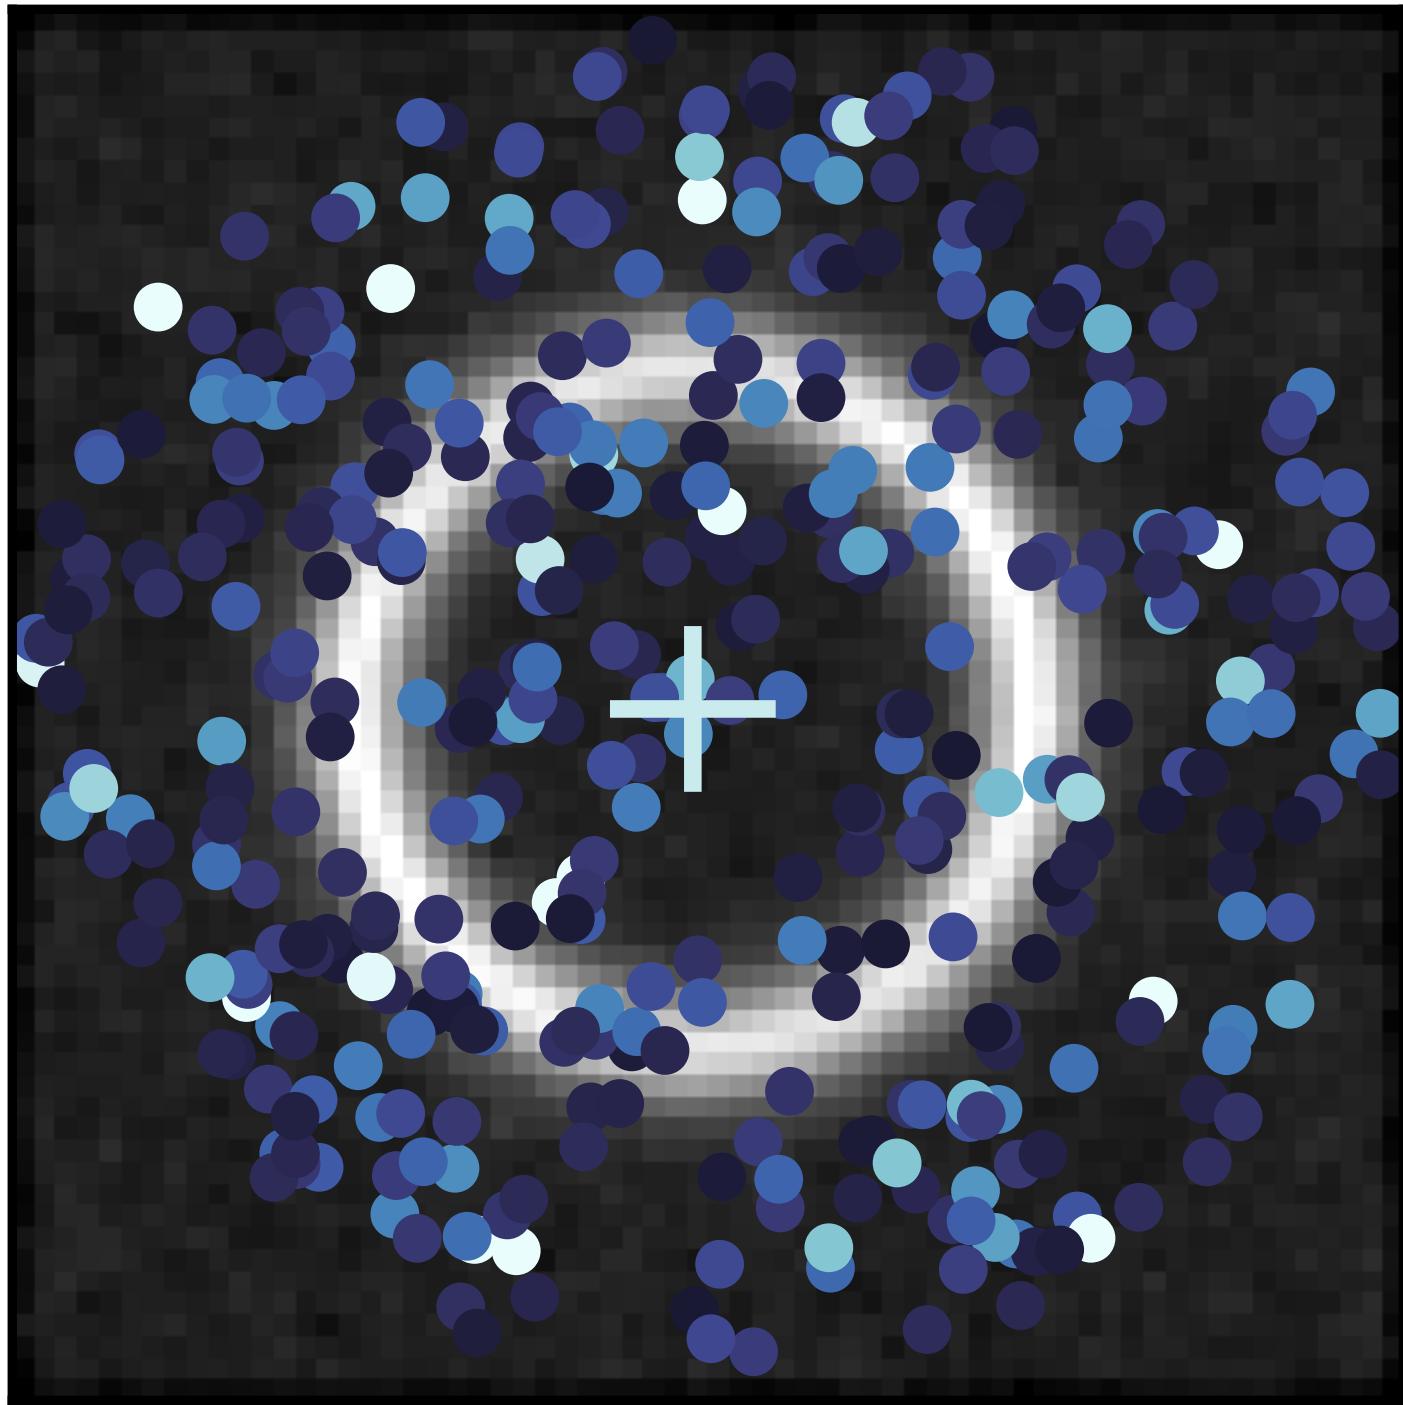
All the things we didn't do

- More involved subhalo mass function
 - Warm DM with DM mass as parameter
- Realistic simulators
 - More diverse source and host galaxies (e.g. data-driven)
 - Realistic subhalo modelling (tidal disruption, redshift dependence...)
 - Line-of-sight substructure
 - Realistic observation model (variable exposure / PSF, multiple bands...)
- Use auxiliary information during inference
- Evaluation on real data

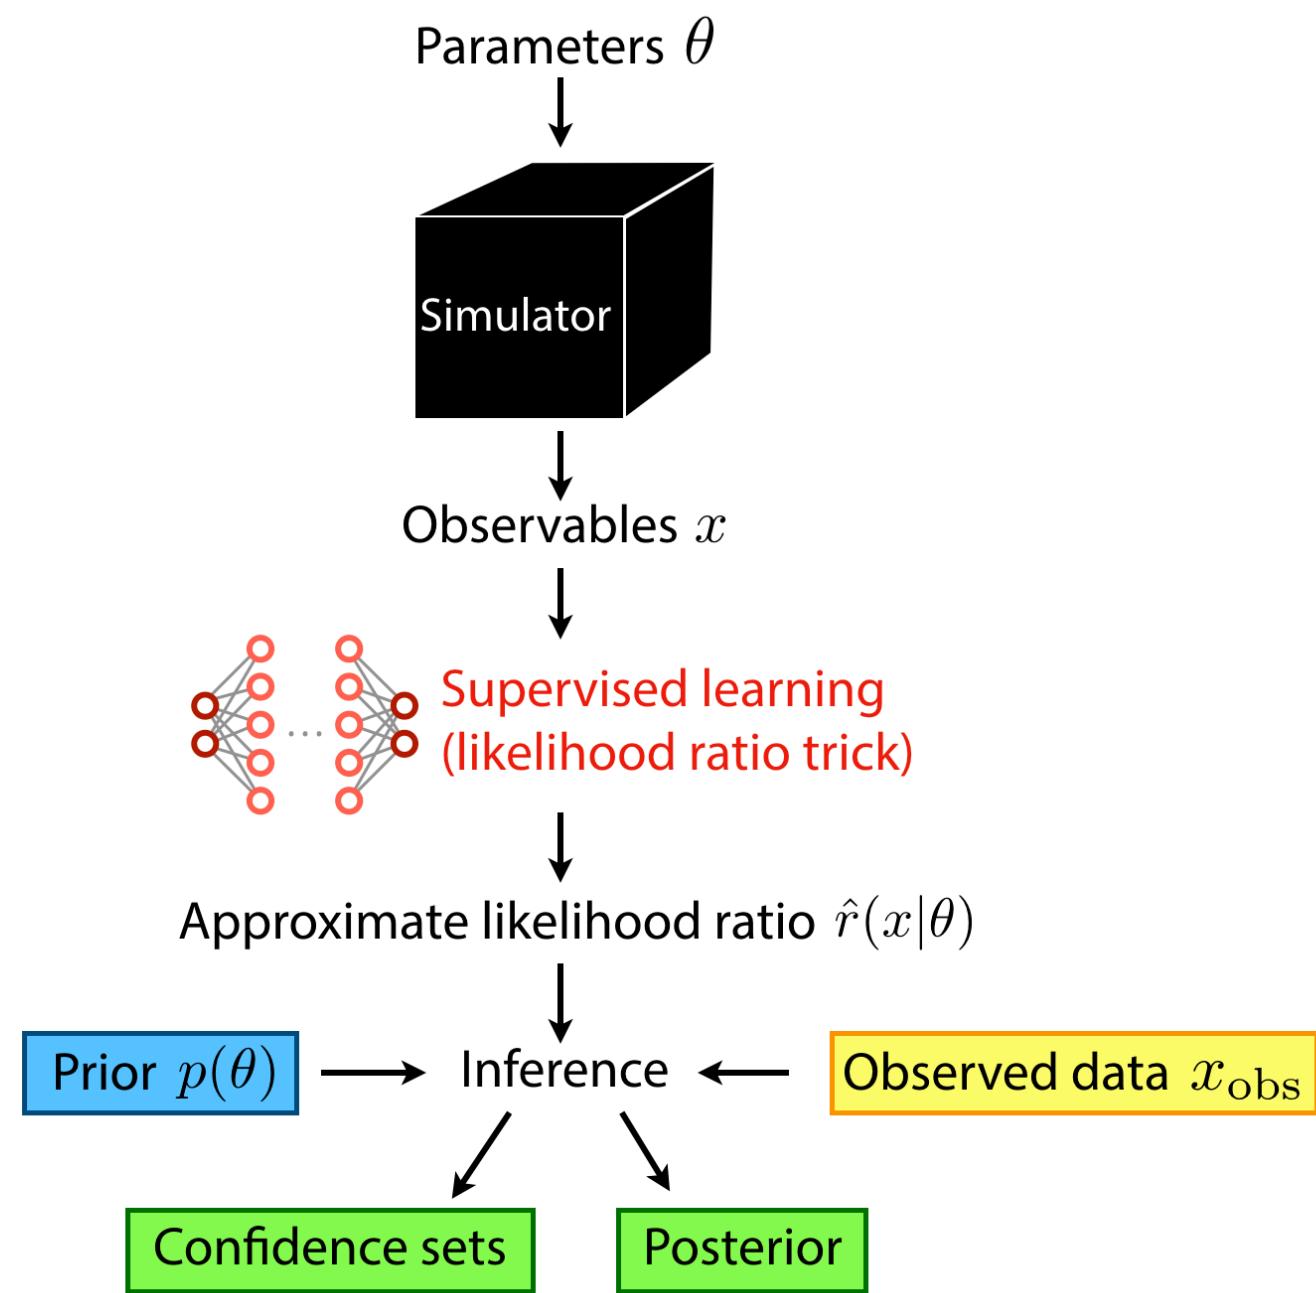


[ESA/Hubble/NASA]

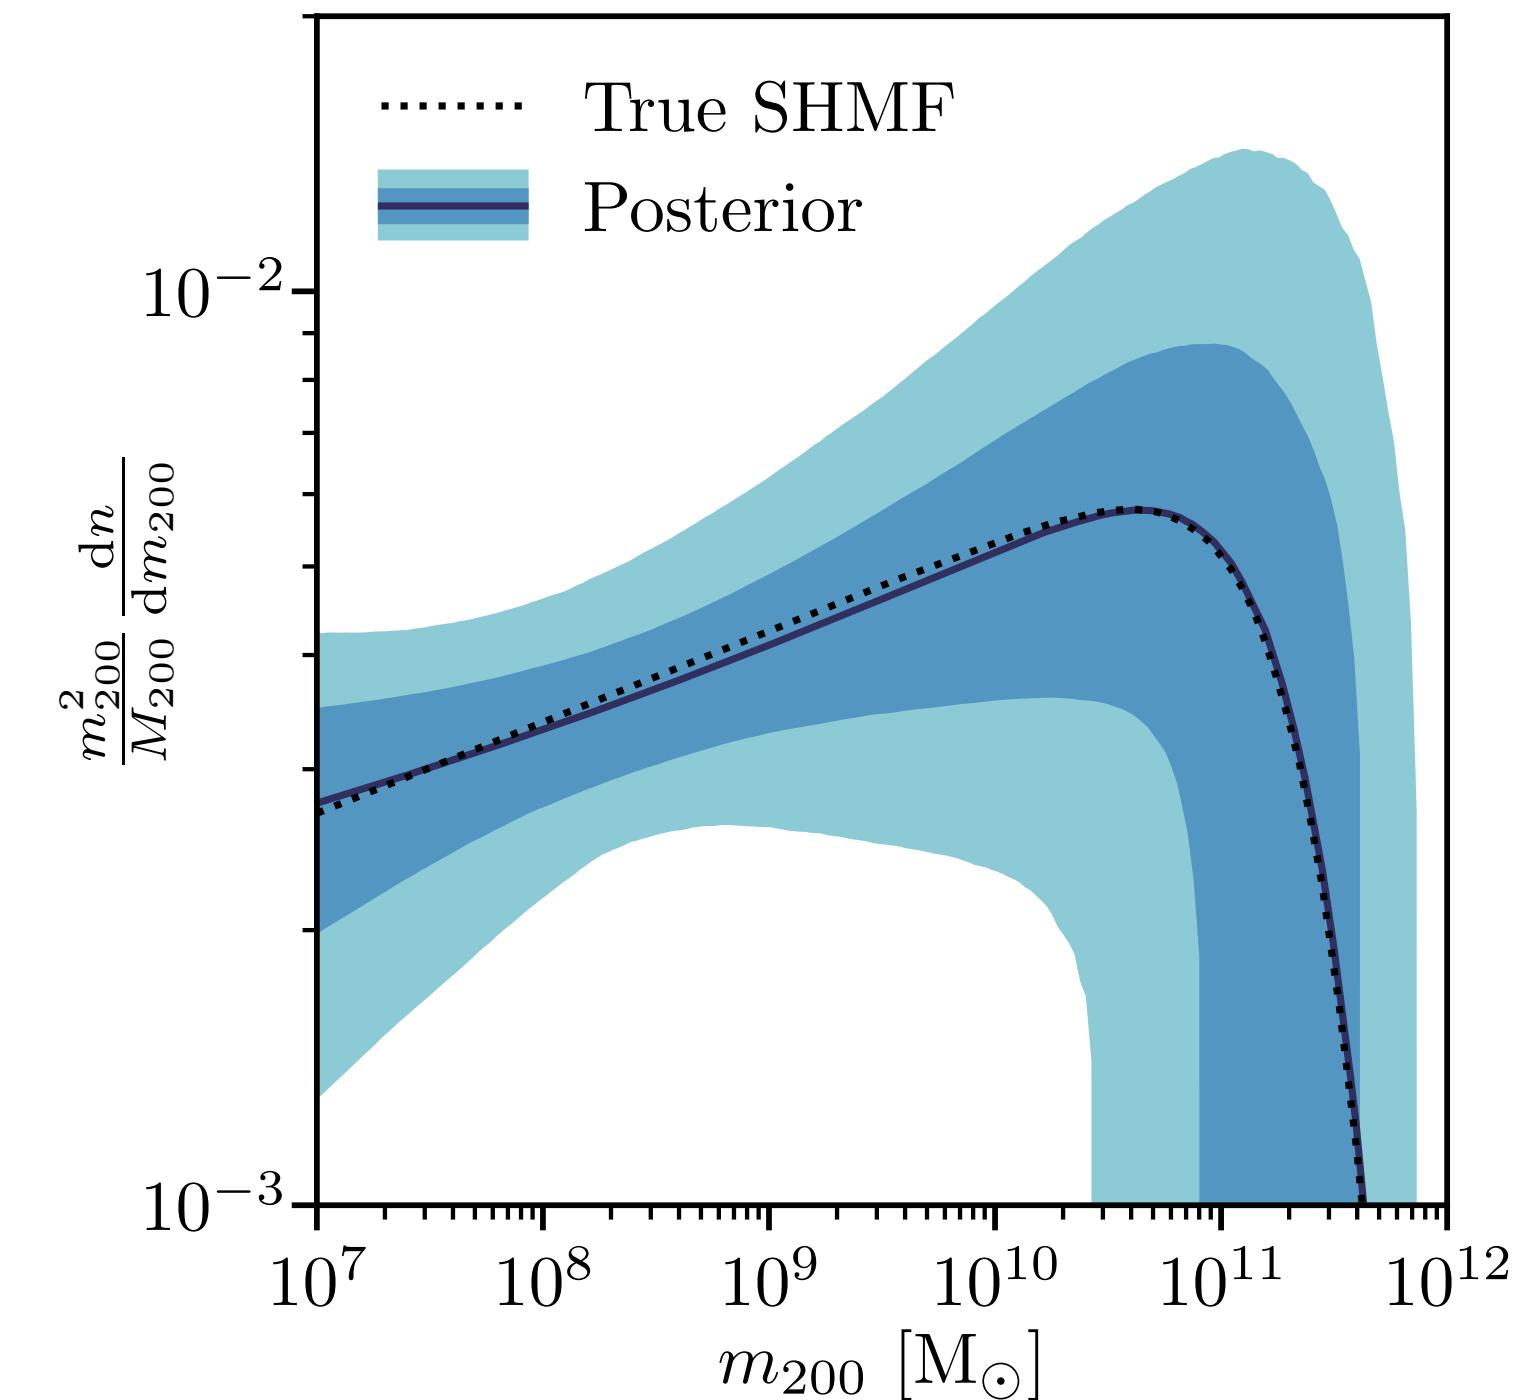
⇒ Our method should scale to a realistic setting, but will require more simulations and careful sanity checks



The small-scale structure of Dark Matter affects strong lensing, but teasing out this information is challenging.



New simulation-based inference algorithms based on machine learning can solve this problem in a scalable way.



In a proof-of-concept study, they allow us to constrain Dark Matter substructure from Euclid-like lensing observations.

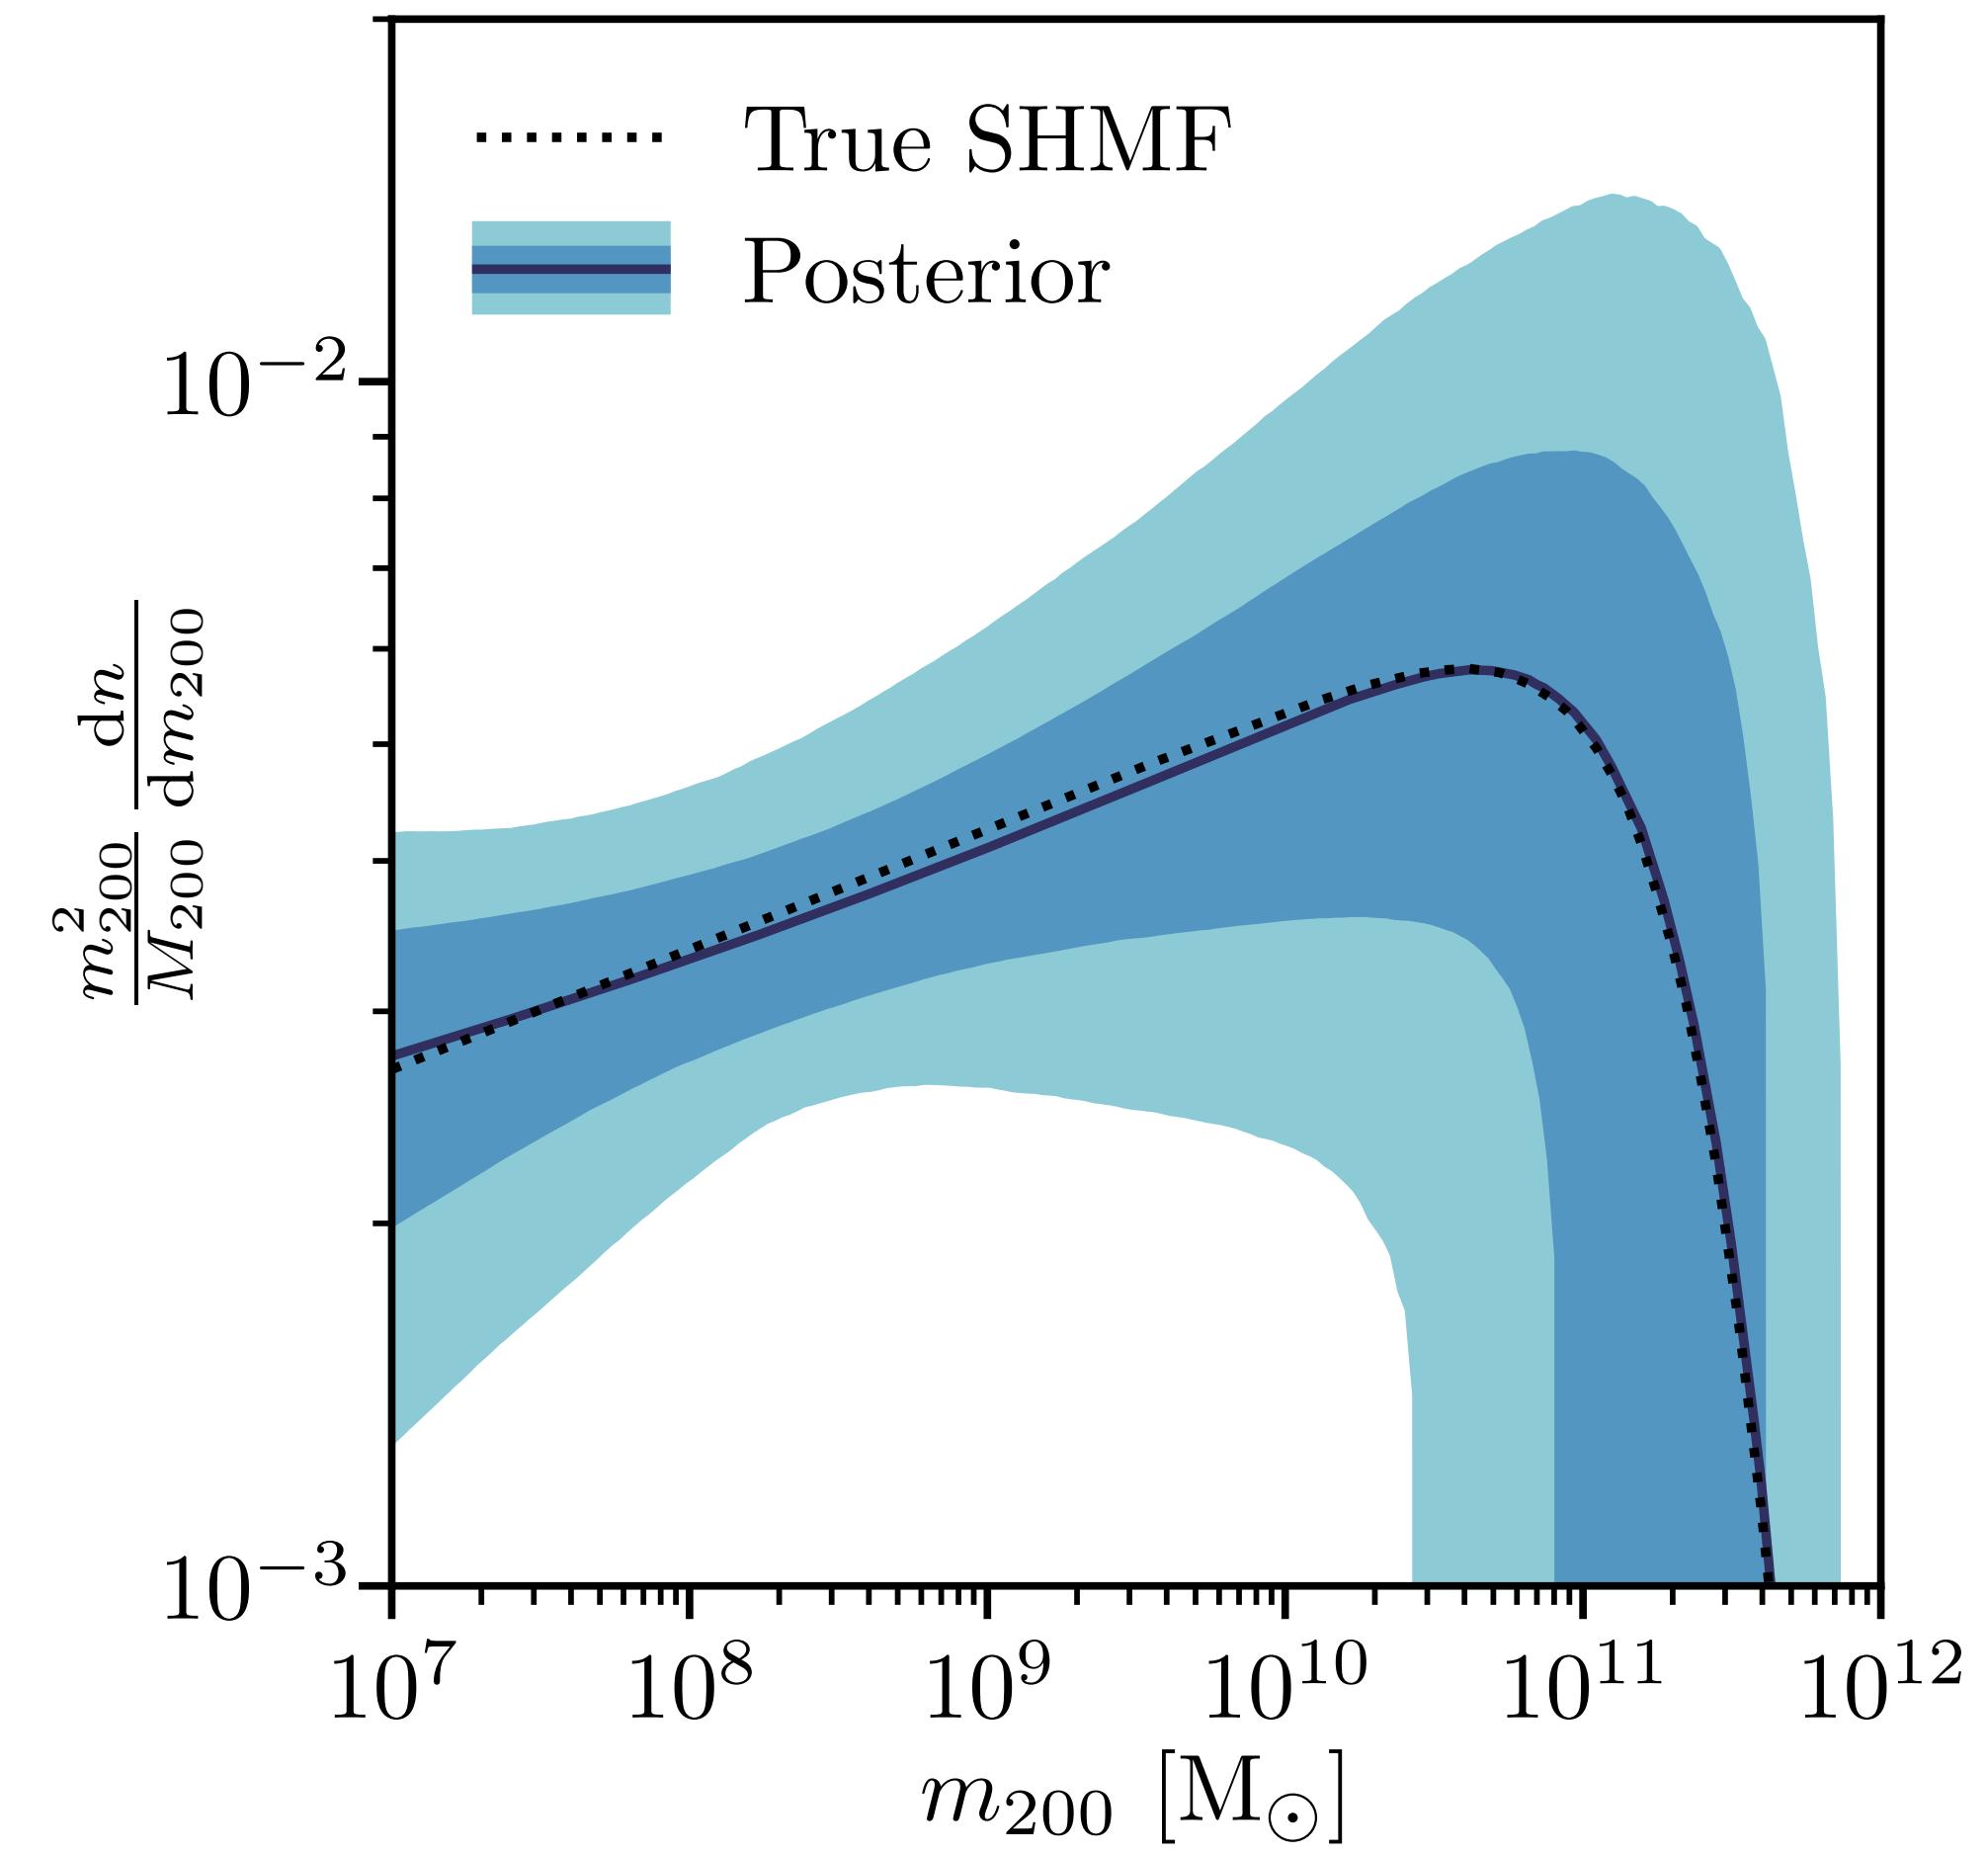
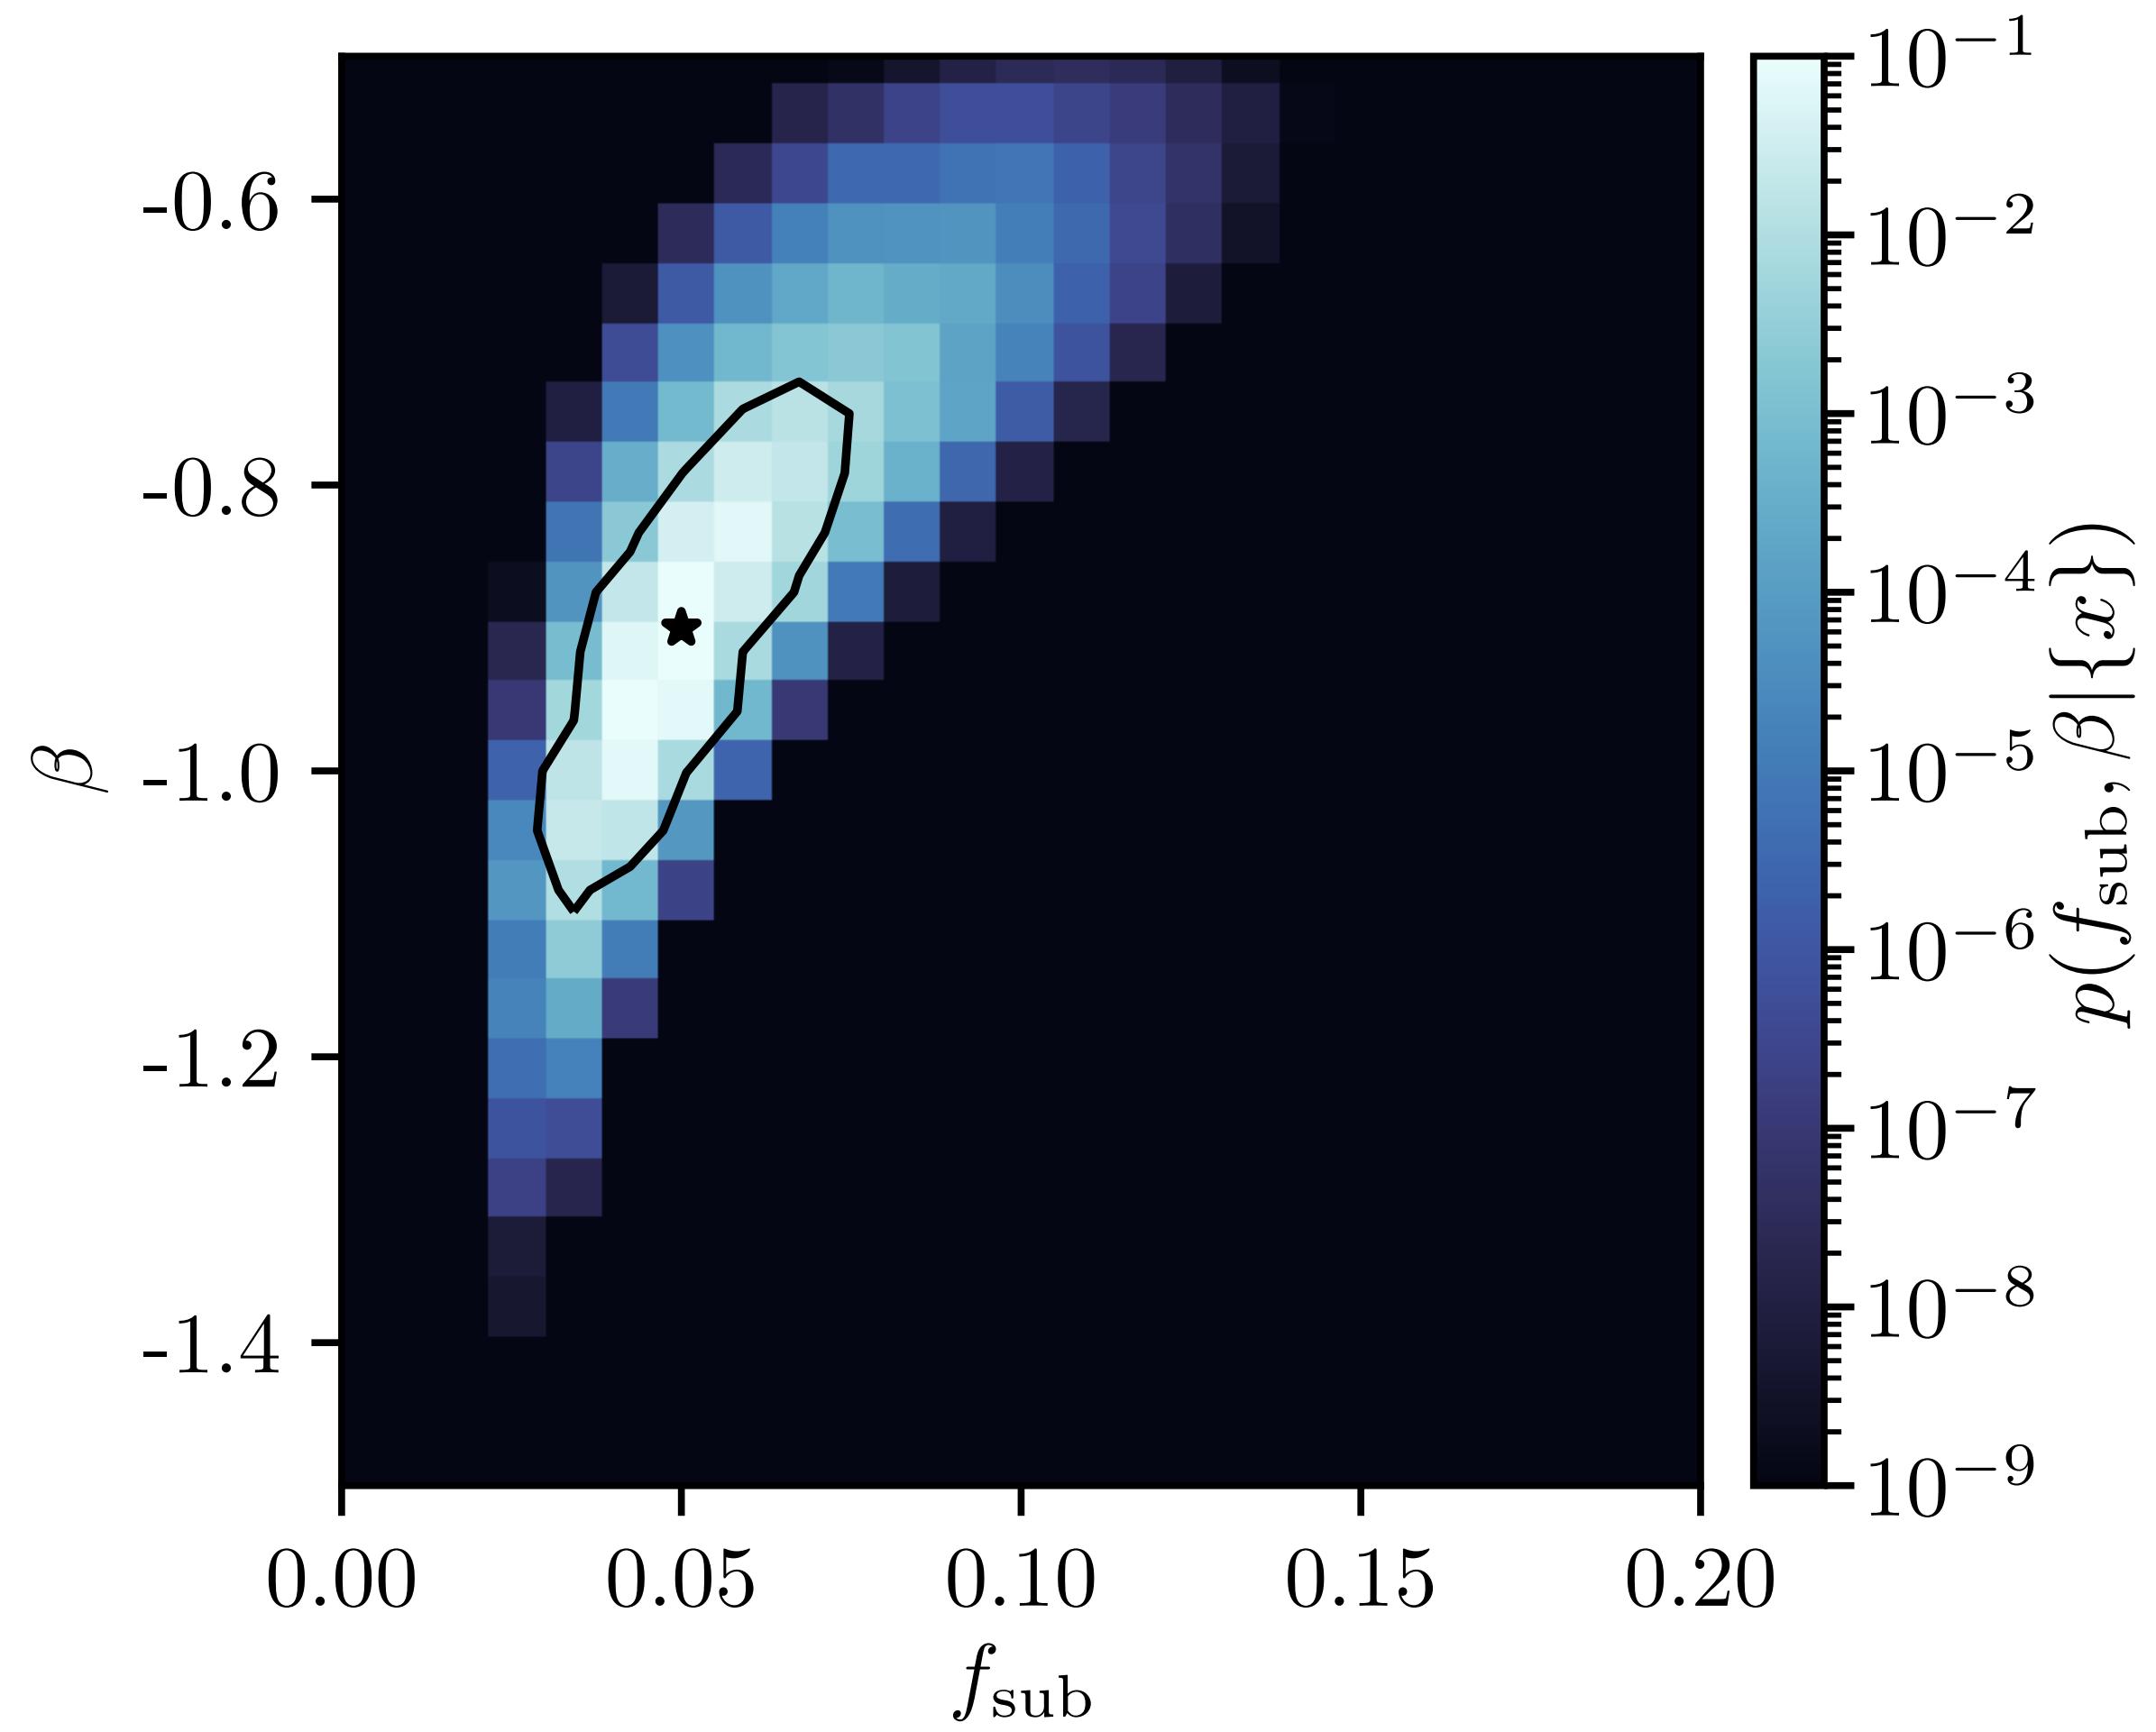
Strong lensing: JB, S. Mishra-Sharma, J. Hermans, G. Louppe, K. Cranmer [1909.02005](#)

Broader review of simulation-based inference: K. Cranmer, JB, G. Louppe [1911.01429](#)

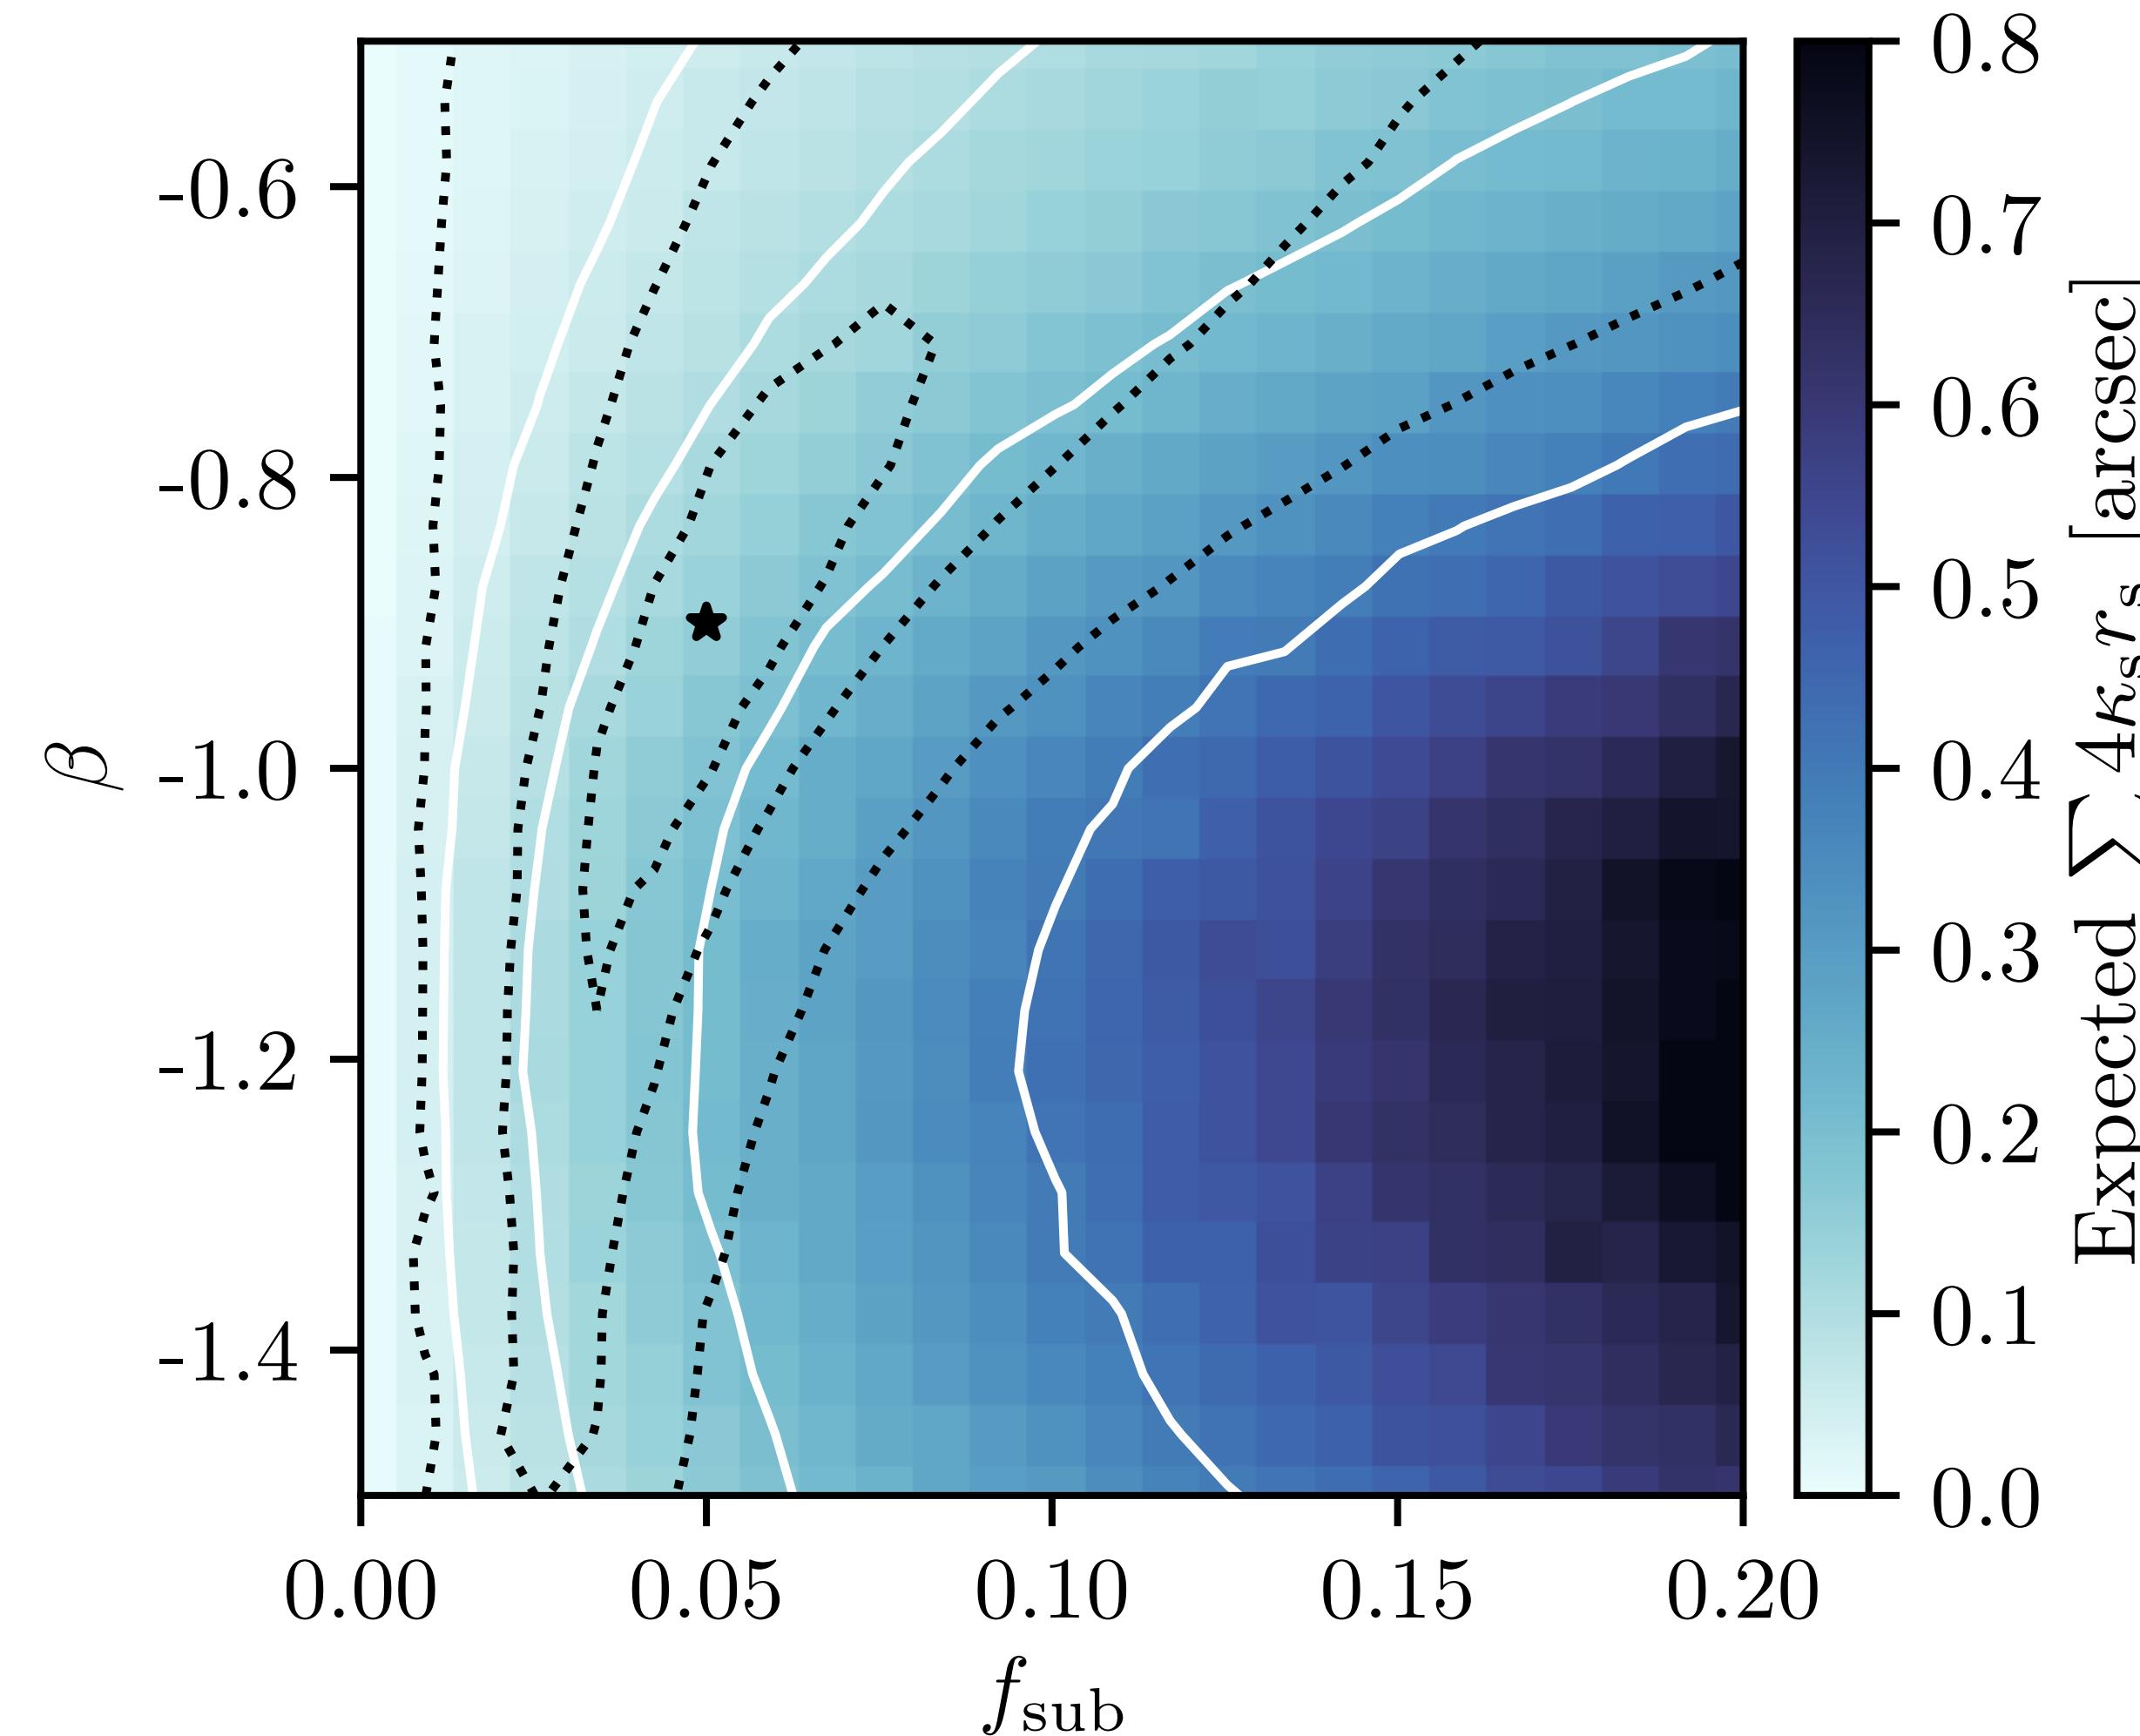
“Hooray! Question mark?”

— Todd Chavez

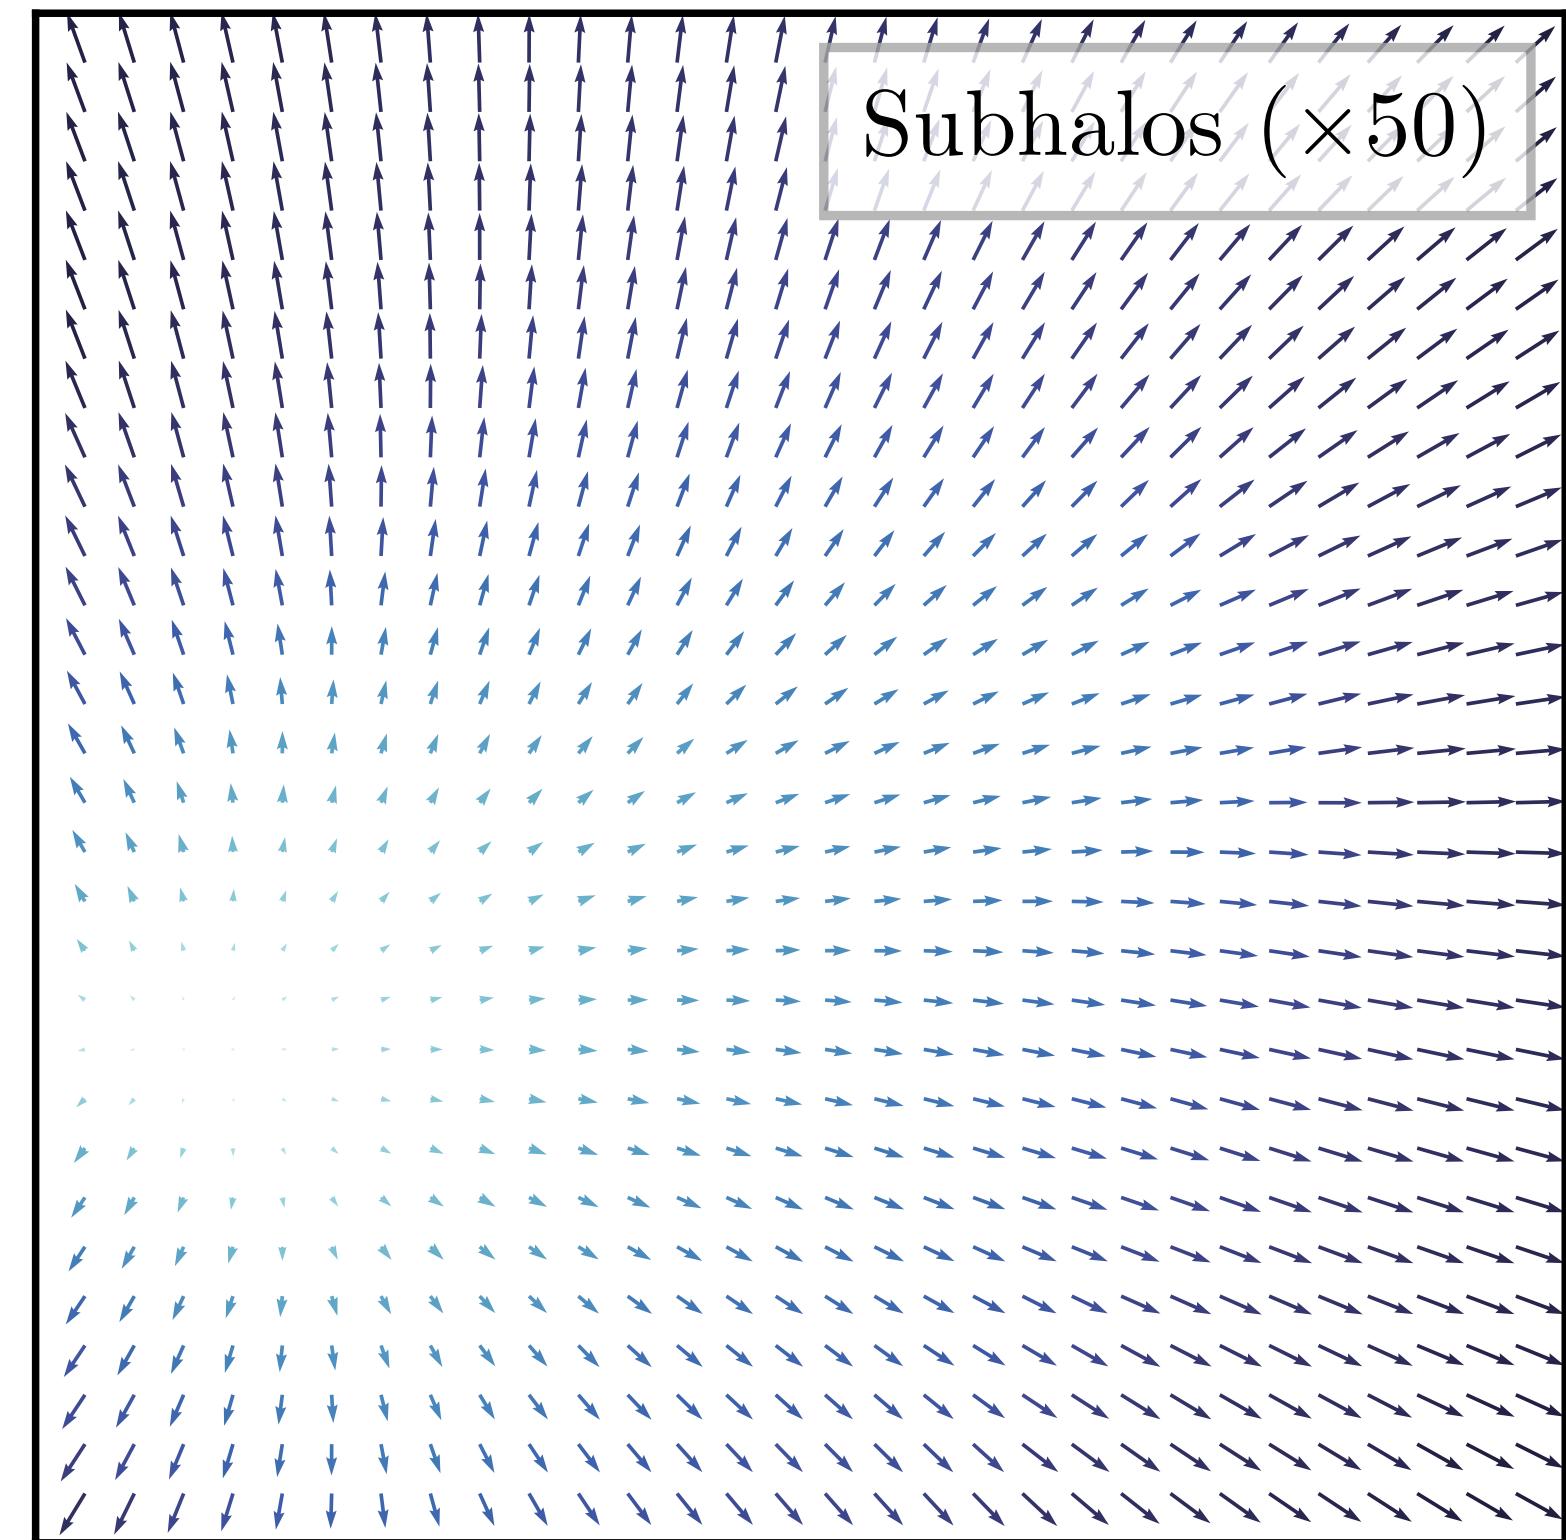
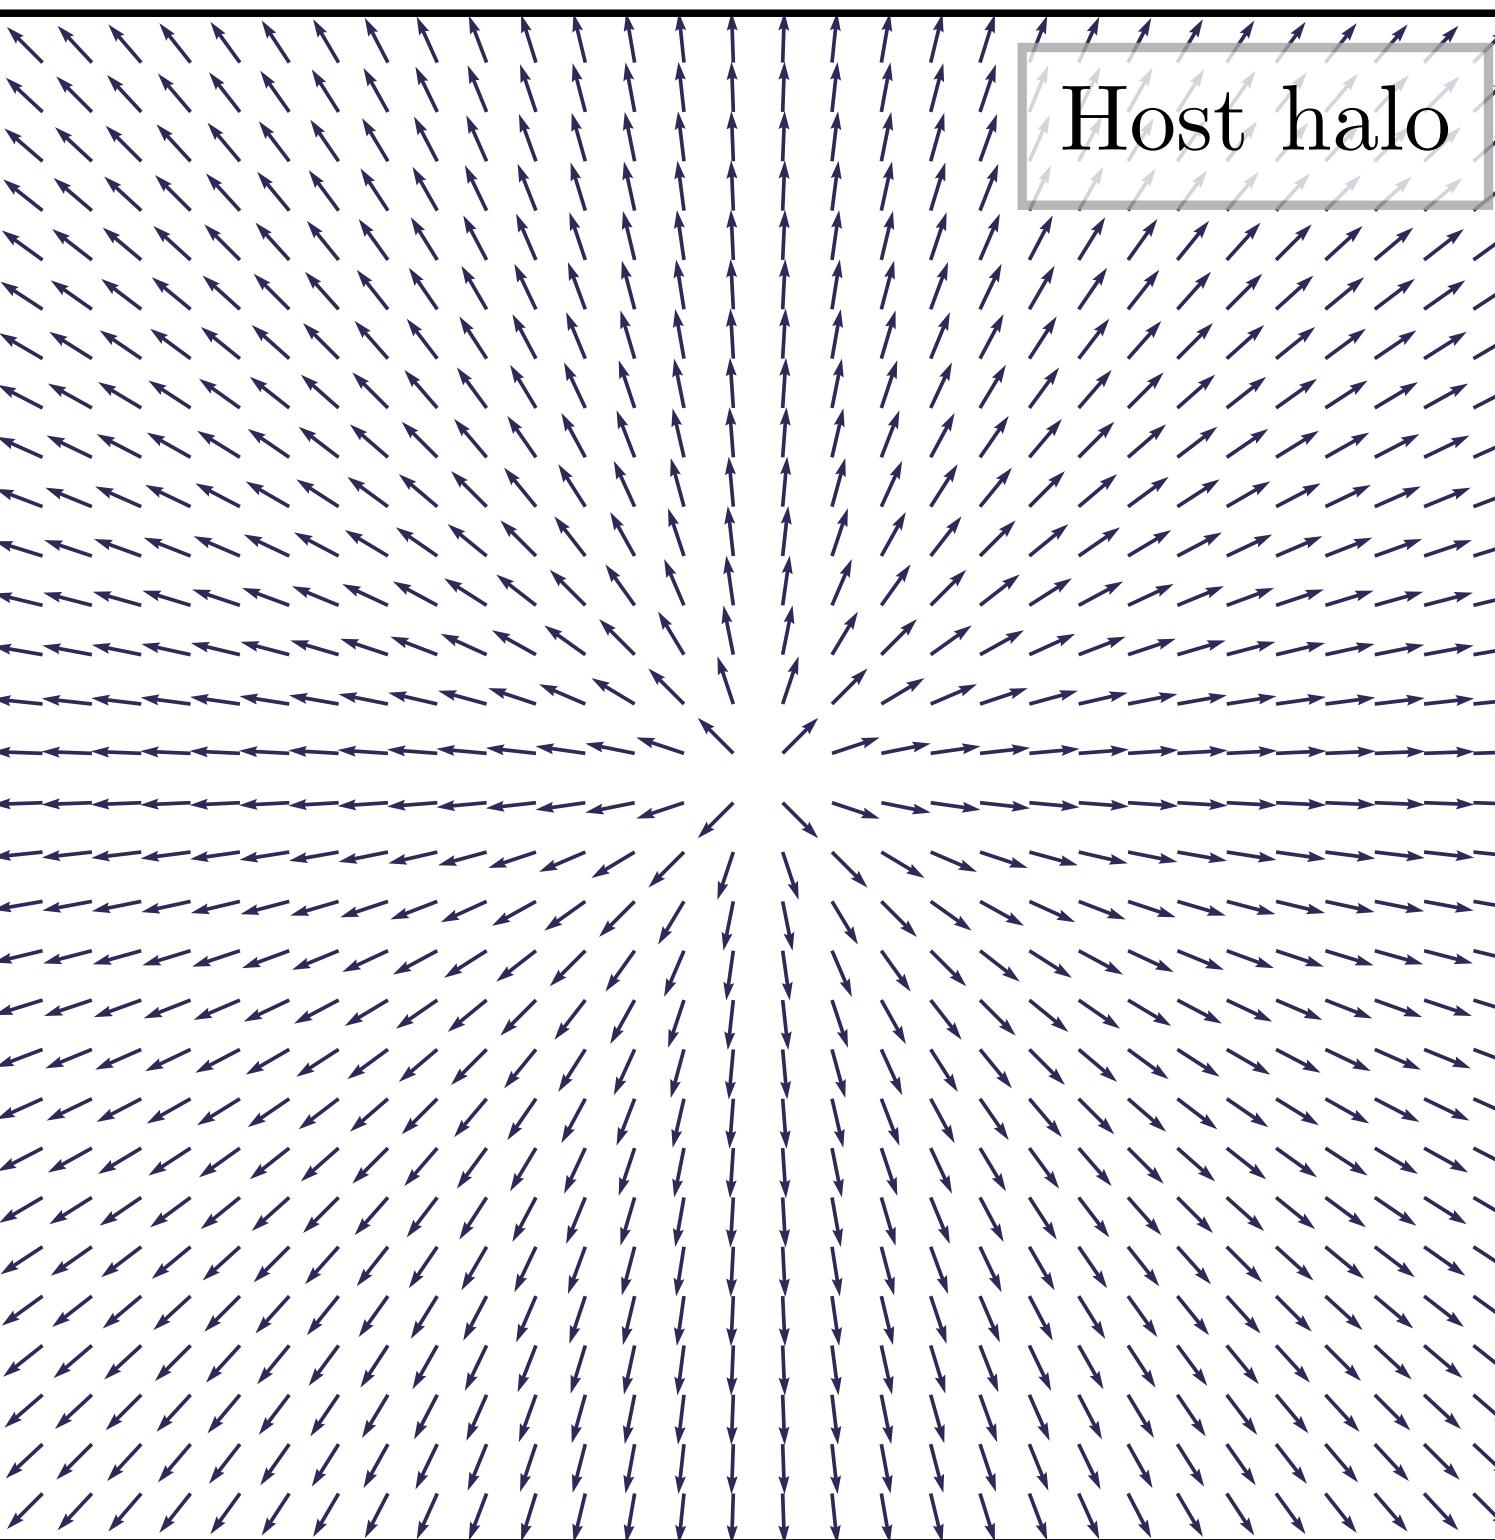
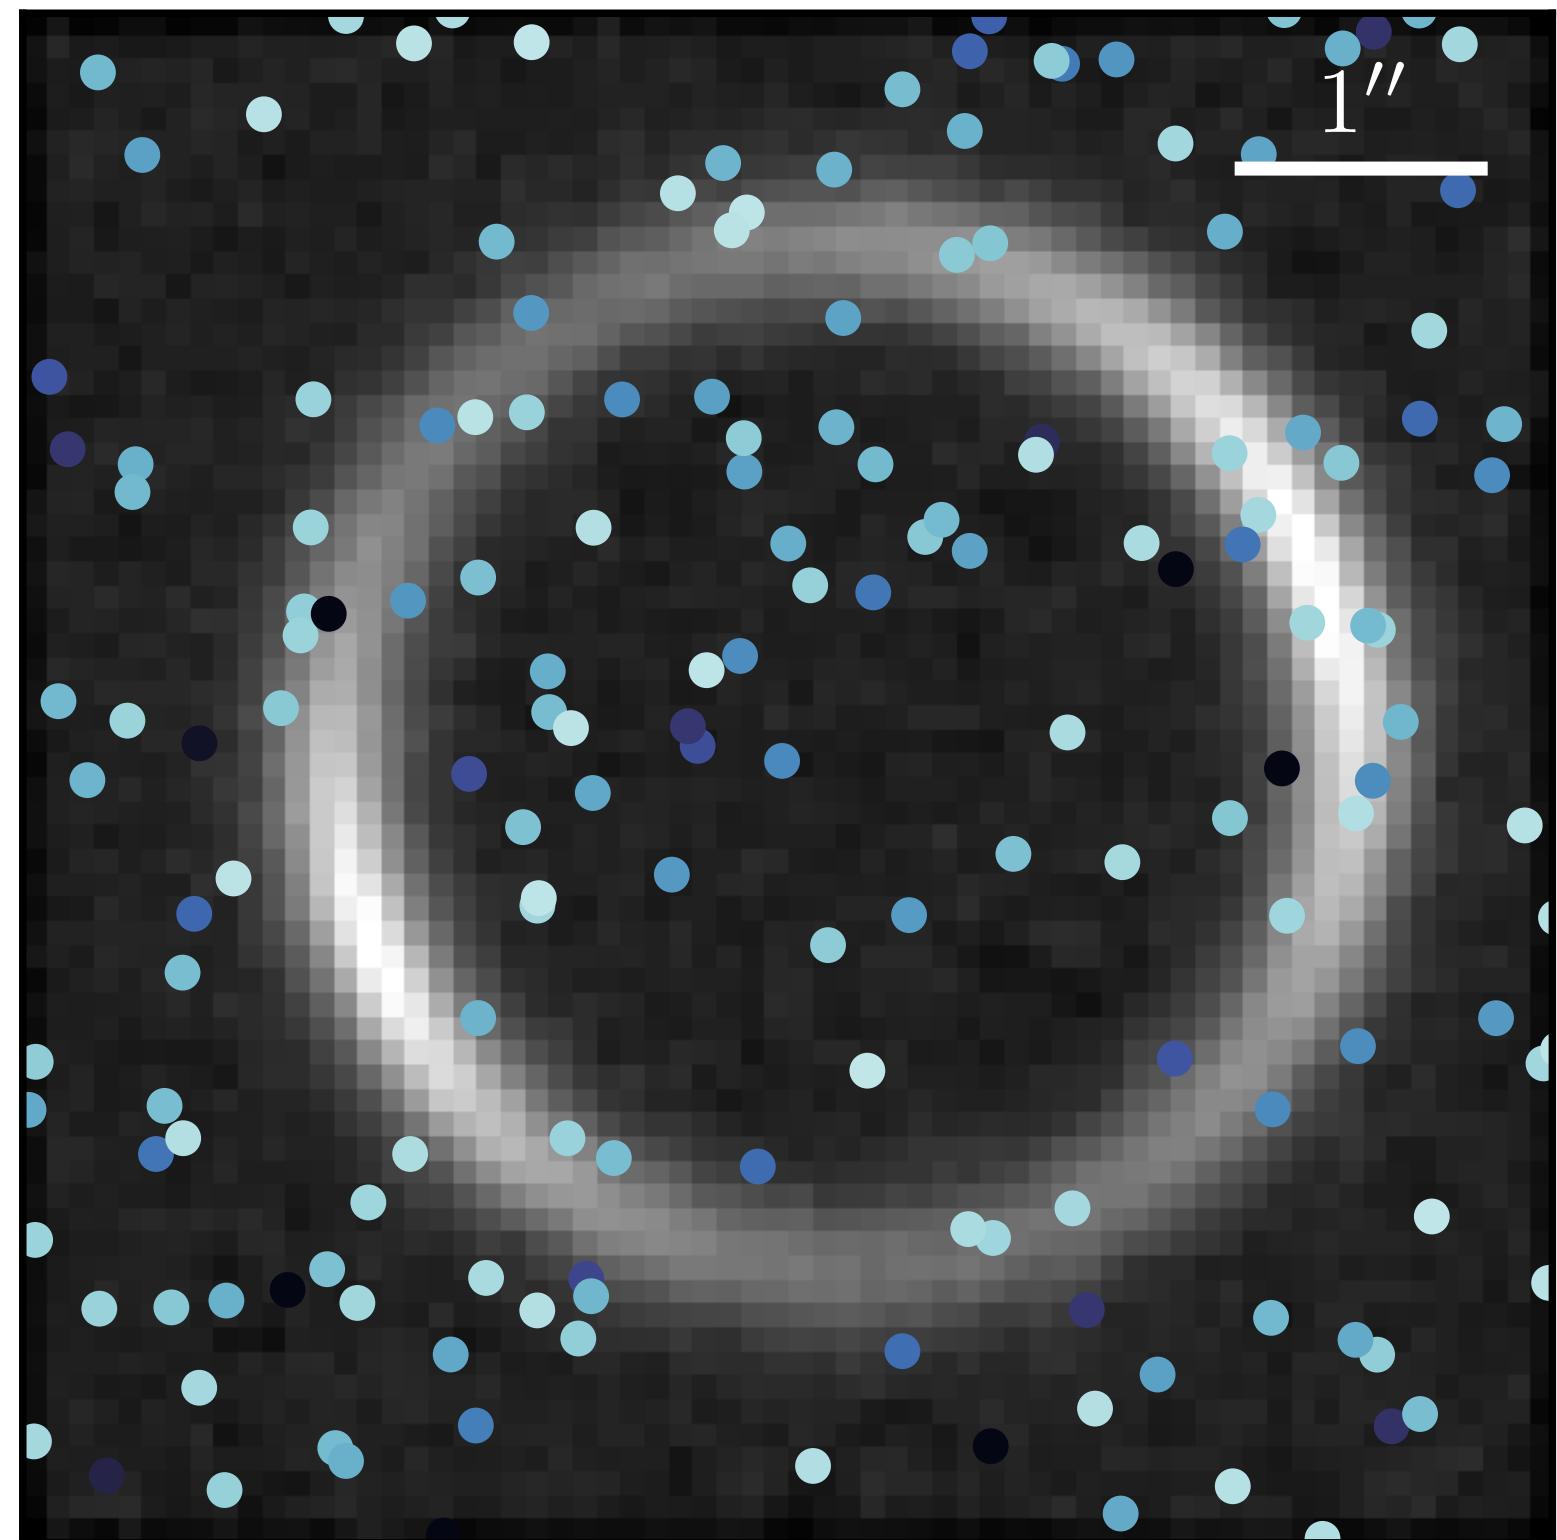
Expected posterior



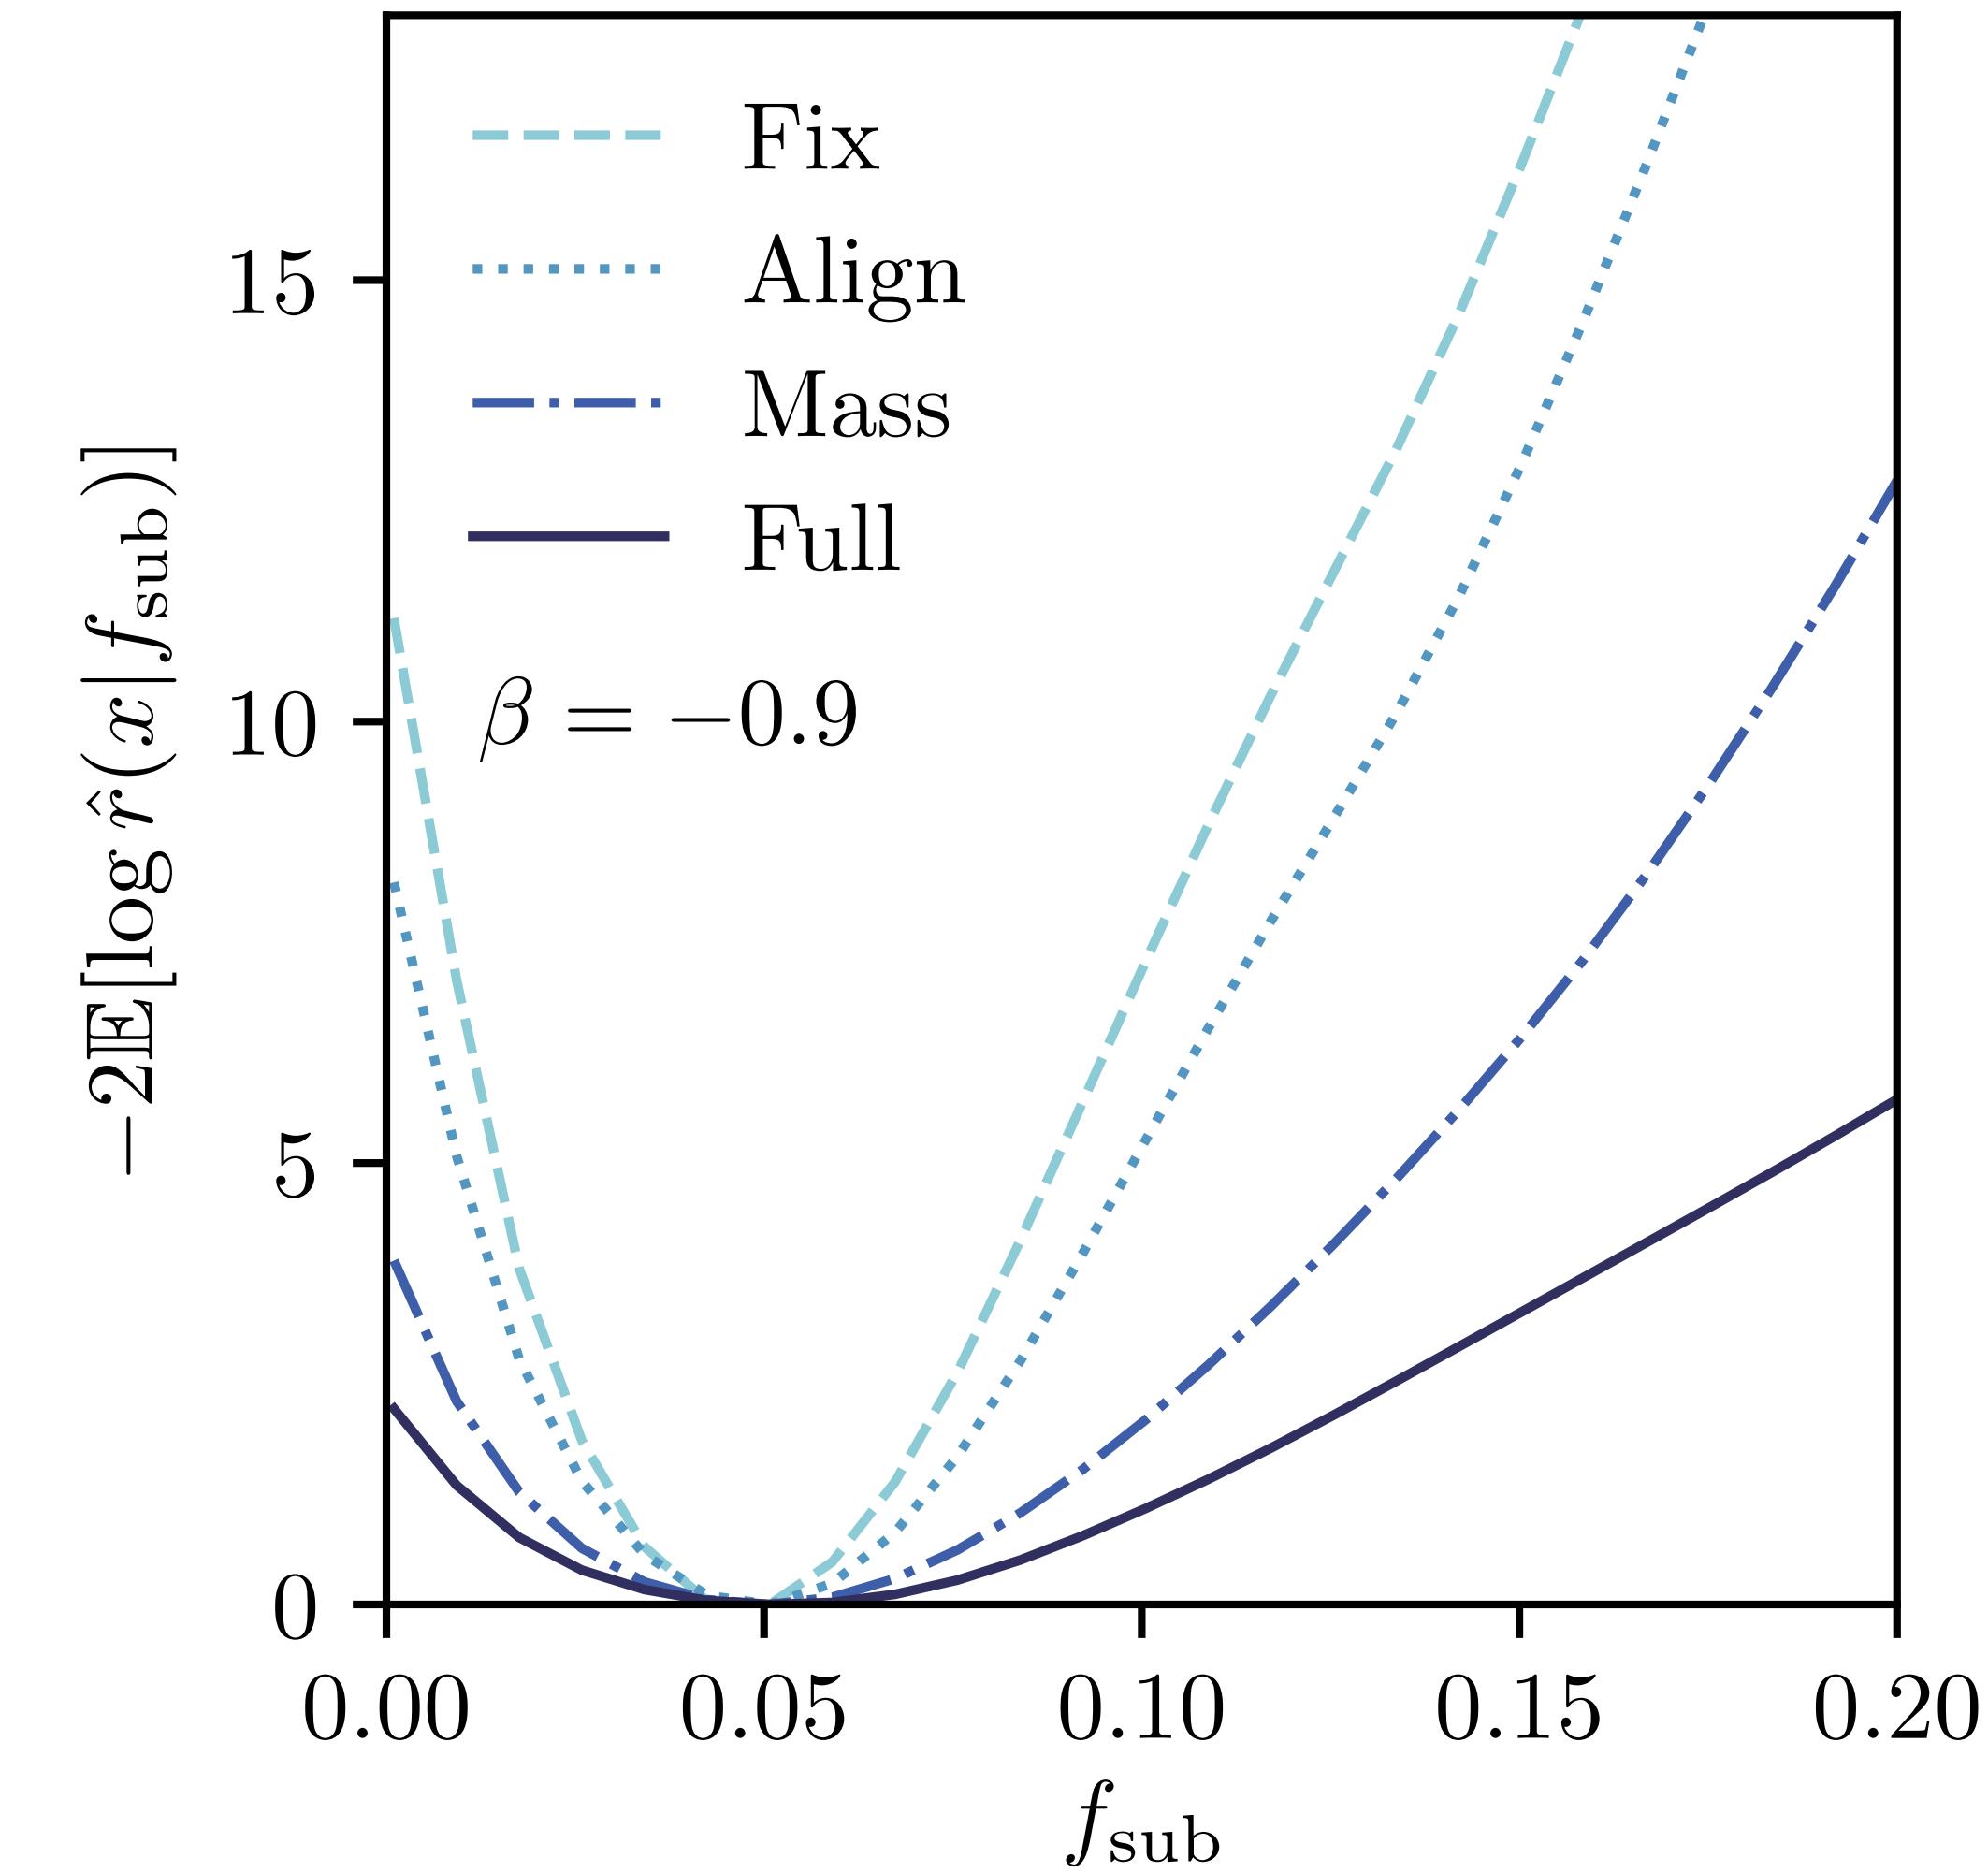
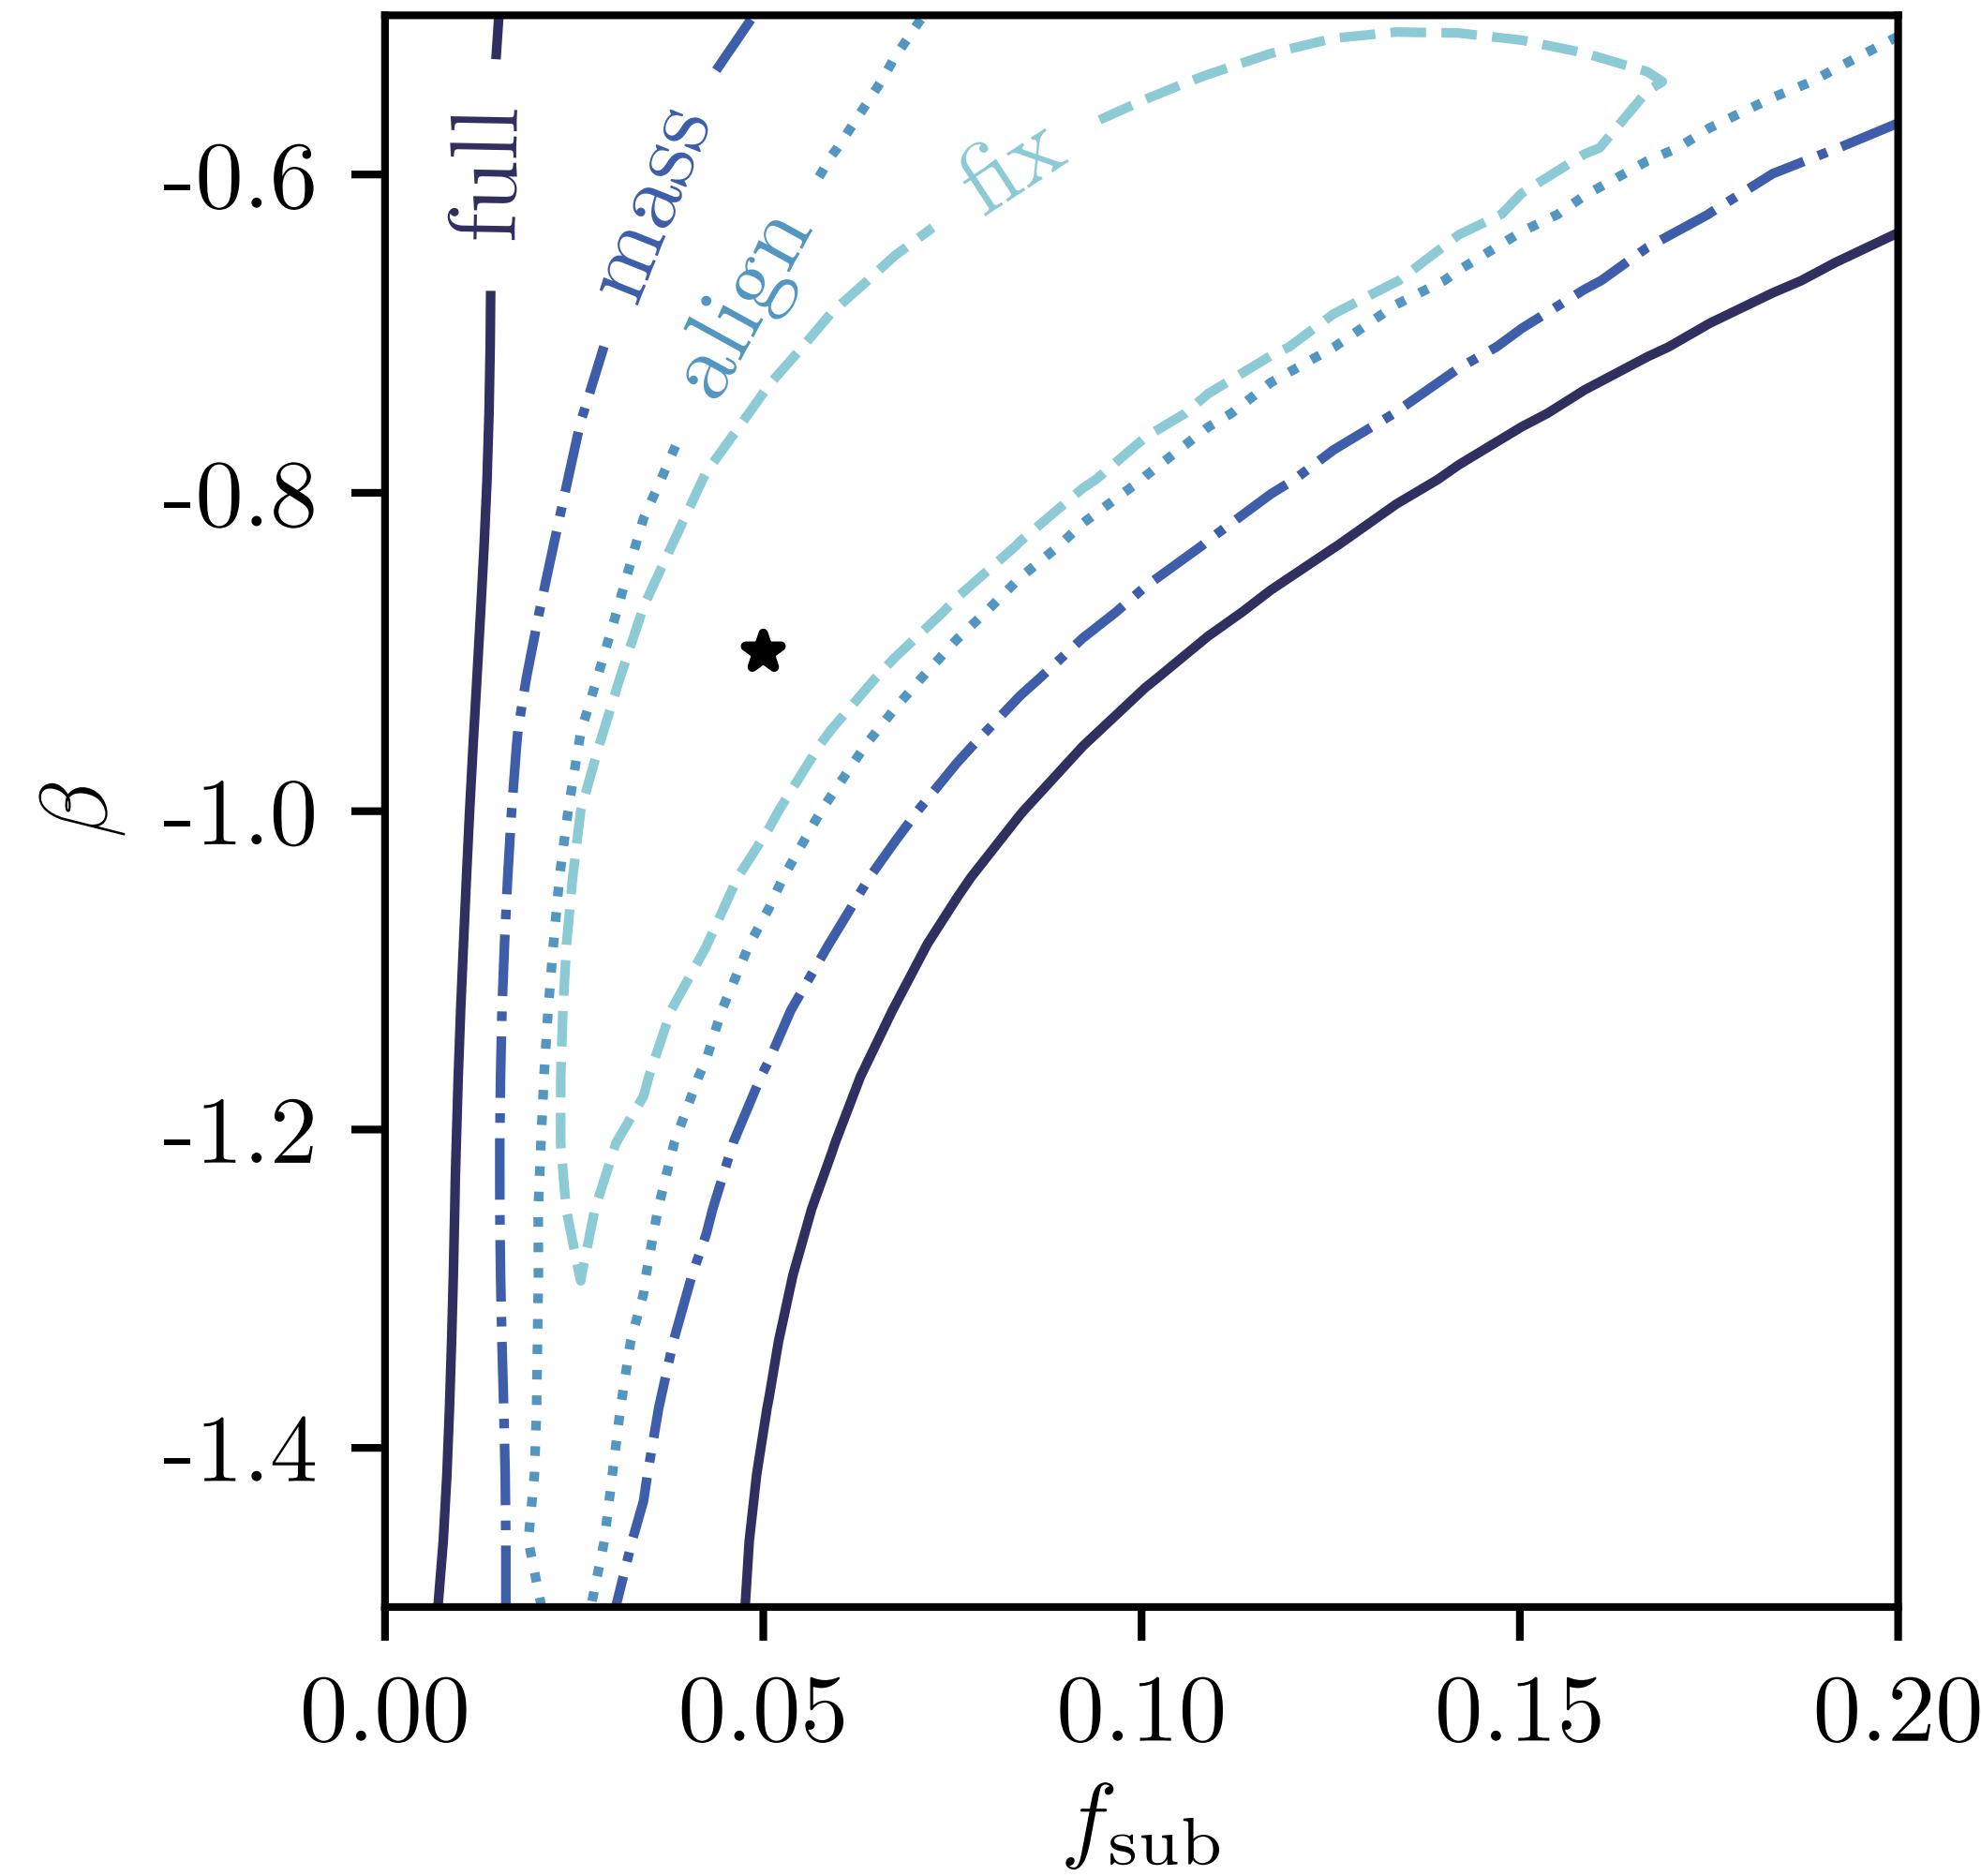
Likelihood vs total subhalo deflection



Deflection maps



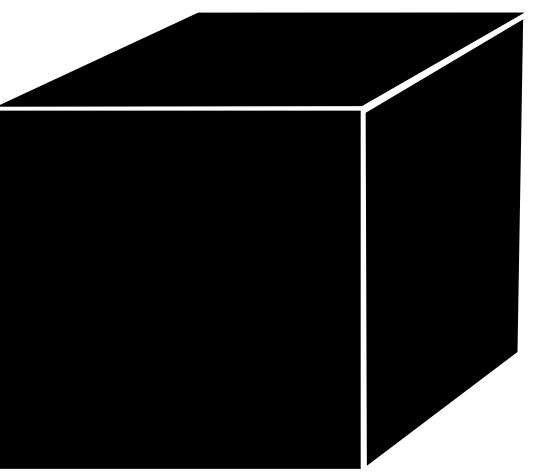
Simplified scenarios



Formalizing the problem

You are given

- a simulator that lets you generate N samples $x_i \sim p(x_i|\theta_i)$ for parameter points θ_i of your choice
- observed data $x_{\text{obs}} \sim p(x_{\text{obs}}|\theta_{\text{true}})$
- prior belief $p(\theta)$



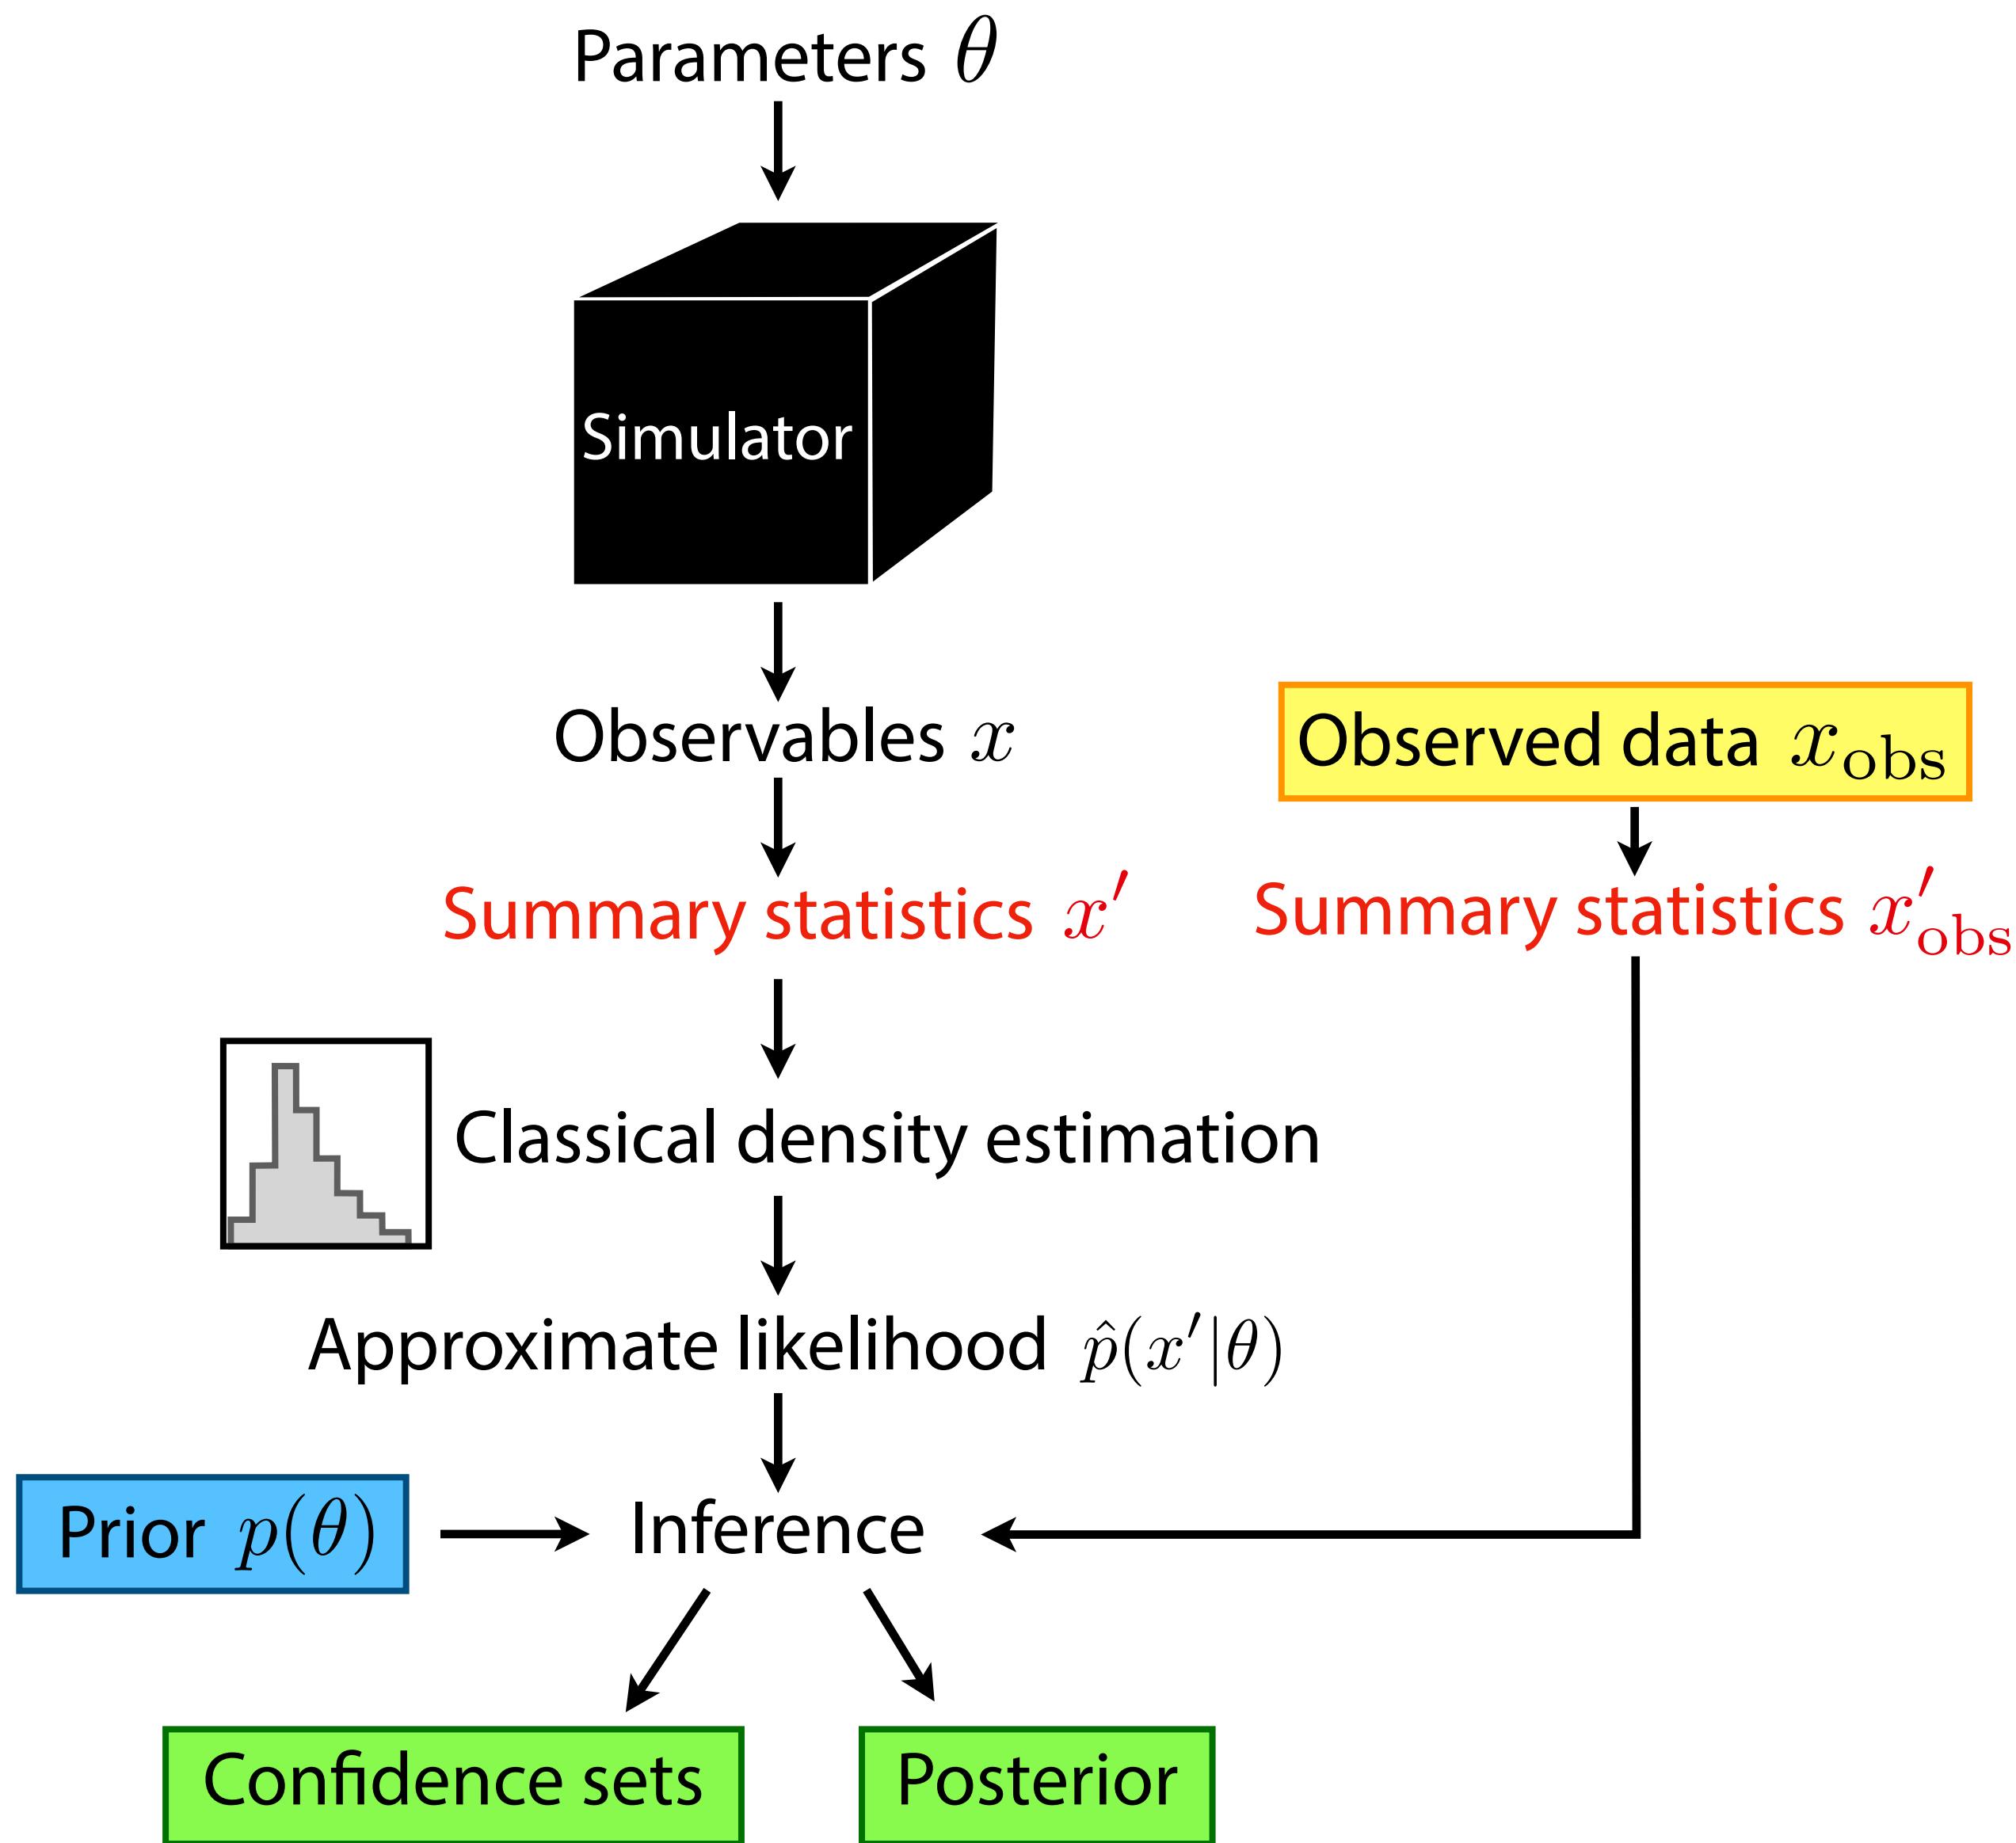
Goals: estimate either...

- true parameters $\hat{\theta}_{\text{true}}$
- confidence sets based on likelihood $\hat{p}(x_{\text{obs}}|\theta)$
- posterior $\hat{p}(\theta|x_{\text{obs}}) = \frac{\hat{p}(x_{\text{obs}}|\theta) p(\theta)}{\int d\theta' \hat{p}(x_{\text{obs}}|\theta') p(\theta')}$
or samples from posterior $\theta \sim \hat{p}(\theta|x_{\text{obs}})$

... depending on domain conventions

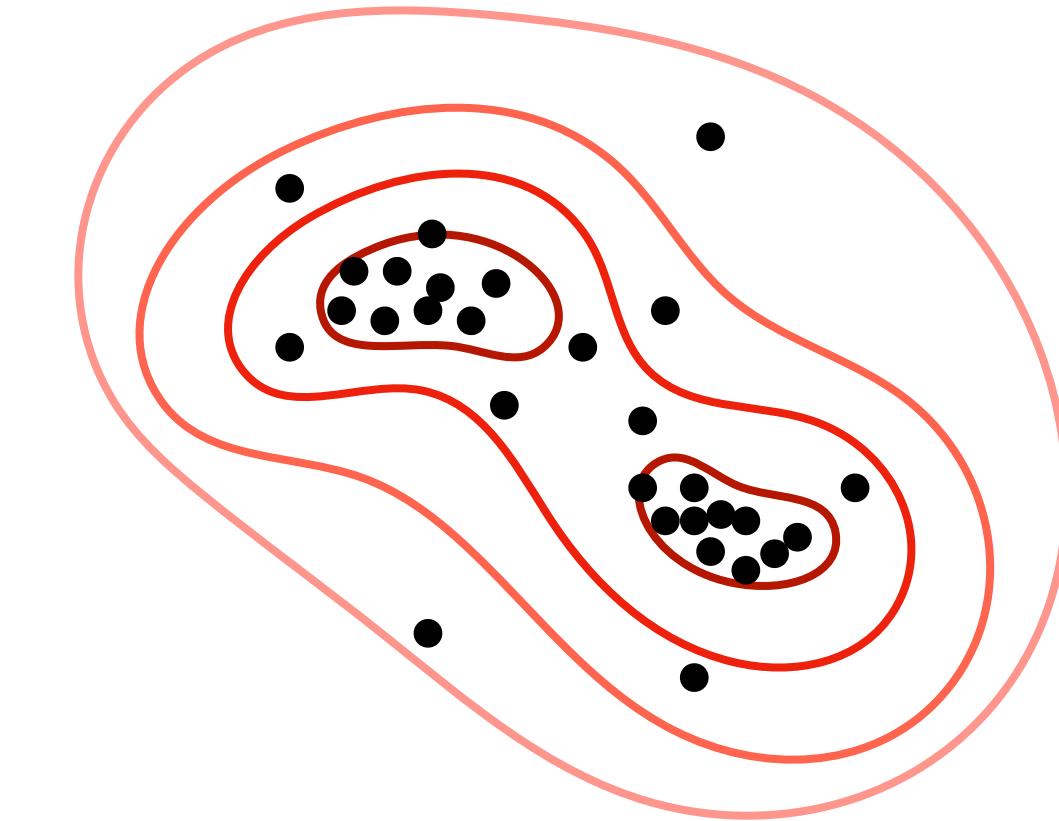
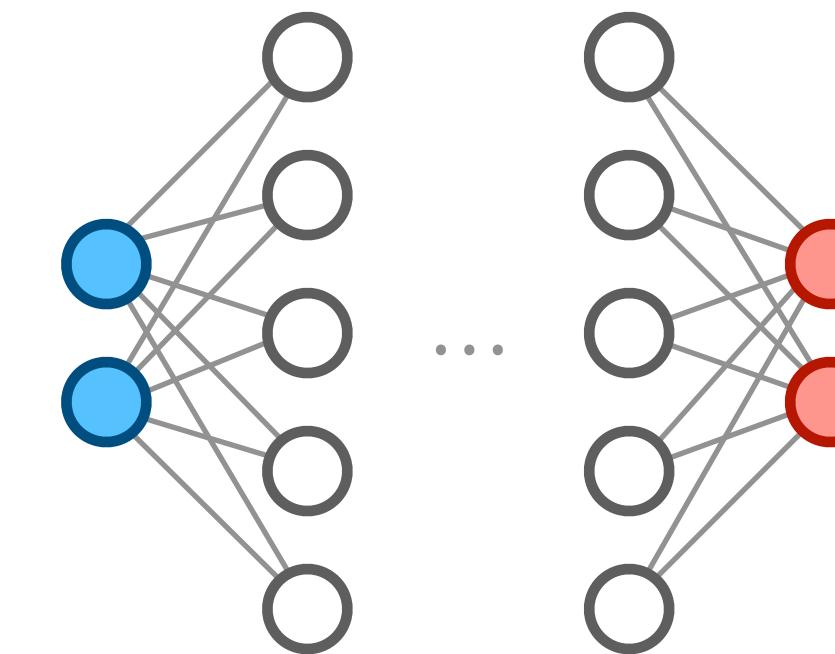
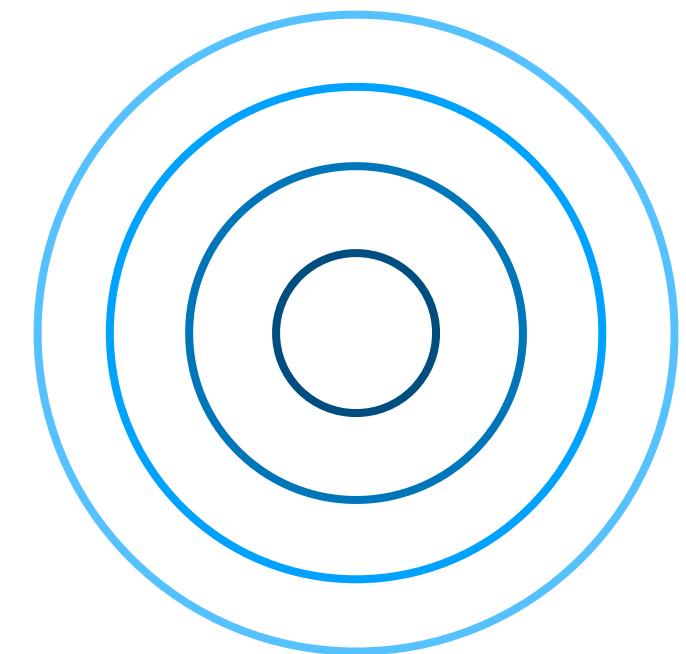
Inference by estimating the likelihood

[e.g. P. Diggle, R. Gratton 1984]



- Compression to summary statistics reduces quality of inference
- Curse of dimensionality: does not scale to more than a few summary statistics

High-dimensional density estimation with normalizing flows



Simple base density

$$u \sim \pi(u)$$

NN: transformation $x = f(u)$

- one-to-one and invertible
- differentiable
- f^{-1} and $\det \nabla f$ are tractable

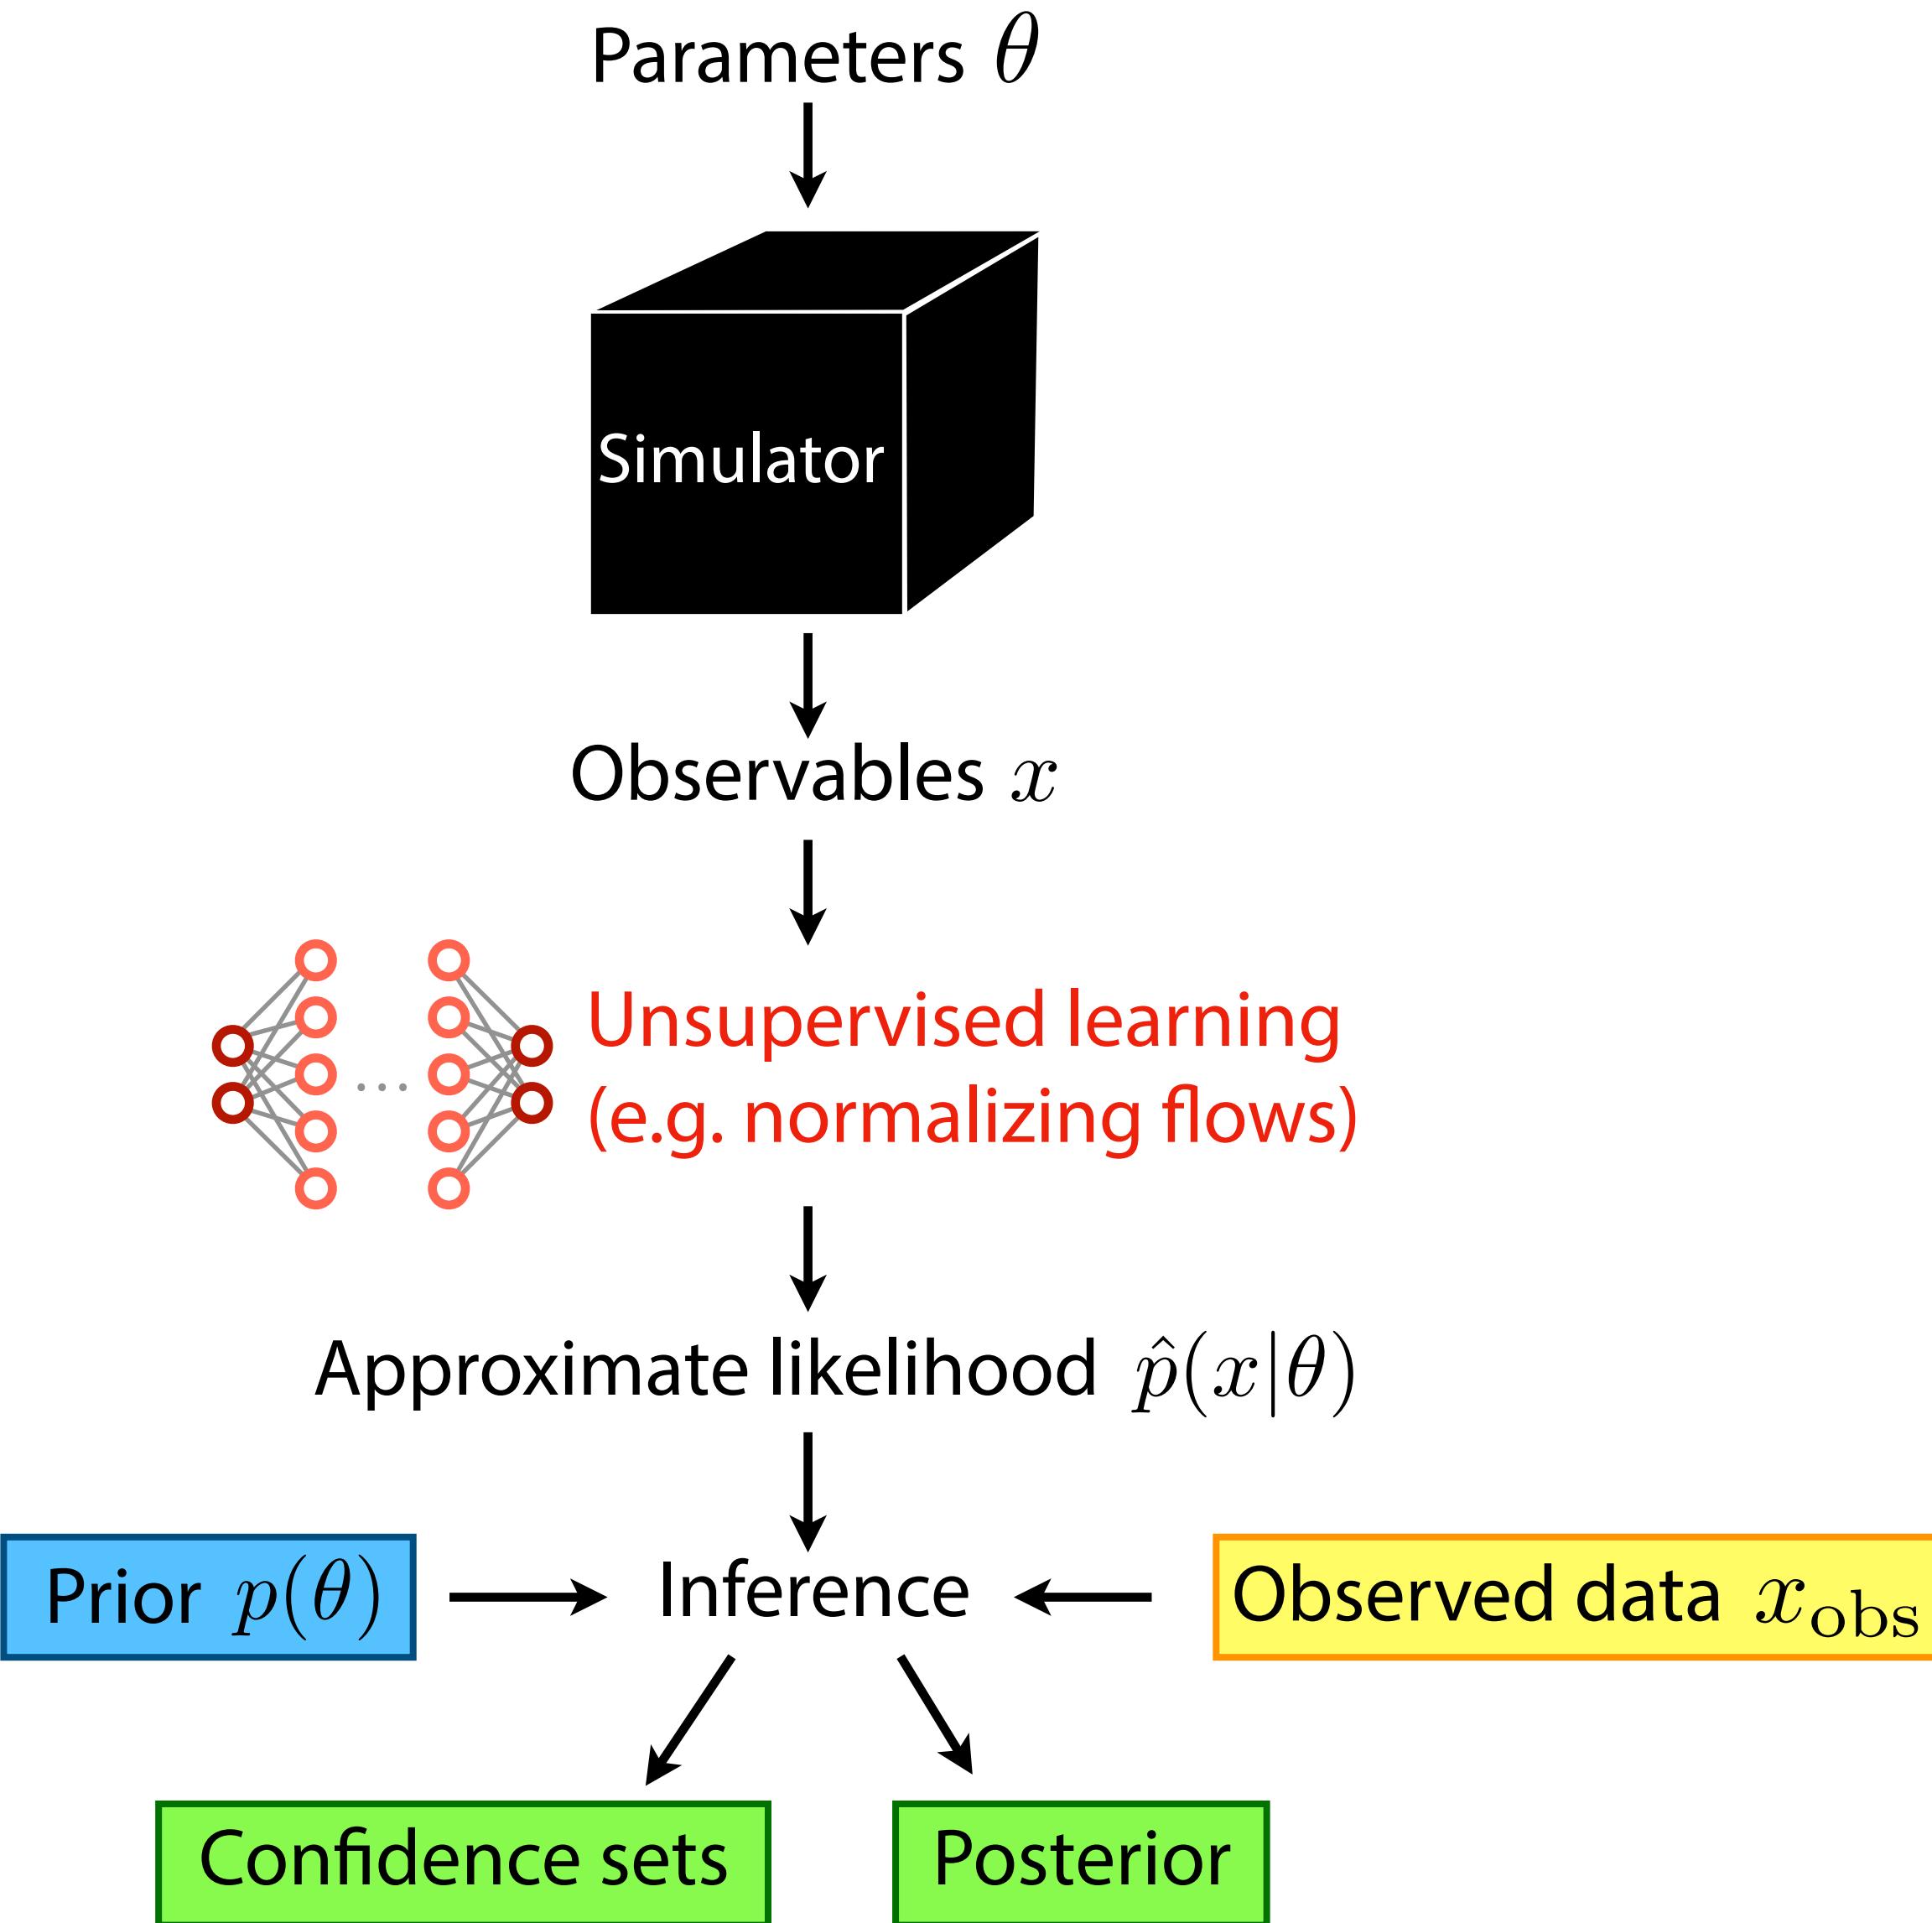
Target density is given by

$$\hat{p}(x) = \pi(f^{-1}(x)) |\det \nabla f|^{-1}$$

Train transformation by
maximizing $\log \hat{p}(x)$

Transformation can depend on θ
to model conditional density $\log \hat{p}(x|\theta)$

Inference with neural likelihood estimation

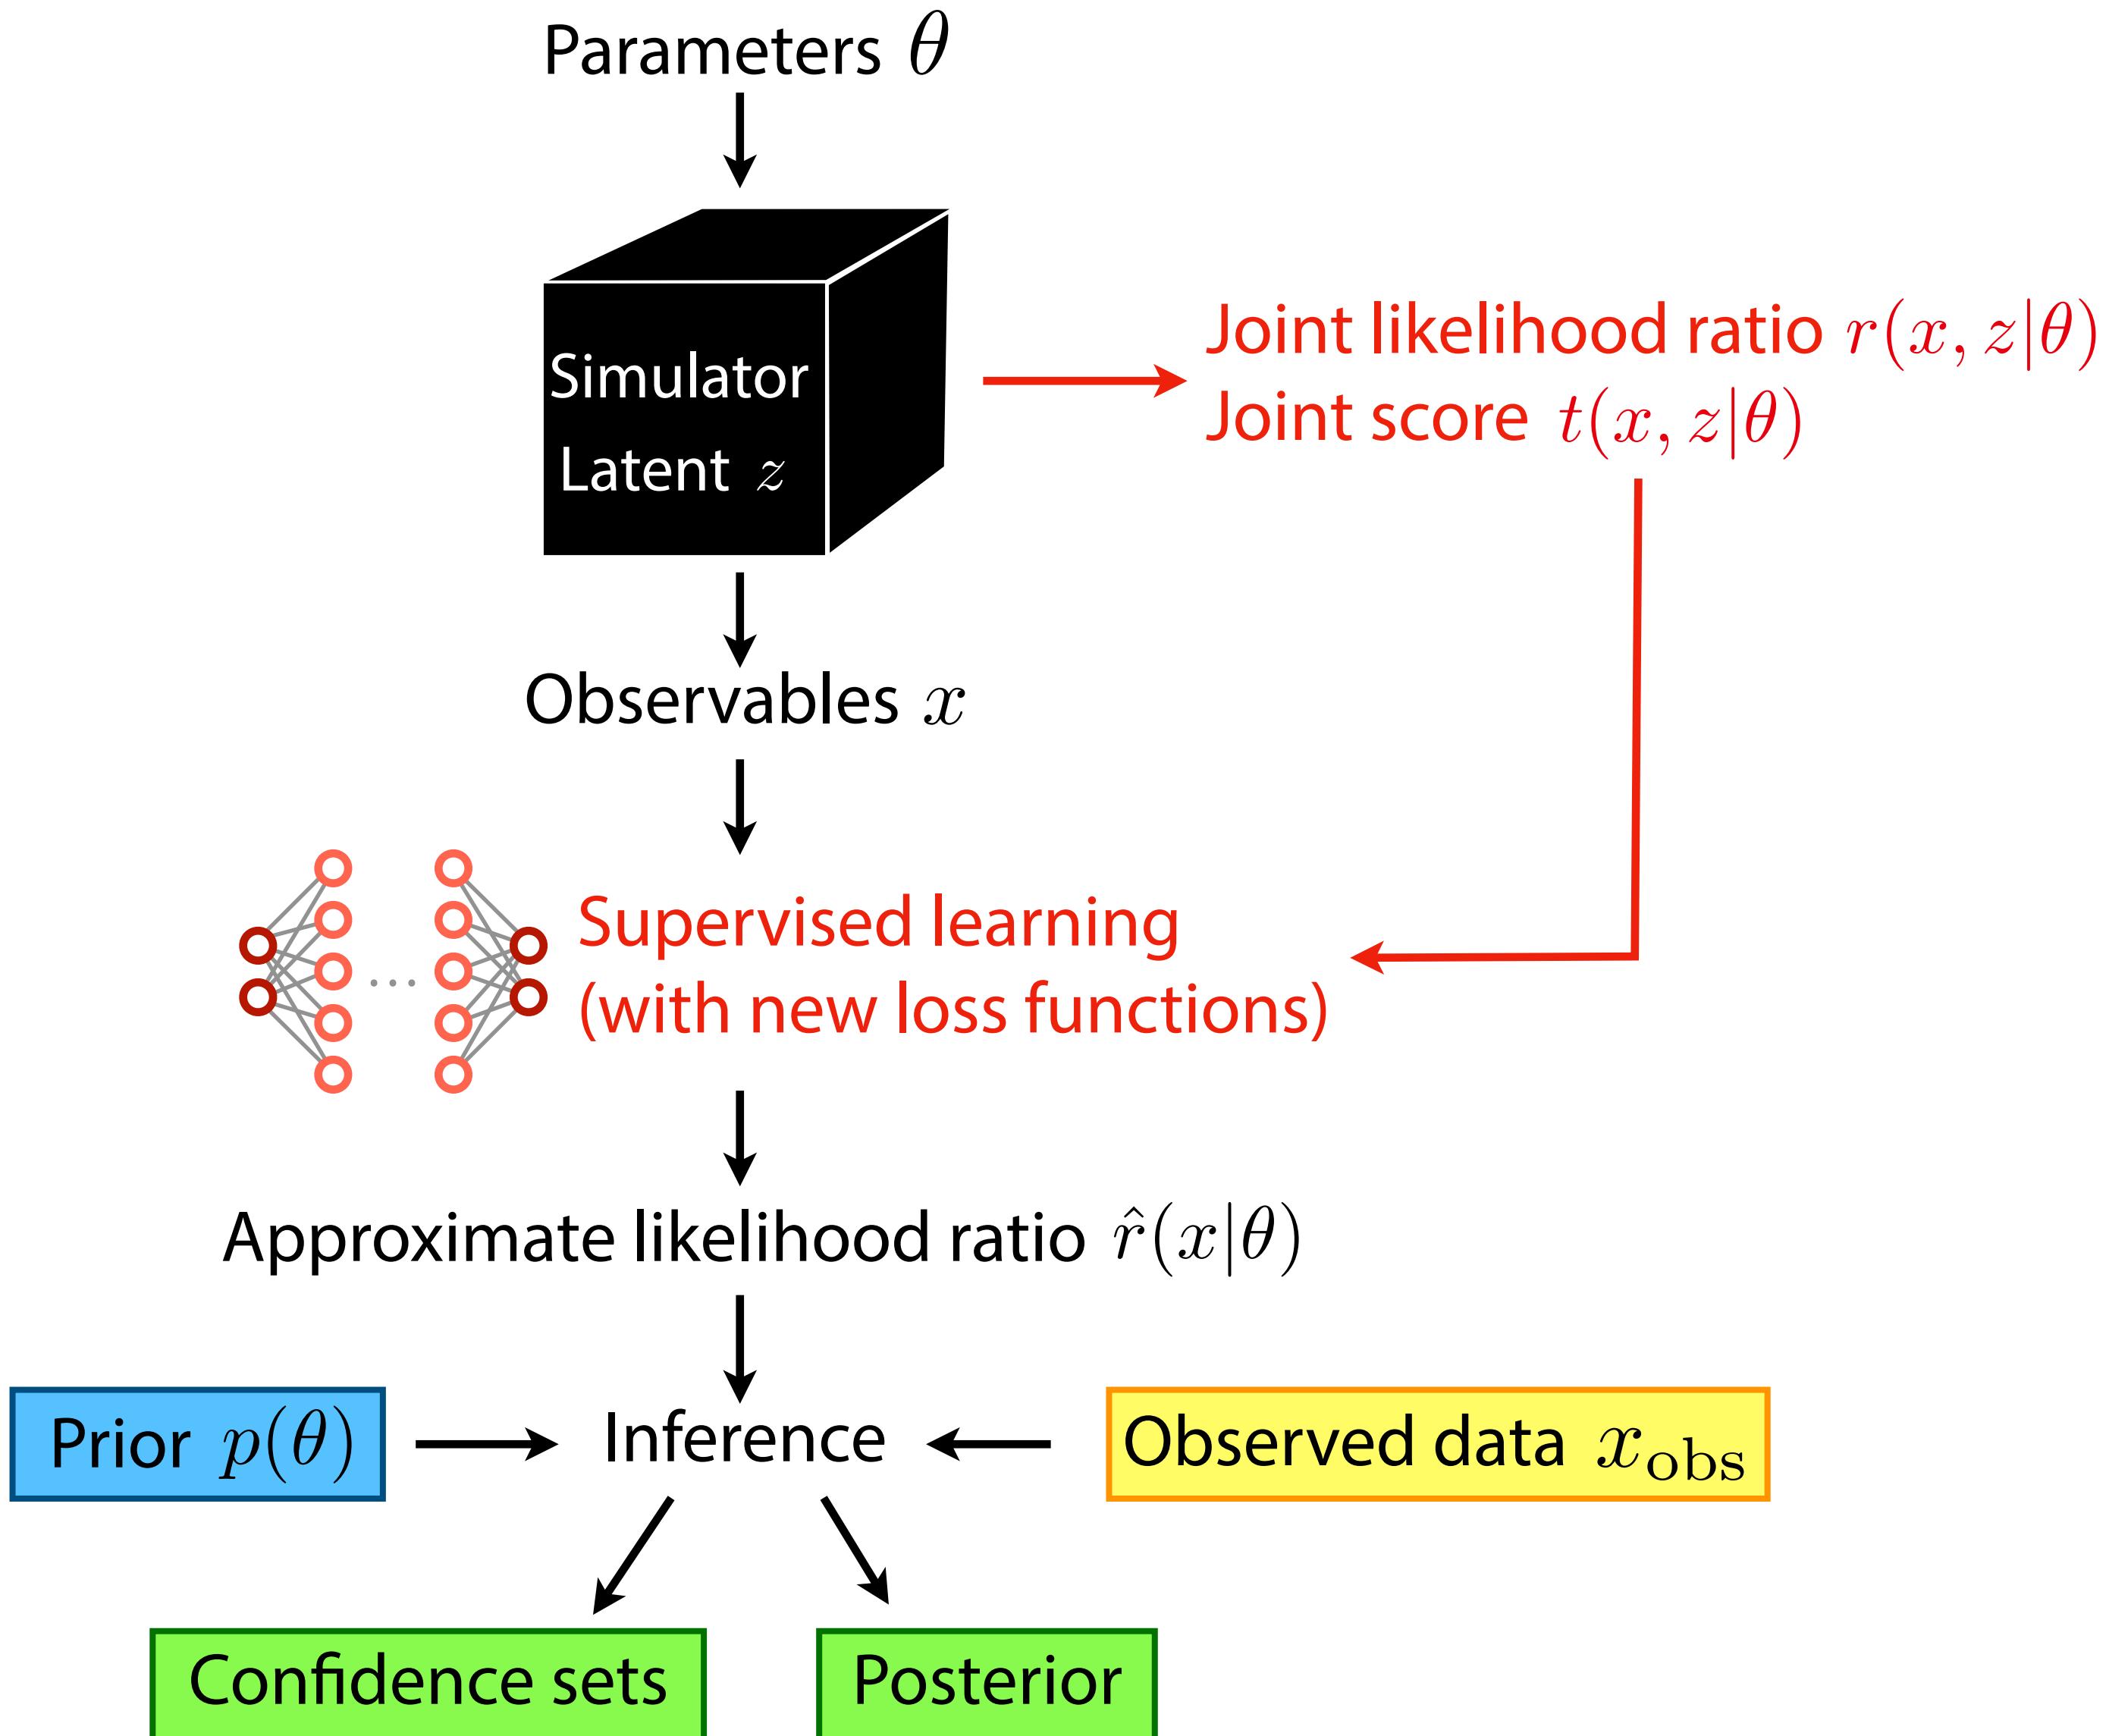


[G. Papamakarios, D. Sterratt, I. Murray 1805.07226;
J.-M. Lueckmann, G. Bassetto, T. Karaletsos, J. Macke 1805.09294]

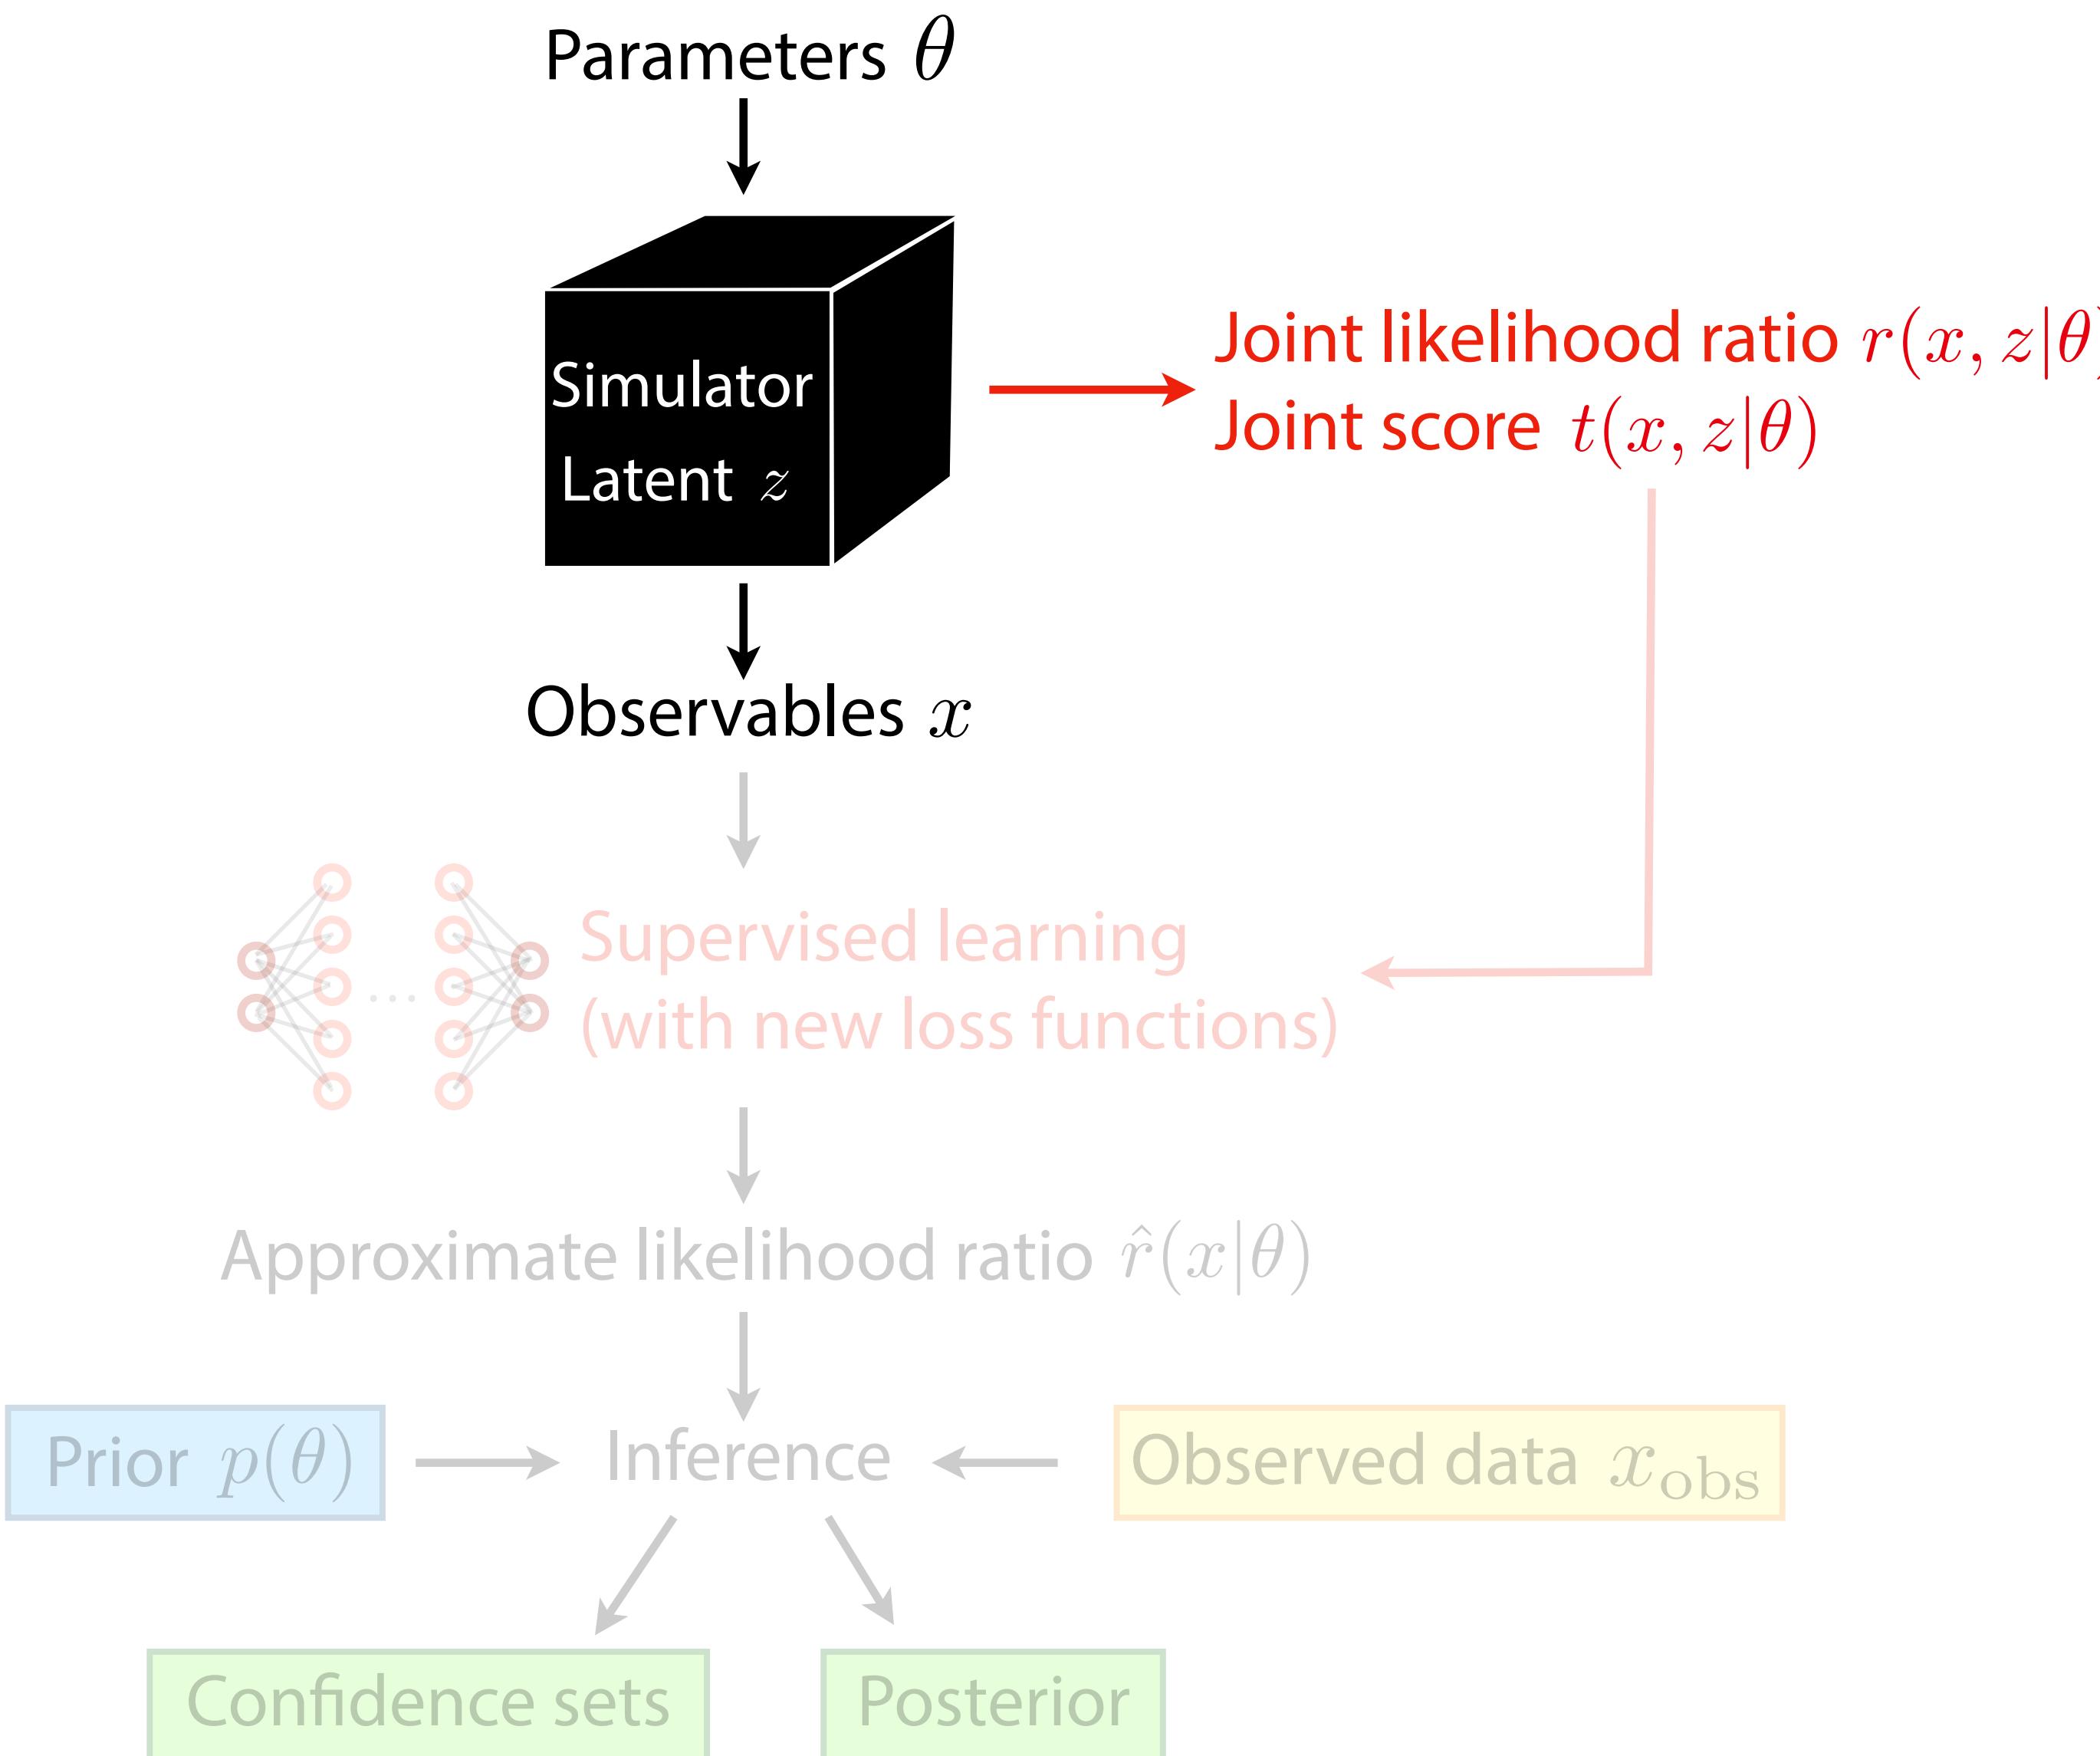
- Conditional neural density estimator (e.g. normalizing flow) as tractable surrogate for simulator likelihood
- Scales well to high-dimensional data (no compression to summary stats necessary)
- Amortized: After upfront simulation + training phase, inference is efficient for new data or prior
- Related alternative: learn posterior $\hat{p}(\theta|x_{\text{obs}})$

[G. Papamakarios et al 1605.06376;
J.-M. Lueckmann et al 1711.01861]

Mining gold

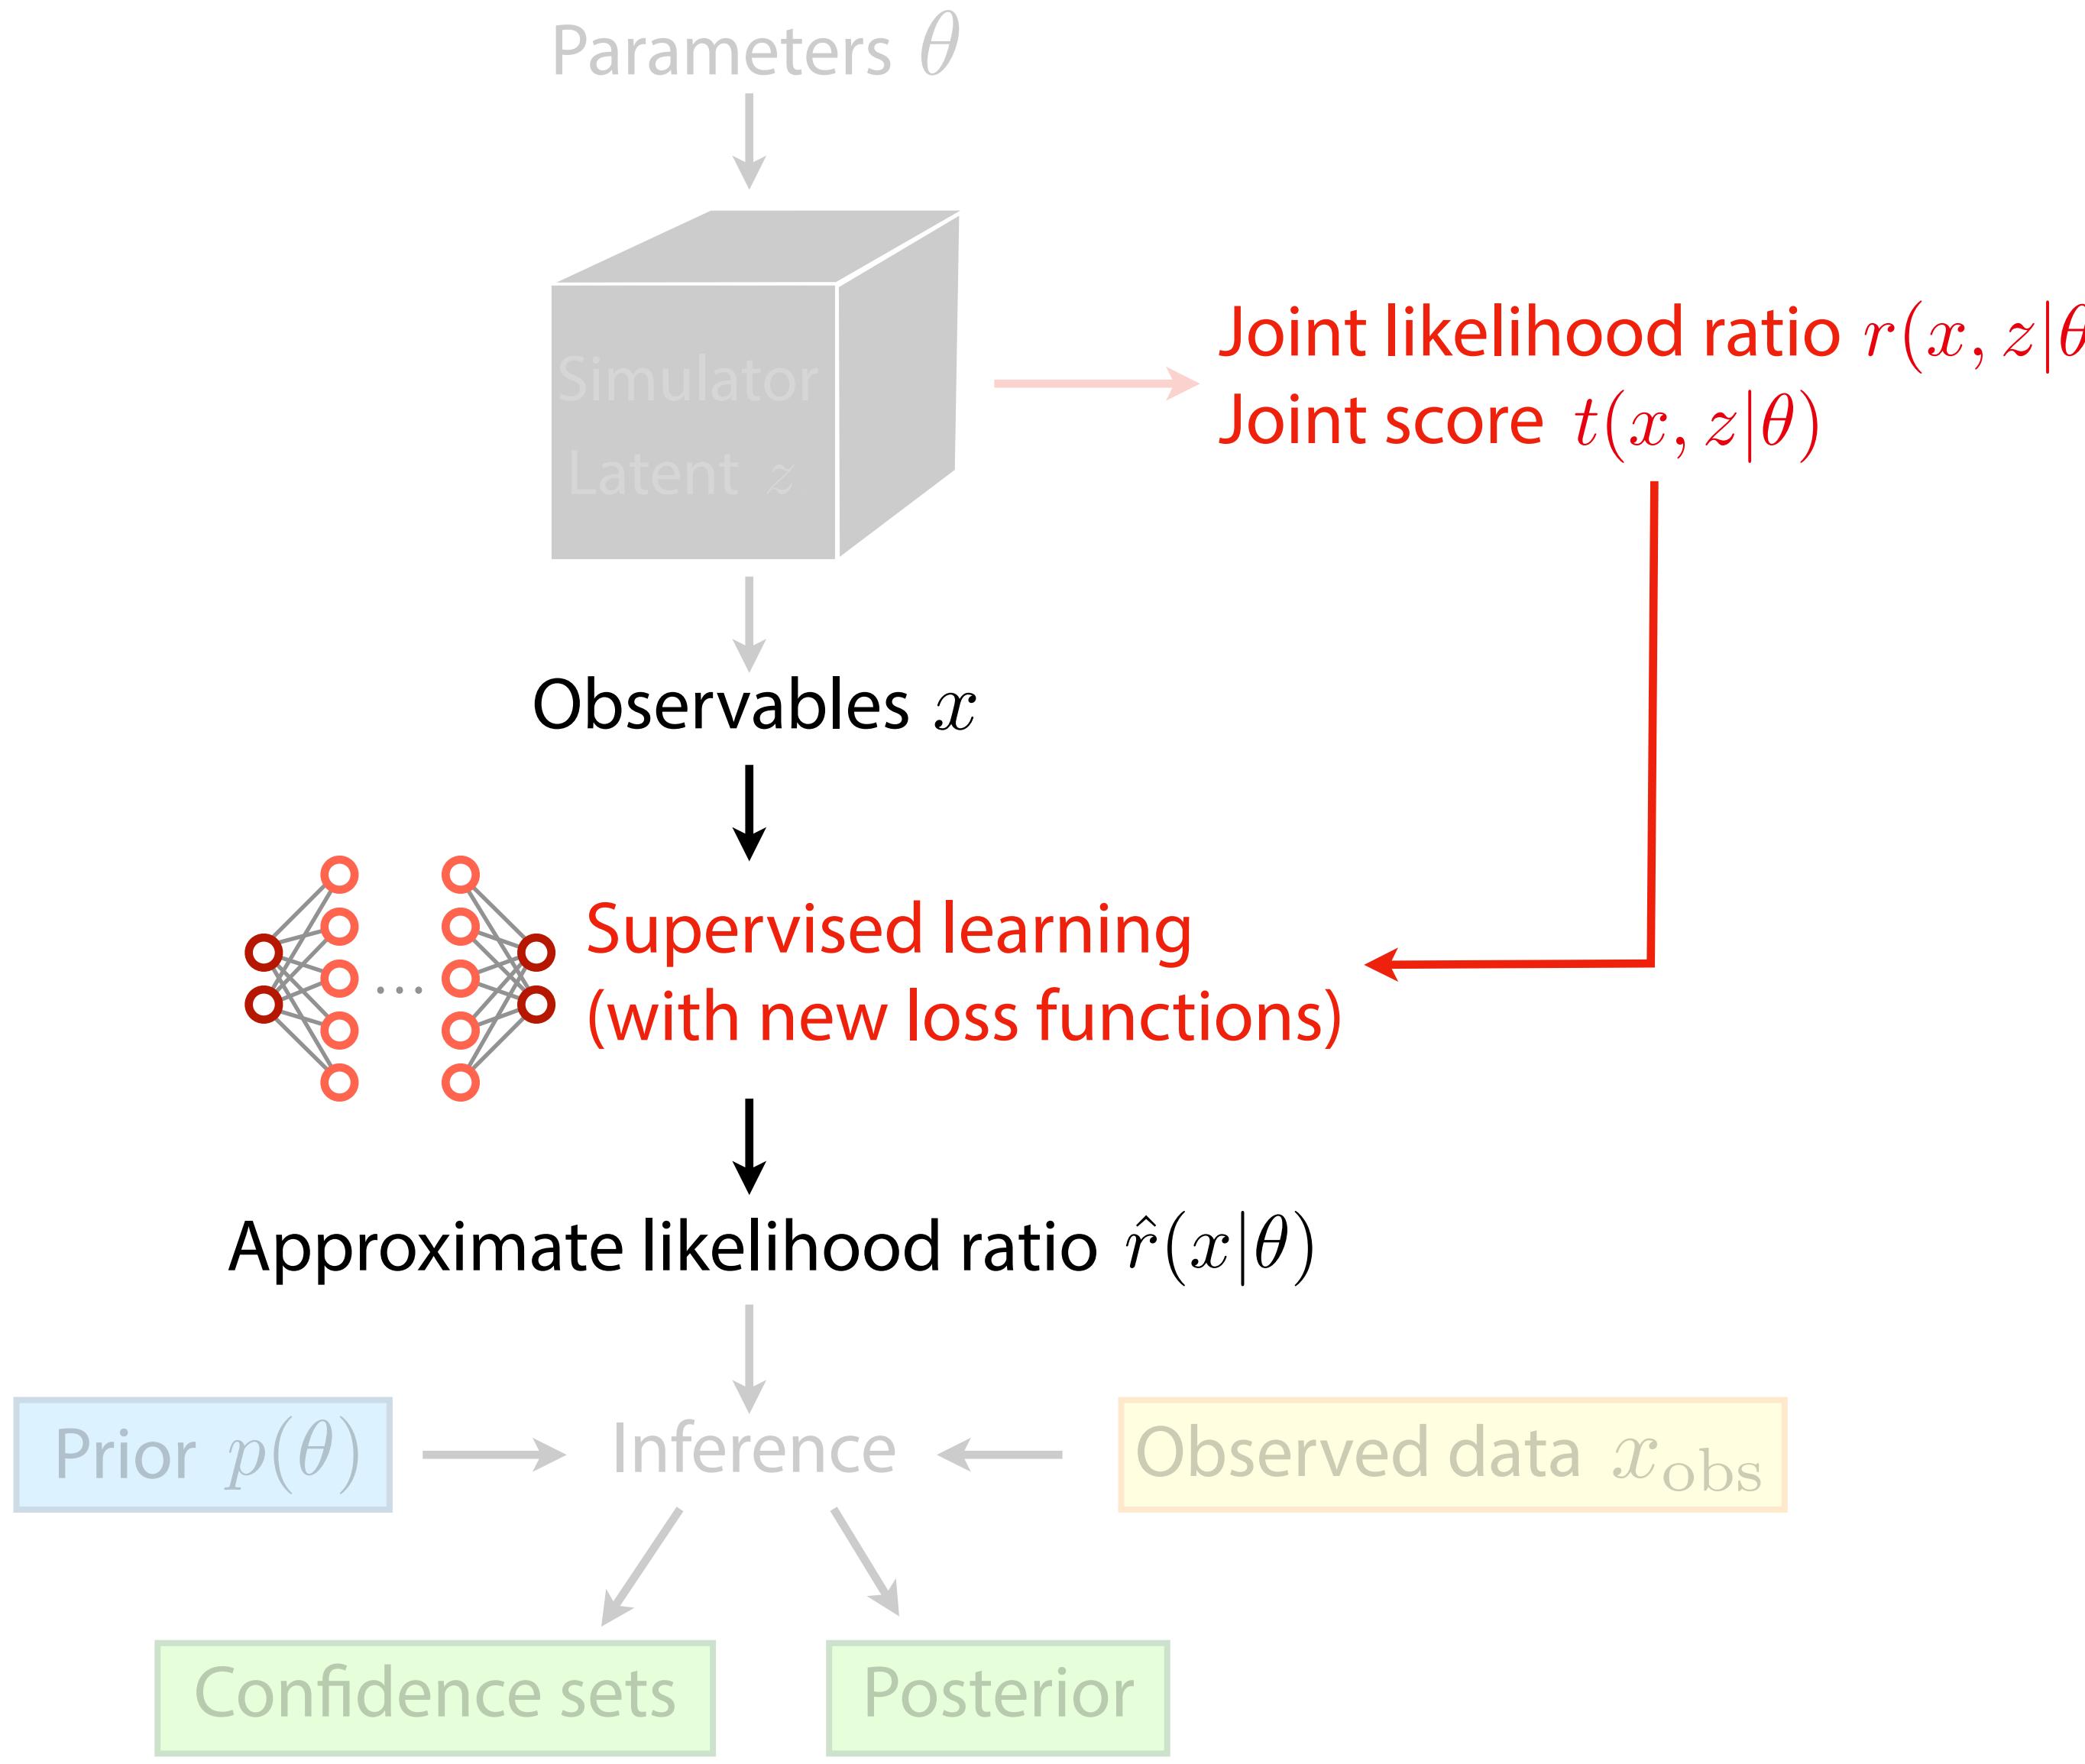


Step 1: Extracting more information from simulations



- For each simulated event, calculate
 - joint likelihood ratio $r(x, z|\theta) = \frac{p(x, z|\theta)}{p_{\text{ref}}(x, z)}$
 - joint score $t(x, z|\theta) = \nabla_{\theta} \log p(x, z|\theta)$
- (How much more or less likely would this simulated event (fixing all latent variables) be when changing the theory parameters?)

Step 2: Machine learning



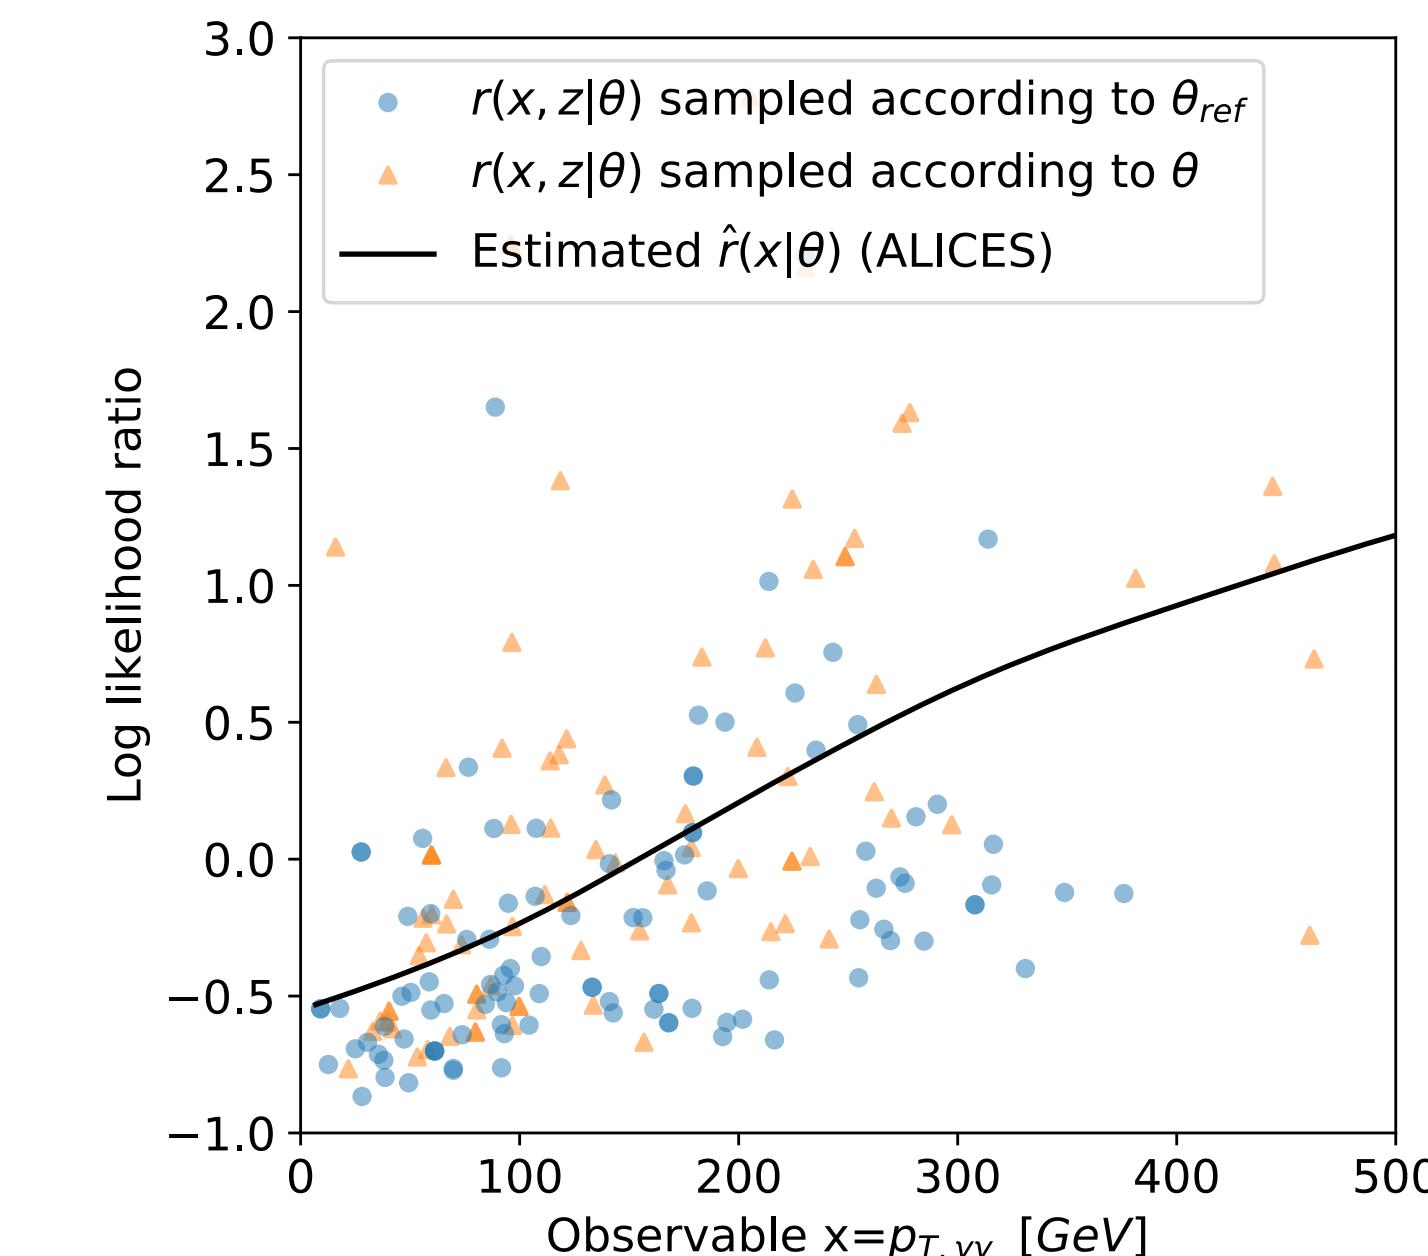
- Train a neural network $g(x, \theta)$ on loss functionals like

$$L[g] = \frac{1}{N} \sum_i |g(x_i, \theta_i) - r(x_i, z_i|\theta_i)|^2$$

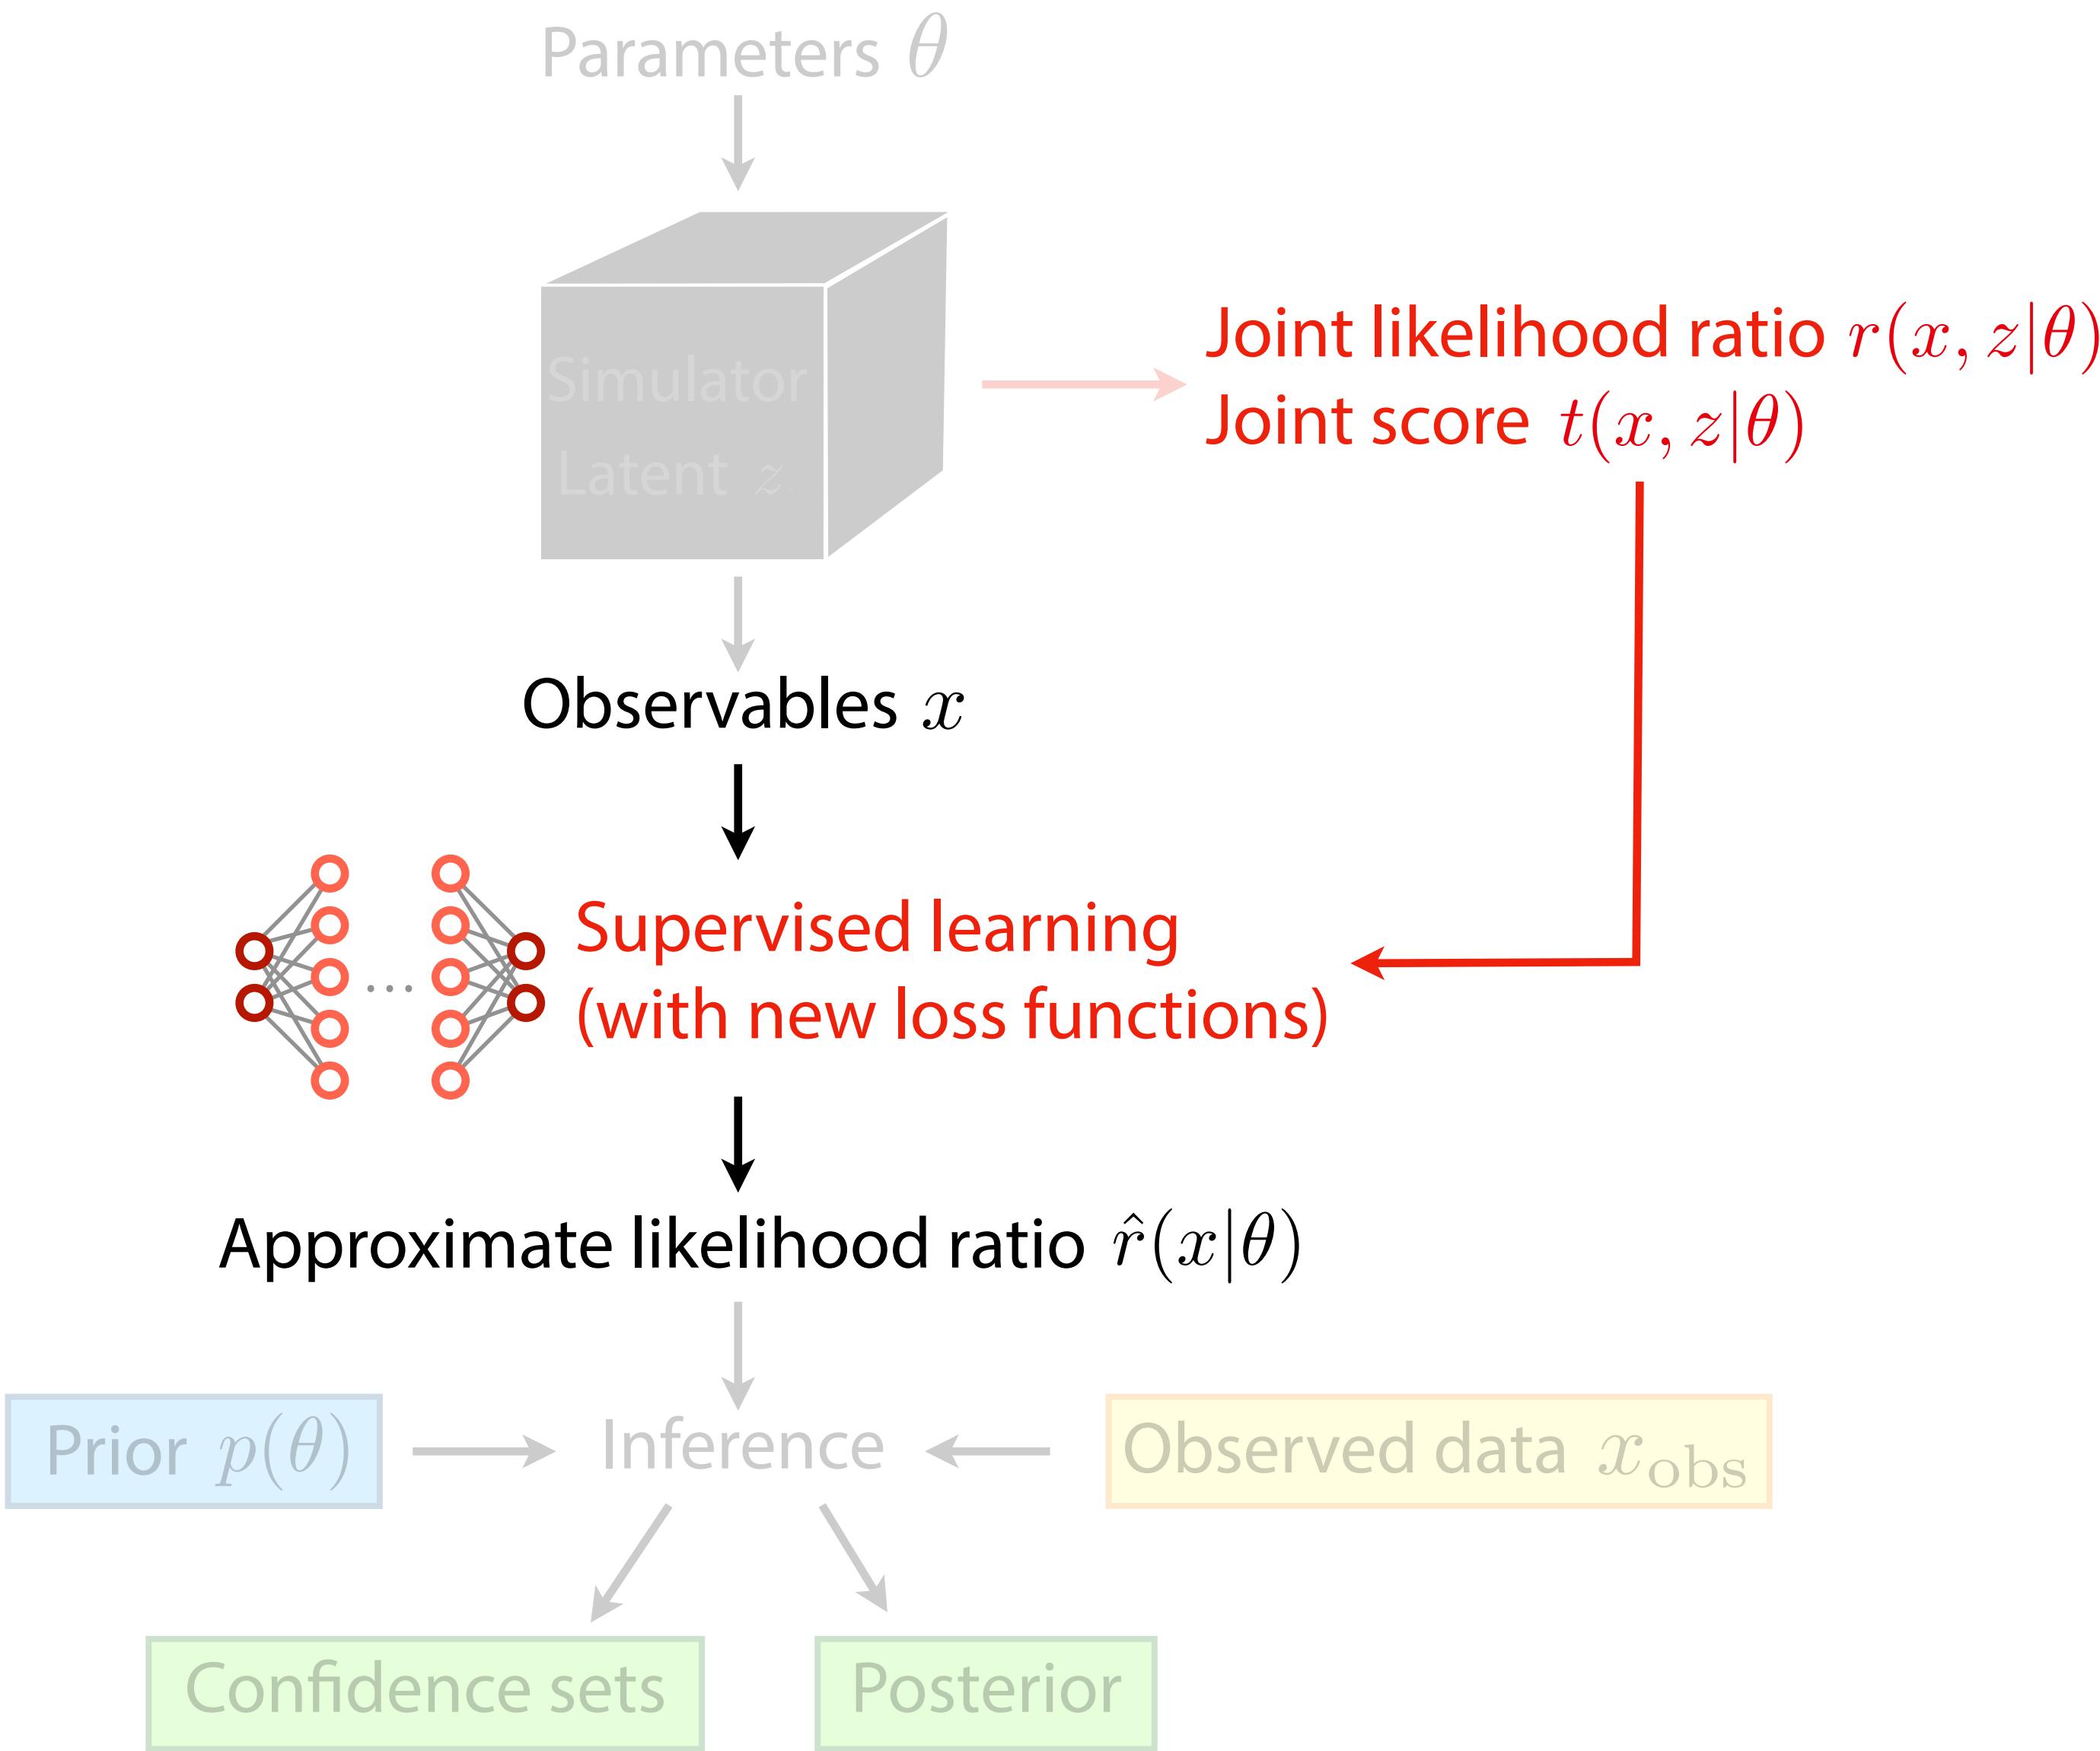
- The network will converge to

$$g(x, \theta) \rightarrow \arg \min L[g] = r(x|\theta) = \frac{p(x|\theta)}{p_{\text{ref}}(x)} !$$

(for sufficient training data, NN capacity, efficient optimization)



Step 2: Machine learning



- Train a neural network $g(x, \theta)$ on loss functionals like

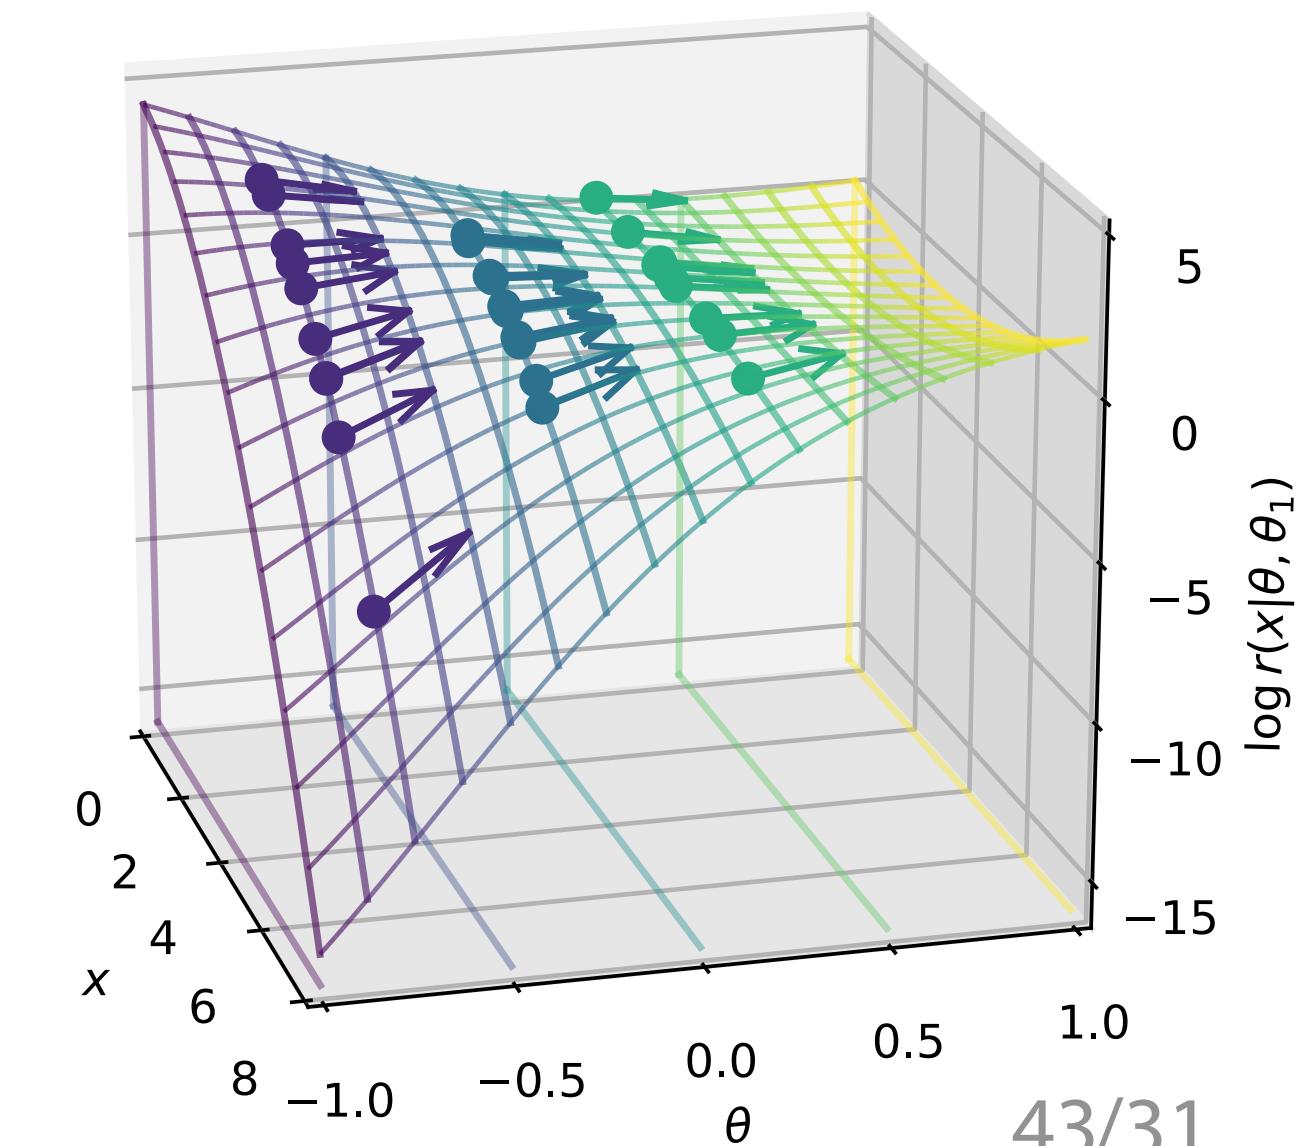
$$L[g] = \frac{1}{N} \sum_i |g(x_i, \theta_i) - r(x_i, z_i | \theta_i)|^2$$

- The network will converge to

$$g(x, \theta) \rightarrow \arg \min L[g] = r(x | \theta) = \frac{p(x | \theta)}{p_{\text{ref}}(x)} !$$

(for sufficient training data, NN capacity, efficient optimization)

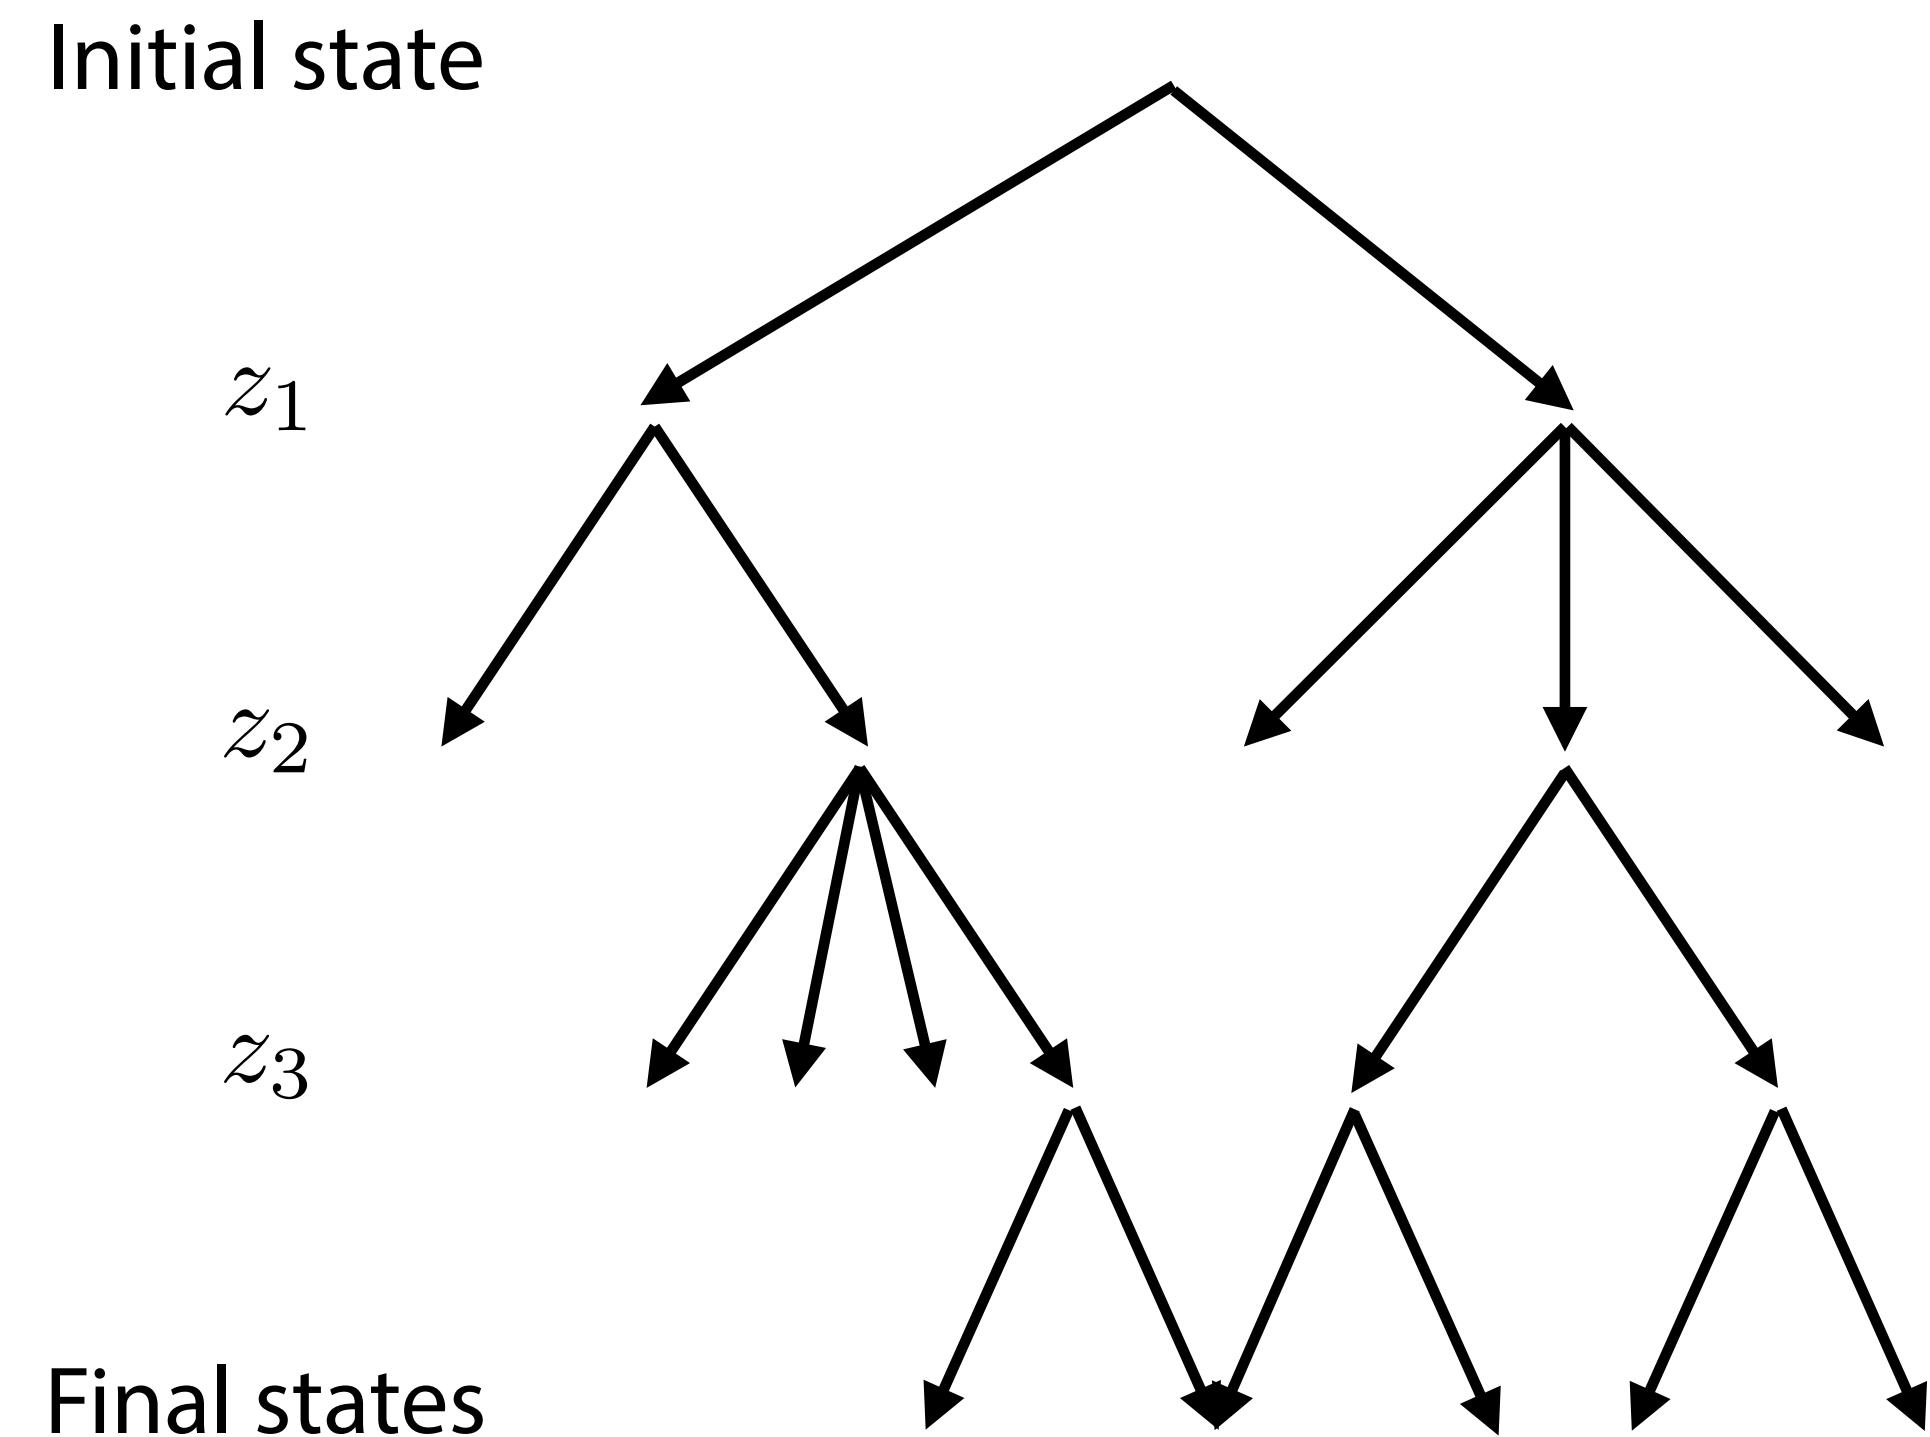
- RASCAL:
Joint score adds gradient information
 \Rightarrow three orthogonal pieces of information



Mining gold from any simulation

- Computer simulation typically evolve along a tree-like structure of successive random branchings
- The probabilities of each branching $p_i(z_i|z_{i-1}, \theta)$ are often clearly defined in the code:

```
if random() > 0.1 + 2.5 * model_parameter:  
    do_one_thing()  
else:  
    do_another_thing()
```



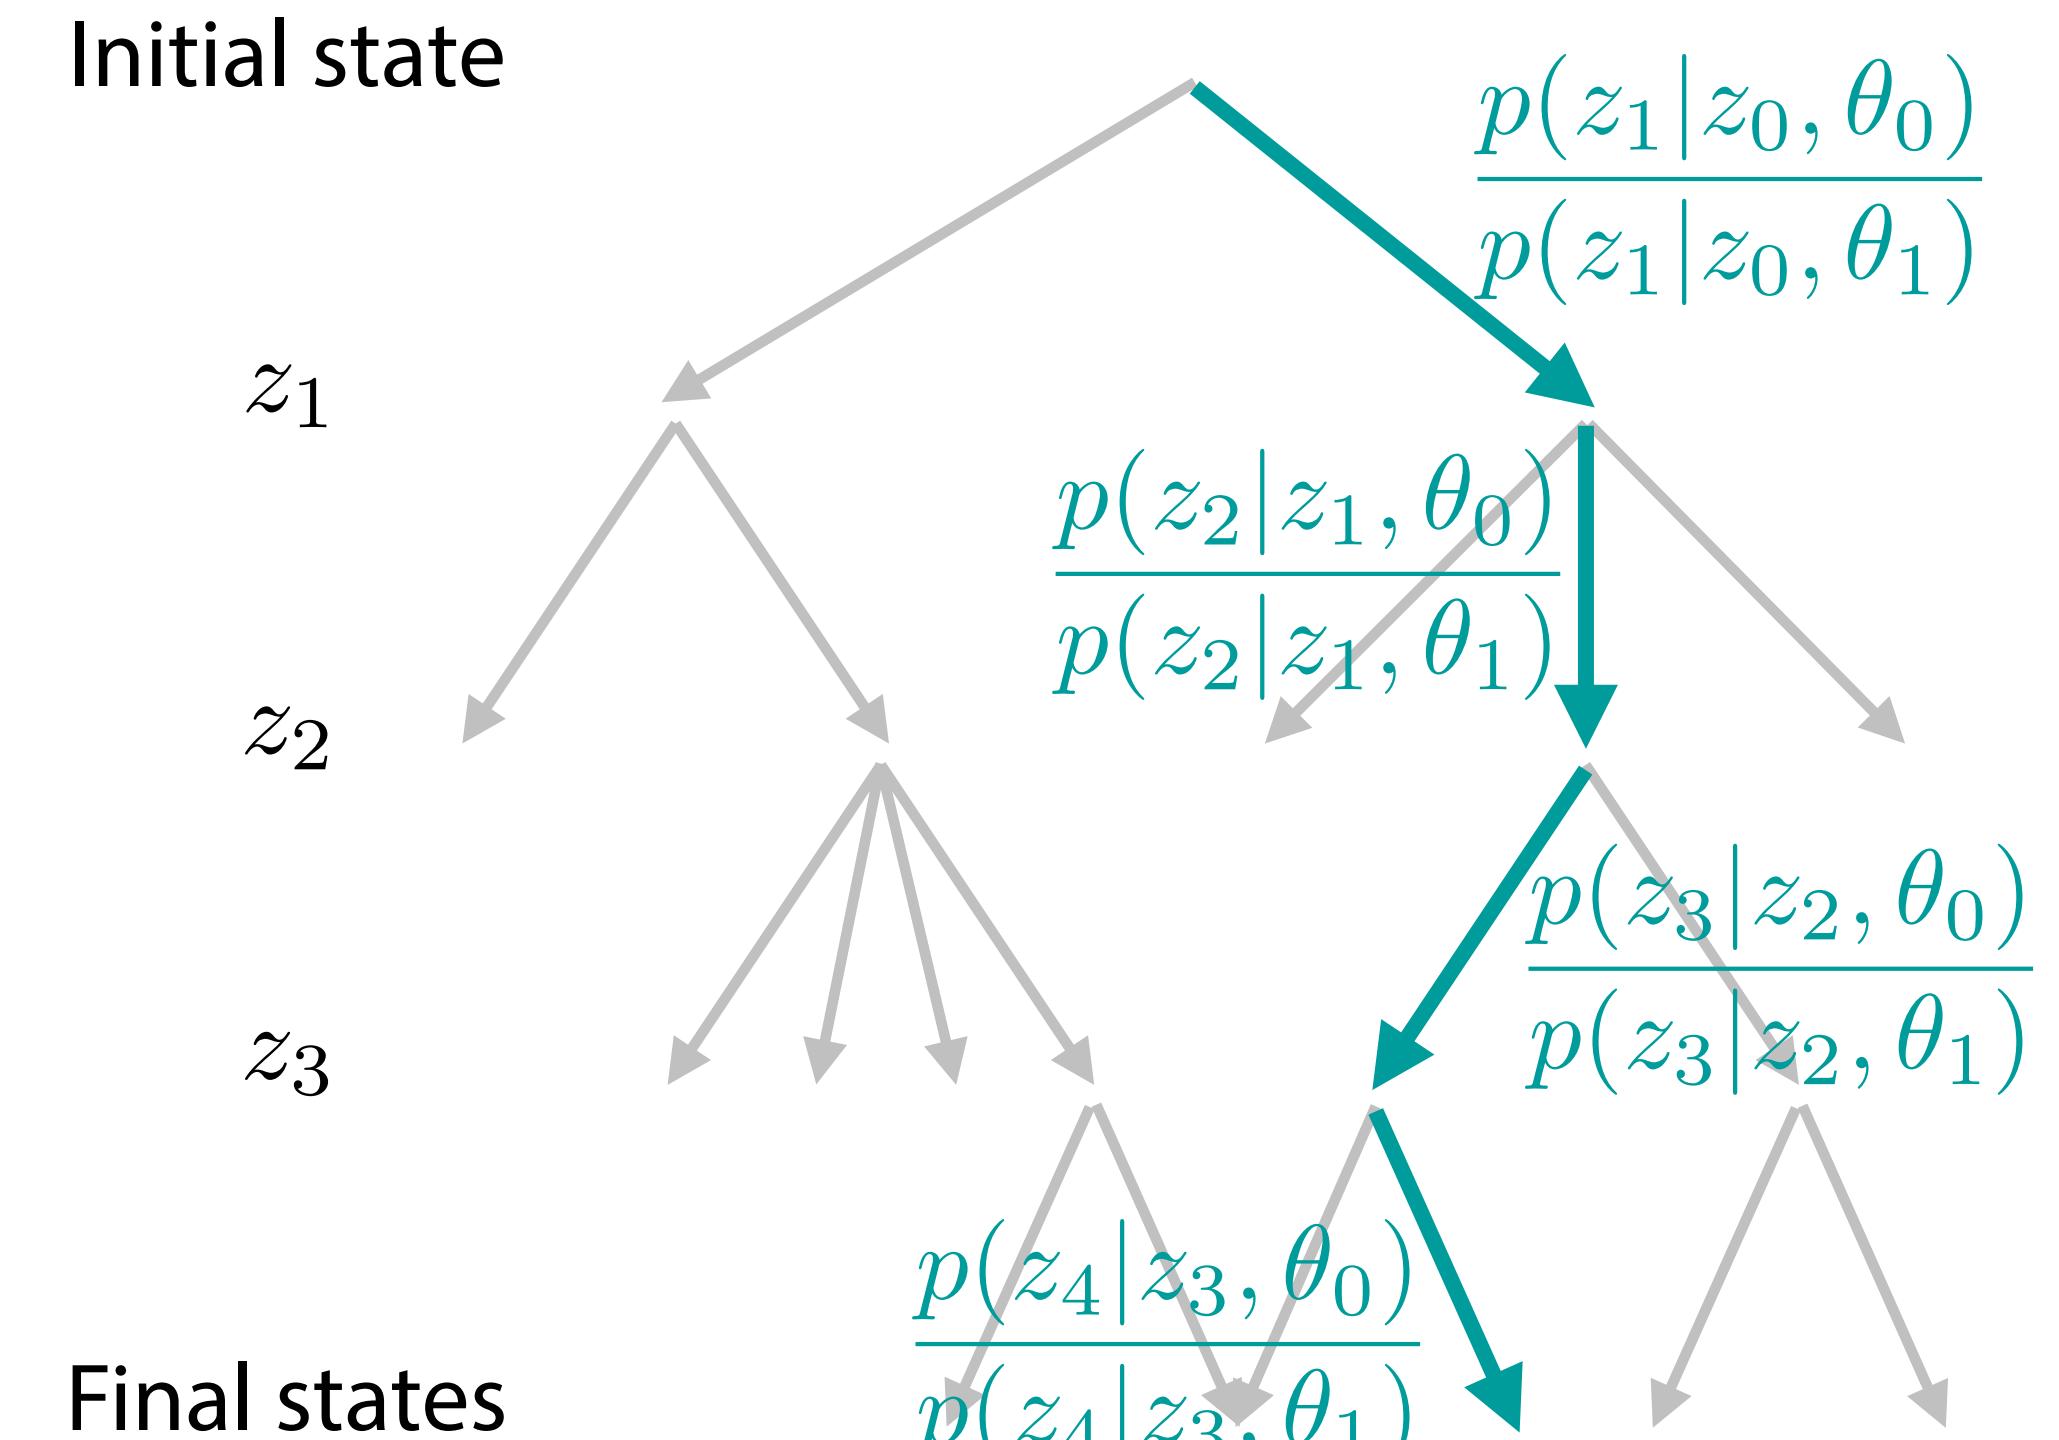
Mining gold from any simulation

- Computer simulation typically evolve along a tree-like structure of successive random branchings
- The probabilities of each branching $p_i(z_i|z_{i-1}, \theta)$ are often clearly defined in the code:

```
if random() > 0.1 + 2.5 * model_parameter:  
    do_one_thing()  
else:  
    do_another_thing()
```

- For each run of the simulator, we can calculate the probability **of the chosen path** for different values of the parameters, and the “**joint likelihood ratio**”:

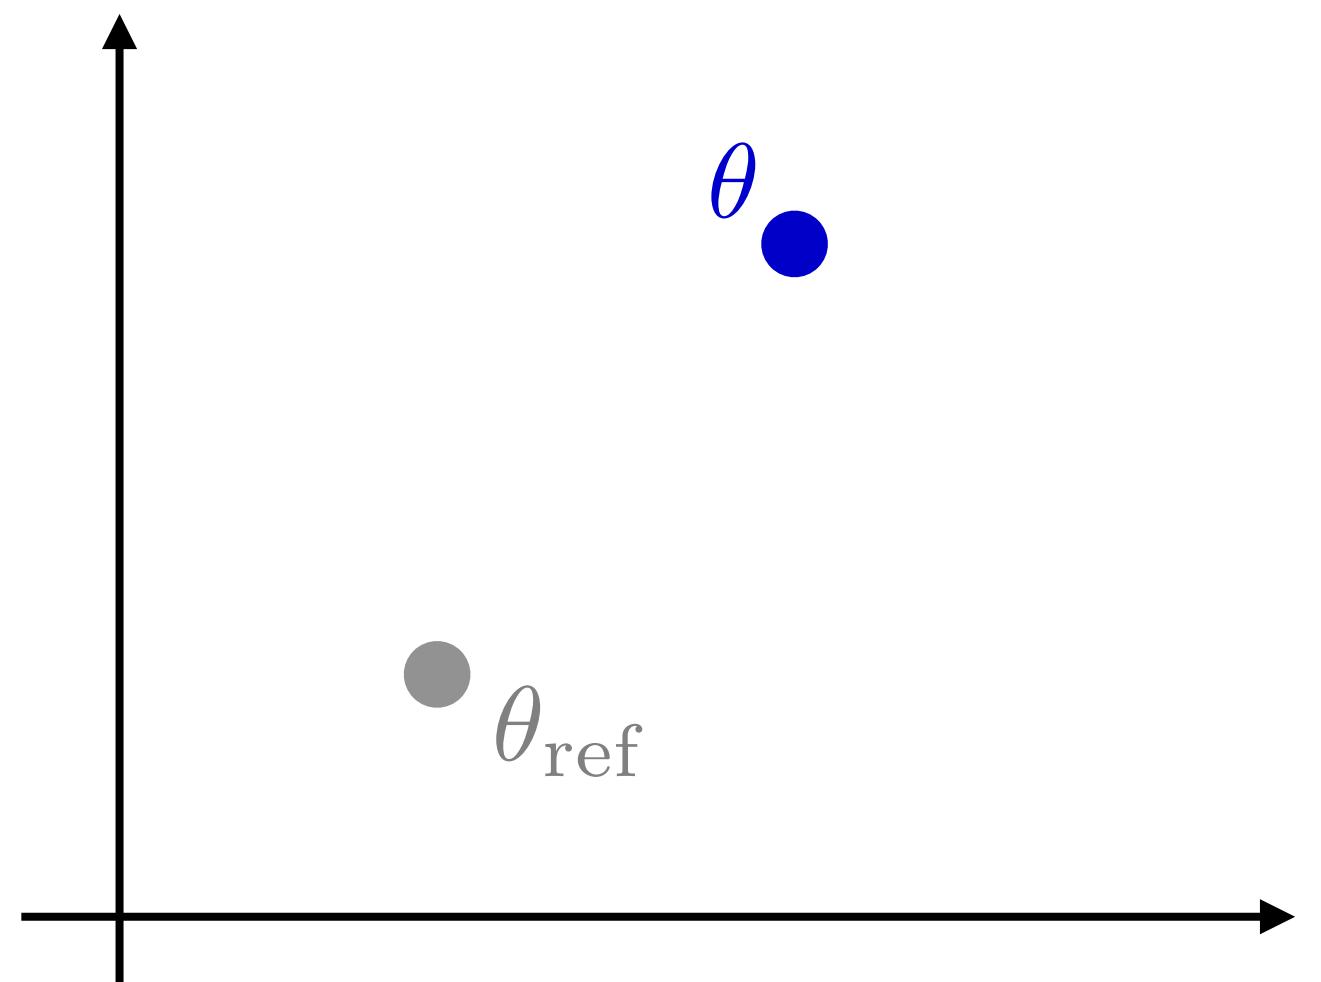
$$r(x, z|\theta_0, \theta_1) = \frac{p(x, z|\theta_0)}{p(x, z|\theta_1)} = \prod_i \frac{p(z_i|z_{i-1}, \theta_0)}{p(z_i|z_{i-1}, \theta_1)}$$



Frequentist inference

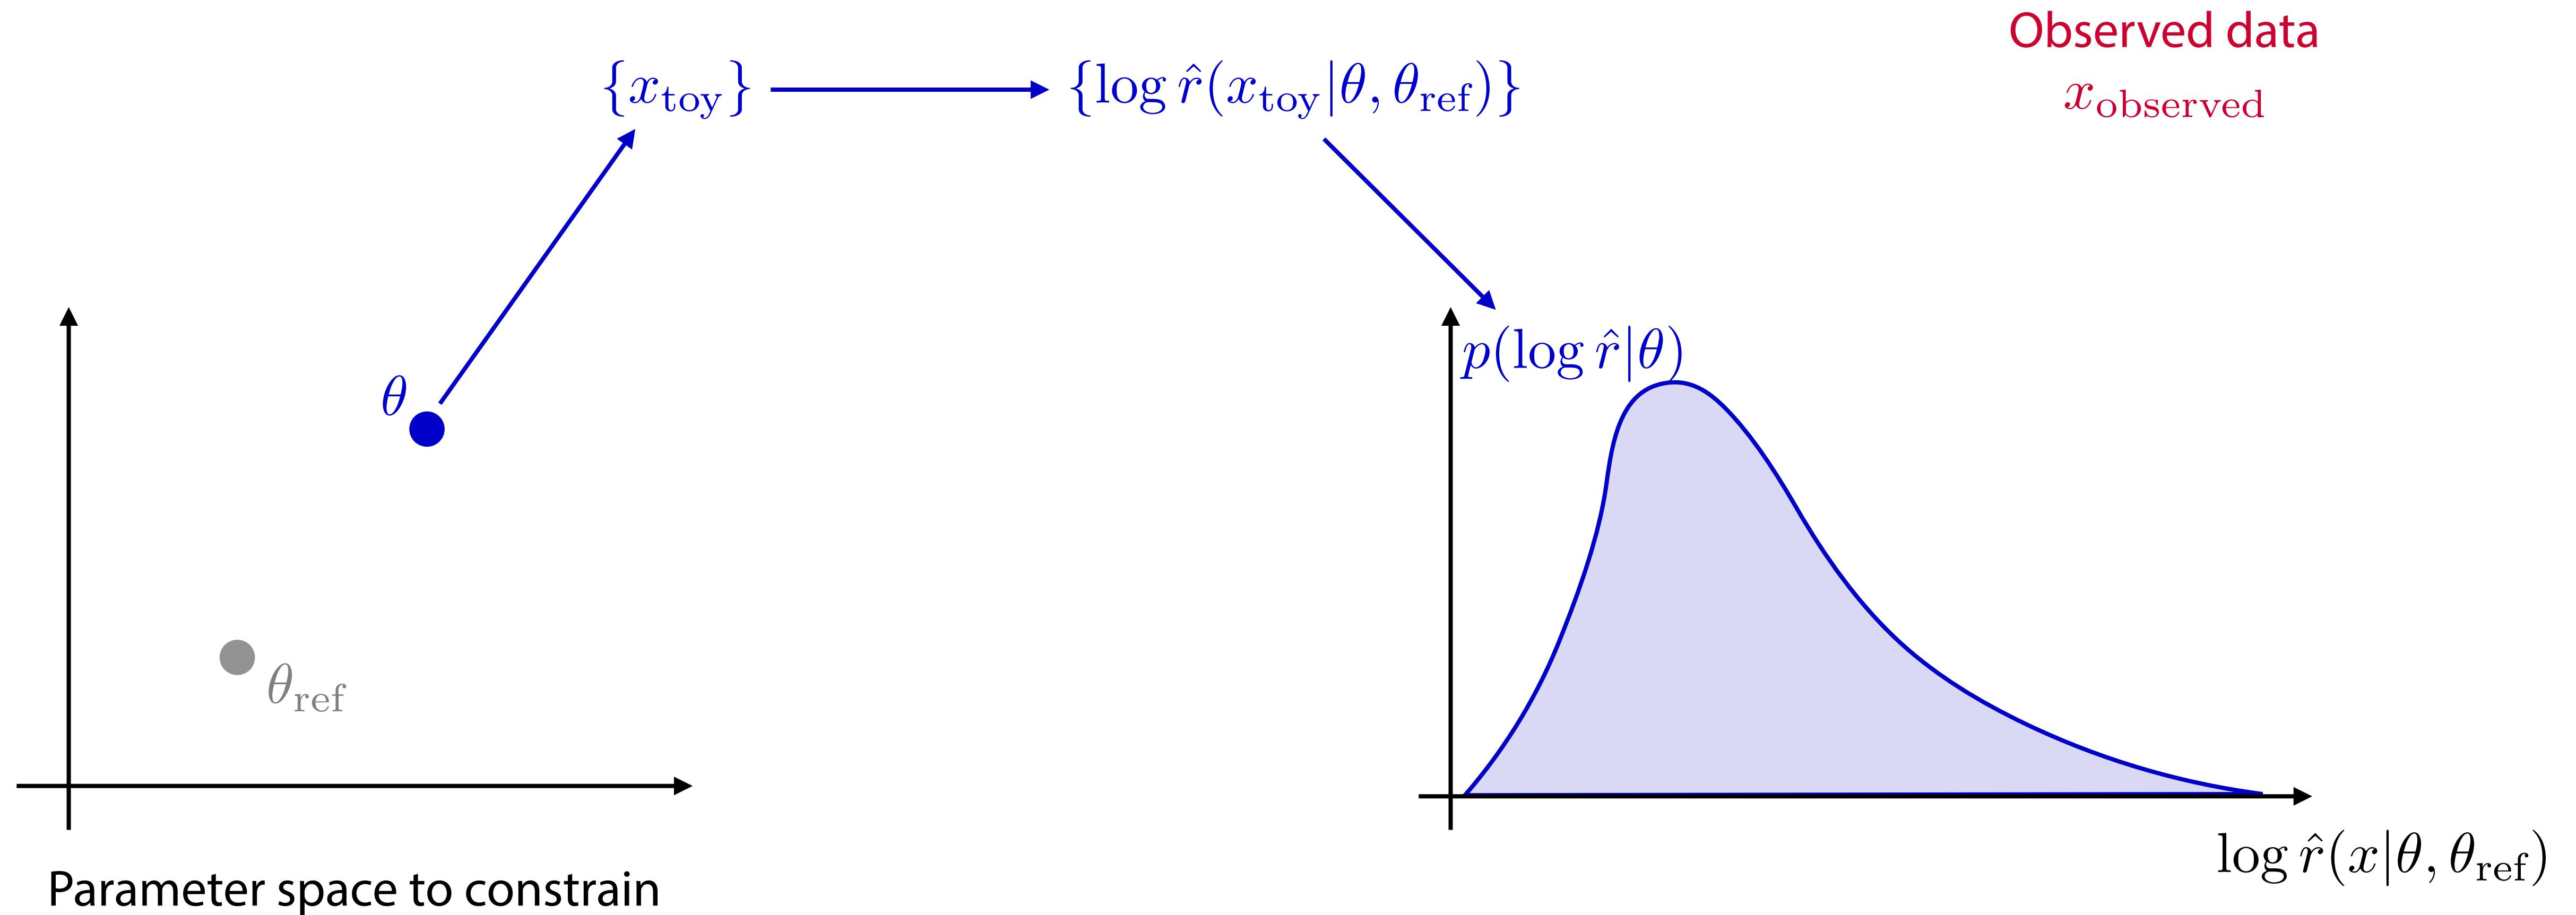
Observed data

x_{observed}

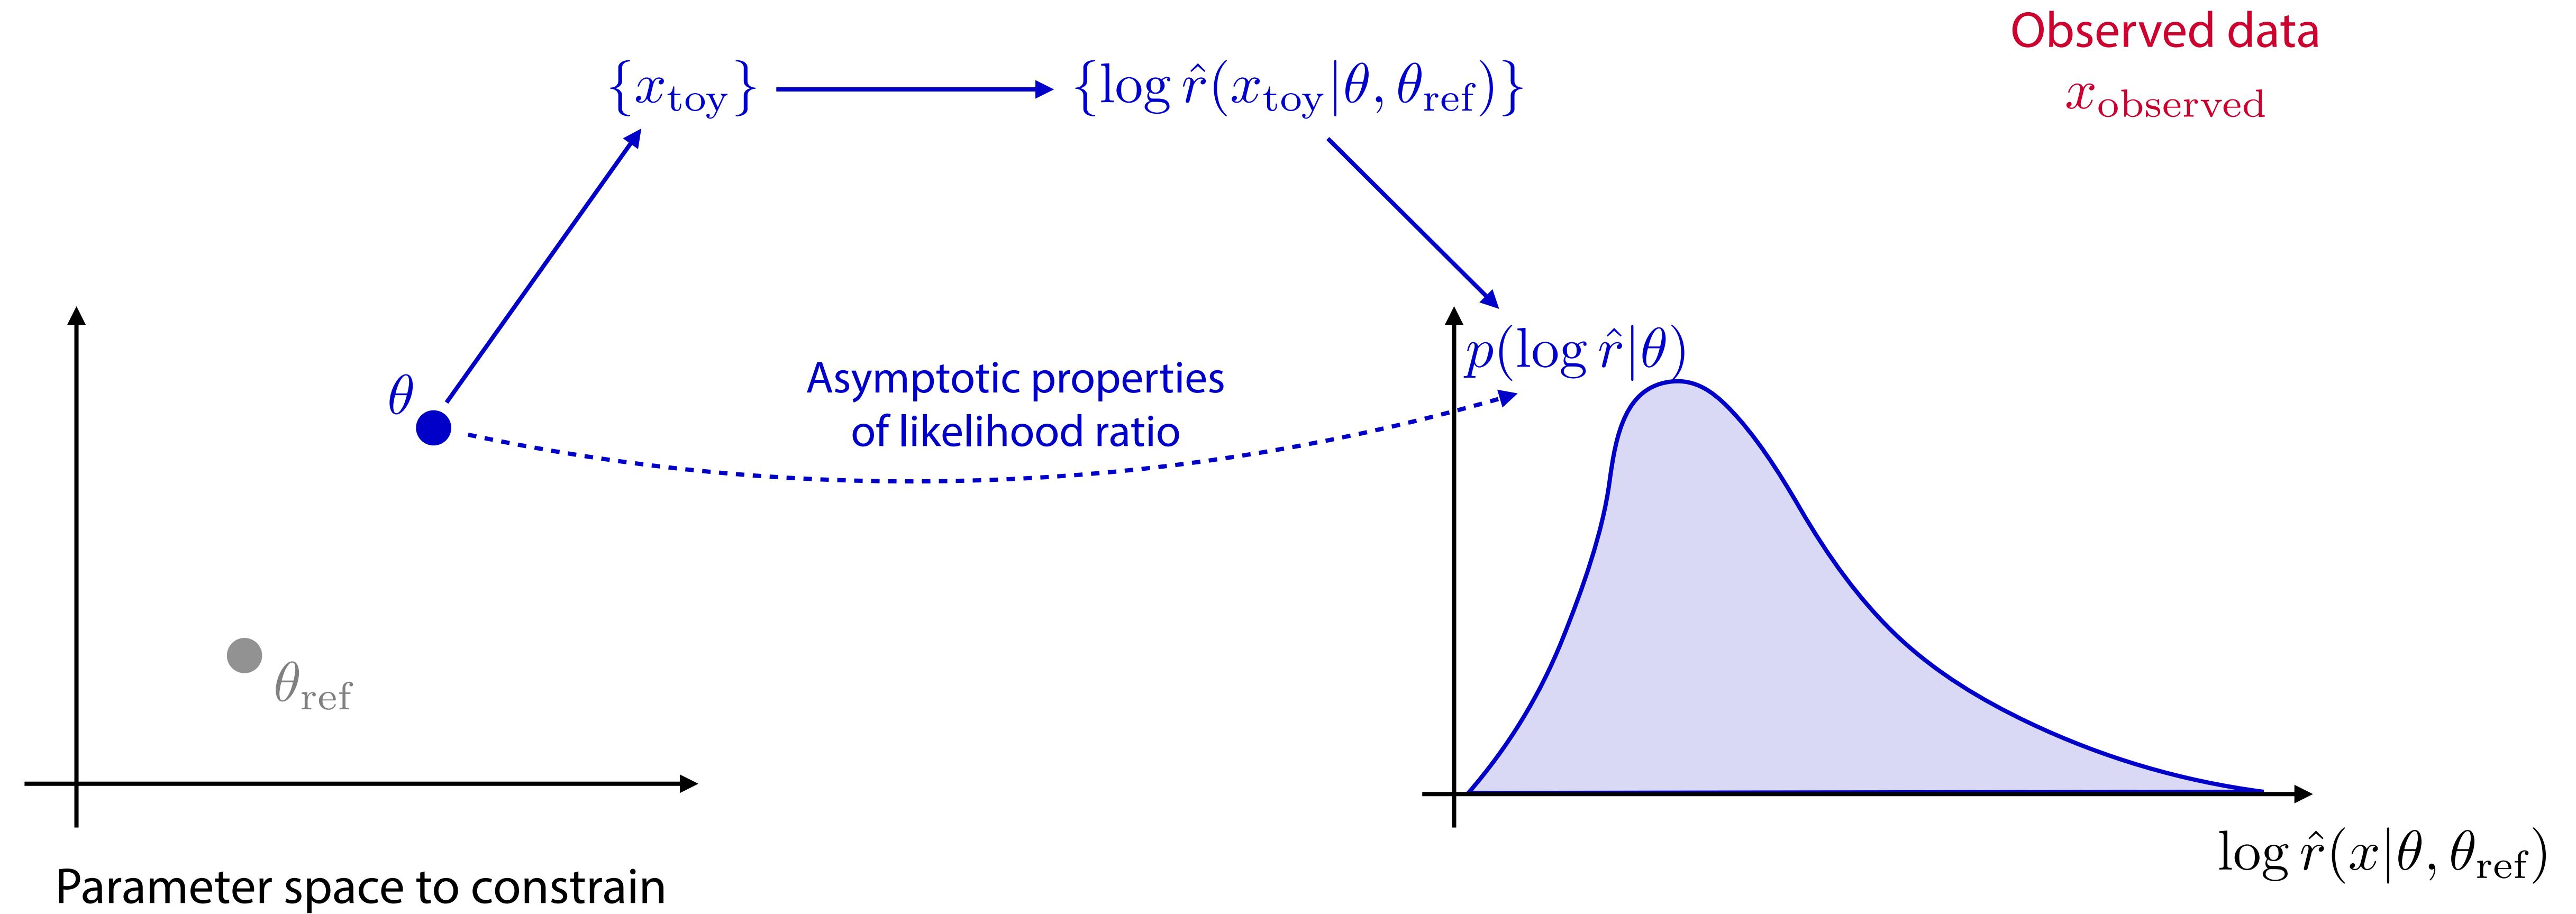


Parameter space to constrain

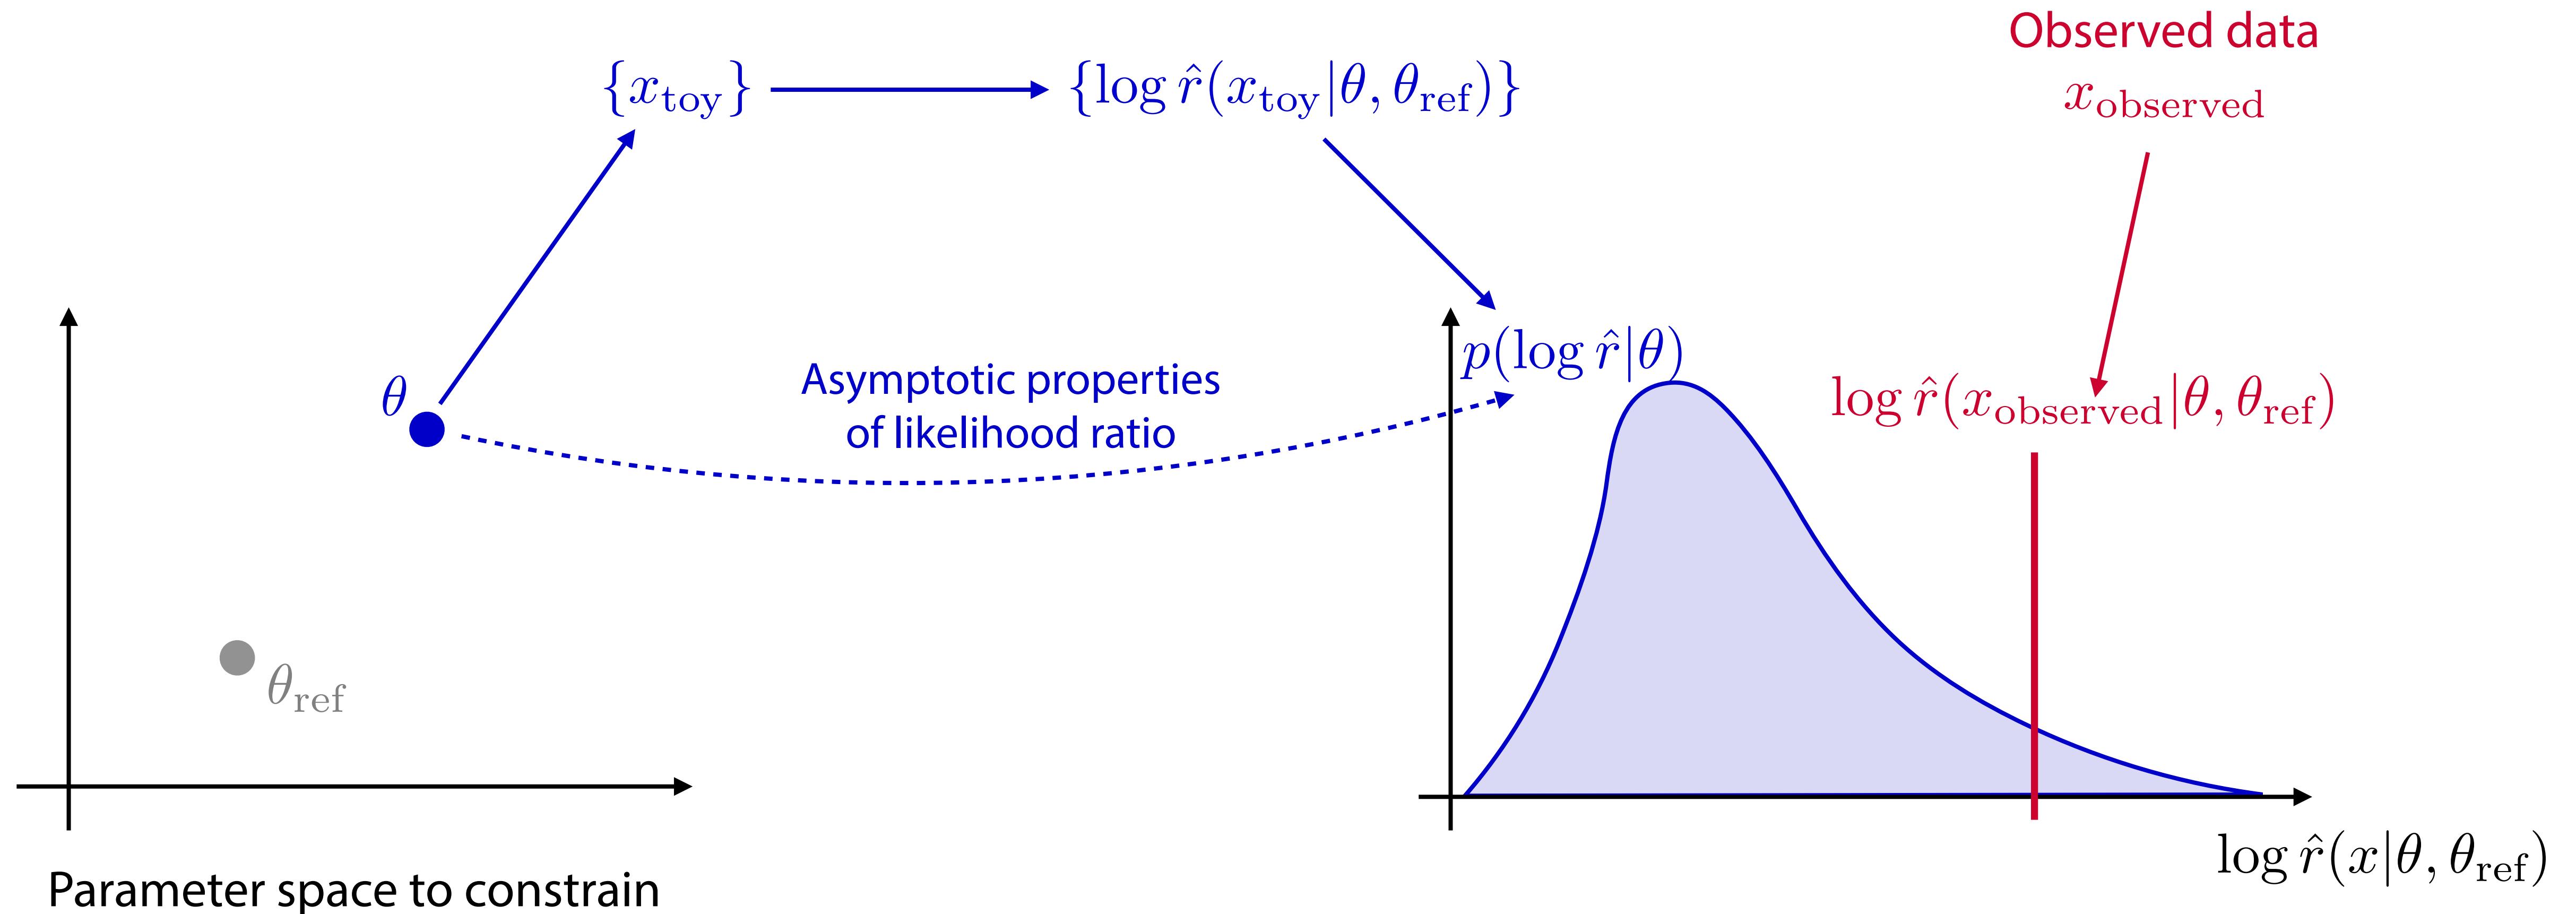
Frequentist inference



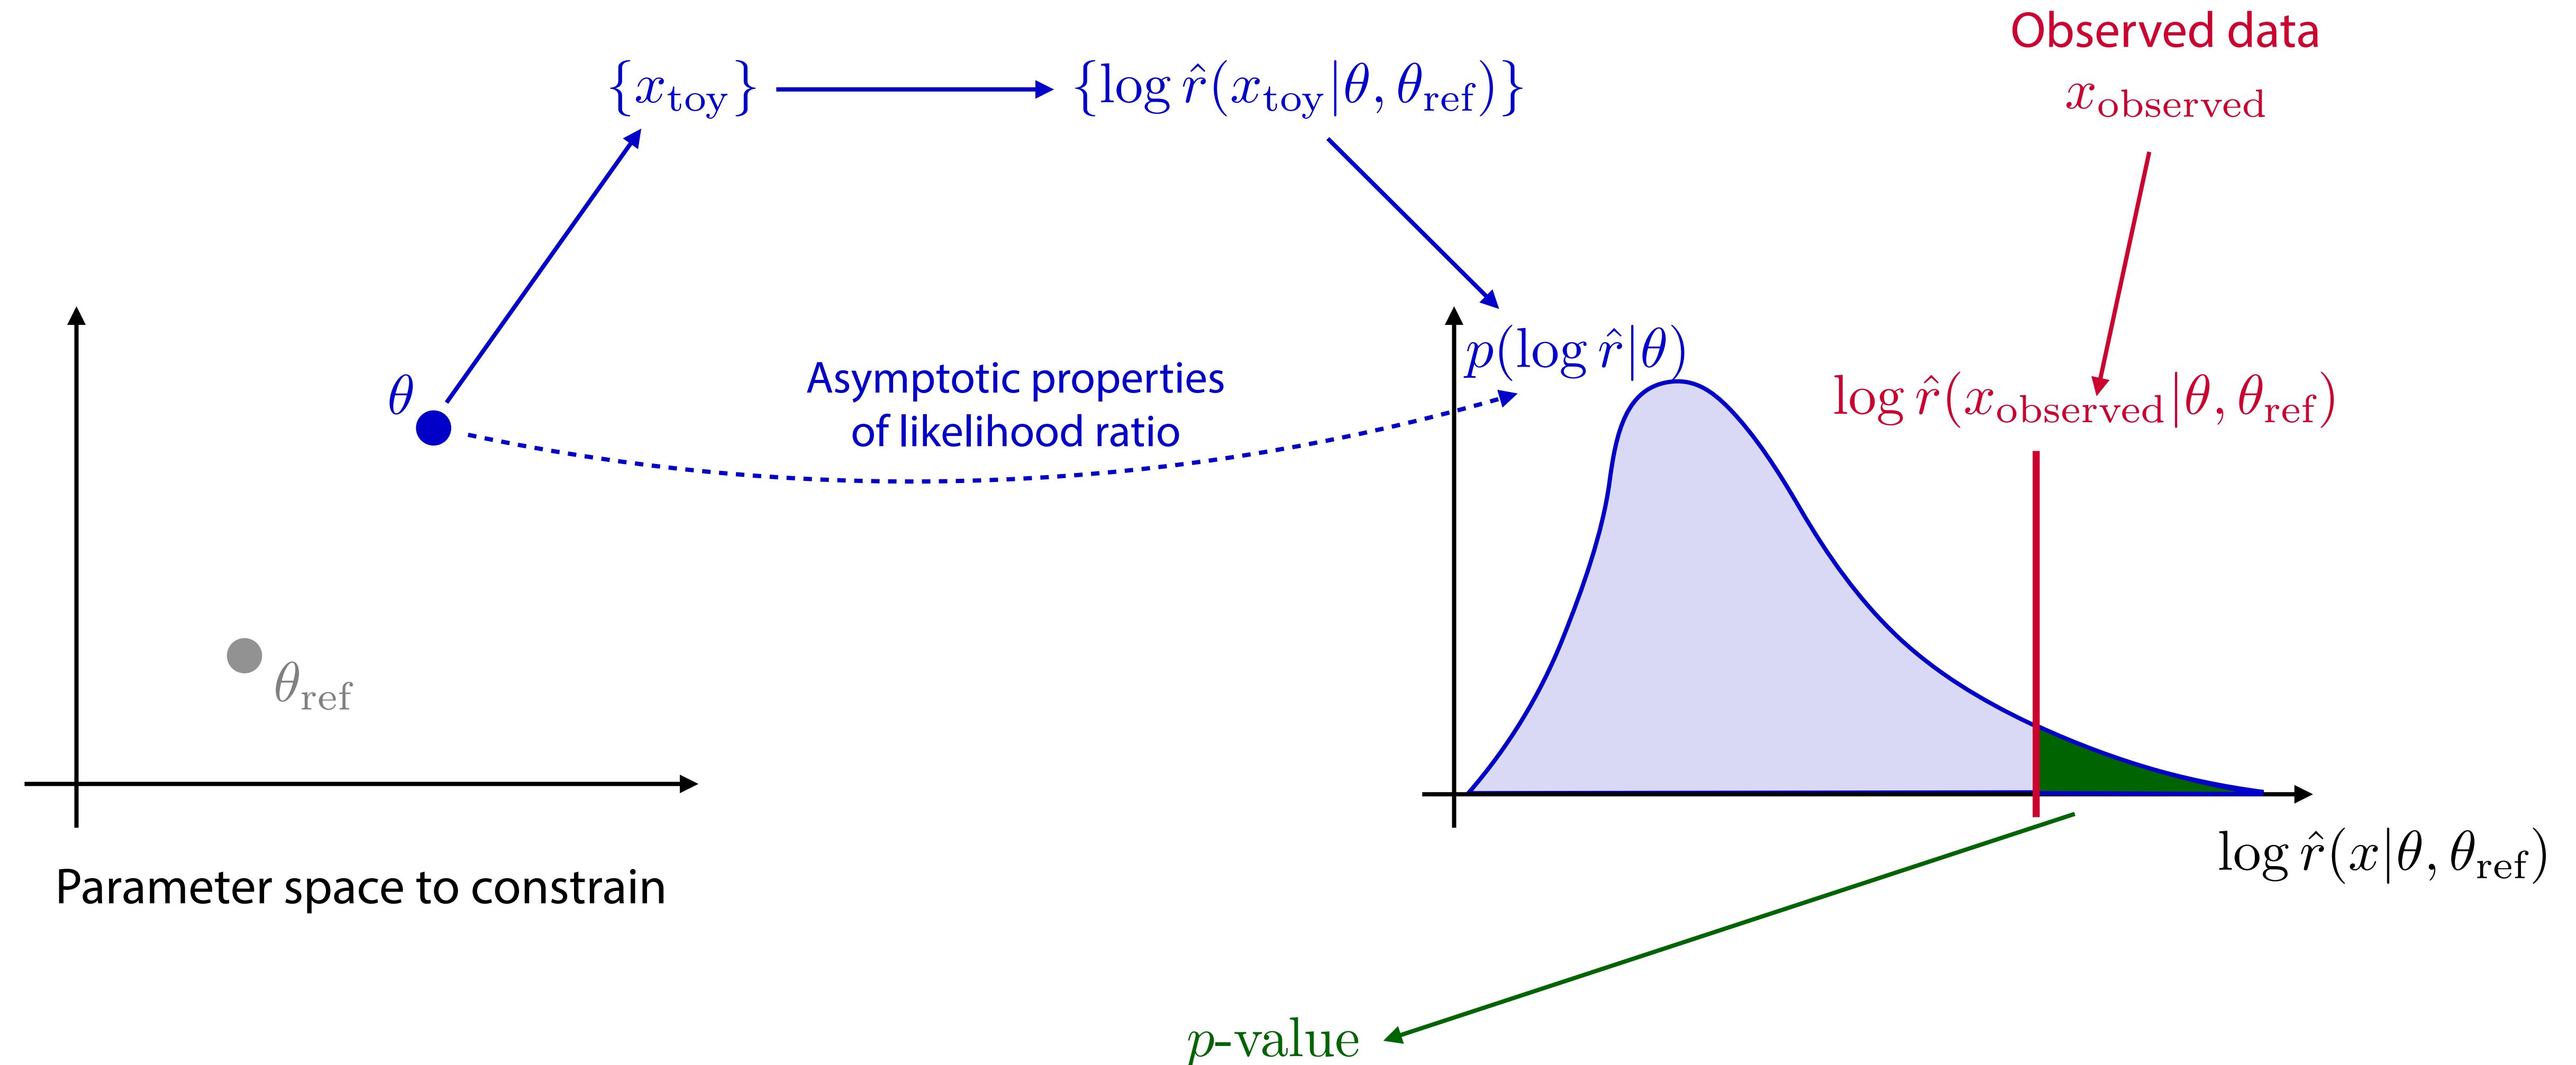
Frequentist inference



Frequentist inference



Frequentist inference



Frequentist inference

